Reconnection Protocol

There are situations in which it is desirable to move one or both
ends of a communication path from one host to another. This note
describes several situations in which the ability to reconnect is
useful, presents a mechanism to achieve reconnection, sketches how
the mechanism could be added to Host-Host or TELNET protocol, and
recommends a place for the mechanism in the protocol hierarchy.

1. Some Examples:

A. Consider the case of an executive program which TIP users could use
   to get network status information, send messages, link to other
   users, etc. Due to the TIP's limited resources the executive program
   would probably not run on the TIP itself but rather would run on one
   or more larger hosts who would be willing to share some of their
   resources with the TIP (see Figure 1).

   The TIP user could access the executive by typing a command such as
   "@ EXEC"; the TIP would then ICP to Host1's executive port. After
   obtaining the latest network news and perhaps sending a few messages,
   the user would be ready to log into Host2 (in general not the same as
   Host1) and do some work. At that point he would like to tell the
   executive program that he is ready to use Host2 and have executive
   hand him off to Host2. To do this the executive program would first
   interact with Host2, telling it to expect a call from TIP, and then
   would instruct the TIP to reconnect to Host2. When the user logs off
   Host2 he could be passed back to the executive at Host1 prepatory to
   doing more work elsewhere. The reconnection activity would be
   invisible to the TIP user.

```
                                  Reconnection
                                       |
         _____                       |             _____
        |       |                      |            |      |
        | EXEC  |<---------------------+----------->| USER |
        |_____|                      | /          |_____|
         Host1                         V /           TIP
                                        /
          _____                      /
         |       |<------------------ /
         |_____|
          Host2

                            Figure 1
```

B. Imagine a scenario in which a user could use the same name and
   password (and perhaps account) to log into any server on the network.
   For reasons of security and economy it would be undesirable to have
   every name and password stored at every site. A user wanting to use
   a Host that doesn't have his name or password locally would connect
   to it and attempt to log in as usual (See Figure 2). The Host,
   discovering that it doesn't know the user, would hand him off to a
   network authentication service which can determine whether the user
   is who he claims to be. If the user passes the authentication test he
   can be handed back to Host which can then provide him service. The
   idea is that the shuffling of the user back and forth between Host
   and Authenticator should invisible to the user.

```
(a)       _____        for authentication          _____
       |        |                |                  |        |
       |        | <--------------+-------------->   | User   |
       |_____   |                | /                |_____   |
          Host                   |/                    
                                 X
                                /|
             _____           / |
          |        |         /  v
          |        | <---   /
          |_____ |
          Authenticator
```

```
(b)
          _____                     ^              _____
       |        | <--\               |      /-->  | User   |
       |        |     \              |     /       |        |
       |_____   |      \             |    /        |_____   |
          Host           ----------+---/
                                  |  /
                                  | /
                                  |/
                            /     |
                          /       | authentication
             _____      /        | complete
          |        |    /
          |        | <------
          |_____ |
          Authenticator
```
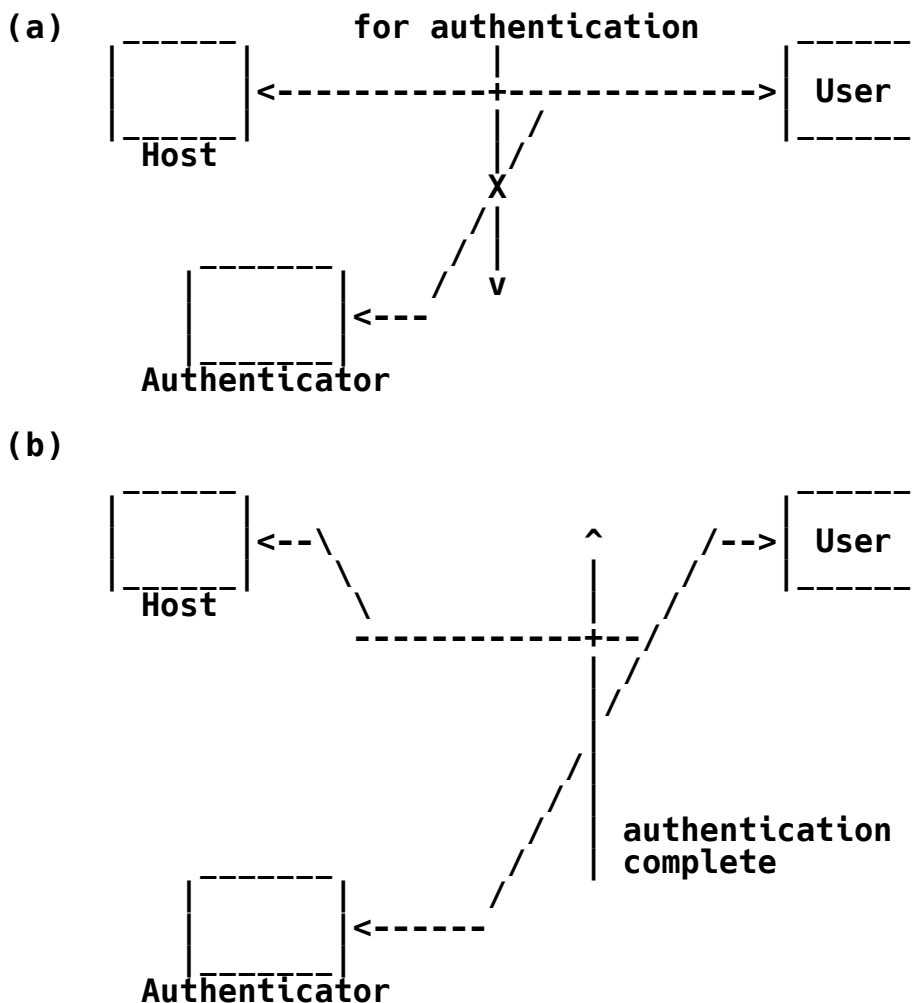
                          Figure 2

If the user doesn't trust the Host and is afraid that it might read
his password rather than pass him off to the authenticator he could
connect directly to the authentication service. After
authentication, the Authenticator can pass him off to the Host.

C. The McROSS air traffic simulation system (see 1972 SJCC paper)
   already supports reconnection. It permits an on-going simulation to
   reconfigure itself by allowing parts to move from computer to
   computer. For example, in a simulation of air traffic in the
   Northeast the program fragment simulating the New York Enroute air
   space could move from Host2 to Host5 (see Figure 3). As part of the
   reconfiguration process the New York Terminal area simulator and
   Boston Enroute area simulators break their connections with New York
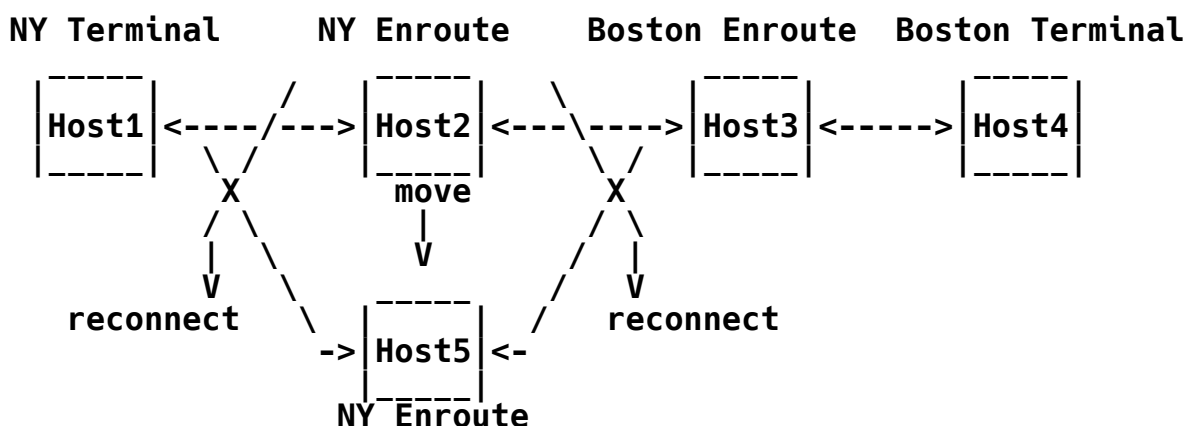   Enroute simulator at Host2 and reconnect to it at Host5.

```
   NY Terminal      NY Enroute     Boston Enroute  Boston Terminal
    _____          / _____      \    _____         _____
   |      |        / |      |      \  |      |        |      |
   |Host1 |<----/--->|Host2 |<---\---->|Host3 |<----->|Host4 |
   |_____|      \ / |_____|      \ / |_____|        |_____|
       X                move          X
      / \                 |          / \
     |   \                |         /   |
     V    \               V        /    V
   reconnect  \          _____   /  reconnect
              \       ->|      |<-
               ->|Host5 |<-
                 |_____|
                 NY Enroute
```

                        Figure 3

2. A Reconnection Mechanism

   The mechanism proposed here could be added to the existing Host-Host
   protocol or to the TELNET protocol. The mechanism is first described
   and then its adaptation to each of the protocols is discussed.

      The reconnection mechanism includes four commands:

          Reconnect Request: RRQ <path>
          Reconnect OK: ROK <path>
          Reconnect No: RNO <path>
          Reconnect Do: RDO <path> <new destination>

   where <path> is the communication path to be redirected to <new
   destination>.

   Assume that H1 wants to move its end of communication path A-C from
   itself to port D at H3 (See figure 4).

```
        (a) situation                    (b) desired situation

       H2            H3                  H2            H3
       ___           ___                 ___           ___
      |   |         |   |               |   |         |   |
      | C|<-+       |D  |               | C|<------>|D  |
      |___|  |      |___|               |___|         |___|
             |
             |       ___                              ___
             |      |   |                             |   |
       +->   |A  |                            |A  |
             |   _|___|                            |_|___|
                   H1                                H1
```

                               Figure 4

The reconnection proceeds by steps:

        a. H1 arranges for the reconnection by sending RRQ to
           H2:
                H1->H2:   RRQ (path A-C)

        b. H2 agrees to reconnect and acknowledges with ROK:

                H2->H1:   ROK (path C-A)

        c. H1 notes that H2 has agreed to reconnect and
           instructs H2 to perform the reconnection:

                H1->H2:   RDO (path A-C) (Host H3, Port D)

        d. H1 breaks paths A-C.
           H2 breaks path C-A and initiates path C-D.

In order for the reconnection to succeed H1 must, of course, have
arranged for H3's cooperation. One way H1 could do this would be to
establish the path B-D and then proceed through the reconnection
protocol exchange with H3 concurrently with its exchange with H2 (See
Figure 5):

        H1->H3: RRQ (path B-D)
        H3->H1: ROK (path D-B)
   H1->H3: RDO (path B-D) (Host H2, Port C)

```
        H2                                          H3
     _____                                      _____
    |  C   |                                    |  D   |
     ---\--                                      -/----
         \            /-->          <--\        /
          \_ _/--- --- --- --- --- ---\_/
           \  /                        \  /
            X                           X
           / \                         / \
   reconnection \                     /  reconnection
             \                       /
              \    _____   /
            ---|A         B|---
               |          |
               |_____|
                   H1
```

                           Figure 5

Either of the parties may use the RNO command to refuse or abort the
reconnection. H2 could respond to H1's RRQ with RNO; H1 can abort
the reconnection by responding to ROK with RNO rather than RDO.

It is easy to insure that messages in transit are not lost during the
reconnection. Receipt of the ROK message by H1 is taken to mean that
no further messages are coming from H2; similarly receipt of RDO from
H1 by H2 is taken to mean that no further messages are coming from
H1.

To complete the specification of the reconnection mechanism consider
the situation in which two "adjacent" entities initiate
reconnections:

   (a) situation   (b) desired situation

   H1   H4   H1   H4
  ____   ____   ____   ____
 | | | | | | | |
 | C | | E | | C |--------| E |
 |____| |____| |____| |____|

   H2   H3   H2   H3
  ____   ____   ____   ____
 | | | | | | | |
 | B |--------| D | | B | | D |
 |____| |____| |____| |____|

H2 and H3 "simultaneously" try to arrange for reconnection:

        H2->H3: RRQ (path B-D)
        H3->H2: RRQ (path D-B)

Thus, H2 sees an RRQ from H3 rather than an ROK or RNO in response to
its RRQ to H3. This "race" situation can be resolved by having the
reconnections proceed in series rather than in parallel: first one
entity (say H2) performs its reconnect and then the other (H3)
performs its reconnect. There are several means that could be used to
decide which gets to go first. Perhaps the simplest is to base the
decision on sockets and site addresses: the entity for which the 40
bit number formed by concatenating the 32 bit socket number with the
8 bit site address is largest gets to go first. Using this mechanism
the rule is the following:

    If H2 receives an RRQ from H3 in response to an RRQ of its own:
   (let NH2,NH3 = the 40 bit numbers corresponding to H2 and H[2])

    a. if NH2>NH3 then both H2 and H3 interpret H3's RRQ as an ROK in
       response to H2's RRQ.

    b. if NH2<NH3 then both interpret H3's RRQ as an RNO in response
       to H2's RRQ. This would be the only case in which it makes
       sense to "ignore" the refusal and try again - of course,
       waiting until completion of the first reconnect before doing
       so.

Once an ordering has been determined the reconnection proceeds as
though there was no conflict.

The following diagram describes the legal protocol command/response
exchange sequences for a reconnection initiated by P:

```
                              ___                         ___
                             | P |---------------| Q |
                             |___|                       |___|
         _____ _____
        | P --> Q || R R Q   |
        |_____||_____|
              +---------+
              |
              V
         _____
        | Q --> P ||   R O K   |  R N O  ----|   R R Q    |
        |         ||           |         | E |            |
        |_____||_____|_____|___|_____|
          +-----------+                       |
          |           |            Yes  +----------+  No
          |           +-----------------| NP > NQ? |------+
          |     |                       +----------+      |
        __V___V_____          v        |
        | P --> Q || R D O ----|  R N O  ----|           |
        |         ||       | E |         | E |           |
        |_____||_____|___|_____|___|           |
          +-----------------------------------------------+
          |
          V
        _____
        | Q --> P || R D O ----|  R N O  ----|
        |         ||       | E |         | E |
        |_____||_____|___|_____|___|
```

NP and NQ are the 40 bit numbers for P and Q; E indicates end of
sequence.

3. Adding the Reconnection Mechanism to Host-Host Protocol

        The four reconnect commands could be included directly in
        Host-Host protocol as follows:

            RRQ <my socket> <your socket>
            ROK <my socket> <your socket>
            RNO <my socket> <your socket>
            RDO <my socket> <your socket> <new host> <new socket>

    The ROK and RDO commands for a send connection should not be sent
    until there are no messages in transit over the connection. The RDO

command is to be interpreted as a CLS.

The reconnection:

```
     H2              H3                         H2              H3
    ___             ___                        __              __
   |   |           |   |                      | C |-------- | D |
   | C |           | D |                      |___|          |___|
   |___|           |___|          ===>
     |               |
     |               |                                ___
     ---| A  B |---                                  |   |
        |_____|                                      |___|
         H1                                           H1
```

  could be accomplished as follows:

```
        H1->H2:   RRQ A C
        H1->H3:   RRQ B D
        H2->H1:   ROK C A
        H3->H1:   ROK D B
        H1->H2:   RDO A C H3 D
        H1->H3:   RDO B D H2 C
        H2->H1:   CLS C A
        H3->H1:   CLS D B
        H2->H3:   STR C D size
        H3->H2:   RTS D C link
```

  Note that it is possible for the STR from H2 to arrive at H3 before
  the RDO from H1. H3 must be prepared to queue such an RFC until it
  gets an RDO or RNO from H1. Stated differently, transmission of an
  ROK for a local socket causes the socket to move from an "open" state
  to a "reconnect pending" state and indicates willingness to queue
  subsequent RFC's until receipt of a "matching" RDO or RNO.

4. Reconnection in TELNET Protocol

  Independently of whether Host-Host protocol directly supports
  reconnection, the reconnection mechanism could be included in TELNET
  with the addition to the TELNET protocol of the five commands:

```
        RRQ
        ROK
        RNO
        RDO <host> <socket>
        RWT <host> <socket>
```

where RRQ, ROK, RNO, RDO, and RWT are appropriately chosen characters
in the range 128 to 255. The first three commands require no
parameters since they refer to the connections they are received on.
For RDO and RWT, <host> is an 8 bit (= 1 TELNET character) host
address and <socket> is a 32 bit (= 4 TELNET characters) number that
specifies a TELNET receive socket at the specified host.

A pending reconnection can be activated with either RDO or RWT. The
response to either is to first break the TELNET connection with the
sender and then reopen the TELNET connection to the host and sockets
specified. For RDO, the connection is to be reopened by sending two
RFC's; for RWT, by waiting for two RFC's.

The RWT command is introduced to avoid races such as the one between
the STR and CLS (RDO) discussed above. In Host-Host protocol the
race is avoided by putting a connection into "reconnect pending"
state upon transmission of ROK. For TELNET the race can be avoided
by the initiator of the reconnection by judicious use of RWT and RDO.
For example the reconnection:

```
   H2                         H3        H2      M      H3
 +---+                      +---+      +---+  +---+
 |   |  ----+        +----->|   |      |   |  |---------->|   |
 | Y | N   |        |     Q | Z | ==>  | Y | | N      | Z |
 |   |  <-+ |    H1  |   +---|   |      |   | |  <-------|   |
 +---+    | | M  +---+ P |  +---+      +---+ +---+      +---+
         | +--->|   |----+
         |      | X |  |                          H1
         +------|   |<-----+                    +---+
              +---+                            |   |
               H1                             | X |
                                              |   |
                                              +---+
```

could be accomplished as follows:

```
        X->Y:  RRQ
        X->Z:  RRQ
        Y->X:  ROK
        Z->X:  ROK
        X->Y:  RWT  H3 P
        X closes connections to Y
        Y closes connections to X
        Y waits for STR and RTS from H3
        X->Z: RDO H2 N
        X closes connections to Z
        Z closes connections to X
        Z sends STR and RTS to H2 which Y answers with
           matching RTS and STR to complete reconnection
```

The reconnection mechanism for TELNET can be made to fit nicely into
the command format suggested by Cosell and Walden in RFC #435.
Consider the addition of three new commands to TELNET:

        Prepare for Reconnect: RCP
        Begin Reconnect by sending RFC's: RCS
        Begin Reconnect by waiting for RFC's: RCW

Using these three command and the DO, DON'T, WILL, WON'T commands of
RFC #435, the commands proposed earlier become:

        RRQ => DO RCP
        ROK => WILL RCP
        RNO => WON'T RCP ;for responses to DO RCP
                DON'T RCP ;for responses to WILL RCP
                           ;i.e. used to cancel an RCP.
        RDO <host> <socket> => DO RCS <host> <socket>
        RWT <host> <socket> => DO RCW <host> <socket>

As RFC #435 notes the nice thing about this structure is that a host
choosing not to implement reconnection does not even have to know
what RCP means. All that it need do in response to DO RCP is to
transmit WON'T RCP.

5. Recommendation

I feel that reconnection is a basic notion and that its proper place
within the protocol hierarchy is at the Host-Host level where it
would be available for use in all higher level protocols. The
difficulty is that placing it there would, of course, require NCP
rewrites. Reluctance to make NCP modifications would probably be
sufficient to kill interest in the proposal.

Therefore, for pragmatic reasons, I recommend that the reconnection
mechanism be included in TELNET as an "option" in the spirit of RFC
#435. This can be accomplished with the addition to the TELNET
protocol of the RCP, RCS, RCW commands as described in Section 4.
Modification of user- and server-TELNET programs to handle these new
commands should be straightforward. If a reconnection option is made
part of TELNET protocol the TENEX hosts will support it. In
addition, the TIP guys (Walden and Cosell) have said that they like
the reconnection mechanism and have agreed, in principle, to
implement it for TIPs if it is accepted as part of TELNET protocol.

Addition of reconnection at the TELNET level rather than the Host-Host level is admittedly a compromise. However, with it, activity of the sort described in Examples A and B of Section 1 will be possible.

6. Additional Remarks

A. Reconnection is not a new notion. An early proposal for Host-Host protocol (RFC #36) included facilities to support reconnection. The reconnection mechanism in RFC #36 supposes a configuration in which entities are "daisy-chained" together by connections:

```
     __         __         __         __         __
 ---|__|----|__|----|__|----|__|----|__|---
    |__|       |__|       |__|       |__|       |__|
```

and specifies how one or more entities can break out of the chain. As suggested above (Figure 5) the mechanism proposed in this note supports that kind of reconnection.

B. In practice one would expect simultaneous initiation of reconnects by adjacent entities to be relatively rare.

C. The approach taken in RFC #36 to handle simultaneous reconnection attempts by entities adjacent in the chain is to accomplish the reconnect as a single, carefully coordinated, global reconnect. I feel that the sequence of locally coordinated reconnects as proposed above is preferable. When N adjacent entities simultaneously attempt reconnection the single, globally coordinated reconnect as outlined in RFC #36 requires $\sim N^2$ control messages whereas the sequential locally coordinated reconnect requires $\sim N$.

D. A word about security is in order. It should be clear that the decision to accept or reject a particular reconnection request is the responsibility of the entity (person at a terminal or process) using the connection. In many cases the entity may choose to delegate that responsibility to its NCP or TELNET (e.g., Example A, Section 1). However, the interface a Host provides to the reconnection mechanism should include means for local entities to exercise control over response to remotely initiated reconnection requests. For example, a user-TELNET might support several modes of operation with respect to remotely initiated reconnections:

   1. transparent: all requested reconnections are to be performed in a way that is invisible to the user;

   2. visible: all requested reconnections are to be performed and the user is to be informed whenever a reconnection occurs;

   3. confirmation: the user is to be informed of each reconnection
      request which he may accept or reject;

   4. rejection: all requested reconnects are to be rejected.

E. Reconnection can be achieved almost trivially within the Message
   Switched Protocol (MSP) proposed by Bressler, Murphy and Walden in
   RFC #333 (within MSP, "reconnection" is probably not the correct
   term). For example use of the following conventions with that MSP
   (expressed in the terminology of RFC #333) support reconnection:

   1. unless a reconnection is in progress, rendezvous is to occur at
      the sending site;

   2. the receiving end of a communication path can be moved by passing
      the names of the rendezvous site and the ports to the new
      receiver;

   3. receipt of an OUT message for which the source site differs from
      the rendezvous site signals that the sending end is being moved;
      the source site should be used as the rendezvous site for
      subsequent IN messages;

   4. the sending end of a communication path can be moved by passing
      the names of the ports to the new sender; to complete the move the
      new sender uses the previous sender's site as rendezvous site for
      its first OUT message and its own site as rendezvous for
      subsequent OUT messages.

   As simple and appealing as this procedure seems, I doubt that it
   would be used in practice if MSP were to be implemented as a
   replacement for or alternative to existing Host-Host protocol. The
   reason is that the ability to pass ports from Host to Host
   (needlessly) complicates port name allocation by requiring
   cooperation among Hosts to manage the allocation (e.g., before a Host
   can safely allocate a port name for use by a local process it must
   not only insure that the port is not in use locally but also that no
   process out in the network is using it.) The inter-Host cooperation
   required to support the passage of ports among Hosts can probably not
   be reliably achieved in practice. Therefore port passage of the sort
   described in RFC #333 should not be supported at the Host-Host
   protocol level. For this reason, I feel that within an MSP
   "reconnection" would be best handled by a mechanism such as the one
   proposed in this note.

        [This RFC was put into machine readable form for entry]
        [into the online RFC archives by Anthony Anderberg 4/99]