DOMAIN NAMES - IMPLEMENTATION and SPECIFICATION

```
+----------------------------------------------------------+
|                                                          |
|    This memo discusses the implementation of domain      |
|    name servers and resolvers, specifies the format of   |
|    transactions, and discusses the use of domain names   |
|    in the context of existing mail systems and other     |
|    network software.                                     |
|                                                          |
|    This memo assumes that the reader is familiar with    |
|    RFC 882, "Domain Names - Concepts and Facilities"     |
|    which discusses the basic principles of domain        |
|    names and their use.                                  |
|                                                          |
|    The algorithms and internal data structures used in   |
|    this memo are offered as suggestions rather than      |
|    requirements; implementers are free to design their   |
|    own structures so long as the same external           |
|    behavior is achieved.                                 |
|                                                          |
+----------------------------------------------------------+



        +--------------------------------------------------+
        |                                                  |
        |             *****  WARNING  *****                |
        |                                                  |
        |   This RFC contains format specifications which  |
        |   are preliminary and are included for purposes  |
        |   of explanation only.  Do not attempt to use    |
        |   this information for actual implementations.   |
        |                                                  |
        +--------------------------------------------------+
```

TABLE OF CONTENTS

INTRODUCTION

  Overview

    The goal of domain names is to provide a mechanism for naming
    resources in such a way that the names are usable in different
    hosts, networks, protocol families, internets, and administrative
    organizations.

    From the user's point of view, domain names are useful as
    arguments to a local agent, called a resolver, which retrieves
    information associated with the domain name. Thus a user might
    ask for the host address or mail information associated with a
    particular domain name. To enable the user to request a
    particular type of information, an appropriate query type is
    passed to the resolver with the domain name. To the user, the
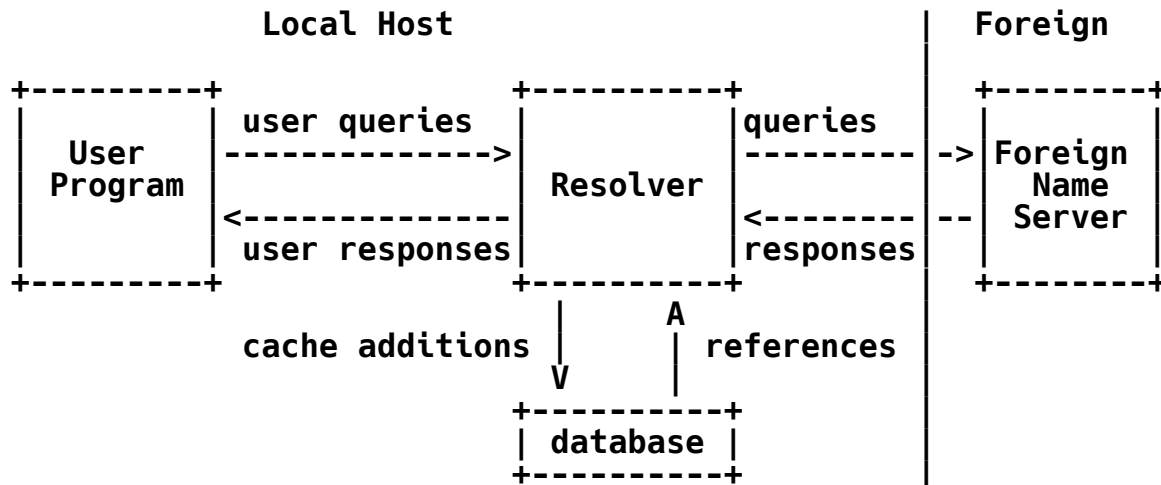   domain tree is a single information space.

    From the resolver's point of view, the database that makes up the
    domain space is distributed among various name servers. Different
    parts of the domain space are stored in different name servers,
    although a particular data item will usually be stored redundantly
    in two or more name servers. The resolver starts with knowledge
    of at least one name server. When the resolver processes a user
    query it asks a known name server for the information; in return,
    the resolver either receives the desired information or a referral
    to another name server. Using these referrals, resolvers learn
    the identities and contents of other name servers. Resolvers are
    responsible for dealing with the distribution of the domain space
    and dealing with the effects of name server failure by consulting
    redundant databases in other servers.

    Name servers manage two kinds of data. The first kind of data
    held in sets called zones; each zone is the complete database for
    a particular subtree of the domain space. This data is called
    authoritative. A name server periodically checks to make sure
    that its zones are up to date, and if not obtains a new copy of
    updated zones from master files stored locally or in another name
    server. The second kind of data is cached data which was acquired
    by a local resolver. This data may be incomplete but improves the
    performance of the retrieval process when non-local data is
    repeatedly accessed. Cached data is eventually discarded by a
    timeout mechanism.

    This functional structure isolates the problems of user interface,
    failure recovery, and distribution in the resolvers and isolates
    the database update and refresh problems in the name servers.

   Implementation components

      A host can participate in the domain name system in a number of
      ways, depending on whether the host runs programs that retrieve
      information from the domain system, name servers that answer
      queries from other hosts, or various combinations of both
      functions. The simplest, and perhaps most typical, configuration
      is shown below:

```
                    Local Host                    |   Foreign
                                                  |
    +---------+               +----------+        |  +--------+
    |         | user queries  |          | queries|  |        |
    |  User   |-------------->|          |--------|->|Foreign |
    | Program |               | Resolver |        |  | Name   |
    |         |<--------------|          |<-------|--| Server |
    |         | user responses|          |responses  |        |
    +---------+               +----------+        |  +--------+
                                   |    A         |
                  cache additions  |    | references
                                   V    |          |
                              +----------+         |
                              | database |         |
                              +----------+         |
```

      User programs interact with the domain name space through
      resolvers; the format of user queries and user responses is
      specific to the host and its operating system. User queries will
      typically be operating system calls, and the resolver and its
      database will be part of the host operating system. Less capable
      hosts may choose to implement the resolver as a subroutine to be
      linked in with every program that needs its services.

      Resolvers answer user queries with information they acquire via
      queries to foreign name servers, and may also cache or reference
      domain information in the local database.

   Note that the resolver may have to make several queries to several
      different foreign name servers to answer a particular user query,
      and hence the resolution of a user query may involve several
      network accesses and an arbitrary amount of time. The queries to
      foreign name servers and the corresponding responses have a
      standard format described in this memo, and may be datagrams.

        Depending on its capabilities, a name server could be a stand
        alone program on a dedicated machine or a process or processes on
        a large timeshared host. A simple configuration might be:

                    Local Host   Foreign
                                                         |
            +--------+ |
           /   /| |
          +--------+ | +----------+ |
          | | | | |responses | +--------+
          | | | | Name |---------- |->| |
          | Master | |--------------->| Server | | | Foreign|
          | files | | | |<-------- |--|Resolver|
          | |/ | | queries | +--------+
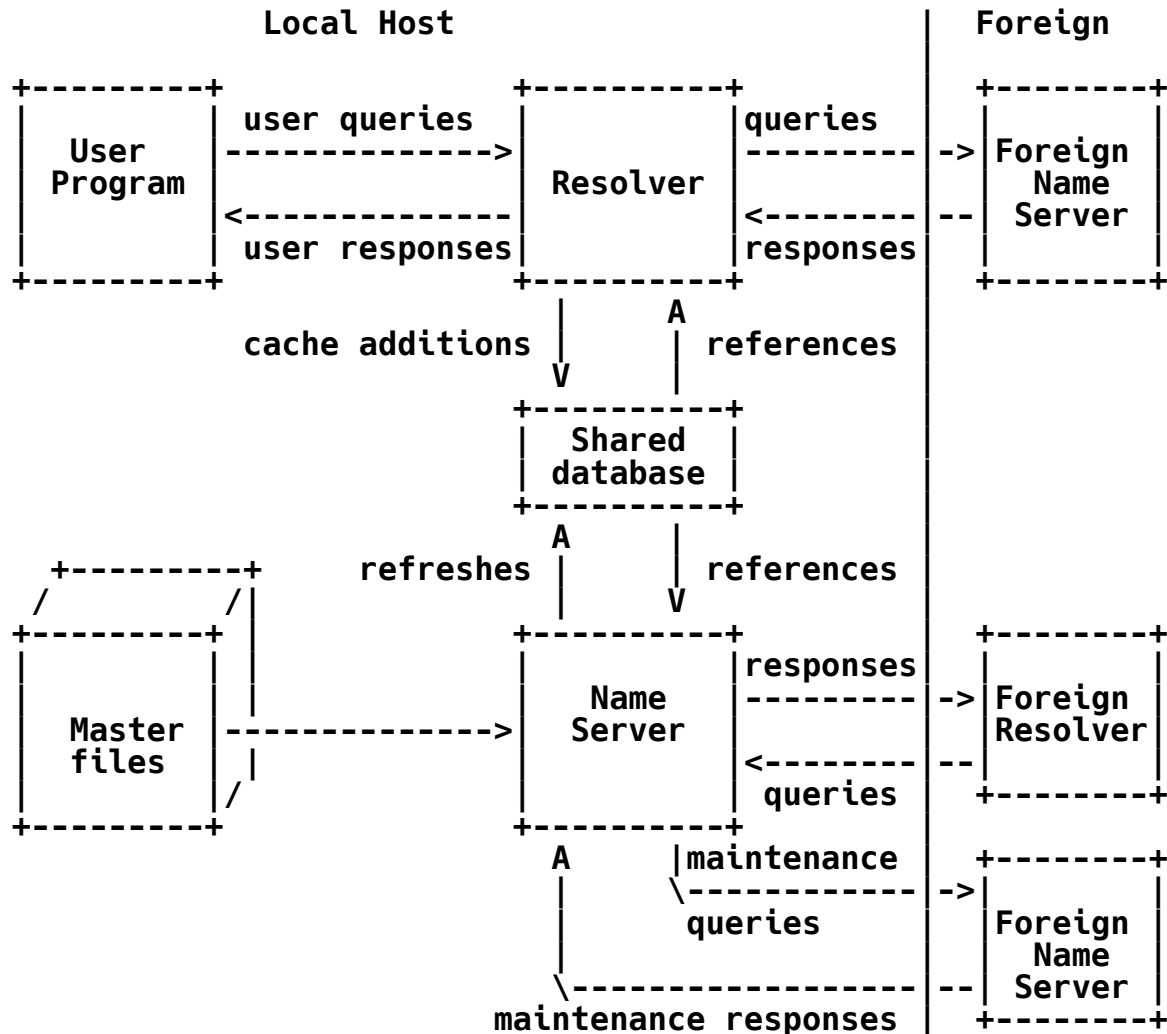   +--------+   +----------+ |

        Here the name server acquires information about one or more zones
        by reading master files from its local file system, and answers
        queries about those zones that arrive from foreign resolvers.

        A more sophisticated name server might acquire zones from foreign
        name servers as well as local master files. This configuration is
        shown below:

                    Local Host   Foreign
                                                         |
            +--------+ |
           /   /| |
          +--------+ | +----------+ |
          | | | | |responses | +--------+
          | | | | Name |---------- |->| |
          | Master | |--------------->| Server | | | Foreign|
          | files | | | |<-------- |--|Resolver|
          | |/ | | queries | +--------+
          +--------+   +----------+ |
                                 A |maintenance | +--------+
                                 | \------------ |->| |
                                 | queries | | Foreign|
                                 | | | Name |
                                 \------------------ |--| Server |
                               maintenance responses | +--------+

        In this configuration, the name server periodically establishes a
        virtual circuit to a foreign name server to acquire a copy of a
        zone or to check that an existing copy has not changed. The
        messages sent for these maintenance activities follow the same
        form as queries and responses, but the message sequences are
        somewhat different.

The information flow in a host that supports all aspects of the
domain name system is shown below:

```
                  Local Host                    |    Foreign
                                                |
    +---------+               +----------+      |      +--------+
    |         | user queries  |          |queries      |        |
    |  User   |-------------->|          |--------- ->|Foreign |
    | Program |               | Resolver |             | Name   |
    |         |<--------------|          |<-------- --| Server |
    |         | user responses|          |responses    |        |
    +---------+               +----------+      |      +--------+
                                  |    A                |
                  cache additions |    | references     |
                                  V    |                |
                              +----------+              |
                              | Shared   |              |
                              | database |              |
                              +----------+              |
                                  A    |                |
    +---------+    refreshes     |    | references      |
   /         /|                  |    V                 |
  +---------+ |                +----------+             +--------+
  |         | |                |          |responses    |        |
  |         | |                |  Name    |--------- ->|Foreign |
  | Master  |-|-------------->|  Server  |             |Resolver|
  | files   | |                |          |<-------- --|        |
  |         | /                |          |queries      +--------+
  +---------+                  +----------+
                                  A    |maintenance   +--------+
                                  |     \-----------  ->|        |
                                  |       queries       |Foreign |
                                  |                      | Name   |
                                  \------------------- --| Server |
                        maintenance responses    |      +--------+
```

The shared database holds domain space data for the local name
server and resolver. The contents of the shared database will
typically be a mixture of authoritative data maintained by the
periodic refresh operations of the name server and cached data
from previous resolver requests. The structure of the domain data
and the necessity for synchronization between name servers and
resolvers imply the general characteristics of this database, but
the actual format is up to the local implementer. This memo
suggests a multiple tree format.

This memo divides the implementation discussion into sections:

NAME SERVER TRANSACTIONS, which discusses the formats for name
servers queries and the corresponding responses.

NAME SERVER MAINTENANCE, which discusses strategies,
algorithms, and formats for maintaining the data residing in
name servers. These services periodically refresh the local
copies of zones that originate in other hosts.

RESOLVER ALGORITHMS, which discusses the internal structure of
resolvers. This section also discusses data base sharing
between a name server and a resolver on the same host.

DOMAIN SUPPORT FOR MAIL, which discusses the use of the domain
system to support mail transfer.

   Conventions

      The domain system has several conventions dealing with low-level,
      but fundamental, issues. While the implementer is free to violate
      these conventions WITHIN HIS OWN SYSTEM, he must observe these
      conventions in ALL behavior observed from other hosts.

            ********** Data Transmission Order **********

      The order of transmission of the header and data described in this
      document is resolved to the octet level. Whenever a diagram shows
      a group of octets, the order of transmission of those octets is
      the normal order in which they are read in English. For example,
      in the following diagram the octets are transmitted in the order
      they are numbered.

                     0   1
                     0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5
                    +-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
                    | 1 | 2 |
                    +-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
                    | 3 | 4 |
                    +-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
                    | 5 | 6 |
                    +-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+

                        Transmission Order of Bytes

      Whenever an octet represents a numeric quantity the left most bit
      in the diagram is the high order or most significant bit. That
      is, the bit labeled 0 is the most significant bit. For example,
      the following diagram represents the value 170 (decimal).

                          0 1 2 3 4 5 6 7
                         +-+-+-+-+-+-+-+-+
                         |1 0 1 0 1 0 1 0|
                         +-+-+-+-+-+-+-+-+

                          Significance of Bits

      Similarly, whenever a multi-octet field represents a numeric
      quantity the left most bit of the whole field is the most
      significant bit. When a multi-octet quantity is transmitted the
      most significant octet is transmitted first.

              ********** Character Case **********

    All comparisons between character strings (e.g. labels, domain
    names, etc.) are done in a case-insensitive manner.

    When data enters the domain system, its original case should be
    preserved whenever possible. In certain circumstances this cannot
    be done. For example, if two domain names x.y and X.Y are entered
    into the domain database, they are interpreted as the same name,
    and hence may have a single representation. The basic rule is
    that case can be discarded only when data is used to define
    structure in a database, and two names are identical when compared
    in a case insensitive manner.

    Loss of case sensitive data must be minimized. Thus while data
    for x.y and X.Y may both be stored under x.y, data for a.x and B.X
    can be stored as a.x and B.x, but not A.x, A.X, b.x, or b.X. In
    general, this prevents the first component of a domain name from
    loss of case information.

    Systems administrators who enter data into the domain database
    should take care to represent the data they supply to the domain
    system in a case-consistent manner if their system is
    case-sensitive. The data distribution system in the domain system
    will ensure that consistent representations are preserved.

    Design philosophy

       The design presented in this memo attempts to provide a base which
       will be suitable for several existing networks. An equally
       important goal is to provide these services within a framework
       that is capable of adjustment to fit the evolution of services in
       early clients as well as to accommodate new networks.

       Since it is impossible to predict the course of these
       developments, the domain system attempts to provide for evolution
       in the form of an extensible framework. This section describes
       the areas in which we expect to see immediate evolution.

       DEFINING THE DATABASE

       This memo defines methods for partitioning the database and data
       for host names, host addresses, gateway information, and mail
       support. Experience with this system will provide guidance for
       future additions.

       While the present system allows for many new RR types, classes,
       etc., we feel that it is more important to get the basic services
       in operation than to cover an exhaustive set of information.
       Hence we have limited the data types to those we felt were
       essential, and would caution designers to avoid implementations
       which are based on the number of existing types and classes.
       Extensibility in this area is very important.

       While the domain system provides techniques for partitioning the
       database, policies for administrating the orderly connection of
       separate domains and guidelines for constructing the data that
       makes up a particular domain will be equally important to the
       success of the system. Unfortunately, we feel that experience
       with prototype systems will be necessary before this question can
       be properly addressed. Thus while this memo has minimal
       discussion of these issues, it is a critical area for development.

       TYING TOGETHER INTERNETS

       Although it is very difficult to characterize the types of
       networks, protocols, and applications that will be clients of the
       domain system, it is very obvious that some of these applications
       will cross the boundaries of network and protocol. At the very
       least, mail is such a service.

       Attempts to unify two such systems must deal with two major
       problems:

       1. Differing formats for environment sensitive data. For example,

network addresses vary in format, and it is unreasonable to
expect to enforce consistent conventions.

2. Connectivity may require intermediaries. For example, it is a
   frequent occurence that mail is sent between hosts that share
   no common protocol.

The domain system acknowledges that these are very difficult
problems, and attempts to deal with both problems through its
CLASS mechanism:

1. The CLASS field in RRs allows data to be tagged so that all
   programs in the domain system can identify the format in use.

2. The CLASS field allows the requestor to identify the format of
   data which can be understood by the requestor.

3. The CLASS field guides the search for the requested data.

The last point is central to our approach. When a query crosses
protocol boundaries, it must be guided though agents capable of
performing whatever translation is required. For example, when a
mailer wants to identify the location of a mailbox in a portion of
the domain system that doesn't have a compatible protocol, the
query must be guided to a name server that can cross the boundary
itself or form one link in a chain that can span the differences.

If query and response transport were the only problem, then this
sort of problem could be dealt with in the name servers
themselves. However, the applications that will use domain
service have similar problems. For example, mail may need to be
directed through mail gateways, and the characteristics of one of
the environments may not permit frequent connectivity between name
servers in all environments.

These problems suggest that connectivity will be achieved through
a variety of measures:

    Translation name servers that act as relays between different
    protocols.

    Translation application servers that translate application
    level transactions.

    Default database entries that route traffic through application
    level forwarders in ways that depend on the class of the
    requestor.

While this approach seems best given our current understanding of

the problem, we realize that the approach of using resource data
that transcends class may be appropriate in future designs or
applications. By not defining class to be directly related to
protocol, network, etc., we feel that such services could be added
by defining a new "universal" class, while the present use of
class will provide immediate service.

This problem requires more thought and experience before solutions
can be discovered. The concepts of CLASS, recursive servers and
other mechanisms are intended as tools for acquiring experience
and not as final solutions.

NAME SERVER TRANSACTIONS

    Introduction

        The primary purpose of name servers is to receive queries from
        resolvers and return responses. The overall model of this service
        is that a program (typically a resolver) asks the name server
        questions (queries) and gets responses that either answer the
        question or refer the questioner to another name server. Other
        functions related to name server database maintenance use similar
        procedures and formats and are discussed in a section later in
        this memo.

        There are three kinds of queries presently defined:

            1. Standard queries that ask for a specified resource attached
               to a given domain name.

            2. Inverse queries that specify a resource and ask for a domain
               name that possesses that resource.

            3. Completion queries that specify a partial domain name and a
               target domain and ask that the partial domain name be
               completed with a domain name close to the target domain.

        This memo uses an unqualified reference to queries to refer to
        either all queries or standard queries when the context is clear.

    Query and response transport

        Name servers and resolvers use a single message format for all
        communications. The message format consists of a variable-length
        octet string which includes binary values.

        The messages used in the domain system are designed so that they
        can be carried using either datagrams or virtual circuits. To
        accommodate the datagram style, all responses carry the query as
        part of the response.

        While the specification allows datagrams to be used in any
        context, some activities are ill suited to datagram use. For
        example, maintenance transactions and recursive queries typically
        require the error control of virtual circuits. Thus datagram use
        should be restricted to simple queries.

        The domain system assumes that a datagram service provides:

            1. A non-reliable (i.e. best effort) method of transporting a
               message of up to 512 octets.

        Hence datagram messages are limited to 512 octets. If a
        datagram message would exceed 512 octets, it is truncated
        and a truncation flag is set in its header.

    2. A message size that gives the number of octets in the
       datagram.

    The main implications for programs accessing name servers via
    datagrams are:

    1. Datagrams should not be used for maintenance transactions
       and recursive queries.

    2. Since datagrams may be lost, the originator of a query must
       perform error recovery (such as retransmissions) as
       appropriate.

    3. Since network or host delay may cause retransmission when a
       datagram has not been lost, the originator of a query must
       be ready to deal with duplicate responses.

    The domain system assumes that a virtual circuit service provides:

    1. A reliable method of transmitting a message of up to 65535
       octets.

    2. A message size that gives the number of octets in the
       message.

       If the virtual circuit service does not provide for message
       boundary detection or limits transmission size to less than
       65535 octets, then messages are prefaced with an unsigned 16
       bit length field and broken up into separate transmissions
       as required. The length field is only prefaced on the first
       message. This technique is used for TCP virtual circuits.

    3. Multiple messages may be sent over a virtual circuit.

    4. A method for closing a virtual circuit.

    5. A method for detecting that the other party has requested
       that the virtual circuit be closed.

    The main implications for programs accessing name servers via
    virtual circuits are:

    1. Either end of a virtual circuit may initiate a close when
       there is no activity in progress. The other end should
       comply.

         The decision to initiate a close is a matter of individual
         site policy; some name servers may leave a virtual circuit
         open for an indeterminate period following a query to allow
         for subsequent queries; other name servers may choose to
         initiate a close following the completion of the first query
         on a virtual circuit. Of course, name servers should not
         close the virtual circuit in the midst of a multiple message
         stream used for zone transfer.

      2. Since network delay may cause one end to erroneously believe
         that no activity is in progress, a program which receives a
         virtual circuit close while a query is in progress should
         close the virtual circuit and resubmit the query on a new
         virtual circuit.

   All messages may use a compression scheme to reduce the space
   consumed by repetitive domain names. The use of the compression
   scheme is optional for the sender of a message, but all receivers
   must be capable of decoding compressed domain names.

 Overall message format

   All messages sent by the domain system are divided into 5 sections
   (some of which are empty in certain cases) shown below:

      +---------------------+
      | Header |
      +---------------------+
      | Question | the question for the name server
      +---------------------+
      | Answer | answering resource records (RRs)
      +---------------------+
      | Authority | RRs pointing toward an authority
      +---------------------+
      | Additional | RRs holding pertinent information
      +---------------------+

   The header section is always present. The header includes fields
   that specify which of the remaining sections are present, and also
   specify whether the message is a query, inverse query, completion
   query, or response.

   The question section contains fields that describe a question to a
   name server. These fields are a query type (QTYPE), a query class
   (QCLASS), and a query domain name (QNAME).

   The last three sections have the same format: a possibly empty
   list of concatenated resource records (RRs). The answer section
   contains RRs that answer the question; the authority section

contains RRs that point toward an authoritative name server; the
additional records section contains RRs which relate to the query,
but are not strictly answers for the question.

The next two sections of this memo illustrate the use of these
message sections through examples; a detailed discussion of data
formats follows the examples.

    The contents of standard queries and responses

        When a name server processes a standard query, it first determines
        whether it is an authority for the domain name specified in the
        query.

        If the name server is an authority, it returns either:

            1. the specified resource information

            2. an indication that the specified name does not exist

            3. an indication that the requested resource information does
               not exist

        If the name server is not an authority for the specified name, it
        returns whatever relevant resource information it has along with
        resource records that the requesting resolver can use to locate an
        authoritative name server.

    Standard query and response example

        The overall structure of a query for retrieving information for
        Internet mail for domain F.ISI.ARPA is shown below:

```
                  +---------------------------------------------+
     Header       |            OPCODE=QUERY, ID=2304            |
                  +---------------------------------------------+
     Question     |QTYPE=MAILA, QCLASS=IN, QNAME=F.ISI.ARPA |
                  +---------------------------------------------+
     Answer       |                   <empty>                   |
                  +---------------------------------------------+
     Authority    |                   <empty>                   |
                  +---------------------------------------------+
     Additional   |                   <empty>                   |
                  +---------------------------------------------+
```

        The header includes an opcode field that specifies that this
        datagram is a query, and an ID field that will be used to
        associate replies with the original query. (Some additional
        header fields have been omitted for clarity.) The question
        section specifies that the type of the query is for mail agent
        information, that only ARPA Internet information is to be
        considered, and that the domain name of interest is F.ISI.ARPA.
        The remaining sections are empty, and would not use any octets in
        a real query.

One possible response to this query might be:

```
              +------------------------------------------+
   Header     |          OPCODE=RESPONSE, ID=2304        |
              +------------------------------------------+
   Question   |QTYPE=MAILA, QCLASS=IN, QNAME=F.ISI.ARPA  |
              +------------------------------------------+
   Answer     |                 <empty>                  |
              +------------------------------------------+
   Authority  |          ARPA NS IN A.ISI.ARPA           |
              |               -------                    |
              |          ARPA NS IN F.ISI.ARPA           |
              +------------------------------------------+
   Additional |         F.ISI.ARPA A IN 10.2.0.52        |
              |               -------                    |
              |         A.ISI.ARPA A IN 10.1.0.22        |
              +------------------------------------------+
```

This type of response would be returned by a name server that was
not an authority for the domain name F.ISI.ARPA. The header field
specifies that the datagram is a response to a query with an ID of
2304. The question section is copied from the question section in
the query datagram.

The answer section is empty because the name server did not have
any information that would answer the query. (Name servers may
happen to have cached information even if they are not
authoritative for the query.)

The best that this name server could do was to pass back
information for the domain ARPA. The authority section specifies
two name servers for the domain ARPA using the Internet family:
A.ISI.ARPA and F.ISI.ARPA. Note that it is merely a coincidence
that F.ISI.ARPA is a name server for ARPA as well as the subject
of the query.

In this case, the name server included in the additional records
section the Internet addresses for the two hosts specified in the
authority section. Such additional data is almost always
available.

Given this response, the process that originally sent the query
might resend the query to the name server on A.ISI.ARPA, with a
new ID of 2305.

The name server on A.ISI.ARPA might return a response:

```
           +------------------------------------------+
Header     |          OPCODE=RESPONSE, ID=2305        |
           +------------------------------------------+
Question   |QTYPE=MAILA, QCLASS=IN, QNAME=F.ISI.ARPA  |
           +------------------------------------------+
Answer     |          F.ISI.ARPA MD IN F.ISI.ARPA     |
           |                  -------                  |
           |          F.ISI.ARPA MF IN A.ISI.ARPA     |
           +------------------------------------------+
Authority  |                  <empty>                 |
           +------------------------------------------+
Additional |          F.ISI.ARPA A IN 10.2.0.52       |
           |                  -------                  |
           |          A.ISI.ARPA A IN 10.1.0.22       |
           +------------------------------------------+
```

This query was directed to an authoritative name server, and hence
the response includes an answer but no authority records. In this
case, the answer section specifies that mail for F.ISI.ARPA can
either be delivered to F.ISI.ARPA or forwarded to A.ISI.ARPA. The
additional records section specifies the Internet addresses of
these hosts.

The contents of inverse queries and responses

Inverse queries reverse the mappings performed by standard query
operations; while a standard query maps a domain name to a
resource, an inverse query maps a resource to a domain name. For
example, a standard query might bind a domain name to a host
address; the corresponding inverse query binds the host address to
a domain name.

Inverse query mappings are not guaranteed to be unique or complete
because the domain system does not have any internal mechanism for
determining authority from resource records that parallels the
capability for determining authority as a function of domain name.
In general, resolvers will be configured to direct inverse queries
to a name server which is known to have the desired information.

Name servers are not required to support any form of inverse
queries; it is anticipated that most name servers will support
address to domain name conversions, but no other inverse mappings.
If a name server receives an inverse query that it does not
support, it returns an error response with the "Not Implemented"
error set in the header. While inverse query support is optional,
all name servers must be at least able to return the error
response.

When a name server processes an inverse query, it either returns:

    1. zero, one, or multiple domain names for the specified
       resource

    2. an error code indicating that the name server doesn't
       support inverse mapping of the specified resource type.

  Inverse query and response example

    The overall structure of an inverse query for retrieving the
    domain name that corresponds to Internet address 10.2.0.52 is
    shown below:

```
                +------------------------------------------+
     Header     |         OPCODE=IQUERY, ID=997            |
                +------------------------------------------+
     Question   |                 <empty>                  |
                +------------------------------------------+
     Answer     |          <anyname> A IN 10.2.0.52        |
                +------------------------------------------+
     Authority  |                 <empty>                  |
                +------------------------------------------+
     Additional |                 <empty>                  |
                +------------------------------------------+
```

    This query asks for a question whose answer is the Internet style
    address 10.2.0.52. Since the owner name is not known, any domain
    name can be used as a placeholder (and is ignored). The response
    to this query might be:

```
                +------------------------------------------+
     Header     |        OPCODE=RESPONSE, ID=997           |
                +------------------------------------------+
     Question   |   QTYPE=A, QCLASS=IN, QNAME=F.ISI.ARPA   |
                +------------------------------------------+
     Answer     |       F.ISI.ARPA A IN 10.2.0.52          |
                +------------------------------------------+
     Authority  |                 <empty>                  |
                +------------------------------------------+
     Additional |                 <empty>                  |
                +------------------------------------------+
```

    Note that the QTYPE in a response to an inverse query is the same
    as the TYPE field in the answer section of the inverse query.
    Responses to inverse queries may contain multiple questions when
    the inverse is not unique.

     Completion queries and responses

          Completion queries ask a name server to complete a partial domain
          name and return a set of RRs whose domain names meet a specified
          set of criteria for "closeness" to the partial input. This type
          of query can provide a local shorthand for domain names or command
          completion similar to that in TOPS-20.

          Implementation of completion query processing is optional in a
          name server. However, a name server must return a "Not
          Implemented" (NI) error response if it does not support
          completion.

          The arguments in a completion query specify:

          1. A type in QTYPE that specifies the type of the desired name.
             The type is used to restrict the type of RRs which will match
             the partial input so that completion queries can be used for
             mailbox names, host names, or any other type of RR in the
             domain system without concern for matches to the wrong type of
             resource.

          2. A class in QCLASS which specifies the desired class of the RR.

          3. A partial domain name that gives the input to be completed.
             All returned RRs will begin with the partial string. The
             search process first looks for names which qualify under the
             assumption that the partial string ends with a full label
             ("whole label match"); if this search fails, the search
             continues under the assumption that the last label in the
             partial sting may be an incomplete label ("partial label
             match"). For example, if the partial string "Smith" was used
             in a mailbox completion, it would match Smith@ISI.ARPA in
             preference to Smithsonian@ISI.ARPA.

             The partial name is supplied by the user through the user
             program that is using domain services. For example, if the
             user program is a mail handler, the string might be "Mockap"
             which the user intends as a shorthand for the mailbox
             Mockapetris@ISI.ARPA; if the user program is TELNET, the user
             might specify "F" for F.ISI.ARPA.

             In order to make parsing of messages consistent, the partial
             name is supplied in domain name format (i.e. a sequence of
             labels terminated with a zero length octet). However, the
             trailing root label is ignored during matching.

          4. A target domain name which specifies the domain which is to be
             examined for matches. This name is specified in the additional

section using a NULL RR. All returned names will end with the
target name.

The user program which constructs the query uses the target
name to restrict the search. For example, user programs
running at ISI might restrict completion to names that end in
ISI.ARPA; user programs running at MIT might restrict
completion to the domain MIT.ARPA.

The target domain name is also used by the resolver to
determine the name server which should be used to process the
query. In general, queries should be directed to a name server
that is authoritative for the target domain name. User
programs which wish to provide completion for a more than one
target can issue multiple completion queries, each directed at
a different target. Selection of the target name and the
number of searches will depend on the goals of the user
program.

5. An opcode for the query. The two types of completion queries
   are "Completion Query - Multiple", or CQUERYM, which asks for
   all RRs which could complete the specified input, and
   "Completion Query - Unique", or CQUERYU, which asks for the
   "best" completion.

   CQUERYM is used by user programs which want to know if
   ambiguities exist or wants to do its own determinations as to
   the best choice of the available candidates.

   CQUERYU is used by user programs which either do not wish to
   deal with multiple choices or are willing to use the closeness
   criteria used by CQUERYU to select the best match.

When a name server receives either completion query, it first
looks for RRs that begin (on the left) with the same labels as are
found in QNAME (with the root deleted), and which match the QTYPE
and QCLASS. This search is called "whole label" matching. If one
or more hits are found the name server either returns all of the
hits (CQUERYM) or uses the closeness criteria described below to
eliminate all but one of the matches (CQUERYU).

If the whole label match fails to find any candidates, then the
name server assumes that the rightmost label of QNAME (after root
deletion) is not a complete label, and looks for candidates that
would match if characters were added (on the right) to the
rightmost label of QNAME. If one or more hits are found the name
server either returns all of the hits (CQUERYM) or uses the
closeness criteria described below to eliminate all but one of the
matches (CQUERYU).

If a CQUERYU query encounters multiple hits, it uses the following
sequence of rules to discard multiple hits:

1. Discard candidates that have more labels than others. Since
   all candidates start with the partial name and end with the
   target name, this means that we select those entries that
   require the fewest number of added labels. For example, a host
   search with a target of "ISI.ARPA" and a partial name of "A"
   will select A.ISI.ARPA in preference to A.IBM-PCS.ISI.ARPA.

2. If partial label matching was used, discard those labels which
   required more characters to be added. For example, a mailbox
   search for partial "X" and target "ISI.ARPA" would prefer
   XX@ISI.ARPA to XYZZY@ISI.ARPA.

If multiple hits are still present, return all hits.

Completion query mappings are not guaranteed to be unique or
complete because the domain system does not have any internal
mechanism for determining authority from a partial domain name
that parallels the capability for determining authority as a
function of a complete domain name. In general, resolvers will be
configured to direct completion queries to a name server which is
known to have the desired information.

When a name server processes a completion query, it either
returns:

1. An answer giving zero, one, or more possible completions.

2. an error response with Not Implemented (NI) set.

   Completion query and response example

      Suppose that the completion service was used by a TELNET program
      to allow a user to specify a partial domain name for the desired
      host. Thus a user might ask to be connected to "B". Assuming
      that the query originated from an ISI machine, the query might
      look like:

```
                     +------------------------------------------+
       Header        |          OPCODE=CQUERYU, ID=409          |
                     +------------------------------------------+
       Question      |        QTYPE=A, QCLASS=IN, QNAME=B        |
                     +------------------------------------------+
       Answer        |                 <empty>                  |
                     +------------------------------------------+
       Authority     |                 <empty>                  |
                     +------------------------------------------+
       Additional    |             ISI.ARPA NULL IN             |
                     +------------------------------------------+
```

      The partial name in the query is "B", the mappings of interest are
      ARPA Internet address records, and the target domain is ISI.ARPA.
      Note that NULL is a special type of NULL resource record that is
      used as a placeholder and has no significance; NULL RRs obey the
      standard format but have no other function.

      The response to this completion query might be:

```
                     +------------------------------------------+
       Header        |          OPCODE=RESPONSE, ID=409         |
                     +------------------------------------------+
       Question      |        QTYPE=A, QCLASS=IN, QNAME=B        |
                     +------------------------------------------+
       Answer        |         B.ISI.ARPA A IN 10.3.0.52        |
                     +------------------------------------------+
       Authority     |                 <empty>                  |
                     +------------------------------------------+
       Additional    |             ISI.ARPA NULL IN             |
                     +------------------------------------------+
```

      This response has completed B to mean B.ISI.ARPA.

Another query might be:

```
              +------------------------------------------+
      Header  |         OPCODE=CQUERYM, ID=410           |
              +------------------------------------------+
    Question  |        QTYPE=A, QCLASS=IN, QNAME=B        |
              +------------------------------------------+
      Answer  |                <empty>                    |
              +------------------------------------------+
   Authority  |                <empty>                    |
              +------------------------------------------+
  Additional  |              ARPA NULL IN                 |
              +------------------------------------------+
```

This query is similar to the previous one, but specifies a target
of ARPA rather than ISI.ARPA. It also allows multiple matches.
In this case the same name server might return:

```
              +------------------------------------------+
      Header  |        OPCODE=RESPONSE, ID=410           |
              +------------------------------------------+
    Question  |       QTYPE=A, QCLASS=IN, QNAME=B         |
              +------------------------------------------+
      Answer  |       B.ISI.ARPA A IN 10.3.0.52          |
              |                    -                     |
              |       B.BBN.ARPA A IN 10.0.0.49          |
              |                    -                     |
              |       B.BBNCC.ARPA A IN 8.1.0.2          |
              +------------------------------------------+
   Authority  |                <empty>                    |
              +------------------------------------------+
  Additional  |              ARPA NULL IN                 |
              +------------------------------------------+
```

This response contains three answers, B.ISI.ARPA, B.BBN.ARPA, and
B.BBNCC.ARPA.

   Recursive Name Service

      Recursive service is an optional feature of name servers.

      When a name server receives a query regarding a part of the name
      space which is not in one of the name server's zones, the standard
      response is a message that refers the requestor to another name
      server. By iterating on these referrals, the requestor eventually
      is directed to a name server that has the required information.

      Name servers may also implement recursive service. In this type
      of service, a name server either answers immediately based on
      local zone information, or pursues the query for the requestor and
      returns the eventual result back to the original requestor.

      A name server that supports recursive service sets the Recursion
      Available (RA) bit in all responses it generates. A requestor
      asks for recursive service by setting the Recursion Desired (RD)
      bit in queries. In some situations where recursive service is the
      only path to the desired information (see below), the name server
      may go recursive even if RD is zero.

      If a query requests recursion (RD set), but the name server does
      not support recursion, and the query needs recursive service for
      an answer, the name server returns a "Not Implemented" (NI) error
      code. If the query can be answered without recursion since the
      name server is authoritative for the query, it ignores the RD bit.

      Because of the difficulty in selecting appropriate timeouts and
      error handling, recursive service is best suited to virtual
      circuits, although it is allowed for datagrams.

      Recursive service is valuable in several special situations:

         In a system of small personal computers clustered around one or
         more large hosts supporting name servers, the recursive
         approach minimizes the amount of code in the resolvers in the
         personal computers. Such a design moves complexity out of the
         resolver into the name server, and may be appropriate for such
         systems.

         Name servers on the boundaries of different networks may wish
         to offer recursive service to create connectivity between
         different networks. Such name servers may wish to provide
         recursive service regardless of the setting of RD.

         Name servers that translate between domain name service and
         some other name service may wish to adopt the recursive style.
         Implicit recursion may be valuable here as well.

These concepts are still under development.

    Header section format

```
            +---------------------------------------------+
            |                                             |
            |             *****  WARNING  *****            |
            |                                             |
            |  The following format is preliminary and is |
            |  included for purposes of explanation only. In |
            |  particular, the size and position of the   |
            |  OPCODE, RCODE fields and the number and    |
            |  meaning of the single bit fields are subject |
            |  to change.                                 |
            |                                             |
            +---------------------------------------------+
```

    The header contains the following fields:

```
                                      1  1  1  1  1  1
      0  1  2  3  4  5  6  7  8  9  0  1  2  3  4  5
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                      ID                       |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |QR|   Opcode  |AA|TC|RD|RA|        |   RCODE   |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                    QDCOUNT                     |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                    ANCOUNT                     |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                    NSCOUNT                     |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                    ARCOUNT                     |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

    where:

    ID   - A 16 bit identifier assigned by the program that
              generates any kind of query. This identifier is copied
              into all replies and can be used by the requestor to
              relate replies to outstanding questions.

    QR   - A one bit field that specifies whether this message is a
              query (0), or a response (1).

    OPCODE - A four bit field that specifies kind of query in this
              message. This value is set by the originator of a query
              and copied into the response. The values are:

                  0   a standard query (QUERY)

                         1   an inverse query (IQUERY)

                         2   an completion query allowing multiple
                              answers (CQUERYM)

                         2   an completion query requesting a single
                              answer (CQUERYU)

                         4-15 reserved for future use

      AA   - Authoritative Answer - this bit is valid in responses,
                 and specifies that the responding name server
                 is an authority for the domain name in the
                 corresponding query.

      TC   - TrunCation - specifies that this message was truncated
                 due to length greater than 512 characters.
                 This bit is valid in datagram messages but not
                 in messages sent over virtual circuits.

      RD   - Recursion Desired - this bit may be set in a query and
                 is copied into the response. If RD is set, it
                 directs the name server to pursue the query
                 recursively. Recursive query support is
                 optional.

      RA   - Recursion Available - this be is set or cleared in a
                 response, and denotes whether recursive query
                 support is available in the name server.

      RCODE   - Response code - this 4 bit field is set as part of
                 responses. The values have the following
                 interpretation:

                         0   No error condition

                         1   Format error - The name server was unable
                              to interpret the query.

                         2   Server failure - The name server was unable
                              to process this query due to a problem with
                              the name server.

                         3   Name Error - Meaningful only for responses
                              from an authoritative name server, this
                              code signifies that the domain name
                              referenced in the query does not exist.

4   Not Implemented - The name server does not
        support the requested kind of query.

5   Refused - The name server refuses to
        perform the specified operation for policy
        reasons. For example, a name server may
        not wish to provide the information to the
        particular requestor, or a name server may
        not wish to perform a particular operation
        (e.g. zone transfer) for particular data.

6-15 Reserved for future use.

QDCOUNT - an unsigned 16 bit integer specifying the number of
          entries in the question section.

ANCOUNT - an unsigned 16 bit integer specifying the number of
          resource records in the answer section.

NSCOUNT - an unsigned 16 bit integer specifying the number of name
          server resource records in the authority records
          section.

ARCOUNT - an unsigned 16 bit integer specifying the number of
          resource records in the additional records section.

Question section format

     The question section is used in all kinds of queries other than
     inverse queries. In responses to inverse queries, this section
     may contain multiple entries; for all other responses it contains
     a single entry. Each entry has the following format:

```
                                     1 1 1 1 1 1
         0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                                               |
        /                     QNAME                     /
        /                                               /
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                     QTYPE                     |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                     QCLASS                    |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

     where:

     QNAME -   a variable number of octets that specify a domain name.
                This field uses the compressed domain name format
                described in the next section of this memo. This field
                can be used to derive a text string for the domain name.
                Note that this field may be an odd number of octets; no
                padding is used.

     QTYPE -   a two octet code which specifies the type of the query.
                The values for this field include all codes valid for a
                TYPE field, together with some more general codes which
                can match more than one type of RR. For example, QTYPE
                might be A and only match type A RRs, or might be MAILA,
                which matches MF and MD type RRs. The values for this
                field are listed in Appendix 2.

     QCLASS -   a two octet code that specifies the class of the query.
                For example, the QCLASS field is IN for the ARPA
                Internet, CS for the CSNET, etc. The numerical values
                are defined in Appendix 2.

Resource record format

     The answer, authority, and additional sections all share the same
     format: a variable number of resource records, where the number of
     records is specified in the corresponding count field in the
     header. Each resource record has the following format:
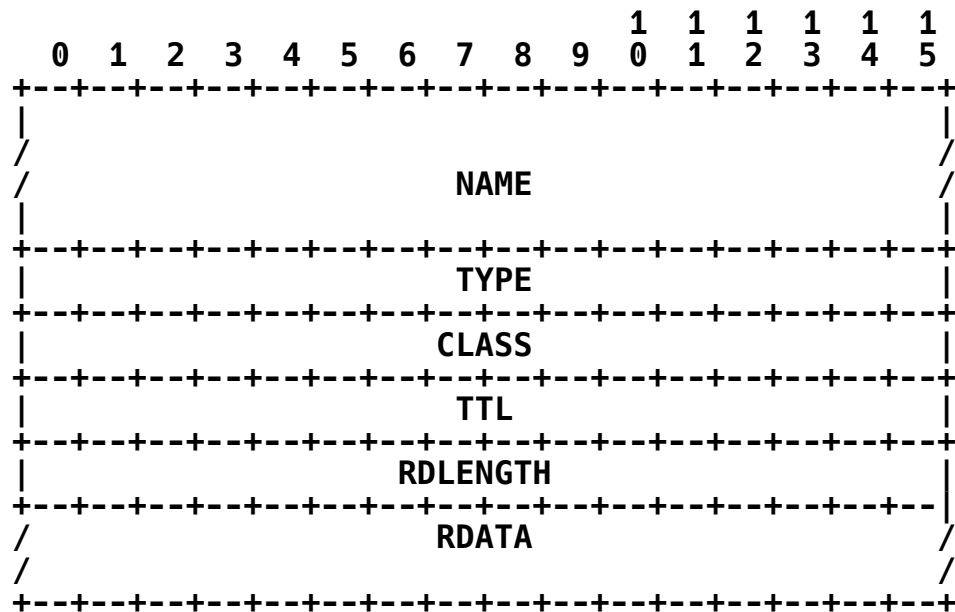
```
                                     1 1 1 1 1 1
         0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                                               |
        /                                               /
        /                      NAME                     /
        |                                               |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                      TYPE                     |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                      CLASS                    |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                      TTL                      |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                      RDLENGTH                 |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--|
        /                      RDATA                    /
        /                                               /
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

     where:

     NAME   - a compressed domain name to which this resource record
                 pertains.

     TYPE   - two octets containing one of the RR type codes defined
                 in Appendix 2. This field specifies the meaning of the
                 data in the RDATA field.

     CLASS   - two octets which specify the class of the data in the
                 RDATA field.

     TTL   - a 16 bit unsigned integer that specifies the time
                 interval (in seconds) that the resource record may be
                 cached before it should be discarded. Zero values are
                 interpreted to mean that the RR can only be used for the
                 transaction in progress, and should not be cached. For
                 example, SOA records are always distributed with a zero
                 TTL to prohibit caching. Zero values can also be used
                 for extremely volatile data.

RDLENGTH- an unsigned 16 bit integer that specifies the length in
          octets of the RDATA field.

RDATA   - a variable length string of octets that describes the
          resource. The format of this information varies
          according to the TYPE and CLASS of the resource record.
          For example, the if the TYPE is A and the CLASS is IN,
          the RDATA field is a 4 octet ARPA Internet address.

Formats for particular resource records are shown in Appendicies 2
and 3.

Domain name representation and compression

Domain names messages are expressed in terms of a sequence of
labels. Each label is represented as a one octet length field
followed by that number of octets. Since every domain name ends
with the null label of the root, a compressed domain name is
terminated by a length byte of zero. The high order two bits of
the length field must be zero, and the remaining six bits of the
length field limit the label to 63 octets or less.

To simplify implementations, the total length of label octets and
label length octets that make up a domain name is restricted to
255 octets or less. Since the trailing root label and its dot are
not printed, printed domain names are 254 octets or less.

Although labels can contain any 8 bit values in octets that make
up a label, it is strongly recommended that labels follow the
syntax described in Appendix 1 of this memo, which is compatible
with existing host naming conventions. Name servers and resolvers
must compare labels in a case-insensitive manner, i.e. A=a, and
hence all character strings must be ASCII with zero parity.
Non-alphabetic codes must match exactly.

Whenever possible, name servers and resolvers must preserve all 8
bits of domain names they process. When a name server is given
data for the same name under two different case usages, this
preservation is not always possible. For example, if a name
server is given data for ISI.ARPA and isi.arpa, it should create a
single node, not two, and hence will preserve a single casing of
the label. Systems with case sensitivity should take special
precautions to insure that the domain data for the system is
created with consistent case.

In order to reduce the amount of space used by repetitive domain
names, the sequence of octets that defines a domain name may be
terminated by a pointer to the length octet of a previously
specified label string. The label string that the pointer

specifies is appended to the already specified label string.
Exact duplication of a previous label string can be done with a
single pointer. Multiple levels are allowed.

Pointers can only be used in positions in the message where the
format is not class specific. If this were not the case, a name
server that was handling a RR for another class could make
erroneous copies of RRs. As yet, there are no such cases, but
they may occur in future RDATA formats.

If a domain name is contained in a part of the message subject to
a length field (such as the RDATA section of an RR), and
compression is used, the length of the compressed name is used in
the length calculation, rather than the length of the expanded
name.

Pointers are represented as a two octet field in which the high
order 2 bits are ones, and the low order 14 bits specify an offset
from the start of the message. The 01 and 10 values of the high
order bits are reserved for future use and should not be used.

Programs are free to avoid using pointers in datagrams they
generate, although this will reduce datagram capacity. However
all programs are required to understand arriving messages that
contain pointers.

For example, a datagram might need to use the domain names
F.ISI.ARPA, FOO.F.ISI.ARPA, ARPA, and the root. Ignoring the
other fields of the message, these domain names might be
represented as:

```
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     20 |              1            |              F            |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     22 |              3            |              I            |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     24 |              S            |              I            |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     26 |              4            |              A            |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     28 |              R            |              P            |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     30 |              A            |              0            |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     40 |              3            |              F            |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     42 |              0            |              0            |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     44 | 1  1|                     20                         |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     64 | 1  1|                     26                         |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
     92 |              0            |                           |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

The domain name for F.ISI.ARPA is shown at offset 20. The domain
name FOO.F.ISI.ARPA is shown at offset 40; this definition uses a
pointer to concatenate a label for FOO to the previously defined
F.ISI.ARPA. The domain name ARPA is defined at offset 64 using a
pointer to the ARPA component of the name F.ISI.ARPA at 20; note
that this reference relies on ARPA being the last label in the
string at 20. The root domain name is defined by a single octet
of zeros at 92; the root domain name has no labels.

Organization of the Shared database

    While name server implementations are free to use any internal
    data structures they choose, the suggested structure consists of
    several separate trees. Each tree has structure corresponding to
    the domain name space, with RRs attached to nodes and leaves.
    Each zone of authoritative data has a separate tree, and one tree
    holds all non-authoritative data. All of the trees corresponding
    to zones are managed identically, but the non-authoritative or
    cache tree has different management procedures.

Data stored in the database can be kept in whatever form is
convenient for the name server, so long as it can be transformed
back into the format needed for messages. In particular, the
database will probably use structure in place of expanded domain
names, and will also convert many of the time intervals used in
the domain systems to absolute local times.

Each tree corresponding to a zone has complete information for a
"pruned" subtree of the domain space. The top node of a zone has
a SOA record that marks the start of the zone. The bottom edge of
the zone is delimited by nodes containing NS records signifying
delegation of authority to other zones, or by leaves of the domain
tree. When a name server contains abutting zones, one tree will
have a bottom node containing a NS record, and the other tree will
begin with a tree location containing a SOA record.

Note that there is one special case that requires consideration
when a name server is implemented. A node that contains a SOA RR
denoting a start of zone will also have NS records that identify
the name servers that are expected to have a copy of the zone.
Thus a name server will usually find itself (and possibly other
redundant name servers) referred to in NS records occupying the
same position in the tree as SOA records. The solution to this
problem is to never interpret a NS record as delimiting a zone
started by a SOA at the same point in the tree. (The sample
programs in this memo deal with this problem by processing SOA
records only after NS records have been processed.)

Zones may also overlap a particular part of the name space when
they are of different classes.

Other than the abutting and separate class cases, trees are always
expected to be disjoint. Overlapping zones are regarded as a
non-fatal error. The scheme described in this memo avoids the
overlap issue by maintaining separate trees; other designs must
take the appropriate measures to defend against possible overlap.

Non-authoritative data is maintained in a separate tree. This
tree is unlike the zone trees in that it may have "holes". Each
RR in the cache tree has its own TTL that is separately managed.
The data in this tree is never used if authoritative data is
available from a zone tree; this avoids potential problems due to
cached data that conflicts with authoritative data.

The shared database will also contain data structures to support
the processing of inverse queries and completion queries if the
local system supports these optional features. Although many
schemes are possible, this memo describes a scheme that is based
on tables of pointers that invert the database according to key.

        Each kind of retrieval has a separate set of tables, with one
        table per zone. When a zone is updated, these tables must also be
        updated. The contents of these tables are discussed in the
        "Inverse query processing" and "Completion query processing"
        sections of this memo.

    The database implementation described here includes two locks that
    are used to control concurrent access and modification of the
    database by name server query processing, name server maintenance
    operations, and resolver access:

        The first lock ("main lock") controls access to all of the
        trees. Multiple concurrent reads are allowed, but write access
        can only be acquired by a single process. Read and write
        access are mutually exclusive. Resolvers and name server
        processes that answer queries acquire this lock in read mode,
        and unlock upon completion of the current message. This lock
        is acquired in write mode by a name server maintenance process
        when it is about to change data in the shared database. The
        actual update procedures are described under "NAME SERVER
        MAINTENANCE" but are designed to be brief.

        The second lock ("cache queue lock") controls access to the
        cache queue. This queue is used by a resolver that wishes to
        add information to the cache tree. The resolver acquires this
        lock, then places the RRs to be cached into the queue. The
        name server maintenance procedure periodically acquires this
        lock and adds the queue information to the cache. The
        rationale for this procedure is that it allows the resolver to
        operate with read-only access to the shared database, and
        allows the update process to batch cache additions and the
        associated costs for inversion calculations. The name server
        maintenance procedure must take appropriate precautions to
        avoid problems with data already in the cache, inversions, etc.

    This organization solves several difficulties:

        When searching the domain space for the answer to a query, a
        name server can restrict its search for authoritative data to
        that tree that matches the most labels on the right side of the
        domain name of interest.

        Since updates to a zone must be atomic with respect to
        searches, maintenance operations can simply acquire the main
        lock, insert a new copy of a particular zone without disturbing
        other zones, and then release the storage used by the old copy.
        Assuming a central table pointing to valid zone trees, this
        operation can be a simple pointer swap.

TTL management of zones can be performed using the SOA record
for the zone. This avoids potential difficulties if individual
RRs in a zone could be timed out separately. This issue is
discussed further in the maintenance section.

Query processing

The following algorithm outlines processing that takes place at a
name server when a query arrives:

1. Search the list of zones to find zones which have the same
   class as the QCLASS field in the query and have a top domain
   name that matches the right end of the QNAME field. If there
   are none, go to step 2. If there are more than one, pick the
   zone that has the longest match and go to step 3.

2. Since the zone search failed, the only possible RRs are
   contained in the non-authoritative tree. Search the cache tree
   for the NS record that has the same class as the QCLASS field
   and the largest right end match for domain name. Add the NS
   record or records to the authority section of the response. If
   the cache tree has RRs that are pertinent to the question
   (domain names match, classes agree, not timed-out, and the type
   field is relevant to the QTYPE), copy these RRs into the answer
   section of the response. The name server may also search the
   cache queue. Go to step 4.

3. Since this zone is the best match, the zone in which QNAME
   resides is either this zone or a zone to which this zone will
   directly or indirectly delegate authority. Search down the
   tree looking for a NS RR or the node specified by QNAME.

      If the node exists and has no NS record, copy the relevant
      RRs to the answer section of the response and go to step 4.

      If a NS RR is found, either matching a part or all of QNAME,
      then QNAME is in a delegated zone outside of this zone. If
      so, copy the NS record or records into the authority section
      of the response, and search the remainder of the zone for an
      A type record corresponding to the NS reference. If the A
      record is found, add it to the additional section. Go to
      step 2.

      If the node is not found and a NS is not found, there is no
      such name; set the Name error bit in the response and exit.

4. When this step is reached, the answer and authority sections
   are complete. What remains is to complete the additional
   section. This procedure is only possible if the name server

knows the data formats implied by the class of records in the
answer and authority sections. Hence this procedure is class
dependent. Appendix 3 discusses this procedure for Internet
class data.

While this algorithm deals with typical queries and databases,
several additions are required that will depend on the database
supported by the name server:

QCLASS=*

Special procedures are required when the QCLASS of the query is
"*". If the database contains several classes of data, the
query processing steps above are performed separately for each
CLASS, and the results are merged into a single response. The
name error condition is not meaningful for a QCLASS=* query.
If the requestor wants this information, it must test each
class independently.

If the database is limited to data of a particular class, this
operation can be performed by simply reseting the authoritative
bit in the response, and performing the query as if QCLASS was
the class used in the database.

* labels in database RRs

Some zones will contain default RRs that use * to match in
cases where the search fails for a particular domain name. If
the database contains these records then a failure must be
retried using * in place of one or more labels of the search
key. The procedure is to replace labels from the left with
"*"s looking for a match until either all labels have been
replaced, or a match is found. Note that these records can
never be the result of caching, so a name server can omit this
processing for zones that don't contain RRs with * in labels,
or can omit this processing entirely if * never appears in
local authoritative data.

Inverse query processing

Name servers that support inverse queries can support these
operations through exhaustive searches of their databases, but
this becomes impractical as the size of the database increases.
An alternative approach is to invert the database according to the
search key.

For name servers that support multiple zones and a large amount of
data, the recommended approach is separate inversions for each

zone. When a particular zone is changed during a refresh, only
its inversions need to be redone.

Support for transfer of this type of inversion may be included in
future versions of the domain system, but is not supported in this
version.

Completion query processing

Completion query processing shares many of the same problems in
data structure design as are found in inverse queries, but is
different due to the expected high rate of use of top level labels
(ie., ARPA, CSNET). A name server that wishes to be efficient in
its use of memory may well choose to invert only occurrences of
ARPA, etc. that are below the top level, and use a search for the
rare case that top level labels are used to constrain a
completion.

NAME SERVER MAINTENANCE

    Introduction

        Name servers perform maintenance operations on their databases to
        insure that the data they distribute is accurate and timely. The
        amount and complexity of the maintenance operations that a name
        server must perform are related to the size, change rate, and
        complexity of the database that the name server manages.

        Maintenance operations are fundamentally different for
        authoritative and non-authoritative data. A name server actively
        attempts to insure the accuracy and timeliness of authoritative
        data by refreshing the data from master copies. Non-authoritative
        data is merely purged when its time-to-live expires; the name
        server does not attempt to refresh it.

        Although the refreshing scheme is fairly simple to implement, it
        is somewhat less powerful than schemes used in other distributed
        database systems. In particular, an update to the master does not
        immediately update copies, and should be viewed as gradually
        percolating though the distributed database. This is adequate for
        the vast majority of applications. In situations where timliness
        is critical, the master name server can prohibit caching of copies
        or assign short timeouts to copies.

    Conceptual model of maintenance operations

        The vast majority of information in the domain system is derived
        from master files scattered among hosts that implement name
        servers; some name servers will have no master files, other name
        servers will have one or more master files. Each master file
        contains the master data for a single zone of authority rather
        than data for the whole domain name space. The administrator of a
        particular zone controls that zone by updating its master file.

        Master files and zone copies from remote servers may include RRs
        that are outside of the zone of authority when a NS record
        delegates authority to a domain name that is a descendant of the
        domain name at which authority is delegated. These forward
        references are a problem because there is no reasonable method to
        guarantee that the A type records for the delegatee are available
        unless they can somehow be attached to the NS records.

        For example, suppose the ARPA zone delegates authority at
        MIT.ARPA, and states that the name server is on AI.MIT.ARPA. If a
        resolver gets the NS record but not the A type record for
        AI.MIT.ARPA, it might try to ask the MIT name server for the
        address of AI.MIT.ARPA.

The solution is to allow type A records that are outside of the
zone of authority to be copied with the zone. While these records
won't be found in a search for the A type record itself, they can
be protected by the zone refreshing system, and will be passed
back whenever the name server passes back a referral to the
corresponding NS record. If a query is received for the A record,
the name server will pass back a referral to the name server with
the A record in the additional section, rather than answer
section.

The only exception to the use of master files is a small amount of
data stored in boot files. Boot file data is used by name servers
to provide enough resource records to allow zones to be imported
from foreign servers (e.g. the address of the server), and to
establish the name and address of root servers. Boot file records
establish the initial contents of the cache tree, and hence can be
overridden by later loads of authoritative data.

The data in a master file first becomes available to users of the
domain name system when it is loaded by the corresponding name
server. By definition, data from a master file is authoritative.

Other name servers which wish to be authoritative for a particular
zone do so by transferring a copy of the zone from the name server
which holds the master copy using a virtual circuit. These copies
include parameters which specify the conditions under which the
data in the copy is authoritative. In the most common case, the
conditions specify a refresh interval and policies to be followed
when the refresh operation cannot be performed.

A name server may acquire multiple zones from different name
servers and master files, but the name server must maintain each
zone separately from others and from non-authoritative data.

When the refresh interval for a particular zone copy expires, the
name server holding the copy must consult the name server that
holds the master copy. If the data in the zone has not changed,
the master name server instructs the copy name server to reset the
refresh interval. If the data has changed, the master passes a
new copy of the zone and its associated conditions to the copy
name server. Following either of these transactions, the copy
name server begins a new refresh interval.

Copy name servers must also deal with error conditions under which
they are unable to communicate with the name server that holds the
master copy of a particular zone. The policies that a copy name
server uses are determined by other parameters in the conditions
distributed with every copy. The conditions include a retry
interval and a maximum holding time. When a copy name server is

unable to establish communications with a master or is unable to
complete the refresh transaction, it must retry the refresh
operation at the rate specified by the retry interval. This retry
interval will usually be substantially shorter than the refresh
interval. Retries continue until the maximum holding time is
reached. At that time the copy name server must assume that its
copy of the data for the zone in question is no longer
authoritative.

Queries must be processed while maintenance operations are in
progress because a zone transfer can take a long time. However,
to avoid problems caused by access to partial databases, the
maintenance operations create new copies of data rather than
directly modifying the old copies. When the new copy is complete,
the maintenance process locks out queries for a short time using
the main lock, and switches pointers to replace the old data with
the new. After the pointers are swapped, the maintenance process
unlocks the main lock and reclaims the storage used by the old
copy.

Name server data structures and top level logic

The name server must multiplex its attention between multiple
activities. For example, a name server should be able to answer
queries while it is also performing refresh activities for a
particular zone. While it is possible to design a name server
that devotes a separate process to each query and refresh activity
in progress, the model described in this memo is based on the
assumption that there is a single process performing all
maintenance operations, and one or more processes devoted to
handling queries. The model also assumes the existence of shared
memory for several control structures, the domain database, locks,
etc.

The model name server uses the following files and shared data
structures:

   1. A configuration file that describes the master and boot
      files which the name server should load and the zones that
      the name server should attempt to load from foreign name
      servers. This file establishes the initial contents of the
      status table.

   2. Domain data files that contain master and boot data to be
      loaded.

   3. A status table that is derived from the configuration file.
      Each entry in this table describes a source of data. Each
      entry has a zone number. The zone number is zero for

non-authoritative sources; authoritative sources are
assigned separate non-zero numbers.

4. The shared database that holds the domain data. This
   database is assumed to be organized in some sort of tree
   structure paralleling the domain name space, with a list of
   resource records attached to each node and leaf in the tree.
   The elements of the resource record list need not contain
   the exact data present in the corresponding output format,
   but must contain data sufficient to create the output
   format; for example, these records need not contain the
   domain name that is associated with the resource because
   that name can be derived from the tree structure. Each
   resource record also internal data that the name server uses
   to organize its data.

5. Inversion data structures that allow the name server to
   process inverse queries and completion queries. Although
   many structures could be used, the implementation described
   in this memo supposes that there is one array for every
   inversion that the name server can handle. Each array
   contains a list of pointers to resource records such that
   the order of the inverted quantities is sorted.

6. The main and cache queue locks

7. The cache queue

The maintenance process begins by loading the status table from
the configuration file. It then periodically checks each entry,
to see if its refresh interval has elapsed. If not, it goes on to
the next entry. If so, it performs different operations depending
on the entry:

   If the entry is for zone 0, or the cache tree, the maintenance
   process checks to see if additions or deletions are required.
   Additions are acquired from the cache queue using the cache
   queue lock. Deletions are detected using TTL checks. If any
   changes are required, the maintenance process recalculates
   inversion data structures and then alters the cache tree under
   the protection of the main lock. Whenever the maintenance
   process modifies the cache tree, it resets the refresh interval
   to the minimum of the contained TTLs and the desired time
   interval for cache additions.

   If the entry is not zone 0, and the entry refers to a local
   file, the maintenance process checks to see if the file has
   been modified since its last load. If so the file is reloaded
   using the procedures specified under "Name server file

loading". The refresh interval is reset to that specified in
the SOA record if the file is a master file.

If the entry is for a remote master file, the maintenance
process checks for a new version using the procedure described
in "Names server remote zone transfer".

Name server file loading

Master files are kept in text form for ease of editing by system
maintainers. These files are not exchanged by name servers; name
servers use the standard message format when transferring zones.

Organizations that want to have a domain, but do not want to run a
name server, can use these files to supply a domain definition to
another organization that will run a name server for them. For
example, if organization X wants a domain but not a name server,
it can find another organization, Y, that has a name server and is
willing to provide service for X. Organization X defines domain X
via the master file format and ships a copy of the master file to
organization Y via mail, FTP, or some other method. A system
administrator at Y configures Y's name server to read in X's file
and hence support the X domain. X can maintain the master file
using a text editor and send new versions to Y for installation.

These files have a simple line-oriented format, with one RR per
line. Fields are separated by any combination of blanks and tab
characters. Tabs are treated the same as spaces; in the following
discussion the term "blank" means either a tab or a blank. A line
can be either blank (and ignored), a RR, or a $INCLUDE line.

If a RR line starts with a domain name, that domain name is used
to specify the location in the domain space for the record, i.e.
the owner. If a RR line starts with a blank, it is loaded into
the location specified by the most recent location specifier.

The location specifiers are assumed to be relative to some origin
that is provided by the user of a file unless the location
specifier contains the root label. This provides a convenient
shorthand notation, and can also be used to prevent errors in
master files from propagating into other zones. This feature is
particularly useful for master files imported from other sites.

An include line begins with $INCLUDE, starting at the first line
position, and is followed by a local file name and an optional
offset modifier. The filename follows the appropriate local
conventions. The offset is one or more labels that are added to
the offset in use for the file that contained the $INCLUDE. If
the offset is omitted, the included file is loaded using the

offset of the file that contained the $INCLUDE command. For
example, a file being loaded at offset ARPA might contain the
following lines:

        $INCLUDE <subsys>isi.data ISI
        $INCLUDE <subsys>addresses.data

The first line would be interpreted to direct loading of the file
<subsys>isi.data at offset ISI.ARPA. The second line would be
interpreted as a request to load data at offset ARPA.

Note that $INCLUDE commands do not cause data to be loaded into a
different zone or tree; they are simply ways to allow data for a
given zone to be organized in separate files. For example,
mailbox data might be kept separately from host data using this
mechanism.

Resource records are entered as a sequence of fields corresponding
to the owner name, TTL, CLASS, TYPE and RDATA components. (Note
that this order is different from the order used in examples and
the order used in the actual RRs; the given order allows easier
parsing and defaulting.)

    The owner name is derived from the location specifier.

    The TTL field is optional, and is expressed as a decimal
    number. If omitted TTL defaults to zero.

    The CLASS field is also optional; if omitted the CLASS defaults
    to the most recent value of the CLASS field in a previous RR.

    The RDATA fields depend on the CLASS and TYPE of the RR. In
    general, the fields that make up RDATA are expressed as decimal
    numbers or as domain names. Some exceptions exist, and are
    documented in the RDATA definitions in Appendicies 2 and 3 of
    this memo.

Because CLASS and TYPE fields don't contain any common
identifiers, and because CLASS and TYPE fields are never decimal
numbers, the parse is always unique.

Because these files are text files several special encodings are
necessary to allow arbitrary data to be loaded. In particular:

    .   A free standing dot is used to refer to the current domain
          name.

    @   A free standing @ is used to denote the current origin.

..   Two free standing dots represent the null domain name of
        the root.

\X   where X is any character other than a digit (0-9), is used
        to quote that character so that its special meaning does
        not apply. For example, "\." can be used to place a dot
        character in a label.

\DDD   where each D is a digit is the octet corresponding to the
        decimal number described by DDD. The resulting octet is
        assumed to be text and is not checked for special meaning.

( )   Parentheses are used to group data that crosses a line
        boundary. In effect, line terminations are not recognized
        within parentheses.

;   Semicolon is used to start a comment; the remainder of the
        line is ignored.

Name server file loading example

A name server for F.ISI.ARPA , serving as an authority for the
ARPA and ISI.ARPA domains, might use a boot file and two master
files. The boot file initializes some non-authoritative data, and
would be loaded without an origin:

```
..                9999999 IN      NS      B.ISI.ARPA
                  9999999 CS      NS      UDEL.CSNET
B.ISI.ARPA        9999999 IN      A       10.3.0.52
UDEL.CSNET        9999999 CS      A       302-555-0000
```

This file loads non-authoritative data which provides the
identities and addresses of root name servers. The first line
contains a NS RR which is loaded at the root; the second line
starts with a blank, and is loaded at the most recent location
specifier, in this case the root; the third and fourth lines load
RRs at B.ISI.ARPA and UDEL.CSNET, respectively. The timeouts are
set to high values (9999999) to prevent this data from being
discarded due to timeout.

The first master file loads authoritative data for the ARPA
domain. This file is designed to be loaded with an origin of
ARPA, which allows the location specifiers to omit the trailing
.ARPA labels.

```
@   IN  SOA     F.ISI.ARPA          Action.E.ISI.ARPA (
                                    20      ; SERIAL
                                    3600    ; REFRESH
                                    600     ; RETRY
                                    3600000; EXPIRE
                                    60)     ; MINIMUM
            NS      F.ISI.ARPA ; F.ISI.ARPA is a name server for ARPA
            NS      A.ISI.ARPA ; A.ISI.ARPA is a name server for ARPA
MIT         NS      AI.MIT.ARPA; delegation to MIT name server
ISI         NS      F.ISI.ARPA ; delegation to ISI name server

UDEL        MD      UDEL.ARPA
            A       10.0.0.96
NBS         MD      NBS.ARPA
            A       10.0.0.19
DTI         MD      DTI.ARPA
            A       10.0.0.12

AI.MIT      A       10.2.0.6
F.ISI       A       10.2.0.52
```

The first group of lines contains the SOA record and its
parameters, and identifies name servers for this zone and for
delegated zones. The Action.E.ISI.ARPA field is a mailbox
specification for the responsible person for the zone, and is the
domain name encoding of the mail destination Action@E.ISI.ARPA.
The second group specifies data for domain names within this zone.
The last group has forward references for name server address
resolution for AI.MIT.ARPA and F.ISI.ARPA. This data is not
technically within the zone, and will only be used for additional
record resolution for NS records used in referrals. However, this
data is protected by the zone timeouts in the SOA, so it will
persist as long as the NS references persist.

The second master file defines the ISI.ARPA environment, and is
loaded with an origin of ISI.ARPA:

```
@   IN  SOA     F.ISI.ARPA          Action\.ISI.E.ISI.ARPA (
                                    20      ; SERIAL
                                    7200    ; REFRESH
                                    600     ; RETRY
                                    3600000; EXPIRE
                                    60)     ; MINIMUM
            NS      F.ISI.ARPA ; F.ISI.ARPA is a name server
A           A       10.1.0.32
            MD      A.ISI.ARPA
            MF      F.ISI.ARPA
B           A       10.3.0.52
            MD      B.ISI.ARPA
```

```
        MF        F.ISI.ARPA
F       A         10.2.0.52
        MD        F.ISI.ARPA
        MF        A.ISI.ARPA
$INCLUDE <SUBSYS>ISI-MAILBOXES.TXT
```

   Where the file <SUBSYS>ISI-MAILBOXES.TXT is:

```
MOE      MB        F.ISI.ARPA
LARRY    MB        A.ISI.ARPA
CURLEY   MB        B.ISI.ARPA
STOOGES  MB        B.ISI.ARPA
         MG        MOE.ISI.ARPA
         MG        LARRY.ISI.ARPA
         MG        CURLEY.ISI.ARPA
```

   Note the use of the \ character in the SOA RR to specify the
   responsible person mailbox "Action.ISI@E.ISI.ARPA".

 Name server remote zone transfer

   When a name server needs to make an initial copy of a zone or test
   to see if a existing zone copy should be refreshed, it begins by
   attempting to open a virtual circuit to the foreign name server.

   If this open attempt fails, and this was an initial load attempt,
   it schedules a retry and exits. If this was a refresh operation,
   the name server tests the status table to see if the maximum
   holding time derived from the SOA EXPIRE field has elapsed. If
   not, the name server schedules a retry. If the maximum holding
   time has expired, the name server invalidates the zone in the
   status table, and scans all resource records tagged with this zone
   number. For each record it decrements TTL fields by the length of
   time since the data was last refreshed. If the new TTL value is
   negative, the record is deleted. If the TTL value is still
   positive, it moves the RR to the cache tree and schedules a retry.

   If the open attempt succeeds, the name server sends a query to the
   foreign name server in which QTYPE=SOA, QCLASS is set according to
   the status table information from the configuration file, and
   QNAME is set to the domain name of the zone of interest.

   The foreign name server will return either a SOA record indicating
   that it has the zone or an error. If an error is detected, the
   virtual circuit is closed, and the failure is treated in the same
   way as if the open attempt failed.

   If the SOA record is returned and this was a refresh, rather than
   an initial load of the zone, the name server compares the SERIAL

field in the new SOA record with the SERIAL field in the SOA
record of the existing zone copy. If these values match, the zone
has not been updated since the last copy and hence there is no
reason to recopy the zone. In this case the name server resets
the times in the existing SOA record and closes the virtual
circuit to complete the operation.

If this is initial load, or the SERIAL fields were different, the
name server requests a copy of the zone by sending the foreign
name server an AXFR query which specifies the zone by its QCLASS
and QNAME fields.

When the foreign name server receives the AXFR request, it sends
each node from the zone to the requestor in a separate message.
It begins with the node that contains the SOA record, walks the
tree in breadth-first order, and completes the transfer by
resending the node containing the SOA record.

Several error conditions are possible:

    If the AXFR request cannot be matched to a SOA, the foreign
    name server will return a single message in response that does
    not contain the AXFR request. (The normal SOA query preceding
    the AXFR is designed to avoid this condition, but it is still
    possible.)

    The foreign name server can detect an internal error or detect
    some other condition (e.g. system going down, out of resources,
    etc.) that forces the transfer to be aborted. If so, it sends
    a message with the "Server failure" condition set. If the AXFR
    can be immediately retried with some chance of success, it
    leaves the virtual open; otherwise it initiates a close.

    If the foreign name server doesn't wish to perform the
    operation for policy reasons (i.e. the system administrator
    wishes to forbid zone copies), the foreign server returns a
    "Refused" condition.

The requestor receives these records and builds a new tree. This
tree is not yet in the status table, so its data are not used to
process queries. The old copy of the zone, if any, may be used to
satisfy request while the transfer is in progress.

When the requestor receives the second copy of the SOA node, it
compares the SERIAL field in the first copy of the SOA against the
SERIAL field in the last copy of the SOA record. If these don't
match, the foreign server updated its zone while the transfer was
in progress. In this case the requestor repeats the AXFR request
to acquire the newer version.

If the AXFR transfer eventually succeeds, the name server closes
the virtual circuit and and creates new versions of inversion data
structures for this zone. When this operation is complete, the
name server acquires the main lock in write mode and then replaces
any old copy of the zone and inversion data structures with new
ones. The name server then releases the main lock, and can
reclaim the storage used by the old copy.

If an error occurs during the AXFR transfer, the name server can
copy any partial information into its cache tree if it wishes,
although it will not normally do so if the zone transfer was a
refresh rather than an initial load.

RESOLVER ALGORITHMS

   Operations

      Resolvers have a great deal of latitude in the semantics they
      allow in user calls. For example, a resolver might support
      different user calls that specify whether the returned information
      must be from and authoritative name server or not. Resolvers are
      also responsible for enforcement of any local restrictions on
      access, etc.

      In any case, the resolver will transform the user query into a
      number of shared database accesses and queries to remote name
      servers. When a user requests a resource associated with a
      particular domain name, the resolver will execute the following
      steps:

      1. The resolver first checks the local shared database, if any,
         for the desired information. If found, it checks the
         applicable timeout. If the timeout check succeeds, the
         information is used to satisfy the user request. If not, the
         resolver goes to step 2.

      2. In this step, the resolver consults the shared database for the
         name server that most closely matches the domain name in the
         user query. Multiple redundant name servers may be found. The
         resolver goes to step 3.

      3. In this step the resolver chooses one of the available name
         servers and sends off a query. If the query fails, it tries
         another name server. If all fail, an error indication is
         returned to the user. If a reply is received the resolver adds
         the returned RRs to its database and goes to step 4.

      4. In this step, the resolver interprets the reply. If the reply
         contains the desired information, the resolver returns the
         information to the user. The the reply indicates that the
         domain name in the user query doesn't exist, then the resolver
         returns an error to the user. If the reply contains a
         transient name server failure, the resolver can either wait and
         retry the query or go back to step 3 and try a different name
         server. If the reply doesn't contain the desired information,
         but does contain a pointer to a closer name server, the
         resolver returns to step 2, where the closer name servers will
         be queried.

      Several modifications to this algorithm are possible. A resolver
      may not support a local cache and instead only cache information
      during the course of a single user request, discarding it upon

completion. The resolver may also find that a datagram reply was
truncated, and open a virtual circuit so that the complete reply
can be recovered.

Inverse and completion queries must be treated in an
environment-sensitive manner, because the domain system doesn't
provide a method for guaranteeing that it can locate the correct
information. The typical choice will be to configure a resolver
to use a particular set of known name servers for inverse queries.

DOMAIN SUPPORT FOR MAIL

    Introduction

        Mail service is a particularly sensitive issue for users of the
        domain system because of the lack of a consistent system for
        naming mailboxes and even hosts, and the need to support continued
        operation of existing services. This section discusses an
        evolutionary approach for adding consistent domain name support
        for mail.

        The crucial issue is deciding on the types of binding to be
        supported. Most mail systems specify a mail destination with a
        two part construct such as X@Y. The left hand side, X, is an
        string, often a user or account, and Y is a string, often a host.
        This section refers to the part on the left, i.e. X, as the local
        part, and refers to the part on the right, i.e. Y, as the global
        part.

        Most existing mail systems route mail based on the global part; a
        mailer with mail to deliver to X@Y will decide on the host to be
        contacted using only Y. We refer to this type of binding as
        "agent binding".

            For example, mail addressed to Mockapetris@ISIF is delivered to
            host USC-ISIF (USC-ISIF is the official name for the host
            specified by nickname ISIF).

        More sophisticated mail systems use both the local and global
        parts, i.e. both X and Y to determine which host should receive
        the mail. These more sophisticated systems usually separate the
        binding of the destination to the host from the actual delivery.
        This allows the global part to be a generic name rather than
        constraining it to a single host. We refer to this type of
        binding as "mailbox binding".

            For example, mail addressed to Mockapetris@ISI might be bound
            to host F.ISI.ARPA, and subsequently delivered to that host,
            while mail for Cohen@ISI might be bound to host B.ISI.ARPA.

        The domain support for mail consists of two levels of support,
        corresponding to these two binding models.

            The first level, agent binding, is compatible with existing
            ARPA Internet mail procedures and uses maps a global part onto
            one or more hosts that will accept the mail. This type of
            binding uses the MAILA QTYPE.

            The second level, mailbox binding, offers extended services

that map a local part and a global part onto one or more sets
of data via the MAILB QTYPE. The sets of data include hosts
that will accept the mail, mailing list members (mail groups),
and mailboxes for reporting errors or requests to change a mail
group.

The domain system encodes the global part of a mail destination as
a domain name and uses dots in the global part to separate labels
in the encoded domain name. The domain system encodes the local
part of a mail destination as a single label, and any dots in this
part are simply copied into the label. The domain system forms a
complete mail destination as the local label concatenated to the
domain string for the global part. We call this a mailbox.

For example, the mailbox Mockapetris@F.ISI.ARPA has a global
domain name of three labels, F.ISI.ARPA. The domain name
encoding for the whole mailbox is Mockapetris.F.ISI.ARPA. The
mailbox Mockapetris.cad@F.ISI.ARPA has the same domain name for
the global part and a 4 label domain name for the mailbox of
Mockapetris\.cad.F.ISI.ARPA (the \ is not stored in the label,
its merely used to denote the "quoted" dot).

It is anticipated that the Internet system will adopt agent
binding as part of the initial implementation of the domain
system, and that mailbox binding will eventually become the
preferred style as organizations convert their mail systems to the
new style. To facilitate this approach, the domain information
for these two binding styles is organized to allow a requestor to
determine which types of support are available, and the
information is kept in two disjoint classes.

Agent binding

In agent binding, a mail system uses the global part of the mail
destination as a domain name, with dots denoting structure. The
domain name is resolved using a MAILA query which return MF and MD
RRs to specify the domain name of the appropriate host to receive
the mail. MD (Mail delivery) RRs specify hosts that are expected
to have the mailbox in question; MF (Mail forwarding) RRs specify
hosts that are expected to be intermediaries willing to accept the
mail for eventual forwarding. The hosts are hints, rather than
definite answers, since the query is made without the full mail
destination specification.

For example, mail for MOCKAPETRIS@F.ISI.ARPA would result in a
query with QTYPE=MAILA and QNAME=F.ISI.ARPA, which might return
two RRs:

                        F.ISI.ARPA MD IN F.ISI.ARPA
                        F.ISI.ARPA MF IN A.ISI.ARPA

        The mailer would interpret these to mean that the mail agent on
        F.ISI.ARPA should be able to deliver the mail directly, but that
        A.ISI.ARPA is willing to accept the mail for probable forwarding.

        Using this system, an organization could implement a system that
        uses organization names for global parts, rather than the usual
        host names, but all mail for the organization would be routed the
        same, regardless of its local part. Hence and organization with
        many hosts would expect to see many forwarding operations.

    Mailbox binding

        In mailbox binding, the mailer uses the entire mail destination
        specification to construct a domain name. The encoded domain name
        for the mailbox is used as the QNAME field in a QTYPE=MAILB query.

        Several outcomes are possible for this query:

        1. The query can return a name error indicating that the mailbox
           does not exist as a domain name.

           In the long term this would indicate that the specified mailbox
           doesn't exist. However, until the use of mailbox binding is
           universal, this error condition should be interpreted to mean
           that the organization identified by the global part does not
           support mailbox binding. The appropriate procedure is to
           revert to agent binding at this point.

        2. The query can return a Mail Rename (MR) RR.

           The MR RR carries new mailbox specification in its RDATA field.
           The mailer should replace the old mailbox with the new one and
           retry the operation.

        3. The query can return a MB RR.

           The MB RR carries a domain name for a host in its RDATA field.
           The mailer should deliver the message to that host via whatever
           protocol is applicable, e.g. SMTP.

        4. The query can return one or more Mail Group (MG) RRs.

           This condition means that the mailbox was actually a mailing
           list or mail group, rather than a single mailbox. Each MG RR
           has a RDATA field that identifies a mailbox that is a member of

        the group. The mailer should deliver a copy of the message to
        each member.

    5. The query can return a MB RR as well as one or more MG RRs.

        This condition means the the mailbox was actually a mailing
        list. The mailer can either deliver the message to the host
        specified by the MB RR, which will in turn do the delivery to
        all members, or the mailer can use the MG RRs to do the
        expansion itself.

    In any of these cases, the response may include a Mail Information
    (MINFO) RR. This RR is usually associated with a mail group, but
    is legal with a MB. The MINFO RR identifies two mailboxes. One
    of these identifies a responsible person for the original mailbox
    name. This mailbox should be used for requests to be added to a
    mail group, etc. The second mailbox name in the MINFO RR
    identifies a mailbox that should receive error messages for mail
    failures. This is particularly appropriate for mailing lists when
    errors in member names should be reported to a person other than
    the one who sends a message to the list. New fields may be added
    to this RR in the future.

Appendix 1 - Domain Name Syntax Specification

   The preferred syntax of domain names is given by the following BNF
   rules. Adherence to this syntax will result in fewer problems with
   many applications that use domain names (e.g., mail, TELNET). Note
   that some applications use domain names containing binary information
   and hence do not follow this syntax.

      <domain> ::= <subdomain> | " "

      <subdomain> ::= <label> | <subdomain> "." <label>

      <label> ::= <letter> [[<ldh-str>] <let-dig>]

      <ldh-str> ::= <let-dig-hyp> | <let-dig-hyp> <ldh-str>

      <let-dig-hyp> ::= <let-dig> | "-"

      <let-dig> ::= <letter> | <digit>

      <letter> ::= any one of the 52 alphabetic characters A through Z
      in upper case and a through z in lower case

      <digit> ::= any one of the ten digits 0 through 9

   Note that while upper and lower case letters are allowed in domain
   names no significance is attached to the case. That is, two names
   with the same spelling but different case are to be treated as if
   identical.

   The labels must follow the rules for ARPANET host names. They must
   start with a letter, end with a letter or digit, and have as interior
   characters only letters, digits, and hyphen. There are also some
   restrictions on the length. Labels must be 63 characters or less.

   For example, the following strings identify hosts in the ARPA
   Internet:

      F.ISI.ARPA   LINKABIT-DCN5.ARPA   UCL-TAC.ARPA

Appendix 2 - Field formats and encodings

```
+----------------------------------------------------+
|                                                    |
|                 *****  WARNING  *****               |
|                                                    |
|    The following formats are preliminary and       |
|  are included for purposes of explanation only.    |
|  In particular, new RR types will be added,        |
|  and the size, position, and encoding of           |
|  fields are subject to change.                     |
|                                                    |
+----------------------------------------------------+
```

   TYPE values

      TYPE fields are used in resource records. Note that these types
      are not the same as the QTYPE fields used in queries, although the
      functions are often similar.

      TYPE value meaning

      A   1   a host address

      NS   2   an authoritative name server

      MD   3   a mail destination

      MF   4   a mail forwarder

      CNAME 5   the canonical name for an alias

      SOA   6   marks the start of a zone of authority

      MB   7   a mailbox domain name

      MG   8   a mail group member

      MR   9   a mail rename domain name

      NULL 10   a null RR

      WKS   11   a well known service description

      PTR   12   a domain name pointer

      HINFO 13   host information

      MINFO 14   mailbox or mail list information

QTYPE values

   QTYPE fields appear in the question part of a query. They include
   the values of TYPE with the following additions:

   AXFR   252 A request for a transfer of an entire zone of authority

   MAILB 253 A request for mailbox-related records (MB, MG or MR)

   MAILA 254 A request for mail agent RRs (MD and MF)

   *   255 A request for all records

CLASS values

   CLASS fields appear in resource records

   CLASS value meaning

   IN   1   the ARPA Internet

   CS   2   the computer science network (CSNET)

QCLASS values

   QCLASS fields appear in the question section of a query. They
   include the values of CLASS with the following additions:

   *   255 any class

Standard resource record formats

All RRs have the same top level format shown below:
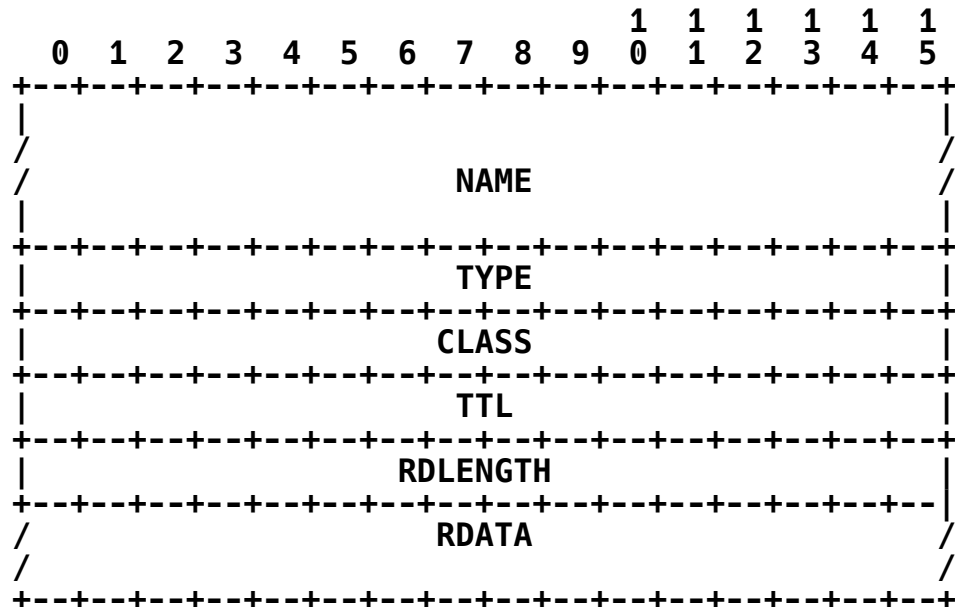
```
                                      1 1 1 1 1 1
      0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                                               |
    /                                               /
    /                     NAME                      /
    |                                               |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                     TYPE                      |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                     CLASS                     |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                     TTL                       |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    |                   RDLENGTH                    |
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--|
    /                     RDATA                     /
    /                                               /
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

where:

NAME   - a compressed domain name to which this resource
           record pertains.

TYPE   - two octets containing one of the RR type codes
           defined in Appendix 2. This field specifies the
           meaning of the data in the RDATA field.

CLASS   - two octets which specifies the class of the data in
           the RDATA field.

TTL   - a 16 bit signed integer that specifies the time
           interval that the resource record may be cached
           before the source of the information should again be
           consulted. Zero values are interpreted to mean that
           the RR can only be used for the transaction in
           progress, and should not be cached. For example, SOA
           records are always distributed with a zero TTL to
           prohibit caching. Zero values can also be used for
           extremely volatile data.

RDLENGTH- an unsigned 16 bit integer that specifies the length
           in octets of the RDATA field.

RDATA - a variable length string of octets that describes the
         resource. The format of this information varies
         according to the TYPE and CLASS of the resource
         record.

The format of the RDATA field is standard for all classes for the
RR types NS, MD, MF, CNAME, SOA, MB, MG, MR, PTR, HINFO, MINFO and
NULL. These formats are shown below together with the appropriate
additional section RR processing.

CNAME RDATA format

```
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    /                     CNAME                     /
    /                                               /
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

where:

CNAME   - A compressed domain name which specifies that the
          domain name of the RR is an alias for a canonical
          name specified by CNAME.

CNAME records cause no additional section processing. The
RDATA section of a CNAME line in a master file is a standard
printed domain name.

HINFO RDATA format

```
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    /                      CPU                      /
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    /                      OS                       /
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

where:

CPU   - A character string which specifies the CPU type. The
         character string is represented as a single octet
         length followed by that number of characters.   The
         following standard strings are defined:.

```
PDP-11/70    C/30          C/70          VAX-11/780
H-316        H-516         DEC-2060      DEC-1090T
ALTO         IBM-PC        IBM-PC/XT     PERQ
IBM-360/67   IBM-370/145
```

OS   - A character string which specifies the operating system
type. The character string is represented as a single octet

length followed by that number of characters.   The following
standard types are defined:.

```
    ASP          AUGUST       BKY          CCP
    DOS/360      ELF          EPOS         EXEC-8
    GCOS         GPOS         ITS          INTERCOM
    KRONOS       MCP          MOS          MPX-RT
    MULTICS      MVT          NOS          NOS/BE
    OS/MVS       OS/MVT       RIG          RSX11
    RSX11M       RT11         SCOPE        SIGNAL
    SINTRAN      TENEX        TOPS10       TOPS20
    TSS          UNIX         VM/370       VM/CMS
    VMS          WAITS
```

HINFO records cause no additional section processing.

HINFO records are used to acquire general information about a
host. The main use is for protocols such as FTP that can use
special procedures when talking between machines or operating
systems of the same type.

MB RDATA format

```
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    /                     MADNAME                   /
    /                                               /
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

where:

MADNAME - A compressed domain name which specifies a host which
          has the specified mailbox.

MB records cause additional section processing which looks up
an A type record corresponding to MADNAME. The RDATA section
of a MB line in a master file is a standard printed domain
name.

MD RDATA format

```
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
    /                     MADNAME                   /
    /                                               /
    +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

where:

MADNAME - A compressed domain name which specifies a host which

                         has a mail agent for the domain which should be able
                         to deliver mail for the domain.

     MD records cause additional section processing which looks up
     an A type record corresponding to MADNAME. The RDATA section
     of a MD line in a master file is a standard printed domain
     name.

  MF RDATA format

         +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
         /   MADNAME   /
         /   /
         +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

     where:

     MADNAME - A compressed domain name which specifies a host which
               has a mail agent for the domain which will accept
               mail for forwarding to the domain.

     MF records cause additional section processing which looks up
     an A type record corresponding to MADNAME. The RDATA section
     of a MF line in a master file is a standard printed domain
     name.

  MG RDATA format

         +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
         /   MGMNAME   /
         /   /
         +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

     where:

     MGMNAME - A compressed domain name which specifies a mailbox
               which is a member of the mail group specified by the
               domain name.

     MF records cause no additional section processing. The RDATA
     section of a MF line in a master file is a standard printed
     domain name.

     MINFO RDATA format

         +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
         /   RMAILBX   /
         +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
         /   EMAILBX   /
         +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

     where:

     RMAILBX - A compressed domain name which specifies a mailbox
               which is responsible for the mailing list or mailbox.
               If this domain name names the root, the owner of the
               MINFO RR is responsible for itself. Note that many
               existing mailing lists use a mailbox X-request for
               the RMAILBX field of mailing list X, e.g.
               Msgroup-request for Msgroup. This field provides a
               more general mechanism.

     EMAILBX - A compressed domain name which specifies a mailbox
               which is to receive error messages related to the
               mailing list or mailbox specified by the owner of the
               MINFO RR (similar to the ERRORS-TO: field which has
               been proposed). If this domain name names the root,
               errors should be returned to the sender of the
               message.

     MINFO records cause no additional section processing. Although
     these records can be associated with a simple mailbox, they are
     usually used with a mailing list. The MINFO section of a MF
     line in a master file is a standard printed domain name.

     MR RDATA format

         +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
         /   NEWNAME   /
         /   /
         +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

     where:

     NEWNAME - A compressed domain name which specifies a mailbox
               which is the proper rename of the specified mailbox.

     MR records cause no additional section processing. The RDATA
     section of a MR line in a master file is a standard printed
     domain name.

NULL RDATA format

```
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
/                    <anything>                 /
/                                               /
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

Anything at all may be in the RDATA field so long as it is
65535 octets or less.

NULL records cause no additional section processing. NULL RRs
are not allowed in master files.

NS RDATA format

```
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
/                     NSDNAME                   /
/                                               /
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

where:

NSDNAME - A compressed domain name which specifies a host which
          has a name server for the domain.

NS records cause both the usual additional section processing
to locate a type A record, and a special search of the zone in
which they reside. The RDATA section of a NS line in a master
file is a standard printed domain name.

PTR RDATA format

```
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
/                     PTRDNAME                  /
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

where:

PTRDNAME - A compressed domain name which points to some
           location in the domain name space.

PTR records cause no additional section processing. These RRs
are used in special domains to point to some other location in
the domain space. These records are simple data, and don't
imply any special processing similar to that performed by
CNAME, which identifies aliases. Appendix 3 discusses the use
of these records in the ARPA Internet address domain.

        SOA RDATA format

```
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        /                      MNAME                    /
        /                                               /
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        /                      RNAME                    /
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                      SERIAL                   |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                      REFRESH                  |
        |                                               |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                      RETRY                    |
        |                                               |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                      EXPIRE                   |
        |                                               |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
        |                      MINIMUM                  |
        +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

    where:

    MNAME   - The domain name of the name server that was the
                original source of data for this zone.

    RNAME   - A domain name which specifies the mailbox of the
                person responsible for this zone.

    SERIAL   - The unsigned 16 bit version number of the of the
                original copy of the zone. This value wraps and
                should be compared using sequence space arithmetic.

    REFRESH   - The unsigned 32 bit time interval before the zone
                should be refreshed.

    RETRY   - The unsigned 32 bit time interval that should elapse
                before a failed refresh should be retried.

    EXPIRE   - A 32 bit time value that specifies the upper limit on
                the time interval that can elapse before the zone is
                no longer authoritative.

    MINIMUM   - The unsigned 16 bit minimum TTL field that should be
                exported with any RR from this zone (other than the
                SOA itself).

    SOA records cause no additional section processing. The RDATA

        section of a SOA line in a master file is a standard printed
        domain name for MNAME, a standard X@Y mailbox specification for
        RNAME, and decimal numbers for the remaining parameters.

        All times are in units of seconds.

        Most of these fields are pertinent only for name server
        maintenance operations. However, MINIMUM is used in all query
        operations that retrieve RRs from a zone. Whenever a RR is
        sent in a response to a query, the TTL field is set to the
        maximum of the TTL field from the RR and the MINIMUM field in
        the appropriate SOA. Thus MINIMUM is a lower bound on the TTL
        field for all RRs in a zone. RRs in a zone are never discarded
        due to timeout unless the whole zone is deleted. This prevents
        partial copies of zones.

Appendix 3 - Internet specific field formats and operations

    Message transport

        The Internet supports name server access using TCP [10] on server
        port 53 (decimal) as well as datagram access using UDP [11] on UDP
        port 53 (decimal). Messages sent over TCP virtual circuits are
        preceded by an unsigned 16 bit length field which describes the
        length of the message, excluding the length field itself.

            +---+
            | |
            | ***** WARNING ***** |
            | |
            | The following formats are preliminary and |
            | are included for purposes of explanation only. |
            | In particular, new RR types will be added, |
            | and the size, position, and encoding of |
            | fields are subject to change. |
            | |
            +---+

    A RDATA format

            +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
            | ADDRESS |
            +--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

        where:

        ADDRESS   - A 32 bit ARPA internet address

        Hosts that have multiple ARPA Internet addresses will have
        multiple A records.

        A records cause no additional section processing. The RDATA
        section of an A line in a master file is an Internet address
        expressed as four decimal numbers separated by dots without any
        imbedded spaces (e.g., "10.2.0.52" or "192.0.5.6").

WKS RDATA format

```
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|                   ADDRESS                     |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|    PROTOCOL       |                           |
+--+--+--+--+--+--+--+--+                        |
|                                               |
/                  <BIT MAP>                    /
/                                               /
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

where:

ADDRESS   - An 32 bit ARPA Internet address

PROTOCOL   - An 8 bit IP protocol number

<BIT MAP> - A variable length bit map. The bit map must be a
           multiple of 8 bits long.

The WKS record is used to describe the well known services
supported by a particular protocol on a particular internet
address. The PROTOCOL field specifies an IP protocol number, and
the bit map has one bit per port of the specified protocol. The
first bit corresponds to port 0, the second to port 1, etc. If
less than 256 bits are present, the remainder are assumed to be
zero. The appropriate values for ports and protocols are
specified in [13].

For example, if PROTOCOL=TCP (6), the 26th bit corresponds to TCP
port 25 (SMTP). If this bit is set, a SMTP server should be
listening on TCP port 25; if zero, SMTP service is not supported
on the specified address.

The anticipated use of WKS RRs is to provide availability
information for servers for TCP and UDP. If a server supports
both TCP and UDP, or has multiple Internet addresses, then
multiple WKS RRs are used.

WKS RRs cause no additional section processing. The RDATA section
of a WKS record consists of a decimal protocol number followed by
mnemonic identifiers which specify bits to be set to 1.

IN-ADDR special domain

The ARPA internet uses a special domain to support gateway
location and ARPA Internet address to host mapping. The intent of
this domain is to allow queries to locate all gateways on a

        particular network in the ARPA Internet, and also to provide a
        guaranteed method to perform host address to host name mapping.

        Note that both of these services are similar to functions that
        could be performed by inverse queries; the difference is that this
        part of the domain name space is structured according to address,
        and hence can guarantee that the appropriate data can be located
        without an exhaustive search of the domain space. It is
        anticipated that the special tree will be used by ARPA Internet
        resolvers for all gateway location services, but that address to
        name resolution will be performed by first trying the inverse
        query on the local name server database followed by a query in the
        special space if the inverse query fails.

        The domain is a top level domain called IN-ADDR whose substructure
        follows the ARPA Internet addressing structure.

        Domain names in the IN-ADDR domain are defined to have up to four
        labels in addition to the IN-ADDR label. Each label is a
        character string which expresses a decimal value in the range
        0-255 (with leading zeros omitted except in the case of a zero
        octet which is represented by a single zero). These labels
        correspond to the 4 octets of an ARPA Internet address.

        Host addresses are represented by domain names that have all four
        labels specified. Thus data for ARPA Internet address 10.2.0.52
        is located at domain name 52.0.2.10.IN-ADDR. The reversal, though
        awkward to read, allows zones to follow the natural grouping of
        hosts within networks. For example, 10.IN-ADDR can be a zone
        containing data for the ARPANET, while 26.IN-ADDR can be a
        separate zone for MILNET. Address nodes are used to hold pointers
        to primary host names in the normal domain space.

        Network addresses correspond to some of the non-terminal nodes in
        the IN-ADDR tree, since ARPA Internet network numbers are either
        1, 2, or 3 octets. Network nodes are used to hold pointers to
        primary host names (which happen to be gateways) in the normal
        domain space. Since a gateway is, by definition, on more than one
        network, it will typically have two or more network nodes that
        point at the gateway. Gateways will also have host level pointers
        at their fully qualified addresses.

        Both the gateway pointers at network nodes and the normal host
        pointers at full address nodes use the PTR RR to point back to the
        primary domain names of the corresponding hosts.

        For example, part of the IN-ADDR domain will contain information
        about the ISI to MILNET and MIT gateways, and hosts F.ISI.ARPA and
        MULTICS.MIT.ARPA. Assuming that ISI gateway has addresses

10.2.0.22 and 26.0.0.103, and a name MILNET-GW.ISI.ARPA, and the
MIT gateway has addresses 10.0.0.77 and 18.10.0.4 and a name
GW.MIT.ARPA, the domain database would contain:

```
10.IN-ADDR          PTR  IN MILNET-GW.ISI.ARPA
10.IN-ADDR          PTR  IN GW.MIT.ARPA
18.IN-ADDR          PTR  IN GW.MIT.ARPA
26.IN-ADDR          PTR  IN MILNET-GW.ISI.ARPA
22.0.2.10.IN-ADDR   PTR  IN MILNET-GW.ISI.ARPA
103.0.0.26.IN-ADDR  PTR  IN MILNET-GW.ISI.ARPA
77.0.0.10.IN-ADDR   PTR  IN GW.MIT.ARPA
4.0.10.18.IN-ADDR   PTR  IN GW.MIT.ARPA
52.0.2.10.IN-ADDR   PTR  IN F.ISI.ARPA
6.0.0.10.IN-ADDR    PTR  IN MULTICS.MIT.ARPA
```

Thus a program which wanted to locate gateways on net 10 would
originate a query of the form QTYPE=PTR, QCLASS=IN,
QNAME=10.IN-ADDR. It would receive two RRs in response:

```
10.IN-ADDR          PTR  IN MILNET-GW.ISI.ARPA
10.IN-ADDR          PTR  IN GW.MIT.ARPA
```

The program could then originate QTYPE=A, QCLASS=IN queries for
MILNET-GW.ISI.ARPA and GW.MIT.ARPA to discover the ARPA Internet
addresses of these gateways.

A resolver which wanted to find the host name corresponding to
ARPA Internet host address 10.0.0.6 might first try an inverse
query on the local name server, but find that this information
wasn't available. It could then try a query of the form
QTYPE=PTR, QCLASS=IN, QNAME=6.0.0.10.IN-ADDR, and would receive:

```
6.0.0.10.IN-ADDR     PTR  IN MULTICS.MIT.ARPA
```

Several cautions apply to the use of these services:

   Since the IN-ADDR special domain and the normal domain for a
   particular host or gateway will be in different zones, the
   possibility exists that that the data may be inconsistent.

   Gateways will often have two names in separate domains, only
   one of which can be primary.

   Systems that use the domain database to initialize their
   routing tables must start with enough gateway information to
   guarantee that they can access the appropriate name server.

   The gateway data only reflects the existence of a gateway in a

         manner equivalent to the current HOSTS.TXT file. It doesn't
         replace the dynamic availability information from GGP or EGP.

REFERENCES and BIBLIOGRAPHY

    [1] E. Feinler, K. Harrenstien, Z. Su, and V. White, "DOD Internet
         Host Table Specification", RFC 810, Network Information Center,
         SRI International, March 1982.

    [2] J. Postel, "Computer Mail Meeting Notes", RFC 805,
         USC/Information Sciences Institute, February 1982.

    [3] Z. Su, and J. Postel, "The Domain Naming Convention for Internet
         User Applications", RFC 819, Network Information Center, SRI
         International, August 1982.

    [4] Z. Su, "A Distributed System for Internet Name Service",
         RFC 830, Network Information Center, SRI International,
         October 1982.

    [5] K. Harrenstien, and V. White, "NICNAME/WHOIS", RFC 812, Network
         Information Center, SRI International, March 1982.

    [6]   M. Solomon, L. Landweber, and D. Neuhengen, "The CSNET Name
         Server", Computer Networks, vol 6, nr 3, July 1982.

    [7] K. Harrenstien, "NAME/FINGER", RFC 742, Network Information
         Center, SRI International, December 1977.

    [8] J. Postel, "Internet Name Server", IEN 116, USC/Information
         Sciences Institute, August 1979.

    [9] K. Harrenstien, V. White, and E. Feinler, "Hostnames Server",
         RFC 811, Network Information Center, SRI International,
         March 1982.

    [10] J. Postel, "Transmission Control Protocol", RFC 793,
         USC/Information Sciences Institute, September 1981.

    [11] J. Postel, "User Datagram Protocol", RFC 768, USC/Information
         Sciences Institute, August 1980.

    [12] J. Postel, "Simple Mail Transfer Protocol", RFC 821,
         USC/Information Sciences Institute, August 1980.

    [13] J. Reynolds, and J. Postel, "Assigned Numbers", RFC 870,
         USC/Information Sciences Institute, October 1983.

    [14] P. Mockapetris, "Domain names - Concepts and Facilities,"
         RFC 882, USC/Information Sciences Institute, November 1983.

INDEX