

Network Working Group
Request for Comments: 3259
Category: Informational

J. Ott
TZI, Universitaet Bremen
C. Perkins
USC Information Sciences Institute
D. Kutscher
TZI, Universitaet Bremen
April 2002

A Message Bus for Local Coordination

Status of this Memo

This memo provides information for the Internet community. It does not specify an Internet standard of any kind. Distribution of this memo is unlimited.

Copyright Notice

Copyright (C) The Internet Society (2002). All Rights Reserved.

Abstract

The local Message Bus (Mbus) is a light-weight message-oriented coordination protocol for group communication between application components. The Mbus provides automatic location of communication peers, subject based addressing, reliable message transfer and different types of communication schemes. The protocol is layered on top of IP multicast and is specified for IPv4 and IPv6. The IP multicast scope is limited to link-local multicast. This document specifies the Mbus protocol, i.e., message syntax, addressing and transport mechanisms.

Table of Contents

1.	Introduction	3
1.1	Mbus Overview	3
1.2	Purpose of this Document	5
1.3	Areas of Application	5
1.4	Terminology for requirement specifications	6
2.	Common Formal Syntax Rules	6
3.	Message Format	7
4.	Addressing	9
4.1	Mandatory Address Elements	10
5.	Message Syntax	11
5.1	Message Encoding	11
5.2	Message Header	11
5.3	Command Syntax	12

6.	Transport	13
6.1	Local Multicast/Broadcast	14
6.1.1	Mbus multicast groups for IPv4	15
6.1.2	Mbus multicast groups for IPv6	15
6.1.3	Use of Broadcast	16
6.1.4	Mbus UDP Port Number	16
6.2	Directed Unicast	16
7.	Reliability	18
8.	Awareness of other Entities	20
8.1	Hello Message Transmission Interval	21
8.1.1	Calculating the Interval for Hello Messages	22
8.1.2	Initialization of Values	23
8.1.3	Adjusting the Hello Message Interval when the Number of Entities increases	23
8.1.4	Adjusting the Hello Message Interval when the Number of Entities decreases	23
8.1.5	Expiration of hello timers	23
8.2	Calculating the Timeout for Mbus Entities	24
9.	Messages	24
9.1	mbus.hello	24
9.2	mbus.bye	25
9.3	mbus.ping	25
9.4	mbus.quit	26
9.5	mbus.waiting	26
9.6	mbus.go	27
10.	Constants	27
11.	Mbus Security	28
11.1	Security Model	28
11.2	Encryption	28
11.3	Message Authentication	29
11.4	Procedures for Senders and Receivers	30
12.	Mbus Configuration	31
12.1	File based parameter storage	33
12.2	Registry based parameter storage	34
13.	Security Considerations	34
14.	IANA Considerations	35
15.	References	35
A.	About References	37
B.	Limitations and Future Work	37
	Authors' Addresses	38
	Full Copyright Statement	39

1. Introduction

The implementation of multiparty multimedia conferencing systems is one example where a simple coordination infrastructure can be useful: In a variety of conferencing scenarios, a local communication channel can provide conference-related information exchange between co-located but otherwise independent application entities, for example those taking part in application sessions that belong to the same conference. In loosely coupled conferences such a mechanism allows for coordination of application entities, e.g., to implement synchronization between media streams or to configure entities without user interaction. It can also be used to implement tightly coupled conferences enabling a conference controller to enforce conference wide control within an end system.

Conferencing systems such as IP telephones can also be viewed as components of a distributed system and can thus be integrated into a group of application modules: For example, an IP telephony call that is conducted with a stand-alone IP telephone can be dynamically extended to include media engines for other media types using the coordination function of an appropriate coordination mechanism. Different individual conferencing components can thus be combined to build a coherent multimedia conferencing system for a user.

Other possible scenarios include the coordination of application components that are distributed on different hosts in a network, for example, so-called Internet appliances.

1.1 Mbus Overview

Local coordination of application components requires a number of different interaction models: some messages (such as membership information, floor control notifications, dissemination of conference state changes, etc.) may need to be sent to all local application entities. Messages may also be targeted at a certain application class (e.g., all whiteboards or all audio tools) or agent type (e.g., all user interfaces rather than all media engines). Or there may be any (application- or message-specific) subgrouping defining the intended recipients, e.g., messages related to media synchronization. Finally, there may be messages that are directed at a single entity: for example, specific configuration settings that a conference controller sends to a particular application entity, or query-response exchanges between any local server and its clients.

The Mbus protocol as defined here satisfies these different communication needs by defining different message transport mechanisms (defined in Section 6) and by providing a flexible addressing scheme (defined in Section 4).

Furthermore, Mbus messages exchanged between application entities may have different reliability requirements (which are typically derived from their semantics). Some messages will have a rather transient character conveying ephemeral state information (which is refreshed/updated periodically), such as the volume meter level of an audio receiver entity to be displayed by its user interface agent. Certain Mbus messages (such as queries for parameters or queries to local servers) may require a response from the peer(s), thereby providing an explicit acknowledgment at the semantic level on top of the Mbus. Other messages will modify the application or conference state and hence it is crucial that they do not get lost. The latter type of message has to be delivered reliably to the recipient, whereas messages of the first type do not require reliability mechanisms at the Mbus transport layer. For messages confirmed at the application layer it is up to the discretion of the application whether or not to use a reliable transport underneath.

In some cases, application entities will want to tailor the degree of reliability to their needs, others will want to rely on the underlying transport to ensure delivery of the messages -- and this may be different for each Mbus message. The Mbus message passing mechanism specified in this document provides a maximum of flexibility by providing reliable transmission achieved through transport-layer acknowledgments (in case of point-to-point communications only) as well as unreliable message passing (for unicast, local multicast, and local broadcast). We address this topic in Section 4.

Finally, accidental or malicious disturbance of Mbus communications through messages originated by applications from other users needs to be prevented. Accidental reception of Mbus messages from other users may occur if either two users share the same host for using Mbus applications or if they are using Mbus applications that are spread across the same network link: in either case, the used Mbus multicast address and the port number may be identical leading to reception of the other party's Mbus messages in addition to the user's own ones. Malicious disturbance may happen because of applications multicasting (e.g., at a global scope) or unicasting Mbus messages. To eliminate the possibility of processing unwanted Mbus messages, the Mbus protocol contains message digests for authentication. Furthermore, the Mbus allows for encryption to ensure privacy and thus enable using the Mbus for local key distribution and other functions potentially sensitive to eavesdropping. This document defines the framework for configuring Mbus applications with regard to security parameters in Section 12.

1.2 Purpose of this Document

Three components constitute the message bus: the low level message passing mechanisms, a command syntax and naming hierarchy, and the addressing scheme.

The purpose of this document is to define the protocol mechanisms of the lower level Mbus message passing mechanism which is common to all Mbus implementations. This includes the specification of

- o the generic Mbus message format;
- o the addressing concept for application entities (note that concrete addressing schemes are to be defined by application-specific profiles);
- o the transport mechanisms to be employed for conveying messages between (co-located) application entities;
- o the security concept to prevent misuse of the Message Bus (such as taking control of another user's conferencing environment);
- o the details of the Mbus message syntax; and
- o a set of mandatory application independent commands that are used for bootstrapping Mbus sessions.

1.3 Areas of Application

The Mbus protocol can be deployed in many different application areas, including but not limited to:

Local conference control: In the Mbone community a model has arisen whereby a set of loosely coupled tools are used to participate in a conference. A typical scenario is that audio, video, and shared workspace functionality is provided by three separate tools (although some combined tools exist). This maps well onto the underlying RTP [8] (as well as other) media streams, which are also transmitted separately. Given such an architecture, it is useful to be able to perform some coordination of the separate media tools. For example, it may be desirable to communicate playout-point information between audio and video tools, in order to implement lip-synchronization, to arbitrate the use of shared resources (such as input devices), etc.

A refinement of this architecture relies on the presence of a number of media engines which perform protocol functions as well as capturing and playout of media. In addition, one (or more)

(separate) user interface agents exist that interact with and control their media engine(s). Such an approach allows flexibility in the user-interface design and implementation, but obviously requires some means by which the various involved agents may communicate with one another. This is particularly desirable to enable a coherent response to a user's conference-related actions (such as joining or leaving a conference).

Although current practice in the Mbone community is to work with a loosely coupled conference control model, situations arise where this is not appropriate and a more tightly coupled wide-area conference control protocol must be employed. In such cases, it is highly desirable to be able to re-use the existing tools (media engines) available for loosely coupled conferences and integrate them with a system component implementing the tight conference control model. One appropriate means to achieve this integration is a communication channel that allows a dedicated conference control entity to "remotely" control the media engines in addition to or instead of their respective user interfaces.

Control of device groups in a network: A group of devices that are connected to a local network, e.g., home appliances in a home network, require a local coordination mechanism. Minimizing manual configuration and the the possibility to deploy group communication will be useful in this application area as well.

1.4 Terminology for requirement specifications

In this document, the key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" are to be interpreted as described in RFC 2119 [1] and indicate requirement levels for compliant Mbus implementations.

2. Common Formal Syntax Rules

This section contains definitions of common ABNF [13] syntax elements that are later referenced by other definitions in this document:

```
base64           = base64_terminal /
                  ( 1*(4base64_CHAR) [base64_terminal] )

base64_char      = UPALPHA / LOALPHA / DIGIT / "+" / "/"
                  ;; Case-sensitive

base64_terminal = (2base64_char "==") / (3base64_char "=")

UPALPHA          = %x41-5A                ;; Uppercase: A-Z
```

LOALPHA	= %x61-7A	;; Lowercase: a-z
ALPHA	= %x41-5A / %x61-7A	; A-Z / a-z
CHAR	= %x01-7E	; any 7-bit US-ASCII character, excluding NUL and delete
OCTET	= %x00-FF	; 8 bits of data
CR	= %x0D	; carriage return
CRLF	= CR LF	; Internet standard newline
DIGIT	= %x30-39	; 0-9
DQUOTE	= %x22	; " (Double Quote)
HTAB	= %x09	; horizontal tab
LF	= %x0A	; linefeed
LWSP	= *(WSP / CRLF WSP)	; linear white space (past newline)
SP	= %x20	; space
WSP	= SP / HTAB	; white space

Taken from RFC 2234 [13] and RFC 2554 [14].

3. Message Format

An Mbus message comprises a header and a body. The header is used to indicate how and where a message should be delivered and the body provides information and commands to the destination entity. The following pieces of information are included in the header:

A fixed ProtocolID field identifies the version of the message bus protocol used. The protocol defined in this document is "mbus/1.0" (case-sensitive).

A sequence number (SeqNum) is contained in each message. The first message sent by a source SHOULD set SeqNum to zero, and it MUST increment by one for each message sent by that source. A single sequence number is used for all messages from a source, irrespective of the intended recipients and the reliability mode selected. The value range of a sequence number is (0,4294967295). An implementation MUST re-set its sequence number to 0 after reaching 4294967295. Implementations MUST take into account that the SeqNum of other entities may wrap-around.

SeqNums are decimal numbers in ASCII representation.

The TimeStamp field is also contained in each message and SHOULD contain a decimal number representing the time of the message construction in milliseconds since 00:00:00, UTC, January 1, 1970.

A MessageType field indicates the kind of message being sent. The value "R" indicates that the message is to be transmitted reliably and MUST be acknowledged by the recipient, "U" indicates an unreliable message which MUST NOT be acknowledged.

The SrcAddr field identifies the sender of a message. This MUST be a complete address, with all address elements specified. The addressing scheme is described in Section 4.

The DestAddr field identifies the intended recipient(s) of the message. This field MAY be wildcarded by omitting address elements and hence address any number (including zero) of application entities. The addressing scheme is described in Section 4.

The AckList field comprises a list of SeqNums for which this message is an acknowledgment. See Section 7 for details.

The header is followed by the message body which contains zero or more commands to be delivered to the destination entity. The syntax for a complete message is given in Section 5.

If multiple commands are contained within the same Mbus message payload, they MUST to be delivered to the Mbus application in the same sequence in which they appear in the message payload.

4. Addressing

Each entity in the message has a unique Mbus address that is used to identify the entity. Mbus addresses are sequences of address elements that are tag/value pairs. The tag and the value are separated by a colon and tag/value pairs are separated by whitespace, like this:

```
(tag:value tag:value ...)
```

The formal ABNF syntax definition for Mbus addresses and their elements is as follows:

```
mbus_address      = "(" *WSP *1address_list *WSP ")"
address_list      = address_element
                  / address_element 1*WSP address_list

address_element    = address_tag ":" address_value

address_tag        = 1*32(ALPHA)

address_value      = 1*64(%x21-27 / %x2A-7E)
                  ; any 7-bit US-ASCII character
                  ; excluding white space, delete,
                  ; control characters, "(" and ")"
```

Note that this and other ABNF definitions in this document use the non-terminal symbols defined in Section 2.

An address_tag MUST be unique within an Mbus address, i.e., it MUST only occur once.

Each entity has a fixed sequence of address elements constituting its address and MUST only process messages sent to addresses that either match all elements or consist of a subset of its own address elements. The order of address elements in an address sequence is not relevant. Two address elements match if both their tags and their values are equivalent. Equivalence for address element and address value strings means that each octet in the one string has the same value as the corresponding octet in the second string. For example, an entity with an address of:

```
(conf:test media:audio module:engine app:rat id:4711-1@192.168.1.1)
```

will process messages sent to

```
(media:audio module:engine)
```

and

(module:engine)

but must ignore messages sent to

(conf:test media:audio module:engine app:rat id:123-4@192.168.1.1
foo:bar)

and

(foo:bar)

A message that should be processed by all entities requires an empty set of address elements.

4.1 Mandatory Address Elements

Each Mbus entity **MUST** provide one mandatory address element that allows it to identify the entity. The element tag is "id" and the value **MUST** be composed of the following components:

- o The IP address of the interface that is used for sending messages to the Mbus. For IPv4 this is the address in dotted decimal notation. For IPv6 the interface-ID-part of the node's link-local address in textual representation as specified in RFC 2373 [3] **MUST** be used.

In this specification, this part is called the "host-ID".

- o An identifier ("entity-ID") that is unique within the scope of a single host-ID. The entity comprises two parts. For systems where the concept of a process ID is applicable it is **RECOMMENDED** that this identifier be composed using a process-ID and a per-process disambiguator for different Mbus entities of a process. If a process ID is not available, this part of the entity-ID may be randomly chosen (it is recommended that at least a 32 bit random number is chosen). Both numbers are represented in decimal textual form and **MUST** be separated by a '-' (ASCII x2d) character.

Note that the entity-ID cannot be the port number of the endpoint used for sending messages to the Mbus because implementations **MAY** use the common Mbus port number for sending to and receiving from the multicast group (as specified in Section 6).

The complete syntax definition for the entity identifier is as follows:

id-element = "id:" id-value
id-value = entity-id "@" host-id
entity-id = 1*10DIGIT "-" 1*5DIGIT
host-id = (IPv4address / IPv6address)

Please refer to [3] for the productions of IPv4address and IPv6address.

An example for an id element:

id:4711-99@192.168.1.1

5. Message Syntax

5.1 Message Encoding

All messages **MUST** use the UTF-8 character encoding. Note that US ASCII is a subset of UTF-8 and requires no additional encoding, and that a message encoded with UTF-8 will not contain zero bytes.

Each Message **MAY** be encrypted using a secret key algorithm as defined in Section 11.

5.2 Message Header

The fields in the header are separated by white space characters, and followed by CRLF. The format of the header is as follows:

```
msg_header = "mbus/1.0" 1*WSP SeqNum 1*WSP TimeStamp 1*WSP  
             MessageType 1*WSP SrcAddr 1*WSP DestAddr 1*WSP AckList
```

The header fields are explained in Message Format (Section 3). Here are the ABNF syntax definitions for the header fields:

SeqNum = 1*10DIGIT ; numeric range 0 - 2³²-1
TimeStamp = 1*13DIGIT
MessageType = "R" / "U"
SrcAddr = mbus_address
DestAddr = mbus_address

AckList = "(" *WSP *1(1*DIGIT *(1*WSP 1*10DIGIT)) *WSP ")"

See Section 4 for a definition of "mbus_address".

The syntax definition of a complete message is as follows:

mbus_message = msg_header *1(CRLF msg_payload)

msg_payload = mbus_command *(CRLF mbus_command)

The definition of production rules for an Mbus command is given in Section 5.3.

5.3 Command Syntax

The header is followed by zero, one, or more, commands to be delivered to the Mbus entities indicated by the DestAddr field. Each command consists of a command name that is followed by a list of zero, or more parameters and is terminated by a newline.

command (parameter parameter ...)

Syntactically, the command name **MUST** be a 'symbol' as defined in the following table. The parameters **MAY** be any data type drawn from the following table:

val	= Integer / Float / String / List / Symbol / Data
Integer	= *1"- " 1*DIGIT
Float	= *1"- " 1*DIGIT "." 1*DIGIT
String	= DQUOTE *CHAR DQUOTE ; see below for escape characters
List	= "(" *WSP *1(val *(1*WSP val)) *WSP ")"
Symbol	= ALPHA *(ALPHA / DIGIT / "_" / "-" / ".")
Data	= "<" *base64 ">"

Boolean values are encoded as an integer, with the value of zero representing false, and non-zero representing true.

String parameters in the payload **MUST** be enclosed in the double quote (") character. Within strings, the escape character is the backslash (\), and the following escape sequences are defined:

Escape Sequence	Meaning
\\	\
\"	"
\n	newline

List parameters do not have to be homogeneous lists, i.e., they can contain parameters of different types.

Opaque data is represented as Base64-encoded (see RFC 1521 [7]) character strings surrounded by "< " and "> "

The ABNF syntax definition for Mbus commands is as follows:

`mbus_command = command_name arglist`

`command_name = Symbol`

`arglist = List`

Command names **SHOULD** be constructed hierarchically to group conceptually related commands under a common hierarchy. The delimiter between names in the hierarchy **MUST** be "." (dot). Application profiles **MUST NOT** define commands starting with "mbus.".

The Mbus addressing scheme defined in Section 4 allows specifying incomplete addresses by omitting certain elements of an address element list, enabling entities to send commands to a group of Mbus entities. Therefore, all command names **SHOULD** be unambiguous in a way that it is possible to interpret or ignore them without considering the message's address.

A set of commands within a certain hierarchy that **MUST** be understood by every entity is defined in Section 9.

6. Transport

All messages are transmitted as UDP messages, with two possible alternatives:

1. Local multicast/broadcast:

This transport class **MUST** be used for all messages that are not sent to a fully qualified target address. It **MAY** also be used for messages that are sent to a fully qualified target address. It **MUST** be provided by conforming implementations. See Section 6.1 for details.

2. Directed unicast:

This transport class **MAY** be used for messages that are sent to a fully qualified destination address. It is **OPTIONAL** and does not have to be provided by conforming implementations.

A fully qualified target address is an Mbus address of an existing Mbus entity in an Mbus session. An implementation can identify an Mbus address as fully qualified by maintaining a list of known entities within an Mbus session. Each known entity has its own unique, fully qualified Mbus address.

Messages are transmitted in UDP datagrams, a maximum message size of 64 KBytes **MUST NOT** be exceeded. It is **RECOMMENDED** that applications using a non host-local scope do not exceed a message size of the link MTU.

Note that "unicast", "multicast" and "broadcast" mean IP Unicast, IP Multicast and IP Broadcast respectively. It is possible to send an Mbus message that is addressed to a single entity using IP Multicast.

This specification deals with both Mbus over UDP/IPv4 and Mbus over UDP/IPv6.

6.1 Local Multicast/Broadcast

In general, the Mbus uses multicast with a limited scope for message transport. Two different Mbus multicast scopes are defined, either of which can be configured to be used with an Mbus session:

1. host-local
2. link-local

Participants of an Mbus session have to know the multicast address in advance -- it cannot be negotiated during the session since it is already needed for initial communication between the Mbus entities during the bootstrapping phase. It also cannot be allocated prior to an Mbus session because there would be no mechanism to announce the allocated address to all potential Mbus entities. Therefore, the multicast address has to be assigned statically. This document defines the use of statically assigned addresses and also provides a

specification of how an Mbus session can be configured to use non-standard, unassigned addresses (see Section 12).

The following sections specify the use of multicast addresses for IPv4 and IPv6.

6.1.1 Mbus multicast groups for IPv4

For IPv4, a statically assigned, scope-relative multicast address as defined by RFC 2365 [11] is used. The offset for the scope relative address for Mbus is 8 (MBUS, see <http://www.iana.org/assignments/multicast-addresses> [19]).

Different scopes are defined by RFC 2365 [11]. The IPv4 Local Scope (239.255.0.0/16) is the minimal enclosing scope for administratively scoped multicast (as defined by RFC 2365 [11]) and not further divisible -- its exact extent is site dependent.

For the IPv4 Local Scope, applying the rules of RFC 2365 [11] and using the assigned offset of 8, the Mbus multicast address is therefore 239.255.255.247.

For IPv4, the different defined Mbus scopes (host-local and link-local) are to be realized as follows:

host-local multicast: Unless configured otherwise, the assigned scope-relative Mbus address in the Local Scope (239.255.255.247 as of RFC 2365 [11]) MUST be used. Mbus UDP datagrams SHOULD be sent with a TTL of 0.

link-local multicast: Unless configured otherwise, the assigned scope-relative Mbus address in the Local Scope (239.255.255.247 as of RFC 2365 [11]) MUST be used. Mbus UDP datagrams SHOULD be sent with a TTL of 1.

6.1.2 Mbus multicast groups for IPv6

IPv6 has different address ranges for different multicast scopes and distinguishes node local and link local scopes, that are implemented as a set of address prefixes for the different address ranges (RFC 2373 [3]). The link-local prefix is FF02, the node-local prefix is FF01. A permanently assigned multicast address will be used for Mbus multicast communication, i.e., an address that is independent of the scope value and that can be used for all scopes. Implementations for IPv6 MUST use the scope-independent address and the appropriate prefix for the selected scope. For host-local Mbus communication the IPv6 node-local scope prefix MUST be used, for link-local Mbus communication the IPv6 link-local scope prefix MUST be used.

The permanent IPv6 multicast address for Mbus/Ipv6 is FF0X:0:0:0:0:0:0:300.

FF0X:0:0:0:0:0:0:300 SHOULD be used for Mbus/IPv6 where the X in FF0X indicates that the scope is not fixed because this is an all scope address. This means, for node-local scope, FF01:0:0:0:0:0:0:300 SHOULD be used and for link-local scope FF02:0:0:0:0:0:0:300 SHOULD be used. See RFC 2375 [4] for IPv6 multicast address assignments.

If a single application system is distributed across several co-located hosts, link local scope SHOULD be used for multicasting Mbus messages that potentially have recipients on the other hosts. The Mbus protocol is not intended (and hence deliberately not designed) for communication between hosts not on the same link. See Section 12 for specifications of Mbus configuration mechanisms.

6.1.3 Use of Broadcast

In situations where multicast is not available, broadcast MAY be used instead. In these cases an IP broadcast address for the connected network SHOULD be used for sending. The node-local broadcast address for IPv6 is FF01:0:0:0:0:0:0:1, the link-local broadcast address for IPv6 is FF02:0:0:0:0:0:0:1. For IPv4, the generic broadcast address (for link-local broadcast) is 255.255.255.255. It is RECOMMENDED that IPv4-implementations use the generic broadcast address and a TTL of zero for host-local broadcast.

Broadcast MUST NOT be used in situations where multicast is available and supported by all systems participating in an Mbus session.

See Section 12 for configuring the use of broadcast.

6.1.4 Mbus UDP Port Number

The registered Mbus UDP port number is 47000.

6.2 Directed Unicast

Directed unicast (via UDP) to the port of a specific application is an alternative transport class to multicast. Directed unicast is an OPTIONAL optimization and MAY be used by Mbus implementations for delivering messages addressed to a single application entity only -- the address of which the Mbus implementation has learned from other message exchanges before. Note that the DestAddr field of such messages MUST be filled in properly nevertheless. Every Mbus entity SHOULD use a single unique endpoint address for sending messages to the Mbus multicast group or to individual receiving entities. A

unique endpoint address is a tuple consisting of the entity's IP address and a UDP source port number, where the port number is different from the standard Mbus port number.

Messages **MUST** only be sent via unicast if the Mbus target address is unique and if the sending entity can verify that the receiving entity uses a unique endpoint address. The latter can be verified by considering the last message received from that entity.

Note that several Mbus entities, say within the same process, may share a common transport address; in this case, the contents of the destination address field is used to further dispatch the message. Given the definition of "unique endpoint address" above, the use of a shared endpoint address and a dispatcher still allows other Mbus entities to send unicast messages to one of the entities that share the endpoint address. So this can be considered an implementation detail.

Messages with an empty target address list **MUST** always be sent to all Mbus entities (via multicast if available).

The following algorithm can be used by sending entities to determine whether an Mbus address is unique considering the current set of Mbus entities:

```
let ta=the target address;
iterate through the set of all
currently known Mbus addresses {
  let ti=the address in each iteration;
  count the addresses for which
  the predicate isSubsetOf(ta,ti) yields true;
}
```

If the count of matching addresses is exactly 1 the address is unique. The following algorithm can be used for the predicate `isSubsetOf`, that checks whether the second message matches the first according to the rules specified in Section 4. (A match means that a receiving entity that uses the second Mbus address must also process received messages with the first address as a target address.)

```
isSubsetOf(addr a1,a2) yields true, iff
  every address element of a1 is contained
  in a2's address element list.
```

An address element `a1` is contained in an address element list if the list contains an element that is equal to `a1`. An address element is considered equal to another address element if it has the same values for both of the two address element fields (tag and value).

7. Reliability

While most messages are expected to be sent using unreliable transport, it may be necessary to deliver some messages reliably. Reliability can be selected on a per message basis by means of the `MessageType` field. Reliable delivery is supported for messages with a single recipient only; i.e., to a fully qualified Mbus address. An entity can thus only send reliable messages to known addresses, i.e., it can only send reliable messages to entities that have announced their existence on the Mbus (e.g., by means of `mbus.hello()` messages as defined in Section 9.1). A sending entity **MUST NOT** send a message reliably if the target address is not unique. (See Section 6 for the specification of an algorithm to determine whether an address is unique.) A receiving entity **MUST** only process and acknowledge a reliable message if the destination address exactly matches its own source address (the destination address **MUST NOT** be a subset of the source address).

Disallowing reliable message delivery for messages sent to multiple destinations is motivated by simplicity of the implementation as well as the protocol. The desired effect can be achieved at the application layer by sending individual reliable messages to each fully qualified destination address, if the membership information for the Mbus session is available.

Each message is tagged with a message sequence number. If the `MessageType` is "R", the sender expects an acknowledgment from the recipient within a short period of time. If the acknowledgment is not received within this interval, the sender **MUST** retransmit the message (with the same message sequence number), increase the timeout, and restart the timer. Messages **MUST** be retransmitted a small number of times (see below) before the transmission or the recipient are considered to have failed. If the message is not delivered successfully, the sending application is notified. In this case, it is up to the application to determine the specific actions (if any) to be taken.

Reliable messages **MUST** be acknowledged by adding their SeqNum to the AckList field of a message sent to the originator of the reliable message. This message **MUST** be sent to a fully qualified Mbus target address. Multiple acknowledgments **MAY** be sent in a single message. Implementations **MAY** either piggy-back the AckList onto another message sent to the same destination, or **MAY** send a dedicated acknowledgment message, with no commands in the message payload part.

The precise procedures are as follows:

Sender: A sender A of a reliable message M to receiver B **MUST** transmit the message either via IP-multicast or via IP-unicast, keep a copy of M, initialize a retransmission counter N to '1', and start a retransmission timer T (initialized to T_r). If an acknowledgment is received from B, timer T **MUST** be cancelled and the copy of M is discarded. If T expires, the message M **MUST** be retransmitted, the counter N **MUST** be incremented by one, and the timer **MUST** be restarted (set to $N * T_r$). If N exceeds the retransmission threshold N_r , the transmission is assumed to have failed, further retransmission attempts **MUST NOT** be undertaken, the copy of M **MUST** be discarded, and the sending application **SHOULD** be notified.

Receiver: A receiver B of a reliable message from a sender A **MUST** acknowledge reception of the message within a time period $T_c < T_r$. This **MAY** be done by means of a dedicated acknowledgment message or by piggy-backing the acknowledgment on another message addressed only to A.

Receiver optimization: In a simple implementation, B may choose to immediately send a dedicated acknowledgment message. However, for efficiency, it could add the SeqNum of the received message to a sender-specific list of acknowledgments; if the added SeqNum is the first acknowledgment in the list, B **SHOULD** start an acknowledgment timer TA (initialized to T_c). When the timer expires, B **SHOULD** create a dedicated acknowledgment message and send it to A. If B is to transmit another Mbus message addressed only to A, it should piggy-back the acknowledgments onto this message and cancel TA. In either case, B should store a copy of the acknowledgment list as a single entry in the per-sender copy list, keep this entry for a period T_k , and empty the acknowledgment list. In case any of the messages kept in an entry of the copy list is received again from A, the entire acknowledgment list stored in this entry is scheduled for (re-) transmission following the above rules.

Constants and Algorithms: The following constants and algorithms **SHOULD** be used by implementations:

$T_r=100\text{ms}$

$N_r=3$

$T_c=70\text{ms}$

$T_k=((N_r)*(N_r+1)/2)*T_r$

8. Awareness of other Entities

Before Mbus entities can communicate with one another, they need to mutually find out about their existence. After this bootstrap procedure that each Mbus entity goes through all other entities listening to the same Mbus know about the newcomer and the newcomer has learned about all the other entities. Furthermore, entities need to be able to notice the failure (or leaving) of other entities.

Any Mbus entity **MUST** announce its presence (on the Mbus) after starting up. This is to be done repeatedly throughout its lifetime to address the issues of startup sequence: Entities should always become aware of other entities independent of the order of starting.

Each Mbus entity **MUST** maintain the number of Mbus session members and continuously update this number according to any observed changes. The mechanisms of how the existence and the leaving of other entities can be detected are dedicated Mbus messages for entity awareness: `mbus.hello` (Section 9.1) and `mbus.bye` (Section 9.2). Each Mbus protocol implementation **MUST** periodically send `mbus.hello` messages that are used by other entities to monitor the existence of that entity. If an entity has not received `mbus.hello` messages for a certain time (see Section 8.2) from an entity, the respective entity is considered to have left the Mbus and **MUST** be excluded from the set of currently known entities. Upon the reception of a `mbus.bye` message the respective entity is considered to have left the Mbus as well and **MUST** be excluded from the set of currently known entities immediately.

Each Mbus entity **MUST** send hello messages to the Mbus after startup. After transmission of the hello message, it **MUST** start a timer after the expiration of which the next hello message is to be transmitted. Transmission of hello messages **MUST NOT** be stopped unless the entity detaches from the Mbus. The interval for sending hello messages is dependent on the current number of entities in an Mbus group and can thus change dynamically in order to avoid congestion due to many entities sending hello messages at a constant high rate.

Section 8.1 specifies the calculation of hello message intervals that **MUST** be used by protocol implementations. Using the values that are calculated for obtaining the current hello message timer, the timeout for received hello messages is calculated in Section 8.2. Section 9 specifies the command synopsis for the corresponding Mbus messages.

8.1 Hello Message Transmission Interval

Since the number of entities in an Mbus session may vary, care must be taken to allow the Mbus protocol to automatically scale over a wide range of group sizes. The average rate at which hello messages are received would increase linearly to the number of entities in a session if the sending interval was set to a fixed value. Given an interval of 1 second this would mean that an entity taking part in an Mbus session with n entities would receive n hello messages per second. Assuming all entities resided on one host, this would lead to $n*n$ messages that have to be processed per second -- which is obviously not a viable solution for larger groups. It is therefore necessary to deploy dynamically adapted hello message intervals, taking varying numbers of entities into account. In the following, we specify an algorithm that **MUST** be used by implementors to calculate the interval for hello messages considering the observed number of Mbus entities.

The algorithm features the following characteristics:

- o The number of hello messages that are received by a single entity in a certain time unit remains approximately constant as the number of entities changes.
- o The effective interval that is used by a specific Mbus entity is randomized in order to avoid unintentional synchronization of hello messages within an Mbus session. The first hello message of an entity is also delayed by a certain random amount of time.
- o A timer reconsideration mechanism is deployed in order to adapt the interval more appropriately in situations where a rapid change of the number of entities is observed. This is useful when an entity joins an Mbus session and is still learning of the existence of other entities or when a larger number of entities leaves the Mbus at once.

8.1.1 Calculating the Interval for Hello Messages

The following variable names are used in the calculation specified below (all time values in milliseconds):

hello_p: The last time a hello message has been sent by a Mbus entity.

hello_now: The current time

hello_d: The deterministic calculated interval between hello messages.

hello_e: The effective (randomized) interval between hello messages.

hello_n: The time for the next scheduled transmission of a hello message.

entities_p: The numbers of entities at the time hello_n has been last recomputed.

entities: The number of currently known entities.

The interval between hello messages **MUST** be calculated as follows:

The number of currently known entities is multiplied by **c_hello_factor**, yielding the interval between hello messages in milliseconds. This is the deterministic calculated interval, denoted **hello_d**. The minimum value for **hello_d** is **c_hello_min** which yields

$$\text{hello_d} = \max(\text{c_hello_min}, \text{c_hello_factor} * \text{entities} * 1\text{ms}).$$

Section 8 provides a specification of how to obtain the number of currently known entities. Section 10 provides values for the constants **c_hello_factor** and **c_hello_min**.

The effective interval **hello_e** that is to be used by individual entities is calculated by multiplying **hello_d** with a randomly chosen number between **c_hello_dither_min** and **c_hello_dither_max** as follows:

$$\text{hello_e} = \text{c_hello_dither_min} + \text{RND} * (\text{c_hello_dither_max} - \text{c_hello_dither_min})$$

with **RND** being a random function that yields an even distribution between 0 and 1. See also Section 10.

hello_n, the time for the next hello message in milliseconds is set to **hello_e + hello_now**.

8.1.2 Initialization of Values

Upon joining an Mbus session a protocol implementation sets `hello_p=0`, `hello_now=0` and `entities=1`, `entities_p=1` (the Mbus entity itself) and then calculates the time for the next hello message as specified in Section 8.1.1. The next hello message is scheduled for transmission at `hello_n`.

8.1.3 Adjusting the Hello Message Interval when the Number of Entities increases

When the existence of a new entity is observed by a protocol implementation the number of currently known entities is updated. No further action concerning the calculation of the hello message interval is required. The reconsideration of the timer interval takes place when the current timer for the next hello message expires (see Section 8.1.5).

8.1.4 Adjusting the Hello Message Interval when the Number of Entities decreases

Upon realizing that an entity has left the Mbus the number of currently known entities is updated and the following algorithm should be used to reconsider the timer interval for hello messages:

1. The value for `hello_n` is updated by setting $\text{hello_n} = \text{hello_now} + (\text{entities}/\text{entities_p}) * (\text{hello_n} - \text{hello_now})$
2. The value for `hello_p` is updated by setting $\text{hello_p} = \text{hello_now} - (\text{entities}/\text{entities_p}) * (\text{hello_now} - \text{hello_p})$
3. The currently active timer for the next hello messages is cancelled and a new timer is started for `hello_n`.
4. `entities_p` is set to `entities`.

8.1.5 Expiration of hello timers

When the hello message timer expires, the protocol implementation **MUST** perform the following operations:

The hello interval `hello_e` is computed as specified in Section 8.1.1.

1. IF $\text{hello_e} + \text{hello_p} \leq \text{hello_now}$ THEN a hello message is transmitted. `hello_p` is set to `hello_now`, `hello_e` is calculated again as specified in Section 8.1.1 and `hello_n` is set to $\text{hello_e} + \text{hello_now}$.

2. ELSE IF `hello_e + hello_p > hello_now` THEN `hello_n` is set to `hello_e + hello_p`. A new timer for the next hello message is started to expire at `hello_n`. No hello message is transmitted.

`entities_p` is set to `entities`.

8.2 Calculating the Timeout for Mbus Entities

Whenever an Mbus entity has not heard for a time span of `c_hello_dead*(hello_d*c_hello_dither_max)` milliseconds from another Mbus entity it may consider this entity to have failed (or have quit silently). The number of the currently known entities MUST be updated accordingly. See Section 8.1.4 for details. Note that no need for any further action is necessarily implied from this observation.

Section 8.1.1 specifies how to obtain `hello_d`. Section 10 defines values for the constants `c_hello_dead` and `c_hello_dither_max`.

9. Messages

This section defines some basic application-independent messages that MUST be understood by all implementations; these messages are required for proper operation of the Mbus. This specification does not contain application-specific messages. These are to be defined outside of the basic Mbus protocol specification in separate Mbus profiles.

9.1 `mbus.hello`

Syntax:
`mbus.hello()`

Parameters: - none -

`mbus.hello` messages MUST be sent unreliably to all Mbus entities.

Each Mbus entity learns about other Mbus entities by observing their `mbus.hello` messages and tracking the sender address of each message and can thus calculate the current number of entities.

`mbus.hello` messages MUST be sent periodically in dynamically calculated intervals as specified in Section 8.

Upon startup the first `mbus.hello` message MUST be sent after a delay `hello_delay`, where `hello_delay` be a randomly chosen number between 0 and `c_hello_min` (see Section 10).

9.2 mbus.bye

Syntax: mbus.bye()

Parameters: - none -

An Mbus entity that is about to terminate (or "detach" from the Mbus) SHOULD announce this by transmitting an mbus.bye message. The mbus.bye message MUST be sent unreliably to all entities.

9.3 mbus.ping

Syntax: mbus.ping()

Parameters: - none -

mbus.ping can be used to solicit other entities to signal their existence by replying with an mbus.hello message. Each protocol implementation MUST understand mbus.ping and reply with an mbus.hello message. The reply hello message MUST be delayed for hello_delay milliseconds, where hello_delay be a randomly chosen number between 0 and c_hello_min (see Section 10). Several mbus.ping messages MAY be answered by a single mbus.hello: if one or more further mbus.ping messages are received while the entity is waiting hello_delay time units before transmitting the mbus.hello message, no extra mbus.hello message need be scheduled for those additional mbus.ping messages.

As specified in Section 9.1 hello messages MUST be sent unreliably to all Mbus entities. This is also the case for replies to ping messages. An entity that replies to mbus.ping with mbus.hello SHOULD stop any outstanding timers for hello messages after sending the hello message and schedule a new timer event for the subsequent hello message. (Note that using the variables and the algorithms of Section 8.1.1 this can be achieved by setting hello_p to hello_now.)

mbus.ping allows a new entity to quickly check for other entities without having to wait for the regular individual hello messages. By specifying a target address the new entity can restrict the solicitation for hello messages to a subset of entities it is interested in.

9.4 mbus.quit

Syntax:
mbus.quit()

Parameters: - none -

The mbus.quit message is used to request other entities to terminate themselves (and detach from the Mbus). Whether this request is honoured by receiving entities or not is application specific and not defined in this document.

The mbus.quit message can be multicast or sent reliably via unicast to a single Mbus entity or a group of entities.

9.5 mbus.waiting

Syntax:
mbus.waiting(condition)

Parameters:

symbol condition

The condition parameter is used to indicate that the entity transmitting this message is waiting for a particular event to occur.

An Mbus entity SHOULD be able to indicate that it is waiting for a certain event to happen (similar to a P() operation on a semaphore but without creating external state somewhere else). In conjunction with this, an Mbus entity SHOULD be capable of indicating to another entity that this condition is now satisfied (similar to a semaphore's V() operation).

The mbus.waiting message MAY be broadcast to all Mbus entities, MAY be multicast to an arbitrary subgroup, or MAY be unicast to a particular peer. Transmission of the mbus.waiting message MUST be unreliable and hence MUST be repeated at an application-defined interval (until the condition is satisfied).

If an application wants to indicate that it is waiting for several conditions to be met, several mbus.waiting messages are sent (possibly included in the same Mbus payload). Note that mbus.hello and mbus.waiting messages may also be transmitted in a single Mbus payload.

9.6 mbus.go

Syntax:
mbus.go(condition)

Parameters:

symbol condition

This parameter specifies which condition is met.

The mbus.go message is sent by an Mbus entity to "unblock" another Mbus entity -- which has indicated that it is waiting for a certain condition to be met. Only a single condition can be specified per mbus.go message. If several conditions are satisfied simultaneously multiple mbus.go messages MAY be combined in a single Mbus payload.

The mbus.go message MUST be sent reliably via unicast to the Mbus entity to unblock.

10. Constants

The following values for timers and counters mentioned in this document SHOULD be used by implementations:

Timer / Counter	Value	Unit
c_hello_factor	200	-
c_hello_min	1000	milliseconds
c_hello_dither_min	0.9	-
c_hello_dither_max	1.1	-
c_hello_dead	5	-

$T_r=100\text{ms}$

$N_r=3$

$T_c=70\text{ms}$

$T_k=((N_r)*(N_r+1)/2)*T_r$

11. Mbus Security

11.1 Security Model

In order to prevent accidental or malicious disturbance of Mbus communications through messages originated by applications from other users, message authentication is deployed (Section 11.3). For each message, a digest **MUST** be calculated based on the value of a shared secret key value. Receivers of messages **MUST** check if the sender belongs to the same Mbus security domain by re-calculating the digest and comparing it to the received value. The messages **MUST** only be processed further if both values are equal. In order to allow different simultaneous Mbus sessions at a given scope and to compensate defective implementations of host local multicast, message authentication **MUST** be provided by conforming implementations.

Privacy of Mbus message transport can be achieved by optionally using symmetric encryption methods (Section 11.2). Each message **MAY** be encrypted using an additional shared secret key and a symmetric encryption algorithm. Encryption is **OPTIONAL** for applications, i.e., it is allowed to configure an Mbus domain not to use encryption. But conforming implementations **MUST** provide the possibility to use message encryption (see below).

Message authentication and encryption can be parameterized: the algorithms to apply, the keys to use, etc. These and other parameters are defined in an Mbus configuration object that is accessible by all Mbus entities that participate in an Mbus session. In order to achieve interoperability conforming implementations **SHOULD** use the values provided by such an Mbus configuration. Section 12 defines the mandatory and optional parameters as well as storage procedures for different platforms. Only in cases where none of the options mentioned in Section 12 is applicable alternative methods of configuring Mbus protocol entities **MAY** be deployed.

The algorithms and procedures for applying encryption and authentication techniques are specified in the following sections.

11.2 Encryption

Encryption of messages is **OPTIONAL**, that means, an Mbus **MAY** be configured not to use encryption.

Implementations can choose between different encryption algorithms. Every conforming implementation **MUST** provide the AES [18] algorithm. In addition, the following algorithms **SHOULD** be supported: DES [16], 3DES (triple DES) [16] and IDEA [20].

For algorithms requiring en/decryption data to be padded to certain boundaries octets with a value of 0 **SHOULD** be used for padding characters.

The length of the encryption keys is determined by the currently used encryption algorithm. This means, the configured encryption key **MUST NOT** be shorter than the native key length for the currently configured algorithm.

DES implementations **MUST** use the DES Cipher Block Chaining (CBC) mode. DES keys (56 bits) **MUST** be encoded as 8 octets as described in RFC 1423 [12], resulting in 12 Base64-encoded characters. IDEA uses 128-bit keys (24 Base64-encoded characters). AES can use either 128-bit, 192-bit or 256-bit keys. For Mbus encryption using AES only 128-bit keys (24 Base64-encoded characters) **MUST** be used.

11.3 Message Authentication

For authentication of messages, hashed message authentication codes (HMACs) as described in RFC 2104 [5] are deployed. In general, implementations can choose between a number of digest algorithms. For Mbus authentication, the HMAC algorithm **MUST** be applied in the following way:

The keyed hash value is calculated using the HMAC algorithm specified in RFC 2104 [5]. The specific hash algorithm and the secret hash key **MUST** be obtained from the Mbus configuration (see Section 12).

The keyed hash values (see RFC 2104 [5]) **MUST** be truncated to 96 bits (12 octets).

Subsequently, the resulting 12 octets **MUST** be Base64-encoded, resulting in 16 Base64-encoded characters (see RFC 1521 [7]).

Either MD5 [15] or SHA-1 [17] **SHOULD** be used for message authentication codes (MACs). An implementation **MAY** provide MD5, whereas SHA-1 **MUST** be implemented.

The length of the hash keys is determined by the selected hashing algorithm. This means, the configured hash key **MUST NOT** be shorter than the native key length for the currently configured algorithm.

11.4 Procedures for Senders and Receivers

The algorithms that **MUST** be provided by implementations are AES and SHA-1.

See Section 12 for a specification of notations for Base64-strings.

A sender **MUST** apply the following operations to a message that is to be sent:

1. If encryption is enabled, the message **MUST** be encrypted using the configured algorithm and the configured encryption key. Padding (adding extra-characters) for block-ciphers **MUST** be applied as specified in Section 11.2. If encryption is not enabled, the message is left unchanged.
2. Subsequently, a message authentication code (MAC) for the (encrypted) message **MUST** be calculated using the configured HMAC-algorithm and the configured hash key.
3. The MAC **MUST** then be converted to Base64 encoding, resulting in 16 Base64-characters as specified in Section 11.3.
4. At last, the sender **MUST** construct the final message by placing the (encrypted) message after the base64-encoded MAC and a CRLF. The ABNF definition for the final message is as follows:

```
final_msg = MsgDigest CRLF encr_msg
```

```
MsgDigest = base64
```

```
encr_msg  = *OCTET
```

A receiver **MUST** apply the following operations to a message that it has received:

1. Separate the base64-encoded MAC from the (encrypted) message and decode the MAC.
2. Re-calculate the MAC for the message using the configured HMAC-algorithm and the configured hash key.
3. Compare the original MAC with re-calculated MAC. If they differ, the message **MUST** be discarded without further processing.
4. If encryption is enabled, the message **MUST** be decrypted using the configured algorithm and the configured encryption key. Trailing octets with a value of 0 **MUST** be deleted. If the message does not

start with the string "mbus/" the message **MUST** be discarded without further processing.

12. Mbus Configuration

An implementation **MUST** be configurable by the following parameters:

Configuration version

The version number of the given configuration entity. Version numbers allow implementations to check if they can process the entries of a given configuration entity. Version number are integer values. The version number for the version specified here is 1.

Encryption key

The secret key used for message encryption.

Hash key

The hash key used for message authentication.

Scope

The multicast scope to be used for sent messages.

The above parameters are mandatory and **MUST** be present in every Mbus configuration entity.

The following parameters are optional. When they are present they **MUST** be honored. When they are not present implementations **SHOULD** fall back to the predefined default values (as defined in Transport (Section 6)):

Address

The non-standard multicast address to use for message transport.

Use of Broadcast

It can be specified whether broadcast should be used. If broadcast has been configured implementations **SHOULD** use the network broadcast address (as specified in Section 6.1.3) instead of the standard multicast address.

Port Number

The non-standard UDP port number to use for message transport.

Two distinct facilities for parameter storage are considered: For Unix-like systems a per-user configuration file **SHOULD** be used and for Windows-95/98/NT/2000/XP systems a set of registry entries is defined that **SHOULD** be used. For other systems it is **RECOMMENDED** that the file-based configuration mechanism is used.

The syntax of the values for the respective parameter entries remains the same for both configuration facilities. The following defines a set of ABNF (see RFC 2234 [13]) productions that are later re-used for the definitions for the configuration file syntax and registry entries:

algo-id	=	"NOENCR" / "AES" / "DES" / "3DES" / "IDEA" / "HMAC-MD5-96" / "HMAC-SHA1-96"
scope	=	"HOSTLOCAL" / "LINKLOCAL"
key	=	base64
version_number	=	1*10DIGIT
key_value	=	(" algo-id "," key ")
address	=	IPv4address / IPv6address / "BROADCAST"
port	=	1*5DIGIT ; values from 0 through 65535

Given the definition above, a key entry **MUST** be specified using this notation:

("algo-id","base64string")

algo-id is one of the character strings specified above. For algo-id=="NOENCR" the other fields are ignored. The delimiting commas **MUST** always be present though.

A Base64 string consists of the characters defined in the Base64 char-set (see RFC 1521 [7]) including all possible padding characters, i.e., the length of a Base64-string is always a multiple of 4.

The scope parameter is used to configure an IP-Multicast scope and may be set to either "HOSTLOCAL" or "LINKLOCAL". Implementations **SHOULD** choose an appropriate IP-Multicast scope depending on the

value of this parameter and construct an effective IP-Address considering the specifications of Section 6.1.

The use of broadcast is configured by providing the value "BROADCAST" for the address field. If broadcast has been configured, implementations SHOULD use the network broadcast address for the used IP version instead of the standard multicast address.

The version_number parameter specifies a version number for the used configuration entity.

12.1 File based parameter storage

The file name for an Mbus configuration file is ".mbus" in the user's home-directory. If an environment variable called MBUS is defined implementations SHOULD interpret the value of this variable as a fully qualified file name that is to be used for the configuration file. Implementations MUST ensure that this file has appropriate file permissions that prevent other users to read or write it. The file MUST exist before a conference is initiated. Its contents MUST be UTF-8 encoded and MUST comply to the following syntax definition:

```
mbus-file      =   mbus-topic LF *(entry LF)
mbus-topic     =   "[MBUS]"
entry          =   1*(version_info / hashkey_info
                    / encryptionkey_info / scope_info
                    / port_info / address_info)
version_info   =   "CONFIG_VERSION=" version_number
hashkey_info   =   "HASHKEY=" key_value
encrkey_info   =   "ENCRYPTIONKEY=" key_value
scope_info     =   "SCOPE=" scope
port_info      =   "PORT=" port
address_info   =   "ADDRESS=" address
```

The following entries are defined: CONFIG_VERSION, HASHKEY, ENCRYPTIONKEY, SCOPE, PORT, ADDRESS.

The entries CONFIG_VERSION, HASHKEY and ENCRYPTIONKEY are mandatory, they MUST be present in every Mbus configuration file. The order of entries is not significant.

An example for an Mbus configuration file:

```
[MBUS]
CONFIG_VERSION=1
HASHKEY=(HMAC-MD5-96,MTIzMTU2MTg5MTEy)
ENCRYPTIONKEY=(DES,MTIzMTU2MQ==)
SCOPE=HOSTLOCAL
ADDRESS=224.255.222.239
PORT=47000
```

12.2 Registry-based parameter storage

For systems lacking the concept of a user's home-directory as a place for configuration files the suggested database for configuration settings (e.g., the Windows9x, Windows NT, Windows 2000, Windows XP registry) SHOULD be used. The hierarchy for Mbus related registry entries is as follows:

HKEY_CURRENT_USER\Software\Mbus

The entries in this hierarchy section are:

Name	Type	ABNF production
CONFIG_VERSION	DWORD	version_number
HASHKEY	String	key_value
ENCRYPTIONKEY	String	key_value
SCOPE	String	scope
ADDRESS	String	address
PORT	DWORD	port

The same syntax for key values as for the file based configuration facility MUST be used.

13. Security Considerations

The Mbus security mechanisms are specified in Section 11.1.

It should be noted that the Mbus transport specification defines a mandatory baseline set of algorithms that have to be supported by implementations. This baseline set is intended to provide reasonable security by mandating algorithms and key lengths that are considered to be cryptographically strong enough at the time of writing.

However, in order to allow for efficiency it is allowable to use cryptographically weaker algorithms, for example HMAC-MD5 instead of

HMAC-SHA1. Furthermore, encryption can be turned off completely if privacy is provided by other means or not considered important for a certain application.

Users of the Mbus should therefore be aware of the selected security configuration and should check if it meets the security demands for a given application. Since every implementation **MUST** provide the cryptographically strong algorithm it should always be possible to configure an Mbus in a way that secure communication with authentication and privacy is ensured.

In any way, application developers should be aware of incorrect IP implementations that do not conform to RFC 1122 [2] and do send datagrams with TTL values of zero, resulting in Mbus messages sent to the local network link although a user might have selected host local scope in the Mbus configuration. When using administratively scoped multicast, users cannot always assume the presence of correctly configured boundary routers. In these cases the use of encryption **SHOULD** be considered if privacy is desired.

14. IANA Considerations

The IANA has assigned a scope-relative multicast address with an offset of 8 for Mbus/IPv4. The IPv6 permanent multicast address is FF0X:0:0:0:0:0:300.

The registered Mbus UDP port number is 47000.

15. References

- [1] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, March 1997.
- [2] Braden, R., "Requirements for Internet Hosts -- Communication Layers", STD 3, RFC 1122, October 1989.
- [3] Hinden, R. and S. Deering, "IP Version 6 Addressing Architecture", RFC 2373, July 1998.
- [4] Hinden, R. and S. Deering, "IPv6 Multicast Address Assignments", RFC 2375, July 1998.
- [5] Krawczyk, H., Bellare, M. and R. Canetti, "HMAC: Keyed-Hashing for Message Authentication", RFC 2104, February 1997.
- [6] Resnick, P., Editor, "Internet Message Format", RFC 2822, April 2001.

- [7] Borenstein, N. and N. Freed, "MIME (Multipurpose Internet Mail Extensions) Part One: Mechanisms for Specifying and Describing the Format of Internet Message Bodies", RFC 1521, September 1993.
- [8] Schulzrinne, H., Casner, S., Frederick, R. and V. Jacobsen, "RTP: A Transport Protocol for Real-Time Applications", RFC 1889, January 1996.
- [9] Handley, M., Schulzrinne, H., Schooler, E. and J. Rosenberg, "SIP: Session Initiation Protocol", RFC 2543, March 1999.
- [10] Handley, M. and V. Jacobsen, "SDP: Session Description Protocol", RFC 2327, April 1998.
- [11] Meyer, D., "Administratively Scoped IP Multicast", BCP 23, RFC 2365, July 1998.
- [12] Balenson, D., "Privacy Enhancement for Internet Electronic Mail: Part III: Algorithms, Modes, and Identifiers", RFC 1423, February 1993.
- [13] Crocker, D. and P. Overell, "Augmented BNF for Syntax Specifications: ABNF", RFC 2234, November 1997.
- [14] Myers, J., "SMTP Service Extension for Authentication", RFC 2554, March 1999.
- [15] Rivest, R., "The MD5 Message-Digest Algorithm", RFC 1321, April 1992.
- [16] U.S. DEPARTMENT OF COMMERCE/National Institute of Standards and Technology, "Data Encryption Standard (DES)", FIPS PUB 46-3, Category Computer Security, Subcategory Cryptography, October 1999.
- [17] U.S. DEPARTMENT OF COMMERCE/National Institute of Standards and Technology, "Secure Hash Standard", FIPS PUB 180-1, April 1995.
- [18] Daemen, J.D. and V.R. Rijmen, "AES Proposal: Rijndael", March 1999.
- [19] IANA, "Internet Multicast Addresses", URL <http://www.iana.org/assignments/multicast-addresses>, May 2001.
- [20] Schneier, B., "Applied Cryptography", Edition 2, Publisher John Wiley & Sons, Inc., status: non-normative, 1996.

Appendix A. About References

Please note that the list of references contains normative as well as non-normative references. Each Non-normative references is marked as "status: non-normative". All unmarked references are normative.

Appendix B. Limitations and Future Work

The Mbus is a light-weight local coordination mechanism and deliberately not designed for larger scope coordination. It is expected to be used on a single node or -- at most -- on a single network link.

Therefore the Mbus protocol does not contain features that would be required to qualify it for the use over the global Internet:

There are no mechanisms to provide congestion control. The issue of congestion control is a general problem for multicast protocols. The Mbus allows for un-acknowledged messages that are sent unreliably, for example as event notifications, from one entity to another. Since negative acknowledgements are not defined there is no way the sender could realize that it is flooding another entity or congesting a low bandwidth network segment.

The reliability mechanism, i.e., the retransmission timers, are designed to provide effective, responsive message transport on local links but are not suited to cope with larger delays that could be introduced from router queues etc.

Some experiments are currently underway to test the applicability of bridges between different distributed Mbus domains without changing the basic protocol semantics. Since the use of such bridges should be orthogonal to the basic Mbus protocol definitions and since these experiments are still work in progress there is no mention of this concept in this specification.

Authors' Addresses

Joerg Ott
TZI, Universitaet Bremen
Bibliothekstr. 1
Bremen 28359
Germany

Phone: +49.421.201-7028
Fax: +49.421.218-7000
EMail: jo@tzi.uni-bremen.de

Colin Perkins
USC Information Sciences Institute
3811 N. Fairfax Drive #200
Arlington VA 22203
USA

EMail: csp@isi.edu

Dirk Kutscher
TZI, Universitaet Bremen
Bibliothekstr. 1
Bremen 28359
Germany

Phone: +49.421.218-7595
Fax: +49.421.218-7000
EMail: dku@tzi.uni-bremen.de

Full Copyright Statement

Copyright (C) The Internet Society (2002). All Rights Reserved.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this paragraph are included on all such copies and derivative works. However, this document itself may not be modified in any way, such as by removing the copyright notice or references to the Internet Society or other Internet organizations, except as needed for the purpose of developing Internet standards in which case the procedures for copyrights defined in the Internet Standards process must be followed, or as required to translate it into languages other than English.

The limited permissions granted above are perpetual and will not be revoked by the Internet Society or its successors or assigns.

This document and the information contained herein is provided on an "AS IS" basis and THE INTERNET SOCIETY AND THE INTERNET ENGINEERING TASK FORCE DISCLAIMS ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Acknowledgement

Funding for the RFC Editor function is currently provided by the Internet Society.