

Network Working Group
Request for Comments: 5354
Category: Experimental

R. Stewart
Q. Xie
The Resource Group
M. Stillman
Nokia
M. Tuexen
Muenster Univ. of Applied Sciences
September 2008

Aggregate Server Access Protocol (ASAP) and Endpoint Handlespace Redundancy Protocol (ENRP) Parameters

Status of This Memo

This memo defines an Experimental Protocol for the Internet community. It does not specify an Internet standard of any kind. Discussion and suggestions for improvement are requested. Distribution of this memo is unlimited.

Abstract

This document details the parameters of the Aggregate Server Access Protocol (ASAP) and Endpoint Handlespace Redundancy Protocol (ENRP) defined within the Reliable Server Pooling (RSerPool) architecture.

Table of Contents

1. Introduction	3
1.1. Conventions	3
2. Parameters in General	3
3. ENRP-ASAP Common Parameters	3
3.1. IPv4 Address Parameter	6
3.2. IPv6 Address Parameter	6
3.3. DCCP Transport Parameter	7
3.4. SCTP Transport Parameter	8
3.5. TCP Transport Parameter	9
3.6. UDP Transport Parameter	9
3.7. UDP-Lite Transport Parameter	10
3.8. Pool Member Selection Policy Parameter	11
3.9. Pool Handle Parameter	12
3.10. Pool Element Parameter	12
3.11. Server Information Parameter	13
3.12. Operation Error Parameter	14
3.12.1. Unspecified Error	15
3.12.2. Unrecognized Parameter Error	15
3.12.3. Unrecognized Message Error	15
3.12.4. Invalid Values Error	16
3.12.5. Non-Unique PE Identifier Error	16
3.12.6. Inconsistent Pool Policy Error	16
3.12.7. Lack of Resources Error	16
3.12.8. Inconsistent Transport Type Error	16
3.12.9. Inconsistent Data/Control Configuration Error	16
3.12.10. Rejected Due to Security Considerations	16
3.12.11. Unknown Pool Handle Error	17
3.13. Cookie Parameter	17
3.14. PE Identifier Parameter	17
3.15. PE Checksum Parameter	18
3.16. Opaque Transport Parameter	18
4. Common Message Formats	18
5. IANA Considerations	20
5.1. A New Table for RSerPool Parameter Types	20
5.2. A New Table for RSerPool Error Causes	21
6. Security Considerations	21
7. Normative References	21

1. Introduction

The Aggregate Server Access Protocol (ASAP) [RFC5352], in conjunction with the Endpoint Handlespace Redundancy Protocol (ENRP) [RFC5353], provides a high-availability, data-transfer mechanism over IP networks.

Both protocols work together and so share many common parameters used in message formats. This document details the common message parameters shared between the two protocols. This document provides parameter formats only; for procedures and message composition, please refer to the respective [RFC5352] and [RFC5353] documents.

1.1. Conventions

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC2119].

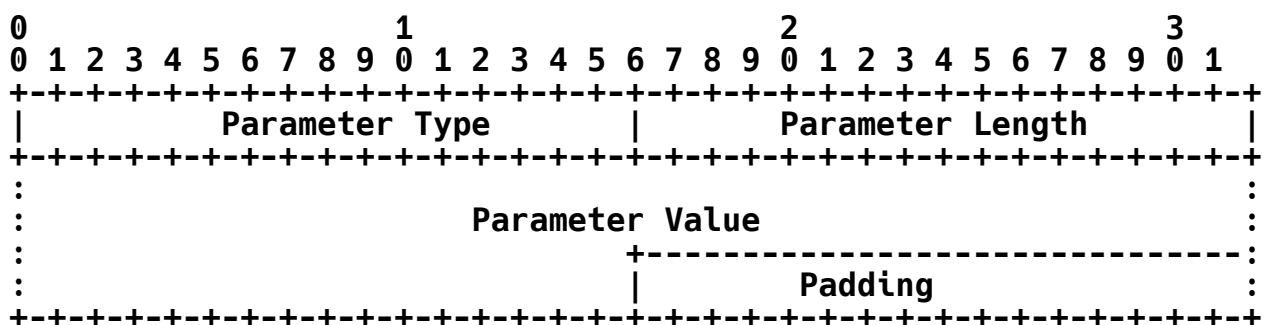
2. Parameters in General

All parameters described below MUST be in network byte order (aka Big Endian, i.e., the most significant byte first) during transmission.

Please note that messages in both ENRP and ASAP are often composed of multiple parameters. These parameters may also be nested. In such a case, a nested parameter will include the length of the padding between the nested parameters but not the last padding.

3. ENRP-ASAP Common Parameters

Parameters are defined in the following Type-Length-Value (TLV) format:



Parameter Type: 16 bits (unsigned integer)

The Type field is a 16-bit identifier of the type of parameter. It takes a value of 0 to 65534.

The value of 65535 is reserved for IETF-defined extensions.

Values, other than those defined in the specific ENRP parameter description, are reserved by IETF. (Additional types, when needed, will be defined in the future through appropriate IETF/IANA procedures.)

The Parameter Types are encoded such that the two bits of the highest-order specify the action that must be taken if the processing endpoint does not recognize the Parameter Type.

- 00 Stop processing this ENRP or ASAP message and discard it; do not process any further parameters within it.
- 01 Stop processing this ENRP or ASAP message and discard it; do not process any further parameters within it, and report the unrecognized parameter in an 'Unrecognized Parameter' error (see Section 3.12).
- 10 Skip this parameter and continue processing.
- 11 Skip this parameter and continue processing, but report the unrecognized parameter in an 'Unrecognized Parameter' error (see Section 3.12).

The values of parameter types are defined as follows:

Value	Parameter Type
0x0	(Reserved by IETF)
0x1	IPv4 Address
0x2	IPv6 Address
0x3	DCCP Transport
0x4	SCTP Transport
0x5	TCP Transport
0x6	UDP Transport
0x7	UDP-Lite
0x8	Pool Member Selection Policy
0x9	Pool Handle
0xa	Pool Element
0xb	Server Information
0xc	Operation Error
0xd	Cookie
0xe	PE Identifier
0xf	PE Checksum
0x10	Opaque Transport
0x11-0xfffffffffe	(Available for assignment)
0xfffffffff	IETF-defined extensions

Table 1

Parameter Length: 16 bits (unsigned integer)

The Parameter Length field contains the size of the parameter in bytes, including the Parameter Type, Parameter Length, and Parameter Value fields. Thus, a parameter with a zero-length Parameter Value field would have a Length field of 4.

The total length of a parameter (including Type, Parameter Length and Value fields) **MUST** be a multiple of 4 bytes. If the length of the parameter is not a multiple of 4 bytes, the sender **MUST** pad the parameter at the end (i.e., after the Parameter Value field) with all zero bytes. The length of this padding is not included in the Parameter Length field. A sender **MUST NOT** pad with more than 3 bytes. The receiver **MUST** ignore the padding bytes.

Parameter Value: variable length.


The Parameter Value field contains the actual information to be transferred in the parameter.

Parameter Padding: variable length.

The Parameter Padding, as described above.

3.1. IPv4 Address Parameter

This parameter defines a TLV that carries an IPv4 address.

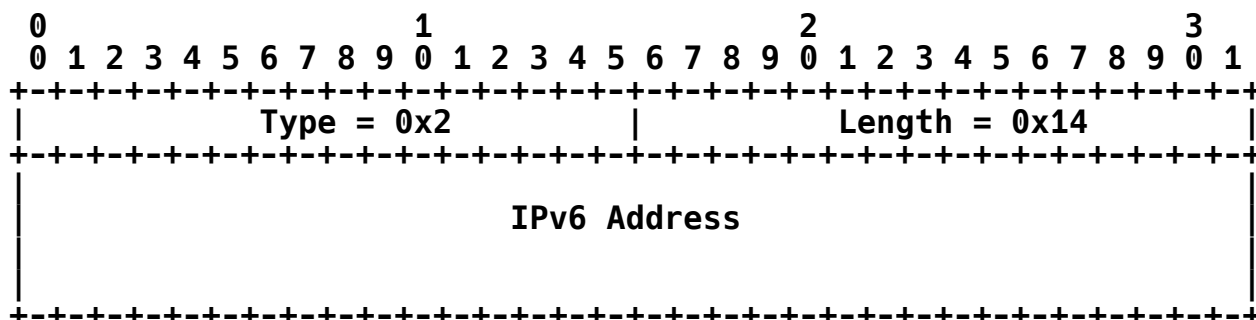


IPv4 Address: 32 bits (unsigned integer)

Contains an IPv4 address. It is binary encoded.

3.2. IPv6 Address Parameter

This parameter defines a TLV that carries an IPv6 address.




IPv6 Address: 128 bits (unsigned integer)

Contains an IPv6 address. It is binary encoded.

3.3. DCCP Transport Parameter

This parameter defines a TLV that describes a user transport using Datagram Congestion Control Protocol (DCCP).



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of octets, including the Type, Length, DCCP port, reserved fields, and IP Address Parameter.

DCCP Port: 16 bits (unsigned integer)

The DCCP port number signed to this DCCP user transport.

DCCP Service Code: 32 bits (unsigned integer)

The DCCP service code signed to this DCCP user transport.

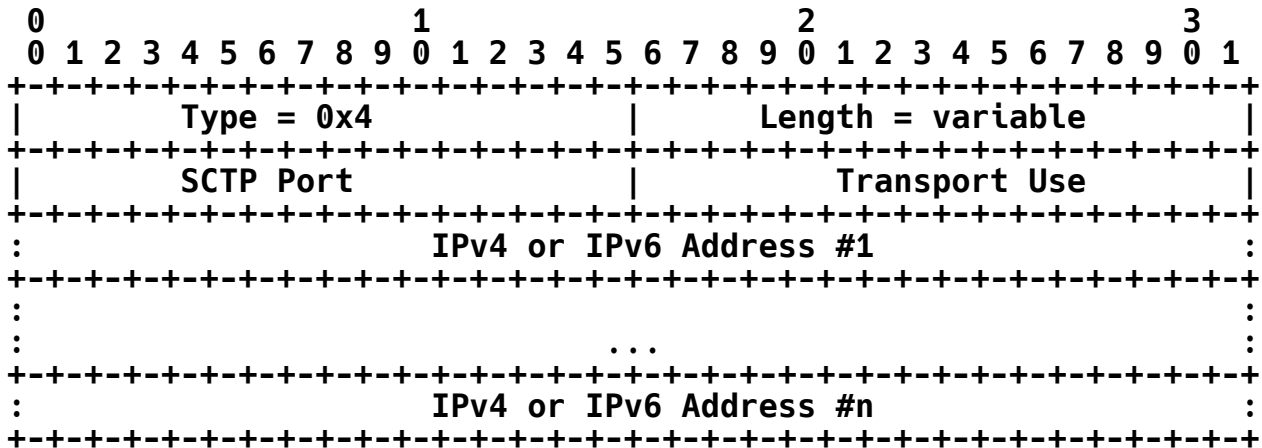
IPv4 or IPv6 Address

Indicates an IPv4 or IPv6 address parameter (as defined above in Section 3.1 and Section 3.2) assigned to this DCCP user transport. Unlike in an SCTP Transport parameter, only one IP address parameter can be present in a DCCP Transport parameter.

Note: The DCCP Port MUST NOT be used for control information. For this reason, no Transport Use field is provided. DCCP MUST always be treated as a "Data Only" type transport use.

3.4. SCTP Transport Parameter

This parameter defines a TLV that describes a user transport using Stream Control Transport Protocol (SCTP).



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of octets, including the Type, Length, SCTP port, reserved fields, and all IP Address Parameters present.

SCTP Port: 16 bits (unsigned integer)

The SCTP port number signed to this SCTP user transport.

Transport Use: 16 bits (unsigned integer)

This field represents how the pool element intends this transport address to be used. The field **MUST** be populated with one of the following values:

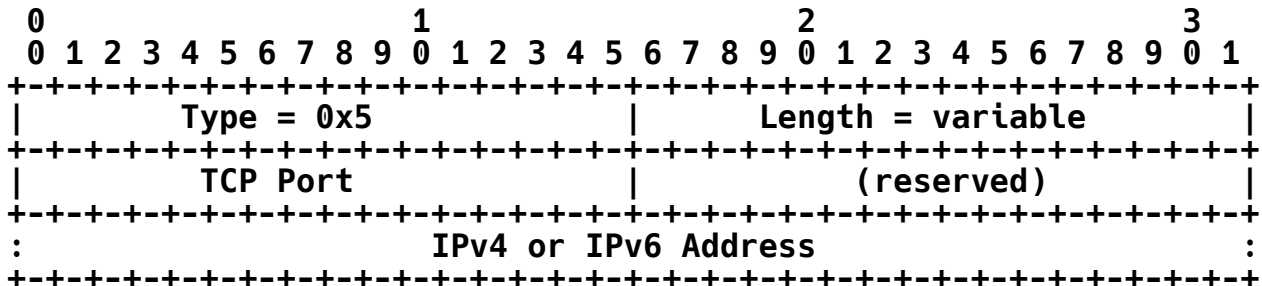
Type	Value
DATA ONLY	0x0000
DATA plus CONTROL	0x0001

IPv4 or IPv6 Address #1 - #n

Each indicates an IPv4 or IPv6 address parameter (as defined above in Section 3.1 and Section 3.2) assigned to this SCTP user transport. An SCTP Transport parameter may have a mixed list of IPv4 and IPv6 addresses and at least one IP address parameter **MUST** be present in an SCTP Transport parameter.

3.5. TCP Transport Parameter

This parameter defines a TLV that describes a user transport using TCP protocol.



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of octets, including the Type, Length, TCP port, reserved fields, and IP Address Parameter.

TCP Port: 16 bits (unsigned integer)

The TCP port number signed to this TCP user transport.


IPv4 or IPv6 Address

Indicates an IPv4 or IPv6 address parameter (as defined above in Section 3.1 and Section 3.2) assigned to this TCP user transport. Unlike in an SCTP Transport parameter, only one IP Address parameter can be present in a TCP Transport parameter.

Note: The TCP Port MUST NOT be used for control information. For this reason, no Transport Use field is provided. TCP MUST always be treated as a "Data Only" type transport use.

3.6. UDP Transport Parameter

This parameter defines a TLV that describes a user transport using UDP protocol.



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of octets, including the Type, Length, UDP port, reserved fields, and IP Address Parameter.

UDP Port: 16 bits (unsigned integer)

The UDP port number signed to this UDP user transport.


IPv4 or IPv6 Address

Indicates an IPv4 or IPv6 address parameter (as defined above in Section 3.1 and Section 3.2) assigned to this UDP user transport. Unlike in an SCTP Transport parameter, only one IP Address parameter can be present in a UDP Transport parameter.

Note: The UDP Port MUST NOT be used for control information. For this reason, no Transport Use field is provided. UDP MUST always be treated as a "Data Only" type transport use.

3.7. UDP-Lite Transport Parameter

This parameter defines a TLV that describes a user transport using UDP-Lite protocol.



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of octets, including the Type, Length, UDP-Lite port, reserved fields, and IP Address Parameter.

UDP Port: 16 bits (unsigned integer)

The UDP-Lite port number signed to this UDP-Lite user transport.

IPv4 or IPv6 Address

Indicates an IPv4 or IPv6 address parameter (as defined above in Section 3.1 and Section 3.2) assigned to this UDP-Lite user transport. Unlike in an SCTP Transport parameter, only one IP address parameter can be present in a UDP-Lite transport parameter.

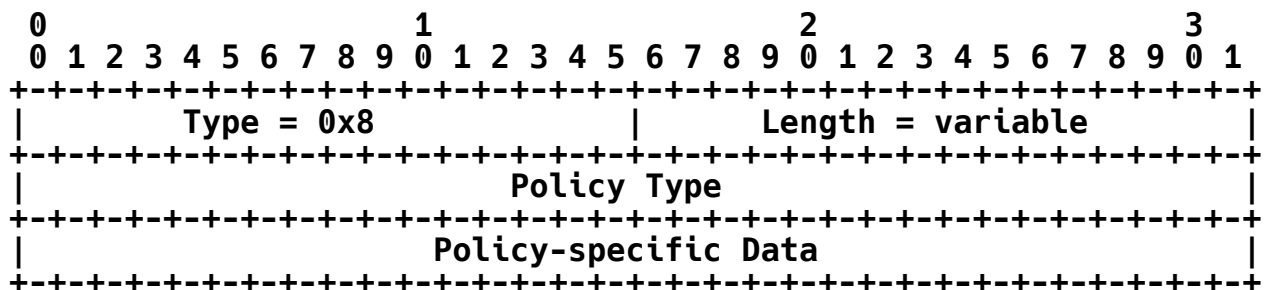
Note: The UDP-Lite Port **MUST NOT** be used for control information. For this reason, no Transport Use field is provided. UDP-Lite **MUST** always be treated as a "Data Only" type transport use.

3.8. Pool Member Selection Policy Parameter

This parameter defines a pool member selection policy. RSerPool supports multiple pool member selection policies and also allows the definition of new selection policies in the future.

The enforcement rules and handling procedures of all the policies are defined in [RFC5352].

All pool member selection policies, both present and future, **MUST** use the following general parameter format:



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of octets, including the Type, Length, Policy Type, and the Policy-specific Data fields.

Note, the Length field value will **NOT** include any padding at the end of the parameter.

Policy Type: 32 bits (unsigned integer)


Specifies the type of selection policy. The values are defined in [RFC5356].

Policy-specific Data:

The structure and fields for each presently defined policy type are described in detail in [RFC5356].

3.9. Pool Handle Parameter

This parameter holds a pool handle.



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of octets, including the Type, Length, and Pool Handle string.

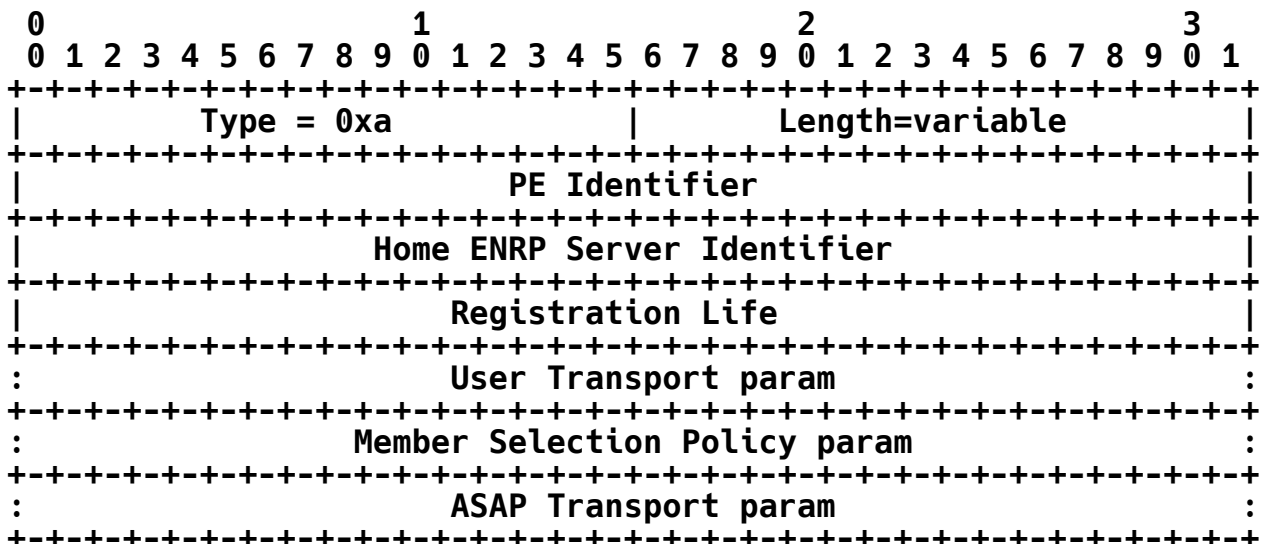
Note, the value in the Length field will NOT cover any padding at the end of the parameter.

Pool Handle

Defined as a sequence of (Length - 4) bytes.

3.10. Pool Element Parameter

This parameter is used in multiple ENRP messages to represent an ASAP endpoint (i.e., a Pool Element (PE) in a pool) and the associated information, such as its transport address, selection policy, and other operational or status information of the PE.



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of octets, including the Type, Length, PE Identifier, Registration Life, User Transport, and Member Selection Policy parameters.

Note, the value in the Length field will NOT cover any padding at the end of this Pool Element parameter.

PE Identifier: 32 bits (unsigned integer)

Uniquely identifies the PE in the pool. The PE picks its identifier when it starts up.

Home ENRP Server Identifier: 32 bits (unsigned integer)

Indicates the current Home ENRP server of this PE. Set to all 0s if the PE's Home ENRP server is undetermined.

Registration Life: 32 bits (signed integer)

Indicates the lifetime of the registration in number of seconds. A value of -1 indicates infinite lifetime.

User Transport

This can be either an DCCP, SCTP, TCP, UDP, UDP-Lite, or Opaque Transport parameter (see Section 3.3, Section 3.4, Section 3.5, Section 3.6, Section 3.7, and Section 3.16). A PE MUST have one and only one User Transport.

Member Selection Policy


Contains one of the defined member selection policy parameters (see Section 3.8).

ASAP Transport

This indicates the ASAP transport address of the PE and MUST be an SCTP type transport parameter (see Section 3.4).

3.11. Server Information Parameter

This parameter is used in ENRP to pass basic information of an ENRP server.



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of bytes.
Note, the value in the Length field will NOT cover any padding at the end of the parameter.

Server ID: 32 bits (unsigned integer)

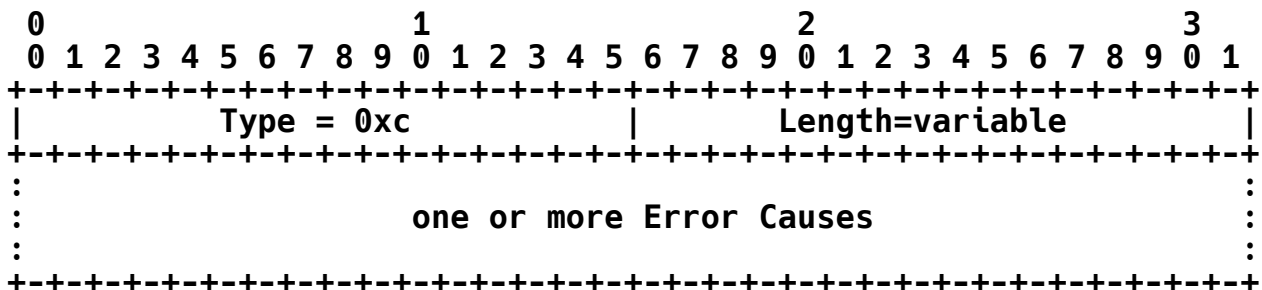
This is the ID of the ENRP server, as defined in [RFC5353].

Server Transport:

This is an SCTP Transport Parameter, as defined in Section 3.4, that contains the network access address(es), SCTP port number, etc. of the ENRP server.

3.12. Operation Error Parameter

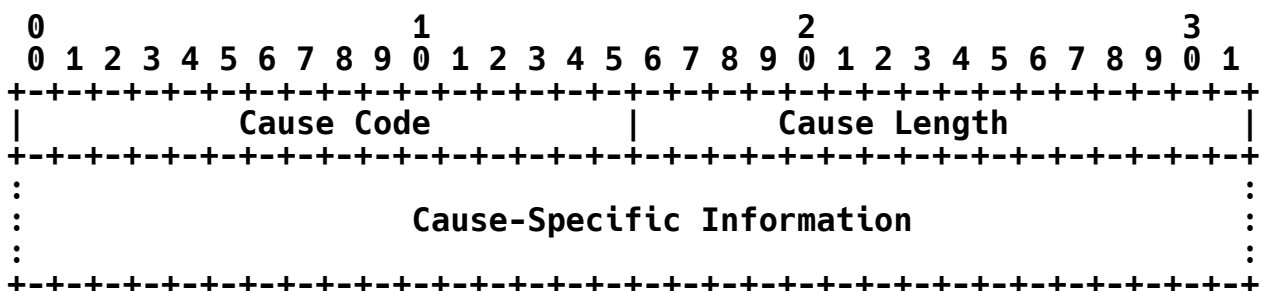
This parameter is used in both ENRP and ASAP for a message sender to report an error(s) to a message receiver.



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of bytes.
Note, the value in the Length field will NOT cover any padding at the end of the parameter.

Error causes are defined as variable-length parameters using the following format:



Cause Code: 16 bits (unsigned integer)

Defines the type of error condition being reported.

Cause Code Value	Cause Code
0x0	Unspecified Error
0x1	Unrecognized Parameter
0x2	Unrecognized Message
0x3	Invalid Values
0x4	Non-unique PE Identifier
0x5	Inconsistent Pooling Policy
0x6	Lack of Resources
0x7	Inconsistent Transport Type
0x8	Inconsistent Data/Control Configuration
0x9	Unknown Pool Handle
0xa	Rejected due to security considerations
0xb - 0xffff	(Available for assignment)

Table 2

Cause Length: 16 bits (unsigned integer)

Set to the size of the parameter in bytes, including the Cause Code, Cause Length, and Cause-Specific Information fields, but not including any padding at the end of this error cause TLV.

Cause-specific Information: variable length

This field carries the details of the error condition.

The following subsections (Section 3.12.1 - Section 3.12.9) define specific error causes.

3.12.1. Unspecified Error

This error cause is used to report an unspecified error by the sender. There is no cause specific information.

3.12.2. Unrecognized Parameter Error

This error cause is used to report an unrecognized parameter. The complete, unrecognized parameter TLV is included as cause-specific information. If a message contains multiple unrecognized parameters, multiple error causes are used.

3.12.3. Unrecognized Message Error

This error cause is used to report an unrecognized message. The unrecognized message TLV is included as cause-specific information.

3.12.4. Invalid Values Error

This error cause is used to report one or more invalid values found in a received parameter. The offending TLV that contains the invalid value(s) is included as cause-specific information.

3.12.5. Non-Unique PE Identifier Error

This error cause is used by an ENRP server to indicate to a registering PE that the PE Identifier it chooses has already been used by another PE in the pool. There is no cause-specific information.

3.12.6. Inconsistent Pool Policy Error

This error cause is used by an ENRP server to indicate to a registering PE that the pool policy it chooses does not match the overall policy of the pool. A Pool Member Selection Policy TLV (see Section 3.8) that indicates the overall pool policy is included as cause-specific information.

3.12.7. Lack of Resources Error

This error cause is used to indicate that the sender does not have certain resources to perform a requested function. There is no cause specific information.

3.12.8. Inconsistent Transport Type Error

This error cause is used by an ENRP server to indicate to a registering PE that the User Transport it chooses does not match the overall user transport of the pool. A Transport TLV that indicates the overall pool user transport type is included as cause-specific information.

3.12.9. Inconsistent Data/Control Configuration Error

This error cause is used by an ENRP server to indicate to a registering PE that the Transport Use field in the User Transport it sent in its registration is inconsistent to the pool's overall data/control channel configuration. There is no cause-specific information.

3.12.10. Rejected Due to Security Considerations

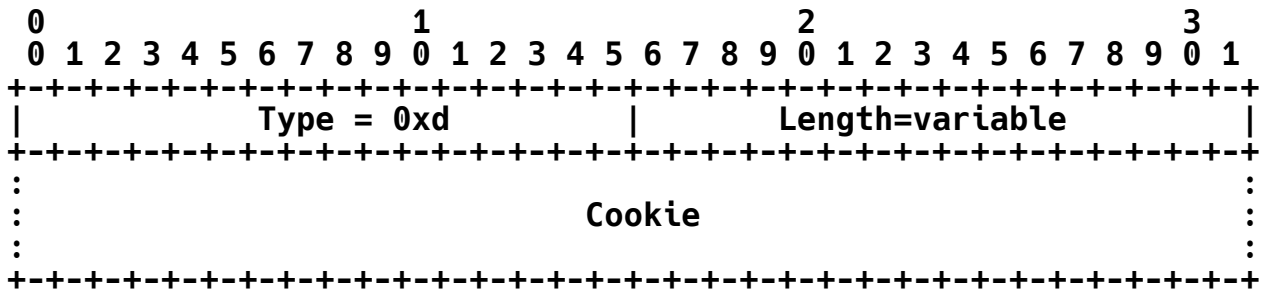
This error cause is used by any endpoint to indicate a rejection of a request due to a failure in security credentials or authorizations.

3.12.11. Unknown Pool Handle Error

This error cause is used by an ENRP server to indicate to a PE or PU that the requested pool is unknown by the server. There is no cause-specific information.

3.13. Cookie Parameter

This parameter defines a TLV that carries a Cookie.



Length: 16 bits (unsigned integer)


Indicates the entire length of the parameter in number of bytes, including the Type, Length, and Cookie.

Cookie: variable length

The Cookie is an arbitrary byte string of (Length - 4) bytes.

3.14. PE Identifier Parameter

This parameter defines a TLV that carries a PE Identifier.




PE Identifier: 32 bits (unsigned integer)

Uniquely identifies the PE in the pool. The PE picks its identifier when it starts up. See [RFC5352] for recommendations on PE identifier generation.

3.15. PE Checksum Parameter

This parameter defines a TLV that carries a PE Checksum.




PE Checksum: 16 bits (unsigned integer)

An overall checksum of all PEs in the current handlespace owned by an ENRP server (which is normally the sender of this TLV). The definition and calculation of this checksum is defined in [RFC5353].

3.16. Opaque Transport Parameter

This parameter defines a TLV that carries opaque transport information.



Length: 16 bits (unsigned integer)

Indicates the entire length of the parameter in number of bytes, including the Type, Length, and Opaque Transport Data.

Opaque Transport Data: variable length

The Opaque Transport Data is an arbitrary byte string of (Length - 4) bytes.

4. Common Message Formats

The figure below illustrates the common format for all ASAP and ENRP messages. Each message is formatted with a Message Type field, a message-specific Flag field, a Message Length field, and a Value field.

message with all zero bytes and this padding is not included in the Message Length field. The sender should never pad with more than 3 bytes. The receiver MUST ignore the padding bytes.

5. IANA Considerations

This document (RFC 5354) is the reference for all registrations described in this section. All registrations have been listed on the RSerPool Parameters page.

5.1. A New Table for RSerPool Parameter Types

RSerPool Parameter Types are maintained by IANA. Thirteen initial values have been assigned by IANA, as described in Table 1. IANA created a new table, "RSerPool Parameter Types":

Value	Parameter Type
0x0	(Reserved by IETF)
0x1	IPv4 Address
0x2	IPv6 Address
0x3	DCCP Transport
0x4	SCTP Transport
0x5	TCP Transport
0x6	UDP Transport
0x7	UDP-Lite
0x8	Pool Member Selection Policy
0x9	Pool Handle
0xa	Pool Element
0xb	Server Information
0xc	Operation Error
0xd	Cookie
0xe	PE Identifier
0xf	PE Checksum
0x10	Opaque Transport
0xffffffff	IETF-defined extensions
others	(Reserved by IETF)

Requests to register an RSerPool Parameter Type in this table should be sent to IANA. The number must be unique. The "Specification Required" policy of [RFC5226] MUST be applied.

5.2. A New Table for RSerPool Error Causes

RSerPool Error Causes are maintained by IANA. Eleven initial values have been assigned by IANA, as described in Table 2. IANA created a new table, "RSerPool Error Causes":

Cause Code Value	Cause Code
0x0	Unspecified Error
0x1	Unrecognized Parameter
0x2	Unrecognized Message
0x3	Invalid Values
0x4	Non-Unique PE Identifier
0x5	Inconsistent Pooling Policy
0x6	Lack of Resources
0x7	Inconsistent Transport Type
0x8	Inconsistent Data/Control Configuration
0x9	Unknown Pool Handle
0xa	Rejected Due to Security Considerations
others	(Reserved by IETF)

Requests to register an RSerPool Error Cause in this table should be sent to IANA. The number must be unique. The "Specification Required" policy of [RFC5226] MUST be applied.

6. Security Considerations

This document contains common parameter formats only. As such, it specifies no new security constraints on either ENRP or ASAP. Details on ENRP and ASAP security constraints are addressed in [RFC5353] and [RFC5352].

7. Normative References

- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, March 1997.
- [RFC5226] Narten, T. and H. Alvestrand, "Guidelines for Writing an IANA Considerations Section in RFCs", BCP 26, RFC 5226, May 2008.
- [RFC5352] Stewart, R., Xie, Q., Stillman, M., and M. Tuexen, "Aggregate Server Access Protocol (ASAP)", RFC 5352, September 2008.

[RFC5353] Xie, Q., Stewart, R., Stillman, M., Tuexen, M., and A. Silverton, "Endpoint Handlespace Redundancy Protocol (ENRP)", RFC 5353, September 2008.

[RFC5356] Dreibholz, T. and M. Tuexen, "Reliable Server Pooling Policies", RFC 5356, September 2008.

Authors' Addresses

Randall R. Stewart
The Resource Group
1700 Pennsylvania Ave NW
Suite 560
Washington, DC 20006
USA

Phone:
EMail: randall.stewart@trgworld.com

Qiaobing Xie
The Resource Group
1700 Pennsylvania Ave NW
Suite 560
Washington, D.C., 20006
USA

Phone: +1 224-465-5954
EMail: Qiaobing.Xie@gmail.com

Maureen Stillman
Nokia
1167 Peachtree Ct.
Naperville, IL 60540
USA

EMail: maureen.stillman@nokia.com

Michael Tuexen
Muenster Univ. of Applied Sciences
Stegerwaldstr. 39
48565 Steinfurt
Germany

EMail: tuexen@fh-muenster.de

Full Copyright Statement

Copyright (C) The IETF Trust (2008).

This document is subject to the rights, licenses and restrictions contained in BCP 78, and except as set forth therein, the authors retain all their rights.

This document and the information contained herein are provided on an "AS IS" basis and THE CONTRIBUTOR, THE ORGANIZATION HE/SHE REPRESENTS OR IS SPONSORED BY (IF ANY), THE INTERNET SOCIETY, THE IETF TRUST AND THE INTERNET ENGINEERING TASK FORCE DISCLAIM ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Intellectual Property

The IETF takes no position regarding the validity or scope of any Intellectual Property Rights or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; nor does it represent that it has made any independent effort to identify any such rights. Information on the procedures with respect to rights in RFC documents can be found in BCP 78 and BCP 79.

Copies of IPR disclosures made to the IETF Secretariat and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the IETF on-line IPR repository at <http://www.ietf.org/ipr>.

The IETF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights that may cover technology that may be required to implement this standard. Please address the information to the IETF at ietf-ipr@ietf.org.