Internet Engineering Task Force (IETF)
Request for Comments: 5867
Category: Informational
ISSN: 2070-1721

J. Martocci, Ed.
Johnson Controls Inc.
P. De Mil
Ghent University - IBCN
N. Riou
Schneider Electric
W. Vermeylen
Arts Centre Vooruit

June 2010

Building Automation Routing Requirements in Low-Power and Lossy Networks

Abstract

The Routing Over Low-Power and Lossy (ROLL) networks Working Group has been chartered to work on routing solutions for Low-Power and Lossy Networks (LLNs) in various markets: industrial, commercial (building), home, and urban networks. Pursuant to this effort, this document defines the IPv6 routing requirements for building automation.

Status of This Memo

This document is not an Internet Standards Track specification; it is published for informational purposes.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Not all documents approved by the IESG are a candidate for any level of Internet Standard; see Section 2 of RFC 5741.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at http://www.rfc-editor.org/info/rfc5867.

Copyright Notice

Table of Contents

1. Introduction

The Routing Over Low-Power and Lossy (ROLL) networks Working Group has been chartered to work on routing solutions for Low-Power and Lossy Networks (LLNs) in various markets: industrial, commercial (building), home, and urban networks. Pursuant to this effort, this document defines the IPv6 routing requirements for building automation.

Commercial buildings have been fitted with pneumatic, and subsequently electronic, communication routes connecting sensors to their controllers for over one hundred years. Recent economic and technical advances in wireless communication allow facilities to increasingly utilize a wireless solution in lieu of a wired solution, thereby reducing installation costs while maintaining highly reliant communication.

The cost benefits and ease of installation of wireless sensors allow customers to further instrument their facilities with additional sensors, providing tighter control while yielding increased energy savings.

Wireless solutions will be adapted from their existing wired counterparts in many of the building applications including, but not limited to, heating, ventilation, and air conditioning (HVAC); lighting; physical security; fire; and elevator/lift systems. These devices will be developed to reduce installation costs while increasing installation and retrofit flexibility, as well as increasing the sensing fidelity to improve efficiency and building service quality.

Sensing devices may be battery-less, battery-powered, or mains-powered. Actuators and area controllers will be mains-powered. Due to building code and/or device density (e.g., equipment room), it is envisioned that a mix of wired and wireless sensors and actuators will be deployed within a building.

Building management systems (BMSs) are deployed in a large set of vertical markets including universities, hospitals, government facilities, kindergarten through high school (K-12), pharmaceutical manufacturing facilities, and single-tenant or multi-tenant office buildings. These buildings range in size from 100K-sq.-ft. structures (5-story office buildings), to 1M-sq.-ft. skyscrapers (100-story skyscrapers), to complex government facilities such as the Pentagon. The described topology is meant to be the model to be used in all of these types of environments but clearly must be tailored to the building class, building tenant, and vertical market being served.

Section 3 describes the necessary background to understand the context of building automation including the sensor, actuator, area controller, and zone controller layers of the topology; typical device density; and installation practices.

Section 4 defines the traffic flow of the aforementioned sensors, actuators, and controllers in commercial buildings.

Section 5 defines the full set of IPv6 routing requirements for commercial buildings.

Appendix A documents important commercial building requirements that are out of scope for routing yet will be essential to the final acceptance of the protocols used within the building.

Section 3 and Appendix A are mainly included for educational purposes.

The expressed aim of this document is to provide the set of IPv6 routing requirements for LLNs in buildings, as described in Section 5.

2. Terminology

For a description of the terminology used in this specification, please see [ROLL-TERM].

2.1. Requirements Language

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC2119].

3. Overview of Building Automation Networks

3.1. Introduction

To understand the network systems requirements of a building management system in a commercial building, this document uses a framework to describe the basic functions and composition of the system. A BMS is a hierarchical system of sensors, actuators, controllers, and user interface devices that interoperate to provide a safe and comfortable environment while constraining energy costs.

A BMS is divided functionally across different but interrelated building subsystems such as heating, ventilation, and air conditioning (HVAC); fire; security; lighting; shutters; and elevator/lift control systems, as denoted in Figure 1.

Much of the makeup of a BMS is optional and installed at the behest of the customer. Sensors and actuators have no standalone functionality. All other devices support partial or complete standalone functionality. These devices can optionally be tethered to form a more cohesive system. The customer requirements dictate the level of integration within the facility. This architecture provides excellent fault tolerance since each node is designed to operate in an independent mode if the higher layers are unavailable.

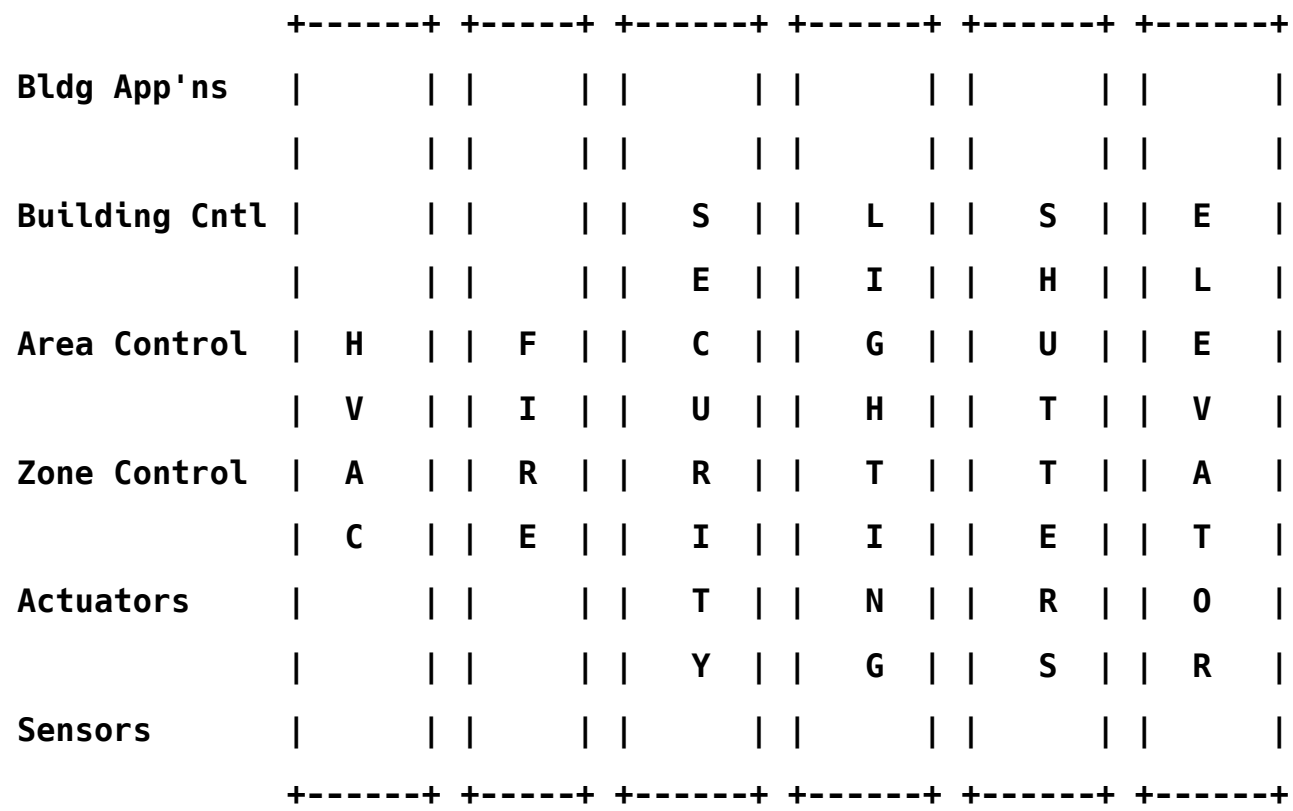

Figure 1: Building Systems and Devices

3.2. Building Systems Equipment

3.2.1. Sensors/Actuators

As Figure 1 indicates, a BMS may be composed of many functional stacks or silos that are interoperably woven together via building applications. Each silo has an array of sensors that monitor the environment and actuators that modify the environment, as determined by the upper layers of the BMS topology. The sensors typically are at the edge of the network structure, providing environmental data for the system. The actuators are the sensors' counterparts, modifying the characteristics of the system, based on the sensor data and the applications deployed.

3.2.2. Area Controllers

An area describes a small physical locale within a building, typically a room. HVAC (temperature and humidity) and lighting (room lighting, shades, solar loads) vendors oftentimes deploy area controllers. Area controllers are fed by sensor inputs that monitor

the environmental conditions within the room. Common sensors found in many rooms that feed the area controllers include temperature, occupancy, lighting load, solar load, and relative humidity. Sensors found in specialized rooms (such as chemistry labs) might include air flow, pressure, and CO2 and CO particle sensors. Room actuation includes temperature setpoint, lights, and blinds/curtains.

3.2.3. Zone Controllers

Zone controllers support a similar set of characteristics to area controllers, albeit for an extended space. A zone is normally a logical grouping or functional division of a commercial building. A zone may also coincidentally map to a physical locale such as a floor.

Zone controllers may have direct sensor inputs (smoke detectors for fire), controller inputs (room controllers for air handlers in HVAC), or both (door controllers and tamper sensors for security). Like area/room controllers, zone controllers are standalone devices that operate independently or may be attached to the larger network for more synergistic control.

3.3. Equipment Installation Methods

A BMS is installed very differently from most other IT networks. IT networks are typically installed as an overlay onto the existing environment and are installed from the inside out. That is, the network wiring infrastructure is installed; the switches, routers, and servers are connected and made operational; and finally, the endpoints (e.g., PCs, VoIP phones) are added.

BMSs, on the other hand, are installed from the outside in. That is, the endpoints (thermostats, lights, smoke detectors) are installed in the spaces first; local control is established in each room and tested for proper operation. The individual rooms are later lashed together into a subsystem (e.g., lighting). The individual subsystems (e.g., lighting, HVAC) then coalesce. Later, the entire system may be merged onto the enterprise network.

The rationale for this is partly due to the different construction trades having access to a building under construction at different times. The sheer size of a building often dictates that even a single trade may have multiple independent teams working simultaneously. Furthermore, the HVAC, lighting, and fire systems must be fully operational before the building can obtain its occupancy permit. Hence, the BMS must be in place and configured well before any of the IT servers (DHCP; Authentication, Authorization, and Accounting (AAA); DNS; etc.) are operational.

This implies that the BMS cannot rely on the availability of the IT network infrastructure or application servers. Rather, the BMS installation should be planned to dovetail into the IT system once the IT system is available for easy migration onto the IT network. Front-end planning of available switch ports, cable runs, access point (AP) placement, firewalls, and security policies will facilitate this adoption.

3.4. Device Density

Device density differs, depending on the application and as dictated by the local building code requirements. The following subsections detail typical installation densities for different applications.

3.4.1. HVAC Device Density

HVAC room applications typically have sensors/actuators and controllers spaced about 50 ft. apart. In most cases, there is a 3:1 ratio of sensors/actuators to controllers. That is, for each room there is an installed temperature sensor, flow sensor, and damper actuator for the associated room controller.

HVAC equipment room applications are quite different. An air handler system may have a single controller with up to 25 sensors and actuators within 50 ft. of the air handler. A chiller or boiler is also controlled with a single equipment controller instrumented with 25 sensors and actuators. Each of these devices would be individually addressed since the devices are mandated or optional as defined by the specified HVAC application. Air handlers typically serve one or two floors of the building. Chillers and boilers may be installed per floor, but many times they service a wing, building, or the entire complex via a central plant.

These numbers are typical. In special cases, such as clean rooms, operating rooms, pharmaceutical facilities, and labs, the ratio of sensors to controllers can increase by a factor of three. Tenant installations such as malls would opt for packaged units where much of the sensing and actuation is integrated into the unit; here, a single device address would serve the entire unit.

3.4.2. Fire Device Density

Fire systems are much more uniformly installed, with smoke detectors installed about every 50 ft. This is dictated by local building codes. Fire pull boxes are installed uniformly about every 150 ft. A fire controller will service a floor or wing. The fireman's fire panel will service the entire building and typically is installed in the atrium.

3.4.3. Lighting Device Density

Lighting is also very uniformly installed, with ballasts installed approximately every 10 ft. A lighting panel typically serves 48 to 64 zones. Wired systems tether many lights together into a single zone. Wireless systems configure each fixture independently to increase flexibility and reduce installation costs.

3.4.4. Physical Security Device Density

Security systems are non-uniformly oriented, with heavy density near doors and windows and lighter density in the building's interior space.

The recent influx of interior and perimeter camera systems is increasing the security footprint. These cameras are atypical endpoints requiring up to 1 megabit/second (Mbit/s) data rates per camera, as contrasted by the few kbit/s needed by most other BMS sensing equipment. Previously, camera systems had been deployed on proprietary wired high-speed networks. More recent implementations utilize wired or wireless IP cameras integrated into the enterprise LAN.

4. Traffic Pattern

The independent nature of the automation subsystems within a building can significantly affect network traffic patterns. Much of the real-time sensor environmental data and actuator control stays within the local LLN environment, while alarms and other event data will percolate to higher layers.

Each sensor in the LLN unicasts point to point (P2P) about 200 bytes of sensor data to its associated controller each minute and expects an application acknowledgment unicast returned from the destination. Each controller unicasts messages at a nominal rate of 6 kbit/minute to peer or supervisory controllers. Thirty percent of each node's packets are destined for other nodes within the LLN. Seventy percent of each node's packets are destined for an aggregation device (multipoint to point (MP2P)) and routed off the LLN. These messages also require a unicast acknowledgment from the destination. The above values assume direct node-to-node communication; meshing and error retransmissions are not considered.

Multicasts (point to multipoint (P2MP)) to all nodes in the LLN occur for node and object discovery when the network is first commissioned. This data is typically a one-time bind that is henceforth persisted. Lighting systems will also readily use multicasting during normal operations to turn banks of lights "on" and "off" simultaneously.

BMSs may be either polled or event-based. Polled data systems will generate a uniform and constant packet load on the network. Polled architectures, however, have proven not to be scalable. Today, most vendors have developed event-based systems that pass data on event. These systems are highly scalable and generate low data on the network at quiescence. Unfortunately, the systems will generate a heavy load on startup since all initial sensor data must migrate to the controller level. They also will generate a temporary but heavy load during firmware upgrades. This latter load can normally be mitigated by performing these downloads during off-peak hours.

Devices will also need to reference peers periodically for sensor data or to coordinate operation across systems. Normally, though, data will migrate from the sensor level upwards through the local and area levels, and then to the supervisory level. Traffic bottlenecks will typically form at the funnel point from the area controllers to the supervisory controllers.

Initial system startup after a controlled outage or unexpected power failure puts tremendous stress on the network and on the routing algorithms. A BMS is comprised of a myriad of control algorithms at the room, area, zone, and enterprise layers. When these control algorithms are at quiescence, the real-time data rate is small, and the network will not saturate. An overall network traffic load of 6 kbit/s is typical at quiescence. However, upon any power loss, the control loops and real-time data quickly atrophy. A short power disruption of only 10 minutes may have a long-term deleterious impact on the building control systems, taking many hours to regain proper control. Control applications that cannot handle this level of disruption (e.g., hospital operating rooms) must be fitted with a secondary power source.

Power disruptions are unexpected and in most cases will immediately impact lines-powered devices. Power disruptions, however, are transparent to battery-powered devices. These devices will continue to attempt to access the LLN during the outage. Battery-powered devices designed to buffer data that has not been delivered will further stress network operations when power returns.

Upon restart, lines-powered devices will naturally dither due to primary equipment delays or variance in the device self-tests. However, most lines-powered devices will be ready to access the LLN network within 10 seconds of power-up. Empirical testing indicates that routes acquired during startup will tend to be very oblique since the available neighbor lists are incomplete. This demands an adaptive routing protocol to allow for route optimization as the network stabilizes.

5. Building Automation Routing Requirements

Following are the building automation routing requirements for networks used to integrate building sensor, actuator, and control products. These requirements are written not presuming any preordained network topology, physical media (wired), or radio technology (wireless).

5.1. Device and Network Commissioning

Building control systems typically are installed and tested by electricians having little computer knowledge and no network communication knowledge whatsoever. These systems are often installed during the building construction phase, before the drywall and ceilings are in place. For new construction projects, the building enterprise IP network is not in place during installation of the building control system. For retrofit applications, the installer will still operate independently from the IP network so as not to affect network operations during the installation phase.

In traditional wired systems, correct operation of a light switch/ballast pair was as simple as flipping on the light switch. In wireless applications, the tradesperson has to assure the same operation, yet be sure the operation of the light switch is associated with the proper ballast.

System-level commissioning will later be deployed using a more computer savvy person with access to a commissioning device (e.g., a laptop computer). The completely installed and commissioned enterprise IP network may or may not be in place at this time. Following are the installation routing requirements.

5.1.1. Zero-Configuration Installation

It MUST be possible to fully commission network devices without requiring any additional commissioning device (e.g., a laptop). From the ROLL perspective, "zero configuration" means that a node can obtain an address and join the network on its own, without human intervention.

5.1.2. Local Testing

During installation, the room sensors, actuators, and controllers SHOULD be able to route packets amongst themselves and to any other device within the LLN, without requiring any additional routing infrastructure or routing configuration.

5.1.3. Device Replacement

To eliminate the need to reconfigure the application upon replacing a failed device in the LLN, the replaced device must be able to advertise the old IP address of the failed device in addition to its new IP address. The routing protocols MUST support hosts and routers that advertise multiple IPv6 addresses.

5.2. Scalability

Building control systems are designed for facilities from 50,000 sq. ft. to 1M+ sq. ft. The networks that support these systems must cost-effectively scale accordingly. In larger facilities, installation may occur simultaneously on various wings or floors, yet the end system must seamlessly merge. Following are the scalability requirements.

5.2.1. Network Domain

The routing protocol MUST be able to support networks with at least 2,000 nodes, where 1,000 nodes would act as routers and the other 1,000 nodes would be hosts. Subnetworks (e.g., rooms, primary equipment) within the network must support up to 255 sensors and/or actuators.

5.2.2. Peer-to-Peer Communication

The data domain for commercial BMSs may sprawl across a vast portion of the physical domain. For example, a chiller may reside in the facility's basement due to its size, yet the associated cooling towers will reside on the roof. The cold-water supply and return pipes snake through all of the intervening floors. The feedback control loops for these systems require data from across the facility.

A network device MUST be able to communicate in an end-to-end manner with any other device on the network. Thus, the routing protocol MUST provide routes between arbitrary hosts within the appropriate administrative domain.

5.3. Mobility

Most devices are affixed to walls or installed on ceilings within buildings. Hence, the mobility requirements for commercial buildings are few. However, in wireless environments, location tracking of occupants and assets is gaining favor. Asset-tracking applications, such as tracking capital equipment (e.g., wheelchairs) in medical

facilities, require monitoring movement with granularity of a minute; however, tracking babies in a pediatric ward would require latencies less than a few seconds.

The following subsections document the mobility requirements in the routing layer for mobile devices. Note, however, that mobility can be implemented at various layers of the system, and the specific requirements depend on the chosen layer. For instance, some devices may not depend on a static IP address and are capable of re-establishing application-level communications when given a new IP address. Alternatively, mobile IP may be used, or the set of routers in a building may give an impression of a building-wide network and allow devices to retain their addresses regardless of where they are, handling routing between the devices in the background.

5.3.1. Mobile Device Requirements

To minimize network dynamics, mobile devices while in motion should not be allowed to act as forwarding devices (routers) for other devices in the LLN. Network configuration should allow devices to be configured as routers or hosts.

5.3.1.1. Device Mobility within the LLN

An LLN typically spans a single floor in a commercial building. Mobile devices may move within this LLN. For example, a wheelchair may be moved from one room on the floor to another room on the same floor.

A mobile LLN device that moves within the confines of the same LLN SHOULD re-establish end-to-end communication with a fixed device also in the LLN within 5 seconds after it ceases movement. The LLN network convergence time should be less than 10 seconds once the mobile device stops moving.

5.3.1.2. Device Mobility across LLNs

A mobile device may move across LLNs, such as a wheelchair being moved to a different floor.

A mobile device that moves outside of its original LLN SHOULD re-establish end-to-end communication with a fixed device also in the new LLN within 10 seconds after the mobile device ceases movement. The network convergence time should be less than 20 seconds once the mobile device stops moving.

5.4. Resource Constrained Devices

Sensing and actuator device processing power and memory may be 4 orders of magnitude less (i.e., 10,000x) than many more traditional client devices on an IP network. The routing mechanisms must therefore be tailored to fit these resource constrained devices.

5.4.1. Limited Memory Footprint on Host Devices

The software size requirement for non-routing devices (e.g., sleeping sensors and actuators) SHOULD be implementable in 8-bit devices with no more than 128 KB of memory.

5.4.2. Limited Processing Power for Routers

The software size requirements for routing devices (e.g., room controllers) SHOULD be implementable in 8-bit devices with no more than 256 KB of flash memory.

5.4.3. Sleeping Devices

Sensing devices will, in some cases, utilize battery power or energy harvesting techniques for power and will operate mostly in a sleep mode to maintain power consumption within a modest budget. The routing protocol MUST take into account device characteristics such as power budget.

Typically, sensor battery life (2,000 mAh) needs to extend for at least 5 years when the device is transmitting its data (200 octets) once per minute over a low-power transceiver (25 mA) and expecting an application acknowledgment. In this case, the transmitting device must leave its receiver in a high-powered state, awaiting the return of the application ACK. To minimize this latency, a highly efficient routing protocol that minimizes hops, and hence end-to-end communication, is required. The routing protocol MUST take into account node properties, such as "low-powered node", that produce efficient low-latency routes that minimize radio "on" time for these devices.

Sleeping devices MUST be able to receive inbound data. Messages sent to battery-powered nodes MUST be buffered by the last-hop router for a period of at least 20 seconds when the destination node is currently in its sleep cycle.

5.5. Addressing

Building management systems require different communication schemes to solicit or post network information. Multicasts or anycasts need to be used to decipher unresolved references within a device when the device first joins the network.

As with any network communication, multicasting should be minimized. This is especially a problem for small embedded devices with limited network bandwidth. Multicasts are typically used for network joins and application binding in embedded systems. Routing MUST support anycast, unicast, and multicast.

5.6. Manageability

As previously noted in Section 3.3, installation of LLN devices within a BMS follows an "outside-in" work flow. Edge devices are installed first and tested for communication and application integrity. These devices are then aggregated into islands, then LLNs, and later affixed onto the enterprise network.

The need for diagnostics most often occurs during the installation and commissioning phase, although at times diagnostic information may be requested during normal operation. Battery-powered wireless devices typically will have a self-diagnostic mode that can be initiated via a button press on the device. The device will display its link status and/or end-to-end connectivity when the button is pressed. Lines-powered devices will continuously display communication status via a bank of LEDs, possibly denoting signal strength and end-to-end application connectivity.

The local diagnostics noted above oftentimes are suitable for defining room-level networks. However, as these devices aggregate, system-level diagnostics may need to be executed to ameliorate route vacillation, excessive hops, communication retries, and/or network bottlenecks.

In operational networks, due to the mission-critical nature of the application, the LLN devices will be temporally monitored by the higher layers to assure that communication integrity is maintained. Failure to maintain this communication will result in an alarm being forwarded to the enterprise network from the monitoring node for analysis and remediation.

In addition to the initial installation and commissioning of the system, it is equally important for the ongoing maintenance of the system to be simple and inexpensive. This implies a straightforward device swap when a failed device is replaced, as noted in Section 5.1.3.

5.6.1. Diagnostics

To improve diagnostics, the routing protocol SHOULD be able to be placed in and out of "verbose" mode. Verbose mode is a temporary debugging mode that provides additional communication information including, at least, the total number of routed packets sent and received, the number of routing failures (no route available), neighbor table members, and routing table entries. The data provided in verbose mode should be sufficient that a network connection graph could be constructed and maintained by the monitoring node.

Diagnostic data should be kept by the routers continuously and be available for solicitation at any time by any other node on the internetwork. Verbose mode will be activated/deactivated via unicast, multicast, or other means. Devices having available resources may elect to support verbose mode continuously.

5.6.2. Route Tracking

Route diagnostics SHOULD be supported, providing information such as route quality, number of hops, and available alternate active routes with associated costs. Route quality is the relative measure of "goodness" of the selected source to destination route as compared to alternate routes. This composite value may be measured as a function of hop count, signal strength, available power, existing active routes, or any other criteria deemed by ROLL as the route cost differentiator.

5.7. Route Selection

Route selection determines reliability and quality of the communication among the devices by optimizing routes over time and resolving any nuances developed at system startup when nodes are asynchronously adding themselves to the network.

5.7.1. Route Cost

The routing protocol MUST support a metric of route quality and optimize selection according to such metrics within constraints established for links along the routes. These metrics SHOULD reflect metrics such as signal strength, available bandwidth, hop count, energy availability, and communication error rates.

5.7.2. Route Adaptation

Communication routes MUST be adaptive and converge toward optimality of the chosen metric (e.g., signal quality, hop count) in time.

5.7.3. Route Redundancy

The routing layer SHOULD be configurable to allow secondary and tertiary routes to be established and used upon failure of the primary route.

5.7.4. Route Discovery Time

Mission-critical commercial applications (e.g., fire, security) require reliable communication and guaranteed end-to-end delivery of all messages in a timely fashion. Application-layer time-outs must be selected judiciously to cover anomalous conditions such as lost packets and/or route discoveries, yet not be set too large to over-damp the network response. If route discovery occurs during packet transmission time (reactive routing), it SHOULD NOT add more than 120 ms of latency to the packet delivery time.

5.7.5. Route Preference

The routing protocol SHOULD allow for the support of manually configured static preferred routes.

5.7.6. Real-Time Performance Measures

A node transmitting a "request with expected reply" to another node must send the message to the destination and receive the response in not more than 120 ms. This response time should be achievable with 5 or less hops in each direction. This requirement assumes network quiescence and a negligible turnaround time at the destination node.

5.7.7. Prioritized Routing

Network and application packet routing prioritization must be supported to assure that mission-critical applications (e.g., fire detection) cannot be deferred while less critical applications access the network. The routing protocol MUST be able to provide routes with different characteristics, also referred to as Quality of Service (QoS) routing.

5.8. Security Requirements

This section sets forth specific requirements that are placed on any protocols developed or used in the ROLL building environment, in order to ensure adequate security and retain suitable flexibility of use and function of the protocol.

Due to the variety of buildings and tenants, the BMSs must be completely configurable on-site.

Due to the quantity of the BMS devices (thousands) and their inaccessibility (oftentimes above ceilings), security configuration over the network is preferred over local configuration.

Wireless encryption and device authentication security policies need to be considered in commercial buildings, while keeping in mind the impact on the limited processing capabilities and additional latency incurred on the sensors, actuators, and controllers.

BMSs are typically highly configurable in the field, and hence the security policy is most often dictated by the type of building to which the BMS is being installed. Single-tenant owner-occupied office buildings installing lighting or HVAC control are candidates for implementing a low level of security on the LLN, especially when the LLN is not connected to an external network. Antithetically, military or pharmaceutical facilities require strong security policies. As noted in the installation procedures described in Sections 3.3 and 5.2, security policies MUST support dynamic configuration to allow for a low level of security during the installation phase (prior to building occupancy, when it may be appropriate to use only diagnostic levels of security), yet to make it possible to easily raise the security level network-wide during the commissioning phase of the system.

5.8.1. Building Security Use Case

LLNs for commercial building applications should always implement and use encrypted packets. However, depending on the state of the LLN, the security keys may either be:

1) a key obtained from a trust center already operable on the LLN;

2) a pre-shared static key as defined by the general contractor or its designee; or

3) a well-known default static key.

Unless a node entering the network had previously received its credentials from the trust center, the entering node will try to solicit the trust center for the network key. If the trust center is accessible, the trust center will MAC-authenticate the entering node and return the security keys. If the trust center is not available, the entering node will check to determine if it has been given a network key by an off-band means and use it to access the network. If no network key has been configured in the device, it will revert to the default network key and enter the network. If neither of these keys were valid, the device would signal via a fault LED.

This approach would allow for independent simplified commissioning, yet centralized authentication. The building owner or building type would then dictate when the trust center would be deployed. In many cases, the trust center need not be deployed until all of the local room commissioning is complete. Yet, at the province of the owner, the trust center may be deployed from the onset, thereby trading installation and commissioning flexibility for tighter security.

5.8.2. Authentication

Authentication SHOULD be optional on the LLN. Authentication SHOULD be fully configurable on-site. Authentication policy and updates MUST be routable over-the-air. Authentication SHOULD occur upon joining or rejoining a network. However, once authenticated, devices SHOULD NOT need to reauthenticate with any other devices in the LLN. Packets may need authentication at the source and destination nodes; however, packets routed through intermediate hops should not need reauthentication at each hop.

These requirements mean that at least one LLN routing protocol solution specification MUST include support for authentication.

5.8.3. Encryption

5.8.3.1. Encryption Types

Data encryption of packets MUST be supported by all protocol solution specifications. Support can be provided by use of a network-wide key and/or an application key. The network key would apply to all devices in the LLN. The application key would apply to a subset of devices in the LLN.

The network key and application key would be mutually exclusive. The routing protocol MUST allow routing a packet encrypted with an application key through forwarding devices without requiring each node in the route to have the application key.

5.8.3.2. Packet Encryption

The encryption policy MUST support either encryption of the payload only or of the entire packet. Payload-only encryption would eliminate the decryption/re-encryption overhead at every hop, providing more real-time performance.

5.8.4. Disparate Security Policies

Due to the limited resources of an LLN, the security policy defined within the LLN MUST be able to differ from that of the rest of the IP network within the facility, yet packets MUST still be able to route to or through the LLN from/to these networks.

5.8.5. Routing Security Policies to Sleeping Devices

The routing protocol MUST gracefully handle routing temporal security updates (e.g., dynamic keys) to sleeping devices on their "awake" cycle to assure that sleeping devices can readily and efficiently access the network.

6. Security Considerations

The requirements placed on the LLN routing protocol in order to provide the correct level of security support are presented in Section 5.8.

LLNs deployed in a building environment may be entirely isolated from other networks, attached to normal IP networks within the building yet physically disjoint from the wider Internet, or connected either directly or through other IP networks to the Internet. Additionally, even where no wired connectivity exists outside of the building, the use of wireless infrastructure within the building means that physical connectivity to the LLN is possible for an attacker.

Therefore, it is important that any routing protocol solution designed to meet the requirements included in this document addresses the security features requirements described in Section 5.8. Implementations of these protocols will be required in the protocol specifications to provide the level of support indicated in Section 5.8, and will be encouraged to make the support flexibly configurable to enable an operator to make a judgment of the level of security that they want to deploy at any time.

As noted in Section 5.8, use/deployment of the different security features is intended to be optional. This means that, although the protocols developed must conform to the requirements specified, the operator is free to determine the level of risk and the trade-offs

against performance. An implementation must not make those choices on behalf of the operator by avoiding implementing any mandatory-to-implement security features.

This informational requirements specification introduces no new security concerns.

7. Acknowledgments

In addition to the authors, JP. Vasseur, David Culler, Ted Humpal, and Zach Shelby are gratefully acknowledged for their contributions to this document.

8. References

8.1. Normative References

[RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, March 1997.

8.2. Informative References

[ROLL-TERM] Vasseur, JP., "Terminology in Low power And Lossy Networks", Work in Progress, March 2010.

Appendix A. Additional Building Requirements

Appendix A contains additional building requirements that were deemed out of scope for ROLL, yet provided ancillary substance for the reader.

A.1. Additional Commercial Product Requirements

A.1.1. Wired and Wireless Implementations

Vendors will likely not develop a separate product line for both wired and wireless networks. Hence, the solutions set forth must support both wired and wireless implementations.

A.1.2. World-Wide Applicability

Wireless devices must be supportable unlicensed bands.

A.2. Additional Installation and Commissioning Requirements

A.2.1. Unavailability of an IP Network

Product commissioning must be performed by an application engineer prior to the installation of the IP network (e.g., switches, routers, DHCP, DNS).

A.3. Additional Network Requirements

A.3.1. TCP/UDP

Connection-based and connectionless services must be supported.

A.3.2. Interference Mitigation

The network must automatically detect interference and seamlessly switch the channel to improve communication. Channel changes, and the nodes' responses to a given channel change, must occur within 60 seconds.

A.3.3. Packet Reliability

In building automation, it is required that the network meet the following minimum criteria:

<1% MAC-layer errors on all messages, after no more than three retries;

<0.1% network-layer errors on all messages, after no more than three additional retries;

<0.01% application-layer errors on all messages.

Therefore, application-layer messages will fail no more than once every 100,000 messages.

A.3.4. Merging Commissioned Islands

Subsystems are commissioned by various vendors at various times during building construction. These subnetworks must seamlessly merge into networks and networks must seamlessly merge into internetworks since the end user wants a holistic view of the system.

A.3.5. Adjustable Routing Table Sizes

The routing protocol must allow constrained nodes to hold an abbreviated set of routes. That is, the protocol should not mandate that the node routing tables be exhaustive.

A.3.6. Automatic Gain Control

For wireless implementations, the device radios should incorporate automatic transmit power regulation to maximize packet transfer and minimize network interference, regardless of network size or density.

A.3.7. Device and Network Integrity

Commercial-building devices must all be periodically scanned to assure that each device is viable and can communicate data and alarm information as needed. Routers should maintain previous packet flow information temporally to minimize overall network overhead.

A.4. Additional Performance Requirements

A.4.1. Data Rate Performance

An effective data rate of 20 kbit/s is the lowest acceptable operational data rate on the network.

A.4.2. Firmware Upgrades

To support high-speed code downloads, routing should support transports that provide parallel downloads to targeted devices, yet guarantee packet delivery. In cases where the spatial position of the devices requires multiple hops, the algorithm should recurse through the network until all targeted devices have been serviced. Devices receiving a download may cease normal operation, but upon completion of the download must automatically resume normal operation.

A.4.3. Route Persistence

To eliminate high network traffic in power-fail or brown-out conditions, previously established routes should be remembered and invoked prior to establishing new routes for those devices re-entering the network.

Authors' Addresses

Jerry Martocci
Johnson Controls Inc.
507 E. Michigan Street
Milwaukee, WI 53202
USA
Phone: +1 414 524 4010
EMail: jerald.p.martocci@jci.com

Pieter De Mil
Ghent University - IBCN
G. Crommenlaan 8 bus 201
Ghent 9050
Belgium
Phone: +32 9331 4981
Fax: +32 9331 4899
EMail: pieter.demil@intec.ugent.be

Nicolas Riou
Schneider Electric
Technopole 38TEC T3
37 quai Paul Louis Merlin
38050 Grenoble Cedex 9
France
Phone: +33 4 76 57 66 15
EMail: nicolas.riou@fr.schneider-electric.com

Wouter Vermeylen
Arts Centre Vooruit
Ghent 9000
Belgium
EMail: wouter@vooruit.be