A Recommendation for IPv6 Address Text Representation

Abstract

   As IPv6 deployment increases, there will be a dramatic increase in
   the need to use IPv6 addresses in text. While the IPv6 address
   architecture in Section 2.2 of RFC 4291 describes a flexible model
   for text representation of an IPv6 address, this flexibility has been
   causing problems for operators, system engineers, and users. This
   document defines a canonical textual representation format. It does
   not define a format for internal storage, such as within an
   application or database. It is expected that the canonical format
   will be followed by humans and systems when representing IPv6
   addresses as text, but all implementations must accept and be able to
   handle any legitimate RFC 4291 format.

Status of This Memo

   This is an Internet Standards Track document.

   This document is a product of the Internet Engineering Task Force
   (IETF). It represents the consensus of the IETF community. It has
   received public review and has been approved for publication by the
   Internet Engineering Steering Group (IESG). Further information on
   Internet Standards is available in Section 2 of RFC 5741.

   Information about the current status of this document, any errata,
   and how to provide feedback on it may be obtained at
   http://www.rfc-editor.org/info/rfc5952.

Copyright Notice

Table of Contents

1. Introduction

   A single IPv6 address can be text represented in many ways. Examples
   are shown below.

      2001:db8:0:0:1:0:0:1

      2001:0db8:0:0:1:0:0:1

      2001:db8::1:0:0:1

      2001:db8::0:1:0:0:1

      2001:0db8::1:0:0:1

      2001:db8:0:0:1::1

      2001:db8:0000:0:1::1

      2001:DB8:0:0:1::1

   All of the above examples represent the same IPv6 address. This
   flexibility has caused many problems for operators, systems
   engineers, and customers. The problems are noted in Section 3. A
   canonical representation format to avoid problems is introduced in
   Section 4.

1.1. Requirements Language

   The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT",
   "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this
   document are to be interpreted as described in [RFC2119].

2. Text Representation Flexibility of RFC 4291

   Examples of flexibility in Section 2.2 of [RFC4291] are described
   below.

2.1. Leading Zeros in a 16-Bit Field

      'It is not necessary to write the leading zeros in an individual
      field.'

   Conversely, it is also not necessary to omit leading zeros. This
   means that it is possible to select from representations such as
   those in the following example. The final 16-bit field is different,
   but all of these addresses represent the same address.

        2001:db8:aaaa:bbbb:cccc:dddd:eeee:0001

        2001:db8:aaaa:bbbb:cccc:dddd:eeee:001

        2001:db8:aaaa:bbbb:cccc:dddd:eeee:01

        2001:db8:aaaa:bbbb:cccc:dddd:eeee:1

2.2. Zero Compression

        'A special syntax is available to compress the zeros. The use of
        "::" indicates one or more groups of 16 bits of zeros.'

    It is possible to select whether or not to omit just one 16-bit 0
    field.

        2001:db8:aaaa:bbbb:cccc:dddd::1

        2001:db8:aaaa:bbbb:cccc:dddd:0:1

    In cases where there is more than one field of only zeros, there is a
    choice of how many fields can be shortened.

        2001:db8:0:0:0::1

        2001:db8:0:0::1

        2001:db8:0::1

        2001:db8::1

    In addition, Section 2.2 of [RFC4291] notes,

        'The "::" can only appear once in an address.'

    This gives a choice on where in a single address to compress the
    zero.

        2001:db8::aaaa:0:0:1

        2001:db8:0:0:aaaa::1

2.3. Uppercase or Lowercase

   [RFC4291] does not mention any preference of uppercase or lowercase.

      2001:db8:aaaa:bbbb:cccc:dddd:eeee:aaaa

      2001:db8:aaaa:bbbb:cccc:dddd:eeee:AAAA

      2001:db8:aaaa:bbbb:cccc:dddd:eeee:AaAa

3. Problems Encountered with the Flexible Model

3.1. Searching

3.1.1. General Summary

   A search of an IPv6 address if conducted through a UNIX system is
   usually case sensitive and extended options that allow for regular
   expression use will come in handy. However, there are many
   applications in the Internet today that do not provide this
   capability. When searching for an IPv6 address in such systems, the
   system engineer will have to try each and every possibility to search
   for an address. This has critical impacts, especially when trying to
   deploy IPv6 over an enterprise network.

3.1.2. Searching Spreadsheets and Text Files

   Spreadsheet applications and text editors on GUI systems rarely have
   the ability to search for text using regular expression. Moreover,
   there are many non-engineers (who are not aware of case sensitivity
   and regular expression use) that use these applications to manage IP
   addresses. This has worked quite well with IPv4 since text
   representation in IPv4 has very little flexibility. There is no
   incentive to encourage these non-engineers to change their tool or
   learn regular expression when they decide to go dual-stack. If the
   entry in the spreadsheet reads, 2001:db8::1:0:0:1, but the search was
   conducted as 2001:db8:0:0:1::1, this will show a result of no match.
   One example where this will cause a problem is, when the search is
   being conducted to assign a new address from a pool, and a check is
   being done to see if it is not in use. This may cause problems for
   the end-hosts or end-users. This type of address management is very
   often seen in enterprise networks and ISPs.

3.1.3. Searching with Whois

   The "whois" utility is used by a wide range of people today. When a
   record is set to a database, one will likely check the output to see
   if the entry is correct. If an entity was recorded as 2001:db8::/48,

but the whois output showed 2001:0db8:0000::/48, most non-engineers
would think that their input was wrong and will likely retry several
times or make a frustrated call to the database hostmaster. If there
was a need to register the same prefix on different systems, and each
system showed a different text representation, this would confuse
people even more. Although this document focuses on addresses rather
than prefixes, it is worth mentioning the prefix problems because the
problems encountered with addresses and prefixes are mostly equal.

3.1.4. Searching for an Address in a Network Diagram

Network diagrams and blueprints often show what IP addresses are
assigned to a system devices. In times of trouble shooting there may
be a need to search through a diagram to find the point of failure
(for example, if a traceroute stopped at 2001:db8::1, one would
search the diagram for that address). This is a technique quite
often in use in enterprise networks and managed services. Again, the
different flavors of text representation will result in a time-
consuming search leading to longer mean times to restoration (MTTR)
in times of trouble.

3.2. Parsing and Modifying

3.2.1. General Summary

With all the possible methods of text representation, each
application must include a module, object, link, etc. to a function
that will parse IPv6 addresses in a manner such that no matter how it
is represented, they will mean the same address. Many system
engineers who integrate complex computer systems for corporate
customers will have difficulties finding that their favorite tool
will not have this function, or will encounter difficulties such as
having to rewrite their macros or scripts for their customers.

3.2.2. Logging

If an application were to output a log summary that represented the
address in full (such as 2001:0db8:0000:0000:1111:2222:3333:4444),
the output would be highly unreadable compared to the IPv4 output.
The address would have to be parsed and reformed to make it useful
for human reading. Sometimes logging for critical systems is done by
mirroring the same traffic to two different systems. Care must be
taken so that no matter what the log output is, the logs should be
parsed so they are equivalent.

3.2.3. Auditing: Case 1

   When a router or any other network appliance machine configuration is
   audited, there are many methods to compare the configuration
   information of a node. Sometimes auditing will be done by just
   comparing the changes made each day. In this case, if configuration
   was done such that 2001:db8::1 was changed to 2001:0db8:0000:0000:
   0000:0000:0000:0001 just because the new engineer on the block felt
   it was better, a simple diff will show that a different address was
   configured. If this was done on a wide scale network, people will be
   focusing on 'why the extra zeros were put in' instead of doing any
   real auditing. Lots of tools are just plain diffs that do not take
   into account address representation rules.

3.2.4. Auditing: Case 2

   Node configurations will be matched against an information system
   that manages IP addresses. If output notation is different, there
   will need to be a script that is implemented to cover for this. The
   result of an SNMP GET operation, converted to text and compared to a
   textual address written by a human is highly unlikely to match on the
   first try.

3.2.5. Verification

   Some protocols require certain data fields to be verified. One
   example of this is X.509 certificates. If an IPv6 address field in a
   certificate was incorrectly verified by converting it to text and
   making a simple textual comparison to some other address, the
   certificate may be mistakenly shown as being invalid due to a
   difference in text representation methods.

3.2.6. Unexpected Modifying

   Sometimes, a system will take an address and modify it as a
   convenience. For example, a system may take an input of
   2001:0db8:0::1 and make the output 2001:db8::1. If the zeros were
   input for a reason, the outcome may be somewhat unexpected.

3.3. Operating

3.3.1. General Summary

   When an operator sets an IPv6 address of a system as 2001:db8:0:0:1:
   0:0:1, the system may take the address and show the configuration
   result as 2001:DB8::1:0:0:1. Someone familiar with IPv6 address
   representation will know that the right address is set, but not
   everyone may understand this.

3.3.2. Customer Calls

When a customer calls to inquire about a suspected outage, IPv6
address representation should be handled with care. Not all
customers are engineers, nor do they have a similar skill level in
IPv6 technology. The network operations center will have to take
extra steps to humanly parse the address to avoid having to explain
to the customers that 2001:db8:0:1::1 is the same as
2001:db8::1:0:0:0:1. This is one thing that will never happen in
IPv4 because IPv4 addresses cannot be abbreviated.

3.3.3. Abuse

Network abuse reports generally include the abusing IP address. This
'reporting' could take any shape or form of the flexible model. A
team that handles network abuse must be able to tell the difference
between a 2001:db8::1:0:1 and 2001:db8:1::0:1. Mistakes in the
placement of the "::" will result in a critical situation. A system
that handles these incidents should be able to handle any type of
input and parse it in a correct manner. Also, incidents are reported
over the phone. It is unnecessary to report if the letter is
uppercase or lowercase. However, when a letter is spelled uppercase,
people tend to specify that it is uppercase, which is unnecessary
information.

3.4. Other Minor Problems

3.4.1. Changing Platforms

When an engineer decides to change the platform of a running service,
the same code may not work as expected due to the difference in IPv6
address text representation. Usually, a change in a platform (e.g.,
Unix to Windows, Cisco to Juniper) will result in a major change of
code anyway, but flexibility in address representation will increase
the work load.

3.4.2. Preference in Documentation

A document that is edited by more than one author may become harder
to read.

3.4.3. Legibility

Capital case D and 0 can be quite often misread. Capital B and 8 can
also be misread.

4. A Recommendation for IPv6 Text Representation

   A recommendation for a canonical text representation format of IPv6
   addresses is presented in this section. The recommendation in this
   document is one that complies fully with [RFC4291], is implemented by
   various operating systems, and is human friendly. The recommendation
   in this section SHOULD be followed by systems when generating an
   address to be represented as text, but all implementations MUST
   accept and be able to handle any legitimate [RFC4291] format. It is
   advised that humans also follow these recommendations when spelling
   an address.

4.1. Handling Leading Zeros in a 16-Bit Field

   Leading zeros MUST be suppressed. For example, 2001:0db8::0001 is
   not acceptable and must be represented as 2001:db8::1. A single 16-
   bit 0000 field MUST be represented as 0.

4.2. "::" Usage

4.2.1. Shorten as Much as Possible

   The use of the symbol "::" MUST be used to its maximum capability.
   For example, 2001:db8:0:0:0:0:2:1 must be shortened to 2001:db8::2:1.
   Likewise, 2001:db8::0:1 is not acceptable, because the symbol "::"
   could have been used to produce a shorter representation 2001:db8::1.

4.2.2. Handling One 16-Bit 0 Field

   The symbol "::" MUST NOT be used to shorten just one 16-bit 0 field.
   For example, the representation 2001:db8:0:1:1:1:1:1 is correct, but
   2001:db8::1:1:1:1:1 is not correct.

4.2.3. Choice in Placement of "::"

   When there is an alternative choice in the placement of a "::", the
   longest run of consecutive 16-bit 0 fields MUST be shortened (i.e.,
   the sequence with three consecutive zero fields is shortened in 2001:
   0:0:1:0:0:0:1). When the length of the consecutive 16-bit 0 fields
   are equal (i.e., 2001:db8:0:0:1:0:0:1), the first sequence of zero
   bits MUST be shortened. For example, 2001:db8::1:0:0:1 is correct
   representation.

4.3. Lowercase

   The characters "a", "b", "c", "d", "e", and "f" in an IPv6 address
   MUST be represented in lowercase.

5. Text Representation of Special Addresses

   Addresses such as IPv4-Mapped IPv6 addresses, ISATAP [RFC5214], and
   IPv4-translatable addresses [ADDR-FORMAT] have IPv4 addresses
   embedded in the low-order 32 bits of the address. These addresses
   have a special representation that may mix hexadecimal and dot
   decimal notations. The decimal notation may be used only for the
   last 32 bits of the address. For these addresses, mixed notation is
   RECOMMENDED if the following condition is met: the address can be
   distinguished as having IPv4 addresses embedded in the lower 32 bits
   solely from the address field through the use of a well-known prefix.
   Such prefixes are defined in [RFC4291] and [RFC2765] at the time of
   this writing. If it is known by some external method that a given
   prefix is used to embed IPv4, it MAY be represented as mixed
   notation. Tools that provide options to specify prefixes that are
   (or are not) to be represented as mixed notation may be useful.

   There is a trade-off here where a recommendation to achieve an exact
   match in a search (no dot decimals whatsoever) and a recommendation
   to help the readability of an address (dot decimal whenever possible)
   does not result in the same solution. The above recommendation is
   aimed at fixing the representation as much as possible while leaving
   the opportunity for future well-known prefixes to be represented in a
   human-friendly manner as tools adjust to newly assigned prefixes.

   The text representation method noted in Section 4 should be applied
   for the leading hexadecimal part (i.e., ::ffff:192.0.2.1 instead of
   0:0:0:0:0:ffff:192.0.2.1).

6. Notes on Combining IPv6 Addresses with Port Numbers

   There are many different ways to combine IPv6 addresses and port
   numbers that are represented in text. Examples are shown below.

   o [2001:db8::1]:80

   o 2001:db8::1:80

   o 2001:db8::1.80

   o 2001:db8::1 port 80

   o 2001:db8::1p80

   o 2001:db8::1#80

   The situation is not much different in IPv4, but the most ambiguous
   case with IPv6 is the second bullet. This is due to the "::"usage in

IPv6 addresses. This style is NOT RECOMMENDED because of its
ambiguity. The [] style as expressed in [RFC3986] SHOULD be
employed, and is the default unless otherwise specified. Other
styles are acceptable when there is exactly one style for the given
context and cross-platform portability does not become an issue. For
URIs containing IPv6 address literals, [RFC3986] MUST be followed, as
well as the rules defined in this document.

7. Prefix Representation

Problems with prefixes are the same as problems encountered with
addresses. The text representation method of IPv6 prefixes should be
no different from that of IPv6 addresses.

8. Security Considerations

This document notes some examples where IPv6 addresses are compared
in text format. The example on Section 3.2.5 is one that may cause a
security risk if used for access control. The common practice of
comparing X.509 data is done in binary format.

9. Acknowledgements

The authors would like to thank Jan Zorz, Randy Bush, Yuichi Minami,
and Toshimitsu Matsuura for their generous and helpful comments in
kick starting this document. We also would like to thank Brian
Carpenter, Akira Kato, Juergen Schoenwaelder, Antonio Querubin, Dave
Thaler, Brian Haley, Suresh Krishnan, Jerry Huang, Roman Donchenko,
Heikki Vatiainen, Dan Wing, and Doug Barton for their input. Also, a
very special thanks to Ron Bonica, Fred Baker, Brian Haberman, Robert
Hinden, Jari Arkko, and Kurt Lindqvist for their support in bringing
this document to light in IETF working groups.

10. References

10.1. Normative References

[RFC2119]   Bradner, S., "Key words for use in RFCs to Indicate
               Requirement Levels", BCP 14, RFC 2119, March 1997.

[RFC2765]   Nordmark, E., "Stateless IP/ICMP Translation Algorithm
               (SIIT)", RFC 2765, February 2000.

[RFC3986]   Berners-Lee, T., Fielding, R., and L. Masinter,
               "Uniform Resource Identifier (URI): Generic Syntax",
               STD 66, RFC 3986, January 2005.

   [RFC4291]   Hinden, R. and S. Deering, "IP Version 6 Addressing
                   Architecture", RFC 4291, February 2006.

10.2. Informative References

   [ADDR-FORMAT]   Bao, C., "IPv6 Addressing of IPv4/IPv6 Translators",
                   Work in Progress, July 2010.

   [RFC4038]   Shin, M-K., Hong, Y-G., Hagino, J., Savola, P., and E.
                   Castro, "Application Aspects of IPv6 Transition",
                   RFC 4038, March 2005.

   [RFC5214]   Templin, F., Gleeson, T., and D. Thaler, "Intra-Site
                   Automatic Tunnel Addressing Protocol (ISATAP)",
                   RFC 5214, March 2008.

Appendix A. For Developers

   We recommend that developers use display routines that conform to
   these rules. For example, the usage of getnameinfo() with flags
   argument NI_NUMERICHOST in FreeBSD 7.0 will give a conforming output,
   except for the special addresses notes in Section 5. The function
   inet_ntop() of FreeBSD7.0 is a good C code reference, but should not
   be called directly. See [RFC4038] for details.

Authors' Addresses

   Seiichi Kawamura
   NEC BIGLOBE, Ltd.
   14-22, Shibaura 4-chome
   Minatoku, Tokyo 108-8558
   JAPAN

   Phone: +81 3 3798 6085
   EMail: kawamucho@mesh.ad.jp

   Masanobu Kawashima
   NEC AccessTechnica, Ltd.
   800, Shimomata
   Kakegawa-shi, Shizuoka 436-8501
   JAPAN

   Phone: +81 537 23 9655
   EMail: kawashimam@necat.nec.co.jp