

Internet Engineering Task Force (IETF)
Request for Comments: 6016
Category: Standards Track
ISSN: 2070-1721

B. Davie
F. Le Faucheur
A. Narayanan
Cisco Systems, Inc.
October 2010

Support for the Resource Reservation Protocol (RSVP) in Layer 3 VPNs

Abstract

RFC 4364 and RFC 4659 define an approach to building provider-provisioned Layer 3 VPNs (L3VPNs) for IPv4 and IPv6. It may be desirable to use Resource Reservation Protocol (RSVP) to perform admission control on the links between Customer Edge (CE) routers and Provider Edge (PE) routers. This document specifies procedures by which RSVP messages traveling from CE to CE across an L3VPN may be appropriately handled by PE routers so that admission control can be performed on PE-CE links. Optionally, admission control across the provider's backbone may also be supported.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 5741.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <http://www.rfc-editor.org/info/rfc6016>.

Copyright Notice

Copyright (c) 2010 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<http://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

This document may contain material from IETF Documents or IETF Contributions published or made publicly available before November 10, 2008. The person(s) controlling the copyright in some of this material may not have granted the IETF Trust the right to allow modifications of such material outside the IETF Standards Process. Without obtaining an adequate license from the person(s) controlling the copyright in such materials, this document may not be modified outside the IETF Standards Process, and derivative works of it may not be created outside the IETF Standards Process, except to format it for publication as an RFC or to translate it into languages other than English.

Table of Contents

1. Introduction	4
1.1. Terminology	5
1.2. Requirements Language	5
2. Problem Statement	5
2.1. Model of Operation	8
3. Admission Control on PE-CE Links	9
3.1. New Objects of Type VPN-IPv4	9
3.2. Path Message Processing at Ingress PE	11
3.3. Path Message Processing at Egress PE	12
3.4. Resv Processing at Egress PE	13
3.5. Resv Processing at Ingress PE	13
3.6. Other RSVP Messages	14
4. Admission Control in Provider's Backbone	14
5. Inter-AS Operation	15
5.1. Inter-AS Option A	15
5.2. Inter-AS Option B	15
5.2.1. Admission Control on ASBR	16
5.2.2. No Admission Control on ASBR	16
5.3. Inter-AS Option C	17
6. Operation with RSVP Disabled	17
7. Other RSVP Procedures	18
7.1. Refresh Overhead Reduction	18
7.2. Cryptographic Authentication	18
7.3. RSVP Aggregation	19
7.4. Support for CE-CE RSVP-TE	19
8. Object Definitions	20
8.1. VPN-IPv4 and VPN-IPv6 SESSION Objects	20
8.2. VPN-IPv4 and VPN-IPv6 SENDER_TEMPLATE Objects	21
8.3. VPN-IPv4 and VPN-IPv6 FILTER_SPEC Objects	22
8.4. VPN-IPv4 and VPN-IPv6 RSVP_HOP Objects	22
8.5. Aggregated VPN-IPv4 and VPN-IPv6 SESSION Objects	24
8.6. AGGREGATE-VPN-IPv4 and AGGREGATE-VPN-IPv6 SENDER_TEMPLATE Objects	26
8.7. AGGREGATE-VPN-IPv4 and AGGREGATE-VPN-IPv6 FILTER_SPEC Objects	27
9. IANA Considerations	28
10. Security Considerations	30
11. Acknowledgments	33
Appendix A. Alternatives Considered	34
A.1. GMPLS UNI Approach	34
A.2. Label Switching Approach	34
A.3. VRF Label Approach	34
A.4. VRF Label Plus VRF Address Approach	35
References	35
Normative References	35
Informative References	36

1. Introduction

[RFC4364] and [RFC4659] define a Layer 3 VPN service known as BGP/MPLS VPNs for IPv4 and for IPv6, respectively. [RFC2205] defines the Resource Reservation Protocol (RSVP), which may be used to perform admission control as part of the Integrated Services (Int-Serv) architecture [RFC1633][RFC2210].

Customers of a Layer 3 VPN service may run RSVP for the purposes of admission control (and associated resource reservation) in their own networks. Since the links between Provider Edge (PE) and Customer Edge (CE) routers in a Layer 3 VPN may often be resource constrained, it may be desirable to be able to perform admission control over those links. In order to perform admission control using RSVP in such an environment, it is necessary that RSVP control messages, such as Path messages and Resv messages, are appropriately handled by the PE routers. This presents a number of challenges in the context of BGP/MPLS VPNs:

- o RSVP Path message processing depends on routers recognizing the Router Alert Option ([RFC2113], [RFC2711]) in the IP header. However, packets traversing the backbone of a BGP/MPLS VPN are MPLS encapsulated, and thus the Router Alert Option may not be visible to the egress PE due to implementation or policy considerations (e.g., if the egress PE processes the message as "pop and go" without examining the IP header).
- o BGP/MPLS VPNs support non-unique addressing of customer networks. Thus, a PE at the ingress or egress of the provider backbone may be called upon to process Path messages from different customer VPNs with non-unique destination addresses within the RSVP message. Current mechanisms for identifying customer context from data packets are incompatible with RSVP message processing rules.
- o A PE at the ingress of the provider's backbone may receive Resv messages corresponding to different customer VPNs from other PEs, and needs to be able to associate those Resv messages with the appropriate customer VPNs.

Further discussion of these issues is presented in Section 2.

This document describes a set of procedures to overcome these challenges and thus to enable admission control using RSVP over the PE-CE links. We note that similar techniques may be applicable to other protocols used for admission control such as the combination of the NSIS Signaling Layer Protocol (NSLP) for Quality-of-Service (QoS) Signaling [RFC5974] and General Internet Signaling Transport (GIST) protocol [RFC5971].

Additionally, it may be desirable to perform admission control over the provider's backbone on behalf of one or more L3VPN customers. Core (P) routers in a BGP/MPLS VPN do not have forwarding entries for customer routes, and thus they cannot natively process RSVP messages for customer flows. Also, the core is a shared resource that carries traffic for many customers, so issues of resource allocation among customers and trust (or lack thereof) also ought to be addressed. This document specifies procedures for supporting such a scenario.

This document deals with establishing reservations for unicast flows only. Because the support of multicast traffic in BGP/MPLS VPNs is still evolving, and raises additional challenges for admission control, we leave the support of multicast flows for further study at this point.

1.1. Terminology

This document draws freely on the terminology defined in [RFC2205] and [RFC4364]. For convenience, we provide a few brief definitions here:

- o Customer Edge (CE) Router: Router at the edge of a customer site that attaches to the network of the VPN provider.
- o Provider Edge (PE) Router: Router at the edge of the service provider's network that attaches to one or more customer sites.
- o VPN Label: An MPLS label associated with a route to a customer prefix in a VPN (also called a VPN route label).
- o VPN Routing and Forwarding (VRF) Table: A PE typically has multiple VRFs, enabling it to be connected to CEs that are in different VPNs.

1.2. Requirements Language

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in RFC 2119 [RFC2119].

2. Problem Statement

The problem space of this document is the support of admission control between customer sites when the customer subscribes to a BGP/MPLS VPN. We subdivide the problem into (a) the problem of admission control on the PE-CE links (in both directions) and (b) the problem of admission control across the provider's backbone.

RSVP Path messages are normally addressed to the destination of a session, and contain the Router Alert Option (RAO) within the IP header. Routers along the path to the destination that are configured to process RSVP messages need to detect the presence of the RAO to allow them to intercept Path messages. However, the egress PEs of a network supporting BGP/MPLS VPNs receive packets destined for customer sites as MPLS-encapsulated packets, and they possibly forward those based only on examination of the MPLS label. In order to process RSVP Path messages, the egress VPN PE would have to pop the VPN label and examine the IP header underneath, before forwarding the packet (based on the VPN label disposition rules), which is not a requirement for data packet processing today. Hence, a Path message would be forwarded without examination of the IP options and would therefore not receive appropriate processing at the PE. Another potential issue is doing Connection Admission Control (CAC) at an Autonomous System Border Router (ASBR). Even an implementation that examines the IP header when removing the VPN label (e.g., PE-CE link) would not be able to do CAC at an Option-B ASBR; that requires examining the (interior) IP header while doing a label swap, which is much less desirable behavior.

In general, there are significant issues with requiring support for IP Router Alert outside of a controlled, "walled-garden" network, as described in [ALERT-USAGE]. The case of a MPLS L3VPN falls under the "Overlay Model" described therein. Fundamental to this model is that providers would seek to eliminate the requirement to process RAO-marked packets from customers, on any routers except the PEs facing those customers. Issues with requiring interior MPLS routers to process RAO-marked packets are also described in [LER-OPTIONS]. The approach for RSVP packet handling described in this document has the advantage of being independent of any data-plane requirements such as IP Router Alert support within the VPN or examining any IP options for MPLS-encapsulated packets. The only requirement for processing IP Router Alert packets is for RSVP packets received from the CE, which do not carry any MPLS encapsulation.

For the PE-CE link subproblem, the most basic challenge is that RSVP control messages contain IP addresses that are drawn from the customer's address space, and PEs need to deal with traffic from many customers who may have non-unique (or overlapping) address spaces. Thus, it is essential that a PE be able, in all cases, to identify the correct VPN context in which to process an RSVP control message. The current mechanism for identifying the customer context is the VPN label, which is carried in an MPLS header outside of the RSVP message. This is divergent from the general RSVP model of session identification ([RFC2205], [RFC2209]), which relies solely on RSVP objects to identify sessions. Further, it is incompatible with protocols like COPS/RSVP (Common Open Policy Service) ([RFC2748],

[RFC2749]), which replace the IP encapsulation of the RSVP message and send only RSVP objects to a COPS server. We believe it is important to retain the model of completely identifying an RSVP session from the contents of RSVP objects. Much of this document deals with this issue.

For the case of making reservations across the provider backbone, we observe that BGP/MPLS VPNs do not create any per-customer forwarding state in the P (provider core) routers. Thus, in order to make reservations on behalf of customer-specified flows, it is clearly necessary to make some sort of aggregated reservation from PE-PE and then map individual, customer-specific reservations onto an aggregate reservation. That is similar to the problem tackled in [RFC3175] and [RFC4804], with the additional complications of handling customer-specific addressing associated with BGP/MPLS VPNs.

Consider the case where an MPLS VPN customer uses RSVP signaling across his sites for resource reservation and admission control. Let's further assume that, initially, RSVP is not processed through the MPLS VPN cloud (i.e., RSVP messages from the sender to the receiver travel transparently from CE to CE). In that case, RSVP allows the establishment of resource reservations and admission control on a subset of the flow path (from sender to ingress CE, and from the RSVP router downstream of the egress CE to the receiver). If admission control is then activated on any of the CE-PE link, the provider's backbone, or PE-CE link (as allowed by the present document), the customer will benefit from an extended coverage of admission control and resource reservation: the resource reservation will now span over a bigger subset of (and possibly the whole) flow path, which in turn will increase the QoS granted to the corresponding flow. Specific flows whose reservation is successful through admission control on the newly enabled segments will indeed benefit from this quality of service enhancement. However, it must be noted that, in case there are not enough resources on one (or more) of the newly enabled segments (e.g., say admission control is enabled on a given PE-->CE link and there is not enough capacity on that link to admit all reservations for all the flows traversing that link), then some flows will not be able to maintain, or establish, their reservation. While this may appear undesirable for these flows, we observe that this only occurs if there is indeed a lack of capacity on a segment, and that in the absence of admission control, all flows would be established but would all suffer from the resulting congestion on the bottleneck segment. We also observe that, in the case of such a lack of capacity, admission control allows enforcement of controlled and flexible policies (so that, for example, more important flows can be granted higher priority at

reserving resources). We note also that flows are given a chance to establish smaller reservations so that the aggregate load can adapt dynamically to the bottleneck capacity.

2.1. Model of Operation

Figure 1 illustrates the basic model of operation with which this document is concerned.

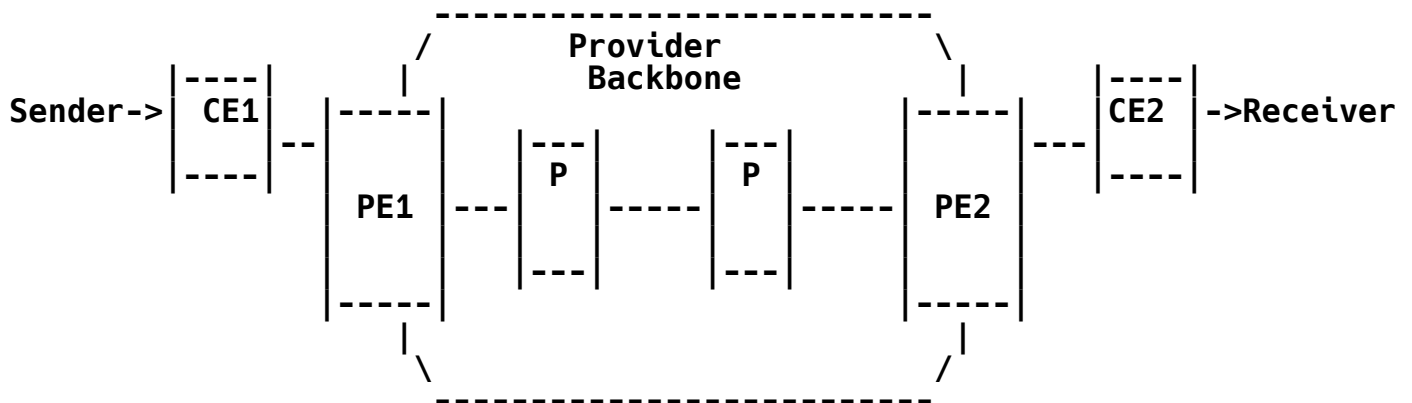


Figure 1. Model of Operation for RSVP-Based Admission Control over MPLS/BGP VPN

To establish a unidirectional reservation for a point-to-point flow from Sender to Receiver that takes account of resource availability on the CE-PE and PE-CE links only, the following steps need to take place:

1. The Sender sends a Path message to an IP address of the Receiver.
2. The Path message is processed by CE1 using normal RSVP procedures and forwarded towards the Receiver along the link CE1-PE1.
3. PE1 processes the Path message and forwards it towards the Receiver across the provider backbone.
4. PE2 processes the Path message and forwards it towards the Receiver along link PE2-CE2.
5. CE2 processes the Path message using normal RSVP procedures and forwards it towards the Receiver.
6. The Receiver sends a Resv message to CE2.

7. CE2 sends the Resv message to PE2.
8. PE2 processes the Resv message (including performing admission control on link PE2-CE2) and sends the Resv message to PE1.
9. PE1 processes the Resv message and sends the Resv message to CE1.
10. CE1 processes the Resv message using normal RSVP procedures, performs admission control on the link CE1-PE1, and sends the Resv message to the Sender if successful.

In each of the steps involving Resv messages (6 through 10) the node sending the Resv message uses the previously established Path state to determine the "RSVP Previous Hop (PHOP)" and sends a Resv message to that address. We note that establishing that Path state correctly at PEs is one of the challenges posed by the BGP/MPLS environment.

3. Admission Control on PE-CE Links

In the following sections, we trace through the steps outlined in Section 2.1 and expand on the details for those steps where standard RSVP procedures need to be extended or modified to support the BGP/MPLS VPN environment. For all the remaining steps described in the preceding section, standard RSVP processing rules apply.

All the procedures described below support both IPv4 and IPv6 addressing. In all cases where IPv4 is referenced, IPv6 can be substituted with identical procedures and results. Object definitions for both IPv4 and IPv6 are provided in Section 8.

3.1. New Objects of Type VPN-IPv4

For RSVP signaling within a VPN, certain RSVP objects need to be extended. Since customer IP addresses need not be unique, the current types of SESSION, SENDER_TEMPLATE, and FILTERSPEC objects are no longer sufficient to globally identify RSVP states in P/PE routers, since they are currently based on IP addresses. We propose new types of SESSION, SENDER_TEMPLATE, and FILTERSPEC objects, which contain globally unique VPN-IPv4 format addresses. The ingress and egress PE nodes translate between the regular IPv4 addresses for messages to and from the CE, and VPN-IPv4 addresses for messages to and from PE routers. The rules for this translation are described in later sections.

The RSVP_HOP object in an RSVP message currently specifies an IP address to be used by the neighboring RSVP hop to reply to the message sender. However, MPLS VPN PE routers (especially those separated by Option-B ASBRs) are not required to have direct IP reachability to each other. To solve this issue, we propose the use of label switching to forward RSVP messages between nodes within an MPLS VPN. This is achieved by defining a new VPN-IPv4 RSVP_HOP object. Use of the VPN-IPv4 RSVP_HOP object enables any two adjacent RSVP hops in an MPLS VPN (e.g., a PE in Autonomous System (AS) 1 and a PE in AS2) to correctly identify each other and send RSVP messages directly to each other.

The VPN-IPv4 RSVP_HOP object carries the IPv4 address of the message sender and a Logical Interface Handle (LIH) as before, but in addition carries a VPN-IPv4 address that also represents the sender of the message. The message sender MUST also advertise this VPN-IPv4 address into BGP, associated with a locally allocated label, and this advertisement MUST be propagated by BGP throughout the VPN and to adjacent ASes if required to provide reachability to this PE. Frames received by the PE marked with this label MUST be given to the local control plane for processing. When a neighboring RSVP hop wishes to reply to a message carrying a VPN-IPv4 RSVP_HOP, it looks for a BGP advertisement of the VPN-IPv4 address contained in that RSVP_HOP. If this address is found and carries an associated label, the neighboring RSVP node MUST encapsulate the RSVP message with this label and send it via MPLS encapsulation to the BGP next hop associated with the route. The destination IP address of the message is taken from the IP address field of the RSVP_HOP object, as described in [RFC2205]. Additionally, the IPv4 address in the RSVP_HOP object continues to be used for all other existing purposes, including neighbor matching between Path/Resv and SRefresh messages [RFC2961], authentication [RFC2747], etc.

The VPN-IPv4 address used in the VPN-IPv4 RSVP_HOP object MAY represent an existing address in the VRF that corresponds to the flow (e.g., a local loopback or PE-CE link address within the VRF for this customer), or it MAY be created specially for this purpose. In the case where the address is specially created for RSVP signaling (and possibly other control protocols), the BGP advertisement MUST NOT be redistributed to, or reachable by, any CEs outside the MPLS VPN. One way to achieve this is by creating a special "control protocols VPN" with VRF state on every PE/ASBR, carrying route targets not imported into customer VRFs. In the case where a customer VRF address is used as the VPN-IPv4 address, a VPN-IPv4 address in one customer VRF MUST NOT be used to signal RSVP messages for a flow in a different VRF.

If a PE/ASBR is sending a Path message to another PE/ASBR within the VPN, and it has any appropriate VPN-IPv4 address for signaling that satisfies the requirements outlined above, it **MUST** use a VPN-IPv4 RSVP_HOP object with this address for all RSVP messages within the VPN. If a PE/ASBR does not have any appropriate VPN-IPv4 address to use for signaling, it **MAY** send the Path message with a regular IPv4 RSVP_HOP object. In this case, the reply will be IP encapsulated. This option is not preferred because there is no guarantee that the neighboring RSVP hop has IP reachability to the sending node. If a PE/ASBR receives or originates a Path message with a VPN-IPv4 RSVP_HOP object, any RSVP_HOP object in corresponding upstream messages for this flow (e.g., Resv, ResvTear) or downstream messages (e.g., ResvError, PathTear) sent by this node within the VPN **MUST** be a VPN-IPv4 RSVP_HOP.

3.2. Path Message Processing at Ingress PE

When a Path message arrives at the ingress PE (step 3 of Section 2.1) the PE needs to establish suitable Path state and forward the Path message on to the egress PE. In the following paragraphs, we described the steps taken by the ingress PE.

The Path message is addressed to the eventual destination (the receiver at the remote customer site) and carries the IP Router Alert Option, in accordance with [RFC2205]. The ingress PE **MUST** recognize the Router Alert Option, intercept these messages and process them as RSVP signaling messages.

As noted above, there is an issue in recognizing Path messages as they arrive at the egress PE (PE2 in Figure 1). The approach defined here is to address the Path messages sent by the ingress PE directly to the egress PE, and send it without the IP Router Alert Option; that is, rather than using the ultimate receiver's destination address as the destination address of the Path message, we use the loopback address of the egress PE as the destination address of the Path message. This approach has the advantage that it does not require any new data-plane capabilities for the egress PE beyond those of a standard BGP/MPLS VPN PE. Details of the processing of this message at the egress PE are described below in Section 3.3. The approach of addressing a Path message directly to an RSVP next hop (that may or may not be the next IP hop) is already used in other environments such as those of [RFC4206] and [RFC4804].

The details of operation at the ingress PE are as follows. When the ingress PE (PE1 in Figure 1) receives a Path message from CE1 that is addressed to the receiver, the VRF that is associated with the incoming interface is identified, just as for normal data path operations. The Path state for the session is stored, and is

associated with that VRF, so that potentially overlapping addresses among different VPNs do not appear to belong to the same session. The destination address of the receiver is looked up in the appropriate VRF, and the BGP next hop for that destination is identified. That next hop is the egress PE (PE2 in Figure 1). A new VPN-IPv4 SESSION object is constructed, containing the Route Distinguisher (RD) that is part of the VPN-IPv4 route prefix for this destination, and the IPv4 address from the SESSION. In addition, a new VPN-IPv4 SENDER_TEMPLATE object is constructed, with the original IPv4 address from the incoming SENDER_TEMPLATE plus the RD that is used by this PE to advertise that prefix for this customer into the VPN. A new Path message is constructed with a destination address equal to the address of the egress PE identified above. This new Path message will contain all the objects from the original Path message, replacing the original SESSION and SENDER_TEMPLATE objects with the new VPN-IPv4 type objects. The Path message is sent without the Router Alert Option and contains an RSVP_HOP object constructed as specified in Section 3.1.

3.3. Path Message Processing at Egress PE

When a Path message arrives at the egress PE, (step 4 of Section 2.1) it is addressed to the PE itself, and is handed to RSVP for processing. The router extracts the RD and IPv4 address from the VPN-IPv4 SESSION object, and determines the local VRF context by finding a matching VPN-IPv4 prefix with the specified RD that has been advertised by this router into BGP. The entire incoming RSVP message, including the VRF information, is stored as part of the Path state.

Now the RSVP module can construct a Path message that differs from the Path it received in the following ways:

- a. Its destination address is the IP address extracted from the SESSION object;
- b. The SESSION and SENDER_TEMPLATE objects are converted back to IPv4-type by discarding the attached RD;
- c. The RSVP_HOP Object contains the IP address of the outgoing interface of the egress PE and a Logical Interface Handle (LIH), as per normal RSVP processing.

The router then sends the Path message on towards its destination over the interface identified above. This Path message carries the Router Alert Option as required by [RFC2205].

3.4. Resv Processing at Egress PE

When a receiver at the customer site originates a Resv message for the session, normal RSVP procedures apply until the Resv, making its way back towards the sender, arrives at the "egress" PE (step 8 of Section 2.1). Note that this is the "egress" PE with respect to the direction of data flow, i.e., PE2 in Figure 1. On arriving at PE2, the SESSION and FILTER_SPEC objects in the Resv, and the VRF in which the Resv was received, are used to find the matching Path state stored previously. At this stage, admission control can be performed on the PE-CE link.

Assuming admission control is successful, the PE constructs a Resv message to send to the RSVP previous hop stored in the Path state, i.e., the ingress PE (PE1 in Figure 1). The IPv4 SESSION object is replaced with the same VPN-IPv4 SESSION object received in the Path. The IPv4 FILTER_SPEC object is replaced with a VPN-IPv4 FILTER_SPEC object, which copies the VPN-IPv4 address from the SENDER_TEMPLATE received in the matching Path message. The RSVP_HOP in the Resv message MUST be constructed as specified in Section 3.1. The Resv message MUST be addressed to the IP address contained within the RSVP_HOP object in the Path message. If the Path message contained a VPN-IPv4 RSVP_HOP object, the Resv MUST be MPLS encapsulated using the label associated with that VPN-IPv4 address in BGP, as described in Section 3.1. If the Path message contained an IPv4 RSVP_HOP object, the Resv is simply IP encapsulated and addressed directly to the IP address in the RSVP_HOP object.

If admission control is not successful on the egress PE, a ResvError message is sent towards the receiver as per normal RSVP processing.

3.5. Resv Processing at Ingress PE

Upon receiving a Resv message at the ingress PE (step 8 of Section 2.1) with respect to data flow (i.e., PE1 in Figure 1), the PE determines the local VRF context and associated Path state for this Resv by decoding the received SESSION and FILTER_SPEC objects. It is now possible to generate a Resv message to send to the appropriate CE. The Resv message sent to the ingress CE will contain IPv4 SESSION and FILTER_SPEC objects, derived from the appropriate Path state. Since we assume, in this section, that admission control over the provider's backbone is not needed, the ingress PE does not perform any admission control for this reservation.

3.6. Other RSVP Messages

Processing of PathError, PathTear, ResvError, ResvTear, and ResvConf messages is generally straightforward and follows the rules of [RFC2205]. These additional rules MUST be observed for messages transmitted within the VPN (i.e., between the PEs):

- o The SESSION, SENDER_TEMPLATE, and FILTER_SPEC objects MUST be converted from IPv4 to VPN-IPv4 form and back in the same manner as described above for Path and Resv messages.
- o The appropriate type of RSVP_HOP object (VPN-IPv4 or IPv4) MUST be used as described above.
- o Depending on the type of RSVP_HOP object received from the neighbor, the message MUST be MPLS encapsulated or IP encapsulated as described above.
- o The matching state and VRF MUST be determined by decoding the RD and IPv4 addresses in the SESSION and FILTER_SPEC objects.
- o The message MUST be directly addressed to the appropriate PE, without using the Router Alert Option.

4. Admission Control in Provider's Backbone

The preceding section outlines how per-customer reservations can be made over the PE-CE links. This may be sufficient in many situations where the backbone is well engineered with ample capacity and there is no need to perform any sort of admission control in the backbone. However, in some cases where excess capacity cannot be relied upon (e.g., during failures or unanticipated periods of overload), it may be desirable to be able to perform admission control in the backbone on behalf of customer traffic.

Because of the fact that routes to customer addresses are not present in the P routers, along with the concerns of scalability that would arise if per-customer reservations were allowed in the P routers, it is clearly necessary to map the per-customer reservations described in the preceding section onto some sort of aggregate reservations. Furthermore, customer data packets need to be tunneled across the provider backbone just as in normal BGP/MPLS VPN operation.

Given these considerations, a feasible way to achieve the objective of admission control in the backbone is to use the ideas described in [RFC4804]. MPLS-TE tunnels can be established between PEs as a means to perform aggregate admission control in the backbone.

An MPLS-TE tunnel from an ingress PE to an egress PE can be thought of as a virtual link of a certain capacity. The main change to the procedures described above is that when a Resv is received at the ingress PE, an admission control decision can be performed by checking whether sufficient capacity of that virtual link remains available to admit the new customer reservation. We note also that [RFC4804] uses the IF_ID RSVP_HOP object to identify the tunnel across the backbone, rather than the simple RSVP_HOP object described in Section 3.2. The procedures of [RFC4804] should be followed here as well.

To achieve effective admission control in the backbone, there needs to be some way to separate the data-plane traffic that has a reservation from that which does not. We assume that packets that are subject to admission control on the core will be given a particular MPLS EXP value, and that no other packets will be allowed to enter the core with this value unless they have passed admission control. Some fraction of link resources will be allocated to queues on core links for packets bearing that EXP value, and the MPLS-TE tunnels will use that resource pool to make their constraint-based routing and admission control decisions. This is all consistent with the principles of aggregate RSVP reservations described in [RFC3175].

5. Inter-AS Operation

[RFC4364] defines three modes of inter-AS operation for MPLS/BGP VPNs, referred to as Options A, B, and C. In the following sections we describe how the scheme described above can operate in each inter-AS environment.

5.1. Inter-AS Option A

Operation of RSVP in Inter-AS Option A is quite straightforward. Each ASBR operates like a PE, and the ASBR-ASBR links can be viewed as PE-CE links in terms of admission control. If the procedures defined in Section 3 are enabled on both ASBRs, then admission control may be performed on the inter-ASBR links. In addition, the operator of each AS can independently decide whether or not to perform admission control across his backbone. The new objects described in this document MUST NOT be sent in any RSVP message between two Option-A ASBRs.

5.2. Inter-AS Option B

To support inter-AS Option B, we require some additional processing of RSVP messages on the ASBRs. Recall that, when packets are forwarded from one AS to another in Option B, the VPN label is swapped by each ASBR as a packet goes from one AS to another. The

BGP next hop seen by the ingress PE will be the ASBR, and there need not be IP visibility between the ingress and egress PEs. Hence, when the ingress PE sends the Path message to the BGP next hop of the VPN-IPv4 route towards the destination, it will be received by the ASBR. The ASBR determines the next hop of the route in a similar way as the ingress PE -- by finding a matching BGP VPN-IPv4 route with the same RD and a matching prefix.

The provider(s) who interconnect ASes using Option B may or may not desire to perform admission control on the inter-AS links. This choice affects the detailed operation of ASBRs. We describe the two modes of operation -- with and without admission control at the ASBRs -- in the following sections.

5.2.1. Admission Control on ASBR

In this scenario, the ASBR performs full RSVP signaling and admission control. The RSVP database is indexed on the ASBR using the VPN-IPv4 SESSION, SENDER_TEMPLATE, and FILTER_SPEC objects (which uniquely identify RSVP sessions and flows as per the requirements of [RFC2205]). These objects are forwarded unmodified in both directions by the ASBR. All other procedures of RSVP are performed as if the ASBR was an RSVP hop. In particular, the RSVP_HOP objects sent in Path and Resv messages contain IP addresses of the ASBR, which MUST be reachable by the neighbor to whom the message is being sent. Note that since the VPN-IPv4 SESSION, SENDER_TEMPLATE, and FILTER_SPEC objects satisfy the uniqueness properties required for an RSVP database implementation as per [RFC2209], no customer VRF awareness is required on the ASBR.

5.2.2. No Admission Control on ASBR

If the ASBR is not doing admission control, it is desirable that per-flow state not be maintained on the ASBR. This requires adjacent RSVP hops (i.e., the ingress and egress PEs of the respective ASes) to send RSVP messages directly to each other. This is only possible if they are MPLS encapsulated. The use of the VPN-IPv4 RSVP_HOP object described in Section 3.1 is REQUIRED in this case.

When an ASBR that is not installing local RSVP state receives a Path message, it looks up the next hop of the matching BGP route as described in Section 3.2, and sends the Path message to the next hop, without modifying any RSVP objects (including the RSVP_HOP). This process is repeated at subsequent ASBRs until the Path message arrives at a router that is installing local RSVP state (either the ultimate egress PE, or an ASBR configured to perform admission control). This router receives the Path and processes it as described in Section 3.3 if it is a PE, or Section 5.2.1 if it is an

ASBR performing admission control. When this router sends the Resv upstream, it looks up the routing table for a next hop+label for the VPN-IPv4 address in the PHOP, encapsulates the Resv with that label, and sends it upstream. This message will be received for control processing directly on the upstream RSVP hop (that last updated the RSVP_HOP field in the Path message), without any involvement of intermediate ASBRs.

The ASBR is not expected to process any other RSVP messages apart from the Path message as described above. The ASBR also does not need to store any RSVP state. Note that any ASBR along the path that wishes to do admission control or insert itself into the RSVP signaling flow may do so by writing its own RSVP_HOP object with IPv4 and VPN-IPv4 addresses pointing to itself.

If an Option-B ASBR that receives an RSVP Path message with an IPv4 RSVP_HOP does not wish to perform admission control but is willing to install local state for this flow, the ASBR MUST process and forward RSVP signaling messages for this flow as described in Section 5.2.1 (with the exception that it does not perform admission control). If an Option-B ASBR receives an RSVP Path message with an IPv4 RSVP_HOP, but does not wish to install local state or perform admission control for this flow, the ASBR MUST NOT forward the Path message. In addition, the ASBR SHOULD send a PathError message of Error Code "RSVP over MPLS Problem" and Error Value "RSVP_HOP not reachable across VPN" (see Section 9) signifying to the upstream RSVP hop that the supplied RSVP_HOP object is insufficient to provide reachability across this VPN. This failure condition is not expected to be recoverable.

5.3. Inter-AS Option C

Operation of RSVP in Inter-AS Option C is also quite straightforward, because there exists an LSP directly from ingress PE to egress PE. In this case, there is no significant difference in operation from the single AS case described in Section 3. Furthermore, if it is desired to provide admission control from PE to PE, it can be done by building an inter-AS TE tunnel and then using the procedures described in Section 4.

6. Operation with RSVP Disabled

It is often the case that RSVP will not be enabled on the PE-CE links. In such an environment, a customer may reasonably expect that RSVP messages sent into the L3 VPN network should be forwarded just like any other IP datagrams. This transparency is useful when the customer wishes to use RSVP within his own sites or perhaps to perform admission control on the CE-PE links (in CE->PE direction

only), without involvement of the PEs. For this reason, a PE SHOULD NOT discard or modify RSVP messages sent towards it from a CE when RSVP is not enabled on the PE-CE links. Similarly a PE SHOULD NOT discard or modify RSVP messages that are destined for one of its attached CEs, even when RSVP is not enabled on those links. Note that the presence of the Router Alert Option in some RSVP messages may cause them to be forwarded outside of the normal forwarding path, but that the guidance of this paragraph still applies in that case. Note also that this guidance applies regardless of whether RSVP-TE is used in some, all, or none of the L3VPN network.

7. Other RSVP Procedures

This section describes modifications to other RSVP procedures introduced by MPLS VPNs.

7.1. Refresh Overhead Reduction

The following points ought to be noted regarding RSVP refresh overhead reduction [RFC2961] across an MPLS VPN:

- o The hop between the ingress and egress PE of a VPN is to be considered as traversing one or more non-RSVP hops. As such, the procedures described in Section 5.3 of [RFC2961] relating to non-RSVP hops SHOULD be followed.
- o The source IP address of a SRefresh message MUST match the IPv4 address signaled in the RSVP_HOP object contained in the corresponding Path or Resv message. The IPv4 address in any received VPN-IPv4 RSVP_HOP object MUST be used as the source address of that message for this purpose.

7.2. Cryptographic Authentication

The following points ought to be noted regarding RSVP cryptographic authentication ([RFC2747]) across an MPLS VPN:

- o The IPv4 address in any received VPN-IPv4 RSVP_HOP object MUST be used as the source address of that message for purposes of identifying the security association.
- o Forwarding of Challenge and Response messages MUST follow the same rules as described above for hop-by-hop messages. Specifically, if the originator of a Challenge/Response message has received a VPN-IPv4 RSVP_HOP object from the corresponding neighbor, it MUST use the label associated with that VPN-IPv4 address in BGP to forward the Challenge/Response message.

7.3. RSVP Aggregation

[RFC3175] and [RFC4860] describe mechanisms to aggregate multiple individual RSVP reservations into a single larger reservation on the basis of a common Differentiated Services Code Point/Per-Hop Behavior (DSCP/PHB) for traffic classification. The following points ought to be noted in this regard:

- o The procedures described in this section apply only in the case where the Aggregator and Deaggregator nodes are C/CE devices, and the entire MPLS VPN lies within the Aggregation Region. The case where the PE is also an Aggregator/Deaggregator is more complex and not considered in this document.
- o Support of Aggregate RSVP sessions is OPTIONAL. When supported:
 - * Aggregate RSVP sessions MUST be treated in the same way as regular IPv4 RSVP sessions. To this end, all the procedures described in Sections 3 and 4 MUST be followed for aggregate RSVP sessions. The corresponding new SESSION, SENDER_TEMPLATE, and FILTERSPEC objects are defined in Section 8.
 - * End-To-End (E2E) RSVP sessions are passed unmodified through the MPLS VPN. These RSVP messages SHOULD be identified by their IP protocol (RSVP-E2E-IGNORE, 134). When the ingress PE receives any RSVP message with this IP protocol, it MUST process this frame as if it is regular customer traffic and ignore any Router Alert Option. The appropriate VPN and transport labels are applied to the frame and it is forwarded towards the remote CE. Note that this message will not be received or processed by any other P or PE node.
 - * Any SESSION-OF-INTEREST object (defined in [RFC4860]) MUST be conveyed unmodified across the MPLS VPN.

7.4. Support for CE-CE RSVP-TE

[RFC5824] describes a set of requirements for the establishment for CE-CE MPLS LSPs across networks offering an L3VPN service. The requirements specified in that document are similar to those addressed by this document, in that both address the issue of handling RSVP requests from customers in a VPN context. It is possible that the solution described here could be adapted to meet the requirements of [RFC5824]. To the extent that this document uses signaling extensions described in [RFC3473] that have already been used for GMPLS/TE, we expect that CE-CE RSVP/TE will be incremental work built on these extensions. These extensions will be considered in a separate document.

8. Object Definitions

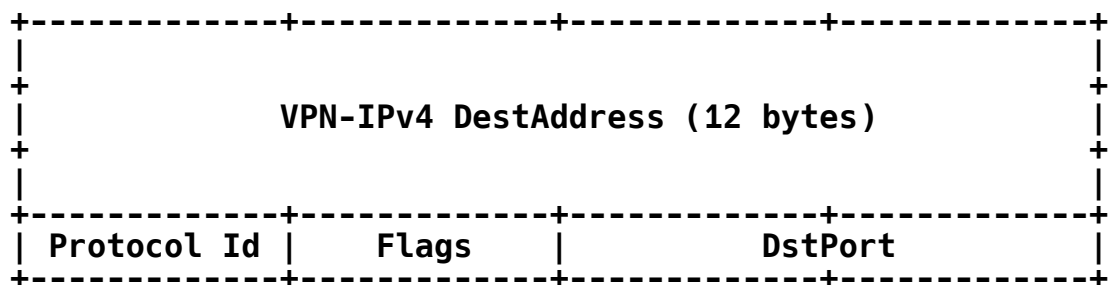
8.1. VPN-IPv4 and VPN-IPv6 SESSION Objects

The usage of the VPN-IPv4 (or VPN-IPv6) SESSION object is described in Sections 3.2 to 3.6. The VPN-IPv4 (or VPN-IPv6) SESSION object appears in RSVP messages that ordinarily contain a SESSION object and are sent between ingress PE and egress PE in either direction. The object **MUST NOT** be included in any RSVP messages that are sent outside of the provider's backbone (except in the inter-AS Option-B and Option-C cases, as described above, when it may appear on inter-AS links).

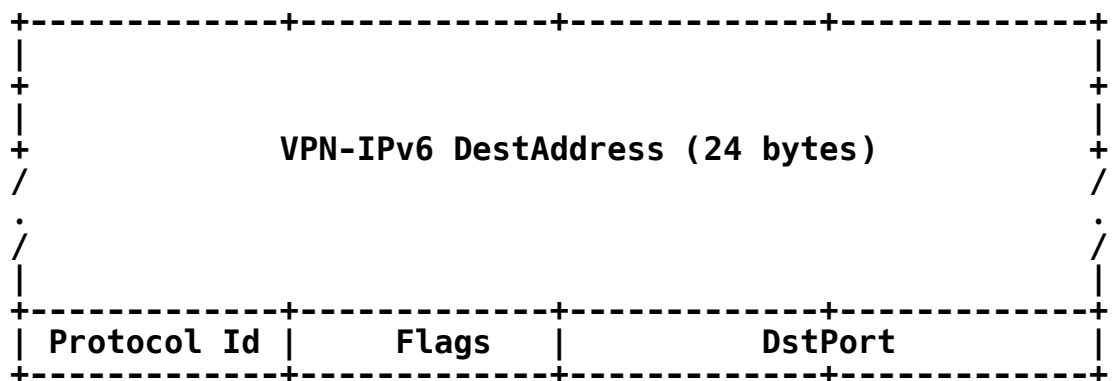
The VPN-IPv6 SESSION object is analogous to the VPN-IPv4 SESSION object, using an VPN-IPv6 address ([RFC4659]) instead of an VPN-IPv4 address ([RFC4364]).

The formats of the objects are as follows:

- o VPN-IPv4 SESSION object: Class = 1, C-Type = 19



- o VPN-IPv6 SESSION object: Class = 1, C-Type = 20




The VPN-IPv4 DestAddress (respectively, VPN-IPv6 DestAddress) field contains an address of the VPN-IPv4 (respectively, VPN-IPv6) address family encoded as specified in [RFC4364] (respectively, [RFC4659]). The content of this field is discussed in Sections 3.2 and 3.3.

The protocol ID, flags, and DstPort are identical to the same fields in the IPv4 and IPv6 SESSION objects ([RFC2205]).


8.2. VPN-IPv4 and VPN-IPv6 SENDER_TEMPLATE Objects

The usage of the VPN-IPv4 (or VPN-IPv6) SENDER_TEMPLATE object is described in Sections 3.2 and 3.3. The VPN-IPv4 (or VPN-IPv6) SENDER_TEMPLATE object appears in RSVP messages that ordinarily contain a SENDER_TEMPLATE object and are sent between ingress PE and egress PE in either direction (such as Path, PathError, and PathTear). The object MUST NOT be included in any RSVP messages that are sent outside of the provider's backbone (except in the inter-AS Option-B and Option-C cases, as described above, when it may appear on inter-AS links). The format of the object is as follows:

- ```
0 VPN-IPv4 SENDER_TEMPLATE object: Class = 11, C-Type = 14
```



- ```
0 VPN-IPv6 SENDER_TEMPLATE object: Class = 11, C-Type = 15
```



The VPN-IPv4 SrcAddress (respectively, VPN-IPv6 SrcAddress) field contains an address of the VPN-IPv4 (respectively, VPN-IPv6) address family encoded as specified in [RFC4364] (respectively, [RFC4659]). The content of this field is discussed in Sections 3.2 and 3.3.

The SrcPort is identical to the SrcPort field in the IPv4 and IPv6 SENDER_TEMPLATE objects ([RFC2205]).

The Reserved field MUST be set to zero on transmit and ignored on receipt.

8.3. VPN-IPv4 and VPN-IPv6 FILTER_SPEC Objects

The usage of the VPN-IPv4 (or VPN-IPv6) FILTER_SPEC object is described in Sections 3.4 and 3.5. The VPN-IPv4 (or VPN-IPv6) FILTER_SPEC object appears in RSVP messages that ordinarily contain a FILTER_SPEC object and are sent between ingress PE and egress PE in either direction (such as Resv, ResvError, and ResvTear). The object MUST NOT be included in any RSVP messages that are sent outside of the provider's backbone (except in the inter-AS Option-B and Option-C cases, as described above, when it may appear on inter-AS links).

- o VPN-IPv4 FILTER_SPEC object: Class = 10, C-Type = 14

Definition same as VPN-IPv4 SENDER_TEMPLATE object.

- o VPN-IPv6 FILTER_SPEC object: Class = 10, C-Type = 15

Definition same as VPN-IPv6 SENDER_TEMPLATE object.

The content of the VPN-IPv4 SrcAddress (or VPN-IPv6 SrcAddress) field is discussed in Sections 3.4 and 3.5.

The SrcPort is identical to the SrcPort field in the IPv4 and IPv6 SENDER_TEMPLATE objects ([RFC2205]).

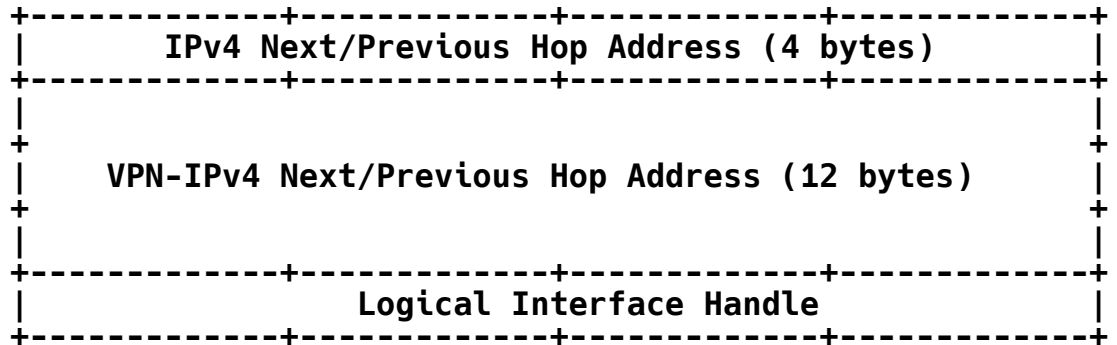
The Reserved field MUST be set to zero on transmit and ignored on receipt.

8.4. VPN-IPv4 and VPN-IPv6 RSVP_HOP Objects

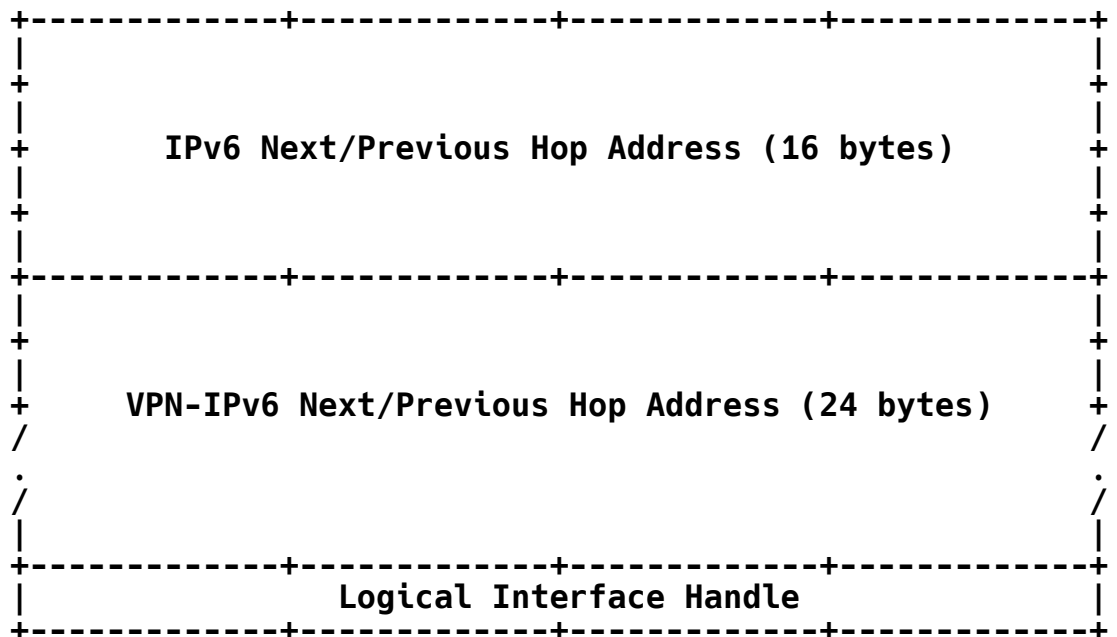
Usage of the VPN-IPv4 (or VPN-IPv6) RSVP_HOP object is described in Sections 3.1 and 5.2.2. The VPN-IPv4 (VPN-IPv6) RSVP_HOP object is used to establish signaling reachability between RSVP neighbors separated by one or more Option-B ASBRs. This object may appear in RSVP messages that carry an RSVP_HOP object, and that travel between the ingress and egress PEs. It MUST NOT be included in any RSVP

messages that are sent outside of the provider's backbone (except in the inter-AS Option-B and Option-C cases, as described above, when it may appear on inter-AS links). The format of the object is as follows:

- o VPN-IPv4 RSVP_HOP object: Class = 3, C-Type = 5



- o VPN-IPv6 RSVP_HOP object: Class = 3, C-Type = 6



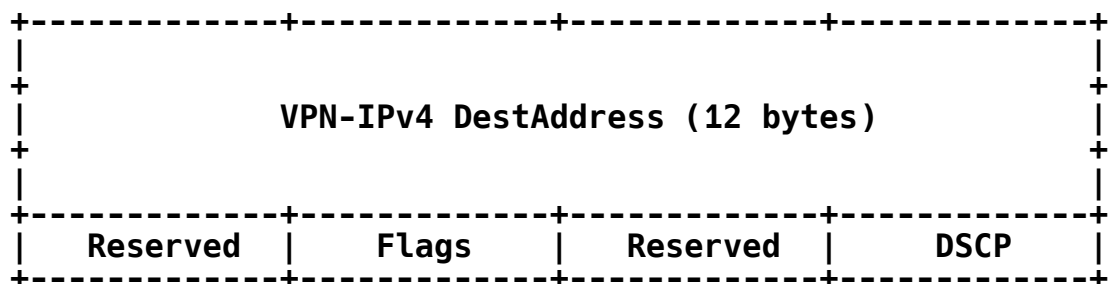
The IPv4 Next/Previous Hop Address, IPv6 Next/Previous Hop Address, and the Logical Interface Handle fields are identical to those of the RSVP_HOP object ([RFC2205]).

The VPN-IPv4 Next/Previous Hop Address (respectively, VPN-IPv6 Next/Previous Hop Address) field contains an address of the VPN-IPv4 (respectively, VPN-IPv6) address family encoded as specified in [RFC4364] (respectively, [RFC4659]). The content of this field is discussed in Section 3.1.

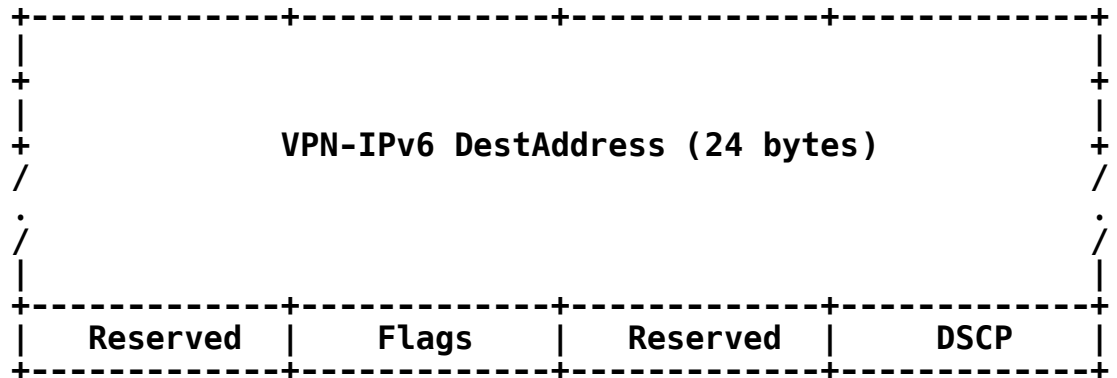
8.5. Aggregated VPN-IPv4 and VPN-IPv6 SESSION Objects

The usage of Aggregated VPN-IPv4 (or VPN-IPv6) SESSION object is described in Section 7.3. The AGGREGATE-VPN-IPv4 (respectively, AGGREGATE-IPv6-VPN) SESSION object appears in RSVP messages that ordinarily contain a AGGREGATE-IPv4 (respectively, AGGREGATE-IPv6) SESSION object as defined in [RFC3175] and are sent between ingress PE and egress PE in either direction. The GENERIC-AGGREGATE-VPN-IPv4 (respectively, AGGREGATE-VPN-IPv6) SESSION object should appear in all RSVP messages that ordinarily contain a GENERIC-AGGREGATE-IPv4 (respectively, GENERIC-AGGREGATE-IPv6) SESSION object as defined in [RFC4860] and are sent between ingress PE and egress PE in either direction. These objects MUST NOT be included in any RSVP messages that are sent outside of the provider's backbone (except in the inter-AS Option-B and Option-C cases, as described above, when it may appear on inter-AS links). The processing rules for these objects are otherwise identical to those of the VPN-IPv4 (respectively, VPN-IPv6) SESSION object defined in Section 8.1. The format of the object is as follows:

- o AGGREGATE-VPN-IPv4 SESSION object: Class = 1, C-Type = 21



- ```
o AGGREGATE-VPN-IPv6 SESSION object: Class = 1, C-Type = 22
```

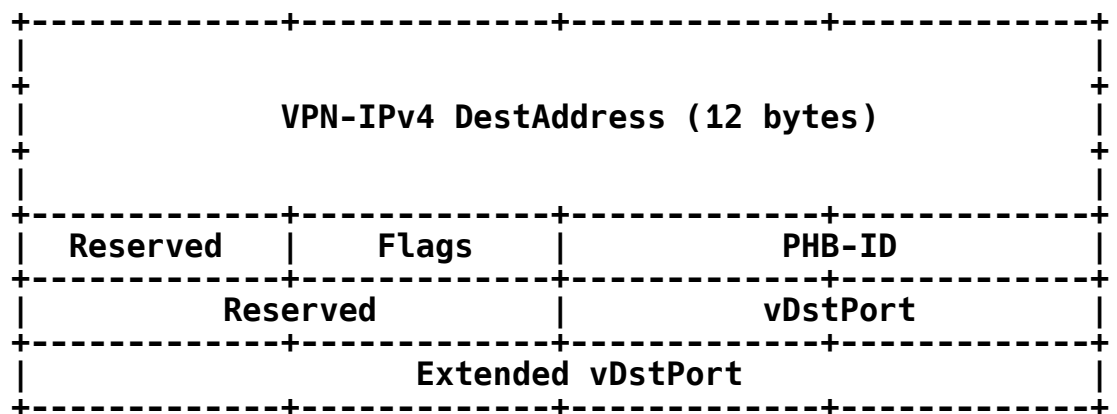


The VPN-IPv4 DestAddress (respectively, VPN-IPv6 DestAddress) field contains an address of the VPN-IPv4 (respectively, VPN-IPv6) address family encoded as specified in [RFC4364] (respectively, [RFC4659]). The content of this field is discussed in Sections 3.2 and 3.3.

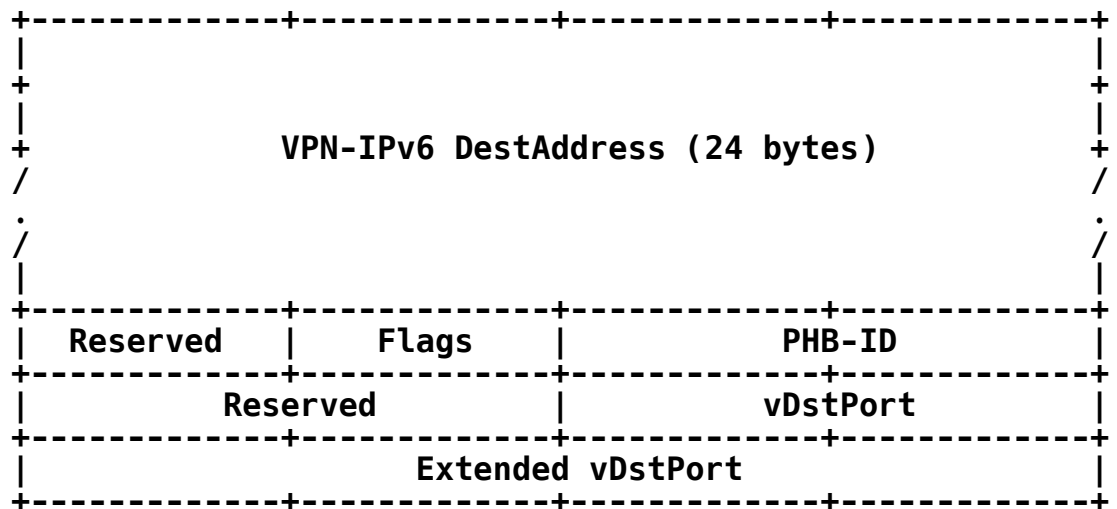
The flags and DSCP are identical to the same fields of the AGGREGATE-IPv4 and AGGREGATE-IPv6 SESSION objects ([RFC3175]).

**The Reserved field MUST be set to zero on transmit and ignored on receipt.**

- ```
0  GENERIC-AGGREGATE-VPN-IPv4 SESSION object:
    Class = 1, C-Type = 23
```



- o **GENERIC-AGGREGATE-VPN-IPv6 SESSION object:**
Class = 1, C-Type = 24



The VPN-IPv4 DestAddress (respectively, VPN-IPv6 DestAddress) field contains an address of the VPN-IPv4 (respectively, VPN-IPv6) address family encoded as specified in [RFC4364] (respectively, [RFC4659]). The content of this field is discussed in Sections 3.2 and 3.3.

The flags, PHB-ID, vDstPort, and Extended vDstPort are identical to the same fields of the GENERIC-AGGREGATE-IPv4 and GENERIC-AGGREGATE-IPv6 SESSION objects ([RFC4860]).

The Reserved field MUST be set to zero on transmit and ignored on receipt.


8.6. AGGREGATE-VPN-IPv4 and AGGREGATE-VPN-IPv6 SENDER_TEMPLATE Objects

The usage of Aggregated VPN-IPv4 (or VPN-IPv6) SENDER_TEMPLATE object is described in Section 7.3. The AGGREGATE-VPN-IPv4 (respectively, AGGREGATE-VPN-IPv6) SENDER_TEMPLATE object appears in RSVP messages that ordinarily contain a AGGREGATE-IPv4 (respectively, AGGREGATE-IPv6) SENDER_TEMPLATE object as defined in [RFC3175] and [RFC4860], and are sent between ingress PE and egress PE in either direction. These objects MUST NOT be included in any RSVP messages that are sent outside of the provider's backbone (except in the inter-AS Option-B and Option-C cases, as described above, when it may appear on inter-AS links). The processing rules for these objects are otherwise identical to those of the VPN-IPv4 (respectively, VPN-IPv6) SENDER_TEMPLATE object defined in Section 8.2. The format of the object is as follows:

- o **AGGREGATE-VPN-IPv4 SENDER_TEMPLATE object:**
Class = 11, C-Type = 16



- o **AGGREGATE-VPN-IPv6 SENDER_TEMPLATE object:**
Class = 11, C-Type = 17



The VPN-IPv4 AggregatorAddress (respectively, VPN-IPv6 AggregatorAddress) field contains an address of the VPN-IPv4 (respectively, VPN-IPv6) address family encoded as specified in [RFC4364] (respectively, [RFC4659]). The content and processing rules for these objects are similar to those of the VPN-IPv4 SENDER_TEMPLATE object defined in Section 8.2.

The flags and DSCP are identical to the same fields of the AGGREGATE-IPv4 and AGGREGATE-IPv6 SESSION objects.

8.7. AGGREGATE-VPN-IPv4 and AGGREGATE-VPN-IPv6 FILTER_SPEC Objects

The usage of Aggregated VPN-IPv4 FILTER_SPEC object is described in Section 7.3. The AGGREGATE-VPN-IPv4 FILTER_SPEC object appears in RSVP messages that ordinarily contain a AGGREGATE-IPv4 FILTER_SPEC object as defined in [RFC3175] and [RFC4860], and are sent between ingress PE and egress PE in either direction. These objects MUST NOT be included in any RSVP messages that are sent outside of the provider's backbone (except in the inter-AS Option-B and Option-C cases, as described above, when it may appear on inter-AS links).

The processing rules for these objects are otherwise identical to those of the VPN-IPv4 FILTER_SPEC object defined in Section 8.3. The format of the object is as follows:

- o AGGREGATE-VPN-IPv4 FILTER_SPEC object:

Class = 10, C-Type = 16

Definition same as AGGREGATE-VPN-IPv4 SENDER_TEMPLATE object.

- o AGGREGATE-VPN-IPv6 FILTER_SPEC object:

Class = 10, C-Type = 17

Definition same as AGGREGATE-VPN-IPv6 SENDER_TEMPLATE object.

9. IANA Considerations

Section 8 defines new objects. Therefore, IANA has modified the RSVP parameters registry, 'Class Names, Class Numbers, and Class Types' subregistry, and:

- o assigned six new C-Types under the existing SESSION Class (Class number 1), as follows:

Class Number	Class Name	Reference
-----	-----	-----
1	SESSION	[RFC2205]
Class Types or C-Types:		
19	VPN-IPv4	[RFC6016]
20	VPN-IPv6	[RFC6016]
21	AGGREGATE-VPN-IPv4	[RFC6016]
22	AGGREGATE-VPN-IPv6	[RFC6016]
23	GENERIC-AGGREGATE-VPN-IPv4	[RFC6016]
24	GENERIC-AGGREGATE-VPN-IPv6	[RFC6016]

- o assigned four new C-Types under the existing SENDER_TEMPLATE Class (Class number 11), as follows:

Class Number	Class Name	Reference
-----	-----	-----

11	SENDER_TEMPLATE	[RFC2205]
----	-----------------	-----------

Class Types or C-Types:

14	VPN-IPv4	[RFC6016]
15	VPN-IPv6	[RFC6016]
16	AGGREGATE-VPN-IPv4	[RFC6016]
17	AGGREGATE-VPN-IPv6	[RFC6016]

- o assigned four new C-Types under the existing FILTER_SPEC Class (Class number 10), as follows:

Class Number	Class Name	Reference
-----	-----	-----

10	FILTER_SPEC	[RFC2205]
----	-------------	-----------

Class Types or C-Types:

14	VPN-IPv4	[RFC6016]
15	VPN-IPv6	[RFC6016]
16	AGGREGATE-VPN-IPv4	[RFC6016]
17	AGGREGATE-VPN-IPv6	[RFC6016]

- o assigned two new C-Types under the existing RSVP_HOP Class (Class number 3), as follows:

Class Number	Class Name	Reference
-----	-----	-----

3	RSVP_HOP	[RFC2205]
---	----------	-----------

Class Types or C-Types:

5	VPN-IPv4	[RFC6016]
6	VPN-IPv6	[RFC6016]

In addition, a new PathError code/value is required to identify a signaling reachability failure and the need for a VPN-IPv4 or VPN-IPv6 RSVP_HOP object as described in Section 5.2.2. Therefore, IANA has modified the RSVP parameters registry, 'Error Codes and Globally-Defined Error Value Sub-Codes' subregistry, and:

- o assigned a new Error Code and sub-code, as follows:

37 RSVP over MPLS Problem [RFC6016]

This Error Code has the following globally-defined Error Value sub-codes:

1 = RSVP_HOP not reachable across VPN [RFC6016]

10. Security Considerations

[RFC4364] addresses the security considerations of BGP/MPLS VPNs in general. General RSVP security considerations are discussed in [RFC2205]. To ensure the integrity of RSVP, the RSVP Authentication mechanisms defined in [RFC2747] and [RFC3097] SHOULD be supported. Those protect RSVP message integrity hop-by-hop and provide node authentication as well as replay protection, thereby protecting against corruption and spoofing of RSVP messages. [RSVP-KEYING] discusses applicability of various keying approaches for RSVP Authentication. First, we note that the discussion about applicability of group keying to an intra-provider environment where RSVP hops are not IP hops is relevant to securing of RSVP among PEs of a given Service Provider deploying the solution specified in the present document. We note that the RSVP signaling in MPLS VPN is likely to spread over multiple administrative domains (e.g., the service provider operating the VPN service, and the customers of the service). Therefore the considerations in [RSVP-KEYING] about inter-domain issues are likely to apply.

Since RSVP messages travel through the L3VPN cloud directly addressed to PE or ASBR routers (without IP Router Alert Option), P routers remain isolated from RSVP messages signaling customer reservations. Providers MAY choose to block PEs from sending datagrams with the Router Alert Option to P routers as a security practice, without impacting the functionality described herein.

Beyond those general issues, four specific issues are introduced by this document: resource usage on PEs, resource usage in the provider backbone, PE route advertisement outside the AS, and signaling exposure to ASBRs and PEs. We discuss these in turn.

A customer who makes resource reservations on the CE-PE links for his sites is only competing for link resources with himself, as in standard RSVP, at least in the common case where each CE-PE link is dedicated to a single customer. Thus, from the perspective of the CE-PE links, the present document does not introduce any new security issues. However, because a PE typically serves multiple customers, there is also the possibility that a customer might attempt to use excessive computational resources on a PE (CPU cycles, memory, etc.) by sending large numbers of RSVP messages to a PE. In the extreme, this could represent a form of denial-of-service attack. In order to prevent such an attack, a PE **SHOULD** support mechanisms to limit the fraction of its processing resources that can be consumed by any one CE or by the set of CEs of a given customer. For example, a PE might implement a form of rate limiting on RSVP messages that it receives from each CE. We observe that these security risks and measures related to PE resource usage are very similar for any control-plane protocol operating between CE and PE (e.g., RSVP, routing, multicast).

The second concern arises only when the service provider chooses to offer resource reservation across the backbone, as described in Section 4. In this case, the concern may be that a single customer might attempt to reserve a large fraction of backbone capacity, perhaps with a coordinated effort from several different CEs, thus denying service to other customers using the same backbone. [RFC4804] provides some guidance on the security issues when RSVP reservations are aggregated onto MPLS tunnels, which are applicable to the situation described here. We note that a provider **MAY** use local policy to limit the amount of resources that can be reserved by a given customer from a particular PE, and that a policy server could be used to control the resource usage of a given customer across multiple PEs if desired. It is **RECOMMENDED** that an implementation of this specification support local policy on the PE to control the amount of resources that can be reserved by a given customer/CE.

Use of the VPN-IPv4 RSVP_HOP object requires exporting a PE VPN-IPv4 route to another AS, and potentially could allow unchecked access to remote PEs if those routes were indiscriminately redistributed. However, as described in Section 3.1, no route that is not within a customer's VPN should ever be advertised to (or be reachable from) that customer. If a PE uses a local address already within a customer VRF (like PE-CE link address), it **MUST NOT** send this address in any RSVP messages in a different customer VRF. A "control-plane" VPN **MAY** be created across PEs and ASBRs and addresses in this VPN can be used to signal RSVP sessions for any customers, but these routes **MUST NOT** be advertised to, or made reachable from, any customer. An implementation of the present document **MAY** support such operation using a "control-plane" VPN. Alternatively, ASBRs **MAY** implement the

signaling procedures described in Section 5.2.1, even if admission control is not required on the inter-AS link, as these procedures do not require any direct P/PE route advertisement out of the AS.

Finally, certain operations described herein (Section 3) require an ASBR or PE to receive and locally process a signaling packet addressed to the BGP next hop address advertised by that router. This requirement does not strictly apply to MPLS/BGP VPNs [RFC4364]. This could be viewed as opening ASBRs and PEs to being directly addressable by customer devices where they were not open before, and could be considered a security issue. If a provider wishes to mitigate this situation, the implementation MAY support the "control protocol VPN" approach described above. That is, whenever a signaling message is to be sent to a PE or ASBR, the address of the router in question would be looked up in the "control protocol VPN", and the message would then be sent on the LSP that is found as a result of that lookup. This would ensure that the router address is not reachable by customer devices.

[RFC4364] mentions use of IPsec both on a CE-CE basis and PE-PE basis:

Cryptographic privacy is not provided by this architecture, nor by Frame Relay or ATM VPNs. These architectures are all compatible with the use of cryptography on a CE-CE basis, if that is desired.

The use of cryptography on a PE-PE basis is for further study.

The procedures specified in the present document for admission control on the PE-CE links (Section 3) are compatible with the use of IPsec on a PE-PE basis. The optional procedures specified in the present document for admission control in the Service Provider's backbone (Section 4) are not compatible with the use of IPsec on a PE-PE basis, since those procedures depend on the use of PE-PE MPLS TE Tunnels to perform aggregate reservations through the Service Provider's backbone.

[RFC4923] describes a model for RSVP operation through IPsec Gateways. In a nutshell, a form of hierarchical RSVP reservation is used where an RSVP reservation is made for the IPsec tunnel and then individual RSVP reservations are admitted/aggregated over the tunnel reservation. This model applies to the case where IPsec is used on a CE-CE basis. In that situation, the procedures defined in the present document would simply apply "as is" to the reservation established for the IPsec tunnel(s).

11. Acknowledgments

Thanks to Ashwini Dahiya, Prashant Srinivas, Yakov Rekhter, Eric Rosen, Dan Tappan, and Lou Berger for their many contributions to solving the problems described in this document. Thanks to Ferit Yegenoglu for his useful comments. We also thank Stefan Santesson, Vijay Gurbani, and Alexey Melnikov for their review comments. We thank Richard Woundy for his very thorough review and comments including those that resulted in additional text discussing scenarios of admission control reject in the MPLS VPN cloud. Also, we thank Adrian Farrel for his detailed review and contributions.

Appendix A. Alternatives Considered

At this stage, a number of alternatives to the approach described above have been considered. We document some of the approaches considered here to assist future discussion. None of these have been shown to improve upon the approach described above, and the first two seem to have significant drawbacks relative to the approach described above.

Appendix A.1. GMPLS UNI Approach

[RFC4208] defines the GMPLS UNI. In Section 7, the operation of the GMPLS UNI in a VPN context is briefly described. This is somewhat similar to the problem tackled in the current document. The main difference is that the GMPLS UNI is primarily aimed at the problem of allowing a CE device to request the establishment of a Label Switched Path (LSP) across the network on the other side of the UNI. Hence, the procedures in [RFC4208] would lead to the establishment of an LSP across the VPN provider's network for every RSVP request received, which is not desired in this case.

To the extent possible, the approach described in this document is consistent with [RFC4208], while filling in more of the details and avoiding the problem noted above.

Appendix A.2. Label Switching Approach

Implementations that always look at IP headers inside the MPLS label on the egress PE can intercept Path messages and determine the correct VRF and RSVP state by using a combination of the encapsulating VPN label and the IP header. In our view, this is an undesirable approach for two reasons. Firstly, it imposes a new MPLS forwarding requirement for all data packets on the egress PE. Secondly, it requires using the encapsulating MPLS label to identify RSVP state, which runs counter to existing RSVP principle and practice where all information used to identify RSVP state is included within RSVP objects. RSVP extensions such as COPS/RSVP [RFC2749] which re-encapsulate RSVP messages are incompatible with this change.

Appendix A.3. VRF Label Approach

Another approach to solving the problems described here involves the use of label switching to ensure that Path, Resv, and other RSVP messages are directed to the appropriate VRF on the next RSVP hop (e.g., egress PE). One challenge with such an approach is that [RFC4364] does not require labels to be allocated for VRFs, only for customer prefixes, and that there is no simple, existing method for

advertising the fact that a label is bound to a VRF. If, for example, an ingress PE sent a Path message labelled with a VPN label that was advertised by the egress PE for the prefix that matches the destination address in the Path, there is a risk that the egress PE would simply label-switch the Path directly on to the CE without performing RSVP processing.

A second challenge with this approach is that an IP address needs to be associated with a VRF and used as the PHOP address for the Path message sent from ingress PE to egress PE. That address needs to be reachable from the egress PE, and to exist in the VRF at the ingress PE. Such an address is not always available in today's deployments, so this represents at least a change to existing deployment practices.

Appendix A.4. VRF Label Plus VRF Address Approach

It is possible to create an approach based on that described in the previous section that addresses the main challenges of that approach. The basic approach has two parts: (a) define a new BGP Extended Community to tag a route (and its associated MPLS label) as pointing to a VRF; (b) allocate a "dummy" address to each VRF, specifically to be used for routing RSVP messages. The dummy address (which could be anything, e.g., a loopback of the associated PE) would be used as a PHOP for Path messages and would serve as the destination for Resv messages but would not be imported into VRFs of any other PE.

References

Normative References

- [RFC2113] Katz, D., "IP Router Alert Option", RFC 2113, February 1997.
- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, March 1997.
- [RFC2205] Braden, B., Zhang, L., Berson, S., Herzog, S., and S. Jamin, "Resource ReSerVation Protocol (RSVP) -- Version 1 Functional Specification", RFC 2205, September 1997.
- [RFC2711] Partridge, C. and A. Jackson, "IPv6 Router Alert Option", RFC 2711, October 1999.
- [RFC3175] Baker, F., Iturralde, C., Le Faucheur, F., and B. Davie, "Aggregation of RSVP for IPv4 and IPv6 Reservations", RFC 3175, September 2001.

- [RFC4364] Rosen, E. and Y. Rekhter, "BGP/MPLS IP Virtual Private Networks (VPNs)", RFC 4364, February 2006.
- [RFC4659] De Clercq, J., Ooms, D., Carugi, M., and F. Le Faucheur, "BGP-MPLS IP Virtual Private Network (VPN) Extension for IPv6 VPN", RFC 4659, September 2006.
- [RFC4804] Le Faucheur, F., "Aggregation of Resource ReSerVation Protocol (RSVP) Reservations over MPLS TE/DS-TE Tunnels", RFC 4804, February 2007.

Informative References

- [ALERT-USAGE] Le Faucheur, F., Ed., "IP Router Alert Considerations and Usage", Work in Progress, July 2010.
- [LER-OPTIONS] Smith, D., Mullooly, J., Jaeger, W., and T. Scholl, "Requirements for Label Edge Router Forwarding of IPv4 Option Packets", Work in Progress, May 2010.
- [RFC1633] Braden, B., Clark, D., and S. Shenker, "Integrated Services in the Internet Architecture: an Overview", RFC 1633, June 1994.
- [RFC2209] Braden, B. and L. Zhang, "Resource ReSerVation Protocol (RSVP) -- Version 1 Message Processing Rules", RFC 2209, September 1997.
- [RFC2210] Wroclawski, J., "The Use of RSVP with IETF Integrated Services", RFC 2210, September 1997.
- [RFC2747] Baker, F., Lindell, B., and M. Talwar, "RSVP Cryptographic Authentication", RFC 2747, January 2000.
- [RFC2748] Durham, D., Boyle, J., Cohen, R., Herzog, S., Rajan, R., and A. Sastry, "The COPS (Common Open Policy Service) Protocol", RFC 2748, January 2000.
- [RFC2749] Herzog, S., Boyle, J., Cohen, R., Durham, D., Rajan, R., and A. Sastry, "COPS usage for RSVP", RFC 2749, January 2000.
- [RFC2961] Berger, L., Gan, D., Swallow, G., Pan, P., Tommasi, F., and S. Molendini, "RSVP Refresh Overhead Reduction Extensions", RFC 2961, April 2001.

- [RFC3097] Braden, R. and L. Zhang, "RSVP Cryptographic Authentication -- Updated Message Type Value", RFC 3097, April 2001.
- [RFC3473] Berger, L., "Generalized Multi-Protocol Label Switching (GMPLS) Signaling Resource ReserVation Protocol-Traffic Engineering (RSVP-TE) Extensions", RFC 3473, January 2003.
- [RFC4206] Kompella, K. and Y. Rekhter, "Label Switched Paths (LSP) Hierarchy with Generalized Multi-Protocol Label Switching (GMPLS) Traffic Engineering (TE)", RFC 4206, October 2005.
- [RFC4208] Swallow, G., Drake, J., Ishimatsu, H., and Y. Rekhter, "Generalized Multiprotocol Label Switching (GMPLS) User-Network Interface (UNI): Resource ReserVation Protocol-Traffic Engineering (RSVP-TE) Support for the Overlay Model", RFC 4208, October 2005.
- [RFC4860] Le Faucheur, F., Davie, B., Bose, P., Christou, C., and M. Davenport, "Generic Aggregate Resource ReSerVation Protocol (RSVP) Reservations", RFC 4860, May 2007.
- [RFC4923] Baker, F. and P. Bose, "Quality of Service (QoS) Signaling in a Nested Virtual Private Network", RFC 4923, August 2007.
- [RFC5824] Kumaki, K., Zhang, R., and Y. Kamite, "Requirements for Supporting Customer Resource ReSerVation Protocol (RSVP) and RSVP Traffic Engineering (RSVP-TE) over a BGP/MPLS IP-VPN", RFC 5824, April 2010.
- [RFC5971] Schulzrinne, H. and R. Hancock, "GIST: General Internet Signalling Transport", RFC 5971, October 2010.
- [RFC5974] Manner, J., Karagiannis, G., and A. McDonald, "NSIS Signaling Layer Protocol (NSLP) for Quality-of-Service Signaling", RFC 5974, October 2010.
- [RSVP-KEYING] Behringer, M., Faucheur, F., and B. Weis, "Applicability of Keying Methods for RSVP Security", Work in Progress, September 2010.

Authors' Addresses

Bruce Davie
Cisco Systems, Inc.
1414 Mass. Ave.
Boxborough, MA 01719
USA

EMail: bsd@cisco.com

Francois Le Faucheur
Cisco Systems, Inc.
Village d'Entreprise Green Side - Batiment T3
400, Avenue de Roumanille
Biot Sophia-Antipolis 06410
France

EMail: flefauch@cisco.com

Ashok Narayanan
Cisco Systems, Inc.
1414 Mass. Ave.
Boxborough, MA 01719
USA

EMail: ashokn@cisco.com