

Internet Engineering Task Force (IETF)
Request for Comments: 7178
Category: Standards Track
ISSN: 2070-1721

D. Eastlake 3rd
Huawei
V. Manral
Ionos Corp.
Y. Li
S. Aldrin
Huawei
D. Ward
Cisco
May 2014

Transparent Interconnection of Lots of Links (TRILL): RBridge Channel Support

Abstract

This document specifies a general channel mechanism for sending messages, such as Bidirectional Forwarding Detection (BFD) messages, between Routing Bridges (RBridges) and between RBridges and end stations in an RBridge campus through extensions to the Transparent Interconnection of Lots of Links (TRILL) protocol.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 5741.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <http://www.rfc-editor.org/info/rfc7178>.

Copyright Notice

Copyright (c) 2014 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<http://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	3
1.1. RBridge Channel Requirements	3
1.2. Relation to the MPLS Generic Associated Channel	4
1.3. Terminology	4
2. Inter-RBridge Channel Messages	4
2.1. RBridge Channel Message Inner Frame	6
2.1.1. RBridge Channel Header	6
2.1.2. Inner Ethernet Header	8
2.1.3. Inner.VLAN Tag	8
2.2. TRILL Header for RBridge Channel Messages	9
2.3. Ethernet Link Header and Trailer	10
2.4. Special Transmission and Rate Considerations	10
3. Processing RBridge Channel TRILL Data Messages	11
3.1. Processing the RBridge Channel Header	12
3.2. RBridge Channel Errors	13
4. Native RBridge Channel Frames	14
5. Indicating Support for RBridge Channel Protocols	16
6. Congestion Considerations	16
7. Allocation Considerations	17
7.1. IANA Considerations	17
7.2. IEEE Registration Authority Considerations	18
8. Security Considerations	18
9. References	19
9.1. Normative References	19
9.2. Informative References	20
10. Acknowledgments	20

1. Introduction

RBridge campuses provide transparent least-cost forwarding using the Transparent Interconnection of Lots of Links (TRILL) protocol that builds on Intermediate System to Intermediate System (IS-IS) routing [IS-IS] [RFC1195] [RFC7176]. Devices that implement TRILL are called Routing Bridges (RBridges) or TRILL Switches. However, the TRILL base protocol standard [RFC6325] provides only for TRILL Data messages and TRILL IS-IS messages.

This document specifies a general channel mechanism for the transmission of other messages within an RBridge campus, such as BFD [RFC5880] messages, (1) between RBridges and end stations that are directly connected on the same link and (2) between RBridges. This mechanism supports a requirement to be able to operate with minimal configuration.

1.1. RBridge Channel Requirements

It is anticipated that various protocols operating at the TRILL layer will be desired in RBridge campuses. For example, there is a need for rapid-response continuity checking with a protocol such as BFD [RFC5880] [RFC5882] and for a variety of optional reporting.

To avoid the requirement to design and specify a way to carry each such protocol, this document specifies a general channel for sending messages between RBridges in a campus at the TRILL level by extending the TRILL protocol. To accommodate a wide variety of protocols, this RBridge Channel facility accommodates all the regular modes of TRILL Data transmission including single- and multiple-hop unicast as well as VLAN-scoped multi-destination distribution.

To minimize any unnecessary burden on transit RBridges and to provide a more realistic test of network continuity and the like, RBridge Channel messages are designed to look like TRILL Data frames and, in the case of multi-hop messages, can normally be handled by transit RBridges as if they were TRILL Data frames; however, to enable processing at transit RBridges when required by particular messages, they may optionally use the RBridge Channel Alert TRILL extended header flags [RFC7179] that causes a transit RBridge implementing the flag to more closely examine a flagged frame.

This document also specifies a format for sending RBridge Channel messages between RBridges and end stations that are directly connected over a link, in either direction, when provided for by the protocol involved. For the most part, this format is the same as the format that is encapsulated by TRILL Data for inter-RBridge Channel messages.

Each particular protocol using the RBridge Channel facility will likely use only a subset of the facilities specified herein.

1.2. Relation to the MPLS Generic Associated Channel

The RBridge Channel is similar to the MPLS Generic Associated Channel specified in [RFC5586]. Instead of using a special MPLS label to indicate a special channel message, an RBridge Channel message is indicated by a special multicast Inner.MacDA and inner Ethertype (see Section 2.1).

1.3. Terminology

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC2119].

The terminology and acronyms of [RFC6325] are used in this document with the additions listed below.

BFD - Bidirectional Forwarding Detection

CHV - Channel Header Version

MH - Multi-Hop

NA - Native

SL - Silent

2. Inter-RBridge Channel Messages

Channel messages between RBridges are transmitted as TRILL Data frames. (For information on channel messages that can be transmitted between RBridges and end stations that are directly connected by a link, see Section 4.) Inter-RBridge Channel messages are identified as such by their Inner.MacDA, which is the All-Egress-RBridges multicast address, together with their inner Ethertype, which is the RBridge-Channel Ethertype. This Ethertype is part of and starts the RBridge Channel Header.

The diagram below shows the overall structure of an RBridge Channel Message frame on a link between two RBridges:

Frame Structure	Section of This Document

+-----+ Link Header	Section 2.3 if Ethernet link
+-----+ TRILL Header	Section 2.2
+-----+ Inner Ethernet Header	Section 2.1.2
+-----+ RBridge Channel Header	Section 2.1.1
+-----+ Protocol-Specific Payload	See specific channel protocol
+-----+ Link Trailer (FCS if Ethernet)	
+-----+	

Figure 1: RBridge Channel Frame Structure

Optionally, some channel messages may require examination of the frame by transit RBridges that support the RBridge Channel feature, to determine if they need to take any action. To indicate this, such messages use an RBridge Channel Alert extended TRILL Header flag as further described in Section 3 below.

Sections 2.1 and 2.2 describe the inner frame and the TRILL Header for frames sent in an RBridge Channel. As always, the outer Link Header and Link Trailer are whatever is needed to get a TRILL Data frame to the next-hop RBridge, depending on the technology of the link, and can change with each hop for multi-hop messages. Section 2.3 describes the outer Link Header for Ethernet links, and Section 2.4 discusses some special considerations for the first hop transmission of RBridge Channel messages.

Section 3 describes some details of RBridge Channel message processing. Section 4 provides the specifications for native RBridge Channel frames between RBridges and end stations that are directly connected over a link. Section 5 describes how support for RBridge Channel protocols is indicated. And Sections 6, 7, and 8 give congestion, allocation (IANA and IEEE), and security considerations respectively.

2.1. RBridge Channel Message Inner Frame

The encapsulated inner frame within an RBridge Channel message frame is as shown below.

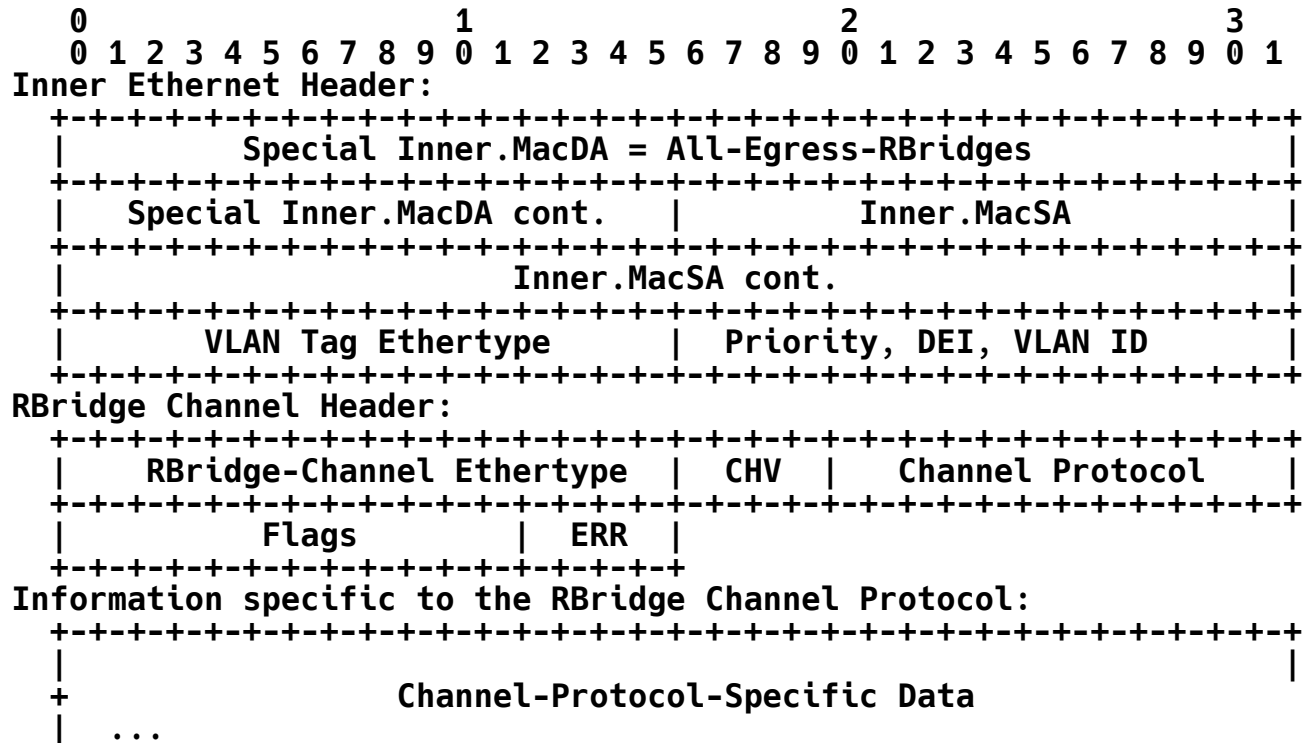


Figure 2: RBridge Channel Inner Frame Header Fields

The Channel-Protocol-Specific Data contains the information related to the specific channel protocol used in the channel message. Details of that data are outside the scope of this document, except in the case of the RBridge Channel Error protocol specified in Section 3.2.

2.1.1. RBridge Channel Header

As shown in Figure 2, the RBridge Channel Header starts with the RBridge-Channel Ethertype (see Section 7.2). Following that is a four-byte quantity with four sub-fields as follows:

CHV: A 4-bit field that gives the RBridge Channel Header Version. This document specifies version zero.

Channel Protocol: A 12-bit unsigned integer that specifies the particular RBridge Channel protocol to which the message applies.

Flags: Provides 12 bits of flags as described below.

ERR: A 4-bit unsigned integer used in connection with error reporting at the RBridge Channel level as described in Section 3.

The flag bits are numbered from 0 to 11 as shown below.




Figure 3: Channel Header Flag Bits

Bit 0: The SL or Silent bit, the high-order bit in network order. If it is a one, it suppresses RBridge Channel Error messages (see Section 3).

Bit 1: The MH or Multi-Hop bit. It is used to inform the destination RBridge protocol that the message may be multi-hop (MH=1) or was intended to be one-hop only (MH=0).

Bit 2: The NA or Native bit. It is used as described in Section 4.

Reserved: Bits reserved for future specification that MUST be sent as zero and ignored on receipt.

The RBridge Channel Protocol field specifies the protocol that the channel message relates to. The initial defined value is listed below.

Protocol	Name - Section of This Document
0x001	RBridge Channel Error - Section 3

IANA Considerations for RBridge Channel protocol numbers are provided in Section 7. These include provisions for Private Use protocol numbers. Because different uses of Private Use RBridge Channel protocol numbers may conflict, such use MUST be within a private network. It is the responsibility of the private network manager to avoid conflicting use of these code points and unacceptable burdens within the private network from their use.

2.1.2. Inner Ethernet Header

The special Inner.MacDA is the All-Egress-RBridges multicast Media Access Control (MAC) address to signal that the frame is intended for the egress (decapsulating) RBridge itself (or the egress RBridges themselves if the frame is multi-destination). (This address is called the All-ESADI-RBridges address in [RFC6325].) The RBridge-Channel Ethertype indicates that the frame is an RBridge Channel message. The only other Ethertype currently specified for use with the All-Egress-RBridges Inner.MacDA is L2-IS-IS to indicate an ESADI frame [RFC6325]. In the future, additional Ethernets may be specified for use with the All-Egress-RBridges multicast address.

The RBridge originating the channel message selects the Inner.MacSA. The Inner.MacSA MUST be set by the originating RBridge to a MAC address unique within the campus owned by the originating RBridge. This MAC address can be considered, in effect, the MAC address of a virtual internal end station that handles the RBridge Channel frames originated by or destined for that RBridge. It MAY be the same as the Inner.MacSA used by the RBridge when it originates ESADI frames [RFC6325].

2.1.3. Inner.VLAN Tag

As with all frames formatted to be processed as a TRILL Data frame, an Inner.VLAN tag is present. Use of a VLAN tag Ethertype other than 0x8100 or stacked tags is beyond the scope of this document but is an obvious extension.

Multi-destination RBridge Channel messages are, like all multi-destination TRILL Data messages, VLAN scoped so the Inner.VLAN ID MUST be set to the VLAN of interest. To the extent that distribution tree pruning is in effect in the campus, such channel messages may only reach RBridges advertising that they have connectivity to that VLAN.

For channel messages sent as known unicast TRILL Data frames, the default value for the Inner.VLAN ID is VLAN 1, but particular RBridge Channel protocols MAY specify other values.

The Inner.VLAN also specifies a three-bit frame priority for which the following recommendations apply:

1. For one-hop channel messages critical to network connectivity, such as one-hop BFD for rapid link-failure detection in support of TRILL IS-IS, the RECOMMENDED priority is 7.

2. For single and multi-hop unicast channel messages important to network operation but not critical for connectivity, the RECOMMENDED priority is 6.
3. For other unicast channel messages and all multi-destination channel messages, it is RECOMMENDED that the default priority zero be used. In any case, priorities higher than 5 SHOULD NOT be used for such frames.

There is one additional bit in a VLAN tag value between the 12-bit VLAN ID and 3-bit priority, the Drop Eligibility Indicator (DEI) [RFC7180]. It is RECOMMENDED that this bit be zero for the first two categories of channel messages listed immediately above. The setting of this bit for channel messages in the third category may be dependent on the channel protocol and no general recommendation is made for that case.

2.2. TRILL Header for RBridge Channel Messages

After the outer Link Header (that, for an Ethernet link, ends with the TRILL Ethertype) and before the encapsulated frame, the channel message's TRILL Header initially appears as follows:




Figure 4: RBridge Channel TRILL Header Fields

The TRILL Header version (V) MUST be zero; the R bits are reserved; the M bit is set appropriately as the channel message is to be forwarded as known destination unicast (M=0) or multi-destination (M=1), regardless of the fact that the Inner.MacDA is always the All-Egress-RBridges multicast address; and Op-Len is set appropriately for the length of the TRILL Header extensions area, if any, all as specified in [RFC6325].

When an RBridge Channel message is originated, the Hop Count field defaults to the maximum value, 0x3F, but particular RBridge Channel protocols MAY specify other values. For messages sent a known number of hops, such as one-hop messages or a two-hop self-addressed message intended to loop back through an immediate neighbor RBridge, setting the Hop Count field in the TRILL Header to the maximum value and checking its value on receipt provides an additional validity check

as discussed in [RFC5082], where this type of field is referred to as "TTL" or "Hop Limit".

The RBridge originating a channel message places a nickname that it holds in the Ingress Nickname field.

There are several cases for the Egress Nickname field. If the channel message is multi-destination, then the Egress Nickname designates the distribution tree to use. If the channel message is a multi-hop unicast message, then the Egress Nickname is a nickname of the target RBridge; this includes the special case of a message intended to loop back from an immediate neighbor where the originator places one of its own nicknames in both the Ingress Nickname and Egress Nickname fields. If the channel message is a one-hop unicast message, there are two possibilities for the Egress Nickname.

- o The Egress Nickname can be set to a nickname of the target neighbor RBridge.
- o The special nickname Any-RBridge may be used. RBridges supporting the RBridge Channel facility MUST recognize the Any-RBridge special nickname and accept TRILL Data frames having that value in the Egress Nickname field as being sent to them as the egress. Thus, for such RBridges, using this egress nickname guarantees processing by an immediate neighbor regardless of the state of nicknames.

2.3. Ethernet Link Header and Trailer

An RBridge Channel frame has the usual Link Header and Link Trailer for a TRILL Data frame depending on the type of link on which it is sent.

For an Ethernet link [RFC6325], the Outer.MacSA is the MAC address of the port from which the frame is sent. The Outer.MacDA is the MAC address of the next-hop RBridge port for unicast RBridge Channel messages or the All-RBridges multicast address for multi-destination RBridge Channel messages. The Outer.VLAN tag specifies the designated VLAN for that hop, and the priority should be the same as in the Inner.VLAN tag; however, the output port may have been configured to strip VLAN tags, in which case no Outer.VLAN tag appears on the wire. And the Link Trailer is the Ethernet FCS.

2.4. Special Transmission and Rate Considerations

If a multi-hop RBridge Channel message is received by an RBridge, the criteria and method of forwarding it are the same as for any TRILL Data frame. If it is so forwarded, it will be on a link that was

included in the routing topology because it was in the Report state as specified in [RFC7177].

However, special considerations apply to single-hop messages because, for some RBridge Channel protocols, it may be desirable to send RBridge Channel messages over a link that is not yet fully up. In particular, it is permissible, if specified by the particular channel protocol, for the source RBridge that has created an RBridge Channel message to attempt to transmit it to a next-hop RBridge when the link is in the Detect or 2-Way state, as specified in [RFC7177], as well as when it is in the Report state. Such messages can also be sent on point-to-point links that are not in the Up state.

RBridge Channel messages represent a burden on the RBridges, and links in a campus and should be rate limited, especially if they are sent as high priority, multi-destination, or multi-hop frames or have an RBridge Channel Alert extended header flag set. See Section 6, "Congestion Considerations".

3. Processing RBridge Channel TRILL Data Messages

RBridge Channel TRILL Data messages are designed to look like and, to the extent practical, be forwarded as regular TRILL Data frames. On receiving a channel message, an RBridge performs the usual initial tests on the frame and makes the same forwarding and/or decapsulation decisions as for a regular TRILL Data frame [RFC6325] with the following exceptions for RBridges implementing the RBridge Channel facility:

1. An RBridge implementing the RBridge Channel facility **MUST** recognize the Any-RBridge egress nickname in TRILL Data frames, decapsulating such frames if they meet other checks. (Such a frame cannot be a valid multi-destination frame because the Any-RBridge nickname is not a valid distribution tree root.)
2. If an RBridge Channel Alert extended header flag is set, then the RBridge **MUST** process the RBridge Channel message as described below even if it is not egressing the frame. If it is egressing the frame, then no additional processing beyond egress processing is needed even if an RBridge Channel Alert flag is set.
3. On decapsulation, the special Inner.MacDA value of All-Egress-RBridges **MUST** be recognized to trigger checking the Inner.Ethertype and processing as an RBridge Channel message if that Ethertype is RBridge-Channel.

RBridge Channel messages **SHOULD** only be sent to RBridges that advertise support for the channel protocol involved as described in Section 5.

All RBridges supporting the RBridge Channel facility **MUST** recognize the RBridge-Channel inner Ethertype.

3.1. Processing the RBridge Channel Header

Knowing that it has an RBridge Channel message, the egress RBridge, and any transit RBridge if an RBridge Channel Alert bit is set in the TRILL Header, looks at the CHV (RBridge Channel Header Version) and Channel Protocol fields.

If any of the following conditions occur at an egress RBridge, the frame is not processed, an error may be generated as specified in Section 3.2, and the frame is discarded. The behavior is the same if the frame is being processed at a transit RBridge because the RBridge Critical Channel Alert flag is set [RFC7179]. However, if these conditions are detected at a transit RBridge examining the message because the RBridge Non-critical Channel Alert flag is set [RFC7179] but the RBridge Critical Channel Alert flag is not set, no error is generated, and the frame is still forwarded normally.

Error Conditions:

1. The Ethertype is not RBridge-Channel and not any other Ethertype known to the RBridge as usable with the All-Egress-RBridges Inner.MacDA, or the frame is so short that the Ethertype is truncated.
2. The CHV field is non-zero, or the frame is so short that the version zero Channel Header is truncated.
3. The Channel Protocol field is a reserved value or a value unknown to the processing RBridge.
4. The ERR field is non-zero, and Channel Protocol is a value other than 0x001.
5. The RBridge Channel Header NA flag is set to one, indicating that the frame should have been received as a native frame rather than a TRILL Data frame.

If the CHV field and NA flag are zero and the processing RBridge recognizes the Channel Protocol value, it processes the message in accordance with that channel protocol. The processing model is as if the received frame starting with and including the TRILL Header is

delivered to the Channel protocol along with a flag indicating whether this is (a) transit RBridge processing due to an RBridge Channel Alert flag being set or (b) egress processing.

Errors within a recognized Channel Protocol are handled by that channel protocol itself and do not produce RBridge Channel Error frames.

3.2. RBridge Channel Errors

A variety of problems at the RBridge Channel level cause the return of an RBridge Channel Error frame unless one of the following apply: (a) the "SL" (Silent) flag is a one in the channel message for which the problem was detected, (b) the processing is due to the RBridge Non-critical Channel Alert flag being set, (c) the frame in error appears, itself, to be an RBridge Channel Error frame (has a non-zero ERR field or a Channel Protocol of 0x001), or (d) the error is suppressed due to rate limiting.

An RBridge Channel Error frame is a multi-hop unicast RBridge Channel message with the Ingress Nickname set to a nickname of the RBridge detecting the error and the Egress Nickname set to the value of the Ingress Nickname in the channel message for which the error was detected. No per-hop transit processing is specified for such error frames, so the RBridge Channel Alert extended header flags SHOULD, if an extension is present, be set to zero. The SL and MH flags SHOULD be set to one; the NA flag MUST be zero; and the ERR field MUST be non-zero as described below. For the protocol-specific data area, an RBridge Channel Message Error frame has at least the first 256 bytes (or less if less are available) of the erroneous decapsulated channel message starting with the TRILL Header. (Note: The TRILL Header does not include the TRILL Ethertype that is part of the Link Header on Ethernet links.)

The following values for ERR are specified:

ERR	RBridge Channel Error Code Meaning
---	-----
0	No error
1	Frame too short (truncated Ethertype or Channel Header)
2	Unrecognized Ethertype
3	Unimplemented value of CHV
4	Wrong value of NA flag
5	Channel Protocol is reserved or unimplemented
6-14	Unassigned (see Section 7)
15	Reserved (see Note)

Note: Intended to be allocated by Standards Action for an error code expansion feature when it appears likely that all other available error codes are being allocated.

All RBridges implementing the RBridge Channel feature **MUST** recognize the RBridge Channel Error protocol value (0x001). They **MUST NOT** generate an RBridge Channel Error message in response to an RBridge Channel Error message, that is, a channel message with a protocol value of 0x001 or with a non-zero ERR field.

4. Native RBridge Channel Frames

Other sections of this document specify non-native RBridge Channel messages and their processing, that is, RBridge Channel messages formatted as TRILL Data frames and sent between RBridges. This section specifies the differences for native RBridge Channel messages.

If provided for by the channel protocol involved, native RBridge Channel messages may be sent between end stations and RBridges that are directly connected over a link, in either direction. On an Ethernet link, such native frames have the RBridge-Channel Ethertype and are like the encapsulated frame inside an RBridge Channel message except as follows:

1. TRILL does not require the presence of a VLAN tag on such native RBridge Channel frames. However, port configuration, link characteristics, or the channel protocol involved may require such tagging.
2. If the frame is unicast, the destination MAC address is the unicast MAC address of the RBridge or end-station port that is its intended destination. If the frame is multicast by an end station to all the RBridges on a link that support an RBridge Channel protocol using this transport, the destination MAC address is the All-Edge-RBridges multicast address (see Section 7). A native RBridge Channel frame received at an ingress RBridge is discarded if its destination MAC address is neither the unicast address of the port nor the multicast address All-Edge-RBridges. If the frame is multicast by an RBridge to all the devices that TRILL considers to be end stations on a link and that support an RBridge Channel protocol using this transport, the destination MAC address is the TRILL-End-Stations multicast address (see Section 7). A native RBridge Channel frame received at an end station is discarded if its destination MAC address is neither the unicast address of the port nor the multicast address TRILL-End-Stations.

3. The RBridge-Channel outer Ethertype must be present. In the future, there may be other protocols using the All-Edge-RBridges and/or TRILL-End-Stations multicast addresses on native frames distinguished by different Ethernets.
4. The NA or Native bit in the RBridge Channel Header flags MUST be a one.
5. There might be additional tags present between the Outer.MacDA, Outer.MacSA pair, and the RBridge-Channel Ethertype.

The RBridge Channel protocol number space for native RBridge Channel messages and TRILL Data formatted RBridge Channel messages is the same. If provided for by the channel protocol involved, the receipt of a native RBridge Channel frame MAY lead to the generation and transmission of one or more Inter-RBridge Channel frames. The decapsulation and processing of a TRILL Data RBridge Channel frame MAY, if provided for by the channel protocol involved, result in the sending of one or more native RBridge Channel frames to one or more end stations. Thus, there could be an RBridge Channel protocol that involved an RBridge Channel message sent (1) from an origin RBridge where the message is created, (2) through one or more transit RBridges, and (3) from a final RBridge as a native RBridge Channel message to an end station (or the reverse of such a three-part path); however, to do this, the RBridge Channel protocol involved must be implemented at the RBridge where the channel message is changed between a native frame and a TRILL Data format frame, and that RBridge must change the channel message itself, at a minimum complementing the NA flag in the Channel Header and making appropriate MAC address changes.

An erroneous native channel message results in a native RBridge Channel Error message under the same conditions for which a TRILL Data RBridge Channel message would result in a TRILL Data RBridge Channel Error message. However, in a native RBridge Channel Error message, the NA flag MUST be one. Also, since there is no TRILL Header in native RBridge Channel protocol frames, the beginning part of the frame in which the error was detected that is included in native RBridge Channel Error frames starts with the RBridge Channel Header (including the RBridge-Channel Ethertype). The destination MAC address of such error messages is set to the source MAC address of the native RBridge Channel message that was in error.

There is no mechanism to stop end stations from directly exchanging native RBridge Channel messages, but such usage is beyond the scope of this document.

5. Indicating Support for RBridge Channel Protocols

Support for RBridge Channel protocols is indicated by the presence of one or more TLVs and/or sub-TLVs in an RBridge's Link State PDU (LSP) as documented in [RFC7176].

RBridge Channel protocols 0 and 0xFF are reserved, and protocol 1, the RBridge Channel Error protocol, **MUST** be implemented as part of the RBridge Channel feature. Thus, if an RBridge supports the RBridge Channel feature, it should be advertising support for protocol 1 and not advertising support for protocols 0 or 0xFF in its LSP. However, indication of support or non-support for RBridge Channel protocol 1 is ignored on receipt, and support for it is always assumed if support for any RBridge Channel is indicated in the RBridge's LSP.

6. Congestion Considerations

The bandwidth resources used by RBridge Channel protocols are recommended to be small compared to the total bandwidth of the links they traverse. When doing network planning, the bandwidth requirements for TRILL Data, TRILL IS-IS, TRILL ESADI, RBridge Channel protocol traffic, and any other link-local traffic need to be taken into account.

Specifications for particular RBridge Channel protocols **MUST** consider congestion and bandwidth usage implications and provide guidance on bandwidth or packet-frequency management. RBridge Channel protocols can have built-in bandwidth management in their protocols. Outgoing channel messages **SHOULD** be rate-limited, by configuring the underlying protocols or otherwise, to prevent aggressive connectivity verification or other applications consuming excessive bandwidth, causing congestion, or becoming denial-of-service attacks.

If these conditions cannot be followed, an adaptive loss-based scheme **SHOULD** be applied to congestion-control outgoing RBridge Channel traffic, so that it competes fairly, taking into account packet priorities and drop eligibility as indicated in the Inner.VLAN, with TCP or similar traffic within an order of magnitude. One method of determining an acceptable bandwidth for RBridge Channel traffic is described in [RFC5348]; other methods exist. For example, bandwidth or packet-frequency management can include any of the following: a negotiation of transmission interval/rate such as that provided in BFD [RFC5880], a throttled transmission rate on "congestion detected" situations, a gradual ramp-up after shutdown due to congestion and until basic connectivity is verified, and other mechanisms.

Connectivity-checking applications such as BFD [RFC5880] SHOULD be rate-limited to below 5% of the bitrate of the associated link or links. For this purpose, the mean or sustained bitrate of the link or links is used.

Incoming RBridge Channel messages MAY be rate-limited as a protection against denial-of-service attacks. This throttling of incoming messages SHOULD honor packet priorities and drop eligibility indications as indicated in the Inner.VLAN, preferentially discarding drop-eligible and lower-priority packets.

7. Allocation Considerations

The following subsections give IANA and IEEE allocation considerations. In this document, the allocation procedure specifications are as defined in [RFC5226].

7.1. IANA Considerations

IANA has allocated a previously unassigned TRILL Nickname as follows:

Any-RBridge	0xFFC0
-------------	--------

IANA has added "All-Egress-RBridges" to the TRILL Parameter Registry as an alternative name for the "All-ESADI-RBridges" multicast address.

IANA has allocated two previously unassigned TRILL multicast addresses as follows:

TRILL-End-Stations	01-80-C2-00-00-45
All-Edge-RBridges	01-80-C2-00-00-46

IANA has created an additional sub-registry in the TRILL Parameter Registry for RBridge Channel Protocols, with initial contents as follows:

Protocol -----	Description -----	Reference -----
0x000	Reserved; not to be allocated	(This document)
0x001	RBridge Channel Error	(This document)
0x002-0x0FF	Unassigned (1)	
0x100-0xFF7	Unassigned (2)	
0xFF8-0xFFE	Reserved for Private Use	
0xFFF	Reserved; not to be allocated	(This document)

- (1) RBridge Channel protocol code points from 0x002 to 0x0FF require a Standards Action, as modified by [RFC7120], for allocation.
- (2) RBridge Channel protocol code points from 0x100 to 0xFF7 are RFC Required to allocate a single value or IESG Approval to allocate multiple values.

IANA has created an additional sub-registry in the TRILL Parameter Registry for RBridge Channel Header Flags with initial contents as follows:

Flag Bit	Mnemonic	Allocation
-----	-----	-----
0	SL	Silent
1	MH	Multi-hop
2	NA	Native
3-11	-	Unassigned

Allocation of an RBridge Channel Header Flag is based on IETF Review.

IANA has created an additional sub-registry in the TRILL Parameter Registry for RBridge Channel Error Codes with initial contents as listed in Section 3.2 above and with available values allocated by Standards Action as modified by [RFC7120].

7.2. IEEE Registration Authority Considerations

The IEEE Registration Authority has assigned the Ethertype 0x8946 for TRILL RBridge Channel.

8. Security Considerations

No general integrity, authentication, or encryption mechanisms are provided herein for RBridge Channel messages. If these services are required for a particular RBridge Channel protocol, they MUST be supplied by that channel protocol. See, for example, the BFD Authentication mechanism [RFC5880].

See [RFC6325] for general TRILL security considerations. As stated therein, no protection is provided by TRILL against forging of the Ingress Nickname in a TRILL Data formatted channel message or the Outer.MacSA in a native RBridge Channel frame on an Ethernet link. This may result in misdirected return responses or error messages. However, link-level security protocols may be used to authenticate the origin station on a link and protect against attacks on links. See also Section 6 concerning congestion.

If indications of RBridge Channel Protocol support are improperly absent from an RBridge's LSP, it could deny all RBridge Channel services, for example, some BFD services, for the RBridge in question. If a particular RBridge Channel protocol is incorrectly not advertised as supported, it could deny the service of that channel protocol to the RBridge in question.

Incorrect indication of RBridge Channel Protocol support or incorrect assertion of support for a channel protocol could encourage RBridge Channel messages to be sent to an RBridge that does not support the channel feature or the particular channel protocol used. The inner frame of such messages could be decapsulated and that inner frame could be sent out all ports that are Appointed Forwarders for the frame's Inner.VLAN. However, this is unlikely to cause much harm; in particular, there are two possibilities as follows: (a) if end stations do not recognize the RBridge-Channel Ethertype of the frame, they will drop it, and (b) if end stations do recognize the RBridge-Channel Ethertype and the channel protocol indicated in the frame, they should refuse to process the frame due to an incorrect value of the RBridge Channel Header NA flag.

9. References

9.1. Normative References

- [IS-IS] International Organization for Standardization, "Intermediate System to Intermediate System intra-domain routing information exchange protocol for use in conjunction with the protocol for providing the connectionless-mode network service (ISO 8473)", Second Edition, November 2002.
- [RFC1195] Callon, R., "Use of OSI IS-IS for routing in TCP/IP and dual environments", RFC 1195, December 1990.
- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, March 1997.
- [RFC5226] Narten, T. and H. Alvestrand, "Guidelines for Writing an IANA Considerations Section in RFCs", BCP 26, RFC 5226, May 2008.
- [RFC5348] Floyd, S., Handley, M., Padhye, J., and J. Widmer, "TCP Friendly Rate Control (TFRC): Protocol Specification", RFC 5348, September 2008.

- [RFC6325] Perlman, R., Eastlake 3rd, D., Dutt, D., Gai, S., and A. Ghanwani, "Routing Bridges (RBridges): Base Protocol Specification", RFC 6325, July 2011.
- [RFC7120] Cotton, M., "Early IANA Allocation of Standards Track Code Points", BCP 100, RFC 7120, January 2014.
- [RFC7176] Eastlake 3rd, D., Senevirathne, T., Ghanwani, A., Dutt, D., and A. Banerjee, "Transparent Interconnection of Lots of Links (TRILL) Use of IS-IS", RFC 7176, May 2014.
- [RFC7177] Eastlake 3rd, D., Perlman, R., Ghanwani, A., Yang, H., and V. Manral, "Transparent Interconnection of Lots of Links (TRILL): Adjacency", RFC 7177, May 2014.
- [RFC7179] Eastlake 3rd, D., Ghanwani, A., Manral, V., Li, Y., and C. Bestler, "Transparent Interconnection of Lots of Links (TRILL): Header Extension", RFC 7179, May 2014.

9.2. Informative References

- [RFC5082] Gill, V., Heasley, J., Meyer, D., Savola, P., Ed., and C. Pignataro, "The Generalized TTL Security Mechanism (GTSM)", RFC 5082, October 2007.
- [RFC5586] Bocci, M., Ed., Vigoureux, M., Ed., and S. Bryant, Ed., "MPLS Generic Associated Channel", RFC 5586, June 2009.
- [RFC5880] Katz, D. and D. Ward, "Bidirectional Forwarding Detection (BFD)", RFC 5880, June 2010.
- [RFC5882] Katz, D. and D. Ward, "Generic Application of Bidirectional Forwarding Detection (BFD)", RFC 5882, June 2010.
- [RFC7180] Eastlake 3rd, D., Zhang, M., Ghanwani, A., Manral, V., and A. Banerjee, "Transparent Interconnection of Lots of Links (TRILL): Clarifications, Corrections, and Updates", RFC 7180, May 2014.

10. Acknowledgments

The authors gratefully acknowledge the comments and contributions of the follows, listed in alphabetic order: Stewart Bryant, Somnath Chatterjee, Adrian Farrel, Stephen Farrell, Miguel A. Garcia, Anoop Ghanwani, Brian Haberman, Rakesh Kumar, Barry Leiba, and Tissa Senevirathne.

Authors' Addresses

Donald Eastlake 3rd
Huawei R&D USA
155 Beaver Street
Milford, MA 01757
USA
Phone: +1-508-333-2270
EMail: d3e3e3@gmail.com

Vishwas Manral
Ionos Corp.
4100 Moorpark Ave.
San Jose, CA 95117
USA
EMail: vishwas@ionosnetworks.com

Yizhou Li
Huawei Technologies
101 Software Avenue,
Nanjing 210012
China
Phone: +86-25-56622310
EMail: liyizhou@huawei.com

Sam Aldrin
Huawei Technologies
2330 Central Expressway
Santa Clara, CA 95050
USA
Phone: +1-408-330-5000
EMail: sam.aldrin@huawei.com

Dave Ward
Cisco Systems
170 W. Tasman Drive
San Jose, CA 95134
USA
EMail: dward@cisco.com