

Internet Engineering Task Force (IETF)
Request for Comments: 8349
Obsoletes: 8022
Category: Standards Track
ISSN: 2070-1721

L. Lhotka
CZ.NIC
A. Lindem
Cisco Systems
Y. Qu
Huawei
March 2018

A YANG Data Model for Routing Management (NMDA Version)

Abstract

This document specifies three YANG modules and one submodule. Together, they form the core routing data model that serves as a framework for configuring and managing a routing subsystem. It is expected that these modules will be augmented by additional YANG modules defining data models for control-plane protocols, route filters, and other functions. The core routing data model provides common building blocks for such extensions -- routes, Routing Information Bases (RIBs), and control-plane protocols.

The YANG modules in this document conform to the Network Management Datastore Architecture (NMDA). This document obsoletes RFC 8022.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 7841.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <https://www.rfc-editor.org/info/rfc8349>.

Copyright Notice

Copyright (c) 2018 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	4
2. Terminology and Notation	4
2.1. Glossary of New Terms	5
2.2. Tree Diagrams	6
2.3. Prefixes in Data Node Names	6
3. Objectives	6
4. The Design of the Core Routing Data Model	7
4.1. System-Controlled and User-Controlled List Entries	8
5. Basic Building Blocks	9
5.1. Routes	9
5.2. Routing Information Base (RIB)	10
5.3. Control-Plane Protocol	11
5.3.1. Routing Pseudo-Protocols	11
5.3.2. Defining New Control-Plane Protocols	11
5.4. Parameters of IPv6 Router Advertisements	12
6. Interactions with Other YANG Modules	13
6.1. Module "ietf-interfaces"	13
6.2. Module "ietf-ip"	14
7. Routing Management YANG Module	15
8. IPv4 Unicast Routing Management YANG Module	29
9. IPv6 Unicast Routing Management YANG Module	37
9.1. IPv6 Router Advertisements Submodule	45
10. IANA Considerations	56
11. Security Considerations	57
12. References	58
12.1. Normative References	58
12.2. Informative References	60
Appendix A. The Complete Schema Tree	61
Appendix B. Minimum Implementation	66
Appendix C. Example: Adding a New Control-Plane Protocol	67
Appendix D. Data Tree Example	70
Appendix E. NETCONF Get Data Reply Example	77
Acknowledgments	80
Authors' Addresses	80

1. Introduction

This document specifies the following YANG modules:

- o The "ietf-routing" module provides generic components of a routing data model.
- o The "ietf-ipv4-unicast-routing" module augments the "ietf-routing" module with additional data specific to IPv4 unicast.
- o The "ietf-ipv6-unicast-routing" module augments the "ietf-routing" module with additional data specific to IPv6 unicast. Its submodule, "ietf-ipv6-router-advertisements", also augments the "ietf-interfaces" [RFC8343] and "ietf-ip" [RFC8344] modules with IPv6 router configuration variables required by [RFC4861].

These modules together define the core routing data model, which is intended as a basis for future data model development covering more-sophisticated routing systems. While these three modules can be directly used for simple IP devices with static routing (see Appendix B), their main purpose is to provide essential building blocks for more-complicated data models involving multiple control-plane protocols, multicast routing, additional address families, and advanced functions such as route filtering or policy routing. To this end, it is expected that the core routing data model will be augmented by numerous modules developed by various IETF working groups.

The YANG modules in this document conform to the Network Management Datastore Architecture (NMDA) [RFC8342]. This document obsoletes RFC 8022 [RFC8022].

2. Terminology and Notation

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in BCP 14 [RFC2119] [RFC8174] when, and only when, they appear in all capitals, as shown here.

The following terms are defined in [RFC8342]:

- o client
- o server
- o configuration

- o system state
- o operational state
- o intended configuration

The following terms are defined in [RFC7950]:

- o action
- o augment
- o container
- o data model
- o data node
- o feature
- o leaf
- o list
- o mandatory node
- o module
- o presence container
- o schema tree
- o RPC (Remote Procedure Call) operation

2.1. Glossary of New Terms

core routing data model: YANG data model comprising "ietf-routing", "ietf-ipv4-unicast-routing", and "ietf-ipv6-unicast-routing" modules.

direct route: A route to a directly connected network.

Routing Information Base (RIB): An object containing a list of routes, together with other information. See Section 5.2 for details.

system-controlled entry: An entry in a list in the operational state ("config false") that is created by the system independently of what has been explicitly configured. See Section 4.1 for details.

user-controlled entry: An entry in a list in the operational state ("config false") that is created and deleted as a direct consequence of certain configuration changes. See Section 4.1 for details.

2.2. Tree Diagrams

Tree diagrams used in this document follow the notation defined in [RFC8340].

2.3. Prefixes in Data Node Names

In this document, names of data nodes, actions, and other data model objects are often used without a prefix, as long as it is clear from the context in which YANG module each name is defined. Otherwise, names are prefixed using the standard prefix associated with the corresponding YANG module, as shown in Table 1.

Prefix	YANG module	Reference
if	ietf-interfaces	[RFC8343]
ip	ietf-ip	[RFC8344]
rt	ietf-routing	Section 7
v4ur	ietf-ipv4-unicast-routing	Section 8
v6ur	ietf-ipv6-unicast-routing	Section 9
yang	ietf-yang-types	[RFC6991]
inet	ietf-inet-types	[RFC6991]

Table 1: Prefixes and Corresponding YANG Modules

3. Objectives

The initial design of the core routing data model was driven by the following objectives:

- o The data model should be suitable for the common address families -- in particular, IPv4 and IPv6 -- and for unicast and multicast routing, as well as Multiprotocol Label Switching (MPLS).
- o A simple IP routing system, such as one that uses only static routing, should be configurable in a simple way, ideally without any need to develop additional YANG modules.

- o On the other hand, the core routing framework must allow for complicated implementations involving multiple RIBs and multiple control-plane protocols, as well as controlled redistributions of routing information.
- o Because device vendors will want to map the data models built on this generic framework to their proprietary data models and configuration interfaces, the framework should be flexible enough to facilitate such mapping and accommodate data models with different logic.

4. The Design of the Core Routing Data Model

The core routing data model consists of three YANG modules and one submodule. The first module, "ietf-routing", defines the generic components of a routing system. The other two modules -- "ietf-ipv4-unicast-routing" and "ietf-ipv6-unicast-routing" -- augment the "ietf-routing" module with additional data nodes that are needed for IPv4 and IPv6 unicast routing, respectively. The "ietf-ipv6-unicast-routing" module has a submodule, "ietf-ipv6-router-advertisements", that augments the "ietf-interfaces" [RFC8343] and "ietf-ip" [RFC8344] modules with configuration variables for IPv6 Router Advertisements as required by [RFC4861].

Figure 1 shows abridged views of the hierarchies. See Appendix A for the complete data trees.

```

+--rw routing
  +--rw router-id?                yang:dotted-quad
  +--ro interfaces
  |   +--ro interface*    if:interface-ref
  +--rw control-plane-protocols
  |   +--rw control-plane-protocol* [type name]
  |   |   +--rw type                identityref
  |   |   +--rw name                string
  |   |   +--rw description?        string
  |   |   +--rw static-routes
  |   |   |   +--rw v4ur:ipv4
  |   |   |   |   ...
  |   |   |   +--rw v6ur:ipv6
  |   |   |   |   ...
  |   |   |   ...
  +--rw ribs
  |   +--rw rib* [name]
  |   |   +--rw name                string
  |   |   +--rw address-family?    identityref
  |   |   +--ro default-rib?        boolean {multiple-ribs}?
  |   |   +--ro routes
  |   |   |   +--ro route*
  |   |   |   |   ...
  |   |   +---x active-route
  |   |   |   +---w input
  |   |   |   |   +---w v4ur:destination-address?    inet:ipv4-address
  |   |   |   |   +---w v6ur:destination-address?    inet:ipv6-address
  |   |   |   +--ro output
  |   |   +---rw description?        string

```

Figure 1: Data Hierarchy

As can be seen from Figure 1, the core routing data model introduces several generic components of a routing framework: routes, RIBs containing lists of routes, and control-plane protocols. Section 5 describes these components in more detail.

4.1. System-Controlled and User-Controlled List Entries

The core routing data model defines several lists in the schema tree, such as "rib", that have to be populated with at least one entry in any properly functioning device, and additional entries may be configured by a client.

In such a list, the server creates the required item as a "system-controlled entry" in the operational state, i.e., inside read-only lists in the "routing" container.

An example can be seen in Appendix D: the `"/routing/ribs/rib"` list has two system-controlled entries -- `"ipv4-master"` and `"ipv6-master"`.

Additional entries called "user-controlled entries" may be created in the configuration by a client, e.g., via the Network Configuration Protocol (NETCONF). If the server accepts a configured user-controlled entry, then this entry also appears in the operational state version of the list.

Corresponding entries in both versions of the list (in the intended configuration and the operational state) [RFC8342] have the same value of the list key.

A client may also provide supplemental configuration of system-controlled entries. To do so, the client creates a new entry in the configuration with the desired contents. In order to bind this entry to the corresponding entry in the operational state, the key of the configuration entry has to be set to the same value as the key of the operational state entry.

Deleting a user-controlled entry from the intended configuration results in the removal of the corresponding entry in the operational state list. In contrast, if a client deletes a system-controlled entry from the intended configuration, only the extra configuration specified in that entry is removed; the corresponding operational state entry is not removed.

5. Basic Building Blocks

This section describes the essential components of the core routing data model.

5.1. Routes

Routes are basic elements of information in a routing system. The core routing data model defines only the following minimal set of route attributes:

- o "destination-prefix": address prefix specifying the set of destination addresses for which the route may be used. This attribute is mandatory.

- o "route-preference": an integer value (also known as "administrative distance") that is used for selecting a preferred route among routes with the same destination prefix. A lower value indicates a route that is more preferred.
- o "next-hop": determines the outgoing interface and/or next-hop address(es), or a special operation to be performed on a packet.

Routes are primarily system state and appear as entries in RIBs (Section 5.2), but they may also be found in configuration data -- for example, as manually configured static routes. In the latter case, configurable route attributes are generally a subset of attributes defined for RIB routes.

5.2. Routing Information Base (RIB)

Every implementation of the core routing data model manages one or more RIBs. A RIB is a list of routes complemented with administrative data. Each RIB contains only routes of one address family. An address family is represented by an identity derived from the "rt:address-family" base identity.

In the core routing data model, RIBs are represented as entries in the list "/routing/ribs/rib" in the operational state. The contents of RIBs are controlled and manipulated by control-plane protocol operations that may result in route additions, removals, and modifications. This also includes manipulations via the "static" and/or "direct" pseudo-protocols; see Section 5.3.1.

For every supported address family, exactly one RIB MUST be marked as the "default RIB", in which control-plane protocols place their routes by default.

Simple router implementations that do not advertise the "multiple-ribs" feature will typically create one system-controlled RIB per supported address family and mark it as the default RIB.

More-complex router implementations advertising the "multiple-ribs" feature support multiple RIBs per address family that can be used for policy routing and other purposes.

The following action (see Section 7.15 of [RFC7950]) is defined for the "rib" list:

- o active-route -- return the active RIB route for the destination address that is specified as the action's input parameter.

5.3. Control-Plane Protocol

The core routing data model provides an open-ended framework for defining multiple control-plane protocol instances, e.g., for Layer 3 routing protocols. Each control-plane protocol instance **MUST** be assigned a type, which is an identity derived from the "rt:control-plane-protocol" base identity. The core routing data model defines two identities for the "direct" and "static" pseudo-protocols (Section 5.3.1).

Multiple control-plane protocol instances of the same type **MAY** be configured.

5.3.1. Routing Pseudo-Protocols

The core routing data model defines two special routing protocol types -- "direct" and "static". Both are in fact pseudo-protocols, which means that they are confined to the local device and do not exchange any routing information with adjacent routers.

Every implementation of the core routing data model **MUST** provide exactly one instance of the "direct" pseudo-protocol type. It is the source of direct routes for all configured address families. Direct routes are normally supplied by the operating system kernel, based on the configuration of network interface addresses; see Section 6.2.

A pseudo-protocol of the type "static" allows for specifying routes manually. It **MAY** be configured in zero or multiple instances, although a typical configuration will have exactly one instance.

5.3.2. Defining New Control-Plane Protocols

It is expected that future YANG modules will create data models for additional control-plane protocol types. Such new modules will have to define the protocol-specific data nodes, and they will have to integrate into the core routing framework in the following way:

- o A new identity **MUST** be defined for the control-plane protocol, and its base identity **MUST** be set to "rt:control-plane-protocol" or to an identity derived from "rt:control-plane-protocol".

- o Additional route attributes MAY be defined, preferably in one place by means of defining a YANG grouping. The new attributes have to be inserted by augmenting the definitions of the node

`/rt:routing/rt:ribs/rt:rib/rt:routes/rt:route`

and possibly other places in the schema tree.

- o Data nodes for the new protocol can be defined by augmenting the "control-plane-protocol" data node under "/routing".

By using a "when" statement, the augmented data nodes specific to the new protocol SHOULD be made conditional and valid only if the value of "rt:type" or "rt:source-protocol" is equal to (or derived from) the new protocol's identity.

It is also RECOMMENDED that protocol-specific data nodes be encapsulated in an appropriately named container with presence. Such a container may contain mandatory data nodes that are otherwise forbidden at the top level of an augment.

The above steps are implemented by the example YANG module for the Routing Information Protocol (RIP); see Appendix C.

5.4. Parameters of IPv6 Router Advertisements

The YANG module "ietf-ipv6-router-advertisements" (Section 9.1), which is a submodule of the "ietf-ipv6-unicast-routing" module, augments the schema tree of IPv6 interfaces with definitions of the following variables as required by Section 6.2.1 of [RFC4861]:

- o send-advertisements
- o max-rtr-adv-interval
- o min-rtr-adv-interval
- o managed-flag
- o other-config-flag
- o link-mtu
- o reachable-time
- o retrans-timer
- o cur-hop-limit

- o default-lifetime
- o prefix-list: a list of prefixes to be advertised.

The following parameters are associated with each prefix in the list:

- * valid-lifetime
- * on-link-flag
- * preferred-lifetime
- * autonomous-flag

NOTES:

1. The "IsRouter" flag, which is also required by [RFC4861], is implemented in the "ietf-ip" module [RFC8344] (leaf "ip:forwarding").
2. The Neighbor Discovery specification [RFC4861] allows the implementations to decide whether the "valid-lifetime" and "preferred-lifetime" parameters remain the same in consecutive advertisements or decrement in real time. However, the latter behavior seems problematic because the values might be reset again to the (higher) configured values after a configuration is reloaded. Moreover, no implementation is known to use the decrementing behavior. The "ietf-ipv6-router-advertisements" submodule therefore stipulates the former behavior with constant values.

6. Interactions with Other YANG Modules

The semantics of the core routing data model also depends on several configuration parameters that are defined in other YANG modules.

6.1. Module "ietf-interfaces"

The following boolean switch is defined in the "ietf-interfaces" YANG module [RFC8343]:

`/if:interfaces/if:interface/if:enabled`

If this switch is set to "false" for a network-layer interface, then all routing and forwarding functions **MUST** be disabled on this interface.

6.2. Module "ietf-ip"

The following boolean switches are defined in the "ietf-ip" YANG module [RFC8344]:

`/if:interfaces/if:interface/ip:ipv4/ip:enabled`

If this switch is set to "false" for a network-layer interface, then all IPv4 routing and forwarding functions **MUST** be disabled on this interface.

`/if:interfaces/if:interface/ip:ipv4/ip:forwarding`

If this switch is set to "false" for a network-layer interface, then the forwarding of IPv4 datagrams through this interface **MUST** be disabled. However, the interface **MAY** participate in other IPv4 routing functions, such as routing protocols.

`/if:interfaces/if:interface/ip:ipv6/ip:enabled`

If this switch is set to "false" for a network-layer interface, then all IPv6 routing and forwarding functions **MUST** be disabled on this interface.

`/if:interfaces/if:interface/ip:ipv6/ip:forwarding`

If this switch is set to "false" for a network-layer interface, then the forwarding of IPv6 datagrams through this interface **MUST** be disabled. However, the interface **MAY** participate in other IPv6 routing functions, such as routing protocols.

In addition, the "ietf-ip" module allows for configuring IPv4 and IPv6 addresses and network prefixes or masks on network-layer interfaces. Configuration of these parameters on an enabled interface **MUST** result in an immediate creation of the corresponding direct route. The destination prefix of this route is set according to the configured IP address and network prefix/mask, and the interface is set as the outgoing interface for that route.

7. Routing Management YANG Module

```
<CODE BEGINS> file "ietf-routing@2018-03-13.yang"

module ietf-routing {
  yang-version "1.1";
  namespace "urn:ietf:params:xml:ns:yang:ietf-routing";
  prefix "rt";

  import ietf-yang-types {
    prefix "yang";
  }

  import ietf-interfaces {
    prefix "if";
    description
      "An 'ietf-interfaces' module version that is compatible with
       the Network Management Datastore Architecture (NMDA)
       is required.";
  }

  organization
    "IETF NETMOD (Network Modeling) Working Group";
  contact
    "WG Web:  <https://datatracker.ietf.org/wg/netmod/>
     WG List: <mailto:rtgw@ietf.org>

     Editor:  Ladislav Lhotka
              <mailto:lhotka@nic.cz>
              Acee Lindem
              <mailto:acee@cisco.com>
              Yingzhen Qu
              <mailto:yingzhen.qu@huawei.com>";

  description
    "This YANG module defines essential components for the management
     of a routing subsystem. The model fully conforms to the Network
     Management Datastore Architecture (NMDA).

     Copyright (c) 2018 IETF Trust and the persons
     identified as authors of the code. All rights reserved.

     Redistribution and use in source and binary forms, with or
     without modification, is permitted pursuant to, and subject
     to the license terms contained in, the Simplified BSD License
     set forth in Section 4.c of the IETF Trust's Legal Provisions
     Relating to IETF Documents
     (https://trustee.ietf.org/license-info).
```

This version of this YANG module is part of RFC 8349; see the RFC itself for full legal notices.";

```
revision 2018-03-13 {
  description
    "Network Management Datastore Architecture (NMDA) revision.";
  reference
    "RFC 8349: A YANG Data Model for Routing Management
      (NMDA Version)";
}

revision 2016-11-04 {
  description
    "Initial revision.";
  reference
    "RFC 8022: A YANG Data Model for Routing Management";
}

/* Features */
feature multiple-ribs {
  description
    "This feature indicates that the server supports
    user-defined RIBs.

    Servers that do not advertise this feature SHOULD provide
    exactly one system-controlled RIB per supported address family
    and also make it the default RIB. This RIB then appears as an
    entry in the list '/routing/ribs/rib'.";
}

feature router-id {
  description
    "This feature indicates that the server supports an explicit
    32-bit router ID that is used by some routing protocols.

    Servers that do not advertise this feature set a router ID
    algorithmically, usually to one of the configured IPv4
    addresses. However, this algorithm is implementation
    specific.";
}

/* Identities */
identity address-family {
  description
    "Base identity from which identities describing address
    families are derived.";
}
```



```
identity ipv4 {
  base address-family;
  description
    "This identity represents an IPv4 address family.";
}

identity ipv6 {
  base address-family;
  description
    "This identity represents an IPv6 address family.";
}

identity control-plane-protocol {
  description
    "Base identity from which control-plane protocol identities are
    derived.";
}

identity routing-protocol {
  base control-plane-protocol;
  description
    "Identity from which Layer 3 routing protocol identities are
    derived.";
}

identity direct {
  base routing-protocol;
  description
    "Routing pseudo-protocol that provides routes to directly
    connected networks.";
}

identity static {
  base routing-protocol;
  description
    "'Static' routing pseudo-protocol.";
}

/* Type Definitions */

typedef route-preference {
  type uint32;
  description
    "This type is used for route preferences.";
}

/* Groupings */
```

```
grouping address-family {
  description
    "This grouping provides a leaf identifying an address
    family.";
  leaf address-family {
    type identityref {
      base address-family;
    }
    mandatory true;
    description
      "Address family.";
  }
}

grouping router-id {
  description
    "This grouping provides a router ID.";
  leaf router-id {
    type yang:dotted-quad;
    description
      "A 32-bit number in the form of a dotted quad that is used by
      some routing protocols identifying a router.";
    reference
      "RFC 2328: OSPF Version 2";
  }
}

grouping special-next-hop {
  description
    "This grouping provides a leaf with an enumeration of special
    next hops.";
  leaf special-next-hop {
    type enumeration {
      enum blackhole {
        description
          "Silently discard the packet.";
      }
      enum unreachable {
        description
          "Discard the packet and notify the sender with an error
          message indicating that the destination host is
          unreachable.";
      }
      enum prohibit {
        description
          "Discard the packet and notify the sender with an error
          message indicating that the communication is
          administratively prohibited.";
      }
    }
  }
}
```

```

    }
    enum receive {
        description
            "The packet will be received by the local system.";
    }
}
description
    "Options for special next hops.";
}
}

grouping next-hop-content {
    description
        "Generic parameters of next hops in static routes.";
    choice next-hop-options {
        mandatory true;
        description
            "Options for next hops in static routes.

            It is expected that further cases will be added through
            augments from other modules.";
        case simple-next-hop {
            description
                "This case represents a simple next hop consisting of the
                next-hop address and/or outgoing interface.

                Modules for address families MUST augment this case with a
                leaf containing a next-hop address of that address
                family.";
            leaf outgoing-interface {
                type if:interface-ref;
                description
                    "Name of the outgoing interface.";
            }
        }
        case special-next-hop {
            uses special-next-hop;
        }
        case next-hop-list {
            container next-hop-list {
                description
                    "Container for multiple next hops.";
                list next-hop {
                    key "index";
                    description
                        "An entry in a next-hop list.

                        Modules for address families MUST augment this list

```



```

    }
    case next-hop-list {
      container next-hop-list {
        description
          "Container for multiple next hops.";
        list next-hop {
          description
            "An entry in a next-hop list.

            Modules for address families MUST augment this list
            with a leaf containing a next-hop address of that
            address family.";
          leaf outgoing-interface {
            type if:interface-ref;
            description
              "Name of the outgoing interface.";
          }
        }
      }
    }
  }
}

grouping route-metadata {
  description
    "Common route metadata.";
  leaf source-protocol {
    type identityref {
      base routing-protocol;
    }
    mandatory true;
    description
      "Type of the routing protocol from which the route
      originated.";
  }
  leaf active {
    type empty;
    description
      "The presence of this leaf indicates that the route is
      preferred among all routes in the same RIB that have the
      same destination prefix.";
  }
  leaf last-updated {
    type yang:date-and-time;
    description
      "Timestamp of the last modification of the route. If the
      route was never modified, it is the time when the route was
      inserted into the RIB.";
  }
}

```

```
    }  
  }  
  
  /* Data nodes */  
  
  container routing {  
    description  
      "Configuration parameters for the routing subsystem.";  
    uses router-id {  
      if-feature "router-id";  
      description  
        "Support for the global router ID. Routing protocols  
        that use a router ID can use this parameter or override it  
        with another value.";  
    }  
    container interfaces {  
      config false;  
      description  
        "Network-layer interfaces used for routing.";  
      leaf-list interface {  
        type if:interface-ref;  
        description  
          "Each entry is a reference to the name of a configured  
          network-layer interface.";  
      }  
    }  
    container control-plane-protocols {  
      description  
        "Support for control-plane protocol instances.";  
      list control-plane-protocol {  
        key "type name";  
        description  
          "Each entry contains a control-plane protocol instance.";  
        leaf type {  
          type identityref {  
            base control-plane-protocol;  
          }  
          description  
            "Type of the control-plane protocol -- an identity  
            derived from the 'control-plane-protocol'  
            base identity.";  
        }  
        leaf name {  
          type string;  
          description  
            "An arbitrary name of the control-plane protocol  
            instance.";  
        }  
      }  
    }  
  }  
}
```

```
leaf description {
  type string;
  description
    "Textual description of the control-plane protocol
    instance.";
}
container static-routes {
  when "derived-from-or-self(..../type, 'rt:static')" {
    description
      "This container is only valid for the 'static' routing
      protocol.";
  }
  description
    "Support for the 'static' pseudo-protocol.

    Address-family-specific modules augment this node with
    their lists of routes.";
}
}
}
container ribs {
  description
    "Support for RIBs.";
  list rib {
    key "name";
    description
      "Each entry contains a configuration for a RIB identified
      by the 'name' key.

      Entries having the same key as a system-controlled entry
      in the list '/routing/ribs/rib' are used for
      configuring parameters of that entry. Other entries
      define additional user-controlled RIBs.";
  }
  leaf name {
    type string;
    description
      "The name of the RIB.

      For system-controlled entries, the value of this leaf
      must be the same as the name of the corresponding entry
      in the operational state.

      For user-controlled entries, an arbitrary name can be
      used.";
  }
  uses address-family {
    description
      "The address family of the system-controlled RIB.";
  }
}
```

```
}  
leaf default-rib {  
  if-feature "multiple-ribs";  
  type boolean;  
  default "true";  
  config false;  
  description  
    "This flag has the value of 'true' if and only if the RIB  
    is the default RIB for the given address family.  
  
    By default, control-plane protocols place their routes  
    in the default RIBs.";  
}  
container routes {  
  config false;  
  description  
    "Current contents of the RIB.";  
  list route {  
    description  
      "A RIB route entry. This data node MUST be augmented  
      with information specific to routes of each address  
      family.";  
    leaf route-preference {  
      type route-preference;  
      description  
        "This route attribute, also known as 'administrative  
        distance', allows for selecting the preferred route  
        among routes with the same destination prefix. A  
        smaller value indicates a route that is  
        more preferred.";  
    }  
    container next-hop {  
      description  
        "Route's next-hop attribute.";  
      uses next-hop-state-content;  
    }  
    uses route-metadata;  
  }  
}  
action active-route {  
  description  
    "Return the active RIB route that is used for the  
    destination address.  
  
    Address-family-specific modules MUST augment input  
    parameters with a leaf named 'destination-address'.";  
  output {
```



```

    container route {
      description
        "The active RIB route for the specified destination.

        If no route exists in the RIB for the destination
        address, no output is returned.

        Address-family-specific modules MUST augment this
        container with appropriate route contents.";
      container next-hop {
        description
          "Route's next-hop attribute.";
        uses next-hop-state-content;
      }
      uses route-metadata;
    }
  }
}
leaf description {
  type string;
  description
    "Textual description of the RIB.";
}
}
}
}

/*
 * The subsequent data nodes are obviated and obsoleted
 * by the Network Management Datastore Architecture
 * as described in RFC 8342.
 */
container routing-state {
  config false;
  status obsolete;
  description
    "State data of the routing subsystem.";
  uses router-id {
    status obsolete;
    description
      "Global router ID.

      It may be either configured or assigned algorithmically by
      the implementation.";
  }
  container interfaces {
    status obsolete;
    description

```

```
"Network-layer interfaces used for routing.";
leaf-list interface {
  type if:interface-state-ref;
  status obsolete;
  description
    "Each entry is a reference to the name of a configured
    network-layer interface.";
}
}
container control-plane-protocols {
  status obsolete;
  description
    "Container for the list of routing protocol instances.";
  list control-plane-protocol {
    key "type name";
    status obsolete;
    description
      "State data of a control-plane protocol instance.

      An implementation MUST provide exactly one
      system-controlled instance of the 'direct'
      pseudo-protocol. Instances of other control-plane
      protocols MAY be created by configuration.";
    leaf type {
      type identityref {
        base control-plane-protocol;
      }
      status obsolete;
      description
        "Type of the control-plane protocol.";
    }
    leaf name {
      type string;
      status obsolete;
      description
        "The name of the control-plane protocol instance.

        For system-controlled instances, this name is
        persistent, i.e., it SHOULD NOT change across
        reboots.";
    }
  }
}
}
container ribs {
  status obsolete;
  description
    "Container for RIBs.";
  list rib {
```

```
key "name";
min-elements 1;
status obsolete;
description
  "Each entry represents a RIB identified by the 'name'
  key. All routes in a RIB MUST belong to the same address
  family.
```

An implementation SHOULD provide one system-controlled default RIB for each supported address family.";

```
leaf name {
  type string;
  status obsolete;
  description
    "The name of the RIB.";
}
uses address-family {
  status obsolete;
  description
    "The address family of the RIB.";
}
```

```
leaf default-rib {
  if-feature "multiple-ribs";
  type boolean;
  default "true";
  status obsolete;
  description
    "This flag has the value of 'true' if and only if the
    RIB is the default RIB for the given address family.
```

By default, control-plane protocols place their routes in the default RIBs.";

```
}
container routes {
  status obsolete;
  description
    "Current contents of the RIB.";
  list route {
    status obsolete;
    description
      "A RIB route entry. This data node MUST be augmented
      with information specific to routes of each address
      family.";
    leaf route-preference {
      type route-preference;
      status obsolete;
      description
        "This route attribute, also known as 'administrative
```

```

        distance', allows for selecting the preferred route
        among routes with the same destination prefix. A
        smaller value indicates a route that is
        more preferred.";
    }
    container next-hop {
        status obsolete;
        description
            "Route's next-hop attribute.";
        uses next-hop-state-content {
            status obsolete;
            description
                "Route's next-hop attribute operational state.";
        }
    }
    uses route-metadata {
        status obsolete;
        description
            "Route metadata.";
    }
}

```

```

action active-route {
    status obsolete;
    description
        "Return the active RIB route that is used for the
        destination address.

```

Address-family-specific modules MUST augment input parameters with a leaf named 'destination-address'.";

```

output {
    container route {
        status obsolete;
        description
            "The active RIB route for the specified
            destination.

```

If no route exists in the RIB for the destination address, no output is returned.

Address-family-specific modules MUST augment this container with appropriate route contents.";

```

    container next-hop {
        status obsolete;
        description
            "Route's next-hop attribute.";
        uses next-hop-state-content {
            status obsolete;

```

```
        description  
            "Active route state data.";  
    }  
}  
uses route-metadata {  
    status obsolete;  
    description  
        "Active route metadata.";  
}  
}  
}  
}  
}  
}  
}
```

<CODE ENDS>

8. IPv4 Unicast Routing Management YANG Module

```
<CODE BEGINS> file "ietf-ipv4-unicast-routing@2018-03-13.yang"
```

```

module ietf-ipv4-unicast-routing {
  yang-version "1.1";
  namespace
    "urn:ietf:params:xml:ns:yang:ietf-ipv4-unicast-routing";
  prefix "v4ur";

  import ietf-routing {
    prefix "rt";
    description
      "An 'ietf-routing' module version that is compatible with
       the Network Management Datastore Architecture (NMDA)
       is required.";
  }

  import ietf-inet-types {
    prefix "inet";
  }

  organization
    "IETF NETMOD (Network Modeling) Working Group";
  contact
    "WG Web:    <https://datatracker.ietf.org/wg/netmod/>
     WG List:   <mailto:rtgwg@ietf.org>

     Editor:    Ladislav Lhotka
                <mailto:lhotka@nic.cz>

```

Acee Lindem
<mailto:acee@cisco.com>
Yingzhen Qu
<mailto:yingzhen.qu@huawei.com>;

description

"This YANG module augments the 'ietf-routing' module with basic parameters for IPv4 unicast routing. The model fully conforms to the Network Management Datastore Architecture (NMDA).

Copyright (c) 2018 IETF Trust and the persons identified as authors of the code. All rights reserved.

Redistribution and use in source and binary forms, with or without modification, is permitted pursuant to, and subject to the license terms contained in, the Simplified BSD License set forth in Section 4.c of the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>).

This version of this YANG module is part of RFC 8349; see the RFC itself for full legal notices."

```
revision 2018-03-13 {
  description
    "Network Management Datastore Architecture (NMDA) revision.";
  reference
    "RFC 8349: A YANG Data Model for Routing Management
      (NMDA Version)";
}

revision 2016-11-04 {
  description
    "Initial revision.";
  reference
    "RFC 8022: A YANG Data Model for Routing Management";
}

/* Identities */

identity ipv4-unicast {
  base rt:ipv4;
  description
    "This identity represents the IPv4 unicast address family.";
}

augment "/rt:routing/rt:ribs/rt:rib/rt:routes/rt:route" {
  when "derived-from-or-self(..../rt:address-family, "
```

```
    + "'v4ur:ipv4-unicast')" {
      description
        "This augment is valid only for IPv4 unicast.";
    }
    description
      "This leaf augments an IPv4 unicast route.";
    leaf destination-prefix {
      type inet:ipv4-prefix;
      description
        "IPv4 destination prefix.";
    }
  }
}

augment "/rt:routing/rt:ribs/rt:rib/rt:routes/rt:route/"
  + "rt:next-hop/rt:next-hop-options/rt:simple-next-hop" {
  when "derived-from-or-self(..../rt:address-family, "
    + "'v4ur:ipv4-unicast')" {
    description
      "This augment is valid only for IPv4 unicast.";
  }
  description
    "Augments the 'simple-next-hop' case in IPv4 unicast routes.";
  leaf next-hop-address {
    type inet:ipv4-address;
    description
      "IPv4 address of the next hop.";
  }
}

augment "/rt:routing/rt:ribs/rt:rib/rt:routes/rt:route/"
  + "rt:next-hop/rt:next-hop-options/rt:next-hop-list/"
  + "rt:next-hop-list/rt:next-hop" {
  when "derived-from-or-self(..../rt:address-family, "
    + "'v4ur:ipv4-unicast')" {
    description
      "This augment is valid only for IPv4 unicast.";
  }
  description
    "This leaf augments the 'next-hop-list' case of IPv4 unicast
      routes.";
  leaf address {
    type inet:ipv4-address;
    description
      "IPv4 address of the next hop.";
  }
}
}
```

augment

```
"/rt:routing/rt:ribs/rt:rib/rt:active-route/rt:input" {
  when "derived-from-or-self(..../rt:address-family, "
    + "'v4ur:ipv4-unicast')" {
    description
      "This augment is valid only for IPv4 unicast RIBs.";
  }
  description
    "This augment adds the input parameter of the 'active-route'
    action.";
  leaf destination-address {
    type inet:ipv4-address;
    description
      "IPv4 destination address.";
  }
}

augment "/rt:routing/rt:ribs/rt:rib/rt:active-route/"
  + "rt:output/rt:route" {
  when "derived-from-or-self(..../rt:address-family, "
    + "'v4ur:ipv4-unicast')" {
    description
      "This augment is valid only for IPv4 unicast.";
  }
  description
    "This augment adds the destination prefix to the reply of the
    'active-route' action.";
  leaf destination-prefix {
    type inet:ipv4-prefix;
    description
      "IPv4 destination prefix.";
  }
}

augment "/rt:routing/rt:ribs/rt:rib/rt:active-route/"
  + "rt:output/rt:route/rt:next-hop/rt:next-hop-options/"
  + "rt:simple-next-hop" {
  when "derived-from-or-self(..../rt:address-family, "
    + "'v4ur:ipv4-unicast')" {
    description
      "This augment is valid only for IPv4 unicast.";
  }
  description
    "Augments the 'simple-next-hop' case in the reply to the
    'active-route' action.";
  leaf next-hop-address {
    type inet:ipv4-address;
    description
      "IPv4 address of the next hop.";
```



```

    }
  }

  augment "/rt:routing/rt:ribs/rt:rib/rt:active-route/"
    + "rt:output/rt:route/rt:next-hop/rt:next-hop-options/"
    + "rt:next-hop-list/rt:next-hop-list/rt:next-hop" {
    when "derived-from-or-self(..../..../..../rt:address-family, "
      + "'v4ur:ipv4-unicast')" {
      description
        "This augment is valid only for IPv4 unicast.";
    }
    description
      "Augments the 'next-hop-list' case in the reply to the
      'active-route' action.";
    leaf next-hop-address {
      type inet:ipv4-address;
      description
        "IPv4 address of the next hop.";
    }
  }

  augment "/rt:routing/rt:control-plane-protocols/"
    + "rt:control-plane-protocol/rt:static-routes" {
    description
      "This augment defines the 'static' pseudo-protocol
      with data specific to IPv4 unicast.";
    container ipv4 {
      description
        "Support for a 'static' pseudo-protocol instance
        consists of a list of routes.";
      list route {
        key "destination-prefix";
        description
          "A list of static routes.";
        leaf destination-prefix {
          type inet:ipv4-prefix;
          mandatory true;
          description
            "IPv4 destination prefix.";
        }
        leaf description {
          type string;
          description
            "Textual description of the route.";
        }
      }
      container next-hop {
        description
          "Support for next-hop.";
      }
    }
  }

```

```

    uses rt:next-hop-content {
      augment "next-hop-options/simple-next-hop" {
        description
          "Augments the 'simple-next-hop' case in IPv4 static
          routes.";
        leaf next-hop-address {
          type inet:ipv4-address;
          description
            "IPv4 address of the next hop.";
        }
      }
      augment "next-hop-options/next-hop-list/next-hop-list/"
        + "next-hop" {
        description
          "Augments the 'next-hop-list' case in IPv4 static
          routes.";
        leaf next-hop-address {
          type inet:ipv4-address;
          description
            "IPv4 address of the next hop.";
        }
      }
    }
  }
}

/*
 * The subsequent data nodes are obviated and obsoleted
 * by the Network Management Datastore Architecture
 * as described in RFC 8342.
 */
augment "/rt:routing-state/rt:ribs/rt:rib/rt:routes/rt:route" {
  when "derived-from-or-self(..../rt:address-family, "
    + "'v4ur:ipv4-unicast')" {
    description
      "This augment is valid only for IPv4 unicast.";
  }
  status obsolete;
  description
    "This leaf augments an IPv4 unicast route.";
  leaf destination-prefix {
    type inet:ipv4-prefix;
    status obsolete;
    description
      "IPv4 destination prefix.";
  }
}

```

```

}
augment "/rt:routing-state/rt:ribs/rt:rib/rt:routes/rt:route/"
  + "rt:next-hop/rt:next-hop-options/rt:simple-next-hop" {
  when "derived-from-or-self(
    ../../../../rt:address-family, 'v4ur:ipv4-unicast')" {
    description
      "This augment is valid only for IPv4 unicast.";
  }
  status obsolete;
  description
    "Augments the 'simple-next-hop' case in IPv4 unicast routes.";
  leaf next-hop-address {
    type inet:ipv4-address;
    status obsolete;
    description
      "IPv4 address of the next hop.";
  }
}
augment "/rt:routing-state/rt:ribs/rt:rib/rt:routes/rt:route/"
  + "rt:next-hop/rt:next-hop-options/rt:next-hop-list/"
  + "rt:next-hop-list/rt:next-hop" {
  when "derived-from-or-self(../../../../../rt:address-family,
    'v4ur:ipv4-unicast')" {
    description
      "This augment is valid only for IPv4 unicast.";
  }
  status obsolete;
  description
    "This leaf augments the 'next-hop-list' case of IPv4 unicast
    routes.";
  leaf address {
    type inet:ipv4-address;
    status obsolete;
    description
      "IPv4 address of the next hop.";
  }
}
augment "/rt:routing-state/rt:ribs/rt:rib/rt:active-route/"
  + "rt:input" {
  when "derived-from-or-self(../../rt:address-family,
    'v4ur:ipv4-unicast')" {
    description
      "This augment is valid only for IPv4 unicast RIBs.";
  }
  status obsolete;
  description
    "This augment adds the input parameter of the 'active-route'
    action.";
}

```

```
    leaf destination-address {
      type inet:ipv4-address;
      status obsolete;
      description
        "IPv4 destination address.";
    }
  }
  augment "/rt:routing-state/rt:ribs/rt:rib/rt:active-route/"
    + "rt:output/rt:route" {
    when "derived-from-or-self(..../rt:address-family,
      'v4ur:ipv4-unicast')" {
      description
        "This augment is valid only for IPv4 unicast.";
    }
    status obsolete;
    description
      "This augment adds the destination prefix to the reply of the
        'active-route' action.";
    leaf destination-prefix {
      type inet:ipv4-prefix;
      status obsolete;
      description
        "IPv4 destination prefix.";
    }
  }
  augment "/rt:routing-state/rt:ribs/rt:rib/rt:active-route/"
    + "rt:output/rt:route/rt:next-hop/rt:next-hop-options/"
    + "rt:simple-next-hop" {
    when "derived-from-or-self(..../rt:address-family,
      'v4ur:ipv4-unicast')" {
      description
        "This augment is valid only for IPv4 unicast.";
    }
    status obsolete;
    description
      "Augments the 'simple-next-hop' case in the reply to the
        'active-route' action.";
    leaf next-hop-address {
      type inet:ipv4-address;
      status obsolete;
      description
        "IPv4 address of the next hop.";
    }
  }
  augment "/rt:routing-state/rt:ribs/rt:rib/rt:active-route/"
    + "rt:output/rt:route/rt:next-hop/rt:next-hop-options/"
    + "rt:next-hop-list/rt:next-hop-list/rt:next-hop" {
    when "derived-from-or-self(..../rt:address-family,
```

```

        'v4ur:ipv4-unicast')" {
    description
        "This augment is valid only for IPv4 unicast.";
    }
    status obsolete;
    description
        "Augments the 'next-hop-list' case in the reply to the
        'active-route' action.";
    leaf next-hop-address {
        type inet:ipv4-address;
        status obsolete;
        description
            "IPv4 address of the next hop.";
    }
}
}
}

```

<CODE ENDS>

9. IPv6 Unicast Routing Management YANG Module

<CODE BEGINS> file "ietf-ipv6-unicast-routing@2018-03-13.yang"

```

module ietf-ipv6-unicast-routing {
    yang-version "1.1";
    namespace
        "urn:ietf:params:xml:ns:yang:ietf-ipv6-unicast-routing";
    prefix "v6ur";

    import ietf-routing {
        prefix "rt";
        description
            "An 'ietf-routing' module version that is compatible with
            the Network Management Datastore Architecture (NMDA)
            is required.";
    }

    import ietf-inet-types {
        prefix "inet";
        description
            "An 'ietf-interfaces' module version that is compatible with
            the Network Management Datastore Architecture (NMDA)
            is required.";
    }

    include ietf-ipv6-router-advertisements {
        revision-date 2018-03-13;
    }
}

```

```
organization
  "IETF NETMOD (Network Modeling) Working Group";
contact
  "WG Web:    <https://datatracker.ietf.org/wg/netmod/>
  WG List:    <mailto:rtgwg@ietf.org>

  Editor:     Ladislav Lhotka
               <mailto:lhotka@nic.cz>
               Acee Lindem
               <mailto:acee@cisco.com>
               Yingzhen Qu
               <mailto:yingzhen.qu@huawei.com>";

description
  "This YANG module augments the 'ietf-routing' module with basic
  parameters for IPv6 unicast routing. The model fully conforms
  to the Network Management Datastore Architecture (NMDA).

  Copyright (c) 2018 IETF Trust and the persons
  identified as authors of the code. All rights reserved.

  Redistribution and use in source and binary forms, with or
  without modification, is permitted pursuant to, and subject
  to the license terms contained in, the Simplified BSD License
  set forth in Section 4.c of the IETF Trust's Legal Provisions
  Relating to IETF Documents
  (https://trustee.ietf.org/license-info).

  This version of this YANG module is part of RFC 8349; see
  the RFC itself for full legal notices.";

revision 2018-03-13 {
  description
    "Network Management Datastore Architecture (NMDA) revision.";
  reference
    "RFC 8349: A YANG Data Model for Routing Management
    (NMDA Version)";
}

/* Identities */

revision 2016-11-04 {
  description
    "Initial revision.";
  reference
    "RFC 8022: A YANG Data Model for Routing Management";
}
```

```
identity ipv6-unicast {
  base rt:ipv6;
  description
    "This identity represents the IPv6 unicast address family.";
}

augment "/rt:routing/rt:ribs/rt:rib/rt:routes/rt:route" {
  when "derived-from-or-self(..../rt:address-family, "
    + "'v6ur:ipv6-unicast')" {
    description
      "This augment is valid only for IPv6 unicast.";
  }
  description
    "This leaf augments an IPv6 unicast route.";
  leaf destination-prefix {
    type inet:ipv6-prefix;
    description
      "IPv6 destination prefix.";
  }
}

augment "/rt:routing/rt:ribs/rt:rib/rt:routes/rt:route/"
  + "rt:next-hop/rt:next-hop-options/rt:simple-next-hop" {
  when "derived-from-or-self(..../rt:address-family, "
    + "'v6ur:ipv6-unicast')" {
    description
      "This augment is valid only for IPv6 unicast.";
  }
  description
    "Augments the 'simple-next-hop' case in IPv6 unicast routes.";
  leaf next-hop-address {
    type inet:ipv6-address;
    description
      "IPv6 address of the next hop.";
  }
}

augment "/rt:routing/rt:ribs/rt:rib/rt:routes/rt:route/"
  + "rt:next-hop/rt:next-hop-options/rt:next-hop-list/"
  + "rt:next-hop-list/rt:next-hop" {
  when "derived-from-or-self(..../rt:address-family, "
    + "'v6ur:ipv6-unicast')" {
    description
      "This augment is valid only for IPv6 unicast.";
  }
  description
    "This leaf augments the 'next-hop-list' case of IPv6 unicast
      routes.";
```

```
    leaf address {
      type inet:ipv6-address;
      description
        "IPv6 address of the next hop.";
    }
  }

  augment
    "/rt:routing/rt:ribs/rt:rib/rt:active-route/rt:input" {
    when "derived-from-or-self(..../rt:address-family, "
      + "'v6ur:ipv6-unicast')" {
      description
        "This augment is valid only for IPv6 unicast RIBs.";
    }
    description
      "This augment adds the input parameter of the 'active-route'
        action.";
    leaf destination-address {
      type inet:ipv6-address;
      description
        "IPv6 destination address.";
    }
  }

  augment "/rt:routing/rt:ribs/rt:rib/rt:active-route/"
    + "rt:output/rt:route" {
    when "derived-from-or-self(..../rt:address-family, "
      + "'v6ur:ipv6-unicast')" {
      description
        "This augment is valid only for IPv6 unicast.";
    }
    description
      "This augment adds the destination prefix to the reply of the
        'active-route' action.";
    leaf destination-prefix {
      type inet:ipv6-prefix;
      description
        "IPv6 destination prefix.";
    }
  }

  augment "/rt:routing/rt:ribs/rt:rib/rt:active-route/"
    + "rt:output/rt:route/rt:next-hop/rt:next-hop-options/"
    + "rt:simple-next-hop" {
    when "derived-from-or-self(..../rt:address-family, "
      + "'v6ur:ipv6-unicast')" {
      description
        "This augment is valid only for IPv6 unicast.";
    }
  }
```



```

    }
    description
      "Augments the 'simple-next-hop' case in the reply to the
      'active-route' action.";
    leaf next-hop-address {
      type inet:ipv6-address;
      description
        "IPv6 address of the next hop.";
    }
  }
}

augment "/rt:routing/rt:ribs/rt:rib/rt:active-route/"
  + "rt:output/rt:route/rt:next-hop/rt:next-hop-options/"
  + "rt:next-hop-list/rt:next-hop-list/rt:next-hop" {
  when "derived-from-or-self(..../..../rt:address-family, "
    + "'v6ur:ipv6-unicast')" {
    description
      "This augment is valid only for IPv6 unicast.";
  }
  description
    "Augments the 'next-hop-list' case in the reply to the
    'active-route' action.";
  leaf next-hop-address {
    type inet:ipv6-address;
    description
      "IPv6 address of the next hop.";
  }
}

/* Data node augmentations */

augment "/rt:routing/rt:control-plane-protocols/"
  + "rt:control-plane-protocol/rt:static-routes" {
  description
    "This augment defines the 'static' pseudo-protocol
    with data specific to IPv6 unicast.";
  container ipv6 {
    description
      "Support for a 'static' pseudo-protocol instance
      consists of a list of routes.";
    list route {
      key "destination-prefix";
      description
        "A list of static routes.";
      leaf destination-prefix {
        type inet:ipv6-prefix;
        mandatory true;
        description

```



```
    }
    status obsolete;
    description
      "This leaf augments an IPv6 unicast route.";
    leaf destination-prefix {
      type inet:ipv6-prefix;
      status obsolete;
      description
        "IPv6 destination prefix.";
    }
  }
  augment "/rt:routing-state/rt:ribs/rt:rib/rt:routes/rt:route/"
    + "rt:next-hop/rt:next-hop-options/rt:simple-next-hop" {
    when "derived-from-or-self(..../rt:address-family,
      'v6ur:ipv6-unicast')" {
      description
        "This augment is valid only for IPv6 unicast.";
    }
    status obsolete;
    description
      "Augments the 'simple-next-hop' case in IPv6 unicast routes.";
    leaf next-hop-address {
      type inet:ipv6-address;
      status obsolete;
      description
        "IPv6 address of the next hop.";
    }
  }
  augment "/rt:routing-state/rt:ribs/rt:rib/rt:routes/rt:route/"
    + "rt:next-hop/rt:next-hop-options/rt:next-hop-list/"
    + "rt:next-hop-list/rt:next-hop" {
    when "derived-from-or-self(..../rt:address-family,
      'v6ur:ipv6-unicast')" {
      description
        "This augment is valid only for IPv6 unicast.";
    }
    status obsolete;
    description
      "This leaf augments the 'next-hop-list' case of IPv6 unicast
        routes.";
    leaf address {
      type inet:ipv6-address;
      status obsolete;
      description
        "IPv6 address of the next hop.";
    }
  }
}
augment "/rt:routing-state/rt:ribs/rt:rib/"
```

```

        + "rt:active-route/rt:input" {
when "derived-from-or-self(..../rt:address-family,
    'v6ur:ipv6-unicast')" {
    description
        "This augment is valid only for IPv6 unicast RIBs.";
}
status obsolete;
description
    "This augment adds the input parameter of the 'active-route'
    action.";
leaf destination-address {
    type inet:ipv6-address;
    status obsolete;
    description
        "IPv6 destination address.";
}
}
augment "/rt:routing-state/rt:ribs/rt:rib/rt:active-route/"
    + "rt:output/rt:route" {
when "derived-from-or-self(..../rt:address-family,
    'v6ur:ipv6-unicast')" {
    description
        "This augment is valid only for IPv6 unicast.";
}
status obsolete;
description
    "This augment adds the destination prefix to the reply of the
    'active-route' action.";
leaf destination-prefix {
    type inet:ipv6-prefix;
    status obsolete;
    description
        "IPv6 destination prefix.";
}
}
augment "/rt:routing-state/rt:ribs/rt:rib/rt:active-route/"
    + "rt:output/rt:route/rt:next-hop/rt:next-hop-options/"
    + "rt:simple-next-hop" {
when "derived-from-or-self(..../rt:address-family,
    'v6ur:ipv6-unicast')" {
    description
        "This augment is valid only for IPv6 unicast.";
}
status obsolete;
description
    "Augments the 'simple-next-hop' case in the reply to the
    'active-route' action.";
leaf next-hop-address {

```

```

        type inet:ipv6-address;
        status obsolete;
        description
            "IPv6 address of the next hop.";
    }
}
augment "/rt:routing-state/rt:ribs/rt:rib/rt:active-route/"
    + "rt:output/rt:route/rt:next-hop/rt:next-hop-options/"
    + "rt:next-hop-list/rt:next-hop-list/rt:next-hop" {
    when "derived-from-or-self(..../..../..../rt:address-family,
        'v6ur:ipv6-unicast')" {
        description
            "This augment is valid only for IPv6 unicast.";
    }
    status obsolete;
    description
        "Augments the 'next-hop-list' case in the reply to the
        'active-route' action.";
    leaf next-hop-address {
        type inet:ipv6-address;
        status obsolete;
        description
            "IPv6 address of the next hop.";
    }
}
}

```

<CODE ENDS>

9.1. IPv6 Router Advertisements Submodule

<CODE BEGINS> file "ietf-ipv6-router-advertisements@2018-03-13.yang"

```

submodule ietf-ipv6-router-advertisements {
    yang-version "1.1";

    belongs-to ietf-ipv6-unicast-routing {
        prefix "v6ur";
    }

    import ietf-inet-types {
        prefix "inet";
    }

    import ietf-interfaces {
        prefix "if";
        description
            "An 'ietf-interfaces' module version that is compatible with

```

```
    the Network Management Datastore Architecture (NMDA)
    is required.";
}

import ietf-ip {
    prefix "ip";
    description
        "An 'ietf-ip' module version that is compatible with
        the Network Management Datastore Architecture (NMDA)
        is required.";
}

organization
    "IETF NETMOD (Network Modeling) Working Group";
contact
    "WG Web:    <https://datatracker.ietf.org/wg/netmod/>
    WG List:    <mailto:rtgwg@ietf.org>

    Editor:     Ladislav Lhotka
                 <mailto:lhotka@nic.cz>
                 Acee Lindem
                 <mailto:acee@cisco.com>
                 Yingzhen Qu
                 <mailto:yingzhen.qu@huawei.com>";

description
    "This YANG module augments the 'ietf-ip' module with
    parameters for IPv6 Router Advertisements. The model fully
    conforms to the Network Management Datastore
    Architecture (NMDA).

    Copyright (c) 2018 IETF Trust and the persons
    identified as authors of the code. All rights reserved.

    Redistribution and use in source and binary forms, with or
    without modification, is permitted pursuant to, and subject
    to the license terms contained in, the Simplified BSD License
    set forth in Section 4.c of the IETF Trust's Legal Provisions
    Relating to IETF Documents
    (https://trustee.ietf.org/license-info).

    This version of this YANG module is part of RFC 8349; see
    the RFC itself for full legal notices.";

reference
    "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)";

revision 2018-03-13 {
```

```
    description
      "Network Management Datastore Architecture (NMDA) revision.";
    reference
      "RFC 8349: A YANG Data Model for Routing Management
        (NMDA Version)";
  }

  revision 2016-11-04 {
    description
      "Initial revision.";
    reference
      "RFC 8022: A YANG Data Model for Routing Management";
  }

  augment "/if:interfaces/if:interface/ip:ipv6" {
    description
      "Augments interface configuration with parameters of IPv6
        Router Advertisements.";
    container ipv6-router-advertisements {
      description
        "Support for IPv6 Router Advertisements.";
      leaf send-advertisements {
        type boolean;
        default "false";
        description
          "A flag indicating whether or not the router sends
            periodic Router Advertisements and responds to
            Router Solicitations.";
        reference
          "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
            - AdvSendAdvertisements";
      }
      leaf max-rtr-adv-interval {
        type uint16 {
          range "4..65535";
        }
        units "seconds";
        default "600";
        description
          "The maximum time allowed between sending unsolicited
            multicast Router Advertisements from the interface.";
        reference
          "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
            - MaxRtrAdvInterval";
      }
      leaf min-rtr-adv-interval {
        type uint16 {
          range "3..1350";
        }
      }
    }
  }
```

```
}
units "seconds";
must ". <= 0.75 * ../max-rtr-adv-interval" {
  description
    "The value MUST NOT be greater than 75% of
    'max-rtr-adv-interval'.";
}
description
  "The minimum time allowed between sending unsolicited
  multicast Router Advertisements from the interface.

  The default value to be used operationally if this
  leaf is not configured is determined as follows:

  - if max-rtr-adv-interval >= 9 seconds, the default
    value is 0.33 * max-rtr-adv-interval;

  - otherwise, it is 0.75 * max-rtr-adv-interval.";
reference
  "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
  - MinRtrAdvInterval";
}
leaf managed-flag {
  type boolean;
  default "false";
  description
    "The value to be placed in the 'Managed address
    configuration' flag field in the Router
    Advertisement.";
  reference
    "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
    - AdvManagedFlag";
}
leaf other-config-flag {
  type boolean;
  default "false";
  description
    "The value to be placed in the 'Other configuration'
    flag field in the Router Advertisement.";
  reference
    "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
    - AdvOtherConfigFlag";
}
leaf link-mtu {
  type uint32;
  default "0";
  description
    "The value to be placed in MTU options sent by the
```



```

        router. A value of zero indicates that no MTU options
        are sent.";
    reference
        "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
        - AdvLinkMTU";
}
leaf reachable-time {
    type uint32 {
        range "0..3600000";
    }
    units "milliseconds";
    default "0";
    description
        "The value to be placed in the Reachable Time field in
        the Router Advertisement messages sent by the router.
        A value of zero means unspecified (by this router).";
    reference
        "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
        - AdvReachableTime";
}
leaf retrans-timer {
    type uint32;
    units "milliseconds";
    default "0";
    description
        "The value to be placed in the Retrans Timer field in
        the Router Advertisement messages sent by the router.
        A value of zero means unspecified (by this router).";
    reference
        "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
        - AdvRetransTimer";
}
leaf cur-hop-limit {
    type uint8;
    description
        "The value to be placed in the Cur Hop Limit field in
        the Router Advertisement messages sent by the router.
        A value of zero means unspecified (by this router).

        If this parameter is not configured, the device SHOULD
        use the IANA-specified value for the default IPv4
        Time to Live (TTL) parameter that was in effect at the
        time of implementation.";
    reference
        "RFC 3232: Assigned Numbers: RFC 1700 is Replaced by
        an On-line Database
        RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
        - AdvCurHopLimit

```

```
        IANA: IP Parameters
            (https://www.iana.org/assignments/ip-parameters)";
    }
    leaf default-lifetime {
        type uint16 {
            range "0..65535";
        }
        units "seconds";
        description
            "The value to be placed in the Router Lifetime field of
            Router Advertisements sent from the interface, in
            seconds. It MUST be either zero or between
            max-rtr-adv-interval and 9000 seconds. A value of zero
            indicates that the router is not to be used as a
            default router. These limits may be overridden by
            specific documents that describe how IPv6 operates over
            different link layers.

            If this parameter is not configured, the device SHOULD
            use a value of 3 * max-rtr-adv-interval.";
        reference
            "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
            - AdvDefaultLifetime";
    }
    container prefix-list {
        description
            "Support for prefixes to be placed in Prefix
            Information options in Router Advertisement messages
            sent from the interface.

            Prefixes that are advertised by default but do not
            have their entries in the child 'prefix' list are
            advertised with the default values of all parameters.

            The link-local prefix SHOULD NOT be included in the
            list of advertised prefixes.";
        reference
            "RFC 4861: Neighbor Discovery for IP version 6 (IPv6)
            - AdvPrefixList";
        list prefix {
            key "prefix-spec";
            description
                "Support for an advertised prefix entry.";
            leaf prefix-spec {
                type inet:ipv6-prefix;
                description
                    "IPv6 address prefix.";
            }
        }
    }
}
```

```
choice control-adv-prefixes {
  default "advertise";
  description
    "Either (1) the prefix is explicitly removed from the
     set of advertised prefixes or (2) the parameters with
     which the prefix is advertised are specified (default
     case).";
  leaf no-advertise {
    type empty;
    description
      "The prefix will not be advertised.

      This can be used for removing the prefix from
      the default set of advertised prefixes.";
  }
  case advertise {
    leaf valid-lifetime {
      type uint32;
      units "seconds";
      default "2592000";
      description
        "The value to be placed in the Valid Lifetime
         in the Prefix Information option. The
         designated value of all 1's (0xffffffff)
         represents infinity.";
      reference
        "RFC 4861: Neighbor Discovery for IP version 6
         (IPv6) - AdvValidLifetime";
    }
    leaf on-link-flag {
      type boolean;
      default "true";
      description
        "The value to be placed in the on-link flag
         ('L-bit') field in the Prefix Information
         option.";
      reference
        "RFC 4861: Neighbor Discovery for IP version 6
         (IPv6) - AdvOnLinkFlag";
    }
    leaf preferred-lifetime {
      type uint32;
      units "seconds";
      must ". <= ../valid-lifetime" {
        description
          "This value MUST NOT be greater than
           valid-lifetime.";
      }
    }
  }
}
```

```

        default "604800";
        description
            "The value to be placed in the Preferred
             Lifetime in the Prefix Information option.
             The designated value of all 1's (0xffffffff)
             represents infinity.";
        reference
            "RFC 4861: Neighbor Discovery for IP version 6
             (IPv6) - AdvPreferredLifetime";
    }
    leaf autonomous-flag {
        type boolean;
        default "true";
        description
            "The value to be placed in the Autonomous Flag
             field in the Prefix Information option.";
        reference
            "RFC 4861: Neighbor Discovery for IP version 6
             (IPv6) - AdvAutonomousFlag";
    }
}
}
}
}
}
}
}
}
}
}

/*
 * The subsequent data nodes are obviated and obsoleted
 * by the Network Management Datastore Architecture
 * as described in RFC 8342.
 */
augment "/if:interfaces-state/if:interface/ip:ipv6" {
    status obsolete;
    description
        "Augments interface state data with parameters of IPv6
         Router Advertisements.";
    container ipv6-router-advertisements {
        status obsolete;
        description
            "Parameters of IPv6 Router Advertisements.";
        leaf send-advertisements {
            type boolean;
            status obsolete;
            description
                "A flag indicating whether or not the router sends
                 periodic Router Advertisements and responds to
                 Router Solicitations.";
        }
    }
}

```

```
}
leaf max-rtr-adv-interval {
  type uint16 {
    range "4..1800";
  }
  units "seconds";
  status obsolete;
  description
    "The maximum time allowed between sending unsolicited
    multicast Router Advertisements from the interface.";
}
leaf min-rtr-adv-interval {
  type uint16 {
    range "3..1350";
  }
  units "seconds";
  status obsolete;
  description
    "The minimum time allowed between sending unsolicited
    multicast Router Advertisements from the interface.";
}
leaf managed-flag {
  type boolean;
  status obsolete;
  description
    "The value that is placed in the 'Managed address
    configuration' flag field in the Router Advertisement.";
}
leaf other-config-flag {
  type boolean;
  status obsolete;
  description
    "The value that is placed in the 'Other configuration' flag
    field in the Router Advertisement.";
}
leaf link-mtu {
  type uint32;
  status obsolete;
  description
    "The value that is placed in MTU options sent by the
    router. A value of zero indicates that no MTU options
    are sent.";
}
leaf reachable-time {
  type uint32 {
    range "0..3600000";
  }
  units "milliseconds";
}
```

```
status obsolete;
description
  "The value that is placed in the Reachable Time field in
  the Router Advertisement messages sent by the router. A
  value of zero means unspecified (by this router).";
}
leaf retrans-timer {
  type uint32;
  units "milliseconds";
  status obsolete;
  description
    "The value that is placed in the Retrans Timer field in the
    Router Advertisement messages sent by the router. A value
    of zero means unspecified (by this router).";
}
leaf cur-hop-limit {
  type uint8;
  status obsolete;
  description
    "The value that is placed in the Cur Hop Limit field in the
    Router Advertisement messages sent by the router. A value
    of zero means unspecified (by this router).";
}
leaf default-lifetime {
  type uint16 {
    range "0..9000";
  }
  units "seconds";
  status obsolete;
  description
    "The value that is placed in the Router Lifetime field of
    Router Advertisements sent from the interface, in seconds.
    A value of zero indicates that the router is not to be
    used as a default router.";
}
container prefix-list {
  status obsolete;
  description
    "A list of prefixes that are placed in Prefix Information
    options in Router Advertisement messages sent from the
    interface.

    By default, these are all prefixes that the router
    advertises via routing protocols as being on-link for the
    interface from which the advertisement is sent.";
  list prefix {
    key "prefix-spec";
    status obsolete;
  }
}
```

```
description
  "Advertised prefix entry and its parameters.";
leaf prefix-spec {
  type inet:ipv6-prefix;
  status obsolete;
  description
    "IPv6 address prefix.";
}
leaf valid-lifetime {
  type uint32;
  units "seconds";
  status obsolete;
  description
    "The value that is placed in the Valid Lifetime in the
    Prefix Information option. The designated value of
    all 1's (0xffffffff) represents infinity.

    An implementation SHOULD keep this value constant in
    consecutive advertisements, except when it is
    explicitly changed in configuration.";
}
leaf on-link-flag {
  type boolean;
  status obsolete;
  description
    "The value that is placed in the on-link flag ('L-bit')
    field in the Prefix Information option.";
}
leaf preferred-lifetime {
  type uint32;
  units "seconds";
  status obsolete;
  description
    "The value that is placed in the Preferred Lifetime in
    the Prefix Information option, in seconds. The
    designated value of all 1's (0xffffffff) represents
    infinity.

    An implementation SHOULD keep this value constant in
    consecutive advertisements, except when it is
    explicitly changed in configuration.";
}
leaf autonomous-flag {
  type boolean;
  status obsolete;
  description
    "The value that is placed in the Autonomous Flag field
    in the Prefix Information option.";
```

```

    }
  }
}

```

<CODE ENDS>

10. IANA Considerations

[RFC8022] registered the following namespace URIs in the "IETF XML Registry" [RFC3688]. IANA has updated the references to refer to this document.

URI: urn:ietf:params:xml:ns:yang:ietf-routing
 Registrant Contact: The IESG.
 XML: N/A; the requested URI is an XML namespace.

URI: urn:ietf:params:xml:ns:yang:ietf-ipv4-unicast-routing
 Registrant Contact: The IESG.
 XML: N/A; the requested URI is an XML namespace.

URI: urn:ietf:params:xml:ns:yang:ietf-ipv6-unicast-routing
 Registrant Contact: The IESG.
 XML: N/A; the requested URI is an XML namespace.

[RFC8022] registered the following YANG modules in the "YANG Module Names" registry [RFC6020]. IANA has updated (1) the modules per this document and (2) the references to refer to this document.

Name: ietf-routing
 Namespace: urn:ietf:params:xml:ns:yang:ietf-routing
 Prefix: rt
 Reference: RFC 8349

Name: ietf-ipv4-unicast-routing
 Namespace: urn:ietf:params:xml:ns:yang:ietf-ipv4-unicast-routing
 Prefix: v4ur
 Reference: RFC 8349

Name: ietf-ipv6-unicast-routing
 Namespace: urn:ietf:params:xml:ns:yang:ietf-ipv6-unicast-routing
 Prefix: v6ur
 Reference: RFC 8349

This document registers the following YANG submodule in the "YANG Module Names" registry [RFC6020]:

Name: ietf-ipv6-router-advertisements
Module: ietf-ipv6-unicast-routing
Reference: RFC 8349

11. Security Considerations

The YANG modules specified in this document define a schema for data that is designed to be accessed via network management protocols such as NETCONF [RFC6241] or RESTCONF [RFC8040]. The lowest NETCONF layer is the secure transport layer, and the mandatory-to-implement secure transport is Secure Shell (SSH) [RFC6242]. The lowest RESTCONF layer is HTTPS, and the mandatory-to-implement secure transport is TLS [RFC5246].

The NETCONF access control model [RFC8341] provides the means to restrict access for particular NETCONF or RESTCONF users to a preconfigured subset of all available NETCONF or RESTCONF protocol operations and content.

There are a number of data nodes defined in these YANG modules that are writable/creatable/deletable (i.e., config true, which is the default). These data nodes may be considered sensitive or vulnerable in some network environments. Write operations (e.g., edit-config) to these data nodes without proper protection can have a negative effect on network operations. These are the subtrees and data nodes and their sensitivity/vulnerability:

/routing/control-plane-protocols/control-plane-protocol: This list specifies the control-plane protocols configured on a device.

/routing/ribs/rib: This list specifies the RIBs configured for the device.

Some of the readable data nodes in these YANG modules may be considered sensitive or vulnerable in some network environments. It is thus important to control read access (e.g., via get, get-config, or notification) to these data nodes. These are the subtrees and data nodes and their sensitivity/vulnerability:

/routing/control-plane-protocols/control-plane-protocol: This list specifies the control-plane protocols configured on a device. Refer to the control-plane models for a list of sensitive information.

`/routing/ribs/rib`: This list specifies the RIBs and their contents for the device. Access to this information may disclose the network topology and/or other information.

Some of the RPC operations in this YANG module may be considered sensitive or vulnerable in some network environments. It is thus important to control access to these operations. These are the operations and their sensitivity/vulnerability:

`/routing/ribs/rib/active-route`: The output from this RPC operation returns the route that is being used for a specified destination. Access to this information may disclose the network topology or relationship (e.g., client/provider). Additionally, the routes used by a network device may be used to mount a subsequent attack on traffic traversing the network device.

12. References

12.1. Normative References

- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<https://www.rfc-editor.org/info/rfc2119>>.
- [RFC3688] Mealling, M., "The IETF XML Registry", BCP 81, RFC 3688, DOI 10.17487/RFC3688, January 2004, <<https://www.rfc-editor.org/info/rfc3688>>.
- [RFC4861] Narten, T., Nordmark, E., Simpson, W., and H. Soliman, "Neighbor Discovery for IP version 6 (IPv6)", RFC 4861, DOI 10.17487/RFC4861, September 2007, <<https://www.rfc-editor.org/info/rfc4861>>.
- [RFC5246] Dierks, T. and E. Rescorla, "The Transport Layer Security (TLS) Protocol Version 1.2", RFC 5246, DOI 10.17487/RFC5246, August 2008, <<https://www.rfc-editor.org/info/rfc5246>>.
- [RFC6020] Bjorklund, M., Ed., "YANG - A Data Modeling Language for the Network Configuration Protocol (NETCONF)", RFC 6020, DOI 10.17487/RFC6020, October 2010, <<https://www.rfc-editor.org/info/rfc6020>>.
- [RFC6241] Enns, R., Ed., Bjorklund, M., Ed., Schoenwaelder, J., Ed., and A. Bierman, Ed., "Network Configuration Protocol (NETCONF)", RFC 6241, DOI 10.17487/RFC6241, June 2011, <<https://www.rfc-editor.org/info/rfc6241>>.

- [RFC6242] Wasserman, M., "Using the NETCONF Protocol over Secure Shell (SSH)", RFC 6242, DOI 10.17487/RFC6242, June 2011, <<https://www.rfc-editor.org/info/rfc6242>>.
- [RFC6991] Schoenwaelder, J., Ed., "Common YANG Data Types", RFC 6991, DOI 10.17487/RFC6991, July 2013, <<https://www.rfc-editor.org/info/rfc6991>>.
- [RFC7950] Bjorklund, M., Ed., "The YANG 1.1 Data Modeling Language", RFC 7950, DOI 10.17487/RFC7950, August 2016, <<https://www.rfc-editor.org/info/rfc7950>>.
- [RFC8022] Lhotka, L. and A. Lindem, "A YANG Data Model for Routing Management", RFC 8022, DOI 10.17487/RFC8022, November 2016, <<https://www.rfc-editor.org/info/rfc8022>>.
- [RFC8040] Bierman, A., Bjorklund, M., and K. Watsen, "RESTCONF Protocol", RFC 8040, DOI 10.17487/RFC8040, January 2017, <<https://www.rfc-editor.org/info/rfc8040>>.
- [RFC8174] Leiba, B., "Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words", BCP 14, RFC 8174, DOI 10.17487/RFC8174, May 2017, <<https://www.rfc-editor.org/info/rfc8174>>.
- [RFC8341] Bierman, A. and M. Bjorklund, "Network Configuration Access Control Model", STD 91, RFC 8341, DOI 10.17487/RFC8341, March 2018, <<https://www.rfc-editor.org/info/rfc8341>>.
- [RFC8342] Bjorklund, M., Schoenwaelder, J., Shafer, P., Watsen, K., and R. Wilton, "Network Management Datastore Architecture (NMDA)", RFC 8342, DOI 10.17487/RFC8342, March 2018, <<https://www.rfc-editor.org/info/rfc8342>>.
- [RFC8343] Bjorklund, M., "A YANG Data Model for Interface Management", RFC 8343, DOI 10.17487/RFC8343, March 2018, <<https://www.rfc-editor.org/info/rfc8343>>.
- [RFC8344] Bjorklund, M., "A YANG Data Model for IP Management", RFC 8344, DOI 10.17487/RFC8344, March 2018, <<https://www.rfc-editor.org/info/rfc8344>>.

[W3C.REC-xml-20081126]

Bray, T., Paoli, J., Sperberg-McQueen, M., Maler, E., and F. Yergeau, "Extensible Markup Language (XML) 1.0 (Fifth Edition)", World Wide Web Consortium Recommendation REC-xml-20081126, November 2008, <<https://www.w3.org/TR/2008/REC-xml-20081126>>.

12.2. Informative References

[RFC7224] Bjorklund, M., "IANA Interface Type YANG Module", RFC 7224, DOI 10.17487/RFC7224, May 2014, <<https://www.rfc-editor.org/info/rfc7224>>.

[RFC7895] Bierman, A., Bjorklund, M., and K. Watsen, "YANG Module Library", RFC 7895, DOI 10.17487/RFC7895, June 2016, <<https://www.rfc-editor.org/info/rfc7895>>.

[RFC7951] Lhotka, L., "JSON Encoding of Data Modeled with YANG", RFC 7951, DOI 10.17487/RFC7951, August 2016, <<https://www.rfc-editor.org/info/rfc7951>>.

[RFC8340] Bjorklund, M. and L. Berger, Ed., "YANG Tree Diagrams", BCP 215, RFC 8340, DOI 10.17487/RFC8340, March 2018, <<https://www.rfc-editor.org/info/rfc8340>>.

[YANG-Guidelines]

Bierman, A., "Guidelines for Authors and Reviewers of YANG Data Model Documents", Work in Progress, draft-ietf-netmod-rfc6087bis-20, March 2018.

Appendix A. The Complete Schema Tree

This appendix presents the complete tree of the core routing data model. See [RFC8340] for an explanation of the symbols used. The data type of every leaf node is shown near the right end of the corresponding line.

```

module: ietf-routing
  +--rw routing
    |
    |   +--rw router-id?                               yang:dotted-quad
    |   +--ro interfaces
    |   |   +--ro interface*   if:interface-ref
    |   +--rw control-plane-protocols
    |   |   +--rw control-plane-protocol* [type name]
    |   |   |   +--rw type                               identityref
    |   |   |   +--rw name                               string
    |   |   |   +--rw description?                       string
    |   |   |   +--rw static-routes
    |   |   |   |   +--rw v4ur:ipv4
    |   |   |   |   |   +--rw v4ur:route* [destination-prefix]
    |   |   |   |   |   |   +--rw v4ur:destination-prefix
    |   |   |   |   |   |   |   inet:ipv4-prefix
    |   |   |   |   |   |   +--rw v4ur:description?           string
    |   |   |   |   |   |   +--rw v4ur:next-hop
    |   |   |   |   |   |   |   +--rw (v4ur:next-hop-options)
    |   |   |   |   |   |   |   |   +--:(v4ur:simple-next-hop)
    |   |   |   |   |   |   |   |   |   +--rw v4ur:outgoing-interface?
    |   |   |   |   |   |   |   |   |   |   if:interface-ref
    |   |   |   |   |   |   |   |   |   +--rw v4ur:next-hop-address?
    |   |   |   |   |   |   |   |   |   |   inet:ipv4-address
    |   |   |   |   |   |   |   |   +--:(v4ur:special-next-hop)
    |   |   |   |   |   |   |   |   |   +--rw v4ur:special-next-hop?
    |   |   |   |   |   |   |   |   |   |   enumeration
    |   |   |   |   |   |   |   |   +--:(v4ur:next-hop-list)
    |   |   |   |   |   |   |   |   |   +--rw v4ur:next-hop-list
    |   |   |   |   |   |   |   |   |   |   +--rw v4ur:next-hop* [index]
    |   |   |   |   |   |   |   |   |   |   |   +--rw v4ur:index
    |   |   |   |   |   |   |   |   |   |   |   |   string
    |   |   |   |   |   |   |   |   |   |   |   +--rw v4ur:outgoing-interface?
    |   |   |   |   |   |   |   |   |   |   |   |   |   if:interface-ref
    |   |   |   |   |   |   |   |   |   |   |   |   +--rw v4ur:next-hop-address?
    |   |   |   |   |   |   |   |   |   |   |   |   |   inet:ipv4-address
    |   |   |   |   |   +--rw v6ur:ipv6
    |   |   |   |   |   |   +--rw v6ur:route* [destination-prefix]
    |   |   |   |   |   |   |   +--rw v6ur:destination-prefix
    |   |   |   |   |   |   |   |   inet:ipv6-prefix
    |   |   |   |   |   |   |   +--rw v6ur:description?           string

```

```

+--rw v6ur:next-hop
  +--rw (v6ur:next-hop-options)
  +--:(v6ur:simple-next-hop)
  |   +--rw v6ur:outgoing-interface?
  |   |   if:interface-ref
  |   +--rw v6ur:next-hop-address?
  |   |   inet:ipv6-address
  +--:(v6ur:special-next-hop)
  |   +--rw v6ur:special-next-hop?
  |   |   enumeration
  +--:(v6ur:next-hop-list)
  |   +--rw v6ur:next-hop-list
  |   |   +--rw v6ur:next-hop* [index]
  |   |   |   +--rw v6ur:index
  |   |   |   |   string
  |   |   +--rw v6ur:outgoing-interface?
  |   |   |   if:interface-ref
  |   |   +--rw v6ur:next-hop-address?
  |   |   |   inet:ipv6-address
+--rw ribs
  +--rw rib* [name]
  |   +--rw name
  |   |   string
  |   +--rw address-family
  |   |   identityref
  |   +--ro default-rib?
  |   |   boolean {multiple-ribs}?
  |   +--ro routes
  |   |   +--ro route*
  |   |   |   +--ro route-preference?
  |   |   |   |   route-preference
  |   |   +--ro next-hop
  |   |   |   +--ro (next-hop-options)
  |   |   |   |   +--:(simple-next-hop)
  |   |   |   |   |   +--ro outgoing-interface?
  |   |   |   |   |   |   if:interface-ref
  |   |   |   |   +--ro v4ur:next-hop-address?
  |   |   |   |   |   |   inet:ipv4-address
  |   |   |   |   +--ro v6ur:next-hop-address?
  |   |   |   |   |   |   inet:ipv6-address
  |   |   |   |   +--:(special-next-hop)
  |   |   |   |   |   +--ro special-next-hop?
  |   |   |   |   |   |   enumeration
  |   |   |   +--:(next-hop-list)
  |   |   |   |   +--ro next-hop-list
  |   |   |   |   |   +--ro next-hop*
  |   |   |   |   |   |   +--ro outgoing-interface?
  |   |   |   |   |   |   |   if:interface-ref
  |   |   |   |   +--ro v4ur:address?
  |   |   |   |   |   |   inet:ipv4-address
  |   |   |   |   +--ro v6ur:address?
  |   |   |   |   |   |   inet:ipv6-address

```

```

    +--ro source-protocol          identityref
    +--ro active?                  empty
    +--ro last-updated?            yang:date-and-time
    +--ro v4ur:destination-prefix? inet:ipv4-prefix
    +--ro v6ur:destination-prefix? inet:ipv6-prefix
  +---x active-route
    +---w input
    |   +---w v4ur:destination-address? inet:ipv4-address
    |   +---w v6ur:destination-address? inet:ipv6-address
    +--ro output
    +--ro route
    |   +--ro next-hop
    |   |   +--ro (next-hop-options)
    |   |   |   +--:(simple-next-hop)
    |   |   |   |   +--ro outgoing-interface?
    |   |   |   |   |   if:interface-ref
    |   |   |   |   +--ro v4ur:next-hop-address?
    |   |   |   |   |   inet:ipv4-address
    |   |   |   |   +--ro v6ur:next-hop-address?
    |   |   |   |   |   inet:ipv6-address
    |   |   |   +--:(special-next-hop)
    |   |   |   |   +--ro special-next-hop?
    |   |   |   |   |   enumeration
    |   |   |   +--:(next-hop-list)
    |   |   |   |   +--ro next-hop-list
    |   |   |   |   |   +--ro next-hop*
    |   |   |   |   |   |   +--ro outgoing-interface?
    |   |   |   |   |   |   |   if:interface-ref
    |   |   |   |   |   |   +--ro v4ur:next-hop-address?
    |   |   |   |   |   |   |   inet:ipv4-address
    |   |   |   |   |   |   +--ro v6ur:next-hop-address?
    |   |   |   |   |   |   |   inet:ipv6-address
    |   |   |   +--ro source-protocol          identityref
    |   |   |   +--ro active?                  empty
    |   |   |   +--ro last-updated?            yang:date-and-time
    |   |   |   |   +--ro v4ur:destination-prefix?
    |   |   |   |   |   inet:ipv4-prefix
    |   |   |   |   +--ro v6ur:destination-prefix?
    |   |   |   |   |   inet:ipv6-prefix
    |   |   +--rw description?                  string
  o--ro routing-state
    o--ro router-id?                          yang:dotted-quad
    o--ro interfaces
    |   o--ro interface*      if:interface-state-ref

```

```

o--ro control-plane-protocols
|   o--ro control-plane-protocol* [type name]
|       o--ro type      identityref
|       o--ro name      string
o--ro ribs
    o--ro rib* [name]
        o--ro name          string
        o--ro address-family identityref
        o--ro default-rib?   boolean {multiple-ribs}?
        o--ro routes
            o--ro route*
                o--ro route-preference?      route-preference
                o--ro next-hop
                    o--ro (next-hop-options)
                        o--:(simple-next-hop)
                            o--ro outgoing-interface?
                                |
                                if:interface-ref
                            o--ro v4ur:next-hop-address?
                                |
                                inet:ipv4-address
                            o--ro v6ur:next-hop-address?
                                |
                                inet:ipv6-address
                        o--:(special-next-hop)
                            o--ro special-next-hop?      enumeration
                        o--:(next-hop-list)
                            o--ro next-hop-list
                                o--ro next-hop*
                                    o--ro outgoing-interface?
                                        |
                                        if:interface-ref
                                    o--ro v4ur:address?
                                        |
                                        inet:ipv4-address
                                    o--ro v6ur:address?
                                        |
                                        inet:ipv6-address
                                o--ro source-protocol      identityref
                                o--ro active?              empty
                                o--ro last-updated?        yang:date-and-time
                                o--ro v4ur:destination-prefix? inet:ipv4-prefix
                                o--ro v6ur:destination-prefix? inet:ipv6-prefix
            o---x active-route
                o---w input
                    |   o---w v4ur:destination-address?    inet:ipv4-address
                    |   o---w v6ur:destination-address?    inet:ipv6-address
                o--ro output

```



```

o--ro route
  o--ro next-hop
    o--ro (next-hop-options)
      o--:(simple-next-hop)
        o--ro outgoing-interface?
          |
          if:interface-ref
        o--ro v4ur:next-hop-address?
          |
          inet:ipv4-address
        o--ro v6ur:next-hop-address?
          |
          inet:ipv6-address
      o--:(special-next-hop)
        o--ro special-next-hop?
          |
          enumeration
      o--:(next-hop-list)
        o--ro next-hop-list
          o--ro next-hop*
            o--ro outgoing-interface?
              |
              if:interface-ref
            o--ro v4ur:next-hop-address?
              |
              inet:ipv4-address
            o--ro v6ur:next-hop-address?
              |
              inet:ipv6-address
    o--ro source-protocol
      |
      identityref
    o--ro active?
      |
      empty
    o--ro last-updated?
      |
      yang:date-and-time
    o--ro v4ur:destination-prefix?
      |
      inet:ipv4-prefix
    o--ro v6ur:destination-prefix?
      |
      inet:ipv6-prefix
module: ietf-ipv6-unicast-routing
augment /if:interfaces/if:interface/ip:ipv6:
  +--rw ipv6-router-advertisements
    +--rw send-advertisements?      boolean
    +--rw max-rtr-adv-interval?     uint16
    +--rw min-rtr-adv-interval?     uint16
    +--rw managed-flag?             boolean
    +--rw other-config-flag?        boolean
    +--rw link-mtu?                 uint32
    +--rw reachable-time?           uint32
    +--rw retrans-timer?            uint32
    +--rw cur-hop-limit?            uint8
    +--rw default-lifetime?         uint16
    +--rw prefix-list
      +--rw prefix* [prefix-spec]
      +--rw prefix-spec             inet:ipv6-prefix

```

```

    +--rw (control-adv-prefixes)?
      +--:(no-advertise)
      |   +--rw no-advertise?          empty
      +--:(advertise)
          +--rw valid-lifetime?         uint32
          +--rw on-link-flag?           boolean
          +--rw preferred-lifetime?     uint32
          +--rw autonomous-flag?        boolean
augment /if:interfaces-state/if:interface/ip:ipv6:
  o--ro ipv6-router-advertisements
    o--ro send-advertisements?          boolean
    o--ro max-rtr-adv-interval?         uint16
    o--ro min-rtr-adv-interval?         uint16
    o--ro managed-flag?                 boolean
    o--ro other-config-flag?            boolean
    o--ro link-mtu?                     uint32
    o--ro reachable-time?               uint32
    o--ro retrans-timer?                uint32
    o--ro cur-hop-limit?                uint8
    o--ro default-lifetime?             uint16
    o--ro prefix-list
      o--ro prefix* [prefix-spec]
        o--ro prefix-spec               inet:ipv6-prefix
        o--ro valid-lifetime?           uint32
        o--ro on-link-flag?             boolean
        o--ro preferred-lifetime?       uint32
        o--ro autonomous-flag?          boolean

```

Appendix B. Minimum Implementation

Some parts and options of the core routing model, such as user-defined RIBs, are intended only for advanced routers. This appendix gives basic non-normative guidelines for implementing a bare minimum of available functions. Such an implementation may be used for hosts or very simple routers.

A minimum implementation does not support the "multiple-ribs" feature. This means that a single system-controlled RIB is available for each supported address family -- IPv4, IPv6, or both. These RIBs are also the default RIBs. No user-controlled RIBs are allowed.

In addition to the mandatory instance of the "direct" pseudo-protocol, a minimum implementation should support configuring instance(s) of the "static" pseudo-protocol.

For hosts that are never intended to act as routers, the ability to turn on sending IPv6 Router Advertisements (Section 5.4) should be removed.

Platforms with severely constrained resources may use deviations for restricting the data model, e.g., limiting the number of "static" control-plane protocol instances.

Appendix C. Example: Adding a New Control-Plane Protocol

This appendix demonstrates how the core routing data model can be extended to support a new control-plane protocol. The YANG module "example-rip" shown below is intended as an illustration rather than a real definition of a data model for the Routing Information Protocol (RIP). For the sake of brevity, this module does not obey all the guidelines specified in [YANG-Guidelines]. See also Section 5.3.2.

```
module example-rip {  
  yang-version "1.1";  
  namespace "http://example.com/rip";  
  prefix "rip";  
  
  import ietf-interfaces {  
    prefix "if";  
  }  
  
  import ietf-routing {  
    prefix "rt";  
  }  
  
  identity rip {  
    base rt:routing-protocol;  
    description  
      "Identity for the Routing Information Protocol (RIP).";  
  }  
  
  typedef rip-metric {  
    type uint8 {  
      range "0..16";  
    }  
  }  
}
```

```
grouping route-content {
  description
    "This grouping defines RIP-specific route attributes.";
  leaf metric {
    type rip-metric;
  }
  leaf tag {
    type uint16;
    default "0";
    description
      "This leaf may be used to carry additional information,
      e.g., an autonomous system (AS) number.";
  }
}

augment "/rt:routing/rt:ribs/rt:rib/rt:routes/rt:route" {
  when "derived-from-or-self(rt:source-protocol, 'rip:rip')" {
    description
      "This augment is only valid for a route whose source
      protocol is RIP.";
  }
  description
    "RIP-specific route attributes.";
  uses route-content;
}

augment "/rt:routing/rt:ribs/rt:rib/rt:active-route/"
  + "rt:output/rt:route" {
  description
    "RIP-specific route attributes in the output of an
    'active-route' RPC.";
  uses route-content;
}

augment "/rt:routing/rt:control-plane-protocols/"
  + "rt:control-plane-protocol" {
  when "derived-from-or-self(rt:type, 'rip:rip')" {
    description
      "This augment is only valid for a routing protocol instance
      of type 'rip'.";
  }
  container rip {
    presence
      "RIP configuration";
    description
      "RIP instance configuration.";
    container interfaces {
```

```

    description
      "Per-interface RIP configuration.";
    list interface {
      key "name";
      description
        "RIP is enabled on interfaces that have an entry in this
        list, unless 'enabled' is set to 'false' for that
        entry.";
      leaf name {
        type if:interface-ref;
      }
      leaf enabled {
        type boolean;
        default "true";
      }
      leaf metric {
        type rip-metric;
        default "1";
      }
    }
  }
  leaf update-interval {
    type uint8 {
      range "10..60";
    }
    units "seconds";
    default "30";
    description
      "Time interval between periodic updates.";
  }
}
}
}

```

Appendix D. Data Tree Example

This section contains an example of an instance data tree from the operational state, in JSON encoding [RFC7951]. (This example includes "iana-if-type", which is defined in [RFC7224].)

The data conforms to a data model that is defined by the following YANG library specification [RFC7895]:

```
{
  "ietf-yang-library:modules-state": {
    "module-set-id": "c2e1f54169aa7f36e1a6e8d0865d441d3600f9c4",
    "module": [
      {
        "name": "ietf-routing",
        "revision": "2018-03-13",
        "feature": [
          "multiple-ribs",
          "router-id"
        ],
        "namespace": "urn:ietf:params:xml:ns:yang:ietf-routing",
        "conformance-type": "implement"
      },
      {
        "name": "ietf-ipv4-unicast-routing",
        "revision": "2018-03-13",
        "namespace":
          "urn:ietf:params:xml:ns:yang:ietf-ipv4-unicast-routing",
        "conformance-type": "implement"
      },
      {
        "name": "ietf-ipv6-unicast-routing",
        "revision": "2018-03-13",
        "namespace":
          "urn:ietf:params:xml:ns:yang:ietf-ipv6-unicast-routing",
        "conformance-type": "implement",
        "submodule": [
          {
            "name": "ietf-ipv6-router-advertisements",
            "revision": "2018-03-13"
          }
        ]
      }
    ],
    {
      "name": "ietf-interfaces",
      "revision": "2018-02-20",
      "namespace": "urn:ietf:params:xml:ns:yang:ietf-interfaces",
      "conformance-type": "implement"
    }
  }
}
```

```
    },
    {
      "name": "ietf-inet-types",
      "namespace": "urn:ietf:params:xml:ns:yang:ietf-inet-types",
      "revision": "2013-07-15",
      "conformance-type": "import"
    },
    {
      "name": "ietf-yang-types",
      "namespace": "urn:ietf:params:xml:ns:yang:ietf-yang-types",
      "revision": "2013-07-15",
      "conformance-type": "import"
    },
    {
      "name": "iana-if-type",
      "namespace": "urn:ietf:params:xml:ns:yang:iana-if-type",
      "revision": "2014-05-08",
      "conformance-type": "implement"
    },
    {
      "name": "ietf-ip",
      "revision": "2018-02-22",
      "namespace": "urn:ietf:params:xml:ns:yang:ietf-ip",
      "conformance-type": "implement"
    }
  ]
}
```

A simple network setup as shown in Figure 2 is assumed: router "A" uses static default routes with the "ISP" router as the next hop. IPv6 Router Advertisements are configured only on the "eth1" interface and disabled on the upstream "eth0" interface.




Figure 2: Example of Network Configuration

The instance data tree could then be as follows:

```

{
  "ietf-interfaces:interfaces": {
    "interface": [
      {
        "name": "eth0",
        "type": "iana-if-type:ethernetCsmacd",
        "description": "Uplink to ISP.",
        "phys-address": "00:0C:42:E5:B1:E9",
        "oper-status": "up",
        "statistics": {
          "discontinuity-time": "2015-10-24T17:11:27+02:00"
        },
        "ietf-ip:ipv4": {
          "forwarding": true,
          "mtu": 1500,
          "address": [
            {
              "ip": "192.0.2.1",
              "prefix-length": 24
            }
          ]
        }
      }
    ]
  }
}
  
```



```

    }
  ]
}
}
ietf-ip:ipv6: {
  "forwarding": true,
  "mtu": 1500,
  "address": [
    {
      "ip": "2001:0db8:0:1::1",
      "prefix-length": 64
    }
  ],
  "autoconf": {
    "create-global-addresses": false
  }
}
ietf-ipv6-unicast-routing:ipv6-router-advertisements: {
  "send-advertisements": false
}
}
},
{
  "name": "eth1",
  "type": "iana-if-type:ethernetCsmacd",
  "description": "Interface to the internal network.",
  "phys-address": "00:0C:42:E5:B1:EA",
  "oper-status": "up",
  "statistics": {
    "discontinuity-time": "2015-10-24T17:11:29+02:00"
  }
}
ietf-ip:ipv4: {
  "forwarding": true,
  "mtu": 1500,
  "address": [
    {
      "ip": "198.51.100.1",
      "prefix-length": 24
    }
  ]
}
}
ietf-ip:ipv6: {
  "forwarding": true,
  "mtu": 1500,
  "address": [
    {
      "ip": "2001:0db8:0:2::1",
      "prefix-length": 64
    }
  ]
},

```

```

    "autoconf": {
      "create-global-addresses": false
    },
    "ietf-ipv6-unicast-routing:ipv6-router-advertisements": {
      "send-advertisements": true,
      "prefix-list": {
        "prefix": [
          {
            "prefix-spec": "2001:db8:0:2::/64"
          }
        ]
      }
    }
  ],
},
"ietf-routing:routing": {
  "router-id": "192.0.2.1",
  "control-plane-protocols": {
    "control-plane-protocol": [
      {
        "type": "ietf-routing:static",
        "name": "st0",
        "description":
          "Static routing is used for the internal network.",
        "static-routes": {
          "ietf-ipv4-unicast-routing:ipv4": {
            "route": [
              {
                "destination-prefix": "0.0.0.0/0",
                "next-hop": {
                  "next-hop-address": "192.0.2.2"
                }
              }
            ]
          },
          "ietf-ipv6-unicast-routing:ipv6": {
            "route": [
              {
                "destination-prefix": "::/0",
                "next-hop": {
                  "next-hop-address": "2001:db8:0:1::2"
                }
              }
            ]
          }
        }
      }
    ]
  }
}

```

```

    }
  }
]
},
"ribs": {
  "rib": [
    {
      "name": "ipv4-master",
      "address-family":
        "ietf-ipv4-unicast-routing:ipv4-unicast",
      "default-rib": true,
      "routes": {
        "route": [
          {
            "ietf-ipv4-unicast-routing:destination-prefix":
              "192.0.2.1/24",
            "next-hop": {
              "outgoing-interface": "eth0"
            },
            "route-preference": 0,
            "source-protocol": "ietf-routing:direct",
            "last-updated": "2015-10-24T17:11:27+02:00"
          },
          {
            "ietf-ipv4-unicast-routing:destination-prefix":
              "198.51.100.0/24",
            "next-hop": {
              "outgoing-interface": "eth1"
            },
            "source-protocol": "ietf-routing:direct",
            "route-preference": 0,
            "last-updated": "2015-10-24T17:11:27+02:00"
          },
          {
            "ietf-ipv4-unicast-routing:destination-prefix":
              "0.0.0.0/0",
            "source-protocol": "ietf-routing:static",
            "route-preference": 5,
            "next-hop": {
              "ietf-ipv4-unicast-routing:next-hop-address":
                "192.0.2.2"
            },
            "last-updated": "2015-10-24T18:02:45+02:00"
          }
        ]
      }
    }
  ]
},
{

```


Appendix E. NETCONF Get Data Reply Example

This section gives an example of an XML [W3C.REC-xml-20081126] reply to the NETCONF <get-data> request for <operational> for a device that implements the example data models above.

```
<rpc-reply
  xmlns="urn:ietf:params:xml:ns:netconf:base:1.0"
  message-id="101">
  <data>
    <routing
      xmlns="urn:ietf:params:xml:ns:yang:ietf-routing"
      xmlns:or="urn:ietf:params:xml:ns:yang:ietf-origin">

      <router-id or:origin="or:intended">192.0.2.1</router-id>
      <control-plane-protocols or:origin="or:intended">
        <control-plane-protocol>
          <type>ietf-routing:static</type>
          <name>static-routing-protocol</name>
          <static-routes>
            <ietf-ipv4-unicast-routing:ipv4>
              <route>
                <destination-prefix>0.0.0.0/0</destination-prefix>
                <next-hop>
                  <next-hop-address>192.0.2.2</next-hop-address>
                </next-hop>
              </route>
            </ietf-ipv4-unicast-routing:ipv4>
            <ietf-ipv6-unicast-routing:ipv6>
              <route>
                <destination-prefix>::/0</destination-prefix>
                <next-hop>
                  <next-hop-address>2001:db8:0:1::2</next-hop-address>
                </next-hop>
              </route>
            </ietf-ipv6-unicast-routing:ipv6>
          </static-routes>
        </control-plane-protocol>
      </control-plane-protocols>

      <ribs>
        <rib or:origin="or:intended">
          <name>ipv4-master</name>
          <address-family>
            ietf-ipv4-unicast-routing:ipv4-unicast
          </address-family>
          <default-rib>true</default-rib>
          <routes>
```

```
<route>
  <ietf-ipv4-unicast-routing:destination-prefix>
    192.0.2.1/24
  </ietf-ipv4-unicast-routing:destination-prefix>
  <next-hop>
    <outgoing-interface>eth0</outgoing-interface>
  </next-hop>
  <route-preference>0</route-preference>
  <source-protocol>ietf-routing:direct</source-protocol>
  <last-updated>2015-10-24T17:11:27+02:00</last-updated>
</route>
<route>
  <ietf-ipv4-unicast-routing:destination-prefix>
    198.51.100.0/24
  </ietf-ipv4-unicast-routing:destination-prefix>
  <next-hop>
    <outgoing-interface>eth1</outgoing-interface>
  </next-hop>
  <route-preference>0</route-preference>
  <source-protocol>ietf-routing:direct</source-protocol>
  <last-updated>2015-10-24T17:11:27+02:00</last-updated>
</route>
<route>
  <ietf-ipv4-unicast-routing:destination-prefix>0.0.0.0/0
</ietf-ipv4-unicast-routing:destination-prefix>
  <next-hop>
    <ietf-ipv4-unicast-routing:next-hop-address>192.0.2.2
    </ietf-ipv4-unicast-routing:next-hop-address>
  </next-hop>
  <route-preference>5</route-preference>
  <source-protocol>ietf-routing:static</source-protocol>
  <last-updated>2015-10-24T18:02:45+02:00</last-updated>
</route>
</routes>
</rib>
<rib or:origin="or:intended">
  <name>ipv6-master</name>
  <address-family>
    ietf-ipv6-unicast-routing:ipv6-unicast
  </address-family>
  <default-rib>true</default-rib>
  <routes>
    <route>
      <ietf-ipv6-unicast-routing:destination-prefix>
        2001:db8:0:1::/64
      </ietf-ipv6-unicast-routing:destination-prefix>
      <next-hop>
        <outgoing-interface>eth0</outgoing-interface>
```

```
    </next-hop>
    <route-preference>0</route-preference>
    <source-protocol>ietf-routing:direct</source-protocol>
    <last-updated>2015-10-24T17:11:27+02:00</last-updated>
  </route>
  <route>
    <ietf-ipv6-unicast-routing:destination-prefix>
      2001:db8:0:2::/64
    </ietf-ipv6-unicast-routing:destination-prefix>
    <next-hop>
      <outgoing-interface>eth1</outgoing-interface>
    </next-hop>
    <route-preference>0</route-preference>
    <source-protocol>ietf-routing:direct</source-protocol>
    <last-updated>2015-10-24T17:11:27+02:00</last-updated>
  </route>
  <route>
    <ietf-ipv6-unicast-routing:destination-prefix>::/0
    </ietf-ipv6-unicast-routing:destination-prefix>
    <next-hop>
      <ietf-ipv6-unicast-routing:next-hop-address>
        2001:db8:0:1::2
      </ietf-ipv6-unicast-routing:next-hop-address>
    </next-hop>
    <route-preference>5</route-preference>
    <source-protocol>ietf-routing:static</source-protocol>
    <last-updated>2015-10-24T18:02:45+02:00</last-updated>
  </route>
</routes>
</rib>
</ribs>
</routing>
</data>
</rpc-reply>
```

Acknowledgments

The authors wish to thank Nitin Bahadur, Martin Bjorklund, Dean Bogdanovic, Joe Clarke, Francis Dupont, Jeff Haas, Joel Halpern, Wes Hardaker, Jia He, Sriganesh Kini, Suresh Krishnan, David Lamparter, Xiang Li, Stephane Litkowski, Andrew McGregor, Jan Medved, Thomas Morin, Tom Petch, Bruno Rijsman, Juergen Schoenwaelder, Phil Shafer, Dave Thaler, Vladimir Vassilev, Rob Wilton, Yi Yang, Derek Man-Kit Yeung, and Jeffrey Zhang for their helpful comments and suggestions.

Authors' Addresses

Ladislav Lhotka
CZ.NIC

Email: lhotka@nic.cz

Acee Lindem
Cisco Systems

Email: acee@cisco.com

Yingzhen Qu
Huawei
2330 Central Expressway
Santa Clara, CA 95050
United States of America

Email: yingzhen.qu@huawei.com