

Operating the Network Service Header (NSH) with Next Protocol "None"

Abstract

This document describes a network that supports Service Function Chaining (SFC) using the Network Service Header (NSH) with no payload data and carrying only metadata. This is achieved by defining a new NSH "Next Protocol" type value of "None".

This document illustrates some of the functions that may be achieved or enhanced by this mechanism, but it does not provide an exhaustive list of use cases, nor is it intended to be definitive about the functions it describes. It is expected that other documents will describe specific use cases in more detail and will define the protocol mechanics for each use case.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 7841.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <https://www.rfc-editor.org/info/rfc8393>.

Copyright Notice

Copyright (c) 2018 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	3
2. Requirements Language	4
3. The Network Service Header	4
3.1. Next Protocol "None"	4
4. Processing Rules	4
5. Backward Compatibility	5
6. Overview of Use Cases	6
6.1. Per-SFP Metadata	6
6.2. Per-Flow Metadata	7
6.3. Coordination between SFC-Aware SFIs	7
6.4. Operations, Administration, and Maintenance (OAM)	8
6.5. Control-Plane and Management-Plane Uses	8
6.6. Non-applicable Use Cases	9
7. Management and Congestion Control Considerations	9
8. Security Considerations	10
9. IANA Considerations	11
10. References	11
10.1. Normative References	11
10.2. Informative References	11
Acknowledgements	12
Contributors	12
Authors' Addresses	12

1. Introduction

An architecture for Service Function Chaining (SFC) is presented in [RFC7665]. That architecture enables packets to be forwarded along Service Function Paths (SFPs) to pass through various Service Functions (SFs) that act on the packets. Each packet is encapsulated with a Network Service Header (NSH) [RFC8300] that identifies the SFP that the packet travels along (by means of a Service Path Identifier -- SPI) and the hop (i.e., the next SF to be executed) along the SFP that the packet has reached (by means of a Service Index -- SI). The SPI and SI are fields encoded in the NSH.

Packets are classified at the SFC network ingress boundaries by classifiers (Section 4.4 of [RFC7665]) and have an NSH applied to them. Such packets are forwarded between Service Function Forwarders (SFFs) using tunnels across the underlay network, and each SFF may hand the packet off to one or more Service Function Instances (SFIs) according to the definition of the SFP.

The SFC classifier or any SFC-aware SFI may wish to share information (possibly state information) about the SFP, the traffic flow, or a specific packet, and they may do this by adding metadata to packets as part of the NSH. Metadata may be used to enhance or enable the function performed by SFC-aware SFs, may enable coordination and data exchange between SFIs, or may be used to assist a network operator in the diagnosis and monitoring of an SFP. The nature of metadata to be supplied and consumed is implementation- and deployment-specific.

This document defines a mechanism for metadata to be carried on an SFP without the need for payload data. This mechanism enables diagnosis and monitoring of SFPs, and coordination between SFC-aware SFIs. The mechanism can be applied without the need for traffic to be flowing; if traffic is flowing, it can be applied without the need to insert what might be substantial amounts of metadata into data packets (an operation that may be costly in some hardware).

This document describes how this function is achieved through the use of a new value for the NSH "Next Protocol" field to indicate "None". Like any NSH packets, such packets are contained within the SFC-enabled domain.

This document illustrates some of the functions that may be achieved or enhanced by this mechanism, but it does not provide an exhaustive list of use cases, nor is it intended to be definitive about the functions it describes (see Section 6).

This document uses the terms defined in [RFC7665] and [RFC8300].

2. Requirements Language

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in BCP 14 [RFC2119] [RFC8174] when, and only when, they appear in all capitals, as shown here.

3. The Network Service Header

The NSH includes a field called "Next Protocol" that is used to indicate the nature of the payload data that follows the NSH. The field can be used by any component that processes the NSH (for example, to understand how to interpret and parse the payload) and by nodes at the end of the SFP that remove the NSH and forward the payload data.

3.1. Next Protocol "None"

This document defines a new value (0x00) for the "Next Protocol" field to indicate that the next protocol is "None", which means that there is no user/payload data following the NSH.

When the next protocol is "None", the rest of the NSH still has meaning; in particular, the metadata carried in the NSH may still be present. It is not intended that a packet with next protocol set to "None" be sent with no metadata (see Section 4). Thus, an SFC-aware node SHOULD NOT create a packet with "Next Protocol" set to "None", Metadata Type set to 0x2, and with an NSH Length of 0x2.

4. Processing Rules

A packet with no payload data may be inserted at the head end of an SFP (such as at a classifier) and may be easily forwarded by an SFF or SFI on the SFP using the processing rules defined in [RFC8300].

A packet with no payload may also be generated by an SFC-aware SFI as a result of processing an incoming packet (i.e., triggered by a condition arising from processing a normal NSH packet with a payload). In such cases, the SPI/SI can be inherited from the original packet or can be set according to information supplied in one of three ways:

- o through the control plane,
- o through the management plane, or
- o through information carried in the metadata of the data packet.

This document does not further specify the triggers to generate an NSH packet with a "Next Protocol" set to "None".

An SFC-aware node wishing to send metadata without a data packet (i.e., a node that conforms to this specification):

- o MUST create a packet carrying an NSH and the desired metadata.
- o MUST set the "Next Protocol" field to 0x00.
- o SHOULD ensure that there are no bytes following the end of the NSH (i.e., that there is no payload data).
- o MUST encapsulate and send the packet as normal for tunneling to the next hop on the SFP as would be done for any NSH packet (i.e., for a data packet following the SFP).

A transit node (SFF, SFI, or classifier) that conforms to this specification and that receives a packet with "Next Protocol" indicating "None" MUST NOT attempt to parse or process beyond the end of the NSH, but it SHOULD process the NSH and the metadata as normal. Processing for nodes that do not support "Next Protocol" set to "None" is described in Section 5. Note, however, that an intermediate node that is instructed to strip all metadata from packets will create a packet with an NSH but no metadata and no payload. Such packets SHOULD NOT continue to be forwarded along the SFP.

A node that is the egress of an SFP would normally strip the NSH and forward the payload according to the setting of the "Next Protocol" field. Such nodes MUST NOT forward packets with "Next Protocol" indicating "None" even if there are some bytes after the NSH.

In deployments where use of next protocol "None" is not desired, administrators SHOULD instruct SFC-aware nodes to not create such packets and to discard packets with next protocol "None".

5. Backward Compatibility

SFC-aware nodes that do not understand the meaning of a value contained in the "Next Protocol" field of the NSH are unable to parse the payload. Such nodes silently drop packets with unknown "Next Protocol" values unless explicitly configured to forward them (Section 2.2 of [RFC8300]).

This means that legacy SFC-aware nodes that are unaware of the meaning of the "Next Protocol" value "None" will act as follows:

- o SFFs can be configured to forward the packets.
- o SFC proxies will drop the packets.
- o SFIs will most likely drop the packets.
- o Classifiers (i.e., nodes performing reclassification) will most likely drop the packets.

SFC-aware nodes at the end of an SFP possibly forward packets with no knowledge of the payload in a "pop and forward" form of processing where the NSH is removed, the packet is simply put on an interface, and the payload protocol is known a priori (or assumed). It is a general processing rule for all packet forwarding engines that they should not attempt to send packets with zero length. Packets with the NSH "Next Protocol" field set to "None" are expected to have zero payload length and so should not be forwarded once the NSH has been stripped. In any case, as noted in Section 4, SFC-aware nodes at the end of an SFP do not forward packets with "Next Protocol" set to "None".

6. Overview of Use Cases

6.1. Per-SFP Metadata

Per-SFP metadata is metadata that applies to an SFP and any data packets on that SFP. It does not need to be transmitted with every packet, but it can be installed at the points of consumption (such as at SFIs) and applied to all packets on the SFP as they pass through this point. It could be installed by inclusion in the NSH of a data packet sent on the SFP by out-of-band control- or management-plane mechanisms, or by separate metadata-only packets using "Next Protocol" set to "None" as described in this document.

Per-SFP metadata-only packets may be sent along the path of an SFP simply by setting the correct SPI in the NSH, and setting the SI to the correct value for the hop of the SFP at which the metadata is to be introduced. SFC-aware nodes (e.g., classifiers) will know the correct SI values to be used from information supplied by the control or management plane as is the case for NSH packets with payload data.

6.2. Per-Flow Metadata

Per-flow metadata is metadata that applies to a subset of the packets on an SFP, such as packets matching a particular 5-tuple of source address, destination address, source port, destination port, and payload protocol. Also, this metadata does not need to be transmitted with every packet, but it can be installed at the SFIs on the SFP and applied to the packets that match the flow description.

If there is just one flow on an SFP, then there is no difference between per-flow metadata and per-SFP metadata as described in Section 6.1.

In normal processing, the flow to which per-flow metadata applies can be deduced by looking at the payload data in the context of the value of the "Next Protocol" field. However, when "Next Protocol" indicates "None", this cannot be done. In this case, the identity of the flow is carried in the metadata itself.

6.3. Coordination between SFC-Aware SFIs

A pair of SFC-aware SFIs (adjacent or not) on an SFP may desire to coordinate state and may do this by sending information encoded in metadata.

To do this using the mechanisms defined in this document:

- o There must be an SFP that passes through the two SFIs in the direction of sender to receiver.
- o The sender must know the correct SPI to use.
- o The sender must know the correct SI to use for the point at which it resides on the SFP.
- o Ideally, the receiver will know to remove the packet from the SFP and not forward it further as this might share metadata wider than desirable and would cause unnecessary packets in the network. Note, however, that continued forwarding of such packets would not be substantially harmful in its own right.

Note that technically (according to the SFC architecture) the process of inserting a packet into an SFP is performed by a classifier. However, a classifier may be co-resident with an SFI so that an implementation of an SF may also be able to generate NSH packets as described in this document.

Note also that a system with SFIs that need to coordinate between each other may be configured so that there is a specific, dedicated SFP between those service functions that is used solely for this purpose. Thus, such an SFI does not need to insert NSH packets onto SFPs used to carry payload data, but it can use (and know the SPI of) this special, dedicated SFP.

6.4. Operations, Administration, and Maintenance (OAM)

Requirements for Operations, Administration, and Maintenance (OAM) in SFC networks are discussed in [SFC-OAM-FRAME]. The NSH definition in [RFC8300] includes an O bit that indicates that the packet contains OAM information.

If OAM information is carried in packets that also include payload data, that information might be carried between the NSH and the payload. Therefore, the mechanism defined in this document can also be used to carry OAM information independent of payload data.

Sending OAM separate from (but interleaved with) packets that carry payload data may have several advantages including:

- o Sending OAM when there is no other traffic flowing.
- o Sending OAM at predictable intervals.
- o Measuring path qualities distinct from behavior of SFIs.
- o Sending OAM without needing to rewrite payload data buffers.
- o Keeping OAM processing components separate from other processing components.

Mechanisms for providing active OAM [RFC7799] in an SFC network have been proposed [OAM-SFC]. This use case is not intended to define another mechanism for active OAM, but it does illustrate a further option for discussion by the working group.

6.5. Control-Plane and Management-Plane Uses

As described in Section 6.3, SFPs can be established specifically to carry metadata-only packets. And as described in Section 6.1, metadata-only packets can be sent down existing SFPs. This means that metadata-only packets can be used to carry control-plane and management-plane messages used to control and manage the SFC network.

In effect, SFPs can be established to serve as a Data Control Network (DCN) or a Management Control Network (MCN). Further details of this process are out of scope for this document, but it should be understood that, just as for OAM, an essential feature of using a control channel is that the various speakers are assigned identifiers (i.e., addresses). In this case, those identifiers could be SPI/SI pairs or could be IP addresses as used in the normal control and management plane of the SFC network.

6.6. Non-applicable Use Cases

Per-packet metadata is metadata that applies specifically to a single payload packet. It informs an SFI how to handle the payload packet and does not apply to any other packet.

The mechanisms described in this document are not applicable to per-packet metadata because, by definition, if the "Next Protocol" indicates "None", then there is no packet following the NSH for the metadata to be associated with.

7. Management and Congestion Control Considerations

The mechanisms described in this document allow SFC-aware nodes in an SFC network to generate additional packets. These are not intended to be sent frequently for any flow, but there is still a risk that they might flood the network. For example, if an attempt is made to use this mechanism for "per-packet metadata" (see Section 6.6) then this might double the number of packets in the network. Similarly, if this mechanism is used for a form of aliveness detection OAM that requires very frequent test messages, then the number of additional messages may be very high. Such additional messages risk causing congestion in the network.

The underlay network (that is, the tunnels across the underlay between SFC nodes) will not distinguish between data-carrying packets and those packets with "Next Protocol" set to "None". All packets will be treated the same and will need to fall within the capabilities of the underlay network to process and forward packets.

Nodes in the SFC overlay network will need to perform special processing on the additional packets according to their roles and according to the application for the metadata. For example, an SFF will likely only have to forward per-SFP metadata, while an SF will need to extract it and process it as it would if the metadata was carried in a packet with user data. On the other hand, metadata might also be used to cause actions at all nodes (see Sections 6.3, 6.4, and 6.5) and could increase the processing load.

In view of these potential issues, all implementations **SHOULD** implement rate limits on the generation of per-SFP packets with "Next Protocol" set to "None". Furthermore, these rate limits **SHOULD** be configurable and applied per SFP and per application so that one application on one SFP does not encumber a different application on this or a different SFP. When an implementation finds that it is unable to generate or send a packet, it **SHOULD** increment a counter that is accessible by the operator and **MAY** raise an alert (although such alerts **SHOULD**, themselves, be rate limited).

Additionally, an SFC node needs to protect itself against another node in the network not applying suitable rate limits. Therefore, implementations **SHOULD** apply incoming rate limits for SFC packets with "Next Protocol" set to "None". Such rate limits **MAY** be application aware, per SFC or interface, and **SHOULD** be configurable, but implementations **MAY** be more subtle if they are aware of internal processing loads and have access to queues/buffers. In any case, when an implementation drops a received packet because of these rate limits, it **SHOULD** increment a counter that is accessible by the operator and **MAY** raise an alert (although such alerts **SHOULD**, themselves, be rate limited).

Suitable default rate limits will restrict an SFC node to not send more than one packet with "Next Protocol" set to "None" per ten data packets on any flow in a unit of time equal to the end-to-end delivery time on the flow.

8. Security Considerations

Metadata-only packets as enabled by this document provide a covert channel. However, this is only different from the metadata feature in the normal NSH in that it can be sent without the presence of a data flow.

Metadata may, of course, contain sensitive data and may also contain information used to control the behavior of SFIs in the network. As such, this data needs to be protected according to its value and according to the perceived vulnerabilities of the network. Protection of metadata may be achieved by using encrypted transport between SFC entities or by encrypting the metadata in its own right, and by authenticating the sender of the metadata. The need to protect the metadata is not modified by this document and forms part of the NSH definition found in [RFC8300].

The mechanism described in this document might be used to introduce packets into the SFC overlay network and might be used to illegitimately introduce false metadata to the nodes on an SFC.

Therefore, measures **SHOULD** be taken to ensure authorization of sources of such packets, and tunneling of such packets into the network **SHOULD** be prevented.

The amount of packets with "Next Protocol" set to "None" on an SFP **SHOULD** be rate limited at each point on the SFP to provide additional network security.

Further discussion of NSH security is presented in [RFC8300].

9. IANA Considerations

IANA maintains a registry called "Network Service Header (NSH) Parameters" with a sub-registry called "NSH Next Protocol". IANA has allocated a new value to the sub-registry as follows:

Next Protocol	Description	Reference
0x00	None	RFC 8393

10. References

10.1. Normative References

- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<https://www.rfc-editor.org/info/rfc2119>>.
- [RFC8174] Leiba, B., "Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words", BCP 14, RFC 8174, DOI 10.17487/RFC8174, May 2017, <<https://www.rfc-editor.org/info/rfc8174>>.
- [RFC8300] Quinn, P., Ed., Elzur, U., Ed., and C. Pignataro, Ed., "Network Service Header (NSH)", RFC 8300, DOI 10.17487/RFC8300, January 2018, <<https://www.rfc-editor.org/info/rfc8300>>.

10.2. Informative References

- [OAM-SFC] Mirsky, G., Meng, W., Khasnabish, B., and C. Wang, "Multi-Layer Active OAM for Service Function Chains in Networks", Work in Progress, draft-wang-sfc-multi-layer-oam-10, September 2017.

- [RFC7665] Halpern, J., Ed. and C. Pignataro, Ed., "Service Function Chaining (SFC) Architecture", RFC 7665, DOI 10.17487/RFC7665, October 2015, <<https://www.rfc-editor.org/info/rfc7665>>.
- [RFC7799] Morton, A., "Active and Passive Metrics and Methods (with Hybrid Types In-Between)", RFC 7799, DOI 10.17487/RFC7799, May 2016, <<https://www.rfc-editor.org/info/rfc7799>>.
- [SFC-OAM-FRAME]
Aldrin, S., Pignataro, C., Kumar, N., Akiya, N., Krishnan, R., and A. Ghanwani, "Service Function Chaining (SFC) Operation, Administration and Maintenance (OAM) Framework", Work in Progress, draft-ietf-sfc-oam-framework-04, March 2018.

Acknowledgements

Thanks to the attendees at the SFC interim meeting in Westford, Massachusetts in January 2017 for discussions that suggested the value of this document.

Thanks to Eric Rosen, Med Boucadair, Greg Mirsky, Dave Dolson, Tal Mizrahi, and Mirja Kuehlewind for valuable review comments.

Contributors

Lucy Yong
Retired

Authors' Addresses

Adrian Farrel
Juniper Networks

Email: afarrel@juniper.net

John Drake
Juniper Networks

Email: jdrake@juniper.net