Time Server

This note describes the Time Server protocol, as currently implemented
on ITS hosts (i.e. MIT-(AI/ML/MC/DMS)). The idea is to provide a
site-independent, machine readable date and time in as efficient and
swift a manner as possible; its motivation arises from the fact that not
all systems have a date/time clock, and all are subject to occasional
human or machine error. The use of time-servers makes it possible to
quickly confirm or correct a system's idea of the time, by making a
brief poll of several independent sites on the network.

In particular the network time server works as follows:

    S: Listen on socket 37 (45 octal).

    U: Connect to socket 37 [not ICP].

    S: Send the time as a 32 bit binary number.

    U: Close the connection.

    S: Close the connection.

Note that this is not a normal ICP sequence. Rather, instead of sending
a 32-bit socket number to reconnect to, the server returns a 32-bit time
value and closes the connection. If the server is unable to determine
the time at its site, it should either refuse the connection or close it
without sending anything.

The time is the number of seconds since 0000 (midnight) 1 January 1900
GMT, such that the time 1 is 12:00:01 am on 1 January 1900 GMT; this
base will serve until the year 2036. As a further example, the most
recent leap year as of this writing began from the time 2,398,291,200
which corresponds to 0000 1 Jan 1976 GMT.