Computer Mail Meeting Notes

Introduction

   A meeting was held on the 11th of January 1982 at USC Information
   Sciences Institute to discuss addressing issues in computer mail.
   The attendees are listed at the end of this memo. The major
   conclusion reached at the meeting is to extend the
   "username@hostname" mailbox format to "username@host.domain", where
   the domain itself can be further structured.

Overview

   The meeting opened with a brief discussion of the objectives of the
   meeting and a review of the agenda.

      The meeting was called to discuss a few specific issues in text
      mail systems for the ARPA Internet. In particular, issues of
      addressing are of major concern as we develop an internet in which
      mail relaying is a common occurance. We need to discuss
      alternatives in the design of the mail system to provide high
      utility at reasonable cost. One scheme suggested is to create
      "mail domains" which are another level of addressing. The ad hoc
      scheme of source routing, while effective for some cases, is seen
      to lead to some problems. A key test of addressing schemes is the
      procedure for sending copies of a reply to a message to the people
      who received copies of the original message. The key reference
      documents for the meeting were RFCs 788, 799, and 801.

   Jon Postel gave a brief review of the NCP-to-TCP transition plan (RFC
   801). The emphasis was on mail, the internet host table, and the
   role of a Host Name Server.

   The major part of the meeting was devoted to a wide ranging
   discussion of the general mailbox identification problem. In
   particular, the notion of a hierarchial structure of name domains was
   discussed, and the issues associated with name servers were discussed
   including the types of information name servers should provide.

Name Domains

   One of the interesting ideas that emerged from this discussion was
   that the "user@host" model of a mailbox identifier should, in

principle, be replaced by a "unique-id@location-id" model, where the
unique-id would be a globally unique id for this mailbox (independent
of location) and the location-id would be advice about where to find
the mailbox. However, it was recognized that the "user@host" model
was well established and that so many different elaborations of the
"user" field were already in use that there was no point in persuing
this "unique-id" idea at this time.

Several alternatives for the structuring and ordering of the
extensions to the "host" field to make it into a general
"location-id" were discussed.

    These basically involved adding more hierarchical name information
    either to the right or the left of the @, with the "higher order"
    portion rightmost or leftmost. It was clear that the information
    content of all these syntactic alternatives was the same, so that
    the one causing least difficulty for existing systems should be
    chosen. Hence it was decided to add all new information on the
    right of the @ sign, leaving the "user" field to the left
    completely to each system to determine (in particular to avoid the
    problem that some systems already use dot (.) internally as part
    of user names).

The conclusion in this area was that the current "user@host" mailbox
identifier should be extended to "user@host.domain" where "domain"
could be a hierarchy of domains.

    In particular, the "host" field would become a "location" field
    and the structure would read (left to right) from the most
    specific to the most general.

        For example: "Postel@F.ISI.IN" might be the mailbox of Jon
        Postel on host F in the ISI complex of the Internet domain.

    Formally, in RFC733, the host-indicator definition rule would
    become:

    host indicator = ("at" / "@") domains

    domains = node / node "." domains

        Note only one "at" or "@" is allowed, and that the domains
        form a hierarchy with the most general in scope last.

        And note that the choice of domain names must be
        administratively controlled and the highest level domain
        names must be globally unique.

The hierarchial domain type naming differs from source routing in that the former gives absolute addressing while the latter gives relative adressing.

Name Servers

The discussion of name servers identified three separate name server functions: "white pages", "unique-id to location-id", and "location-id to address".

The "white pages" service is a way of looking up a user by name and other properties using pattern matching and may return several data base "hits". Each hit must have an associated unique-id.

The "unique-id to location-id" service returns the character string location-id where the unique-id is currently found.

The "location-id to address" service returns a network address (numeric) corresponding to the location-id.

If the location-id is the name of a host in the current domain it is clear that the address returned will be the address to send the mail to, but if the location-id is that of some other domain then the address returned may be either the address to send the mail to, or the address of a name server for that domain, and these two cases must be distinguished.

The conclusion of this discussion was that a location-id to address name service must be defined soon. The other types of name servers were not further discussed, and are not required in the implemenation.

Another aspect of the name server is returning additional information besides the address. In particular, for mail it is important to know which mail procedures the destination implements (NCP/FTP, TCP/SMTP, etc.). Two approaches were discussed: one is coding the information as service names (e.g., NCP/SMTP), and the other is by reference to protocol and port numbers (e.g., PROTOCOL=6, PORT=25). Another suggestion was that the request ought to be "location-id,service" (e.g., "ISIF.IN,MAIL") and the response ought to be the location-id, address, protocol, and port. A different way of getting this information was suggested that instead of (or in addition to) having this information in the name server, one should get this data from the host itself via some sort of query or "who are you" protocol.

Also discussed was the initial provision for name service. It seems useful to start with a text file that can be accessed via FTP, and to have both "Telnet-Like" (i.e., based on TCP) and "Datagram" (i.e.,

based on UDP) access to a query server. This might be possible as an
extension of the IEN-116 name server.

Another issue was the central vs. distributed implementation of the
name look up service. It is recognized that separate servers for
each domain has administrative and maintenance advantages, but that a
central server may be a useful first step. It is also recognized
that each distinct database should be replicated a few times and be
avialiable from distinct servers for robust and reliable service.

An Example:

   Suppose that the new mailbox specification is of the form
   USER@HOST.ORG.DOMAIN.

      e.g., Postel@F.ISI.IN

   A source host sending mail to this address first queries a name
   server for the domain IN (giving the whole location "F.ISI.IN").

   The result of the query is either (1) the final address of the
   destination host (F.ISI), or (2) the address of a name server for
   ISI, or (3) the address of a forwarder for ISI. In cases 1 and 3,
   the source host sends the mail to the address returned. In case
   2, the source host queries the ISI name server and ... (recursive
   call to this paragraph).

Action Items:

   RFC 733 Revision

      To include the hierarchial host and domain naming procedure, and
      to delete the features decommitted at the Computer Mail meeting on
      10-JAN-79.

      By: Dave Crocker

      Due: 15-Feb-82

   Host Name Server Description

      To specify a way to get name to address conversions and to find
      out about services offered. Also how to get info on domain names.

      By: Jon Postel

      Due: 15-Feb-82

Transition Plan Revision

    To include new host and domain names.

    By: Jon Postel

    Due: 15-Feb-82

SMTP Revision

    To include new host and domain names.

    By: Jon Postel

    Due: Unspecified

Mail System Description Revision

    How to do mail systems, including use of SMTP and Host Name
    Server.

    By: Jon Postel

    Due: Unspecified

Conversion of User Programs and Mailer Programs.

    Programs have to handle dots in the "host" field. Many programs
    on many hosts will have to be modified to a greater or lesser
    extent. In many cases the modifications should be quite simple.

    By: A Cast of Thousands

    Due: Unspecified (See the Following Item)

Set a date when it ok to send messages with dots in "host" field.

    The must be a date after which it is ok to send host fields with
    dots throughout the ARPANET and Internet world without the
    recipients complaining.

    By: DARPA (Duane Adams)

    Due: 1-Mar-82

Attendees:

Duane A. Adams	DARPA/IPTO	Adams@ISI	(202) 694-8096
Vint Cerf	DARPA/IPTO	Cerf@ISI	(202) 694-3049
Harry Forsdick	BBN	Forsdick@BBN	(617) 497-3638
Eric Schienbrood	BBN	shienbrood@bbn-unix	(617) 497-3756
Bob Thomas	BBN	BThomas@BBND	(617) 497-3483
Bob Fabry	Berkeley	Fabry@Berkeley	(415) 642-2714
Bill Joy	Berkeley	unj@berkeley	(415) 642-7780
Gene Ball	CMU	Ball@CMUA	(412) 578-2569
Anil Agarwal	COMSAT	Agarwal@ISID	(301) 863-6103
David L. Mills	COMSAT	Mills@ISID	(202) 863-6092
Dave Crocker	Univ. Del	DCrocker@Udel	(302) 738-8913
Ray McFarland	DoD	McFarland@ISIA	(301) 796-6290
Dave Lebling	MIT	PDL@MIT-XX	(617) 253-1440
Paul Mockapetris	ISI	Mockapetris@ISIF	(213) 822-1511
Jon Postel	ISI	Postel@ISIF	(213) 822-1511
Carl Sunshine	ISI	Sunshine@ISIF	(213) 822-1511
Mark Crispin	Stanford U.	Admin.MRC@SCORE	(415) 497-1407
Bob Braden	UCL[A]	braden@ISIA	(uk) (01)387-7050
Steve Kille	UCL	UCL-Netwiz@ISIE	(uk) (01)387-7050
Bill Tuck	UCL	UKSAT@ISIE	(uk) (01)387-7050
Marv Solomon	Univ. Wisc	Solomon@UWisc	
Ed Taft	Xerox Parc	Taft@Parc-Maxc	(415) 494-4419