

A Compact Representation of IPv6 Addresses

Status of this Memo

This memo provides information for the Internet community. This memo does not specify an Internet standard of any kind. Distribution of this memo is unlimited.

1. Abstract

IPv6 addresses, being 128 bits long, need 32 characters to write in the general case, if standard hex representation, is used, plus more for any punctuation inserted (typically about another 7 characters, or 39 characters total). This document specifies a more compact representation of IPv6 addresses, which permits encoding in a mere 20 bytes.

2. Introduction

It is always necessary to be able to write in characters the form of an address, though in actual use it is always carried in binary. For IP version 4 (IP Classic) the well known dotted quad format is used. That is, 10.1.0.23 is one such address. Each decimal integer represents a one octet of the 4 octet address, and consequently has a value between 0 and 255 (inclusive). The written length of the address varies between 7 and 15 bytes.

For IPv6 however, addresses are 16 octets long [IPv6], if the old standard form were to be used, addresses would be anywhere between 31 and 63 bytes, which is, of course, untenable.

Because of that, IPv6 had chosen to represent addresses using hex digits, and use only half as many punctuation characters, which will mean addresses of between 15 and 39 bytes, which is still quite long. Further, in an attempt to save more bytes, a special format was invented, in which a single run of zero octets can be dropped, the two adjacent punctuation characters indicate this has happened, the number of missing zeroes can be deduced from the fixed size of the address.

In most cases, using genuine IPv6 addresses, one may expect the address as written to tend toward the upper limit of 39 octets, as long strings of zeroes are likely to be rare, and most of the other

groups of 4 hex digits are likely to be longer than a single non-zero digit (just as MAC addresses typically have digits spread throughout their length).

This document specifies a new encoding, which can always represent any IPv6 address in 20 octets. While longer than the shortest possible representation of an IPv6 address, this is barely longer than half the longest representation, and will typically be shorter than the representation of most IPv6 addresses.

3. Current formats

[AddrSpec] specifies that the preferred text representation of IPv6 addresses is in one of three conventional forms.

The preferred form is x:x:x:x:x:x:x:x, where the 'x's are the hexadecimal values of the eight 16-bit pieces of the address.

Examples:

FEDC:BA98:7654:3210:FEDC:BA98:7654:3210 (39 characters)

1080:0:0:0:8:800:200C:417A (25 characters)

The second, or zero suppressed, form allows "::" to indicate multiple groups of suppressed zeroes, hence:

1080:0:0:0:8:800:200C:417A

may be represented as

1080::8:800:200C:417A

a saving of just 5 characters from this typical address form, and still leaving 21 characters.

In other cases the saving is more dramatic, in the extreme case, the address:

0:0:0:0:0:0:0:0

that is, the unspecified address, can be written as

::

This is just 2 characters, which is a considerable saving. However such cases will rarely be encountered.

The third possible form mixes the new IPv6 form with the old IPv4 form, and is intended mostly for transition, when IPv4 addresses are embedded into IPv6 addresses. These can be considerably longer than the longest normal IPv6 representation, and will eventually be phased out. Consequently they will not be considered further here.

4. The New Encoding Format

The new standard way of writing IPv6 addresses is to treat them as a 128 bit integer, encode that in base 85 notation, then encode that using 85 ASCII characters.

4.1. Why 85?

2^{128} is 340282366920938463463374607431768211456. 85^{20} is 387595310845143558731231784820556640625, and thus in 20 digits of base 85 representation all possible 2^{128} IPv6 addresses can clearly be encoded.

84^{20} is 305904398238499908683087849324518834176, clearly not sufficient, 21 characters would be needed to encode using base 84, this wastage of notational space cannot be tolerated.

On the other hand, 94^{19} is just 30862366077815087592879016454695419904, also insufficient to encode all 2^{128} different IPv6 addresses, so 20 characters would be needed even with base 94 encoding. As there are just 94 ASCII characters (excluding control characters, space, and del) base 94 is the largest reasonable value that can be used. Even if space were allowed, base 95 would still require 20 characters.

Thus, any value between 85 and 94 inclusive could reasonably be chosen. Selecting 85 allows the use of the smallest possible subset of the ASCII characters, enabling more characters to be retained for other uses, eg, to delimit the address.

4.2. The Character Set

The character set to encode the 85 base85 digits, is defined to be, in ascending order:

'0'..'9', 'A'..'Z', 'a'..'z', '!', '#', '\$', '%', '&', '(', ')', '*', '+', '-', ';', '<', '=', '>', '?', '@', '^', '_', '\'', '{', '|', '}', and '~'.

This set has been chosen with considerable care. From the 94 printable ASCII characters, the following nine were omitted:

'"' and '"', which allow the representation of IPv6 addresses to be quoted in other environments where some of the characters in the chosen character set may, unquoted, have other meanings.

',' to allow lists of IPv6 addresses to conveniently be written, and '.' to allow an IPv6 address to end a sentence without requiring it to be quoted.

'/' so IPv6 addresses can be written in standard CIDR address/length notation, and ':' because that causes problems when used in mail headers and URLs.

'[' and ']', so those can be used to delimit IPv6 addresses when represented as text strings, as they often are for IPv4,

And last, '\\', because it is often difficult to represent in a way where it does not appear to be a quote character, including in the source of this document.

5. Converting an IPv6 address to base 85.

The conversion process is a simple one of division, taking the remainders at each step, and dividing the quotient again, then reading up the page, as is done for any other base conversion.

For example, consider the address shown above

1080:0:0:0:8:800:200C:417A

In decimal, considered as a 128 bit number, that is 21932261930451111902915077091070067066.

As we divide that successively by 85 the following remainders emerge: 51, 34, 65, 57, 58, 0, 75, 53, 37, 4, 19, 61, 31, 63, 12, 66, 46, 70, 68, 4.

Thus in base85 the address is:

4-68-70-46-66-12-63-31-61-19-4-37-53-75-0-58-57-65-34-51.

Then, when encoded as specified above, this becomes:

4)+k&C#VzJ4br>0wv%Yp

This procedure is trivially reversed to produce the binary form of the address from textually encoded format.

6. Additional Benefit

Apart from generally reducing the length of an IPv6 address when encode in a textual format, this scheme also has the benefit of returning IPv6 addresses to a fixed length representation, leading zeroes are never omitted, thus removing the ugly and awkward variable length representation that has previously been recommended.

7. Implementation Issues

Many current processors do not find 128 bit integer arithmetic, as required for this technique, a trivial operation. This is not considered a serious drawback in the representation, but a flaw of the processor designs.

It may be expected that future processors will address this defect, quite possibly before any significant IPv6 deployment has been accomplished.

8. Security Considerations

By encoding addresses in this form, it is less likely that a casual observer will be able to immediately detect the binary form of the address, and thus will find it harder to make immediate use of the address. As IPv6 addresses are not intended to be learned by humans, one reason for which being that they are expected to alter in comparatively short timespan, by human perception, the somewhat challenging nature of the addresses is seen as a feature.

Further, the appearance of the address, as if it may be random gibberish in a compressed file, makes it much harder to detect by a packet sniffer programmed to look for bypassing addresses.

9. References

- [IPv6] Internet Protocol, Version 6 (IPv6) Specification,
 S. Deering, R. Hinden, RFC 1883, January 4, 1996.
- [AddrSpec] IP Version 6 Addressing Architecture,
 R. Hinden, S. Deering, RFC 1884, January 4, 1996.

10. Author's Address

Robert Elz
Computer Science
University of Melbourne
Parkville, Victoria, 3052
Australia

EMail: kre@munniari.OZ.AU