Network Working Group   A. Barbir
Request for Comments: 4593   Nortel
Category: Informational   S. Murphy
                                                            Sparta, Inc.
                                                                Y. Yang
                                                          Cisco Systems
                                                           October 2006

Generic Threats to Routing Protocols

Status of This Memo

   This memo provides information for the Internet community. It does
   not specify an Internet standard of any kind. Distribution of this
   memo is unlimited.

Abstract

   Routing protocols are subject to attacks that can harm individual
   users or network operations as a whole. This document provides a
   description and a summary of generic threats that affect routing
   protocols in general. This work describes threats, including threat
   sources and capabilities, threat actions, and threat consequences, as
   well as a breakdown of routing functions that might be attacked
   separately.

Table of Contents

1. Introduction

   Routing protocols are subject to threats and attacks that can harm
   individual users or the network operations as a whole. The document
   provides a summary of generic threats that affect routing protocols.
   In particular, this work identifies generic threats to routing
   protocols that include threat sources, threat actions, and threat
   consequences. A breakdown of routing functions that might be
   separately attacked is provided.

   This work should be considered a precursor to developing a common set
   of security requirements for routing protocols. While it is well
   known that bad, incomplete, or poor implementations of routing
   protocols may, in themselves, lead to routing problems or failures or
   may increase the risk of a network's being attacked successfully,
   these issues are not considered here. This document only considers

attacks against robust, well-considered implementations of routing
protocols, such as those specified in Open Shortest Path First (OSPF)
[4], Intermediate System to Intermediate System (IS-IS) [5][8], RIP
[6] and BGP [7]. Attacks against implementation-specific weaknesses
and vulnerabilities are out of scope for this document.

The document is organized as follows: Section 2 provides a review of
routing functions. Section 3 defines threats. In Section 4, a
discussion on generally identifiable routing threat actions is
provided. Section 5 addresses security considerations.

2. Routing Functions Overview

This section provides an overview of common functions that are shared
among various routing protocols. In general, routing protocols share
the following functions:

o Transport Subsystem: The routing protocol transmits messages to
   its neighbors using some underlying protocol. For example, OSPF
   uses IP, while other protocols may run over TCP.

o Neighbor State Maintenance: Neighboring relationship formation is
   the first step for topology determination. For this reason,
   routing protocols may need to maintain state information. Each
   routing protocol may use a different mechanism for determining its
   neighbors in the routing topology. Some protocols have distinct
   exchanges through which they establish neighboring relationships,
   e.g., Hello exchanges in OSPF.

o Database Maintenance: Routing protocols exchange network topology
   and reachability information. The routers collect this
   information in routing databases with varying detail. The
   maintenance of these databases is a significant portion of the
   function of a routing protocol.

In a routing protocol, there are message exchanges that are intended
for the control of the state of the protocol. For example, neighbor
maintenance messages carry such information. On the other hand,
there are messages that are used to exchange information that is
intended to be used in the forwarding function, for example, messages
that are used to maintain the database. These messages affect the
data (information) part of the routing protocol.

3. Generic Routing Protocol Threat Model

The model developed in this section can be used to identify threats to any routing protocol.

Routing protocols are subject to threats at various levels. For example, threats can affect the transport subsystem, where the routing protocol can be subject to attacks on its underlying protocol. An attacker may also attack messages that carry control information in a routing protocol to break a neighboring (e.g., peering, adjacency) relationship. This type of attack can impact the network routing behavior in the affected routers and likely the surrounding neighborhood as well. For example, in BGP, if a router receives a CEASE message, it will break its neighboring relationship to its peer and potentially send new routing information to any remaining peers.

An attacker may also attack messages that carry data information in order to break a database exchange between two routers or to affect the database maintenance functionality. For example, the information in the database must be authentic and authorized. An attacker who is able to introduce bogus data can have a strong effect on the behavior of routing in the neighborhood. For example, if an OSPF router sends LSAs with the wrong Advertising Router, the receivers will compute a Shortest Path First (SPF) tree that is incorrect and might not forward the traffic. If a BGP router advertises a Network Layer Reachability Information (NLRI) that it is not authorized to advertise, then receivers might forward that NLRI's traffic toward that router and the traffic would not be deliverable. A Protocol Independent Multicast (PIM) router might transmit a JOIN message to receive multicast data it would otherwise not receive.

3.1. Threat Definitions

In [1], a threat is defined as a potential for violation of security, which exists when there is a circumstance, capability, action, or event that could breach security and cause harm. Threats can be categorized as threat sources, threat actions, threat consequences, threat consequence zones, and threat consequence periods.

3.1.1. Threat Sources

In the context of deliberate attack, a threat source is defined as a motivated, capable adversary. By modeling the motivations (attack goals) and capabilities of the adversaries who are threat sources, one can better understand what classes of attacks these threats may mount and thus what types of countermeasures will be required to deal with these attacks.

3.1.1.1.   Adversary Motivations

We assume that the most common goal of an adversary deliberately
attacking routing is to cause inter-domain routing to malfunction. A
routing malfunction affects data transmission such that traffic
follows a path (sequence of autonomous systems in the case of BGP)
other than one that would have been computed by the routing protocol
if it were operating properly (i.e., if it were not under attack).
As a result of an attack, a route may terminate at a router other
than the one that legitimately represents the destination address of
the traffic, or it may traverse routers other than those that it
would otherwise have traversed. In either case, a routing
malfunction may allow an adversary to wiretap traffic passively, or
to engage in man-in-the-middle (MITM) active attacks, including
discarding traffic (denial of service).

A routing malfunction might be effected for financial gain related to
traffic volume (vs. the content of the routed traffic), e.g., to
affect settlements among ISPs.

Another possible goal for attacks against routing can be damage to
the network infrastructure itself, on a targeted or wide-scale basis.
Thus, for example, attacks that cause excessive transmission of
UPDATE or other management messages, and attendant router processing,
could be motivated by these goals.

Irrespective of the goals noted above, an adversary may or may not be
averse to detection and identification. This characteristic of an
adversary influences some of the ways in which attacks may be
accomplished.

3.1.1.2.   Adversary Capabilities

Different adversaries possess varied capabilities.

o All adversaries are presumed to be capable of directing packets to
   routers from remote locations and can assert a false IP source
   address with each packet (IP address spoofing) in an effort to
   cause the targeted router to accept and process the packet as
   though it emanated from the indicated source. Spoofing attacks
   may be employed to trick routers into acting on bogus messages to
   effect misrouting, or these messages may be used to overwhelm the
   management processor in a router, to effect DoS. Protection from
   such adversaries must not rely on the claimed identity in routing
   packets that the protocol receives.

o Some adversaries can monitor links over which routing traffic is
   carried and emit packets that mimic data contained in legitimate
   routing traffic carried over these links; thus, they can actively
   participate in message exchanges with the legitimate routers.
   This increases the opportunities for an adversary to generate
   bogus routing traffic that may be accepted by a router, to effect
   misrouting or DoS. Retransmission of previously delivered
   management traffic (replay attacks) exemplify this capability. As
   a result, protection from such adversaries ought not to rely on
   the secrecy of unencrypted data in packet headers or payloads.

o Some adversaries can effect MITM attacks against routing traffic,
   e.g., as a result of active wiretapping on a link between two
   routers. This represents the ultimate wiretapping capability for
   an adversary. Protection from such adversaries must not rely on
   the integrity of inter-router links to authenticate traffic,
   unless cryptographic measures are employed to detect unauthorized
   modification.

o Some adversaries can subvert routers, or the management
   workstations used to control these routers. These Byzantine
   failures represent the most serious form of attack capability in
   that they result in emission of bogus traffic by legitimate
   routers. As a result, protection from such adversaries must not
   rely on the correct operation of neighbor routers. Protection
   measures should adopt the principle of least privilege, to
   minimize the impact of attacks of this sort. To counter Byzantine
   attacks, routers ought not to trust management traffic (e.g.,
   based on its source) but rather each router should independently
   authenticate management traffic before acting upon it.

We will assume that any cryptographic countermeasures employed to
secure BGP will employ algorithms and modes that are resistant to
attack, even by sophisticated adversaries; thus, we will ignore
cryptanalytic attacks.

Deliberate attacks are mimicked by failures that are random and
unintentional. In particular, a Byzantine failure in a router may
occur because the router is faulty in hardware or software or is
misconfigured. As described in [3], "A node with a Byzantine failure
may corrupt messages, forge messages, delay messages, or send
conflicting messages to different nodes". Byzantine routers, whether
faulty, misconfigured, or subverted, have the context to provide

believable and very damaging bogus routing information. Byzantine
routers may also claim another legitimate peer's identity. Given
their status as peers, they may even elude the authentication
protections, if those protections can only detect that a source is
one of the legitimate peers (e.g., the router uses the same
cryptographic key to authenticate all peers).

We therefore characterize threat sources into two groups:

Outsiders: These attackers may reside anywhere in the Internet, have
   the ability to send IP traffic to the router, may be able to
   observe the router's replies, and may even control the path for a
   legitimate peer's traffic. These are not legitimate participants
   in the routing protocol.

Byzantine: These attackers are faulty, misconfigured, or subverted
   routers; i.e., legitimate participants in the routing protocol.

3.1.2. Threat Consequences

A threat consequence is a security violation that results from a
threat action [1]. To a routing protocol, a security violation is a
compromise of some aspect of the correct behavior of the routing
system. The compromise can damage the data traffic intended for a
particular network or host or can damage the operation of the routing
infrastructure of the network as a whole.

There are four types of general threat consequences: disclosure,
deception, disruption, and usurpation [1].

o Disclosure: Disclosure of routing information happens when an
   attacker successfully accesses the information without being
   authorized. Outsiders who can observe or monitor a link may cause
   disclosure, if routing exchanges lack confidentiality. Byzantine
   routers can cause disclosure, as long as they are successfully
   involved in the routing exchanges. Although inappropriate
   disclosure of routing information can pose a security threat or be
   part of a later, larger, or higher layer attack, confidentiality
   is not generally a design goal of routing protocols.

o Deception: This consequence happens when a legitimate router
   receives a forged routing message and believes it to be authentic.
   Both outsiders and Byzantine routers can cause this consequence if
   the receiving router lacks the ability to check routing message
   integrity or origin authentication.

o Disruption: This consequence occurs when a legitimate router's
   operation is being interrupted or prevented. Outsiders can cause
   this by inserting, corrupting, replaying, delaying, or dropping
   routing messages, or by breaking routing sessions between
   legitimate routers. Byzantine routers can cause this consequence
   by sending false routing messages, interfering with normal routing
   exchanges, or flooding unnecessary routing protocol messages.
   (DoS is a common threat action causing disruption.)

o Usurpation: This consequence happens when an attacker gains
   control over the services/functions a legitimate router is
   providing to others. Outsiders can cause this by delaying or
   dropping routing exchanges, or fabricating or replaying routing
   information. Byzantine routers can cause this consequence by
   sending false routing information or interfering with routing
   exchanges.

Note: An attacker does not have to control a router directly to
control its services. For example, in Figure 1, Network 1 is dual-
homed through Router A and Router B, and Router A is preferred.
However, Router B is compromised and advertises a better metric.
Consequently, devices on the Internet choose the path through Router
B to reach Network 1. In this way, Router B steals the data traffic,
and Router A loses its control of the services to Router B. This is
depicted in Figure 1.

```
            +--------------+   +-------+
            |   Internet   |---| Rtr A |
            +------+-------+   +---+---+
                   |               |
                   |               |
                   |             *-+-*
            +-------+            /     \
            | Rtr B |----------*  N 1   *
            +-------+            \     /
                                 *---*
```
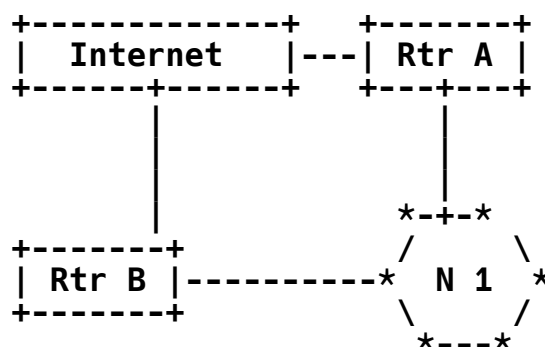
              Figure 1. Dual-homed network

Several threat consequences might be caused by a single threat
action. In Figure 1, there exist at least two consequences: routers
using Router B to reach Network 1 are deceived, and Router A is
usurped.

3.1.2.1. Threat Consequence Scope

   As mentioned above, an attack might damage the data traffic intended
   for a particular network or host or damage the operation of the
   routing infrastructure of the network as a whole. Damage that might
   result from attacks against the network as a whole may include the
   following:

   o Network congestion. More data traffic is forwarded through some
      portion of the network than would otherwise need to carry the
      traffic.

   o Blackhole. Large amounts of traffic are unnecessarily re-directed
      to be forwarded through one router and that router drops
      many/most/all packets.

   o Looping. Data traffic is forwarded along a route that loops, so
      that the data is never delivered (resulting in network
      congestion).

   o Partition. Some portion of the network believes that it is
      partitioned from the rest of the network when it is not.

   o Churn. The forwarding in the network changes (unnecessarily) at a
      rapid pace, resulting in large variations in the data delivery
      patterns (and adversely affecting congestion control techniques).

   o Instability. The protocol becomes unstable so that convergence on
      a global forwarding state is not achieved.

   o Overcontrol. The routing protocol messages themselves become a
      significant portion of the traffic the network carries.

   o Clog. A router receives an excessive number of routing protocol
      messages, causing it to exhaust some resource (e.g., memory, CPU,
      battery).

   The damage that might result from attacks against a particular host
   or network address may include the following:

   o Starvation. Data traffic destined for the network or host is
      forwarded to a part of the network that cannot deliver it.

   o Eavesdrop. Data traffic is forwarded through some router or
      network that would otherwise not see the traffic, affording an
      opportunity to see the data or at least the data delivery pattern.

o Cut. Some portion of the network believes that it has no route to
   the host or network when it is in fact connected.

o Delay. Data traffic destined for the network or host is forwarded
   along a route that is in some way inferior to the route it would
   otherwise take.

o Looping. Data traffic for the network or host is forwarded along
   a route that loops, so that the data is never delivered.

It is important to consider all consequences, because some security
solutions can protect against one consequence but not against others.
It might be possible to design a security solution that protects
against eavesdropping on one destination's traffic without protecting
against churn in the network. Similarly, it is possible to design a
security solution that prevents a starvation attack against one host,
but not a clogging attack against a router. The security
requirements must be clear as to which consequences are being avoided
and which consequences must be addressed by other means (e.g., by
administrative means outside the protocol).

3.1.2.2. Threat Consequence Zone

A threat consequence zone covers the area within which the network
operations have been affected by threat actions. Possible threat
consequence zones can be classified as a single link or router,
multiple routers (within a single routing domain), a single routing
domain, multiple routing domains, or the global Internet. The threat
consequence zone varies based on the threat action and the position
of the target of the attack. Similar threat actions that happen at
different locations may result in totally different threat
consequence zones. For example, when an outsider breaks the routing
session between a distribution router and a stub router, only
reachability to and from the network devices attached to the stub
router will be impaired. In other words, the threat consequence zone
is a single router. In another case, if the outsider is located
between a customer edge router and its corresponding provider edge
router, such an action might cause the whole customer site to lose
its connection. In this case, the threat consequence zone might be a
single routing domain.

3.1.2.3. Threat Consequence Periods

A threat consequence period is defined as the portion of time during
which the network operations are impacted by the threat consequences.
The threat consequence period is influenced by, but not totally
dependent on, the duration of the threat action. In some cases, the
network operations will get back to normal as soon as the threat

action has been stopped. In other cases, however, threat
consequences may persist longer than does the threat action. For
example, in the original Advanced Research Projects Agency Network
(ARPANET) link-state algorithm, some errors in a router introduced
three instances of a Link-State Announcement (LSA). All of them
flooded throughout the network continuously, until the entire network
was power cycled [2].

4. Generally Identifiable Routing Threat Actions

This section addresses generally identifiable and recognized threat
actions against routing protocols. The threat actions are not
necessarily specific to individual protocols but may be present in
one or more of the common routing protocols in use today.

4.1. Deliberate Exposure

Deliberate exposure occurs when an attacker takes control of a router
and intentionally releases routing information to other entities
(e.g., the attacker, a web page, mail posting, other routers) that
otherwise should not receive the exposed information.

The consequence of deliberate exposure is the disclosure of routing
information.

The threat consequence zone of deliberate exposure depends on the
routing information that the attackers have exposed. The more
knowledge they have exposed, the bigger the threat consequence zone.

The threat consequence period of deliberate exposure might be longer
than the duration of the action itself. The routing information
exposed will not be outdated until there is a topology change of the
exposed network.

4.2. Sniffing

Sniffing is an action whereby attackers monitor and/or record the
routing exchanges between authorized routers to sniff for routing
information. Attackers can also sniff data traffic information
(however, this is out of scope of the current work).

The consequence of sniffing is disclosure of routing information.

The threat consequence zone of sniffing depends on the attacker's
location, the routing protocol type, and the routing information that
has been recorded. For example, if the outsider is sniffing a link
that is in an OSPF totally stubby area, the threat consequence zone
should be limited to the whole area. An attacker that is sniffing a

link in an External Border Gateway Protocol (EBGP) session can gain
knowledge of multiple routing domains.

The threat consequence period might be longer than the duration of
the action. If an attacker stops sniffing a link, their acquired
knowledge will not be out-dated until there is a topology change of
the affected network.

4.3. Traffic Analysis

Traffic analysis is an action whereby attackers gain routing
information by analyzing the characteristics of the data traffic on a
subverted link. Traffic analysis threats can affect any data that is
sent over a communication link. This threat is not peculiar to
routing protocols and is included here for completeness.

The consequence of data traffic analysis is the disclosure of routing
information. For example, the source and destination IP addresses of
the data traffic and the type, magnitude, and volume of traffic can
be disclosed.

The threat consequence zone of the traffic analysis depends on the
attacker's location and what data traffic has passed through. An
attacker at the network core should be able to gather more
information than its counterpart at the edge and would therefore have
to be able to analyze traffic patterns in a wider area.

The threat consequence period might be longer than the duration of
the traffic analysis. After the attacker stops traffic analysis, its
knowledge will not be outdated until there is a topology change of
the disclosed network.

4.4. Spoofing

Spoofing occurs when an illegitimate device assumes the identity of a
legitimate one. Spoofing in and of itself is often not the true
attack. Spoofing is special in that it can be used to carry out
other threat actions causing other threat consequences. An attacker
can use spoofing as a means for launching other types of attacks.
For example, if an attacker succeeds in spoofing the identity of a
router, the attacker can send out unrealistic routing information
that might cause the disruption of network services.

There are a few cases where spoofing can be an attack in and of
itself. For example, messages from an attacker that spoof the
identity of a legitimate router may cause a neighbor relationship to
form and deny the formation of the relationship with the legitimate
router.

The consequences of spoofing are as follows:

o The disclosure of routing information. The spoofing router will
   be able to gain access to the routing information.

o The deception of peer relationship. The authorized routers, which
   exchange routing messages with the spoofing router, do not realize
   that they are neighboring with a router that is faking another
   router's identity.

The threat consequence zone is as follows:

o The consequence zone of the fake peer relationship will be limited
   to those routers trusting the attacker's claimed identity.

o The consequence zone of the disclosed routing information depends
   on the attacker's location, the routing protocol type, and the
   routing information that has been exchanged between the attacker
   and its deceived neighbors.

Note: This section focuses on addressing spoofing as a threat on its
own. However, spoofing creates conditions for other threats actions.
The other threat actions are considered falsifications and are
treated in the next section.

4.5. Falsification

Falsification is an action whereby an attacker sends false routing
information. To falsify the routing information, an attacker has to
be either the originator or a forwarder of the routing information.
It cannot be a receiver-only. False routing information describes
the network in an unrealistic fashion, whether or not intended by the
authoritative network administrator.

4.5.1. Falsifications by Originators

An originator of routing information can launch the falsifications
that are described in the next sections.

4.5.1.1. Overclaiming

Overclaiming occurs when a Byzantine router or outsider advertises
its control of some network resources, while in reality it does not,
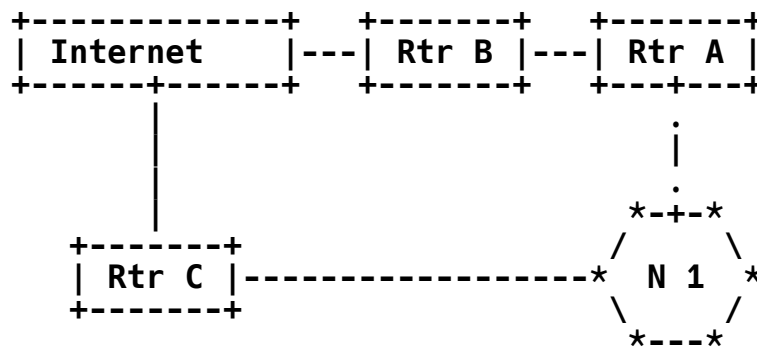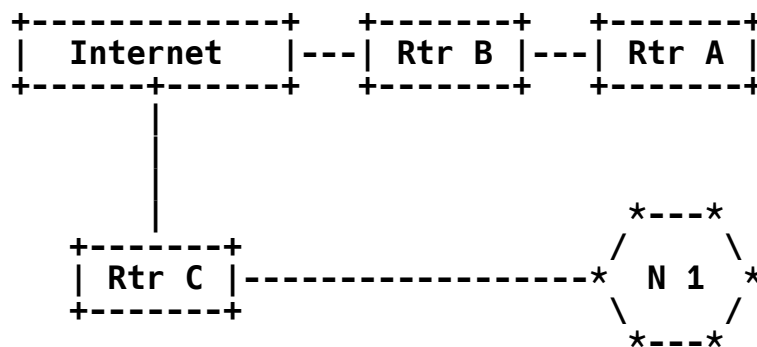or if the advertisement is not authorized. This is given in Figures
2 and 3.

```
           +--------------+   +-------+   +-------+
           |  Internet    |---| Rtr B |---| Rtr A |
           +------+-------+   +-------+   +---+---+
                  |                           :
                  |                           |
                  |                           .
                  |                         *-+-*
           +-------+                       /     \
           | Rtr C |---------------------*   N 1   *
           +-------+                       \     /
                                            *---*
```

                  Figure 2. Overclaiming-1

```
           +--------------+   +-------+   +-------+
           |  Internet    |---| Rtr B |---| Rtr A |
           +------+-------+   +-------+   +-------+
                  |
                  |
                  |
                  |                         *---*
           +-------+                       /     \
           | Rtr C |---------------------*   N 1   *
           +-------+                       \     /
                                            *---*
```

                  Figure 3. Overclaiming-2

   The above figures provide examples of overclaiming. Router A, the
   attacker, is connected to the Internet through Router B. Router C is
   authorized to advertise its link to Network 1. In Figure 2, Router A
   controls a link to Network 1 but is not authorized to advertise it.
   In Figure 3, Router A does not control such a link. But in either
   case, Router A advertises the link to the Internet, through Router B.

   Both Byzantine routers and outsiders can overclaim network resources.
   The consequences of overclaiming include the following:

   o Usurpation of the overclaimed network resources. In Figures 2 and
      3, usurpation of Network 1 can occur when Router B (or other
      routers on the Internet not shown in the figures) believes that
      Router A provides the best path to reach the Network 1. As a
      result, routers forward data traffic destined to Network 1 to
      Router A. The best result is that the data traffic uses an
      unauthorized path, as in Figure 2. The worst case is that the

      data never reaches the destination Network 1, as in Figure 3. The
      ultimate consequence is that Router A gains control over Network
      1's services, by controlling the data traffic.

   o Usurpation of the legitimate advertising routers. In Figures 2
      and 3, Router C is the legitimate advertiser of Network 1. By
      overclaiming, Router A also controls (partially or totally) the
      services/functions provided by the Router C. (This is NOT a
      disruption, as Router C is operating in a way intended by the
      authoritative network administrator.)

   o Deception of other routers. In Figures 2 and 3, Router B, or
      other routers on the Internet, might be deceived into believing
      that the path through Router A is the best.

   o Disruption of data planes on some routers. This might happen to
      routers that are on the path that is used by other routers to
      reach the overclaimed network resources through the attacker. In
      Figures 2 and 3, when other routers on the Internet are deceived,
      they will forward the data traffic to Router B, which might be
      overloaded.

   The threat consequence zone varies based on the consequence:

   o Where usurpation is concerned, the consequence zone covers the
      network resources that are overclaimed by the attacker (Network 1
      in Figures 2 and 3), and the routers that are authorized to
      advertise the network resources but lose the competition against
      the attacker (Router C in Figures 2 and 3).

   o Where deception is concerned, the consequence zone covers the
      routers that do believe the attacker's advertisement and use the
      attacker to reach the claimed networks (Router B and other
      deceived routers on the Internet in Figures 2 and 3).

   o Where disruption is concerned, the consequence zone includes the
      routers that are on the path of misdirected data traffic (Router B
      in Figures 2 and 3 and other routers in the Internet on the path
      of the misdirected traffic).

   The threat consequence will not cease when the attacker stops
   overclaiming and will totally disappear only when the routing tables
   are converged. As a result, the consequence period is longer than
   the duration of the overclaiming.

4.5.1.2. Misclaiming

   A misclaiming threat is defined as an action whereby an attacker is
   advertising some network resources that it is authorized to control,
   but in a way that is not intended by the authoritative network
   administrator. For example, it may be advertising inappropriate link
   costs in an OSPF LSA. An attacker can eulogize or disparage when
   advertising these network resources. Byzantine routers can misclaim
   network resources.

   The threat consequences of misclaiming are similar to the
   consequences of overclaiming.

   The consequence zone and period are also similar to those of
   overclaiming.

4.5.2. Falsifications by Forwarders

   In each routing protocol, routers that forward routing protocol
   messages are expected to leave some fields unmodified and to modify
   other fields in certain circumscribed ways. The fields to be
   modified, the possible new contents of those fields and their
   computation from the original fields, the fields that must remain
   unmodified, etc. are all detailed in the protocol specification.
   They may vary depending on the function of the router or its network
   environment. For example, in RIP, the forwarder must modify the
   routing information by increasing the hop count by 1. On the other
   hand, a forwarder must not modify any field of the type 1 LSA in OSPF
   except the age field. In general, forwarders in distance vector
   routing protocols are authorized to and must modify the routing
   information, while most forwarders in link state routing protocols
   are not authorized to and must not modify most routing information.

   As a forwarder authorized to modify routing messages, an attacker
   might also falsify by not forwarding routing information to other
   authorized routers as required.

4.5.2.1. Misstatement

   This is defined as an action whereby the attacker modifies route
   attributes in an incorrect manner. For example, in RIP, the attacker
   might increase the path cost by two hops instead of one. In BGP, the
   attacker might delete some AS numbers from the AS PATH.

Where forwarding routing information should not be modified, an attacker can launch the following falsifications:

o Deletion. Attacker deletes valid data in the routing message.

o Insertion. Attacker inserts false data in the routing message.

o Substitution. Attacker replaces valid data in the routing message with false data.

A forwarder can also falsify data by replaying out-dated data in the routing message as current data.

All types of attackers, outsiders and Byzantine routers, can falsify the routing information when they forward the routing messages.

The threat consequences of these falsifications by forwarders are similar to those caused by originators: usurpation of some network resources and related routers; deception of routers using false paths; and disruption of data planes of routers on the false paths. The threat consequence zone and period are also similar.

4.6. Interference

Interference is a threat action whereby an attacker inhibits the exchanges by legitimate routers. The attacker can do this by adding noise, by not forwarding packets, by replaying out-dated packets, by inserting or corrupting messages, by delaying responses, by denial of receipts, or by breaking synchronization.

Byzantine routers can slow down their routing exchanges or induce flapping in the routing sessions of legitimate neighboring routers.

The consequence of interference is the disruption of routing operations.

The consequence zone of interference depends on the severity of the interference. If the interference results in consequences at the neighbor maintenance level, then there may be changes in the database, resulting in network-wide consequences.

The threat consequences might disappear as soon as the interference is stopped or might not totally disappear until the networks have converged. Therefore, the consequence period is equal to or longer than the duration of the interference.

4.7. Overload

   Overload is defined as a threat action whereby attackers place excess
   burden on legitimate routers. For example, it is possible for an
   attacker to trigger a router to create an excessive amount of state
   that other routers within the network are not able to handle. In a
   similar fashion, it is possible for an attacker to overload database
   routing exchanges and thus to influence the routing operations.

5. Security Considerations

   This entire document is security related. Specifically, the document
   addresses security of routing protocols as associated with threats to
   those protocols. In a larger context, this work builds upon the
   recognition of the IETF community that signaling and
   control/management planes of networked devices need strengthening.
   Routing protocols can be considered part of that signaling and
   control plane. However, to date, routing protocols have largely
   remained unprotected and open to malicious attacks. This document
   discusses inter- and intra-domain routing protocol threats that are
   currently known and lays the foundation for other documents that will
   discuss security requirements for routing protocols. This document
   is protocol independent.

6. References

6.1. Normative References

   [1] Shirey, R., "Internet Security Glossary", RFC 2828, May 2000.

   [2] Rosen, E., "Vulnerabilities of network control protocols: An
        example", RFC 789, July 1981.

   [3] Perlman, R., "Network Layer Protocols with Byzantine
        Robustness", PhD thesis, MIT LCS TR-429, October 1988.

   [4] Moy, J., "OSPF Version 2", STD 54, RFC 2328, April 1998.

   [5] Callon, R., "Use of OSI IS-IS for routing in TCP/IP and dual
        environments", RFC 1195, December 1990.

   [6] Malkin, G., "RIP Version 2", STD 56, RFC 2453, November 1998.

   [7] Rekhter, Y., Li, T., and S. Hares, "A Border Gateway Protocol 4
        (BGP-4)", RFC 4271, January 2006.

   [8]   ISO 10589, "Intermediate System to Intermediate System intra-
         domain routeing information exchange protocol for use in
         conjunction with the protocol for providing the connectionless-
         mode network service (ISO 8473)", ISO/IEC 10589:2002.

Appendix A.   Acknowledgments

   This document would not have been possible save for the excellent
   efforts and teamwork characteristics of those listed here.

   o Dennis Beard, Nortel
   o Ayman Musharbash, Nortel
   o Jean-Jacques Puig, int-evry, France
   o Paul Knight, Nortel
   o Elwyn Davies, Nortel
   o Ameya Dilip Pandit, Graduate student, University of Missouri
   o Senthilkumar Ayyasamy, Graduate student, University of Missouri
   o Stephen Kent, BBN
   o Tim Gage, Cisco Systems
   o James Ng, Cisco Systems
   o Alvaro Retana, Cisco Systems

Appendix B.   Acronyms

   AS - Autonomous system. Set of routers under a single technical
   administration. Each AS normally uses a single interior gateway
   protocol (IGP) and metrics to propagate routing information within
   the set of routers. Also called routing domain.

   AS-Path - In BGP, the route to a destination. The path consists of
   the AS numbers of all routers a packet must go through to reach a
   destination.

   BGP - Border Gateway Protocol. Exterior gateway protocol used to
   exchange routing information among routers in different autonomous
   systems.

   LSA - Link-State Announcement

   NLRI - Network Layer Reachability Information. Information that is
   carried in BGP packets and is used by MBGP.

   OSPF - Open Shortest Path First. A link-state IGP that makes routing
   decisions based on the shortest-path-first (SPF) algorithm (also
   referred to as the Dijkstra algorithm).

Authors' Addresses

   Abbie Barbir
   Nortel
   3500 Carling Avenue
   Nepean, Ontario K2H 8E9
   Canada

   EMail: abbieb@nortel.com

   Sandy Murphy
   Sparta, Inc.
   7110 Samuel Morse Drive
   Columbia, MD
   USA

   Phone: 443-430-8000
   EMail: sandy@sparta.com

   Yi Yang
   Cisco Systems
   7025 Kit Creek Road
   RTP, NC 27709
   USA

   EMail: yiya@cisco.com