         A Framework for Inter-Domain Multiprotocol Label Switching
                          Traffic Engineering

Status of This Memo

Copyright Notice

Abstract

   This document provides a framework for establishing and controlling
   Multiprotocol Label Switching (MPLS) and Generalized MPLS (GMPLS)
   Traffic Engineered (TE) Label Switched Paths (LSPs) in multi-domain
   networks.

   For the purposes of this document, a domain is considered to be any
   collection of network elements within a common sphere of address
   management or path computational responsibility. Examples of such
   domains include Interior Gateway Protocol (IGP) areas and Autonomous
   Systems (ASes).

Table of Contents

1. Introduction

   The Traffic Engineering Working Group has developed requirements for
   inter-area and inter-AS Multiprotocol Label Switching (MPLS) Traffic
   Engineering in [RFC4105] and [RFC4216].

   Various proposals have subsequently been made to address some or all
   of these requirements through extensions to the Resource Reservation
   Protocol Traffic Engineering extensions (RSVP-TE) and to the Interior
   Gateway Protocols (IGPs) (i.e., Intermediate System to Intermediate
   System (IS-IS) and OSPF).

   This document introduces the techniques for establishing Traffic
   Engineered (TE) Label Switched Paths (LSPs) across multiple domains.
   In this context and within the remainder of this document, we
   consider all source-based and constraint-based routed LSPs and refer
   to them interchangeably as "TE LSPs" or "LSPs".

   The functional components of these techniques are separated into the
   mechanisms for discovering reachability and TE information, for
   computing the paths of LSPs, and for signaling the LSPs. Note that
   the aim of this document is not to detail each of those techniques,
   which are covered in separate documents referenced from the sections
   of this document that introduce the techniques, but rather to propose
   a framework for inter-domain MPLS Traffic Engineering.

   Note that in the remainder of this document, the term "MPLS Traffic
   Engineering" is used equally to apply to MPLS and Generalized MPLS
   (GMPLS) traffic. Specific issues pertaining to the use of GMPLS in
   inter-domain environments (for example, policy implications of the
   use of the Link Management Protocol [RFC4204] on inter-domain links)
    are covered in separate documents such as [GMPLS-AS].

   For the purposes of this document, a domain is considered to be any
   collection of network elements within a common sphere of address
   management or path computational responsibility. Examples of such
   domains include IGP areas and Autonomous Systems. Wholly or
   partially overlapping domains (e.g., path computation sub-domains of
   areas or ASes) are not within the scope of this document.

1.1. Nested Domains

   Nested domains are outside the scope of this document. It may be
   that some domains that are nested administratively or for the
   purposes of address space management can be considered as adjacent
   domains for the purposes of this document; however, the fact that the
   domains are nested is then immaterial. In the context of MPLS TE,
   domain A is considered to be nested within domain B if domain A is

wholly contained in domain B, and domain B is fully or partially
aware of the TE characteristics and topology of domain A.

2. Signaling Options

Three distinct options for signaling TE LSPs across multiple domains
are identified. The choice of which options to use may be influenced
by the path computation technique used (see section 3), although some
path computation techniques may apply to multiple signaling options.
The choice may further depend on the application to which the TE LSPs
are put and the nature, topology, and switching capabilities of the
network.

A comparison of the usages of the different signaling options is
beyond the scope of this document and should be the subject of a
separate applicability statement.

2.1. LSP Nesting

Hierarchical LSPs form a fundamental part of MPLS [RFC3031] and are
discussed in further detail in [RFC4206]. Hierarchical LSPs may
optionally be advertised as TE links. Note that a hierarchical LSP
that spans multiple domains cannot be advertised in this way because
there is no concept of TE information that spans domains.

Hierarchical LSPs can be used in support of inter-domain TE LSPs. In
particular, a hierarchical LSP may be used to achieve connectivity
between any pair of Label Switching Routers (LSRs) within a domain.
The ingress and egress of the hierarchical LSP could be the edge
nodes of the domain in which case connectivity is achieved across the
entire domain, or they could be any other pair of LSRs in the domain.

The technique of carrying one TE LSP within another is termed LSP
nesting. A hierarchical LSP may provide a TE LSP tunnel to transport
(i.e., nest) multiple TE LSPs along a common part of their paths.
Alternatively, a TE LSP may carry (i.e., nest) a single LSP in a
one-to-one mapping.

The signaling trigger for the establishment of a hierarchical LSP may
be the receipt of a signaling request for the TE LSP that it will
carry, or may be a management action to "pre-engineer" a domain to be
crossed by TE LSPs that would be used as hierarchical LSPs by the
traffic that has to traverse the domain. Furthermore, the mapping
(inheritance rules) between attributes of the nested and the
hierarchical LSPs (including bandwidth) may be statically pre-
configured or, for on-demand hierarchical LSPs, may be dynamic

according to the properties of the nested LSPs. Even in the dynamic
case, inheritance from the properties of the nested LSP(s) can be
complemented by local or domain-wide policy rules.

Note that a hierarchical LSP may be constructed to span multiple
domains or parts of domains. However, such an LSP cannot be
advertised as a TE link that spans domains. The end points of a
hierarchical LSP are not necessarily on domain boundaries, so nesting
is not limited to domain boundaries.

Note also that the Interior/Exterior Gateway Protocol (IGP/EGP)
routing topology is maintained unaffected by the LSP connectivity and
TE links introduced by hierarchical LSPs even if they are advertised
as TE links. That is, the routing protocols do not exchange messages
over the hierarchical LSPs, and LSPs are not used to create routing
adjacencies between routers.

During the operation of establishing a nested LSP that uses a
hierarchical LSP, the SENDER_TEMPLATE and SESSION objects remain
unchanged along the entire length of the nested LSP, as do all other
objects that have end-to-end significance.

2.2. Contiguous LSP

A single contiguous LSP is established from ingress to egress in a
single signaling exchange. No further LSPs are required to be
established to support this LSP so that hierarchical or stitched LSPs
are not needed.

A contiguous LSP uses the same Session/LSP ID along the whole of its
path (that is, at each LSR). The notions of "splicing" together
different LSPs or of "shuffling" Session or LSP identifiers are not
considered.

2.3. LSP Stitching

LSP Stitching is described in [STITCH]. In the LSP stitching model,
separate LSPs (referred to as a TE LSP segments) are established and
are "stitched" together in the data plane so that a single end-to-end
Label Switched Path is achieved. The distinction is that the
component LSP segments are signaled as distinct TE LSPs in the
control plane. Each signaled TE LSP segment has a different source
and destination.

LSP stitching can be used in support of inter-domain TE LSPs. In
particular, an LSP segment may be used to achieve connectivity
between any pair of LSRs within a domain. The ingress and egress of
the LSP segment could be the edge nodes of the domain in which case

connectivity is achieved across the entire domain, or they could be any other pair of LSRs in the domain.

The signaling trigger for the establishment of a TE LSP segment may be the establishment of the previous TE LSP segment, the receipt of a setup request for TE LSP that it plans to stitch to a local TE LSP segment, or a management action.

LSP segments may be managed and advertised as TE links.

2.4. Hybrid Methods

There is nothing to prevent the mixture of signaling methods described above when establishing a single, end-to-end, inter-domain TE LSP. It may be desirable in this case for the choice of the various methods to be reported along the path, perhaps through the Record Route Object (RRO).

If there is a desire to restrict which methods are used, this must be signaled as described in the next section.

2.5. Control of Downstream Choice of Signaling Method

Notwithstanding the previous section, an ingress LSR may wish to restrict the signaling methods applied to a particular LSP at domain boundaries across the network. Such control, where it is required, may be achieved by the definition of appropriate new flags in the SESSION-ATTRIBUTE object or the Attributes Flags TLV of the LSP_ATTRIBUTES object [RFC4420]. Before defining a mechanism to provide this level of control, the functional requirement to control the way in which the network delivers a service must be established. Also, due consideration must be given to the impact on interoperability since new mechanisms must be backwards compatible, and care must be taken to avoid allowing standards-conformant implementations that each supports a different functional subset in such a way that they are not capable of establishing LSPs.

3. Path Computation Techniques

The discussion of path computation techniques within this document is limited significantly to the determination of where computation may take place and what components of the full path may be determined.

The techniques used are closely tied to the signaling methodologies described in the previous section in that certain computation techniques may require the use of particular signaling approaches and vice versa.

Any discussion of the appropriateness of a particular path
computation technique in any given circumstance is beyond the scope
of this document and should be described in a separate applicability
statement.

Path computation algorithms are firmly out of the scope of this
document.

3.1. Management Configuration

Path computation may be performed by offline tools or by a network
planner. The resultant path may be supplied to the ingress LSR as
part of the TE LSP or service request, and encoded by the ingress LSR
as an Explicit Route Object (ERO) on the Path message that is sent
out.

There is no reason why the path provided by the operator should not
span multiple domains if the relevant information is available to the
planner or the offline tool. The definition of what information is
needed to perform this operation and how that information is
gathered, is outside the scope of this document.

3.2. Head-End Computation

The head-end, or ingress, LSR may assume responsibility for path
computation when the operator supplies part or none of the explicit
path. The operator must, in any case, supply at least the
destination address (egress) of the LSP.

3.2.1. Multi-Domain Visibility Computation

If the ingress has sufficient visibility of the topology and TE
information for all of the domains across which it will route the LSP
to its destination, then it may compute and provide the entire path.
The quality of this path (that is, its optimality as discussed in
section 3.5) can be better if the ingress has full visibility into
all relevant domains rather than just sufficient visibility to
provide some path to the destination.

Extreme caution must be exercised in consideration of the
distribution of the requisite TE information. See section 4.

3.2.2. Partial Visibility Computation

It may be that the ingress does not have full visibility of the
topology of all domains, but does have information about the
connectedness of the domains and the TE resource availability across
the domains. In this case, the ingress is not able to provide a

fully specified strict explicit path from ingress to egress.
However, for example, the ingress might supply an explicit path that
comprises:

   - explicit hops from ingress to the local domain boundary
   - loose hops representing the domain entry points across the
     network
   - a loose hop identifying the egress.

Alternatively, the explicit path might be expressed as:

   - explicit hops from ingress to the local domain boundary
   - strict hops giving abstract nodes representing each domain in
     turn
   - a loose hop identifying the egress.

These two explicit path formats could be mixed according to the
information available resulting in different combinations of loose
hops and abstract nodes.

This form of explicit path relies on some further computation
technique being applied at the domain boundaries. See section 3.3.

As with the multi-domain visibility option, extreme caution must be
exercised in consideration of the distribution of the requisite TE
information. See section 4.

3.2.3. Local Domain Visibility Computation

A final possibility for ingress-based computation is that the ingress
LSR has visibility only within its own domain, and connectivity
information only as far as determining one or more domain exit points
that may be suitable for carrying the LSP to its egress.

In this case, the ingress builds an explicit path that comprises
just:

   - explicit hops from ingress to the local domain boundary
   - a loose hop identifying the egress.

3.3. Domain Boundary Computation

If the partial explicit path methods described in sections 3.2.2 or
3.2.3 are applied, then the LSR at each domain boundary is
responsible for ensuring that there is sufficient path information
added to the Path message to carry it at least to the next domain
boundary (that is, out of the new domain).

If the LSR at the domain boundary has full visibility to the egress
then it can supply the entire explicit path. Note, however, that the
ERO processing rules of [RFC3209] state that it should only update
the ERO as far as the next specified hop (that is, the next domain
boundary if one was supplied in the original ERO) and, of course,
must not insert ERO subobjects immediately before a strict hop.

If the LSR at the domain boundary has only partial visibility (using
the definitions of section 3.2.2), it will fill in the path as far as
the next domain boundary, and will supply further domain/domain
boundary information if not already present in the ERO.

If the LSR at the domain boundary has only local visibility into the
immediate domain, it will simply add information to the ERO to carry
the Path message as far as the next domain boundary.

Domain boundary path computations are performed independently from
each other. Domain boundary LSRs may have different computation
capabilities, run different path computation algorithms, apply
different sets of constraints and optimization criteria, and so
forth, which might result in path segment quality that is
unpredictable to and out of the control of the ingress LSR. A
solution to this issue lies in enhancing the information signaled
during LSP setup to include a larger set of constraints and to
include the paths of related LSPs (such as diverse protected LSPs) as
described in [GMPLS-E2E].

It is also the case that paths generated on domain boundaries may
produce loops. Specifically, the paths computed may loop back into a
domain that has already been crossed by the LSP. This may or may not
be a problem, and might even be desirable, but could also give rise
to real loops. This can be avoided by using the recorded route (RRO)
to provide exclusions within the path computation algorithm, but in
the case of lack of trust between domains it may be necessary for the
RRO to indicate the previously visited domains. Even this solution
is not available where the RRO is not available on a Path message.
Note that when an RRO is used to provide exclusions, and a loop-free
path is found to be not available by the computation at a downstream
border node, crankback [CRANKBACK] may enable an upstream border node
to select an alternate path.

3.4. Path Computation Element

The computation techniques in sections 3.2 and 3.3 rely on topology
and TE information being distributed to the ingress LSR and those
LSRs at domain boundaries. These LSRs are responsible for computing
paths. Note that there may be scaling concerns with distributing the
required information; see section 4.

An alternative technique places the responsibility for path
computation with a Path Computation Element (PCE) [RFC4655]. There
may be either a centralized PCE, or multiple PCEs (each having local
visibility and collaborating in a distributed fashion to compute an
end-to-end path) across the entire network and even within any one
domain. The PCE may collect topology and TE information from the
same sources as would be used by LSRs in the previous paragraph, or
though other means.

Each LSR called upon to perform path computation (and even the
offline management tools described in section 3.1) may abdicate the
task to a PCE of its choice. The selection of PCE(s) may be driven
by static configuration or the dynamic discovery.

3.4.1. Multi-Domain Visibility Computation

A PCE may have full visibility, perhaps through connectivity to
multiple domains. In this case, it is able to supply a full explicit
path as in section 3.2.1.

3.4.2. Path Computation Use of PCE When Preserving Confidentiality

Note that although a centralized PCE or multiple collaborative PCEs
may have full visibility into one or more domains, it may be
desirable (e.g., to preserve topology confidentiality) that the full
path not be provided to the ingress LSR. Instead, a partial path is
supplied (as in section 3.2.2 or 3.2.3), and the LSRs at each domain
boundary are required to make further requests for each successive
segment of the path.

In this way, an end-to-end path may be computed using the full
network capabilities, but confidentiality between domains may be
preserved. Optionally, the PCE(s) may compute the entire path at the
first request and hold it in storage for subsequent requests, or it
may recompute each leg of the path on each request or at regular
intervals until requested by the LSRs establishing the LSP.

It may be the case that the centralized PCE or the collaboration
between PCEs may define a trust relationship greater than that
normally operational between domains.

3.4.3. Per-Domain Computation Elements

A third way that PCEs may be used is simply to have one (or more) per
domain. Each LSR within a domain that wishes to derive a path across
the domain may consult its local PCE.

This mechanism could be used for all path computations within the
domain, or specifically limited to computations for LSPs that will
leave the domain where external connectivity information can then be
restricted to just the PCE.

3.5. Optimal Path Computation

There are many definitions of an optimal path depending on the
constraints applied to the path computation. In a multi-domain
environment, the definitions are multiplied so that an optimal route
might be defined as the route that would be computed in the absence
of domain boundaries. Alternatively, another constraint might be
applied to the path computation to reduce or limit the number of
domains crossed by the LSP.

It is easy to construct examples that show that partitioning a
network into domains, and the resulting loss or aggregation of
routing information may lead to the computation of routes that are
other than optimal. It is impossible to guarantee optimal routing in
the presence of aggregation / abstraction / summarization of routing
information.

It is beyond the scope of this document to define what is an optimum
path for an inter-domain TE LSP. This debate is abdicated in favor
of requirements documents and applicability statements for specific
deployment scenarios. Note, however, that the meaning of certain
computation metrics may differ between domains (see section 5.6).

4. Distributing Reachability and TE Information

Traffic Engineering information is collected into a TE Database (TED)
on which path computation algorithms operate either directly or by
first constructing a network graph.

The path computation techniques described in the previous section
make certain demands upon the distribution of reachability
information and the TE capabilities of nodes and links within domains
as well as the TE connectivity across domains.

Currently, TE information is distributed within domains by additions
to IGPs [RFC3630], [RFC3784].

In cases where two domains are interconnected by one or more links
(that is, the domain boundary falls on a link rather than on a node),
there should be a mechanism to distribute the TE information
associated with the inter-domain links to the corresponding domains.
This would facilitate better path computation and reduce TE-related
crankbacks on these links.

Where a domain is a subset of an IGP area, filtering of TE
information may be applied at the domain boundary. This filtering
may be one way or two way.

Where information needs to reach a PCE that spans multiple domains,
the PCE may snoop on the IGP traffic in each domain, or play an
active part as an IGP-capable node in each domain. The PCE might
also receive TED updates from a proxy within the domain.

It is possible that an LSR that performs path computation (for
example, an ingress LSR) obtains the topology and TE information of
not just its own domain, but other domains as well. This information
may be subject to filtering applied by the advertising domain (for
example, the information may be limited to Forwarding Adjacencies
(FAs) across other domains, or the information may be aggregated or
abstracted).

Before starting work on any protocols or protocol extensions to
enable cross-domain reachability and TE advertisement in support of
inter-domain TE, the requirements and benefits must be clearly
established. This has not been done to date. Where any cross-domain
reachability and TE information needs to be advertised, consideration
must be given to TE extensions to existing protocols such as BGP, and
how the information advertised may be fed to the IGPs. It must be
noted that any extensions that cause a significant increase in the
amount of processing (such as aggregation computation) at domain
boundaries, or a significant increase in the amount of information
flooded (such as detailed TE information) need to be treated with
extreme caution and compared carefully with the scaling requirements
expressed in [RFC4105] and [RFC4216].

5. Comments on Advanced Functions

This section provides some non-definitive comments on the constraints
placed on advanced MPLS TE functions by inter-domain MPLS. It does
not attempt to state the implications of using one inter-domain
technique or another. Such material is deferred to appropriate
applicability statements where statements about the capabilities of
existing or future signaling, routing, and computation techniques to
deliver the functions listed should be made.

5.1. LSP Re-Optimization

Re-optimization is the process of moving a TE LSP from one path to
another, more preferable path (where no attempt is made in this
document to define "preferable" as no attempt was made to define
"optimal"). Make-before-break techniques are usually applied to
ensure that traffic is disrupted as little as possible. The Shared

Explicit style is usually used to avoid double booking of network
resources.

Re-optimization may be available within a single domain.
Alternatively, re-optimization may involve a change in route across
several domains or might involve a choice of different transit
domains.

Re-optimization requires that all or part of the path of the LSP be
re-computed. The techniques used may be selected as described in
section 3, and this will influence whether the whole or part of the
path is re-optimized.

The trigger for path computation and re-optimization may be an
operator request, a timer, information about a change in availability
of network resources, or a change in operational parameters (for
example, bandwidth) of an LSP. This trigger must be applied to the
point in the network that requests re-computation and controls re-
optimization and may require additional signaling.

Note also that where multiple mutually-diverse paths are applied
end-to-end (i.e., not simply within protection domains; see section
5.5) the point of calculation for re-optimization (whether it is PCE,
ingress, or domain entry point) needs to know all such paths before
attempting re-optimization of any one path. Mutual diversity here
means that a set of computed paths has no commonality. Such
diversity might be link, node, Shared Risk Link Group (SRLG), or even
domain disjointedness according to circumstances and the service
being delivered.

It may be the case that re-optimization is best achieved by
recomputing the paths of multiple LSPs at once. Indeed, this can be
shown to be most efficient when the paths of all LSPs are known, not
simply those LSPs that originate at a particular ingress. While this
problem is inherited from single domain re-optimization and is out of
scope within this document, it should be noted that the problem grows
in complexity when LSPs wholly within one domain affect the re-
optimization path calculations performed in another domain.

5.2. LSP Setup Failure

When an inter-domain LSP setup fails in some domain other than the
first, various options are available for reporting and retrying the
LSP.

In the first instance, a retry may be attempted within the domain
that contains the failure. That retry may be attempted by nodes
wholly within the domain, or the failure may be referred back to the
LSR at the domain boundary.

If the failure cannot be bypassed within the domain where the failure
occurred (perhaps there is no suitable alternate route, perhaps
rerouting is not allowed by domain policy, or perhaps the Path
message specifically bans such action), the error must be reported
back to the previous or head-end domain.

Subsequent repair attempts may be made by domains further upstream,
but will only be properly effective if sufficient information about
the failure and other failed repair attempts is also passed back
upstream [CRANKBACK]. Note that there is a tension between this
requirement and that of topology confidentiality although crankback
aggregation may be applicable at domain boundaries.

Further attempts to signal the failed LSP may apply the information
about the failures as constraints to path computation, or may signal
them as specific path exclusions [EXCLUDE].

When requested by signaling, the failure may also be systematically
reported to the head-end LSR.

5.3. LSP Repair

An LSP that fails after it has been established may be repaired
dynamically by re-routing. The behavior in this case is either like
that for re-optimization, or for handling setup failures (see
previous two sections). Fast Reroute may also be used (see below).

5.4. Fast Reroute

MPLS Traffic Engineering Fast Reroute ([RFC4090]) defines local
protection schemes intended to provide fast recovery (in 10s of
msecs) of fast-reroutable packet-based TE LSPs upon link/SRLG/Node
failure. A backup TE LSP is configured and signaled at each hop, and
activated upon detecting or being informed of a network element
failure. The node immediately upstream of the failure (called the
PLR, or Point of Local Repair) reroutes the set of protected TE LSPs
onto the appropriate backup tunnel(s) and around the failed resource.

In the context of inter-domain TE, there are several different
failure scenarios that must be analyzed. Provision of suitable
solutions may be further complicated by the fact that [RFC4090]
specifies two distinct modes of operation referred to as the "one to
one mode" and the "facility back-up mode".

The failure scenarios specific to inter-domain TE are as follows:

- Failure of a domain edge node that is present in both domains.
  There are two sub-cases:

  - The Point of Local Repair (PLR) and the Merge Point (MP) are in
    the same domain.

  - The PLR and the MP are in different domains.

- Failure of a domain edge node that is only present in one of the
  domains.

- Failure of an inter-domain link.

Although it may be possible to apply the same techniques for Fast
Reroute (FRR) to the different methods of signaling inter-domain LSPs
described in section 2, the results of protection may be different
when it is the boundary nodes that need to be protected, and when
they are the ingress and egress of a hierarchical LSP or stitched LSP
segment. In particular, the choice of PLR and MP may be different,
and the length of the protection path may be greater. These uses of
FRR techniques should be explained further in applicability
statements or, in the case of a change in base behavior, in
implementation guidelines specific to the signaling techniques.

Note that after local repair has been performed, it may be desirable
to re-optimize the LSP (see section 5.1). If the point of re-
optimization (for example, the ingress LSR) lies in a different
domain to the failure, it may rely on the delivery of a PathErr or
Notify message to inform it of the local repair event.

It is important to note that Fast Reroute techniques are only
applicable to packet switching networks because other network
technologies cannot apply label stacking within the same switching
type. Segment protection [GMPLS-SEG] provides a suitable alternative
that is applicable to packet and non-packet networks.

5.5. Comments on Path Diversity

Diverse paths may be required in support of load sharing and/or
protection. Such diverse paths may be required to be node diverse,
link diverse, fully path diverse (that is, link and node diverse), or
SRLG diverse.

Diverse path computation is a classic problem familiar to all graph
theory majors. The problem is compounded when there are areas of
"private knowledge" such as when domains do not share topology

information. The problem can be resolved more efficiently (e.g.,
avoiding the "trap problem") when mutually resource disjoint paths
can be computed "simultaneously" on the fullest set of information.

That being said, various techniques (out of the scope of this
document) exist to ensure end-to-end path diversity across multiple
domains.

Many network technologies utilize "protection domains" because they
fit well with the capabilities of the technology. As a result, many
domains are operated as protection domains. In this model,
protection paths converge at domain boundaries.

Note that the question of SRLG identification is not yet fully
answered. There are two classes of SRLG:

- those that indicate resources that are all contained within one
  domain

- those that span domains.

The former might be identified using a combination of a globally
scoped domain ID, and an SRLG ID that is administered by the domain.
The latter requires a global scope to the SRLG ID. Both schemes,
therefore, require external administration. The former is able to
leverage existing domain ID administration (for example, area and AS
numbers), but the latter would require a new administrative policy.

5.6. Domain-Specific Constraints

   While the meaning of certain constraints, like bandwidth, can be
   assumed to be constant across different domains, other TE constraints
   (such as resource affinity, color, metric, priority, etc.) may have
   different meanings in different domains and this may impact the
   ability to support Diffserv-aware MPLS, or to manage preemption.

   In order to achieve consistent meaning and LSP establishment, this
   fact must be considered when performing constraint-based path
   computation or when signaling across domain boundaries.

   A mapping function can be derived for most constraints based on
   policy agreements between the domain administrators. The details of
   such a mapping function are outside the scope of this document, but
   it is important to note that the default behavior must either be that
   a constant mapping is applied or that any requirement to apply these
   constraints across a domain boundary must fail in the absence of
   explicit mapping rules.

5.7. Policy Control

   Domain boundaries are natural points for policy control. There is
   little to add on this subject except to note that a TE LSP that
   cannot be established on a path through one domain because of a
   policy applied at the domain boundary may be satisfactorily
   established using a path that avoids the demurring domain. In any
   case, when a TE LSP signaling attempt is rejected due to non-
   compliance with some policy constraint, this should be reflected to
   the ingress LSR.

5.8. Inter-Domain Operations and Management (OAM)

   Some elements of OAM may be intentionally confined within a domain.
   Others (such as end-to-end liveness and connectivity testing) clearly
   need to span the entire multi-domain TE LSP. Where issues of
   topology confidentiality are strong, collaboration between PCEs or
   domain boundary nodes might be required in order to provide end-to-
   end OAM, and a significant issue to be resolved is to ensure that the
   end-points support the various OAM capabilities.

   The different signaling mechanisms described above may need
   refinements to [RFC4379], [BFD-MPLS], etc., to gain full end-to-end
   visibility. These protocols should, however, be considered in the
   light of topology confidentiality requirements.

   Route recording is a commonly used feature of signaling that provides
   OAM information about the path of an established LSP. When an LSP
   traverses a domain boundary, the border node may remove or aggregate
   some of the recorded information for topology confidentiality or
   other policy reasons.

5.9. Point-to-Multipoint

   Inter-domain point-to-multipoint (P2MP) requirements are explicitly
   out of the scope of this document. They may be covered by other
   documents dependent on the details of MPLS TE P2MP solutions.

5.10. Applicability to Non-Packet Technologies

   Non-packet switching technologies may present particular issues for
   inter-domain LSPs. While packet switching networks may utilize
   control planes built on MPLS or GMPLS technology, non-packet networks
   are limited to GMPLS.

On the other hand, some problems such as Fast Reroute on domain
boundaries (see section 5.4) may be handled by the GMPLS technique of
segment protection [GMPLS-SEG] that is applicable to both packet and
non-packet switching technologies.

The specific architectural considerations and requirements for
inter-domain LSP setup in non-packet networks are covered in a
separate document [GMPLS-AS].

6. Security Considerations

Requirements for security within domains are unchanged from [RFC3209]
and [RFC3473], and from [RFC3630] and [RFC3784]. That is, all
security procedures for existing protocols in the MPLS context
continue to apply for the intra-domain cases.

Inter-domain security may be considered as a more important and more
sensitive issue than intra-domain security since in inter-domain
traffic engineering control and information may be passed across
administrative boundaries. The most obvious and most sensitive case
is inter-AS TE.

All of the intra-domain security measures for the signaling and
routing protocols are equally applicable in the inter-domain case.
There is, however, a greater likelihood of them being applied in the
inter-domain case.

Security for inter-domain MPLS TE is the subject of a separate
document that analyzes the security deployment models and risks.
This separate document must be completed before inter-domain MPLS TE
solution documents can be advanced.

Similarly, the PCE procedures [RFC4655] are subject to security
measures for the exchange computation information between PCEs and
for LSRs that request path computations from a PCE. The requirements
for this security (set out in [RFC4657]) apply whether the LSR and
PCE (or the cooperating PCEs) are in the same domain or lie across
domain boundaries.

It should be noted, however, that techniques used for (for example)
authentication require coordination of secrets, keys, or passwords
between sender and receiver. Where sender and receiver lie within a
single administrative domain, this process may be simple. But where
sender and receiver lie in different administrative domains, cross-
domain coordination between network administrators will be required
in order to provide adequate security. At this stage, it is not
proposed that this coordination be provided through an automatic
process or through the use of a protocol. Human-to-human

   coordination is more likely to provide the required level of
   confidence in the inter-domain security.

   One new security concept is introduced by inter-domain MPLS TE. This
   is the preservation of confidentiality of topology information. That
   is, one domain may wish to keep secret the way that its network is
   constructed and the availability (or otherwise) of end-to-end network
   resources. This issue is discussed in sections 3.4.2, 5.2, and 5.8
   of this document. When there is a requirement to preserve inter-
   domain topology confidentiality, policy filters must be applied at
   the domain boundaries to avoid distributing such information. This
   is the responsibility of the domain that distributes information, and
   it may be adequately addressed by aggregation of information as
   described in the referenced sections.

   Applicability statements for particular combinations of signaling,
   routing, and path computation techniques to provide inter-domain MPLS
   TE solutions are expected to contain detailed security sections.

7. Acknowledgements

   The authors would like to extend their warmest thanks to Kireeti
   Kompella for convincing them to expend effort on this document.

   Grateful thanks to Dimitri Papadimitriou, Tomohiro Otani, and Igor
   Bryskin for their review and suggestions on the text.

   Thanks to Jari Arkko, Gonzalo Camarillo, Brian Carpenter, Lisa
   Dusseault, Sam Hartman, Russ Housley, and Dan Romascanu for final
   review of the text.

8. Normative References

   [RFC3031]   Rosen, E., Viswanathan, A., and R. Callon,
                  "Multiprotocol Label Switching Architecture", RFC 3031,
                  January 2001.

   [RFC3209]   Awduche, D., Berger, L., Gan, D., Li, T., Srinivasan,
                  V., and G. Swallow, "RSVP-TE: Extensions to RSVP for
                  LSP Tunnels", RFC 3209, December 2001.

   [RFC3473]   Berger, L., Ed., "Generalized Multi-Protocol Label
                  Switching (GMPLS) Signaling Resource ReserVation
                  Protocol-Traffic Engineering (RSVP-TE) Extensions", RFC
                  3473, January 2003.

   [RFC3630]   Katz, D., Kompella, K., and D. Yeung, "Traffic
                  Engineering (TE) Extensions to OSPF Version 2", RFC
                  3630, September 2003.

   [RFC3784]   Smit, H. and T. Li, "Intermediate System to
                  Intermediate System (IS-IS) Extensions for Traffic
                  Engineering (TE)", RFC 3784, June 2004.

9. Informative References

   [BFD-MPLS]   Aggarwal, R., Kompella, K., Nadeau, T., and G. Swallow,
                  "BFD For MPLS LSPs", Work in Progress, June 2006.

   [CRANKBACK]   Farrel, A., et al., "Crankback Signaling Extensions for
                  MPLS Signaling", Work in Progress, May 2005.

   [EXCLUDE]   Lee, CY., Farrel, A., and DeCnodder, "Exclude Routes -
                  Extension to RSVP-TE", Work in Progress, August 2005.

   [RFC4090]   Pan, P., Swallow, G., and A. Atlas, "Fast Reroute
                  Extensions to RSVP-TE for LSP Tunnels", RFC 4090, May
                  2005.

   [GMPLS-AS]   Otani, T., Kumaki, K., Okamoto, S., and W. Imajuku,
                  "GMPLS Inter-domain Traffic Engineering Requirements",
                  Work in Progress, August 2006.

   [GMPLS-E2E]   Lang, J.P., Rekhter, Y., and D. Papadimitriou, Editors,
                  "RSVP-TE Extensions in support of End-to-End
                  Generalized Multi-Protocol Label Switching (GMPLS)-
                  based Recovery", Work in Progress, April 2005.

   [GMPLS-SEG]   Berger, L., Bryskin, I., Papadimitriou, D., and A.
                  Farrel, "GMPLS Based Segment Recovery", Work in
                  Progress, May 2005.

   [RFC4206]   Kompella, K. and Y. Rekhter, "Label Switched Paths
                  (LSP) Hierarchy with Generalized Multi-Protocol Label
                  Switching (GMPLS) Traffic Engineering (TE)", RFC 4206,
                  October 2005.

   [RFC4105]   Le Roux, J.-L., Vasseur, J.-P., and J. Boyle,
                  "Requirements for Inter-Area MPLS Traffic Engineering",
                  RFC 4105, June 2005.

   [RFC4204]   Lang, J., "Link Management Protocol (LMP)", RFC 4204,
                  October 2005.

   [RFC4216]   Zhang, R. and J.-P. Vasseur, "MPLS Inter-Autonomous
                  System (AS) Traffic Engineering (TE) Requirements", RFC
                  4216, November 2005.

   [RFC4379]   Kompella, K. and G. Swallow, "Detecting Multi-Protocol
                  Label Switched (MPLS) Data Plane Failures", RFC 4379,
                  February 2006.

   [RFC4420]   Farrel, A., Papadimitriou, D., Vasseur, J.-P., and A.
                  Ayyangar, "Encoding of Attributes for Multiprotocol
                  Label Switching (MPLS) Label Switched Path (LSP)
                  Establishment Using Resource ReserVation Protocol-
                  Traffic Engineering (RSVP-TE)", RFC 4420, February
                  2006.

   [RFC4655]   Farrel, A., Vasseur, J.-P., and J. Ash, "A Path
                  Computation Element (PCE)-Based Architecture", RFC
                  4655, August 2006.

   [RFC4657]   Ash, J. and J. Le Roux, "Path Computation Element (PCE)
                  Communication Protocol Generic Requirements", RFC 4657,
                  September 2006.

   [STITCH]   Ayyangar, A. and J.-P. Vasseur, "LSP Stitching with
                  Generalized MPLS TE", Work in Progress, September 2005.

Authors' Addresses

   Adrian Farrel
   Old Dog Consulting
   EMail: adrian@olddog.co.uk

   Jean-Philippe Vasseur
   Cisco Systems, Inc
   1414 Massachusetts Avenue
   Boxborough, MA 01719
   USA
   EMail: jpv@cisco.com

   Arthi Ayyangar
   Nuova Systems
   EMail: arthi@nuovasystems.com

Full Copyright Statement

   Copyright (C) The IETF Trust (2006).

   This document is subject to the rights, licenses and restrictions
   contained in BCP 78, and except as set forth therein, the authors
   retain all their rights.

   This document and the information contained herein are provided on an
   "AS IS" basis and THE CONTRIBUTOR, THE ORGANIZATION HE/SHE REPRESENTS
   OR IS SPONSORED BY (IF ANY), THE INTERNET SOCIETY, THE IETF TRUST,
   AND THE INTERNET ENGINEERING TASK FORCE DISCLAIM ALL WARRANTIES,
   EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT
   THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY
   IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR
   PURPOSE.

Intellectual Property

   The IETF takes no position regarding the validity or scope of any
   Intellectual Property Rights or other rights that might be claimed to
   pertain to the implementation or use of the technology described in
   this document or the extent to which any license under such rights
   might or might not be available; nor does it represent that it has
   made any independent effort to identify any such rights. Information
   on the procedures with respect to rights in RFC documents can be
   found in BCP 78 and BCP 79.

   Copies of IPR disclosures made to the IETF Secretariat and any
   assurances of licenses to be made available, or the result of an
   attempt made to obtain a general license or permission for the use of
   such proprietary rights by implementers or users of this
   specification can be obtained from the IETF on-line IPR repository at
   http://www.ietf.org/ipr.

   The IETF invites any interested party to bring to its attention any
   copyrights, patents or patent applications, or other proprietary
   rights that may cover technology that may be required to implement
   this standard. Please address the information to the IETF at
   ietf-ipr@ietf.org.