       Requirements for Multicast Support in Virtual Private LAN Services

Status of This Memo

   This memo provides information for the Internet community.  It does
   not specify an Internet standard of any kind.  Distribution of this
   memo is unlimited.

Copyright Notice

Abstract

   This document provides functional requirements for network solutions
   that support multicast over Virtual Private LAN Service (VPLS).  It
   specifies requirements both from the end user and service provider
   standpoints.  It is intended that potential solutions will use these
   requirements as guidelines.

Table of Contents

## 1.  Introduction

## 1.1.  Background

   VPLS (Virtual Private LAN Service) is a provider service that
   emulates the full functionality of a traditional Local Area Network
   (LAN).  VPLS interconnects several customer LAN segments over a
   packet switched network (PSN) backbone, creating a multipoint-to-
   multipoint Ethernet VPN.  For customers, their remote LAN segments
   behave as one single LAN.

   In a VPLS, the provider network emulates a learning bridge, and
   forwarding takes place based on Ethernet MAC (media access control)
   learning.  Hence, a VPLS requires MAC address learning/aging on a
   per-PW (pseudowire) basis, where forwarding decisions treat the PW as
   a "bridge port".

   VPLS is a Layer-2 (L2) service.  However, it provides two
   applications from the customer's point of view:

   -  LAN Routing application: providing connectivity between customer
      routers

   -  LAN Switching application: providing connectivity between customer
      Ethernet switches

   Thus, in some cases, customers across MAN/WAN have transparent
   Layer-2 connectivity while their main goal is to run Layer-3
   applications within their routing domain.  As a result, different
   requirements arise from their variety of applications.

Originally, PEs (Provider Edges) in VPLS transport broadcast/
multicast Ethernet frames by replicating all multicast/broadcast
frames received from an Attachment Circuit (AC) to all PW's
corresponding to a particular Virtual Switching Instance (VSI).  Such
a technique has the advantage of keeping the P (Provider Router) and
PE devices completely unaware of IP multicast-specific issues.
Obviously, however, it has quite a few scalability drawbacks in terms
of bandwidth consumption, which will lead to increased cost in large-
scale deployment.

Meanwhile, there is a growing need for support of multicast-based
services such as IP TV.  This commercial trend makes it necessary for
most VPLS deployments to support multicast more efficiently than
before.  It is also necessary as customer routers are now likely to
be running IP multicast protocols, and those routers are connected to
switches that will be handling large amounts of multicast traffic.

Therefore, it is desirable to have more efficient techniques to
support IP multicast over VPLS.

## 1.2.  Scope of This Document

This document provides functional requirements for network solutions
that support IP multicast in VPLS [RFC4761] [RFC4762].  It identifies
requirements that MAY apply to the existing base VPLS architecture in
order to optimize IP multicast.  It also complements the generic
L2VPN requirements document [RFC4665], by specifying additional
requirements specific to the deployment of IP multicast in VPLS.

The technical specifications are outside the scope of this document.
In this document, there is no intent to specify either solution-
specific details or application-specific requirements.  Also, this
document does NOT aim to express multicast-inferred requirements that
are not specific to VPLS.  It does NOT aim to express any
requirements for native Ethernet specifications, either.

This document is proposed as a solution guideline and a checklist of
requirements for solutions, by which we will evaluate how each
solution satisfies the requirements.

This document clarifies the needs from both VPLS customer as well as
provider standpoints and formulates the problems that should be
addressed by technical solutions while staying solution agnostic.

A technical solution and corresponding service that supports this
document's requirements are hereinafter called a "multicast VPLS".

2.  Conventions Used in This Document

2.1.  Terminology

   The reader is assumed to be familiar with the terminology, reference
   models, and taxonomy defined in [RFC4664] and [RFC4665].  For
   readability purposes, we repeat some of the terms here.

   Moreover, we also propose some other terms needed when IP multicast
   support in VPLS is discussed.

   -  ASM: Any Source Multicast.  One of the two multicast service
      models where each corresponding service can have an arbitrary
      number of senders.

   -  G: denotes a multicast group.

   -  MDTunnel: Multicast Distribution Tunnel, the means by which the
      customer's multicast traffic will be conveyed across the Service
      Provider (SP) network.  This is meant in a generic way: such
      tunnels can be point-to-point, point-to-multipoint, or multipoint-
      to-multipoint.  Although this definition may seem to assume that
      distribution tunnels are unidirectional, the wording encompasses
      bidirectional tunnels as well.

   -  Multicast Channel: In the multicast SSM (Source Specific
      Multicast) model [RFC4607], a "multicast channel" designates
      traffic from a specific source S to a multicast group G.  Also
      denominated as "(S,G)".

   -  Multicast domain: An area in which multicast data is transmitted.
      In this document, this term has a generic meaning that can refer
      to Layer-2 and Layer-3.  Generally, the Layer-3 multicast domain
      is determined by the Layer-3 multicast protocol used to establish
      reachability between all potential receivers in the corresponding
      domain.  The Layer-2 multicast domain can be the same as the
      Layer-2 broadcast domain (i.e., VLAN), but it may be restricted to
      being smaller than the Layer-2 broadcast domain if an additional
      control protocol is used.

   -  CE: Customer Edge Device.

   -  PE: Provider Edge.

   -  P: Provider Router.

   -  S: denotes a multicast source.

   -  SP: Service Provider.

   -  SSM: Source Specific Multicast.  One of the two multicast service
      models where each corresponding service relies upon the use of a
      single source.

   -  U-PE/N-PE: The device closest to the customer/user is called the
      User-facing PE (U-PE) and the device closest to the core network
      is called the Network-facing PE (N-PE).

   -  VPLS instance: A service entity manageable in VPLS architecture.
      All CE devices participating in a single VPLS instance appear to
      be on the same LAN, composing a VPN across the SP's network.  A
      VPLS instance corresponds to a group of VSIs that are
      interconnected using PWs (pseudowires).

   -  VSI: Virtual Switching Instance.  A VSI is a logical entity in a
      PE that maps multiple ACs (Attachment Circuits) to multiple PWs.
      The VSI is populated in much the same way as a standard bridge
      populates its forwarding table.  Each PE device may have multiple
      VSIs, where each VSI belongs to a different VPLS instance.

## 2.2.  Conventions

   The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT",
   "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this
   document are to be interpreted as described in [RFC2119] .

## 3.  Problem Statements

## 3.1.  Motivation

   Today, many kinds of IP multicast services are becoming available.
   Over their Layer-2 VPN service, particularly over VPLS, customers
   would often like to operate their multicast applications to remote
   sites.  Also, VPN service providers using an IP-based network expect
   that such Layer-2 network infrastructure will efficiently support
   multicast data traffic.

   However, VPLS has a shortcoming as it relates to multicast
   scalability as mentioned below because of the replication mechanisms
   intrinsic to the original architecture.  Accordingly, the primary
   goal for technical solutions is to solve this issue partially or
   completely, and provide efficient ways to support IP multicast
   services over VPLS.

3.2.  Multicast Scalability

   In VPLS, replication occurs at an ingress PE (in the hierarchical
   VPLS (H-VPLS) case, at N-PE) when a CE sends (1) Broadcast, (2)
   Multicast, or (3) Unknown destination unicast.  There are two well-
   known issues with this approach:

   Issue A: Replication to non-member site:

      In cases (1) and (3), the upstream PE has to transmit packets to
      all of the downstream PEs that belong to the common VPLS instance.
      You cannot decrease the number of members, so this is basically an
      inevitable situation for most VPLS deployments.

      In case (2), however, there is an issue that multicast traffic is
      sent to sites with no members.  Usually, this is caused when the
      upstream PE does not maintain downstream membership information.
      The upstream PE simply floods frames to all downstream PEs, and
      the downstream PEs forward them to directly connected CEs;
      however, those CEs might not be the members of any multicast
      group.  From the perspective of customers, they might suffer from
      pressure on their own resources due to unnecessary traffic.  From
      the perspective of SPs, they would not like wasteful over-
      provisioning to cover such traffic.

   Issue B: Replication of PWs on shared physical path:

      In VPLS, a VSI associated with each VPLS instance behaves as a
      logical emulated bridge that can transport Ethernet across the PSN
      backbone using PWs.  In principle, PWs are designed for unicast
      traffic.

      In all cases, (1), (2), and (3), Ethernet frames are replicated on
      one or more PWs that belong to that VSI.  This replication is
      often inefficient in terms of bandwidth usage if those PWs are
      traversing shared physical links in the backbone.

      For instance, suppose there are 20 remote PEs belonging to a
      particular VPLS instance, and all PWs happen to be traversing over
      the same link from one local PE to its next-hop P.  In this case,
      even if a CE sends 50 Mbps to the local PE, the total bandwidth of
      that link will be to 1000 Mbps.

      Note that while traditional 802.1D Ethernet switches replicate
      broadcast/multicast flows once at most per output interface, VPLS
      often needs to transmit one or more flows duplicated over the same
      output interface.

From the perspective of customers, there is no serious issue
because they do not know what happens in the core.  However, from
the perspective of SPs, unnecessary replication brings the risk of
resource exhaustion when the number of PWs increases.

In both Issues A and B, these undesirable situations will become
obvious with the wide-spread use of IP multicast applications by
customers.  Naturally, the problem will become more serious as the
number of sites grows.  In other words, there are concerns over the
scalability of multicast in VPLS today.

## 3.3.  Application Considerations

### 3.3.1.  Two Perspectives of the Service

When it comes to IP multicast over VPLS, there are two different
aspects in terms of service provisioning.  They are closely related
to the functional requirements from two technical standpoints:

Layer-2 and Layer-3.

- Native Ethernet service aspect

   This aspect mainly affects Ethernet network service operators.
   Their main interest is to solve the issue that existing VPLS
   deployments cannot always handle multicast/broadcast frames
   efficiently.

   Today, wide-area Ethernet services are becoming popular, and VPLS
   can be utilized to provide wide-area LAN services.  As customers
   come to use various kinds of content distribution applications
   that use IP multicast (or other protocols that lead to multicast/
   broadcast in the Ethernet layer), the total amount of traffic will
   also grow.  In addition, considerations of Operations,
   Administration, and Management (OAM), security and other related
   points in multicast in view of Layer-2 are important.

   In such circumstances, the native VPLS specification would not
   always be satisfactory if multicast traffic is more dominant in
   total resource utilization than before.  The scalability issues
   mentioned in the previous section are expected to be solved.

- IP multicast service aspect

   This aspect mainly affects both IP service providers and end
   users.  Their main interest is to provide IP multicast services
   transparently but effectively by means of VPLS as a network
   infrastructure.

SPs might expect VPLS as an access/metro network to deliver
multicast traffic (such as Triple-play (Video, Voice, Data) and
Multicast IP VPNs) in an efficient way.

4.  General Requirements

We assume the basic requirements for VPLS written in [RFC4665] are
fulfilled unless otherwise specified in this document.

4.1.  Scope of Transport

4.1.1.  Traffic Types

4.1.1.1.  Multicast and Broadcast

As described before, any solution is expected to have mechanisms for
efficient transport of IP multicast.  Multicast is related to both
Issues A and B (see Section 3.2); however, broadcast is related to
Issue B only because it does not need membership control.

-  A multicast VPLS solution SHOULD attempt to solve both Issues A
   and B, if possible.  However, since some applications prioritize
   solving one issue over the other, the solution MUST identify which
   Issue (A or B) it is attempting to solve.  The solution SHOULD
   provide a basis for evaluating how well it solves the issue(s) it
   is targeting, if it is providing an approximate solution.

4.1.1.2.  Unknown Destination Unicast

Unknown destination MAC unicast requires flooding, but its
characteristics are quite different from multicast/broadcast.  When
the unicast MAC address is learned, the PE changes its forwarding
behavior from flooding over all PWs into sending over one PW.
Thereby, it will require different technical studies from multicast/
broadcast, which is out of scope of this document.

4.1.2.  Multicast Packet Types

Ethernet multicast is used for conveying Layer-3 multicast data.
When IP multicast is encapsulated by an Ethernet frame, the IP
multicast group address is mapped to the Ethernet destination MAC
address.  In IPv4, the mapping uses the lower 23 bits of the (32-bit)
IPv4 multicast address and places them as the lower 23 bits of a
destination MAC address with the fixed header of 01-00-5E in hex.
Since this mapping is ambiguous (i.e., there is a multiplicity of 1
Ethernet address to 32 IPv4 addresses), MAC-based forwarding is not
ideal for IP multicast because some hosts might possibly receive
packets they are not interested in, which is inefficient in traffic

delivery and has an impact on security.  On the other hand, if the
solution tracks IP addresses rather than MAC addresses, this concern
can be prevented.  The drawback of this approach is, however, that
the network administration becomes slightly more complicated.

Ethernet multicast is also used for Layer-2 control frames.  For
example, BPDU (Bridge Protocol Data Unit) for IEEE 802.1D Spanning
Trees uses a multicast destination MAC address (01-80-C2-00-00-00).
Also, some of IEEE 802.1ag [802.1ag] Connectivity Fault Management
(CFM) messages use a multicast destination MAC address dependent on
their message type and application.  From the perspective of IP
multicast, however, it is necessary in VPLS to flood such control
frames to all participating CEs, without requiring any membership
controls.

As for a multicast VPLS solution, it can only use Ethernet-related
information, if you stand by the strict application of the basic
requirement: "a L2VPN service SHOULD be agnostic to customer's Layer
3 traffic" [RFC4665].  This means no Layer-3 information should be
checked for transport.  However, it is obvious this is an impediment
to solve Issue A.

Consequently, a multicast VPLS can be allowed to make use of some
Layer-3-related supplementary information in order to improve
transport efficiency.  In fact, today's LAN-switch implementations
often support such approaches and snoop upper-layer protocols and
examine IP multicast memberships (e.g., Protocol Independent
Multicast (PIM) snooping and IGMP/MLD (Multicast Listener Discovery)
snooping [RFC4541]).  This will implicitly suggest that VPLS may
adopt similar techniques although this document does NOT state
Layer-3 snooping is mandatory.  If such an approach is taken, careful
consideration of Layer-3 state maintenance is necessary.  In
addition, note that snooping approaches sometimes have disadvantages
in the system's transparency; that is, one particular protocol's
snooping solution might hinder other (especially future) protocol's
working (e.g., an IGMPv2-snooping switch vs. a new IGMPv3-snooping
one).  Also, note that there are potential alternatives to snooping:

- Static configuration of multicast Ethernet addresses and ports/
  interfaces.

- Multicast control protocol based on Layer-2 technology that
  signals mappings of multicast addresses to ports/interfaces, such
  as Generic Attribute Registration Protocol / GARP Multicast
  Registration Protocol (GARP/GMRP) [802.1D], Cisco Group Management
  Protocol [CGMP], and Router-port Group Management Protocol (RGMP)
  [RFC3488].

On the basis described above, general requirements about packet types are given as follows:

- A solution SHOULD support a way to facilitate IP multicast forwarding of the customers.  It MAY observe Layer-3 information (i.e., multicast routing protocols and state) to the degree necessary, but any information irrelevant to multicast transport SHOULD NOT be consulted.

- In a solution, Layer-2 control frames (e.g., BPDU, 802.1ag CFM) SHOULD be flooded to all PE/CEs in a common VPLS instance.  A solution SHOULD NOT change or limit the flooding scope to remote PE/CEs in terms of end-point reachability.

- In a solution, Layer-2 frames that encapsulate Layer-3 multicast control packets (e.g., PIM, IGMP (for IPv4), and MLD (for IPv6)) MAY be flooded only to relevant members, with the goal of limiting flooding scope.  However, Layer-2 frames that encapsulate other Layer-3 control packets (e.g., OSPF, IS-IS) SHOULD be flooded to all PE/CEs in a VPLS instance.

## 4.1.3.  MAC Learning Consideration

In a common VPLS architecture, MAC learning is carried out by PEs based on the incoming frame's source MAC address, independently of the destination MAC address (i.e., regardless of whether it is unicast, multicast, or broadcast).  This is the case with the multicast VPLS solution's environment too.  In this document, the improvement of MAC learning scalability is beyond the scope.  It will be covered in future work.

## 4.2.  Static Solutions

A solution SHOULD allow static configuration to account for various operator policies, where the logical multicast topology does not change dynamically in conjunction with a customer's multicast routing.

## 4.3.  Backward Compatibility

A solution SHOULD be backward compatible with the existing VPLS solution.  It SHOULD allow a case where a common VPLS instance is composed of both PEs supporting the solution and PEs not supporting it, and the multicast optimization (both forwarding and receiving) is achieved between the compliant PEs.

Note again that the existing VPLS solutions already have a simple
flooding capability.  Thus, this backward compatibility will give
customers and SPs the improved efficiency of multicast forwarding
incrementally as the solution is deployed.

5.  Customer Requirements

5.1.  CE-PE Protocol

5.1.1.  Layer-2 Aspect

A solution SHOULD allow transparent operation of Ethernet control
protocols employed by customers (e.g., Spanning Tree Protocol
[802.1D]) and their seamless operation with multicast data transport.

Solutions MAY examine Ethernet multicast control frames for the
purpose of efficient dynamic transport (e.g., GARP/GMRP [802.1D]).
However, solutions MUST NOT assume all CEs are always running such
protocols (typically in the case where a CE is a router and is not
aware of Layer-2 details).

A whole Layer-2 multicast frame (whether for data or control) SHOULD
NOT be altered from a CE to CE(s) EXCEPT for the VLAN ID field,
ensuring that it is transparently transported.  If VLAN IDs are
assigned by the SP, they can be altered.  Note, however, when VLAN
IDs are changed, Layer-2 protocols may be broken in some cases, such
as Multiple Spanning Trees [802.1s].  Also, if the Layer-2 frame is
encapsulating a Layer-3 multicast control packet (e.g., PIM/IGMP) and
customers allow it to be regenerated at the PE (aka proxy: see
Section 5.1.2), then the MAC address for that frame MAY be altered to
the minimum necessary (e.g., use PE's own MAC address as a source).

5.1.2.  Layer-3 Aspect

Again, a solution MAY examine the customer's Layer-3 multicast
protocol packets for the purpose of efficient and dynamic transport.
If it does, supported protocols SHOULD include:

o  PIM-SM (Sparse Mode) [RFC4601], PIM-SSM (Source-Specific
   Multicast) [RFC4607], bidirectional PIM [RFC5015], and PIM-DM
   (Dense Mode) [RFC3973].

o  IGMP (v1 [RFC1112], v2 [RFC2236], and v3 [RFC3376]) (for IPv4
   solutions).

o  Multicast Listener Discovery Protocol (MLD) (v1 [RFC2710] and v2
   [RFC3810]) (for IPv6 solutions).

A solution MUST NOT require any special Layer-3 multicast protocol
packet processing by the end users.  However, it MAY require some
configuration changes (e.g., turning explicit tracking on/off in the
PIM).

A whole Layer-3 multicast packet (whether for data or control), which
is encapsulated inside a Layer-2 frame, SHOULD NOT be altered from a
CE to CE(s), ensuring that it is transparently transported.  However,
as for Layer-3 multicast control (like PIM Join/Prune/Hello and IGMP
Query/Report packet), it MAY be altered to the minimum necessary if
such partial non-transparency is acceptable from point of view of the
multicast service.  Similarly, a PE MAY consume such Layer-3
multicast control packets and regenerate an entirely new packet if
partial non-transparency is acceptable with legitimate reason for
customers (aka proxy).

## 5.2.  Multicast Domain

As noted in Section 2.1, the term "multicast domain" is used in a
generic context for Layer-2 and Layer-3.

A solution SHOULD NOT alter the boundaries of customer multicast
domains.  It MUST ensure that the provided Ethernet multicast domain
always encompasses the corresponding customer Layer-3 multicast
domain.

A solution SHOULD optimize those domains' coverage sizes, i.e., a
solution SHOULD ensure that unnecessary traffic is not sent to CEs
with no members.  Ideally, the provided domain size will be close to
that of the customer's Layer-3 multicast membership distribution;
however, it is OPTIONAL to achieve such absolute optimality from the
perspective of Layer-3.

If a customer uses VLANs and a VLAN ID as a service delimiter (i.e.,
each VPLS instance is represented by a unique customer VLAN tag
carried by a frame through the User Network Interface (UNI) port), a
solution MUST provide a separate multicast domain for each VLAN ID.
Note that if VLAN ID translation is provided (i.e., if a customer
VLAN at one site is mapped into a different customer VLAN at a
different site), multicast domains will be created per set of VLAN
IDs that are associated with translation.

If a customer uses VLANs but a VLAN ID is not a service delimiter
(i.e., the VPN disregards customer VLAN IDs), a solution MAY provide
a separate multicast domain for each VLAN ID.  An SP is not
mandatorily required to provide a separate multicast domain for each
VLAN ID, but it may be considered beneficial to do so.

A solution MAY build multicast domains based on Ethernet MAC
addresses.  It MAY also build multicast domains based on the IP
addresses inside Ethernet frames.  That is, PEs in each VPLS instance
might control forwarding behavior and provide different multicast
frame reachability depending on each MAC/IP destination address
separately.  If IP multicast channels are fully considered in a
solution, the provided domain size will be closer to actual channel
reachability.

## 5.3.  Quality of Service (QoS)

Customers require that multicast quality of service MUST be at least
on par with what exists for unicast traffic.  Moreover, as multicast
is often used to deliver high-quality services such as TV broadcast,
delay-, jitter-, and loss-sensitive traffic MUST be supported over
multicast VPLS.

To accomplish this, the solution MAY have additional features to
support high QoS such as bandwidth reservation and flow admission
control.  Also, multicast VPLS deployment SHALL benefit from IEEE
802.1p Class-of-Service (CoS) techniques [802.1D] and Diffserv
[RFC2475] mechanisms.

Moreover, multicast traffic SHOULD NOT affect the QoS that unicast
traffic receives and vice versa.  That is, separation of multicast
and unicast traffic in terms of QoS is necessary.

## 5.4.  SLA Parameters Measurement

Since SLA parameters are part of the service sold to customers, they
simply want to verify their application performance by measuring the
parameters SP(s) provide.

Multicast specific characteristics that may be monitored are, for
instance, multicast statistics per stream (e.g., total/incoming/
outgoing/dropped traffic by period of time), one-way delay, jitter
and group join/leave delay (time to start receiving traffic from a
multicast group across the VPN since the join/leave was issued).  An
operator may also wish to compare the difference in one-way delay for
a solitary multicast group/stream from a single, source PE to
multiple receiver PEs.

A solution SHOULD provide these parameters with Ethernet multicast
group level granularity.  (For example, a multicast MAC address will
be one of those entries for classifying flows with statistics, delay,
and so on.)  However, if a solution is aimed at IP multicast
transport efficiency, it MAY support IP multicast-level granularity.

(For example, multicast IP address/channel will be entries for
latency time.)

In order to monitor them, standard interfaces for statistics
gathering SHOULD also be provided (e.g., standard Simple Network
Management Protocol (SNMP) MIB Modules).

## 5.5.  Security

A solution MUST provide customers with architectures that give the
same level of security both for unicast and multicast.

### 5.5.1.  Isolation from Unicast

Solutions SHOULD NOT affect any forwarding information base,
throughput, or resiliency, etc., of unicast frames; that is, they
SHOULD provide isolation from unicast.

### 5.5.2.  Access Control

A solution MAY filter multicast traffic inside a VPLS, upon the
request of an individual customer, (for example, MAC/VLAN filtering,
IP multicast channel filtering, etc.).

### 5.5.3.  Policing and Shaping on Multicast

A solution SHOULD support policing and shaping multicast traffic on a
per-customer basis and on a per-AC (Attachment Circuit) basis.  This
is intended to prevent multicast traffic from exhausting resources
for unicast inside a common customer's VPN.  This might also be
beneficial for QoS separation (see Section 5.3).

## 5.6.  Access Connectivity

First and foremost, various physical connectivity types described in
[RFC4665] MUST be supported.

## 5.7.  Multi-Homing

A multicast VPLS MUST allow a situation in which a CE is dual-homed
to two different SPs via diverse access networks -- one is supporting
multicast VPLS but the other is not supporting it, (because it is an
existing VPLS or 802.1Q/QinQ network).

## 5.8.  Protection and Restoration

A multicast VPLS infrastructure SHOULD allow redundant paths to
assure high availability.

Multicast forwarding restoration time MUST NOT be greater than the
time it takes a customer's Layer-3 multicast protocols to detect a
failure in the VPLS infrastructure.  For example, if a customer uses
PIM with default configuration, the hello hold timer is 105 seconds,
and solutions are required to restore a failure no later than this
period.  To achieve this, a solution might need to support providing
alternative multicast paths.

Moreover, if multicast forwarding was not successfully restored
(e.g., in case of no redundant paths), a solution MAY raise alarms to
provide outage notification to customers before such a hold timer
expires.

## 5.9.  Minimum MTU

Multicast applications are often sensitive to packet fragmentation
and reassembly, so the requirement to avoid fragmentation might be
stronger than the existing VPLS solution.

A solution SHOULD provide customers with enough committed minimum MTU
(i.e., service MTU) for multicast Ethernet frames to ensure that IP
fragmentation between customer sites never occurs.  It MAY give
different MTU sizes to multicast and unicast.

## 5.10.  Frame Reordering Prevention

A solution SHOULD attempt to prevent frame reordering when delivering
customer multicast traffic.  Likewise, for unicast and unknown
unicast traffic, it SHOULD attempt not to increase the likelihood of
reordering compared with existing VPLS solutions.

It is to be noted that delivery of out-of-order frames is not
avoidable in certain cases.  Specifically, if a solution adopts some
MDTunnels (see Section 6.2) and dynamically selects them for
optimized delivery (e.g., switching from one aggregate tree to
another), end-to-end data delivery is prone to be out of order.  This
fact can be considered a trade-off between bandwidth optimization and
network stability.  Therefore, such a solution is expected to promote
awareness about this kind of drawback.

## 5.11.  Fate-Sharing between Unicast and Multicast

In native Ethernet, multicast and unicast connectivity are often
managed together.  For instance, an 802.1ag CFM Continuity Check
message is forwarded by multicast as a periodic heartbeat, but it is
supposed to check the "whole" traffic continuity regardless of
unicast or multicast, at the same time.  Hence, the aliveness of

unicast and multicast is naturally coupled (i.e., fate-shared) in
this customer's environment.

A multicast VPLS solution may decouple the path that a customer's
unicast and multicast traffic follow through a SP's backbone, in
order to provide the most optimal path for multicast data traffic.
This may cause concern among some multicast VPLS customers who desire
that, during a failure in the SP's network, both unicast and
multicast traffic fail concurrently.

Therefore, there will be an additional requirement that makes both
unicast and multicast connectivity coupled.  This means that if
either one of them have a failure, the other is also disabled.  If
one of the services (either unicast or multicast) becomes
operational, the other is also activated simultaneously.

-  It SHOULD be identified if the solution can provide customers with
   fate-sharing between unicast and multicast connectivity for their
   LAN switching application.  It MAY have a configurable mechanism
   for SPs to provide that on behalf of customers, e.g., aliveness
   synchronization, but its use is OPTIONAL.

This policy will benefit customers.  Some customers would like to
detect failure soon at CE side and restore full connectivity by
switching over to their backup line, rather than to keep poor half
connectivity (i.e., either unicast or multicast being in fail).  Even
if either unicast or multicast is kept alive, it is just
disadvantageous to the customer's application protocols that need
both types of traffic.  Fate-sharing policy contributes to preventing
such a complicated situation.

Note that how serious this issue is depends on each customer's stance
in Ethernet operation.  If all CEs are IP routers, i.e., if VPLS is
provided for a LAN routing application, the customer might not care
about it because both unicast and multicast connectivity is assured
in the IP layer.  If the CE routers are running an IGP (e.g., OSPF/
IS-IS) and a multicast routing protocol (e.g., PIM), then aliveness
of both the unicast and multicast paths will be detected by the CEs.
This does not guarantee that unicast and multicast traffic are to
follow the same path in the SP's backbone network, but does mitigate
this issue to some degree.

6.  Service Provider Network Requirements

6.1.  Scalability

   The existing VPLS architecture has major advantages in scalability.
   For example, P-routers are free from maintaining customers'
   information because customer traffic is encapsulated in PSN tunnels.
   Also, a PW's split-horizon technique can prevent loops, making PE
   routers free from maintaining complicated spanning trees.

   However, a multicast VPLS needs additional scalability considerations
   related to its expected enhanced mechanisms.  [RFC3809] lists common
   L2VPN sizing and scalability requirements and metrics, which are
   applicable in multicast VPLS too.  Accordingly, this section deals
   with specific requirements related to scalability.

6.1.1.  Trade-Off of Optimality and State Resource

   A solution needs to improve the scalability of multicast as is shown
   in Section 3:

      Issue A: Replication to non-member site.

      Issue B: Replication of PWs on shared physical path.

   For both issues, the optimization of physical resources (i.e., link
   bandwidth usage and router duplication performance) will become a
   major goal.  However, there is a trade-off between optimality and
   state resource consumption.

   In order to solve Issue A, a PE might have to maintain multicast
   group information for CEs that was not kept in the existing VPLS
   solutions.  This will present scalability concerns about state
   resources (memory, CPU, etc.) and their maintenance complexity.

   In order to solve Issue B, PE and P routers might have to have
   knowledge of additional membership information for remote PEs, and
   possibly additional tree topology information, when they are using
   point-to-multipoint (P2MP) techniques (PIM tree, P2MP-LSP (Label
   Switched Path), etc.).

   Consequently, the scalability evaluation of multicast VPLS solutions
   needs a careful trade-off analysis between bandwidth optimality and
   state resource consumption.

6.1.2.  Key Metrics for Scalability

     (Note: This part has a number of similar characteristics to
     requirements for Layer-3 Multicast VPN [RFC4834].)

   A multicast VPLS solution MUST be designed to scale well with an
   increase in the number of any of the following metrics:

   -  the number of PEs

   -  the number of VPLS instances (total and per PE)

   -  the number of PEs and sites in any VPLS instance

   -  the number of client VLAN IDs

   -  the number of client Layer-2 MAC multicast groups

   -  the number of client Layer-3 multicast channels (groups or source-
      groups)

   -  the number of PWs and PSN Tunnels (MDTunnels) (total and per PE)

   Each multicast VPLS solution SHALL document its scalability
   characteristics in quantitative terms.  A solution SHOULD quantify
   the amount of state that a PE and a P device has to support.

   The scalability characteristics SHOULD include:

   -  the processing resources required by the control plane in managing
      PWs (neighborhood or session maintenance messages, keepalives,
      timers, etc.)

   -  the processing resources required by the control plane in managing
      PSN tunnels

   -  the memory resources needed for the control plane

   -  the amount of protocol information transmitted to manage a
      multicast VPLS (e.g., signaling throughput)

   -  the amount of Layer-2/Layer-3 multicast information a P/PE router
      consumes (e.g., traffic rate of join/leave, keepalives, etc.)

   -  the number of multicast IP addresses used (if IP multicast in ASM
      mode is proposed as a multicast distribution tunnel)

- other particular elements inherent to each solution that impact
  scalability

Another metric for scalability is operational complexity.  Operations
will naturally become more complicated if the number of managed
objects (e.g., multicast groups) increases, or the topology changes
occur more frequently.  A solution SHOULD note the factors that lead
to additional operational complexity.

## 6.2.  Tunneling Requirements

## 6.2.1.  Tunneling Technologies

An MDTunnel denotes a multicast distribution tunnel.  This is a
generic term for tunneling where customer multicast traffic is
carried over a provider's network.  In the L2VPN service context, it
will correspond to a PSN tunnel.

A solution SHOULD be able to use a range of tunneling technologies,
including point-to-point (unicast oriented) and point-to-multipoint/
multipoint-to-multipoint (multicast oriented).  For example, today
there are many kinds of protocols for tunneling such as L2TP, IP,
(including multicast IP trees), MPLS (including P2MP-LSP [RFC4875],
and P2MP/MP2MP-LSP [LDP-P2MP]), etc.

Note that which variant, point-to-point, point-to-multipoint, or
multipoint-to-multipoint, is used depends largely on the trade-offs
mentioned above and the targeted network and applications.
Therefore, this document does not mandate any specific protocols.  A
solution, however, SHOULD state reasonable criteria if it adopts a
specific kind of tunneling protocol.

## 6.2.2.  MTU of MDTunnel

From the view of an SP, it is not acceptable to have fragmentation/
reassembly so often while packets are traversing a MDTunnel.
Therefore, a solution SHOULD support a method that provides the
minimum path MTU of the MDTunnel in order to accommodate the service
MTU.

## 6.3.  Robustness

Multicast VPLS solutions SHOULD avoid single points of failures or
propose technical solutions that make it possible to implement a
failover mechanism.

## 6.4.  Discovering Related Information

The operation of a multicast VPLS solution SHALL be as light as possible, and providing automatic configuration and discovery SHOULD be considered a high priority.

Therefore, in addition to the L2VPN discovery requirements in [RFC4665], a multicast VPLS solution SHOULD provide a method that dynamically allows multicast membership information to be discovered by PEs if the solution supports (A), as defined in Section 3.2.  This means, a PE needs to discover multicast membership (e.g., join group addresses) that is controlled dynamically from the sites connected to that PE.  In addition, a PE needs to discover such information automatically from other remote PEs as well in order to limit flooding scope across the backbone.

## 6.5.  Operation, Administration, and Maintenance

## 6.5.1.  Activation

The activation of multicast enhancement in a solution MUST be possible:

o  with a VPLS instance granularity.

o  with an Attachment Circuit granularity (i.e., with a PE-CE Ethernet port granularity, or with a VLAN ID granularity when it is a service delimiter).

Also it SHOULD be possible:

o  with a CE granularity (when multiple CEs of the same VPN are associated with a common VPLS instance).

o  with a distinction between multicast reception and emission.

o  with a multicast MAC address granularity.

o  with a customer IP multicast group and/or channel granularity (when Layer-3 information is consulted).

Also it MAY be possible:

o  with a VLAN ID granularity when it is not a service delimiter.

6.5.2.  Testing

   A solution MUST provide a mechanism for testing multicast data
   connectivity and verifying the associated information.  Examples that
   SHOULD be supported that are specific to multicast are:

   -  Testing connectivity per multicast MAC address

   -  Testing connectivity per multicast Layer-3 group/channel

   -  Verifying data plane and control plane integrity (e.g., PW,
      MDTunnel)

   -  Verifying multicast membership-relevant information (e.g.,
      multicast MAC-addresses/PW-ports associations, Layer-3 group
      associations)

   Operators usually want to test if an end-to-end multicast user's
   connectivity is OK before and after activation.  Such end-to-end
   multicast connectivity checking SHOULD enable the end-to-end testing
   of the data path used by that customer's multicast data packets.
   Specifically, end-to-end checking will have a CE-to-CE path test and
   PE-to-PE path test.  A solution MUST support the PE-to-PE path test
   and MAY support the CE-to-CE path test.

   Also, operators will want to make use of a testing mechanism for
   diagnosis and troubleshooting.  In particular, a solution SHOULD be
   able to monitor information describing how client multicast traffic
   is carried over the SP network.  Note that if a solution supports
   frequent dynamic membership changes with optimized transport,
   troubleshooting within the SP's network will tend to be difficult.

6.5.3.  Performance Management

   Mechanisms to monitor multicast-specific parameters and statistics
   MUST be offered to the SP.

      (Note: This part has a number of similar characteristics to
      requirements for Layer 3 Multicast VPN [RFC4834].)

   A solution MUST provide SPs with access to:

   -  Multicast traffic statistics (total traffic forwarded, incoming,
      outgoing, dropped, etc., by period of time).

A solution SHOULD provide access to:

- Information about a customer's multicast resource usage (the
  amount of multicast state and throughput).

- Performance information related to multicast traffic usage, e.g.,
  one-way delay, jitter, loss, delay variations (the difference in
  one-way delay for a solitary multicast group/stream from a single,
  source PE to multiple receiver PEs), etc.

- Alarms when limits are reached on such resources.

- Statistics on decisions related to how client traffic is carried
  on MDTunnels (e.g., "How much traffic was switched onto a
  multicast tree dedicated to such groups or channels").

- Statistics on parameters that could help the provider to evaluate
  its optimality/state trade-off.

All or part of this information SHOULD be made available through
standardized SNMP MIB Modules (Management Information Base).

6.5.4.  Fault Management

A multicast VPLS solution needs to consider those management steps
taken by SPs below:

o  Fault detection

   A solution MUST provide tools that detect group membership/
   reachability failure and traffic looping for multicast
   transport.  It is anticipated that such tools are coordinated
   with the testing mechanisms mentioned in Section 6.5.2.

   In particular, such mechanisms SHOULD be able to detect a
   multicast failure quickly, (on par with unicast cases).  It
   SHOULD also avoid situations where multicast traffic has been
   in a failure state for a relatively long time while unicast
   traffic remains operational.  If such a situation were to
   occur, it would end up causing problems with customer
   applications that depend on a combination of unicast and
   multicast forwarding.

   With multicast, there may be many receivers associated with a
   particular multicast stream/group.  As the number of receivers
   increases, the number of places (typically nearest the
   receivers) required to detect a fault will increase
   proportionately.  This raises concerns over the scalability of

fault detection in large multicast deployments.  Consequently,
a fault detection solution SHOULD scale well; in particular, a
solution should consider key metrics for scalability as
described in Section 6.1.2.

o  Fault notification

A solution MUST also provide fault notification and trouble
tracking mechanisms (e.g., SNMP-trap and syslog).

In case of multicast, one point of failure often affects a
number of downstream routers/receivers that might be able to
raise a notification.  Hence, notification messages MAY be
summarized or compressed for operators' ease of management.

o  Fault isolation

A solution MUST provide diagnostic/troubleshooting tools for
multicast as well.  Also, it is anticipated that such tools are
coordinated with the testing mechanisms mentioned in
Section 6.5.2.

In particular, a solution needs to correctly identify the area
inside a multicast group impacted by the failure.  A solution
SHOULD be able to diagnose if an entire multicast group is
faulty or if some specific destinations are still alive.

## 6.6.  Security

## 6.6.1.  Security Threat Analysis

In multicast VPLS, there is a concern that one or more customer nodes
(presumably untrusted) might cause multicast-related attacks to the
SP network.  There is a danger that it might compromise some
components that belong to the whole system.

This subsection states possible security threats relevant to the
system and whether or not they are protected against.

General security consideration about a base VPLS (as part of L2VPNs)
is referred to in [RFC4665].  The following is the threat analysis
list that is inherent to multicast VPLS.

(a)  Attack by a huge amount of multicast control packets.

There is a threat that a CE joins too many multicast groups and
causes Denial of Service (DoS).  This is caused by sending a
large number of join/prune messages in a short time and/or

putting a large variety of group addresses in join/prune
messages.  This attack will waste PE's control resources (e.g.,
CPU, memory) that examine customer control messages (for solving
Issue A in Section 3.2), and it will not continue expected
services for other trusted customers.

(b)   Attack by invalid/malformed multicast control packets.

There is a threat that a CE sends invalid or malformed control
packets that might corrupt PE, which will cause a DoS attack.
In particular, a CE might be spoofing legitimate source/group IP
multicast addresses in such control packets (in PIM, IGMP, etc.)
and source/destination MAC addresses as Layer-2 frames.

(c)   Attack by rapid state change of multicast.

If a malicious CE changes multicast state by sending control
packets in an extremely short period, this might affect PE's
control resources (e.g., CPU, memory) to follow such state
changes.  Besides, it might also affect PE/P's control resources
if MDTunnel inside the core is dynamically created in
conjunction with customer's multicast group.

(d)   Attack by high volume of multicast/broadcast data traffic.

A malicious CE might send a very high volume of multicast and/or
broadcast data to a PE.  If that PE does not provide any
safeguards, it will cause excessive replication in the SP
network and the bandwidth resources for other trusted customers
might be exhausted.

(e)   Attack by high volume of unknown destination unicast data
traffic.

A malicious CE can send a high volume of unknown unicast to a
PE.  Generally, according to VPLS architecture, that PE must
flood such unknown traffic to all corresponding PEs in the same
VPN.  A variety of unknown destinations and huge amount of such
frames might cause excess traffic in SP network unless there is
an appropriate safeguard provided.

## 6.6.2.  Security Requirements

Based on the analysis in the previous subsection, the security
requirements from the SP's perspective are shown as follows.

An SP network MUST be invulnerable to malformed or maliciously constructed customer traffic.  This applies to both multicast data packets and multicast control packets.

Moreover, because multicast, broadcast, and unknown-unicast need more resources than unicast, an SP network MUST have safeguards against unwanted or malicious multicast traffic.  This applies to both multicast data packets and multicast control packets.

Specifically, a multicast VPLS solution SHOULD have mechanisms to protect an SP network from:

(1)  invalid multicast MAC addresses

(2)  invalid multicast IP addresses

(3)  malformed Ethernet multicast control protocol frames

(4)  malformed IP multicast control protocol packets

(5)  high volumes of

   *  valid/invalid customer control packets

   *  valid/invalid customer data packets (broadcast/multicast/
      unknown-unicast)

Depending on each solution's actual approach to tackle with Issue A, or B, or both (see Section 3.2.), there are relationships to be highlighted about each item's importance listed above.  First off, protection against (3) and (4) becomes significantly important if a solution supports solving Issue A, and PEs are processing customer's Ethernet/IP multicast control messages from CE.  Moreover, protection against (2) should also be much focused because PIM/IGMP snooping will usually require that PE's data forwarding be based on IP addresses.  By contrast, however, if a solution is solving only Issue B, not A, then PEs might never process the customer's multicast control messages at all; they do not perform IP address-based forwarding, but they do perform native Ethernet forwarding.  If so, there is relatively less danger about (2), (3), and (4) compared to the first case.

The following are a few additional guidelines in detail.

   For protecting against threat (a), a solution SHOULD support imposing some bounds on the quantity of state used by a VPN to be imposed in order to prevent state resource exhaustion (i.e., lack of memory, CPU etc.).  In this case, the bounds MUST be

configurable per VPN basis, not the total of various VPNs so that
SP can isolate the resource waste that is caused by any malicious
customer.

For protecting against threat (d) and (e), a solution SHOULD
support performing traffic policing to limit the unwanted data
traffic shown above.  In this case, while policing MAY be
configurable to the sum of unicast, multicast, broadcast, and
unknown unicast traffic, it SHOULD also be configurable to each
such type of traffic individually in order to prevent physical
resource exhaustion (i.e., lack of bandwidth and degradation of
throughput).  If the policing limit is configured on total traffic
only, there will be a concern that one customer's huge multicast
might close other irrelevant unicast traffic.  If it can be
configured individually, this concern will be avoided.  Moreover,
such a policing mechanism MUST be configurable per VPN basis, not
the total of various VPNs to isolate malicious customer's traffic
from others.

For protecting against threat (c), a solution SHOULD be able to
limit frequent changes of group membership by customers.  For
example, PEs might support a dampening mechanism that throttles
their multicast state changes if the customers are changing too
excessively.  Also, if MDTunnel is provided being tightly coupled
to dynamic changes of customer's multicast domain, it is also
effective to delay building the tunnel when customer's state is
changed frequently.

Protecting against threat (b) might not be an easy task.
Generally, checking the legitimacy of a customer's IP multicast
control packets will eventually require the authentication between
PE and CE in Layer-3; however, L2VPN (including VPLS) by its
nature does not usually assume Layer-3-based security mechanism
supported at PE-CE level.

The ramification of this fact is that there remains a possibility
that a PE's control plain might be badly affected by corrupted
multicast control packets that the PE is examining.  Hence, each
PE implementation will need to make an effort to minimize this
impact from malicious customers and isolate it from other trusted
customers as much as possible.

Nevertheless, it is possible to mitigate this threat to some
degree.  For example, a PE MAY support a filter mechanism about
MAC and IP addresses in a Layer-2/Layer-3 header and a filter
mechanism about source/group addresses in the multicast join/prune
messages.  This will help a PE to validate customers' control
messages, to a certain extent.

## 6.7. Hierarchical VPLS support

A VPLS multicast solution SHOULD allow a hierarchical VPLS (H-VPLS) [RFC4762] service model.  In other words, a solution is expected to operate seamlessly with existing hub and spoke PW connectivity.

Note that it is also important to take into account the case of redundant spoke connections between U-PEs and N-PEs.

## 6.8. L2VPN Wholesale

A solution MUST allow a situation where one SP is offering L2VPN services to another SP.  One example here is a wholesale model where one VPLS interconnects other SPs' VPLS or 802.1D network islands. For customer SPs, their multicast forwarding can be optimized by making use of multicast VPLS in the wholesaler SP.

## 7. Security Considerations

Security concerns and requirements for a base VPLS solution are described in [RFC4665].

In addition, there are security considerations specific to multicast VPLS.  Thus, a set of security issues have been identified that MUST be addressed when considering the design and deployment of multicast VPLS.  Such issues have been described in Sections 5.5 and 6.6.

In particular, security requirements from the view of customers are shown in Section 5.5.  Security requirements from the view of providers are shown in Section 6.6.  Section 6.6.1 conducts security threat analysis about the provider's whole system.  Section 6.6.2 explains how each threat can be addressed or mitigated.

## 8. Acknowledgments

The authors thank the contributors of [RFC4834] since the structure and content of this document were, for some sections, largely inspired by [RFC4834].

The authors also thank Yuichi Ikejiri, Jerry Ash, Bill Fenner, Vach Kompella, Shane Amante, Ben Niven-Jenkins, and Venu Hemige for their valuable reviews and feedback.

## 9.  References

### 9.1.  Normative References

   [RFC2119]   Bradner, S., "Key words for use in RFCs to Indicate
               Requirement Levels", BCP 14, RFC 2119, March 1997.

   [RFC4665]   Augustyn, W. and Y. Serbest, "Service Requirements for
               Layer 2 Provider-Provisioned Virtual Private Networks",
               RFC 4665, September 2006.

### 9.2.  Informative References

   [802.1D]    IEEE Std 802.1D-2004, "IEEE Standard for Local and
               Metropolitan Area Networks: Media Access Control (MAC)
               Bridges", 2004.

   [802.1ag]   IEEE Std 802.1ag-2007, "Virtual Bridged Local Area
               Networks - Amendment 5: Connectivity Fault Management",
               2007.

   [802.1s]    IEEE Std 802.1s-2002, "Virtual Bridged Local Area
               Networks - Amendment 3: Multiple Spanning Trees", 2002.

   [CGMP]      Farinacci, D., Tweedly, A., and T. Speakman, "Cisco Group
               Management Protocol (CGMP)", 1996/1997,
               <ftp://ftpeng.cisco.com/ipmulticast/specs/cgmp.txt>.

   [LDP-P2MP]  Minei, I., Ed., Kompella, K., Wijnands, I., and B.
               Thomas, "Label Distribution Protocol Extensions for
               Point-to-Multipoint and Multipoint-to-Multipoint Label
               Switched Paths", Work in Progress, May 2008.

   [RFC1112]   Deering, S., "Host extensions for IP multicasting",
               STD 5, RFC 1112, August 1989.

   [RFC2236]   Fenner, W., "Internet Group Management Protocol, Version
               2", RFC 2236, November 1997.

   [RFC2475]   Blake, S., Black, D., Carlson, M., Davies, E., Wang, Z.,
               and W. Weiss, "An Architecture for Differentiated
               Services", RFC 2475, December 1998.

   [RFC2710]   Deering, S., Fenner, W., and B. Haberman, "Multicast
               Listener Discovery (MLD) for IPv6", RFC 2710,
               October 1999.

   [RFC3376]    Cain, B., Deering, S., Kouvelas, I., Fenner, B., and A.
                Thyagarajan, "Internet Group Management Protocol, Version
                3", RFC 3376, October 2002.

   [RFC3488]    Wu, I. and T. Eckert, "Cisco Systems Router-port Group
                Management Protocol (RGMP)", RFC 3488, February 2003.

   [RFC3809]    Nagarajan, A., "Generic Requirements for Provider
                Provisioned Virtual Private Networks (PPVPN)", RFC 3809,
                June 2004.

   [RFC3810]    Vida, R. and L. Costa, "Multicast Listener Discovery
                Version 2 (MLDv2) for IPv6", RFC 3810, June 2004.

   [RFC3973]    Adams, A., Nicholas, J., and W. Siadak, "Protocol
                Independent Multicast - Dense Mode (PIM-DM): Protocol
                Specification (Revised)", RFC 3973, January 2005.

   [RFC4541]    Christensen, M., Kimball, K., and F. Solensky,
                "Considerations for Internet Group Management Protocol
                (IGMP) and Multicast Listener Discovery (MLD) Snooping
                Switches", RFC 4541, May 2006.

   [RFC4601]    Fenner, B., Handley, M., Holbrook, H., and I. Kouvelas,
                "Protocol Independent Multicast - Sparse Mode (PIM-SM):
                Protocol Specification (Revised)", RFC 4601, August 2006.

   [RFC4607]    Holbrook, H. and B. Cain, "Source-Specific Multicast for
                IP", RFC 4607, August 2006.

   [RFC4664]    Andersson, L. and E. Rosen, "Framework for Layer 2
                Virtual Private Networks (L2VPNs)", RFC 4664,
                September 2006.

   [RFC4761]    Kompella, K. and Y. Rekhter, "Virtual Private LAN Service
                (VPLS) Using BGP for Auto-Discovery and Signaling",
                RFC 4761, January 2007.

   [RFC4762]    Lasserre, M. and V. Kompella, "Virtual Private LAN
                Service (VPLS) Using Label Distribution Protocol (LDP)
                Signaling", RFC 4762, January 2007.

   [RFC4834]    Morin, T., Ed., "Requirements for Multicast in Layer 3
                Provider-Provisioned Virtual Private Networks (PPVPNs)",
                RFC 4834, April 2007.

   [RFC4875]   Aggarwal, R., Papadimitriou, D., and S. Yasukawa,
               "Extensions to Resource Reservation Protocol - Traffic
               Engineering (RSVP-TE) for Point-to-Multipoint TE Label
               Switched Paths (LSPs)", RFC 4875, May 2007.

   [RFC5015]   Handley, M., Kouvelas, I., Speakman, T., and L. Vicisano,
               "Bidirectional Protocol Independent Multicast (BIDIR-
               PIM)", RFC 5015, October 2007.

Authors' Addresses

   Yuji Kamite (editor)
   NTT Communications Corporation
   Granpark Tower
   3-4-1 Shibaura, Minato-ku
   Tokyo  108-8118
   Japan
   EMail: y.kamite@ntt.com

   Yuichiro Wada
   NTT
   3-9-11 Midori-cho
   Musashino-shi
   Tokyo  180-8585
   Japan
   EMail: wada.yuichiro@lab.ntt.co.jp

   Yetik Serbest
   AT&T Labs
   9505 Arboretum Blvd.
   Austin, TX  78759
   USA
   EMail: yetik_serbest@labs.att.com

   Thomas Morin
   France Telecom R&D
   2, avenue Pierre-Marzin
   22307 Lannion Cedex
   France
   EMail: thomas.morin@francetelecom.com

   Luyuan Fang
   Cisco Systems, Inc.
   300 Beaver Brook Road
   Boxborough, MA  01719
   USA
   EMail: lufang@cisco.com