

Internet Engineering Task Force (IETF)
Request for Comments: 6641
Category: Standards Track
ISSN: 2070-1721

C. Everhart
W. Adamson
NetApp
J. Zhang
Google
June 2012

Using DNS SRV to Specify a Global File Namespace with NFS Version 4

Abstract

The NFS version 4 (NFSv4) protocol provides a mechanism for a collection of NFS file servers to collaborate in providing an organization-wide file namespace. The DNS SRV Resource Record (RR) allows a simple way for an organization to publish the root of its file system namespace, even to clients that might not be intimately associated with such an organization. The DNS SRV RR can be used to join these organization-wide file namespaces together to allow construction of a global, uniform NFS file namespace.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 5741.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <http://www.rfc-editor.org/info/rfc6641>.

Copyright Notice

Copyright (c) 2012 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<http://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

This document may contain material from IETF Documents or IETF Contributions published or made publicly available before November 10, 2008. The person(s) controlling the copyright in some of this material may not have granted the IETF Trust the right to allow modifications of such material outside the IETF Standards Process. Without obtaining an adequate license from the person(s) controlling the copyright in such materials, this document may not be modified outside the IETF Standards Process, and derivative works of it may not be created outside the IETF Standards Process, except to format it for publication as an RFC or to translate it into languages other than English.

Table of Contents

1. Background	3
2. Requirements Notation	3
3. Use of the SRV Resource Record in DNS	3
4. Integration with Use of NFS Version 4	5
4.1. Globally Useful Names: Conventional Mount Point	5
4.2. Mount Options	6
4.3. File System Integration Issues	6
4.4. Multicast DNS	7
5. Where Is This Integration Carried Out?	7
6. Security Considerations	7
7. IANA Considerations	9
8. References	9
8.1. Normative References	9
8.2. Informative References	10

1. Background

Version 4 of the NFS protocol [RFC3530] introduced the `fs_locations` attribute. Use of this attribute was elaborated further in the NFSv4 minor version 1 protocol [RFC5661], which also defined an extended version of the attribute as `fs_locations_info`. With the advent of these attributes, NFS servers can cooperate to build a file namespace that crosses server boundaries. The `fs_locations` and `fs_locations_info` attributes are used as referrals, so that a file server may indicate to its client that the file name tree beneath a given name in the server is not present on itself but is represented by a file system in some other set of servers. The mechanism is general, allowing servers to describe any file system as being reachable by requests to any of a set of servers. Thus, starting with a single NFSv4 server, using these referrals, an NFSv4 client could see a large namespace associated with a collection of interrelated NFSv4 file servers. An organization could use this capability to construct a uniform file namespace for itself.

An organization might wish to publish the starting point for this namespace to its clients. In many cases, the organization will want to publish this starting point to a broader set of possible clients. At the same time, it is useful to require that clients know only the smallest amount of information in order to locate the appropriate namespace. Also, that required information should be constant through the life of an organization if the clients are not to require reconfiguration as administrative events change, for instance, a server's name or address.

2. Requirements Notation

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC2119].

3. Use of the SRV Resource Record in DNS

Providing an organization's published file system namespace is a service, and the DNS [RFC1034][RFC1035] provides methods for discovery of that service. This standard defines a mapping from a DNS name to the NFS file system(s) providing the root of the file system namespace associated with that DNS name; such file systems are called "domain root" file systems. From such file systems, like other NFS file systems, an NFS client can use the standard NFS mechanisms to navigate the rest of the NFS file servers that make up the file system namespace for the given domain.

Such domain root file systems are mounted at a conventional point in the NFS client namespace. The mechanism results in a uniform cross-organizational file namespace, similar to that seen in both AFS [AFS][RFC5864] and Distributed Computing Environment / Distributed File System (DCE/DFS) [DFS]. An NFS client need know only the domain name for an organization in order to locate the file namespace published by that organization.

The DNS SRV RR type [RFC2782] is used to locate domain root file servers. The format of the DNS SRV record is as follows:

`_Service._Proto.Name TTL Class SRV Priority Weight Port Target`

The Service name used is `"_nfs-domainroot"`, in conformance with RFC 6335 [RFC6335]. The Protocol name used is `"_tcp"`, for NFS service over TCP. Future NFS services using other protocols MUST use another protocol name. The `"_udp"` label MUST NOT be used to imply use of UDP with NFSv4, as neither RFC 3530 [RFC3530] nor RFC 5661 [RFC5661] defines NFSv4 over UDP. The Target fields give the domain names of the NFS servers that export file systems for the domain's root. An NFS client may then interpret any of the exported root file systems as the root of the file system published by the organization with the given domain name.

The domain root service is not useful for NFS versions prior to version 4, as the `fs_locations` attribute was introduced only in NFSv4 (as described in Section 1).

In order to allow the NFSv4 servers as given to export a variety of file systems, those file servers MUST export the given domain's root file systems at `"/.domainroot/{Name}"` within their pseudo-file systems, where the `"{Name}"` is the name of the organization as used in the SRV RR.

As an example, suppose a client wished to locate the root of the file system published by organization `example.net`. The DNS servers for the domain would publish records like

```
$ORIGIN example.net.  
_nfs-domainroot._tcp IN SRV 0 0 2049 nfs1tr.example.net.  
_nfs-domainroot._tcp IN SRV 1 0 18204 nfs2ex.example.net.
```

The resulting domain names `nfs1tr.example.net` and `nfs2ex.example.net` indicate NFSv4 file servers that export the root of the published namespace for the `example.net` domain. In accordance with RFC 2782 [RFC2782], these records are to be interpreted using the Priority and Weight field values, selecting an appropriate file server with which to begin a network conversation. The two file servers would export file systems that would be found at `"/.domainroot/example.net"` in their pseudo-file systems, which clients would mount. Clients then carry out subsequent accesses in accordance with the ordinary NFSv4 protocol. The first record uses the port number 2049 assigned to NFS, and another port is specified for the second record; the NFS servers would provide NFS service at their indicated port numbers, and NFS clients would connect to the service via the corresponding port numbers on those indicated servers.

Other file system protocols could make use of the same domain root abstraction, but it is necessary to use different Service names not specified here.

4. Integration with Use of NFS Version 4

NFSv4 clients adhering to this specification implement a special directory, analogous to an Automounter [AMD1][AMD2] directory, the entries in which are domain names that have recently been traversed. When an application attempts to traverse a new name in that special directory, the NFSv4 client consults DNS to obtain the SRV data for the given name, and if successful, it mounts the indicated file system(s) in that name in the special directory. The goal is that NFSv4 applications will be able to look up an organization's domain name in the special directory, and the NFSv4 client will be able to discover the file system that the organization publishes. Entries in the special directory will be domain names, and they will each appear to the application as a directory name pointing to the root directory of the file system published by the organization responsible for that domain name.

As noted in Section 3, the domain root service is not useful for NFS versions prior to version 4.

4.1. Globally Useful Names: Conventional Mount Point

In order for the inter-organizational namespace to function as a global file namespace, the client-side mount point for that namespace must be the same on different clients. Conventionally, on Portable Operating System Interface (POSIX) machines, the name `"/nfs4/"` is used so that names on one machine will be directly usable on any machine. Thus, the `example.net` published file system would be accessible as

`/nfs4/example.net/`

on any POSIX client. Using this convention, `"/nfs4/"` is the name of the special directory that is populated with domain names, leading to file servers and file systems that capture the results of SRV record lookups.

4.2. Mount Options

SRV records are necessarily less complete than the information in the existing NFSv4 attributes `fs_locations` [RFC3530] or `fs_locations_info` [RFC5661]. For the `rootpath` field of `fs_location`, or the `flr_fs_root` field of `fs_locations_info`, NFS servers **MUST** use the `"/.domainroot/{Name}"` string. Thus, the servers listed as targets for the SRV RRs **MUST** export the root of the organization's published file system as the directory `"/.domainroot/{Name}"` (for the given organization Name) in their exported NFS namespaces. For example, for organization `example.net`, the directory `"/.domainroot/example.net"` would be used.

Section 11 of the NFSv4.1 document [RFC5661] describes the approach that an NFS client should take to navigate `fs_locations_info` information.

The process of mounting an organization's namespace should permit the use of what is likely to impose the lowest cost on the server. Thus, the NFS client **SHOULD NOT** insist on using a writable copy of the file system if read-only copies exist, or a zero-age copy rather than a copy that may be a little older. The organization's file system representatives can be navigated to provide access to higher-cost properties such as writability or freshness as necessary, but the default use when navigating to the base information for an organization ought to be as low-overhead as possible.

4.3. File System Integration Issues

The result of the DNS search **SHOULD** appear as a (pseudo-)directory in the client namespace. A further refinement is **RECOMMENDED**: that only fully qualified domain names appear as directories. That is, in many environments, DNS names may be abbreviated from their fully qualified form. In such circumstances, multiple names might be given to NFS clients that all resolve to the same DNS SRV RRs. The abbreviated form **SHOULD** be represented in the client's namespace cache as a symbolic link, pointing to the fully qualified name. This will allow pathnames obtained with, say, `getcwd()` to include the DNS name that is most likely to be usable outside the scope of any particular DNS abbreviation convention.

4.4. Multicast DNS

Location of the NFS domain root by this SRV record is intended to be performed with unicast by using the ordinary DNS [RFC1034][RFC1035] protocol.

This document does not define the use of this DNS SRV record format in conjunction with Multicast DNS (mDNS). While mDNS could be used to locate a local domain root via these SRV records, no other domain's root could be discovered. This means that mDNS has too little value to use in locating NFSv4 domain roots.

5. Where Is This Integration Carried Out?

The NFS client is responsible for interpreting SRV records. Using something like Automounter [AMD1] [AMD2] technology, the client interprets names under a particular directory, by first discovering the appropriate file system to mount and then mounting it in the specified place in the client namespace before returning control to the application doing a lookup. The result of the DNS lookup should be cached (obeying Time to Live (TTL)) so that the result could be returned more quickly the next time.

6. Security Considerations

This functionality introduces a new reliance of NFSv4 on the integrity of DNS. Forged SRV records in DNS could cause the NFSv4 client to connect to the file servers of an attacker, rather than the legitimate file servers of an organization. This is similar to attacks that can be made on the base NFSv4 protocol, if server names are given in `fs_location` attributes: the client can be made to connect to the file servers of an attacker, not the file servers intended to be the target for the `fs_location` attributes.

If DNS Security Extensions (DNSSEC) [RFC4033] is available, it SHOULD be used to avoid both such attacks. Domain-based service principal names are an additional mechanism that also apply in this case, and it would be prudent to use them. They provide a mapping from the domain name that the user specified to names of security principals used on the NFSv4 servers that are indicated as the targets in the SRV records (as providing file service for the root file systems).

With domain-based service principal names, the idea is that one wants to authenticate {nfs, domainname, host.fqdn}, not simply {nfs, host.fqdn}, when the server is a domain's root file server obtained through a DNS SRV RR lookup that may or may not have been secure.

The domain administrator can thus ensure that only domain root NFSv4 servers have credentials for such domain-based service principal names.

Domain-based service principal names are defined in RFCs 5178 [RFC5178] and 5179 [RFC5179]. To make use of RFC 5178's domain-based names, the syntax for "domain-based-name" MUST be used with a service of "nfs", a domain matching the name of the organization whose root file system is being sought, and a hostname given in the target of the DNS SRV RR. Thus, in the example above, two file servers (nfs1tr.example.net and nfs2ex.example.net) are located as hosting the root file system for the organization example.net. To communicate with, for instance, the second of the given file servers, Generic Security Service Application Program Interface (GSS-API) is used with the name-type of GSS_C_NT_DOMAINBASED_SERVICE defined in RFC 5178 and with a symbolic name of

nfs@example.net@nfs2ex.example.net

in order to verify that the named server (nfs2ex.example.net) is authorized to provide the root file system for the example.net organization.

NFSv4 itself contains a facility for the negotiation of security mechanisms to be used between NFS clients and NFS servers. Section 3.3 of RFC 3530 [RFC3530] and Section 2.6 of RFC 5661 [RFC5661] both describe how security mechanisms are to be negotiated. As such, there is no need for this document to describe how that negotiation is to be carried out when the NFS client contacts the NFS server for the specified domain root file system(s).

Using SRV records to advertise the locations of NFS servers may expose those NFS servers to attacks. Organizations should carefully consider whether they wish their DNS servers to respond differentially to different DNS clients, perhaps exposing their SRV records to only those DNS requests that originate within a given perimeter, in order to reduce this exposure.

7. IANA Considerations

IANA has assigned a new Service name without an associated port number (as defined in RFC 6335 [RFC6335]) for TCP. For this new Service, the Reference is this document.

Service name: nfs-domainroot

Transport Protocol(s) TCP

Assignee (REQUIRED) IESG (iesg@ietf.org)

Contact (REQUIRED) IETF Chair (chair@ietf.org)

Description (REQUIRED) NFS service for the domain root, the root of an organization's published file namespace.

Reference (REQUIRED) This document

Port Number (OPTIONAL)

Service Code (REQUIRED for DCCP only)

Known Unauthorized Uses (OPTIONAL)

Assignment Notes (OPTIONAL)

8. References

8.1. Normative References

- [RFC1034] Mockapetris, P., "Domain names - concepts and facilities", STD 13, RFC 1034, November 1987.
- [RFC1035] Mockapetris, P., "Domain names - implementation and specification", STD 13, RFC 1035, November 1987.
- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, March 1997.
- [RFC2782] Gulbrandsen, A., Vixie, P., and L. Esibov, "A DNS RR for specifying the location of services (DNS SRV)", RFC 2782, February 2000.
- [RFC3530] Shepler, S., Callaghan, B., Robinson, D., Thurlow, R., Beame, C., Eisler, M., and D. Noveck, "Network File System (NFS) version 4 Protocol", RFC 3530, April 2003.
- [RFC4033] Arends, R., Austein, R., Larson, M., Massey, D., and S. Rose, "DNS Security Introduction and Requirements", RFC 4033, March 2005.
- [RFC5178] Williams, N. and A. Melnikov, "Generic Security Service Application Program Interface (GSS-API) Internationalization and Domain-Based Service Names and Name Type", RFC 5178, May 2008.

- [RFC5179] Williams, N., "Generic Security Service Application Program Interface (GSS-API) Domain-Based Service Names Mapping for the Kerberos V GSS Mechanism", RFC 5179, May 2008.
- [RFC5661] Shepler, S., Ed., Eisler, M., Ed., and D. Noveck, Ed., "Network File System (NFS) Version 4 Minor Version 1 Protocol", RFC 5661, January 2010.
- [RFC5864] Allbery, R., "DNS SRV Resource Records for AFS", RFC 5864, April 2010.
- [RFC6335] Cotton, M., Eggert, L., Touch, J., Westerlund, M., and S. Cheshire, "Internet Assigned Numbers Authority (IANA) Procedures for the Management of the Service Name and Transport Protocol Port Number Registry", BCP 165, RFC 6335, August 2011.

8.2. Informative References

- [AFS] Howard, J., "An Overview of the Andrew File System", Proc. USENIX Winter Tech. Conf. Dallas, February 1988.
- [AMD1] Pendry, J. and N. Williams, "Amd: The 4.4 BSD Automounter Reference Manual", March 1991, <<http://docs.freebsd.org/info/amdref/amdref.pdf>>.
- [AMD2] Crosby, M., "AMD--AutoMount Daemon", Linux Journal, 35es Article 4, March 1997.
- [DFS] Kazar, M., Leverett, B., Anderson, O., Apostolides, V., Bottos, B., Chutani, S., Everhart, C., Mason, W., Tu, S., and E. Zayas, "DEcorum File System Architectural Overview", Proc. USENIX Summer Conf. Anaheim, Calif., June 1990.

Authors' Addresses

Craig Everhart
NetApp
800 Cranberry Woods Drive, Ste. 300
Cranberry Township, PA 16066
USA

Phone: +1 724 741 5101
EMail: everhart@netapp.com

W.A. (Andy) Adamson
NetApp
495 East Java Drive
Sunnyvale, CA 94089
USA

Phone: +1 734 665 1204
EMail: andros@netapp.com

Jiaying Zhang
Google
604 Arizona Avenue
Santa Monica, CA 90401
USA

Phone: +1 310 309 6884
EMail: jiayingz@google.com