

Internet Engineering Task Force (IETF)
Request for Comments: 6905
Category: Informational
ISSN: 2070-1721

T. Senevirathne
Cisco
D. Bond
IBM
S. Aldrin
Y. Li
Huawei
R. Watve
Cisco
March 2013

Requirements for Operations, Administration, and Maintenance (OAM) in Transparent Interconnection of Lots of Links (TRILL)

Abstract

Operations, Administration, and Maintenance (OAM) is a general term used to identify functions and toolsets to troubleshoot and monitor networks. This document presents OAM requirements applicable to the Transparent Interconnection of Lots of Links (TRILL).

Status of This Memo

This document is not an Internet Standards Track specification; it is published for informational purposes.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Not all documents approved by the IESG are a candidate for any level of Internet Standard; see Section 2 of RFC 5741.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <http://www.rfc-editor.org/info/rfc6905>.

Copyright Notice

Copyright (c) 2013 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<http://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	3
1.1. Scope	3
2. Conventions Used in This Document	3
3. Terminology	3
4. OAM Requirements	4
4.1. Data Plane	4
4.2. Connectivity Verification	5
4.2.1. Unicast	5
4.2.2. Distribution Trees	5
4.3. Continuity Check	5
4.4. Path Tracing	6
4.5. General Requirements	6
4.6. Performance Monitoring	7
4.6.1. Packet Loss	7
4.6.2. Packet Delay	7
4.7. ECMP Utilization	8
4.8. Security and Operational Considerations	8
4.9. Fault Indications	8
4.10. Defect Indications	9
4.11. Live Traffic Monitoring	9
5. Security Considerations	9
6. References	9
6.1. Normative References	9
6.2. Informative References	10
7. Acknowledgments	11
8. Contributors	11

1. Introduction

The Operations, Administration, and Maintenance (OAM) generally covers various production aspects of a network. In this document, we use the term OAM as defined in [RFC6291].

The success of network operations depends on the ability to proactively monitor it for faults, performance, etc., as well as the ability to efficiently and quickly troubleshoot defects and failures. A well-defined OAM toolset is a vital requirement for wider adoption of Transparent Interconnection of Lots of Links (TRILL) as the next generation data-forwarding technology in larger networks such as data centers.

In this document, we define the requirements for TRILL OAM. It is assumed that the readers are familiar with the OAM concepts and terminologies defined in other OAM standards such as [8021ag], [RFC5860], and [RFC4377]. This document does not attempt to redefine the terms and concepts specified elsewhere.

1.1. Scope

The scope of this document is OAM between Routing Bridges (RBridges) of a TRILL campus over links selected by TRILL routing.

2. Conventions Used in This Document

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in RFC 2119 [RFC2119]. Although this document is not a protocol specification, the use of this language clarifies the instructions to protocol designers producing solutions that satisfy the requirements set out in this document.

3. Terminology

Section: This term refers to a segment of a path between any two given RBridges. As an example, consider the case where RB1 is connected to RBx via RB2, RB3, and RB4. The segment between RB2 to RB4 is referred to as a section of the path RB1 to RBx. More details of this definition can be found in [RFC5960].

Flow: This term indicates a set of packets that share the same path and per-hop behavior (such as priority). A flow is typically identified by a portion of the inner payload that affects the hop-by-hop forwarding decisions. This may contain Layer 2 through Layer 4 information.

All Selectable Least-Cost Paths: This term refers to a subset of all potentially available least-cost paths to a specified destination RBridge that are available (and usable) for forwarding of frames. It is important to note that in practice, due to limitations in implementations, not all available least-cost paths may be selectable for forwarding.

Connectivity: This term indicates reachability between an arbitrary RBridge RB1 and any other RBridge RB2. The specific path can be either explicit (i.e., associated with a specific flow) or unspecified. Unspecified means that messages used for connectivity verification take whatever path the RBs happen to select. Please refer to [OAMOVER] for details.

Continuity Verification: This term refers to proactive verification of liveness between two RBridges at periodic intervals and the generation of explicit notification when connectivity failures occur. Please refer to [OAMOVER] for details.

Fault: This term refers to an inability to perform a required action, e.g., an unsuccessful attempt to deliver a packet. Please refer to [TERMTTP] for definition.

Defect: This term refers to an interruption in the normal operation, such that over a period of time no packets are delivered successfully. Please refer to [TERMTTP] for definition.

Failure: This term refers to the termination of the required function over a longer period of time. Persistence of a defect for a period of time is interpreted as a failure. Please refer to [TERMTTP] for definition.

Simulated Flow: This term refers to a sequence of OAM-generated packets designed to follow a specific path. The fields of the packets in the simulated flow may or may not be identical to the fields of data packets of an actual flow being simulated. However, the purpose of the simulated flow is to have OAM packets of the simulated flow follow a specific path.

4. OAM Requirements

4.1. Data Plane

OAM frames, utilized for connectivity verification, continuity checks, performance measurements, etc., will by default take whatever path TRILL chooses based on the current topology and per-hop equal-cost path choices. In some cases, it may be required that the OAM

frames utilize specific paths. Thus, it **MUST** be possible to arrange that OAM frames follow the path taken by a specific flow.

R Bridges **MUST** have the ability to identify frames that require OAM processing.

TRILL OAM frames **MUST** remain within a TRILL campus and **MUST NOT** be egressed from a TRILL network as native frames.

OAM **MUST** have the ability to include all Ethernet traffic types carried by TRILL.

4.2. Connectivity Verification

4.2.1. Unicast

From an arbitrary R Bridge RB1, OAM **MUST** have the ability to verify connectivity to any other R Bridge RB2.

From an arbitrary R Bridge RB1, OAM **MUST** have the ability to verify connectivity to any other R Bridge RB2 for a specific flow via the path associated with the specified flow.

4.2.2. Distribution Trees

OAM **MUST** have the ability to verify connectivity from an arbitrary R Bridge RB1 to either a specific set of R Bridges or all member R Bridges, for a specified distribution tree. This functionality is referred to as verification of the unpruned distribution tree.

OAM **MUST** have the ability to verify connectivity from an arbitrary R Bridge RB1 to either a specific set of R Bridges or all member R Bridges, for a specified distribution tree and for a specified flow. This functionality is referred to as verification of the pruned tree.

4.3. Continuity Check

OAM **MUST** provide functions that allow any arbitrary R Bridge RB1 to perform a Continuity Check to any other R Bridge.

OAM **MUST** provide functions that allow any arbitrary R Bridge RB1 to perform a Continuity Check to any other R Bridge using a path associated with a specified flow.

OAM **SHOULD** provide functions that allow any arbitrary R Bridge to perform a Continuity Check to any other R Bridge over any section of any selectable least-cost path.

OAM SHOULD provide the ability to perform a Continuity Check on sections of any selectable path within the network.

OAM SHOULD provide the ability to perform a multicast Continuity Check for specified distribution tree(s), as well as specified combinations of distribution trees and flows. The former is referred to as an unpruned multi-destination tree Continuity Check and the latter is referred to as a pruned tree Continuity Check.

4.4. Path Tracing

OAM MUST provide the ability to trace a path between any two RBridges corresponding to a specified unicast flow.

OAM SHOULD provide the ability to trace all selectable least-cost paths between any two RBridges.

OAM SHOULD provide functionality to trace all branches of a specified distribution tree (unpruned tree).

OAM SHOULD provide functionality to trace all branches of a specified distribution tree for a specified flow (pruned tree).

4.5. General Requirements

OAM MUST provide the ability to initiate and maintain multiple concurrent sessions for multiple OAM functions between any arbitrary RBridge RB1 to any other RBridge. In general, multiple OAM operations will run concurrently. For example, proactive continuity checks may take place between RB1 and RB2 at the same time that an operator decides to test connectivity between the same two RBs. Multiple OAM functions and instances of those functions MUST be able to run concurrently without interfering with each other.

OAM MUST provide a single OAM framework for all TRILL OAM functions within the scope of this document.

OAM, as practical and as possible, SHOULD reuse functional, operational, and semantic elements of existing OAM standards.

OAM MUST maintain related error and operational counters. Such counters MUST be accessible via network management applications (e.g., SNMP).

OAM functions related to continuity and connectivity checks MUST be able to be invoked either proactively or on demand.

OAM MAY be required to provide the ability to specify a desired response mode for a specific OAM message. The desired response mode can be in-band, out-of-band, or none.

The OAM Framework MUST be extensible to include new functionality. For example, the solution needs to include a version number to differentiate older and newer implementations and TLV structures for flexibility to include new information elements.

OAM MAY provide methods to verify control-plane and forwarding-plane alignments.

OAM SHOULD leverage existing OAM technologies, where practical.

4.6. Performance Monitoring

4.6.1. Packet Loss

In this document, the term "packet loss" is used as defined in Section 2.4 of [RFC2680].

OAM SHOULD provide the ability to measure packet loss statistics for a flow from any arbitrary RBridge RB1 to any other RBridge.

OAM SHOULD provide the ability to measure packet loss statistics over a section for a flow between any arbitrary RBridge RB1 to any other RBridge.

OAM SHOULD provide the ability to measure packet loss statistics between any two RBridges over all least-cost paths.

An RBridge SHOULD be able to perform the above packet loss measurement functions either proactively or on demand.

4.6.2. Packet Delay

There are two types of packet delays -- one-way delay and two-way delay (Round-Trip Delay).

One-way delay is defined in [RFC2679] as the time elapsed from the start of transmission of the first bit of a packet by an RBridge until the reception of the last bit of the packet by the destination RBridge.

Two-way delay is also referred to as Round-Trip Delay and is defined similar to [RFC2681]; i.e., the time elapsed from the start of transmission of the first bit of a packet from RB1, receipt of the packet at RB2, RB2 sending a response packet back to RB1, and RB1 receiving the last bit of that response packet.

OAM SHOULD provide functions to measure two-way delay between two RBridges.

OAM MAY provide functions to measure one-way delay between two RBridges for a specified flow.

OAM MAY provide functions to measure one-way delay between two RBridges for a specified flow over a specific section.

4.7. ECMP Utilization

OAM MAY provide functionality to monitor the effectiveness of per-hop Equal-Cost Multipath (ECMP) hashing. For example, individual RBridges could maintain counters that show how packets are being distributed across equal-cost next hops for a specified destination RBridge or RBridges as a result of ECMP hashing.

4.8. Security and Operational Considerations

Methods MUST be provided to protect against exploitation of OAM framework for security and denial-of-service attacks.

Methods MUST be provided to prevent OAM messages from causing congestion in the networks. Periodically generated messages with high frequencies may lead to congestion, hence methods such as shaping or rate limiting SHOULD be utilized.

Certain OAM functions may be utilized to gather operational information such as topology of the network. Methods MUST be provided to prevent unauthorized users accessing OAM functions to gather critical and sensitive information of the network.

OAM packets MUST be limited to within the TRILL campus, and the implementation MUST provide methods to prevent leaking of OAM packets out of the TRILL campus. Additionally, methods MUST be provided to prevent accepting OAM packets from outside the TRILL campus.

4.9. Fault Indications

OAM MUST provide a Fault Indication framework to notify the packet's ingress RBridge or other interested parties (such as syslog servers) about faults.

OAM MUST provide functions to selectively enable or disable different types of Fault Indications.

4.10. Defect Indications

OAM SHOULD provide a framework for Defect Detection and Indication.

OAM Defect Detection and Indication Framework SHOULD provide methods to selectively enable or disable Defect Detection per defect type.

OAM Defect Detection and Indication Framework SHOULD provide methods to configure Defect Detection thresholds per different types of defects.

OAM Defect Detection and Indication Framework SHOULD provide methods to log defect indications to a locally defined archive (such as log buffer) or Simple Network Management Protocol (SNMP) traps.

OAM Defect Detection and Indication Framework SHOULD provide a Remote Defect Indication framework that facilitates notifying the originator/owner of the flow experiencing the defect, which is the ingress RBridge.

Remote Defect Indication MAY be either in-band or out-of-band.

4.11. Live Traffic Monitoring

OAM implementations MAY provide methods to utilize live traffic for troubleshooting and performance monitoring.

5. Security Considerations

Security requirements are specified in Section 4.8. For general TRILL security considerations, please refer to [RFC6325].

6. References

6.1. Normative References

[RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, March 1997.

[RFC6291] Andersson, L., van Helvoort, H., Bonica, R., Romascanu, D., and S. Mansfield, "Guidelines for the Use of the "OAM" Acronym in the IETF", BCP 161, RFC 6291, June 2011.

6.2. Informative References

- [8021ag] IEEE, "Virtual Bridged Local Area Networks Amendment 5: Connectivity Fault Management", IEEE Std 802.1ag-2007, 2007.
- [OAMOVER] Mizrahi, T., Sprecher, N., Bellagamba, E., Y. Weingarten, "An Overview of Operations, Administration, and Maintenance (OAM) Mechanisms", Work in Progress, January 2013.
- [RFC2679] Almes, G., Kalidindi, S., and M. Zekauskas, "A One-way Delay Metric for IPPM", RFC 2679, September 1999.
- [RFC2680] Almes, G., Kalidindi, S., and M. Zekauskas, "A One-way Packet Loss Metric for IPPM", RFC 2680, September 1999.
- [RFC2681] Almes, G., Kalidindi, S., and M. Zekauskas, "A Round-trip Delay Metric for IPPM", RFC 2681, September 1999.
- [RFC4377] Nadeau, T., Morrow, M., Swallow, G., Allan, D., and S. Matsushima, "Operations and Management (OAM) Requirements for Multi-Protocol Label Switched (MPLS) Networks", RFC 4377, February 2006.
- [RFC5860] Vigoureux, M., Ed., Ward, D., Ed., and M. Betts, Ed., "Requirements for Operations, Administration, and Maintenance (OAM) in MPLS Transport Networks", RFC 5860, May 2010.
- [RFC5960] Frost, D., Ed., Bryant, S., Ed., and M. Bocci, Ed., "MPLS Transport Profile Data Plane Architecture", RFC 5960, August 2010.
- [RFC6325] Perlman, R., Eastlake 3rd, D., Dutt, D., Gai, S., and A. Ghanwani, "Routing Bridges (RBriges): Base Protocol Specification", RFC 6325, July 2011.
- [TERMTP] van Helvoort, H., Ed., Andersson, L., Ed., and N. Sprecher, Ed., "A Thesaurus for the Terminology used in Multiprotocol Label Switching Transport Profile (MPLS-TP) drafts/RFCs and ITU-T' Transport Network Recommendations", Work in Progress, February 2013.

7. Acknowledgments

Special acknowledgments to IEEE 802.1 chair, Tony Jeffree, for allowing us to solicit comments from IEEE 802.1 group. Also recognized are the comments received from the IEEE group, IESG, Stewart Bryant, Ralph Droms, Adrian Farrel, Benoit Claise, Ayal Lior, and others.

8. Contributors

Thomas Narten
IBM Corporation
3039 Cornwallis Avenue,
PO Box 12195
Research Triangle Park, NC 27709
USA

EMail: narten@us.ibm.com

Donald Eastlake
Huawei Technologies
155 Beaver Street,
Milford, MA 01757
USA

EMail: d3e3e3@gmail.com

Anoop Ghanwani
Dell
350 Holger Way
San Jose, CA 95134
USA

Phone: +1-408-571-3500
EMail: Anoop@alumni.duke.edu

Jon Hudson
Brocade
120 Holger Way
San Jose, CA 95134
USA

EMail: jon.hudson@gmail.com

Naveen Nimmu
Broadcom
9th Floor, Building no 9, Raheja Mind space
Hi-Tec City, Madhapur,
Hyderabad - 500 081
India

Phone: +1-408-218-8893
EMail: naveen@broadcom.com

Radia Perlman
Intel Labs
2700 156th Ave NE, Suite 300,
Bellevue, WA 98007
USA

Phone: +1-425-881-4824
EMail: radia.perlman@intel.com

Tal Mizrahi
Marvell
6 Hamada St.
Yokneam, 20692
Israel

EMail: talmi@marvell.com

Authors' Addresses

Tissa Senevirathne
Cisco Systems
375 East Tasman Drive
San Jose, CA 95134
USA

Phone: +1-408-853-2291
EMail: tsenevir@cisco.com

David Bond
IBM
4400 North 1st Street
San Jose, CA 95134
USA

Phone: +1-603-339-7575
EMail: mokon@mokon.net

Sam Aldrin
Huawei Technologies
2330 Central Express Way
Santa Clara, CA 95951
USA

EMail: aldrin.ietf@gmail.com

Yizhou Li
Huawei Technologies
101 Software Avenue,
Nanjing 210012
China

Phone: +86-25-56625375
EMail: liyizhou@huawei.com

Rohit Watve
Cisco Systems
375 East Tasman Drive
San Jose, CA 95134
USA

Phone: +1-408-424-2091
EMail: rwatve@cisco.com