

Internet Engineering Task Force (IETF)
Request for Comments: 7496
Category: Standards Track
ISSN: 2070-1721

M. Tuexen
Muenster Univ. of Appl. Sciences
R. Seggelmann
Metafinanz Informationssysteme GmbH
R. Stewart
Netflix, Inc.
S. Loreto
Ericsson
April 2015

Additional Policies for the Partially Reliable Stream Control Transmission Protocol Extension

Abstract

This document defines two additional policies for the Partially Reliable Stream Control Transmission Protocol (PR-SCTP) extension. These policies allow limitation of the number of retransmissions and prioritization of user messages for more efficient usage of the send buffer.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 5741.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <http://www.rfc-editor.org/info/rfc7496>.

Copyright Notice

Copyright (c) 2015 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<http://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	3
2. Conventions	3
3. Additional PR-SCTP Policies	3
3.1. Limited Retransmissions Policy	3
3.2. Priority Policy	4
4. Socket API Considerations	4
4.1. Data Types	5
4.2. Support for Added PR-SCTP Policies	5
4.3. Socket Option for Getting the Stream-Specific PR-SCTP Status (SCTP_PR_STREAM_STATUS)	6
4.4. Socket Option for Getting the Association-Specific PR-SCTP Status (SCTP_PR_ASSOC_STATUS)	7
4.5. Socket Option for Getting and Setting the PR-SCTP Support (SCTP_PR_SUPPORTED)	8
5. Security Considerations	9
6. References	9
6.1. Normative References	9
6.2. Informative References	9
Acknowledgments	10
Authors' Addresses	11

1. Introduction

The Partially Reliable SCTP (PR-SCTP) extension defined in [RFC3758] provides a generic method for senders to abandon user messages. The decision to abandon a user message is sender side only, and the exact condition is called a "PR-SCTP policy" ([RFC3758] refers to them as "PR-SCTP Services"). [RFC3758] also defines one particular PR-SCTP policy, called "Timed Reliability". This allows the sender to specify a timeout for a user message after which the SCTP stack abandons the user message.

This document specifies the following two additional PR-SCTP policies:

Limited Retransmission Policy: Allows limitation of the number of retransmissions.

Priority Policy: Allows removal of lower-priority messages if space for higher-priority messages is needed in the send buffer.

2. Conventions

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC2119].

3. Additional PR-SCTP Policies

This section defines two new PR-SCTP policies, one in each subsection.

Please note that it is REQUIRED to implement [RFC3758], if you want to implement these additional policies. However, these additional policies are OPTIONAL when implementing [RFC3758].

3.1. Limited Retransmissions Policy

Using the Limited Retransmission Policy allows the sender of a user message to specify an upper limit for the number of retransmissions for each DATA chunk of the given user messages. The sender MUST abandon a user message if the number of retransmissions of any of the DATA chunks of the user message would exceed the provided limit. The sender MUST perform all other actions required for processing the retransmission event, such as adapting the congestion window and the retransmission timeout. Please note that the number of retransmissions includes both fast and timer-based retransmissions.

The sender MAY limit the number of retransmissions to 0. This will result in abandoning the message when it would get retransmitted for the first time. The use of this setting provides a service similar to UDP, which also does not perform any retransmissions.

Please note that using this policy does not affect the handling of the thresholds 'Association.Max.Retrans' and 'Path.Max.Retrans' as specified in Section 8 of [RFC4960].

The WebRTC protocol stack (see [DATA-CHAN]) is an example of where the Limited Retransmissions Policy is used.

3.2. Priority Policy

Using the Priority Policy allows the sender of a user message to specify a priority. When storing a user message in the send buffer while there is not enough available space, the SCTP stack at the sender side MAY abandon other user message(s) of the same SCTP association (with the same or a different stream) with a priority lower than the provided one. User messages sent reliably are considered to have a priority higher than all messages sent with the Priority Policy. The algorithm for selecting the message(s) being abandoned is implementation specific.

After lower-priority messages have been abandoned, high-priority messages can be transferred without the send call blocking (if used in blocking mode) or the send call failing (if used in non-blocking mode).

The IP Flow Information Export (IPFIX) protocol stack (see [RFC7011]) is an example of where the Priority Policy can be used. Template records would be sent with full reliability, while flow records related to billing, security, and other monitoring would be sent using the Priority Policy with varying priority. The priority of security-related flow records would be set higher than the priority of monitoring-related flow records.

4. Socket API Considerations

This section describes how the socket API defined in [RFC6458] is extended to support the newly defined PR-SCTP policies, to provide some statistical information, and to control the negotiation of the PR-SCTP extension during the SCTP association setup.

Please note that this section is informational only.

4.1. Data Types

This section uses data types from [IEEE.1003-1G.1997]: `uintN_t` means an unsigned integer of exactly N bits (e.g., `uint16_t`). This is the same as in [RFC6458].

4.2. Support for Added PR-SCTP Policies

As defined in [RFC6458], the PR-SCTP policy is specified and configured by using the following `sctp_prinfo` structure:

```
struct sctp_prinfo {
    uint16_t pr_policy;
    uint32_t pr_value;
};
```

When the Limited Retransmission Policy described in Section 3.1 is used, `pr_policy` has the value `SCTP_PR_SCTP_RTX` and the number of retransmissions is given in `pr_value`.

When using the Priority Policy described in Section 3.2, `pr_policy` has the value `SCTP_PR_SCTP_PRIO`. The priority is given in `pr_value`. The value of zero is the highest priority, and larger numbers in `pr_value` denote lower priorities.

The following table summarizes the possible parameter settings defined in [RFC6458] and this document:

pr_policy	pr_value	Specification
<code>SCTP_PR_SCTP_NONE</code>	Ignored	[RFC6458]
<code>SCTP_PR_SCTP_TTL</code>	Lifetime in ms	[RFC6458]
<code>SCTP_PR_SCTP_RTX</code>	Number of retransmissions	Section 3.1
<code>SCTP_PR_SCTP_PRIO</code>	Priority	Section 3.2

4.3. Socket Option for Getting the Stream-Specific PR-SCTP Status (SCTP_PR_STREAM_STATUS)

This socket option uses `IPPROTO_SCTP` as its level and `SCTP_PR_STREAM_STATUS` as its name. It can only be used with `getsockopt()` but not with `setsockopt()`. The socket option value uses the following structure:

```
struct sctp_prstatus {
    sctp_assoc_t sprstat_assoc_id;
    uint16_t sprstat_sid;
    uint16_t sprstat_policy;
    uint64_t sprstat_abandoned_unsent;
    uint64_t sprstat_abandoned_sent;
};
```

sprstat_assoc_id: This parameter is ignored for one-to-one style sockets. For one-to-many style sockets, this parameter indicates for which association the user wants the information. It is an error to use `SCTP_{CURRENT|ALL|FUTURE}_ASSOC` in `sprstat_assoc_id`.

sprstat_sid: This parameter indicates for which outgoing SCTP stream the user wants the information.

sprstat_policy: This parameter indicates for which PR-SCTP policy the user wants the information. It is an error to use `SCTP_PR_SCTP_NONE` in `sprstat_policy`. If `SCTP_PR_SCTP_ALL` is used, the counters provided are aggregated over all supported policies.

sprstat_abandoned_unsent: The number of user messages that have been abandoned using the policy specified in `sprstat_policy` on the stream specified in `sprstat_sid` for the association specified by `sprstat_assoc_id`, before any part of the user message could be sent.

sprstat_abandoned_sent: The number of user messages that have been abandoned using the policy specified in `sprstat_policy` on the stream specified in `sprstat_sid` for the association specified by `sprstat_assoc_id`, after a part of the user message has been sent.

There are separate counters for unsent and sent user messages because the `SCTP_SEND_FAILED_EVENT` supports a similar differentiation. Please note that an abandoned large user message requiring SCTP-level fragmentation is reported in the `sprstat_abandoned_sent` counter as soon as at least one fragment of it has been sent. Therefore, each abandoned user message is counted in either `sprstat_abandoned_unsent` or `sprstat_abandoned_sent`.

If more detailed information about abandoned user messages is required, the subscription to the `SCTP_SEND_FAILED_EVENT` is recommended. Please note that some implementations might choose not to support this option, since it increases the resources needed for an outgoing SCTP stream. For the same reasons, some implementations might only support using `SCTP_PR_SCTP_ALL` in `sprstat_policy`.

`sctp_opt_info()` needs to be extended to support `SCTP_PR_STREAM_STATUS`.

4.4. Socket Option for Getting the Association-Specific PR-SCTP Status (`SCTP_PR_ASSOC_STATUS`)

This socket option uses `IPPROTO_SCTP` as its level and `SCTP_PR_ASSOC_STATUS` as its name. It can only be used with `getsockopt()`, but not with `setsockopt()`. The socket option value uses the same structure as described in Section 4.3:

```
struct sctp_prstatus {
    sctp_assoc_t sprstat_assoc_id;
    uint16_t sprstat_sid;
    uint16_t sprstat_policy;
    uint64_t sprstat_abandoned_unsent;
    uint64_t sprstat_abandoned_sent;
};
```

`sprstat_assoc_id`: This parameter is ignored for one-to-one style sockets. For one-to-many style sockets, this parameter indicates for which association the user wants the information. It is an error to use `SCTP_{CURRENT|ALL|FUTURE}_ASSOC` in `sprstat_assoc_id`.

`sprstat_sid`: This parameter is ignored.

`sprstat_policy`: This parameter indicates for which PR-SCTP policy the user wants the information. It is an error to use `SCTP_PR_SCTP_NONE` in `sprstat_policy`. If `SCTP_PR_SCTP_ALL` is used, the counters provided are aggregated over all supported policies.

`sprstat_abandoned_unsent`: The number of user messages that have been abandoned using the policy specified in `sprstat_policy` for the association specified by `sprstat_assoc_id`, before any part of the user message could be sent.

`sprstat_abandoned_sent`: The number of user messages that have been abandoned using the policy specified in `sprstat_policy` for the association specified by `sprstat_assoc_id`, after a part of the user message has been sent.

There are separate counters for unsent and sent user messages because the `SCTP_SEND_FAILED_EVENT` supports a similar differentiation. Please note that an abandoned large user message requiring SCTP-level fragmentation is reported in the `sprstat_abandoned_sent` counter as soon as at least one fragment of it has been sent. Therefore, each abandoned user message is counted in either `sprstat_abandoned_unsent` or `sprstat_abandoned_sent`.

If more detailed information about abandoned user messages is required, the usage of the option described in Section 4.3 or the subscription to the `SCTP_SEND_FAILED_EVENT` is recommended.

`sctp_opt_info()` needs to be extended to support `SCTP_PR_ASSOC_STATUS`.

4.5. Socket Option for Getting and Setting the PR-SCTP Support (`SCTP_PR_SUPPORTED`)

This socket option allows the enabling or disabling of the negotiation of PR-SCTP support for future associations. For existing associations, it allows one to query whether or not PR-SCTP support was negotiated on a particular association.

Whether or not PR-SCTP is enabled by default is implementation specific.

This socket option uses `IPPROTO_SCTP` as its level and `SCTP_PR_SUPPORTED` as its name. It can be used with `getsockopt()` and `setsockopt()`. The socket option value uses the following structure defined in [RFC6458]:

```
struct sctp_assoc_value {
    sctp_assoc_t assoc_id;
    uint32_t assoc_value;
};
```

assoc_id: This parameter is ignored for one-to-one style sockets. For one-to-many style sockets, this parameter indicates upon which association the user is performing an action. The special `sctp_assoc_t` `SCTP_FUTURE_ASSOC` can also be used; it is an error to use `SCTP_{CURRENT|ALL}_ASSOC` in `assoc_id`.

assoc_value: A non-zero value encodes the enabling of PR-SCTP, whereas a value of 0 encodes the disabling of PR-SCTP.

`sctp_opt_info()` needs to be extended to support `SCTP_PR_SUPPORTED`.

5. Security Considerations

This document does not add any security considerations to those given in [RFC4960], [RFC3758], and [RFC6458]. As indicated in the Security Considerations of [RFC3758], transport-layer security in the form of TLS over SCTP (see [RFC3436]) can't be used for PR-SCTP. However, DTLS over SCTP (see [RFC6083]) could be used instead. If DTLS over SCTP as specified in [RFC6083] is used, the Security Considerations of [RFC6083] do apply. It should also be noted that using PR-SCTP for an SCTP association doesn't allow that association to behave more aggressively than an SCTP association not using PR-SCTP.

6. References

6.1. Normative References

- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, March 1997, <<http://www.rfc-editor.org/info/rfc2119>>.
- [RFC3758] Stewart, R., Ramalho, M., Xie, Q., Tuexen, M., and P. Conrad, "Stream Control Transmission Protocol (SCTP) Partial Reliability Extension", RFC 3758, May 2004, <<http://www.rfc-editor.org/info/rfc3758>>.
- [RFC4960] Stewart, R., Ed., "Stream Control Transmission Protocol", RFC 4960, September 2007, <<http://www.rfc-editor.org/info/rfc4960>>.

6.2. Informative References

- [RFC3436] Jungmaier, A., Rescorla, E., and M. Tuexen, "Transport Layer Security over Stream Control Transmission Protocol", RFC 3436, December 2002, <<http://www.rfc-editor.org/info/rfc3436>>.
- [RFC6083] Tuexen, M., Seggelmann, R., and E. Rescorla, "Datagram Transport Layer Security (DTLS) for Stream Control Transmission Protocol (SCTP)", RFC 6083, January 2011, <<http://www.rfc-editor.org/info/rfc6083>>.
- [RFC6458] Stewart, R., Tuexen, M., Poon, K., Lei, P., and V. Yasevich, "Sockets API Extensions for the Stream Control Transmission Protocol (SCTP)", RFC 6458, December 2011, <<http://www.rfc-editor.org/info/rfc6458>>.

[RFC7011] Claise, B., Ed., Trammell, B., Ed., and P. Aitken, "Specification of the IP Flow Information Export (IPFIX) Protocol for the Exchange of Flow Information", STD 77, RFC 7011, September 2013, <<http://www.rfc-editor.org/info/rfc7011>>.

[DATA-CHAN] Jesup, R., Loreto, S., and M. Tuexen, "WebRTC Data Channels", Work in Progress, draft-ietf-rtcweb-data-channel-13, January 2015.

[IEEE.1003-1G.1997] IEEE, "Protocol Independent Interfaces", IEEE Standard 1003.1G, March 1997.

Acknowledgments

The authors wish to thank Benoit Claise, Spencer Dawkins, Gorrry Fairhurst, Stephen Farrell, Barry Leiba, Karen Egede Nielsen, Ka-Cheong Poon, Dan Romascanu, Irene Ruengeler, Jamal Hadi Salim, Joseph Salowey, Brian Trammell, and Vlad Yasevich for their invaluable comments.

Authors' Addresses

Michael Tuexen
Muenster University of Applied Sciences
Stegerwaldstrasse 39
48565 Steinfurt
Germany

EMail: tuexen@fh-muenster.de

Robin Seggelmann
Metafinanz Informationssysteme GmbH
Leopoldstrasse 146
80804 Muenchen
Germany

EMail: rfc@robin-seggelmann.com

Randall R. Stewart
Netflix, Inc.
Chapin, SC 29036
United States

EMail: randall@lakerest.net

Salvatore Loreto
Ericsson
Hirsalantie 11
Jorvas 02420
Finland

EMail: Salvatore.Loreto@ericsson.com