

Internet Engineering Task Force (IETF)
Request for Comments: 7715
Category: Standards Track
ISSN: 2070-1721

IJ. Wijnands, Ed.
K. Raza
Cisco Systems, Inc.
A. Atlas
Juniper Networks, Inc.
J. Tantsura
Ericsson
Q. Zhao
Huawei Technology
January 2016

Multipoint LDP (mLDP) Node Protection

Abstract

This document describes procedures to support node protection for Point-to-Multipoint and Multipoint-to-Multipoint Label Switched Paths (P2MP and MP2MP LSPs) that have been built by the Multipoint Label Distribution Protocol (mLDP). In order to protect a node N, the Point of Local Repair (PLR) Label Switching Router (LSR) of N must learn the Merge Point (MPT) LSR(s) of node N such that traffic can be redirected to them in case node N fails. Redirecting the traffic around the failed node N depends on existing Point-to-Point (P2P) Label Switched Paths (LSPs). The pre-established LSPs originate from the PLR LSR and terminate on the MPT LSRs while bypassing LSR N.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 5741.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <http://www.rfc-editor.org/info/rfc7715>.

Copyright Notice

Copyright (c) 2016 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<http://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1.	Introduction	3
1.1.	Conventions Used in This Document	3
1.2.	Terminology	4
2.	PLR Determination	4
2.1.	Transit Node Procedure	5
2.2.	MP2MP Root Node Procedure	6
2.3.	PLR Information Encoding	7
3.	Using the tLDP Session	9
4.	Link or Node Failure	10
4.1.	Reconvergence after Node or Link Failure	11
4.1.1.	Node Failure	12
4.1.2.	Link Failure	12
4.1.3.	Switching to New Primary Path	12
5.	mLDP Capabilities for Node Protection	13
5.1.	PLR Capability	13
5.2.	MPT Capability	14
5.3.	The Protected LSR	14
5.4.	The Node Protection Capability	15
6.	Security Considerations	16
7.	IANA Considerations	16
8.	References	17
8.1.	Normative References	17
8.2.	Informative References	17
	Acknowledgments	18
	Contributors	18
	Authors' Addresses	19

1. Introduction

This document describes procedures to support node protection for Point-to-Multipoint and Multipoint-to-Multipoint Label Switched Paths (P2MP and MP2MP LSPs) that have been built by the Multipoint Label Distribution Protocol (mLDP) [RFC6388]. In order to protect a node N, the Point of Local Repair (PLR) LSR of N must learn the Merge Point (MPT) LSR(s) of node N such that traffic can be redirected to them in case node N fails. Redirecting the traffic around the failed node N depends on existing P2P LSPs. The pre-established LSPs originate from the PLR LSR and terminate on the MPT LSRs while bypassing LSR N. The procedures to set up these P2P LSPs are outside the scope of this document, but one can imagine using techniques based on the Resource Reservation Protocol for Traffic Engineering (RSVP-TE) [RFC5420] or Label Distribution Protocol (LDP) Loop-Free Alternate (LFA) [RFC5286] to accomplish this.

The solution described in this document notifies the PLR(s) of the MPT LSR(s) via signaling using a Targeted LDP (tLDP) session [RFC7060]. By having a tLDP session with the PLR, no additional procedures need to be defined in order to support Make-Before-Break (MBB), Graceful Restart (GR), and Typed Wildcard Forwarding Equivalence Class (FEC). All this is achieved at the expense of having additional tLDP sessions between each MPT and PLR LSR.

In order to allow a node to be protected against failure, the LSRs providing the PLR and the MPT functionality as well as the protected node MUST support the functionality described in this document. LDP capability negotiation [RFC5561] is used to signal the availability of the functionality between the participating nodes; these nodes MUST support capability negotiation.

1.1. Conventions Used in This Document

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in RFC 2119 [RFC2119].

The term "node" is used to refer to an LSR; "node" and "LSR" are used interchangeably in this document. The terms "PLR" and "MPT" are used as shorthand to refer to "PLR LSR" and "MPT LSR", respectively.

1.2. Terminology

mLDP: Multipoint LDP

PLR: Point of Local Repair

The LSR that redirects the traffic to one or more Merge Point LSRs.

MPT: Merge Point

The LSR that merges the backup LSP with the primary LSP. Note, there can be multiple MPT LSRs for a single MP-LSP node protection.

tLDP: Targeted LDP

MP LSP: Multi-Point LSP

Either a P2MP or MP2MP LSP.

root node:

The root of either a P2MP or MP2MP LSP as defined in [RFC6388].


2. PLR Determination

In order for an MPT to establish a tLDP session with a PLR, it first has to learn the PLR for a particular MP LSP. It is the responsibility of the protected node N to advertise the address of the PLR to the MPT. The PLR address for an MP LSP on node N is the address of the upstream LDP peer, but only when node N is NOT the root node of the MP2MP LSP. If the upstream LDP peer is unable to function as PLR, the procedures in this document do not apply and are out of the scope. If node N is the root node, the procedures are slightly different as described in Section 2.2. The procedures that follow assume that all the participating nodes (N, PLRs, MPTs) are enabled (e.g., by a user configuration) to support and implement the PLR determination feature.

The procedures as documented in this RFC require the protected node to be directly connected to the PLR and MPT nodes. This is because mLDP depends on unicast routing to determine the upstream LSR and unicast routing (by default) only has information about the next hop and not beyond that. Support for non-directly connected PLR and MPT nodes is outside the scope of this document.

2.1. Transit Node Procedure

Below are the procedures for when the protected node is a transit node along the path to the root.



N: The node being protected.

...: Backup LSPs from LSR1 to LSR2 and LSR3.


Figure 1

Node N uses the root address of the MP LSP to determine the upstream LSR for a given MP LSP following the procedures as documented in Section 2.4.1.1 of [RFC6388]. The upstream LSR in Figure 1 is LSR1 because it is the first hop along the shortest path to reach the root address. After determining the upstream LSR, node N (which has the node protection feature enabled) MUST advertise the address of LSR1 as the PLR address to the downstream members of the MP LSP (i.e., LSR2 and LSR3) if the given downstream member has announced support for node protection (see Section 5 regarding capability negotiation). For the format and encoding of PLR address information, see Section 2.3.

Note, in order for the protected traffic to reach nodes LSR2 and LSR3, LSR1 MUST have two unidirectional LSPs to LSR2 and LSR3, bypassing node N. The procedures for setting up these LSPs are outside the scope of this document.

2.2. MP2MP Root Node Procedure

Below are the procedures for when the protected node is the root of an MP2MP LSP. Consider figure 2 below.



N: The MP2MP root node being protected.
 ...: Backup LSPs between LSR1, LSR2, and LSR3.

Figure 2

Assume that LSR1, LSR2, and LSR3 are all members of an MP2MP LSP for which N is the root node. Since N is the root of the MP2MP LSP, there is no upstream LSR and no 'single' PLR LSR for protecting node N. In order to protect node N, all the directly connected members of the MP2MP must participate in protecting node N by acting both as PLR and MPT LSR. An LSR will act as MPT for traffic coming from the other LSR(s) and it will act as PLR for traffic it is sending to the other LSR(s). Since node N knows the members of the MP2MP LSP, it will advertise the member list to its directly connected members, excluding the member it is sending to. For example, node N will advertise list {LSR3,LSR1} to LSR2 excluding LSR2 from it. Instead of advertising a single PLR when node N is not the root, a list of PLRs is advertised using the procedures documented in Section 2.3.

It should be noted that the MP2MP root node protection mechanism doesn't replace the Root Node Redundancy (RNR) procedures as described in Section 7 of [RFC6388]. The node protection procedures in this document will help in restoring traffic for the existing MP2MP LSPs after node failure, but a new root node has to be elected eventually in order to allow new MP2MP LSPs to be created.

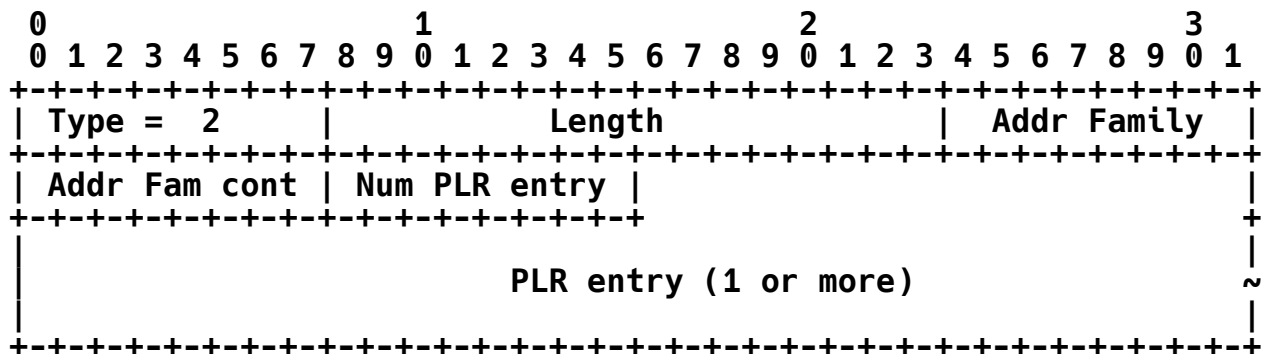
Note, in order for the protected traffic to be exchanged between nodes LSR1, LSR2, and LSR3, bidirectional LSPs have to exist between the LSRs, bypassing node N. The procedures for setting up these LSPs are outside the scope of this document.

2.3. PLR Information Encoding

The upstream LSR address is conveyed via an LDP Notification message with an MP Status TLV, where the MP Status TLV contains a new "PLR Status Value Element" that specifies the address of the PLR.

The new "PLR Status Value Element" is encoded as described below.

PLR Status Element:



Where


Type: PLR Status Value Element (Type 2).

Length: The Length field is an unsigned integer that encodes the length of the Status Value following the Length field. The encoded Length varies based on the Addr Family and the number of PLR entries.

Addr Family: Two-octet quantity containing a value from IANA's "Address Family Numbers" registry [AFI] that encodes the address family for the PLR address encoded in the PLR entry.

Num PLR entry: Element as an unsigned integer followed by the number of "PLR entry" fields in the format specified below.

The format of a "PLR Entry" is as follows:



Where

A bit: 0 = Withdraw, 1 = Add.

Reserved: 15 bits; MUST be zero on transmit and ignored on receipt.

PLR address: PLR address encoded according to the Address Family field encoded in the PLR Status Value Element. Note that the length of the PLR address field is specific to the Address Family that is encoded.

The size of a "PLR Entry" is the 2 octets ("A bit + Reserved") + PLR address length. The length of the PLR address is dependent on the Address Family as encoded in the PLR Status Value Element. The size of a "PLR entry" is 6 octets and 18 octets, respectively, for an IPv4 PLR address and an IPv6 PLR address.

If the PLR address on N changes for a given MP LSP, N needs to trigger a new PLR Status to update the MPT(s). Node N can advertise or withdraw a given PLR from its PLR set by setting the A bit to 1 or 0 respectively in the corresponding PLR entry. Removing a PLR address is likely due to a link failure; see the procedures as documented in Section 4.1. To remove all PLR addresses belonging to the encoded Address Family, an LSR N MUST encode a PLR Status Value Element with no PLR entry and the "Num PLR entry" field MUST be set to zero.

Both the PLR Status and an MP FEC TLV [RFC5036] MUST be included in the LDP Notification message so that a receiver is able to associate the PLR Status with the MP LSP.

3. Using the tLDP Session

The receipt of a PLR MP Status (with PLR addresses) for an MP LSP on a receiving LSR makes it an MPT for node protection. If not already established, the MPT LSR MUST establish a tLDP session with all of the learned PLR addresses using the procedures as documented in [RFC7060].

Using Figure 1 as the reference topology, let us assume that both LSR2 and LSR3 are MPTs and have established a tLDP session with the PLR being LSR1. Assume that both LSR2 and LSR3 have a FEC <R,X> with an upstream LSR N and label Ln assigned to FEC towards N. The MPTs will create a secondary upstream LSR for the FEC <R,X> (using the received PLR address) and assign label Lpx to it. The MPTs will do that for each PLR address that was learned for the MP LSP. In this example, the MPTs will have a FEC <R,X> with two local labels associated with it. Label Ln that was assigned to N using the normal mLDP procedures, and Label Lpx that was assigned to PLR (LSR1) for the purpose of node protection. Note, when the protected node is an MP2MP root node, there will be an upstream LSR for each PLR address that was advertised along with a unique Label Lpx.

The receipt of a FEC Label Mapping alone over the tLDP session from MPT on a PLR conveys the label information but does not convey the node being protected. The information about a protected node is known to the MPT LSR and needs to be communicated to the PLR as well. For this reason, the FEC Label Mapping (FEC <R,X> : Lpx) sent by the MPT over the tLDP session to the PLR MUST include a Status TLV with an MP Status and a new LDP MP Status Value Element called the "Protected Node Status Value Element". This new value element is used to specify the address of the node being protected. The "Protected Node Status Value Element" has the following format:

```

      0               1               2               3
      0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
| Type = 3 |                               Length | Addr Family |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
| Addr Fam cont | Node address ~
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+

```

Type : Protected Node Status Value Element (Type 3).

Length: The Length field is an unsigned integer that encodes the length of the Status Value following the Length field. The encoded Length varies based on the Address Family and is 6 octets for Address Family + IPv4 address and 18 octets for Address Family + IPv6 address.

Addr Family: Two-octet quantity containing a value from IANA's "Address Family Numbers" registry [AFI] that encodes the address family for the node address.


Node address: Protected node address encoded according to the Address Family field.

When a PLR receives a Label Mapping for FEC <R,X> that includes a Protected Node Status, it will only use that label binding once the Node advertised in the Status value becomes unreachable. If the LSP is an MP2MP LSP, the PLR would have assigned a Label Mapping for the upstream MP2MP FEC Element to the MPT ([RFC6388], Section 3) for FEC <R,X>. This label binding on the MPT MUST only be used once node N becomes unreachable.

The procedures to determine if a node is unreachable is a local decision and not spelled out in this document. Typically, link failure or Bidirectional Forwarding Detection (BFD) can be used to determine and detect node unreachability.

4. Link or Node Failure

Consider the following topology:



N: The node being protected.
 M: The backup node to protect link LSR1 - N.
 ...: Backup LSPs from LSR1 to LSR2 and LSR3.

Figure 3


Assume that LSR1 is the PLR for protected node N and that LSR2 and LSR3 are MPTs for node N. When LSR1 discovers that node N is unreachable, it cannot immediately determine whether it is the link from LSR1 to N or the actual node N that has failed. In Figure 3,

the link between LSR1 and N is also protected using Fast Reroute (FRR) [RFC4090] link protection via node M. LSR1 MAY simultaneously invoke both link protection via node M to N using redirection of the traffic and node protection directly to LSR1 and LSR2. If only the link failed, LSR2 and LSR3 will receive the packets twice due to the two protection mechanisms. To prevent duplicate packets being forwarded to the receivers on the tree, LSR2 and LSR3 need to determine from which upstream node they should accept the packets. This can be either from the primary upstream LSR N or from the secondary upstream LSR1, but never both at the same time. The selection between the primary upstream LSR or (one or more) secondary upstream LSRs (on LSR2 and LSR3) is based on the reachability of the protected node N. As long as N is reachable from an MPT, the MPT should accept and forward the MPLS packets from N. Once N becomes unreachable, the LSPs from secondary upstream PLR LSRs (LSR1 in our example) are activated. Note that detecting if N is unreachable is a local decision and not spelled out in this document.

Typically, link failure or BFD can be used to determine and detect node unreachability.

4.1. Reconvergence after Node or Link Failure

Consider the following topology:



N: The node being protected.
 M: The backup node to protect link 'LSR1 - N'.
 P and Q: The nodes on the new primary path after failure of node N.
: P2P backup LSPs.

Figure 4

Assume that LSR1 has detected that node N is unreachable and invoked both the link protection and node protection procedures as described in this example. LSR1 is acting as PLR and sending traffic over both the backup P2P LSP to node N (via M) and the P2P LSPs directly to LSR2 and LSR3, acting as MPT LSRs. The sequence of events is dependent on whether the link from LSR1 to N has failed or node N itself has failed. The nodes downstream from the protected node (and participating in node protection) **MUST** have the capability to determine that the protected node has become unreachable. Otherwise, the procedures below cannot be applied.

4.1.1. Node Failure

If node N failed, both LSR2 and LSR3 will have changed the primary upstream LSR to the secondary upstream LSR (LSR1) due to node N being unreachable. With that, the label bindings previously assigned to LSR1 will be activated on the MPTs (LSR2 and LSR3) and the label binding to N will be disabled. Traffic is now switched over to the label bindings that were installed for node protection.

4.1.2. Link Failure

If the link 'LSR1 - N' has failed, both LSR2 and LSR3 will not change the primary upstream LSR because node N is still reachable. LSR2 and LSR3 will receive traffic over two different bindings, the primary label binding assigned to node N (due to link protection via node M) as well as over the binding assigned to LSR1 for the node protection. Since the secondary upstream LSRs have not been activated, the traffic received due to node protection will be dropped. Node N will reconverge and update LSR2 and LSR3 (Section 2.3) with the information that the PLR address (LSR1) is no longer applicable and must be removed. In response, LSR2 and LSR3 **MUST** send a Label Withdraw to LSR1 to withdraw the label binding. This will stop the traffic being forwarded over the backup P2P LSPs for node protection. LSR1 will respond back with a Label Release as soon as the binding has been removed.

4.1.3. Switching to New Primary Path

The network will eventually reconverge and a new best path to the root will be found by LSR2 and LSR3. LSR2 will find that P is its new primary upstream LSR to reach the root and LSR3 will find Q. Note that although the current active upstream LSR can either be node N or LSR1 (depending on link or node failure), it does not matter for the following procedures. Both LSR2 and LSR3 **SHOULD** use the Make-Before-Break (MBB) procedures as described in Section 8 of [RFC6388] to switch to the new primary upstream node. As soon as the new primary upstream LSRs P and Q are activated, a Label Withdraw message

MUST be sent to the old upstream LSR. Note that an upstream LSR switchover from a tLDP neighbor to a directly connected LDP neighbor is no different compared to switching between two directly connected neighbors. After the Label Withdraw message has been received by LSR1 or node N, forwarding will stop and a Label Release will be sent.

When it is determined that after reconvergence there is no more interest in the tLDP session between the MPT and the PLR, the tLDP session MAY be taken down. It is possible that having no more interest in the tLDP session is temporarily due to link flapping. In order to avoid the tLDP session from flapping, it is RECOMMENDED to apply a delay before tearing down the session. Determining the delay is a local implementation matter. If the operator is not concerned with the tLDP session flapping and/or other procedures are in place to avoid this altogether, there is no need to apply the delay.

5. mLDP Capabilities for Node Protection

In order to describe the capabilities of the participating LSRs, this document is organizing it per role in the network, i.e., Point of Local Repair (PLR), Merge Point (MPT), and protected node (as depicted in Figure 1).

5.1. PLR Capability

A PLR node should handle the following conditions:

1. Accept an incoming tLDP session from the MPT LSR.
2. Support the receipt of a "Protected Node Status Value Element" status in an MP Status TLV over tLDP session.
3. Upon node failure detection, capable of switching traffic towards one or more MPT(s) over a P2P LSP (bypassing N) using the labels previously advertised for MP LSPs over the tLDP session.

An LSR capable of performing these actions will advertise itself as PLR capable in the Node Protection Capability (see Section 5.4). This is a unidirectional capability announced from PLR to the protected LSR.

5.2. MPT Capability

An MPT node should handle the following conditions;

1. Support the receipt of "PLR Status Value Element" in an MP Status TLV from a protected node N.
2. Support to transmit "Protected Node Status Value Element" in an MP Status TLV to a PLR.

An LSR capable of performing these actions will advertise itself as MPT capable in the Node Protection Capability (see Section 5.4).

This is a unidirectional capability from MPT to the protected LSR.

5.3. The Protected LSR

A protected node should handle the following conditions:


1. Determine the PLR and MPT capability for directly connected upstream and downstream LSRs for a given MP FEC.
2. Support transmitting of "PLR Status Value Element" in an MP Status TLV to one or more downstream MPT LSRs.

The protected LSR does not advertise any capability for mLDP Node Protection because it does not need to receive any of the defined MP Status values as described above. However, the protected node does play an important role in the signaling and setup of the node protection. For a given FEC, the protected node can only send PLR information to a downstream LSR if the PLR has signaled the PLR capability and the downstream LSR has signaled the MPT capability. When the downstream LSR (acting as MPT) receives the PLR Status, it can implicitly infer that the advertised LSR(s) are PLR capable. The MPT LSR can now proceed with setting up a tLDP session with the PLR(s) and MP LSP node protection signaling.

5.4. The Node Protection Capability

We define a single capability "MP Node Protection Capability" to announce the PLR and MPT capability.

The format of the capability parameter TLV is as follows:



Where

U/F bits: MUST be set to 1 and 0, respectively (as per [RFC5561]).

Type: MP Node Protection Capability (Type = 0x0972).

Length: Unsigned integer; MUST be set to 2.

S bit: Set to 1 to announce and 0 to withdraw the capability (as per [RFC5561]).

P bit: Set to 1 to indicate the PLR is capable of MP LSP node protection.

M bit: Set to 1 to indicate the MPT is capable of MP LSP node protection.

Reserved: MUST be zero on transmit and ignored on receipt.

The above capability can be sent in an LDP Initialization message to announce capability at the session establishment time, or it can be sent in an LDP Capability message to dynamically update (announce or withdraw) its capability towards its peer using procedures specified in [RFC5561].

An LSR that supports the PLR functionality LSR MAY send this capability to its downstream MP peers with P bit set; whereas, an LSR that supports the MPT functionality MAY send this capability to its upstream peer with M bit set. Moreover, an LSR that supports both the PLR and MPT functionality MAY send this capability to its peers with both P and M bit set.

6. Security Considerations

The procedures in this document add two new TLVs to existing LDP messages. Those TLVs can be protected by the mechanisms that are used to protect LDP messages as described in [RFC6388] and [RFC5920]. If it were possible to attack the mechanisms described in this document, an LSR (a PLR or a MPT) could be induced to support a large number of tLDP sessions and set up an even larger number of LSPs. The security mechanisms described in [RFC6388] and [RFC5920] are believed to be adequate, but an implementation could provide additional protection by counting such protection sessions and LSPs and producing a log message to the operator if a threshold is crossed.

7. IANA Considerations

IANA has allocated the following two new code points from the "LDP MP Status Value Element type" registry within the "Label Distribution Protocol (LDP) Parameters" registry.

Value	Name	Reference
2	PLR Status Value Element	this doc
3	Protected Node Status Value Element	this doc

IANA has assigned the following new code point for a new Capability Parameter TLV. The code point has been assigned from the IETF Consensus range of the "TLV Type Name Space" registry within the "Label Distribution Protocol (LDP) Parameters" registry.

Value	Description	Reference	Notes/Reg Date
0x0972	MP Node Protection Capability	this doc	

8. References

8.1. Normative References

- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<http://www.rfc-editor.org/info/rfc2119>>.
- [RFC5036] Andersson, L., Ed., Minei, I., Ed., and B. Thomas, Ed., "LDP Specification", RFC 5036, DOI 10.17487/RFC5036, October 2007, <<http://www.rfc-editor.org/info/rfc5036>>.
- [RFC6388] Wijnands, IJ., Ed., Minei, I., Ed., Kompella, K., and B. Thomas, "Label Distribution Protocol Extensions for Point-to-Multipoint and Multipoint-to-Multipoint Label Switched Paths", RFC 6388, DOI 10.17487/RFC6388, November 2011, <<http://www.rfc-editor.org/info/rfc6388>>.
- [RFC5561] Thomas, B., Raza, K., Aggarwal, S., Aggarwal, R., and JL. Le Roux, "LDP Capabilities", RFC 5561, DOI 10.17487/RFC5561, July 2009, <<http://www.rfc-editor.org/info/rfc5561>>.
- [RFC7060] Napierala, M., Rosen, E., and IJ. Wijnands, "Using LDP Multipoint Extensions on Targeted LDP Sessions", RFC 7060, DOI 10.17487/RFC7060, November 2013, <<http://www.rfc-editor.org/info/rfc7060>>.
- [AFI] IANA, "Address Family Numbers", <<http://www.iana.org/assignments/address-family-numbers>>.

8.2. Informative References

- [RFC4090] Pan, P., Ed., Swallow, G., Ed., and A. Atlas, Ed., "Fast Reroute Extensions to RSVP-TE for LSP Tunnels", RFC 4090, DOI 10.17487/RFC4090, May 2005, <<http://www.rfc-editor.org/info/rfc4090>>.
- [RFC5286] Atlas, A., Ed., and A. Zinin, Ed., "Basic Specification for IP Fast Reroute: Loop-Free Alternates", RFC 5286, DOI 10.17487/RFC5286, September 2008, <<http://www.rfc-editor.org/info/rfc5286>>.

- [RFC5420] Farrel, A., Ed., Papadimitriou, D., Vasseur, JP., and A. Ayyangarps, "Encoding of Attributes for MPLS LSP Establishment Using Resource Reservation Protocol Traffic Engineering (RSVP-TE)", RFC 5420, DOI 10.17487/RFC5420, February 2009, <<http://www.rfc-editor.org/info/rfc5420>>.
- [RFC5920] Fang, L., Ed., "Security Framework for MPLS and GMPLS Networks", RFC 5920, DOI 10.17487/RFC5920, July 2010, <<http://www.rfc-editor.org/info/rfc5920>>.

Acknowledgments

The authors thank Nagendra Kumar, Duan Hong, Martin Vigoureux, Kenji Fujihira, Loa Andersson, and Ben Campbell for their comments on this document. Also, many thanks to Elwyn Davies and Adrian Farrel for the detailed review and contribution to this document.

Contributors

The following individual contributed to this document:

Eric Rosen
Juniper Networks, Inc.
10 Technology Park Drive
Westford, MA 01886
United States
Email: erosen@juniper.net

Authors' Addresses

IJsbrand Wijnands (editor)
Cisco Systems, Inc.
De kleetlaan 6a
Diegem 1831
Belgium

Email: ice@cisco.com

Kamran Raza
Cisco Systems, Inc.
2000 Innovation Drive
Ottawa, Ontario K2K-3E8
Canada

Email: skraza@cisco.com

Alia Atlas
Juniper Networks, Inc.
10 Technology Park Drive
Westford, MA 01886
United States

Email: akatlas@juniper.net

Jeff Tantsura
Ericsson
300 Holger Way
San Jose, CA 95134
United States

Email: jeff.tantsura@ericsson.com

Quintin Zhao
Huawei Technology
125 Nagog Technology Park
Acton, MA 01719
United States

Email: quintin.zhao@huawei.com