

Internet Engineering Task Force (IETF)
Request for Comments: 8312
Category: Informational
ISSN: 2070-1721

I. Rhee
NCSU
L. Xu
UNL
S. Ha
Colorado
A. Zimmermann

L. Eggert
R. Scheffenegger
NetApp
February 2018

CUBIC for Fast Long-Distance Networks

Abstract

CUBIC is an extension to the current TCP standards. It differs from the current TCP standards only in the congestion control algorithm on the sender side. In particular, it uses a cubic function instead of a linear window increase function of the current TCP standards to improve scalability and stability under fast and long-distance networks. CUBIC and its predecessor algorithm have been adopted as defaults by Linux and have been used for many years. This document provides a specification of CUBIC to enable third-party implementations and to solicit community feedback through experimentation on the performance of CUBIC.

Status of This Memo

This document is not an Internet Standards Track specification; it is published for informational purposes.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Not all documents approved by the IESG are a candidate for any level of Internet Standard; see Section 2 of RFC 7841.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <https://www.rfc-editor.org/info/rfc8312>.

Copyright Notice

Copyright (c) 2018 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	3
2. Conventions	3
3. Design Principles of CUBIC	4
4. CUBIC Congestion Control	6
4.1. Window Increase Function	6
4.2. TCP-Friendly Region	7
4.3. Concave Region	8
4.4. Convex Region	8
4.5. Multiplicative Decrease	8
4.6. Fast Convergence	9
4.7. Timeout	10
4.8. Slow Start	10
5. Discussion	10
5.1. Fairness to Standard TCP	11
5.2. Using Spare Capacity	13
5.3. Difficult Environments	13
5.4. Investigating a Range of Environments	13
5.5. Protection against Congestion Collapse	14
5.6. Fairness within the Alternative Congestion Control Algorithm	14
5.7. Performance with Misbehaving Nodes and Outside Attackers ..	14
5.8. Behavior for Application-Limited Flows	14
5.9. Responses to Sudden or Transient Events	14
5.10. Incremental Deployment	14
6. Security Considerations	15
7. IANA Considerations	15
8. References	15
8.1. Normative References	15
8.2. Informative References	16
Acknowledgements	17
Authors' Addresses	18

1. Introduction

The low utilization problem of TCP in fast long-distance networks is well documented in [K03] and [RFC3649]. This problem arises from a slow increase of the congestion window following a congestion event in a network with a large bandwidth-delay product (BDP). [HKLRX06] indicates that this problem is frequently observed even in the range of congestion window sizes over several hundreds of packets. This problem is equally applicable to all Reno-style TCP standards and their variants, including TCP-RENO [RFC5681], TCP-NewReno [RFC6582] [RFC6675], SCTP [RFC4960], and TFRC [RFC5348], which use the same linear increase function for window growth, which we refer to collectively as "Standard TCP" below.

CUBIC, originally proposed in [HRX08], is a modification to the congestion control algorithm of Standard TCP to remedy this problem. This document describes the most recent specification of CUBIC. Specifically, CUBIC uses a cubic function instead of a linear window increase function of Standard TCP to improve scalability and stability under fast and long-distance networks.

Binary Increase Congestion Control (BIC-TCP) [XHR04], a predecessor of CUBIC, was selected as the default TCP congestion control algorithm by Linux in the year 2005 and has been used for several years by the Internet community at large. CUBIC uses a similar window increase function as BIC-TCP and is designed to be less aggressive and fairer to Standard TCP in bandwidth usage than BIC-TCP while maintaining the strengths of BIC-TCP such as stability, window scalability, and RTT fairness. CUBIC has already replaced BIC-TCP as the default TCP congestion control algorithm in Linux and has been deployed globally by Linux. Through extensive testing in various Internet scenarios, we believe that CUBIC is safe for testing and deployment in the global Internet.

In the following sections, we first briefly explain the design principles of CUBIC, then provide the exact specification of CUBIC, and finally discuss the safety features of CUBIC following the guidelines specified in [RFC5033].

2. Conventions

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in BCP 14 [RFC2119] [RFC8174] when, and only when, they appear in all capitals, as shown here.

3. Design Principles of CUBIC

CUBIC is designed according to the following design principles:

Principle 1: For better network utilization and stability, CUBIC uses both the concave and convex profiles of a cubic function to increase the congestion window size, instead of using just a convex function.

Principle 2: To be TCP-friendly, CUBIC is designed to behave like Standard TCP in networks with short RTTs and small bandwidth where Standard TCP performs well.

Principle 3: For RTT-fairness, CUBIC is designed to achieve linear bandwidth sharing among flows with different RTTs.

Principle 4: CUBIC appropriately sets its multiplicative window decrease factor in order to balance between the scalability and convergence speed.

Principle 1: For better network utilization and stability, CUBIC [HRX08] uses a cubic window increase function in terms of the elapsed time from the last congestion event. While most alternative congestion control algorithms to Standard TCP increase the congestion window using convex functions, CUBIC uses both the concave and convex profiles of a cubic function for window growth. After a window reduction in response to a congestion event is detected by duplicate ACKs or Explicit Congestion Notification-Echo (ECN-Echo) ACKs [RFC3168], CUBIC registers the congestion window size where it got the congestion event as W_{max} and performs a multiplicative decrease of congestion window. After it enters into congestion avoidance, it starts to increase the congestion window using the concave profile of the cubic function. The cubic function is set to have its plateau at W_{max} so that the concave window increase continues until the window size becomes W_{max} . After that, the cubic function turns into a convex profile and the convex window increase begins. This style of window adjustment (concave and then convex) improves the algorithm stability while maintaining high network utilization [CEHRX07]. This is because the window size remains almost constant, forming a plateau around W_{max} where network utilization is deemed highest. Under steady state, most window size samples of CUBIC are close to W_{max} , thus promoting high network utilization and stability. Note that those congestion control algorithms using only convex functions to increase the congestion window size have the maximum increments around W_{max} , and thus introduce a large number of packet bursts around the saturation point of the network, likely causing frequent global loss synchronizations.

Principle 2: CUBIC promotes per-flow fairness to Standard TCP. Note that Standard TCP performs well under short RTT and small bandwidth (or small BDP) networks. There is only a scalability problem in networks with long RTTs and large bandwidth (or large BDP). An alternative congestion control algorithm to Standard TCP designed to be friendly to Standard TCP on a per-flow basis must operate to increase its congestion window less aggressively in small BDP networks than in large BDP networks. The aggressiveness of CUBIC mainly depends on the maximum window size before a window reduction, which is smaller in small BDP networks than in large BDP networks. Thus, CUBIC increases its congestion window less aggressively in small BDP networks than in large BDP networks. Furthermore, in cases when the cubic function of CUBIC increases its congestion window less aggressively than Standard TCP, CUBIC simply follows the window size of Standard TCP to ensure that CUBIC achieves at least the same throughput as Standard TCP in small BDP networks. We call this region where CUBIC behaves like Standard TCP, the "TCP-friendly region".

Principle 3: Two CUBIC flows with different RTTs have their throughput ratio linearly proportional to the inverse of their RTT ratio, where the throughput of a flow is approximately the size of its congestion window divided by its RTT. Specifically, CUBIC maintains a window increase rate independent of RTTs outside of the TCP-friendly region, and thus flows with different RTTs have similar congestion window sizes under steady state when they operate outside the TCP-friendly region. This notion of a linear throughput ratio is similar to that of Standard TCP under high statistical multiplexing environments where packet losses are independent of individual flow rates. However, under low statistical multiplexing environments, the throughput ratio of Standard TCP flows with different RTTs is quadratically proportional to the inverse of their RTT ratio [XHR04]. CUBIC always ensures the linear throughput ratio independent of the levels of statistical multiplexing. This is an improvement over Standard TCP. While there is no consensus on particular throughput ratios of different RTT flows, we believe that under wired Internet, use of a linear throughput ratio seems more reasonable than equal throughputs (i.e., the same throughput for flows with different RTTs) or a higher-order throughput ratio (e.g., a quadratical throughput ratio of Standard TCP under low statistical multiplexing environments).

Principle 4: To balance between the scalability and convergence speed, CUBIC sets the multiplicative window decrease factor to 0.7 while Standard TCP uses 0.5. While this improves the scalability of CUBIC, a side effect of this decision is slower convergence, especially under low statistical multiplexing environments. This design choice is following the observation that the author of

HighSpeed TCP (HSTCP) [RFC3649] has made along with other researchers (e.g., [GV02]): the current Internet becomes more asynchronous with less frequent loss synchronizations with high statistical multiplexing. Under this environment, even strict Multiplicative-Increase Multiplicative-Decrease (MIMD) can converge. CUBIC flows with the same RTT always converge to the same throughput independent of statistical multiplexing, thus achieving intra-algorithm fairness. We also find that under the environments with sufficient statistical multiplexing, the convergence speed of CUBIC flows is reasonable.

4. CUBIC Congestion Control

The unit of all window sizes in this document is segments of the maximum segment size (MSS), and the unit of all times is seconds. Let *cwnd* denote the congestion window size of a flow, and *ssthresh* denote the slow-start threshold.

4.1. Window Increase Function

CUBIC maintains the acknowledgment (ACK) clocking of Standard TCP by increasing the congestion window only at the reception of an ACK. It does not make any change to the fast recovery and retransmit of TCP, such as TCP-NewReno [RFC6582] [RFC6675]. During congestion avoidance after a congestion event where a packet loss is detected by duplicate ACKs or a network congestion is detected by ACKs with ECN-Echo flags [RFC3168], CUBIC changes the window increase function of Standard TCP. Suppose that *W_{max}* is the window size just before the window is reduced in the last congestion event.

CUBIC uses the following window increase function:

$$W_{\text{cubic}}(t) = C \cdot (t - K)^3 + W_{\text{max}} \quad (\text{Eq. 1})$$

where *C* is a constant fixed to determine the aggressiveness of window increase in high BDP networks, *t* is the elapsed time from the beginning of the current congestion avoidance, and *K* is the time period that the above function takes to increase the current window size to *W_{max}* if there are no further congestion events and is calculated using the following equation:

$$K = \text{cubic_root}(W_{\text{max}} \cdot (1 - \text{beta_cubic}) / C) \quad (\text{Eq. 2})$$

where *beta_cubic* is the CUBIC multiplication decrease factor, that is, when a congestion event is detected, CUBIC reduces its *cwnd* to *W_{cubic}(0) = W_{max} * beta_cubic*. We discuss how we set *beta_cubic* in Section 4.5 and how we set *C* in Section 5.

Upon receiving an ACK during congestion avoidance, CUBIC computes the window increase rate during the next RTT period using Eq. 1. It sets $W_{cubic}(t+RTT)$ as the candidate target value of the congestion window, where RTT is the weighted average RTT calculated by Standard TCP.

Depending on the value of the current congestion window size $cwnd$, CUBIC runs in three different modes.

1. The TCP-friendly region, which ensures that CUBIC achieves at least the same throughput as Standard TCP.
2. The concave region, if CUBIC is not in the TCP-friendly region and $cwnd$ is less than W_{max} .
3. The convex region, if CUBIC is not in the TCP-friendly region and $cwnd$ is greater than W_{max} .

Below, we describe the exact actions taken by CUBIC in each region.

4.2. TCP-Friendly Region

Standard TCP performs well in certain types of networks, for example, under short RTT and small bandwidth (or small BDP) networks. In these networks, we use the TCP-friendly region to ensure that CUBIC achieves at least the same throughput as Standard TCP.

The TCP-friendly region is designed according to the analysis described in [FHP00]. The analysis studies the performance of an Additive Increase and Multiplicative Decrease (AIMD) algorithm with an additive factor of α_{aimd} (segments per RTT) and a multiplicative factor of β_{aimd} , denoted by $AIMD(\alpha_{aimd}, \beta_{aimd})$. Specifically, the average congestion window size of $AIMD(\alpha_{aimd}, \beta_{aimd})$ can be calculated using Eq. 3. The analysis shows that $AIMD(\alpha_{aimd}, \beta_{aimd})$ with $\alpha_{aimd}=3*(1-\beta_{aimd})/(1+\beta_{aimd})$ achieves the same average window size as Standard TCP that uses $AIMD(1, 0.5)$.

$$AVG_W_aimd = [\alpha_{aimd} * (1+\beta_{aimd}) / (2*(1-\beta_{aimd})*p)]^{0.5} \text{ (Eq. 3)}$$

Based on the above analysis, CUBIC uses Eq. 4 to estimate the window size W_{est} of $AIMD(\alpha_{aimd}, \beta_{aimd})$ with $\alpha_{aimd}=3*(1-\beta_{cubic})/(1+\beta_{cubic})$ and $\beta_{aimd}=\beta_{cubic}$, which achieves the same average window size as Standard TCP. When receiving an ACK in congestion avoidance ($cwnd$ could be greater than

or less than W_{\max}), CUBIC checks whether $W_{\text{cubic}}(t)$ is less than $W_{\text{est}}(t)$. If so, CUBIC is in the TCP-friendly region and cwnd SHOULD be set to $W_{\text{est}}(t)$ at each reception of an ACK.

$$W_{\text{est}}(t) = W_{\max} \cdot \text{beta_cubic} + [3 \cdot (1 - \text{beta_cubic}) / (1 + \text{beta_cubic})] \cdot (t / \text{RTT}) \quad (\text{Eq. 4})$$

4.3. Concave Region

When receiving an ACK in congestion avoidance, if CUBIC is not in the TCP-friendly region and cwnd is less than W_{\max} , then CUBIC is in the concave region. In this region, cwnd MUST be incremented by $(W_{\text{cubic}}(t + \text{RTT}) - \text{cwnd}) / \text{cwnd}$ for each received ACK, where $W_{\text{cubic}}(t + \text{RTT})$ is calculated using Eq. 1.

4.4. Convex Region

When receiving an ACK in congestion avoidance, if CUBIC is not in the TCP-friendly region and cwnd is larger than or equal to W_{\max} , then CUBIC is in the convex region. The convex region indicates that the network conditions might have been perturbed since the last congestion event, possibly implying more available bandwidth after some flow departures. Since the Internet is highly asynchronous, some amount of perturbation is always possible without causing a major change in available bandwidth. In this region, CUBIC is being very careful by very slowly increasing its window size. The convex profile ensures that the window increases very slowly at the beginning and gradually increases its increase rate. We also call this region the "maximum probing phase" since CUBIC is searching for a new W_{\max} . In this region, cwnd MUST be incremented by $(W_{\text{cubic}}(t + \text{RTT}) - \text{cwnd}) / \text{cwnd}$ for each received ACK, where $W_{\text{cubic}}(t + \text{RTT})$ is calculated using Eq. 1.

4.5. Multiplicative Decrease

When a packet loss is detected by duplicate ACKs or a network congestion is detected by ECN-Echo ACKs, CUBIC updates its W_{\max} , cwnd , and ssthresh as follows. Parameter beta_cubic SHOULD be set to 0.7.

```
W_max = cwnd;           // save window size before reduction
ssthresh = cwnd * beta_cubic; // new slow-start threshold
ssthresh = max(ssthresh, 2); // threshold is at least 2 MSS
cwnd = cwnd * beta_cubic; // window reduction
```


A side effect of setting `beta_cubic` to a value bigger than 0.5 is slower convergence. We believe that while a more adaptive setting of `beta_cubic` could result in faster convergence, it will make the analysis of CUBIC much harder. This adaptive adjustment of `beta_cubic` is an item for the next version of CUBIC.

4.6. Fast Convergence

To improve the convergence speed of CUBIC, we add a heuristic in CUBIC. When a new flow joins the network, existing flows in the network need to give up some of their bandwidth to allow the new flow some room for growth if the existing flows have been using all the bandwidth of the network. To speed up this bandwidth release by existing flows, the following mechanism called "fast convergence" SHOULD be implemented.

With fast convergence, when a congestion event occurs, before the window reduction of the congestion window, a flow remembers the last value of `W_max` before it updates `W_max` for the current congestion event. Let us call the last value of `W_max` to be `W_last_max`.

```
if (W_max < W_last_max){ // should we make room for others
    W_last_max = W_max;           // remember the last W_max
    W_max = W_max*(1.0+beta_cubic)/2.0; // further reduce W_max
} else {
    W_last_max = W_max           // remember the last W_max
}
```

At a congestion event, if the current value of `W_max` is less than `W_last_max`, this indicates that the saturation point experienced by this flow is getting reduced because of the change in available bandwidth. Then we allow this flow to release more bandwidth by reducing `W_max` further. This action effectively lengthens the time for this flow to increase its congestion window because the reduced `W_max` forces the flow to have the plateau earlier. This allows more time for the new flow to catch up to its congestion window size.

The fast convergence is designed for network environments with multiple CUBIC flows. In network environments with only a single CUBIC flow and without any other traffic, the fast convergence SHOULD be disabled.

4.7. Timeout

In case of timeout, CUBIC follows Standard TCP to reduce cwnd [RFC5681], but sets ssthresh using `beta_cubic` (same as in Section 4.5) that is different from Standard TCP [RFC5681].

During the first congestion avoidance after a timeout, CUBIC increases its congestion window size using Eq. 1, where `t` is the elapsed time since the beginning of the current congestion avoidance, `K` is set to 0, and `W_max` is set to the congestion window size at the beginning of the current congestion avoidance.

4.8. Slow Start

CUBIC MUST employ a slow-start algorithm, when the cwnd is no more than ssthresh. Among the slow-start algorithms, CUBIC MAY choose the standard TCP slow start [RFC5681] in general networks, or the limited slow start [RFC3742] or hybrid slow start [HR08] for fast and long-distance networks.

In the case when CUBIC runs the hybrid slow start [HR08], it may exit the first slow start without incurring any packet loss and thus `W_max` is undefined. In this special case, CUBIC switches to congestion avoidance and increases its congestion window size using Eq. 1, where `t` is the elapsed time since the beginning of the current congestion avoidance, `K` is set to 0, and `W_max` is set to the congestion window size at the beginning of the current congestion avoidance.

5. Discussion

In this section, we further discuss the safety features of CUBIC following the guidelines specified in [RFC5033].

With a deterministic loss model where the number of packets between two successive packet losses is always $1/p$, CUBIC always operates with the concave window profile, which greatly simplifies the performance analysis of CUBIC. The average window size of CUBIC can be obtained by the following function:

$$\text{AVG_W_cubic} = [C * (3 + \text{beta_cubic}) / (4 * (1 - \text{beta_cubic}))]^{0.25} * (\text{RTT}^{0.75}) / (p^{0.75}) \quad (\text{Eq. 5})$$

With `beta_cubic` set to 0.7, the above formula is reduced to:

$$\text{AVG_W_cubic} = (C * 3.7 / 1.2)^{0.25} * (\text{RTT}^{0.75}) / (p^{0.75}) \quad (\text{Eq. 6})$$

We will determine the value of `C` in the following subsection using Eq. 6.

5.1. Fairness to Standard TCP

In environments where Standard TCP is able to make reasonable use of the available bandwidth, CUBIC does not significantly change this state.

Standard TCP performs well in the following two types of networks:

1. networks with a small bandwidth-delay product (BDP)
2. networks with a short RTTs, but not necessarily a small BDP

CUBIC is designed to behave very similarly to Standard TCP in the above two types of networks. The following two tables show the average window sizes of Standard TCP, HSTCP, and CUBIC. The average window sizes of Standard TCP and HSTCP are from [RFC3649]. The average window size of CUBIC is calculated using Eq. 6 and the CUBIC TCP-friendly region for three different values of C.

Loss Rate P	Average TCP W	Average HSTCP W	CUBIC (C=0.04)	CUBIC (C=0.4)	CUBIC (C=4)
10 ⁻²	12	12	12	12	12
10 ⁻³	38	38	38	38	59
10 ⁻⁴	120	263	120	187	333
10 ⁻⁵	379	1795	593	1054	1874
10 ⁻⁶	1200	12279	3332	5926	10538
10 ⁻⁷	3795	83981	18740	33325	59261
10 ⁻⁸	12000	574356	105383	187400	333250

Table 1

Table 1 describes the response function of Standard TCP, HSTCP, and CUBIC in networks with RTT = 0.1 seconds. The average window size is in MSS-sized segments.

Loss Rate P	Average TCP W	Average HSTCP W	CUBIC (C=0.04)	CUBIC (C=0.4)	CUBIC (C=4)
10 ⁻²	12	12	12	12	12
10 ⁻³	38	38	38	38	38
10 ⁻⁴	120	263	120	120	120
10 ⁻⁵	379	1795	379	379	379
10 ⁻⁶	1200	12279	1200	1200	1874
10 ⁻⁷	3795	83981	3795	5926	10538
10 ⁻⁸	12000	574356	18740	33325	59261

Table 2

Table 2 describes the response function of Standard TCP, HSTCP, and CUBIC in networks with RTT = 0.01 seconds. The average window size is in MSS-sized segments.

Both tables show that CUBIC with any of these three C values is more friendly to TCP than HSTCP, especially in networks with a short RTT where TCP performs reasonably well. For example, in a network with RTT = 0.01 seconds and $p=10^{-6}$, TCP has an average window of 1200 packets. If the packet size is 1500 bytes, then TCP can achieve an average rate of 1.44 Gbps. In this case, CUBIC with $C=0.04$ or $C=0.4$ achieves exactly the same rate as Standard TCP, whereas HSTCP is about ten times more aggressive than Standard TCP.

We can see that C determines the aggressiveness of CUBIC in competing with other congestion control algorithms for bandwidth. CUBIC is more friendly to Standard TCP, if the value of C is lower. However, we do not recommend setting C to a very low value like 0.04, since CUBIC with a low C cannot efficiently use the bandwidth in long RTT and high-bandwidth networks. Based on these observations and our experiments, we find $C=0.4$ gives a good balance between TCP-friendliness and aggressiveness of window increase. Therefore, C SHOULD be set to 0.4. With C set to 0.4, Eq. 6 is reduced to:

$$\text{AVG_W_cubic} = 1.054 * (\text{RTT}^{0.75}) / (p^{0.75}) \quad (\text{Eq. 7})$$

Eq. 7 is then used in the next subsection to show the scalability of CUBIC.

5.2. Using Spare Capacity

CUBIC uses a more aggressive window increase function than Standard TCP under long RTT and high-bandwidth networks.

The following table shows that to achieve the 10 Gbps rate, Standard TCP requires a packet loss rate of $2.0\text{e-}10$, while CUBIC requires a packet loss rate of $2.9\text{e-}8$.

Throughput(Mbps)	Average W	TCP P	HSTCP P	CUBIC P
1	8.3	$2.0\text{e-}2$	$2.0\text{e-}2$	$2.0\text{e-}2$
10	83.3	$2.0\text{e-}4$	$3.9\text{e-}4$	$2.9\text{e-}4$
100	833.3	$2.0\text{e-}6$	$2.5\text{e-}5$	$1.4\text{e-}5$
1000	8333.3	$2.0\text{e-}8$	$1.5\text{e-}6$	$6.3\text{e-}7$
10000	83333.3	$2.0\text{e-}10$	$1.0\text{e-}7$	$2.9\text{e-}8$

Table 3

Table 3 describes the required packet loss rate for Standard TCP, HSTCP, and CUBIC to achieve a certain throughput. We use 1500-byte packets and an RTT of 0.1 seconds.

Our test results in [HKLRX06] indicate that CUBIC uses the spare bandwidth left unused by existing Standard TCP flows in the same bottleneck link without taking away much bandwidth from the existing flows.

5.3. Difficult Environments

CUBIC is designed to remedy the poor performance of TCP in fast and long-distance networks.

5.4. Investigating a Range of Environments

CUBIC has been extensively studied by using both NS-2 simulation and test-bed experiments covering a wide range of network environments. More information can be found in [HKLRX06].

Same as Standard TCP, CUBIC is a loss-based congestion control algorithm. Because CUBIC is designed to be more aggressive (due to a faster window increase function and bigger multiplicative decrease factor) than Standard TCP in fast and long-distance networks, it can fill large drop-tail buffers more quickly than Standard TCP and

increase the risk of a standing queue [KWAf17]. In this case, proper queue sizing and management [RFC7567] could be used to reduce the packet queuing delay.

5.5. Protection against Congestion Collapse

With regard to the potential of causing congestion collapse, CUBIC behaves like Standard TCP since CUBIC modifies only the window adjustment algorithm of TCP. Thus, it does not modify the ACK clocking and Timeout behaviors of Standard TCP.

5.6. Fairness within the Alternative Congestion Control Algorithm

CUBIC ensures convergence of competing CUBIC flows with the same RTT in the same bottleneck links to an equal throughput. When competing flows have different RTTs, their throughput ratio is linearly proportional to the inverse of their RTT ratios. This is true independent of the level of statistical multiplexing in the link.

5.7. Performance with Misbehaving Nodes and Outside Attackers

This is not considered in the current CUBIC.

5.8. Behavior for Application-Limited Flows

CUBIC does not raise its congestion window size if the flow is currently limited by the application instead of the congestion window. In case of long periods when cwnd has not been updated due to the application rate limit, such as idle periods, t in Eq. 1 MUST NOT include these periods; otherwise, $W_{\text{cubic}}(t)$ might be very high after restarting from these periods.

5.9. Responses to Sudden or Transient Events

If there is a sudden congestion, a routing change, or a mobility event, CUBIC behaves the same as Standard TCP.

5.10. Incremental Deployment

CUBIC requires only the change of TCP senders, and it does not make any changes to TCP receivers. That is, a CUBIC sender works correctly with the Standard TCP receivers. In addition, CUBIC does not require any changes to the routers and does not require any assistance from the routers.

6. Security Considerations

This proposal makes no changes to the underlying security of TCP. More information about TCP security concerns can be found in [RFC5681].

7. IANA Considerations

This document does not require any IANA actions.

8. References

8.1. Normative References

- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<https://www.rfc-editor.org/info/rfc2119>>.
- [RFC3168] Ramakrishnan, K., Floyd, S., and D. Black, "The Addition of Explicit Congestion Notification (ECN) to IP", RFC 3168, DOI 10.17487/RFC3168, September 2001, <<https://www.rfc-editor.org/info/rfc3168>>.
- [RFC3649] Floyd, S., "HighSpeed TCP for Large Congestion Windows", RFC 3649, DOI 10.17487/RFC3649, December 2003, <<https://www.rfc-editor.org/info/rfc3649>>.
- [RFC3742] Floyd, S., "Limited Slow-Start for TCP with Large Congestion Windows", RFC 3742, DOI 10.17487/RFC3742, March 2004, <<https://www.rfc-editor.org/info/rfc3742>>.
- [RFC4960] Stewart, R., Ed., "Stream Control Transmission Protocol", RFC 4960, DOI 10.17487/RFC4960, September 2007, <<https://www.rfc-editor.org/info/rfc4960>>.
- [RFC5033] Floyd, S. and M. Allman, "Specifying New Congestion Control Algorithms", BCP 133, RFC 5033, DOI 10.17487/RFC5033, August 2007, <<https://www.rfc-editor.org/info/rfc5033>>.
- [RFC5348] Floyd, S., Handley, M., Padhye, J., and J. Widmer, "TCP Friendly Rate Control (TFRC): Protocol Specification", RFC 5348, DOI 10.17487/RFC5348, September 2008, <<https://www.rfc-editor.org/info/rfc5348>>.

- [RFC5681] Allman, M., Paxson, V., and E. Blanton, "TCP Congestion Control", RFC 5681, DOI 10.17487/RFC5681, September 2009, <<https://www.rfc-editor.org/info/rfc5681>>.
- [RFC6582] Henderson, T., Floyd, S., Gurtov, A., and Y. Nishida, "The NewReno Modification to TCP's Fast Recovery Algorithm", RFC 6582, DOI 10.17487/RFC6582, April 2012, <<https://www.rfc-editor.org/info/rfc6582>>.
- [RFC6675] Blanton, E., Allman, M., Wang, L., Jarvinen, I., Kojo, M., and Y. Nishida, "A Conservative Loss Recovery Algorithm Based on Selective Acknowledgment (SACK) for TCP", RFC 6675, DOI 10.17487/RFC6675, August 2012, <<https://www.rfc-editor.org/info/rfc6675>>.
- [RFC7567] Baker, F., Ed. and G. Fairhurst, Ed., "IETF Recommendations Regarding Active Queue Management", BCP 197, RFC 7567, DOI 10.17487/RFC7567, July 2015, <<https://www.rfc-editor.org/info/rfc7567>>.
- [RFC8174] Leiba, B., "Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words", BCP 14, RFC 8174, DOI 10.17487/RFC8174, May 2017, <<https://www.rfc-editor.org/info/rfc8174>>.

8.2. Informative References

- [CEHRX07] Cai, H., Eun, D., Ha, S., Rhee, I., and L. Xu, "Stochastic Ordering for Internet Congestion Control and its Applications", In Proceedings of IEEE INFOCOM, DOI 10.1109/INFCOM.2007.111, May 2007.
- [FHP00] Floyd, S., Handley, M., and J. Padhye, "A Comparison of Equation-Based and AIMD Congestion Control", May 2000.
- [GV02] Gorinsky, S. and H. Vin, "Extended Analysis of Binary Adjustment Algorithms", Technical Report TR2002-29, Department of Computer Sciences, The University of Texas at Austin, August 2002.
- [HKLRX06] Ha, S., Kim, Y., Le, L., Rhee, I., and L. Xu, "A Step toward Realistic Performance Evaluation of High-Speed TCP Variants", International Workshop on Protocols for Fast Long-Distance Networks.
- [HR08] Ha, S. and I. Rhee, "Hybrid Slow Start for High-Bandwidth and Long-Distance Networks", International Workshop on Protocols for Fast Long-Distance Networks.

- [HRX08] Ha, S., Rhee, I., and L. Xu, "CUBIC: A New TCP-Friendly High-Speed TCP Variant", ACM SIGOPS Operating System Review, DOI 10.1145/1400097.1400105, July 2008.
- [K03] Kelly, T., "Scalable TCP: Improving Performance in HighSpeed Wide Area Networks", ACM SIGCOMM Computer Communication Review, DOI 10.1145/956981.956989, April 2003.
- [KWF17] Khademi, N., Welzl, M., Armitage, G., and G. Fairhurst, "TCP Alternative Backoff with ECN (ABE)", Work in Progress, draft-ietf-tcpm-alternativebackoff-ecn-05, December 2017.
- [XHR04] Xu, L., Harfoush, K., and I. Rhee, "Binary Increase Congestion Control for Fast, Long Distance Networks", In Proceedings of IEEE INFOCOM, DOI 10.1109/INFCOM.2004.1354672, March 2004.

Acknowledgements

Alexander Zimmermann and Lars Eggert have received funding from the European Union's Horizon 2020 research and innovation program 2014-2018 under grant agreement No. 644866 (SSICLOPS). This document reflects only the authors' views and the European Commission is not responsible for any use that may be made of the information it contains.

The work of Lisong Xu was partially supported by the National Science Foundation (NSF) under Grant No. 1526253. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the authors and do not necessarily reflect the views of the NSF.

Authors' Addresses

Injong Rhee
North Carolina State University
Department of Computer Science
Raleigh, NC 27695-7534
United States of America
Email: rhee@ncsu.edu

Lisong Xu
University of Nebraska-Lincoln
Department of Computer Science and Engineering
Lincoln, NE 68588-0115
United States of America

Email: xu@unl.edu

Sangtae Ha
University of Colorado at Boulder
Department of Computer Science
Boulder, CO 80309-0430
United States of America
Email: sangtae.ha@colorado.edu

Alexander Zimmermann
Phone: +49 175 5766838
Email: alexander.zimmermann@rwth-aachen.de

Lars Eggert
NetApp
Sonnenallee 1
Kirchheim 85551
Germany
Phone: +49 151 12055791
Email: lars@netapp.com

Richard Scheffenegger
NetApp
Am Europlatz 2
Vienna 1120
Austria
Email: rs.ietf@gmx.at