

Internet Engineering Task Force (IETF)
Request for Comments: 8519
Category: Standards Track
ISSN: 2070-1721

M. Jethanandani
VMware
S. Agarwal
Cisco Systems, Inc.
L. Huang
D. Blair
March 2019

YANG Data Model for Network Access Control Lists (ACLs)

Abstract

This document defines a data model for Access Control Lists (ACLs). An ACL is a user-ordered set of rules used to configure the forwarding behavior in a device. Each rule is used to find a match on a packet and define actions that will be performed on the packet.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 7841.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <https://www.rfc-editor.org/info/rfc8519>.

Copyright Notice

Copyright (c) 2019 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	2
1.1. Definitions and Acronyms	3
1.2. Terminology	4
1.3. Tree Diagram	4
2. Problem Statement	4
3. Understanding ACL's Filters and Actions	4
3.1. ACL Modules	5
4. ACL YANG Models	9
4.1. IETF Access Control List Module	9
4.2. IETF Packet Fields Module	24
4.3. ACL Examples	37
4.4. Port Range Usage and Other Examples	39
5. Security Considerations	42
6. IANA Considerations	43
6.1. URI Registration	43
6.2. YANG Module Name Registration	44
7. References	44
7.1. Normative References	44
7.2. Informative References	46
Appendix A. Extending ACL Model Examples	47
A.1. Example of a Company's Proprietary Module	47
A.2. Linux nftables	50
A.3. Ethertypes	51
Acknowledgements	60
Authors' Addresses	60

1. Introduction

An Access Control List (ACL) is one of the basic elements used to configure device-forwarding behavior. It is used in many networking technologies such as Policy-Based Routing (PBR), firewalls, etc.

An ACL is a user-ordered set of rules that is used to filter traffic on a networking device. Each rule is represented by an Access Control Entry (ACE).

Each ACE has a group of match criteria and a group of actions.

The match criteria allow for the definition of packet headers and metadata, the contents of which must match the definitions.

- o Packet header matches apply to fields visible in the packet such as address, Class of Service (CoS), or port number.

- o In case a vendor supports it, metadata matches apply to fields associated with the packet, that are not in the packet header, such as the input interface or length of the packet as received over the wire.

The actions specify what to do with the packet when the matching criteria are met. These actions are any operations that would apply to the packet, such as counting, policing, or simply forwarding. The list of potential actions is unbounded depending on the capabilities of the networking devices.

Access Control List is also widely known as ACL (pronounced as [ak-uh l]) or Access List. In this document, Access Control List, ACL, and Access List are used interchangeably.

The matching of filters and actions in an ACE/ACL is triggered only after the application/attachment of the ACL to an interface, a Virtual Routing and Forwarding (VRF) interface, a vty/tty session, a QoS policy, or routing protocols, amongst various other configuration attachment points. Once attached, it is used for filtering traffic using the match criteria in the ACEs and taking appropriate action(s) that has been configured against that ACE. In order to apply an ACL to any attachment point other than an interface, vendors would have to augment the ACL YANG model.

1.1. Definitions and Acronyms

ACE: Access Control Entry

ACL: Access Control List

CoS: Class of Service

DSCP: Differentiated Services Code Point

ICMP: Internet Control Message Protocol

IP: Internet Protocol

IPv4: Internet Protocol version 4

IPv6: Internet Protocol version 6

MAC: Media Access Control

PBR: Policy-Based Routing

TCP: Transmission Control Protocol

UDP: User Datagram Protocol

1.2. Terminology

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in BCP 14 [RFC2119] [RFC8174] when, and only when, they appear in all capitals, as shown here.

1.3. Tree Diagram

For a reference to the annotations used in the tree diagrams included in this document, please see "YANG Tree Diagrams" [RFC8340].

2. Problem Statement

This document defines a YANG 1.1 data model [RFC7950] for the configuration of ACLs. The model defines matching rules for commonly used protocols such as Ethernet, IPv4, IPv6, TCP, UDP, and ICMP. If more protocols need to be supported in the future, this base model can be augmented. An example of such an augmentation can be seen in Appendix A.

ACL implementations in every device may vary greatly in terms of the filter constructs and actions that they support. Therefore, this document proposes a model that can be augmented by standard extensions and vendor proprietary models.

3. Understanding ACL's Filters and Actions

Although different vendors have different ACL data models, there is a common understanding of what an ACL is. A network system usually has a list of ACLs, and each ACL contains an ordered list of rules, also known as ACEs. Each ACE has a group of match criteria and a group of actions. The match criteria allows for definition of the contents of the packet headers or metadata, if supported by the vendor. Packet header matching applies to fields visible in the packet such as address, CoS, or port number. Metadata matching applies to fields associated with the packet, that are not in the packet header, such as the input interface, packet length, or source or destination prefix length. The actions can be any sort of operation from logging to rate-limiting or dropping to simply forwarding. Actions on the first matching ACE are applied with no processing of subsequent ACEs.

The model also includes a container to hold overall operational state for each ACL and for each ACE. One ACL can be applied to multiple targets within the device, such as the interface of a networking device, applications or features running in the device, etc. When applied to interfaces of a networked device, distinct ACLs are defined for the ingress (input) or egress (output) interface.

This document tries to address the commonalities between all vendors and creates a common model, which can be augmented with proprietary models. The base model is simple in design, and we hope to achieve enough flexibility for each vendor to extend the base model.

The use of feature statements in the model allows vendors to advertise match rules they are capable and willing to support. There are two sets of feature statements a device needs to advertise. The first set of feature statements specifies the capability of the device. These include features such as "Device can support matching on Ethernet headers" or "Device can support matching on IPv4 headers". The second set of feature statements specifies the combinations of headers the device is willing to support. These include features such as "Plain IPv6 ACL supported" or "Ethernet, IPv4 and IPv6 ACL combinations supported".

3.1. ACL Modules

There are two YANG modules in the model. The first module, "ietf-access-control-list", defines generic ACL aspects that are common to all ACLs regardless of their type or vendor. In effect, the module can be viewed as providing a generic ACL "superclass". It imports the second module, "ietf-packet-fields". The match container in "ietf-access-control-list" uses groupings in "ietf-packet-fields" to specify match fields such as port numbers or protocols. The combination of 'if-feature' checks and 'must' statements allows for the selection of relevant match fields that a user can define rules for.

If there is a need to define a new "matches" choice, such as IP Flow Information Export (IPFIX) [RFC7011], the container "matches" can be augmented.

```

module: ietf-access-control-list
  +--rw acls
    +--rw acl* [name]
      +--rw name      string
      +--rw type?    acl-type
      +--rw aces
        +--rw ace* [name]
          +--rw name      string
          +--rw matches
            +--rw (l2)?
              +--:(eth)
                +--rw eth {match-on-eth}?
                  +--rw destination-mac-address?
                    | yang:mac-address
                  +--rw destination-mac-address-mask?
                    | yang:mac-address
                  +--rw source-mac-address?
                    | yang:mac-address
                  +--rw source-mac-address-mask?
                    | yang:mac-address
                  +--rw ethertype?
                    | eth:ethertype
            +--rw (l3)?
              +--:(ipv4)
                +--rw ipv4 {match-on-ipv4}?
                  +--rw dscp?
                    | inet:dscp
                  +--rw ecn?
                    | uint8
                  +--rw length?
                    | uint16
                  +--rw ttl?
                    | uint8
                  +--rw protocol?
                    | uint8
                  +--rw ihl?
                    | uint8
                  +--rw flags?
                    | bits
                  +--rw offset?
                    | uint16
                  +--rw identification?
                    | uint16
                  +--rw (destination-network)?
                    +--:(destination-ipv4-network)
                      +--rw destination-ipv4-network?
                        | inet:ipv4-prefix

```

```

    +---rw (source-network)?
      +---:(source-ipv4-network)
        +---rw source-ipv4-network?
          inet:ipv4-prefix
    +---:(ipv6)
      +---rw ipv6 {match-on-ipv6}?
        +---rw dscp?
          | inet:dscp
        +---rw ecn?
          | uint8
        +---rw length?
          | uint16
        +---rw ttl?
          | uint8
        +---rw protocol?
          | uint8
        +---rw (destination-network)?
          | +---:(destination-ipv6-network)
            | +---rw destination-ipv6-network?
              | inet:ipv6-prefix
        +---rw (source-network)?
          | +---:(source-ipv6-network)
            | +---rw source-ipv6-network?
              | inet:ipv6-prefix
        +---rw flow-label?
          | inet:ipv6-flow-label
    +---rw (l4)?
      +---:(tcp)
        +---rw tcp {match-on-tcp}?
          +---rw sequence-number?          uint32
          +---rw acknowledgement-number?    uint32
          +---rw data-offset?               uint8
          +---rw reserved?                  uint8
          +---rw flags?                     bits
          +---rw window-size?               uint16
          +---rw urgent-pointer?             uint16
          +---rw options?                   binary
          +---rw source-port
            +---rw (source-port)?
              +---:(range-or-operator)
                +---rw (port-range-or-operator)?
                  +---:(range)
                    +---rw lower-port
                      | inet:port-number
                    +---rw upper-port
                      | inet:port-number
                  +---:(operator)

```



```

|
|
|
|      +---rw icmp {match-on-icmp}?
|      |      +---rw type?          uint8
|      |      +---rw code?          uint8
|      |      +---rw rest-of-header? binary
|      +---rw egress-interface?    if:interface-ref
|      +---rw ingress-interface?   if:interface-ref
+---rw actions
|      +---rw forwarding    identityref
|      +---rw logging?      identityref
+---ro statistics {acl-aggregate-stats}?
|      +---ro matched-packets? yang:counter64
|      +---ro matched-octets?  yang:counter64
+---rw attachment-points
|      +---rw interface* [interface-id] {interface-attachment}?
|      |      +---rw interface-id    if:interface-ref
|      |      +---rw ingress
|      |      |      +---rw acl-sets
|      |      |      |      +---rw acl-set* [name]
|      |      |      |      |      +---rw name          -> /acls/acl/name
|      |      |      |      |      +---ro ace-statistics* [name] {interface-stats}?
|      |      |      |      |      |      +---ro name
|      |      |      |      |      |      |      -> /acls/acl/aces/ace/name
|      |      |      |      |      |      +---ro matched-packets? yang:counter64
|      |      |      |      |      |      +---ro matched-octets?  yang:counter64
|      |      |      +---rw egress
|      |      |      |      +---rw acl-sets
|      |      |      |      |      +---rw acl-set* [name]
|      |      |      |      |      |      +---rw name          -> /acls/acl/name
|      |      |      |      |      |      +---ro ace-statistics* [name] {interface-stats}?
|      |      |      |      |      |      |      +---ro name
|      |      |      |      |      |      |      |      -> /acls/acl/aces/ace/name
|      |      |      |      |      |      |      +---ro matched-packets? yang:counter64
|      |      |      |      |      |      |      +---ro matched-octets?  yang:counter64

```

4. ACL YANG Models

4.1. IETF Access Control List Module

The "ietf-access-control-list" module defines the "acls" container that has a list of each "acl". Each "acl" has information identifying the access list by a name ("name") and a list ("aces") of rules associated with the "name". Each of the entries in the list ("aces"), indexed by the string "name", has containers defining "matches" and "actions".

The model defines several ACL types and actions in the form of identities and features. Features are used by implementors to select the ACL types the system can support, and identities are used to validate the types that have been selected. These types are implicitly inherited by the "ace", thus safeguarding against misconfiguration of "ace" types in an "acl".

The "matches" define criteria used to identify patterns in "ietf-packet-fields". The choice statements within the match container allow for the selection of one header within each of "l2", "l3", or "l4" headers. The "actions" define the behavior to undertake once a "match" has been identified. In addition to permit and deny actions, a logging option allows for a match to be logged that can later be used to determine which rule was matched upon. The model also defines the ability for ACLs to be attached to a particular interface.

Statistics in the ACL can be collected for an "ace" or for an "interface". The feature statements defined for statistics can be used to determine whether statistics are being collected per "ace" or per "interface".

This module imports definitions from "Common YANG Data Types" [RFC6991] and "A YANG Data Model for Interface Management" [RFC8343].

<CODE BEGINS> file "ietf-access-control-list@2019-03-04.yang"

```
module ietf-access-control-list {
  yang-version 1.1;
  namespace "urn:ietf:params:xml:ns:yang:ietf-access-control-list";
  prefix acl;

  import ietf-yang-types {
    prefix yang;
    reference
      "RFC 6991 - Common YANG Data Types.";
  }

  import ietf-packet-fields {
    prefix pf;
    reference
      "RFC 8519 - YANG Data Model for Network Access Control
        Lists (ACLs).";
  }

  import ietf-interfaces {
    prefix if;
    reference
```

```
"RFC 8343 - A YANG Data Model for Interface Management.";
}

organization
  "IETF NETMOD (Network Modeling) Working Group.";

contact
  "WG Web: <https://datatracker.ietf.org/wg/netmod/>
  WG List: netmod@ietf.org

  Editor: Mahesh Jethanandani
          mjethanandani@gmail.com
  Editor: Lisa Huang
          huangyi\_99@yahoo.com
  Editor: Sonal Agarwal
          sagarwal12@gmail.com
  Editor: Dana Blair
          dana@blairhome.com";

description
  "This YANG module defines a component that describes the
  configuration and monitoring of Access Control Lists (ACLs).

  The key words 'MUST', 'MUST NOT', 'REQUIRED', 'SHALL',
  'SHALL NOT', 'SHOULD', 'SHOULD NOT', 'RECOMMENDED',
  'NOT RECOMMENDED', 'MAY', and 'OPTIONAL' in this document
  are to be interpreted as described in BCP 14 (RFC 2119)
  (RFC 8174) when, and only when, they appear in all
  capitals, as shown here.

  Copyright (c) 2019 IETF Trust and the persons identified as
  the document authors. All rights reserved.

  Redistribution and use in source and binary forms, with or
  without modification, is permitted pursuant to, and subject
  to the license terms contained in, the Simplified BSD
  License set forth in Section 4.c of the IETF Trust's Legal
  Provisions Relating to IETF Documents
  (http://trustee.ietf.org/license-info).

  This version of this YANG module is part of RFC 8519; see
  the RFC itself for full legal notices.";

revision 2019-03-04 {
  description
    "Initial version.";
  reference
    "RFC 8519: YANG Data Model for Network Access Control
```

```
        Lists (ACLs).";
    }

    /*
     * Identities
     */
    /*
     * Forwarding actions for a packet
     */

    identity forwarding-action {
        description
            "Base identity for actions in the forwarding category.";
    }

    identity accept {
        base forwarding-action;
        description
            "Accept the packet.";
    }

    identity drop {
        base forwarding-action;
        description
            "Drop packet without sending any ICMP error message.";
    }

    identity reject {
        base forwarding-action;
        description
            "Drop the packet and send an ICMP error message to the source.";
    }

    /*
     * Logging actions for a packet
     */

    identity log-action {
        description
            "Base identity for defining the destination for logging
            actions.";
    }

    identity log-syslog {
        base log-action;
        description
            "System log (syslog) the information for the packet.";
    }
}
```

```
identity log-none {
  base log-action;
  description
    "No logging for the packet.";
}

/*
 * ACL type identities
 */

identity acl-base {
  description
    "Base Access Control List type for all Access Control List type
    identifiers.";
}

identity ipv4-acl-type {
  base acl:acl-base;
  if-feature "ipv4";
  description
    "An ACL that matches on fields from the IPv4 header
    (e.g., IPv4 destination address) and Layer 4 headers (e.g., TCP
    destination port). An ACL of type ipv4 does not contain
    matches on fields in the Ethernet header or the IPv6 header.";
}

identity ipv6-acl-type {
  base acl:acl-base;
  if-feature "ipv6";
  description
    "An ACL that matches on fields from the IPv6 header
    (e.g., IPv6 destination address) and Layer 4 headers (e.g., TCP
    destination port). An ACL of type ipv6 does not contain
    matches on fields in the Ethernet header or the IPv4 header.";
}

identity eth-acl-type {
  base acl:acl-base;
  if-feature "eth";
  description
    "An ACL that matches on fields in the Ethernet header,
    like 10/100/1000baseT or a Wi-Fi Access Control List. An ACL
    of type ethernet does not contain matches on fields in the
    IPv4 header, the IPv6 header, or Layer 4 headers.";
}

identity mixed-eth-ipv4-acl-type {
  base acl:eth-acl-type;
}
```

```
    base acl:ipv4-acl-type;
    if-feature "mixed-eth-ipv4";
    description
        "An ACL that contains a mix of entries that match
         on fields in Ethernet headers and in IPv4 headers.
         Matching on Layer 4 header fields may also exist in the
         list.";
}

identity mixed-eth-ipv6-acl-type {
    base acl:eth-acl-type;
    base acl:ipv6-acl-type;
    if-feature "mixed-eth-ipv6";
    description
        "An ACL that contains a mix of entries that match on fields
         in Ethernet headers and in IPv6 headers. Matching
         on Layer 4 header fields may also exist in the list.";
}

identity mixed-eth-ipv4-ipv6-acl-type {
    base acl:eth-acl-type;
    base acl:ipv4-acl-type;
    base acl:ipv6-acl-type;
    if-feature "mixed-eth-ipv4-ipv6";
    description
        "An ACL that contains a mix of entries that
         match on fields in Ethernet headers, IPv4 headers, and IPv6
         headers. Matching on Layer 4 header fields may also exist
         in the list.";
}

/*
 * Features
 */

/*
 * Features supported by device
 */
feature match-on-eth {
    description
        "The device can support matching on Ethernet headers.";
}

feature match-on-ipv4 {
    description
        "The device can support matching on IPv4 headers.";
}
```

```
feature match-on-ipv6 {
  description
    "The device can support matching on IPv6 headers.";
}

feature match-on-tcp {
  description
    "The device can support matching on TCP headers.";
}

feature match-on-udp {
  description
    "The device can support matching on UDP headers.";
}

feature match-on-icmp {
  description
    "The device can support matching on ICMP (v4 and v6) headers.";
}

/*
 * Header classifications combinations supported by
 * device
 */

feature eth {
  if-feature "match-on-eth";
  description
    "Plain Ethernet ACL supported.";
}

feature ipv4 {
  if-feature "match-on-ipv4";
  description
    "Plain IPv4 ACL supported.";
}

feature ipv6 {
  if-feature "match-on-ipv6";
  description
    "Plain IPv6 ACL supported.";
}

feature mixed-eth-ipv4 {
  if-feature "match-on-eth and match-on-ipv4";
  description
    "Ethernet and IPv4 ACL combinations supported.";
}
```

```
feature mixed-eth-ipv6 {
  if-feature "match-on-eth and match-on-ipv6";
  description
    "Ethernet and IPv6 ACL combinations supported.";
}

feature mixed-eth-ipv4-ipv6 {
  if-feature
    "match-on-eth and match-on-ipv4
    and match-on-ipv6";
  description
    "Ethernet, IPv4, and IPv6 ACL combinations supported.";
}

/*
 * Stats Features
 */
feature interface-stats {
  description
    "ACL counters are available and reported only per interface.";
}

feature acl-aggregate-stats {
  description
    "ACL counters are aggregated over all interfaces and reported
    only per ACL entry.";
}

/*
 * Attachment point features
 */
feature interface-attachment {
  description
    "ACLs are set on interfaces.";
}

/*
 * Typedefs
 */
typedef acl-type {
  type identityref {
    base acl-base;
  }
  description
    "This type is used to refer to an ACL type.";
}

/*
```



```
* Groupings
*/
grouping acl-counters {
  description
    "Common grouping for ACL counters.";
  leaf matched-packets {
    type yang:counter64;
    config false;
    description
      "Count of the number of packets matching the current ACL
      entry.

      An implementation should provide this counter on a
      per-interface, per-ACL-entry basis if possible.

      If an implementation only supports ACL counters on a per-
      entry basis (i.e., not broken out per interface), then the
      value should be equal to the aggregate count across all
      interfaces.

      An implementation that provides counters on a per-entry, per-
      interface basis is not required to also provide an aggregate
      count, e.g., per entry -- the user is expected to be able to
      implement the required aggregation if such a count is
      needed.";
  }
  leaf matched-octets {
    type yang:counter64;
    config false;
    description
      "Count of the number of octets (bytes) matching the current
      ACL entry.

      An implementation should provide this counter on a
      per-interface, per-ACL-entry basis if possible.

      If an implementation only supports ACL counters per entry
      (i.e., not broken out per interface), then the value
      should be equal to the aggregate count across all interfaces.

      An implementation that provides counters per entry per
      interface is not required to also provide an aggregate count,
      e.g., per entry -- the user is expected to be able to
      implement the required aggregation if such a count is needed.";
  }
}
```

```
/*
 * Configuration and monitoring data nodes
 */

container acls {
  description
    "This is a top-level container for Access Control Lists.
    It can have one or more acl nodes.";
  list acl {
    key "name";
    description
      "An ACL is an ordered list of ACEs. Each ACE has a
      list of match criteria and a list of actions.
      Since there are several kinds of ACLs implemented
      with different attributes for different vendors,
      this model accommodates customizing ACLs for
      each kind and for each vendor.";
    leaf name {
      type string {
        length "1..64";
      }
      description
        "The name of the access list. A device MAY further
        restrict the length of this name; space and special
        characters are not allowed.";
    }
    leaf type {
      type acl-type;
      description
        "Type of ACL. Indicates the primary intended
        type of match criteria (e.g., Ethernet, IPv4, IPv6, mixed,
        etc.) used in the list instance.";
    }
  }
  container aces {
    description
      "The aces container contains one or more ACE nodes.";
    list ace {
      key "name";
      ordered-by user;
      description
        "List of ACEs.";
      leaf name {
        type string {
          length "1..64";
        }
        description
          "A unique name identifying this ACE.";
      }
    }
  }
}
```

```
container matches {
  description
    "The rules in this set determine what fields will be
    matched upon before any action is taken on them.
    The rules are selected based on the feature set
    defined by the server and the acl-type defined.
    If no matches are defined in a particular container,
    then any packet will match that container. If no
    matches are specified at all in an ACE, then any
    packet will match the ACE.";

  choice l2 {
    container eth {
      when "derived-from-or-self(/acls/acl/type, "
        + "'acl:eth-acl-type')";
      if-feature "match-on-eth";
      uses pf:acl-eth-header-fields;
      description
        "Rule set that matches Ethernet headers.";
    }
    description
      "Match Layer 2 headers, for example, Ethernet
      header fields.";
  }

  choice l3 {
    container ipv4 {
      when "derived-from-or-self(/acls/acl/type, "
        + "'acl:ipv4-acl-type')";
      if-feature "match-on-ipv4";
      uses pf:acl-ip-header-fields;
      uses pf:acl-ipv4-header-fields;
      description
        "Rule set that matches IPv4 headers.";
    }

    container ipv6 {
      when "derived-from-or-self(/acls/acl/type, "
        + "'acl:ipv6-acl-type')";
      if-feature "match-on-ipv6";
      uses pf:acl-ip-header-fields;
      uses pf:acl-ipv6-header-fields;
      description
        "Rule set that matches IPv6 headers.";
    }
    description
      "Choice of either IPv4 or IPv6 headers";
  }
}
```

```
choice l4 {
  container tcp {
    if-feature "match-on-tcp";
    uses pf:acl-tcp-header-fields;
    container source-port {
      choice source-port {
        case range-or-operator {
          uses pf:port-range-or-operator;
          description
            "Source port definition from range or
            operator.";
        }
        description
          "Choice of source port definition using
          range/operator or a choice to support future
          'case' statements, such as one enabling a
          group of source ports to be referenced.";
      }
      description
        "Source port definition.";
    }
    container destination-port {
      choice destination-port {
        case range-or-operator {
          uses pf:port-range-or-operator;
          description
            "Destination port definition from range or
            operator.";
        }
        description
          "Choice of destination port definition using
          range/operator or a choice to support future
          'case' statements, such as one enabling a
          group of destination ports to be referenced.";
      }
      description
        "Destination port definition.";
    }
    description
      "Rule set that matches TCP headers.";
  }
}

container udp {
  if-feature "match-on-udp";
  uses pf:acl-udp-header-fields;
  container source-port {
    choice source-port {
      case range-or-operator {
```

```

        uses pf:port-range-or-operator;
        description
            "Source port definition from range or
            operator.";
    }
    description
        "Choice of source port definition using
        range/operator or a choice to support future
        'case' statements, such as one enabling a
        group of source ports to be referenced.";
    }
    description
        "Source port definition.";
}
container destination-port {
    choice destination-port {
        case range-or-operator {
            uses pf:port-range-or-operator;
            description
                "Destination port definition from range or
                operator.";
        }
        description
            "Choice of destination port definition using
            range/operator or a choice to support future
            'case' statements, such as one enabling a
            group of destination ports to be referenced.";
    }
    description
        "Destination port definition.";
}
description
    "Rule set that matches UDP headers.";
}

container icmp {
    if-feature "match-on-icmp";
    uses pf:acl-icmp-header-fields;
    description
        "Rule set that matches ICMP headers.";
}
description
    "Choice of TCP, UDP, or ICMP headers.";
}

leaf egress-interface {
    type if:interface-ref;
    description

```

```

        "Egress interface. This should not be used if this ACL
        is attached as an egress ACL (or the value should
        equal the interface to which the ACL is attached).";
    }

    leaf ingress-interface {
        type if:interface-ref;
        description
            "Ingress interface. This should not be used if this ACL
            is attached as an ingress ACL (or the value should
            equal the interface to which the ACL is attached).";
    }
}

container actions {
    description
        "Definition of actions for this ace entry.";
    leaf forwarding {
        type identityref {
            base forwarding-action;
        }
        mandatory true;
        description
            "Specifies the forwarding action per ace entry.";
    }

    leaf logging {
        type identityref {
            base log-action;
        }
        default "log-none";
        description
            "Specifies the log action and destination for
            matched packets. Default value is not to log the
            packet.";
    }
}

container statistics {
    if-feature "acl-aggregate-stats";
    config false;
    description
        "Statistics gathered across all attachment points for the
        given ACL.";
    uses acl-counters;
}
}
}
}

```

```

container attachment-points {
  description
    "Enclosing container for the list of
    attachment points on which ACLs are set.";
  /*
   * Groupings
   */
  grouping interface-acl {
    description
      "Grouping for per-interface ingress ACL data.";
    container acl-sets {
      description
        "Enclosing container for the list of ingress ACLs on the
        interface.";
      list acl-set {
        key "name";
        ordered-by user;
        description
          "List of ingress ACLs on the interface.";
        leaf name {
          type leafref {
            path "/acls/acl/name";
          }
          description
            "Reference to the ACL name applied on the ingress.";
        }
        list ace-statistics {
          if-feature "interface-stats";
          key "name";
          config false;
          description
            "List of ACEs.";
          leaf name {
            type leafref {
              path "/acls/acl/aces/ace/name";
            }
            description
              "Name of the ace entry.";
          }
          uses acl-counters;
        }
      }
    }
  }
}

list interface {
  if-feature "interface-attachment";
  key "interface-id";
}

```

```

    description
      "List of interfaces on which ACLs are set.";

    leaf interface-id {
      type if:interface-ref;
      description
        "Reference to the interface id list key.";
    }

    container ingress {
      uses interface-acl;
      description
        "The ACLs applied to the ingress interface.";
    }
    container egress {
      uses interface-acl;
      description
        "The ACLs applied to the egress interface.";
    }
  }
}
}
}
}
}

```

<CODE ENDS>

4.2. IETF Packet Fields Module

The packet fields module defines the necessary groups for matching on fields in the packet including Ethernet, IPv4, IPv6, and transport-layer fields. The "type" node determines which of these fields get included for any given ACL with the exception of TCP, UDP, and ICMP header fields. Those fields can be used in conjunction with any of the above Layer 2 or Layer 3 fields.

Since the number of match criteria are very large, the base specification does not include these directly but references them by the 'uses' statement to keep the base module simple. In case more match conditions are needed, those can be added by augmenting choices within container "matches" in the ietf-access-control-list.yang data model.

This module imports definitions from "Common YANG Data Types" [RFC6991] and references "Internet Protocol" [RFC791], "Internet Control Message Protocol" [RFC792], "Transmission Control Protocol" [RFC793], "Definition of the Differentiated Services Field (DS Field) in the IPv4 and IPv6 Headers" [RFC2474], "The Addition of Explicit Congestion Notification (ECN) to IP" [RFC3168], "IPv6 Scoped Address

Architecture" [RFC4007], "IP Version 6 Addressing Architecture" [RFC4291], "A Recommendation for IPv6 Address Text Representation" [RFC5952], and "Internet Protocol, Version 6 (IPv6) Specification" [RFC8200].

<CODE BEGINS> file "ietf-packet-fields@2019-03-04.yang"

```
module ietf-packet-fields {
  yang-version 1.1;
  namespace "urn:ietf:params:xml:ns:yang:ietf-packet-fields";
  prefix packet-fields;

  import ietf-inet-types {
    prefix inet;
    reference
      "RFC 6991 - Common YANG Data Types.";
  }

  import ietf-yang-types {
    prefix yang;
    reference
      "RFC 6991 - Common YANG Data Types.";
  }

  import ietf-ethertypes {
    prefix eth;
    reference
      "RFC 8519 - YANG Data Model for Network Access Control
        Lists (ACLs).";
  }

  organization
    "IETF NETMOD (Network Modeling) Working Group.";

  contact
    "WG Web: <https://datatracker.ietf.org/wg/netmod/>
    WG List: netmod@ietf.org

    Editor: Mahesh Jethanandani
            mjethanandani@gmail.com
    Editor: Lisa Huang
            huangyi_99@yahoo.com
    Editor: Sonal Agarwal
            sagarwal12@gmail.com
    Editor: Dana Blair
            dana@blairhome.com";
```

description

"This YANG module defines groupings that are used by the ietf-access-control-list YANG module. Their usage is not limited to ietf-access-control-list and can be used anywhere as applicable.

Copyright (c) 2019 IETF Trust and the persons identified as the document authors. All rights reserved.

Redistribution and use in source and binary forms, with or without modification, is permitted pursuant to, and subject to the license terms contained in, the Simplified BSD License set forth in Section 4.c of the IETF Trust's Legal Provisions Relating to IETF Documents (<http://trustee.ietf.org/license-info>).

This version of this YANG module is part of RFC 8519; see the RFC itself for full legal notices.";

```
revision 2019-03-04 {
  description
    "Initial version.";
  reference
    "RFC 8519: YANG Data Model for Network Access Control
      Lists (ACLs).";
}

/*
 * Typedefs
 */
typedef operator {
  type enumeration {
    enum lte {
      description
        "Less than or equal to.";
    }
    enum gte {
      description
        "Greater than or equal to.";
    }
    enum eq {
      description
        "Equal to.";
    }
    enum neq {
      description
        "Not equal to.";
    }
  }
}
```

```
    }
    description
      "The source and destination port range definitions
      can be further qualified using an operator. An
      operator is needed only if the lower-port is specified
      and the upper-port is not specified. The operator
      therefore further qualifies the lower-port only.";
  }

  /*
  * Groupings
  */
  grouping port-range-or-operator {
    choice port-range-or-operator {
      case range {
        leaf lower-port {
          type inet:port-number;
          must '. <= ../upper-port' {
            error-message
              "The lower-port must be less than or equal to
              the upper-port.";
          }
          mandatory true;
          description
            "Lower boundary for a port.";
        }
        leaf upper-port {
          type inet:port-number;
          mandatory true;
          description
            "Upper boundary for a port.";
        }
      }
    }
    case operator {
      leaf operator {
        type operator;
        default "eq";
        description
          "Operator to be applied on the port below.";
      }
      leaf port {
        type inet:port-number;
        mandatory true;
        description
          "Port number along with the operator on which to
          match.";
      }
    }
  }
}
```

```
    description
      "Choice of specifying a port range or a single
       port along with an operator.";
  }
  description
    "Grouping for port definitions in the form of a
     choice statement.";
}

grouping acl-ip-header-fields {
  description
    "IP header fields common to IPv4 and IPv6";
  reference
    "RFC 791: Internet Protocol.";

  leaf dscp {
    type inet:dscp;
    description
      "Differentiated Services Code Point.";
    reference
      "RFC 2474: Definition of the Differentiated Services
       Field (DS Field) in the IPv4 and IPv6
       Headers.";
  }

  leaf ecn {
    type uint8 {
      range "0..3";
    }
    description
      "Explicit Congestion Notification.";
    reference
      "RFC 3168: The Addition of Explicit Congestion
       Notification (ECN) to IP.";
  }

  leaf length {
    type uint16;
    description
      "In the IPv4 header field, this field is known as the Total
       Length. Total Length is the length of the datagram, measured
       in octets, including internet header and data.

       In the IPv6 header field, this field is known as the Payload
       Length, which is the length of the IPv6 payload, i.e., the rest
       of the packet following the IPv6 header, in octets.";
    reference
      "RFC 791: Internet Protocol
```

```
    RFC 8200: Internet Protocol, Version 6 (IPv6) Specification.";
  }
  leaf ttl {
    type uint8;
    description
      "This field indicates the maximum time the datagram is allowed
       to remain in the internet system.  If this field contains the
       value zero, then the datagram must be dropped.

       In IPv6, this field is known as the Hop Limit.";
    reference
      "RFC 791: Internet Protocol
       RFC 8200: Internet Protocol, Version 6 (IPv6) Specification.";
  }
  leaf protocol {
    type uint8;
    description
      "Internet Protocol number.  Refers to the protocol of the
       payload.  In IPv6, this field is known as 'next-header',
       and if extension headers are present, the protocol is
       present in the 'upper-layer' header.";
    reference
      "RFC 791: Internet Protocol
       RFC 8200: Internet Protocol, Version 6 (IPv6) Specification.";
  }
}

grouping acl-ipv4-header-fields {
  description
    "Fields in the IPv4 header.";
  leaf ihl {
    type uint8 {
      range "5..60";
    }
    description
      "In an IPv4 header field, the Internet Header Length (IHL) is
       the length of the internet header in 32-bit words and
       thus points to the beginning of the data.  Note that the
       minimum value for a correct header is 5.";
  }
  leaf flags {
    type bits {
      bit reserved {
        position 0;
        description
          "Reserved.  Must be zero.";
      }
      bit fragment {
```

```
        position 1;
        description
            "Setting the value to 0 indicates may fragment, while
             setting the value to 1 indicates do not fragment.";
    }
    bit more {
        position 2;
        description
            "Setting the value to 0 indicates this is the last fragment,
             and setting the value to 1 indicates more fragments are
             coming.";
    }
}
description
    "Bit definitions for the Flags field in the IPv4 header.";
}
leaf offset {
    type uint16 {
        range "20..65535";
    }
    description
        "The fragment offset is measured in units of 8 octets (64 bits).
         The first fragment has offset zero. The length is 13 bits";
}
leaf identification {
    type uint16;
    description
        "An identifying value assigned by the sender to aid in
         assembling the fragments of a datagram.";
}
}
choice destination-network {
    case destination-ipv4-network {
        leaf destination-ipv4-network {
            type inet:ipv4-prefix;
            description
                "Destination IPv4 address prefix.";
        }
    }
}
description
    "Choice of specifying a destination IPv4 address or
     referring to a group of IPv4 destination addresses.";
}
choice source-network {
    case source-ipv4-network {
        leaf source-ipv4-network {
            type inet:ipv4-prefix;
```

```
        description
            "Source IPv4 address prefix.";
    }
}
description
    "Choice of specifying a source IPv4 address or
    referring to a group of IPv4 source addresses.";
}
}

grouping acl-ipv6-header-fields {
    description
        "Fields in the IPv6 header.";

    choice destination-network {
        case destination-ipv6-network {
            leaf destination-ipv6-network {
                type inet:ipv6-prefix;
                description
                    "Destination IPv6 address prefix.";
            }
        }
        description
            "Choice of specifying a destination IPv6 address
            or referring to a group of IPv6 destination
            addresses.";
    }

    choice source-network {
        case source-ipv6-network {
            leaf source-ipv6-network {
                type inet:ipv6-prefix;
                description
                    "Source IPv6 address prefix.";
            }
        }
        description
            "Choice of specifying a source IPv6 address or
            referring to a group of IPv6 source addresses.";
    }

    leaf flow-label {
        type inet:ipv6-flow-label;
        description
            "IPv6 Flow label.";
    }

    reference
        "RFC 4291: IP Version 6 Addressing Architecture
```

```

    RFC 4007: IPv6 Scoped Address Architecture
    RFC 5952: A Recommendation for IPv6 Address Text
        Representation.";
}

grouping acl-eth-header-fields {
    description
        "Fields in the Ethernet header.";
    leaf destination-mac-address {
        type yang:mac-address;
        description
            "Destination IEEE 802 Media Access Control (MAC)
            address.";
    }
    leaf destination-mac-address-mask {
        type yang:mac-address;
        description
            "Destination IEEE 802 MAC address mask.";
    }
    leaf source-mac-address {
        type yang:mac-address;
        description
            "Source IEEE 802 MAC address.";
    }
    leaf source-mac-address-mask {
        type yang:mac-address;
        description
            "Source IEEE 802 MAC address mask.";
    }
    leaf ethertype {
        type eth:ethertype;
        description
            "The Ethernet Type (or Length) value represented
            in the canonical order defined by IEEE 802.
            The canonical representation uses lowercase
            characters.";
        reference
            "IEEE 802-2014, Clause 9.2.";
    }
    reference
        "IEEE 802: IEEE Standard for Local and Metropolitan
        Area Networks: Overview and Architecture.";
}

grouping acl-tcp-header-fields {
    description
        "Collection of TCP header fields that can be used to
        set up a match filter.";
}
```



```
leaf sequence-number {
  type uint32;
  description
    "Sequence number that appears in the packet.";
}
leaf acknowledgement-number {
  type uint32;
  description
    "The acknowledgement number that appears in the
    packet.";
}
leaf data-offset {
  type uint8 {
    range "5..15";
  }
  description
    "Specifies the size of the TCP header in 32-bit
    words. The minimum size header is 5 words and
    the maximum is 15 words; thus, this gives a
    minimum size of 20 bytes and a maximum of 60
    bytes, allowing for up to 40 bytes of options
    in the header.";
}
leaf reserved {
  type uint8;
  description
    "Reserved for future use.";
}
leaf flags {
  type bits {
    bit cwr {
      position 1;
      description
        "The Congestion Window Reduced (CWR) flag is set
        by the sending host to indicate that it received
        a TCP segment with the ECN-Echo (ECE) flag set
        and had responded in the congestion control
        mechanism.";
      reference
        "RFC 3168: The Addition of Explicit Congestion
        Notification (ECN) to IP.";
    }
    bit ece {
      position 2;
      description
        "ECN-Echo has a dual role, depending on the value
        of the SYN flag. It indicates the following: if
        the SYN flag is set (1), the TCP peer is ECN
```

```
        capable, and if the SYN flag is clear (0), a packet
        with the Congestion Experienced flag set (ECN=11)
        in the IP header was received during normal
        transmission (added to the header by RFC 3168).
        This serves as an indication of network congestion
        (or impending congestion) to the TCP sender.";
    reference
        "RFC 3168: The Addition of Explicit Congestion
        Notification (ECN) to IP.";
}
bit urg {
    position 3;
    description
        "Indicates that the Urgent Pointer field is significant.";
}
bit ack {
    position 4;
    description
        "Indicates that the Acknowledgement field is significant.
        All packets after the initial SYN packet sent by the
        client should have this flag set.";
}
bit psh {
    position 5;
    description
        "Push function. Asks to push the buffered data to the
        receiving application.";
}
bit rst {
    position 6;
    description
        "Reset the connection.";
}
bit syn {
    position 7;
    description
        "Synchronize sequence numbers. Only the first packet
        sent from each end should have this flag set. Some
        other flags and fields change meaning based on this
        flag, and some are only valid for when it is set,
        and others when it is clear.";
}
bit fin {
    position 8;
    description
        "Last package from the sender.";
}
}
```

```
    description
      "Also known as Control Bits.  Contains nine 1-bit flags.";
    reference
      "RFC 793: Transmission Control Protocol.";
  }
  leaf window-size {
    type uint16;
    units "bytes";
    description
      "The size of the receive window, which specifies
       the number of window size units beyond the segment
       identified by the sequence number in the Acknowledgement
       field that the sender of this segment is currently
       willing to receive.";
  }
  leaf urgent-pointer {
    type uint16;
    description
      "This field is an offset from the sequence number
       indicating the last urgent data byte.";
  }
  leaf options {
    type binary {
      length "1..40";
    }
    description
      "The length of this field is determined by the
       Data Offset field.  Options have up to three
       fields: Option-Kind (1 byte), Option-Length
       (1 byte), and Option-Data (variable).  The Option-Kind
       field indicates the type of option and is the
       only field that is not optional.  Depending on
       what kind of option we are dealing with,
       the next two fields may be set: the Option-Length
       field indicates the total length of the option,
       and the Option-Data field contains the value of
       the option, if applicable.";
  }
}

grouping acl-udp-header-fields {
  description
    "Collection of UDP header fields that can be used
     to set up a match filter.";
  leaf length {
    type uint16;
    description
      "A field that specifies the length in bytes of
```

the UDP header and UDP data. The minimum length is 8 bytes because that is the length of the header. The field size sets a theoretical limit of 65,535 bytes (8-byte header plus 65,527 bytes of data) for a UDP datagram. However, the actual limit for the data length, which is imposed by the underlying IPv4 protocol, is 65,507 bytes (65,535 minus 8-byte UDP header minus 20-byte IP header).

In IPv6 jumbograms, it is possible to have UDP packets of a size greater than 65,535 bytes. RFC 2675 specifies that the Length field is set to zero if the length of the UDP header plus UDP data is greater than 65,535.";

```
}
}
```

```
grouping acl-icmp-header-fields {
  description
    "Collection of ICMP header fields that can be
    used to set up a match filter.";
  leaf type {
    type uint8;
    description
      "Also known as control messages.";
    reference
      "RFC 792: Internet Control Message Protocol
      RFC 4443: Internet Control Message Protocol (ICMPv6)
      for Internet Protocol Version 6 (IPv6)
      Specification.";
  }
  leaf code {
    type uint8;
    description
      "ICMP subtype. Also known as control messages.";
    reference
      "RFC 792: Internet Control Message Protocol
      RFC 4443: Internet Control Message Protocol (ICMPv6)
      for Internet Protocol Version 6 (IPv6)
      Specification.";
  }
  leaf rest-of-header {
    type binary;
    description
      "Unbounded in length, the contents vary based on the
      ICMP type and code. Also referred to as 'Message Body'
      in ICMPv6.";
```

```

    reference
      "RFC 792: Internet Control Message Protocol
       RFC 4443: Internet Control Message Protocol (ICMPv6)
        for Internet Protocol Version 6 (IPv6)
        Specification.";
  }
}
}
<CODE ENDS>

```

4.3. ACL Examples

Requirement: Deny tcp traffic from 192.0.2.0/24, destined to 198.51.100.0/24.

Here is the ACL configuration xml for this Access Control List:

[note: '\\' line wrapping for formatting only]

```

<?xml version="1.0" encoding="UTF-8"?>
<config xmlns="urn:ietf:params:xml:ns:netconf:base:1.0">
  <acls
    xmlns="urn:ietf:params:xml:ns:yang:ietf-access-control-list">
    <acl>
      <name>sample-ipv4-acl</name>
      <type>ipv4-acl-type</type>
      <aces>
        <ace>
          <name>rule1</name>
          <matches>
            <ipv4>
              <protocol>6</protocol>
              <destination-ipv4-network>198.51.100.0/24</destination-
-ipv4-network>
              <source-ipv4-network>192.0.2.0/24</source-ipv4-network>
            </ipv4>
          </matches>
          <actions>
            <forwarding>drop</forwarding>
          </actions>
        </ace>
      </aces>
    </acl>
  </acls>
</config>

```

The ACL and ACEs can be described in the command-line interface (CLI) as the following:

```
acl ipv4 sample-ipv4-acl
deny tcp 192.0.2.0/24 198.51.100.0/24
```

Requirement: Accept all DNS traffic destined for 2001:db8::/32 on port 53.

[note: '\\' line wrapping for formatting only]

```
<?xml version="1.0" encoding="UTF-8"?>
<config xmlns="urn:ietf:params:xml:ns:netconf:base:1.0">
  <acls
    xmlns="urn:ietf:params:xml:ns:yang:ietf-access-control-list">
    <acl>
      <name>allow-dns-packets</name>
      <type>ipv6-acl-type</type>
      <aces>
        <ace>
          <name>rule1</name>
          <matches>
            <ipv6>
              <destination-ipv6-network>2001:db8::/32</destination-ipv6-network>
            </ipv6>
            <udp>
              <destination-port>
                <operator>eq</operator>
                <port>53</port>
              </destination-port>
            </udp>
          </matches>
          <actions>
            <forwarding>accept</forwarding>
          </actions>
        </ace>
      </aces>
    </acl>
  </acls>
</config>
```

4.4. Port Range Usage and Other Examples

When a lower-port and an upper-port are both present, it represents a range between the lower-port and upper-port with both the lower-port and upper-port included. When only a port is present, it represents a port, with the operator specifying the range.

The following XML example represents a configuration where TCP traffic from source ports 16384, 16385, 16386, and 16387 is dropped.

```
<?xml version="1.0" encoding="UTF-8"?>
<config xmlns="urn:ietf:params:xml:ns:netconf:base:1.0">
  <acls
    xmlns="urn:ietf:params:xml:ns:yang:ietf-access-control-list">
    <acl>
      <name>sample-port-acl</name>
      <type>ipv4-acl-type</type>
      <aces>
        <ace>
          <name>rule1</name>
          <matches>
            <tcp>
              <source-port>
                <lower-port>16384</lower-port>
                <upper-port>16387</upper-port>
              </source-port>
            </tcp>
          </matches>
          <actions>
            <forwarding>drop</forwarding>
          </actions>
        </ace>
      </aces>
    </acl>
  </acls>
</config>
```

The following XML example represents a configuration where all IPv4 ICMP echo requests are dropped.

```
<?xml version="1.0" encoding="UTF-8"?>
<config xmlns="urn:ietf:params:xml:ns:netconf:base:1.0">
  <acls
    xmlns="urn:ietf:params:xml:ns:yang:ietf-access-control-list">
    <acl>
      <name>sample-icmp-acl</name>
      <aces>
        <ace>
          <name>rule1</name>
          <matches>
            <ipv4>
              <protocol>1</protocol>
            </ipv4>
            <icmp>
              <type>8</type>
              <code>0</code>
            </icmp>
          </matches>
          <actions>
            <forwarding>drop</forwarding>
          </actions>
        </ace>
      </aces>
    </acl>
  </acls>
</config>
```


The following XML example represents a configuration of a single port, port 21, that accepts TCP traffic.

```
<?xml version="1.0" encoding="UTF-8"?>
<config xmlns="urn:ietf:params:xml:ns:netconf:base:1.0">
  <acls
    xmlns="urn:ietf:params:xml:ns:yang:ietf-access-control-list">
    <acl>
      <name>sample-ipv4-acl</name>
      <type>ipv4-acl-type</type>
      <aces>
        <ace>
          <name>rule1</name>
          <matches>
            <tcp>
              <destination-port>
                <operator>eq</operator>
                <port>21</port>
              </destination-port>
            </tcp>
          </matches>
          <actions>
            <forwarding>accept</forwarding>
          </actions>
        </ace>
      </aces>
    </acl>
  </acls>
</config>
```

The following XML example represents a configuration specifying all ports that are not equal to 21 that will drop TCP packets destined for those ports.

```
<?xml version="1.0" encoding="UTF-8"?>
<config xmlns="urn:ietf:params:xml:ns:netconf:base:1.0">
  <acls
    xmlns="urn:ietf:params:xml:ns:yang:ietf-access-control-list">
    <acl>
      <name>sample-ipv4-acl</name>
      <type>ipv4-acl-type</type>
      <aces>
        <ace>
          <name>rule1</name>
          <matches>
            <tcp>
              <destination-port>
                <operator>neq</operator>
                <port>21</port>
              </destination-port>
            </tcp>
          </matches>
          <actions>
            <forwarding>drop</forwarding>
          </actions>
        </ace>
      </aces>
    </acl>
  </acls>
</config>
```

5. Security Considerations

The YANG modules specified in this document define a schema for data that is designed to be accessed via network management protocol such as NETCONF [RFC6241] or RESTCONF [RFC8040]. The lowest NETCONF layer is the secure transport layer, and the mandatory-to-implement secure transport is Secure Shell (SSH) [RFC6242]. The lowest RESTCONF layer is HTTPS, and the mandatory-to-implement secure transport is TLS [RFC8446].

The NETCONF Access Control Model (NACM) [RFC8341] provides the means to restrict access for particular NETCONF or RESTCONF users to a preconfigured subset of all available NETCONF or RESTCONF protocol operations and content.

There are a number of data nodes defined in these YANG modules that are writable/creatable/deletable (i.e., config true, which is the default). These data nodes may be considered sensitive or vulnerable in some network environments. Write operations (e.g., edit-config) to these data nodes without proper protection can have a negative effect on network operations. These are the subtrees and data nodes and their sensitivity/vulnerability:

`/acls/acl/aces`: This list specifies all the configured access control entries on the device. Unauthorized write access to this list can allow intruders to modify the entries so as to permit traffic that should not be permitted, or deny traffic that should be permitted. The former may result in a DoS attack, or compromise the device. The latter may result in a DoS attack. The impact of an unauthorized read access of the list will allow the attacker to determine which rules are in effect, to better craft an attack.

`/acls/acl/aces/ace/actions/logging`: This node specifies ability to log packets that match this ace entry. Unauthorized write access to this node can allow intruders to enable logging on one or many ace entries, overwhelming the server in the process. Unauthorized read access of this node can allow intruders to access logging information, which could be used to craft an attack the server.

6. IANA Considerations

This document registers three URIs and three YANG modules.

6.1. URI Registration

This document registers three URIs in the "IETF XML Registry" [RFC3688] as follows:

URI: `urn:ietf:params:xml:ns:yang:ietf-access-control-list`
URI: `urn:ietf:params:xml:ns:yang:ietf-packet-fields`
URI: `urn:ietf:params:xml:ns:yang:ietf-ethertypes`

Registrant Contact: The IESG.

XML: N/A; the requested URI is an XML namespace.

6.2. YANG Module Name Registration

This document registers three YANG modules in the "YANG Module Names" registry [RFC6020].

Name: ietf-access-control-list
Namespace: urn:ietf:params:xml:ns:yang:ietf-access-control-list
Prefix: acl
Reference: RFC 8519

Name: ietf-packet-fields
Namespace: urn:ietf:params:xml:ns:yang:ietf-packet-fields
Prefix: packet-fields
Reference: RFC 8519

Name: ietf-ethertypes
Namespace: urn:ietf:params:xml:ns:yang:ietf-ethertypes
Prefix: ethertypes
Reference: RFC 8519

7. References

7.1. Normative References

- [RFC791] Postel, J., "Internet Protocol", STD 5, RFC 791, DOI 10.17487/RFC0791, September 1981, <<https://www.rfc-editor.org/info/rfc791>>.
- [RFC792] Postel, J., "Internet Control Message Protocol", STD 5, RFC 792, DOI 10.17487/RFC0792, September 1981, <<https://www.rfc-editor.org/info/rfc792>>.
- [RFC793] Postel, J., "Transmission Control Protocol", STD 7, RFC 793, DOI 10.17487/RFC0793, September 1981, <<https://www.rfc-editor.org/info/rfc793>>.
- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<https://www.rfc-editor.org/info/rfc2119>>.
- [RFC2474] Nichols, K., Blake, S., Baker, F., and D. Black, "Definition of the Differentiated Services Field (DS Field) in the IPv4 and IPv6 Headers", RFC 2474, DOI 10.17487/RFC2474, December 1998, <<https://www.rfc-editor.org/info/rfc2474>>.

- [RFC3168] Ramakrishnan, K., Floyd, S., and D. Black, "The Addition of Explicit Congestion Notification (ECN) to IP", RFC 3168, DOI 10.17487/RFC3168, September 2001, <<https://www.rfc-editor.org/info/rfc3168>>.
- [RFC4007] Deering, S., Haberman, B., Jinmei, T., Nordmark, E., and B. Zill, "IPv6 Scoped Address Architecture", RFC 4007, DOI 10.17487/RFC4007, March 2005, <<https://www.rfc-editor.org/info/rfc4007>>.
- [RFC4291] Hinden, R. and S. Deering, "IP Version 6 Addressing Architecture", RFC 4291, DOI 10.17487/RFC4291, February 2006, <<https://www.rfc-editor.org/info/rfc4291>>.
- [RFC5952] Kawamura, S. and M. Kawashima, "A Recommendation for IPv6 Address Text Representation", RFC 5952, DOI 10.17487/RFC5952, August 2010, <<https://www.rfc-editor.org/info/rfc5952>>.
- [RFC6991] Schoenwaelder, J., Ed., "Common YANG Data Types", RFC 6991, DOI 10.17487/RFC6991, July 2013, <<https://www.rfc-editor.org/info/rfc6991>>.
- [RFC7950] Bjorklund, M., Ed., "The YANG 1.1 Data Modeling Language", RFC 7950, DOI 10.17487/RFC7950, August 2016, <<https://www.rfc-editor.org/info/rfc7950>>.
- [RFC8174] Leiba, B., "Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words", BCP 14, RFC 8174, DOI 10.17487/RFC8174, May 2017, <<https://www.rfc-editor.org/info/rfc8174>>.
- [RFC8200] Deering, S. and R. Hinden, "Internet Protocol, Version 6 (IPv6) Specification", STD 86, RFC 8200, DOI 10.17487/RFC8200, July 2017, <<https://www.rfc-editor.org/info/rfc8200>>.
- [RFC8343] Bjorklund, M., "A YANG Data Model for Interface Management", RFC 8343, DOI 10.17487/RFC8343, March 2018, <<https://www.rfc-editor.org/info/rfc8343>>.

7.2. Informative References

- [RFC3688] Mealling, M., "The IETF XML Registry", BCP 81, RFC 3688, DOI 10.17487/RFC3688, January 2004, <<https://www.rfc-editor.org/info/rfc3688>>.
- [RFC6020] Bjorklund, M., Ed., "YANG - A Data Modeling Language for the Network Configuration Protocol (NETCONF)", RFC 6020, DOI 10.17487/RFC6020, October 2010, <<https://www.rfc-editor.org/info/rfc6020>>.
- [RFC6241] Enns, R., Ed., Bjorklund, M., Ed., Schoenwaelder, J., Ed., and A. Bierman, Ed., "Network Configuration Protocol (NETCONF)", RFC 6241, DOI 10.17487/RFC6241, June 2011, <<https://www.rfc-editor.org/info/rfc6241>>.
- [RFC6242] Wasserman, M., "Using the NETCONF Protocol over Secure Shell (SSH)", RFC 6242, DOI 10.17487/RFC6242, June 2011, <<https://www.rfc-editor.org/info/rfc6242>>.
- [RFC7011] Claise, B., Ed., Trammell, B., Ed., and P. Aitken, "Specification of the IP Flow Information Export (IPFIX) Protocol for the Exchange of Flow Information", STD 77, RFC 7011, DOI 10.17487/RFC7011, September 2013, <<https://www.rfc-editor.org/info/rfc7011>>.
- [RFC8040] Bierman, A., Bjorklund, M., and K. Watsen, "RESTCONF Protocol", RFC 8040, DOI 10.17487/RFC8040, January 2017, <<https://www.rfc-editor.org/info/rfc8040>>.
- [RFC8340] Bjorklund, M. and L. Berger, Ed., "YANG Tree Diagrams", BCP 215, RFC 8340, DOI 10.17487/RFC8340, March 2018, <<https://www.rfc-editor.org/info/rfc8340>>.
- [RFC8341] Bierman, A. and M. Bjorklund, "Network Configuration Access Control Model", STD 91, RFC 8341, DOI 10.17487/RFC8341, March 2018, <<https://www.rfc-editor.org/info/rfc8341>>.
- [RFC8446] Rescorla, E., "The Transport Layer Security (TLS) Protocol Version 1.3", RFC 8446, DOI 10.17487/RFC8446, August 2018, <<https://www.rfc-editor.org/info/rfc8446>>.

Appendix A. Extending ACL Model Examples

A.1. Example of a Company's Proprietary Module

The "example-newco-acl" module is an example of a company's proprietary model that augments the "ietf-acl" module. It shows how to use 'augment' with an XML Path Language (XPath) expression to add additional match criteria, actions, and default actions for when no ACE matches are found. All these are company proprietary extensions or system feature extensions. "example-newco-acl" is just an example, and it is expected that vendors will create their own proprietary models.

```
module example-newco-acl {  
  yang-version 1.1;  
  namespace "http://example.com/ns/example-newco-acl";  
  prefix example-newco-acl;  
  import ietf-access-control-list {  
    prefix acl;  
  }  
  organization  
    "Newco model group.";  
  contact  
    "abc@newco.com";  
  description  
    "This YANG module augments the IETF ACL YANG module.";  
  revision 2019-03-04 {  
    description  
      "Creating NewCo proprietary extensions to the ietf-acl  
      model.";  
    reference  
      "RFC 8519: YANG Data Model for Network Access Control  
      Lists (ACLs).";  
  }  
  augment "/acl:acls/acl:acl/"  
    + "acl:aces/acl:ace/"  
    + "acl:matches" {  
    description  
      "Newco proprietary simple filter matches.";  
  }
```

```
choice protocol-payload-choice {
  description
    "Newco proprietary payload match condition.";
  list protocol-payload {
    key "value-keyword";
    ordered-by user;
    description
      "Match protocol payload.";
    uses match-simple-payload-protocol-value;
  }
}

choice metadata {
  description
    "Newco proprietary interface match condition.";
  leaf packet-length {
    type uint16;
    description
      "Match on packet length.";
  }
}
}

augment "/acl:acls/acl:acl/"
  + "acl:aces/acl:ace/"
  + "acl:actions" {
  description
    "Newco proprietary simple filter actions.";
  choice action {
    description
      "Newco proprietary action choices.";
    case count {
      description
        "Count the packet in the named counter.";
      leaf count {
        type uint32;
        description
          "Count.";
      }
    }
    case policer {
      description
        "Name of policer used to rate-limit traffic.";
      leaf policer {
        type string;
        description
          "Name of the policer.";
      }
    }
  }
}
```



```

    }
    case hierarchical-policer {
      leaf hierarchical-policer {
        type string;
        description
          "Name of the hierarchical policer.";
      }
      description
        "Name of the hierarchical policer used to
        rate-limit traffic.";
    }
  }
}

augment "/acl:acls/acl:acl"
+ "/acl:aces/acl:ace/"
+ "acl:actions" {
  leaf default-action {
    type identityref {
      base acl:forwarding-action;
    }
    default "acl:drop";
    description
      "Actions that occur if no ACE is matched.";
  }
  description
    "Newco proprietary default action.";
}

grouping match-simple-payload-protocol-value {
  description
    "Newco proprietary payload";
  leaf value-keyword {
    type enumeration {
      enum icmp {
        description
          "Internet Control Message Protocol.";
      }
      enum icmp6 {
        description
          "Internet Control Message Protocol
          Version 6.";
      }
      enum range {
        description
          "Range of values.";
      }
    }
  }
}

```

```

        description
            "(null).";
    }
}
}

```

The following figure is the tree diagram of example-newco-acl. In this example, /ietf-acl:acls/ietf-acl:acl/ietf-acl:aces/ietf-acl:ace/ietf-acl:matches are augmented with two new choices: protocol-payload-choice and metadata. The protocol-payload-choice uses a grouping with an enumeration of all supported protocol values. Metadata matches apply to fields associated with the packet, that are not in the packet header, such as overall packet length. In another example, /ietf-acl:acls/ietf-acl:acl/ietf-acl:aces/ietf-acl:ace/ietf-acl:actions are augmented with a new choice of actions.

```

module: example-newco-acl
  augment /acl:acls/acl:acl/acl:aces/acl:ace/acl:matches:
    +--rw (protocol-payload-choice)?
      | +--:(protocol-payload)
      |   +--rw protocol-payload* [value-keyword]
      |   +--rw value-keyword      enumeration
    +--rw (metadata)?
      | +--:(packet-length)
      |   +--rw packet-length?      uint16
  augment /acl:acls/acl:acl/acl:aces/acl:ace/acl:actions:
    +--rw (action)?
      | +--:(count)
      |   +--rw count?              uint32
      | +--:(policer)
      |   +--rw policer?            string
      | +--:(hierarchical-policer)
      |   +--rw hierarchical-policer? string
  augment /acl:acls/acl:acl/acl:aces/acl:ace/acl:actions:
    +--rw default-action? identityref

```

A.2. Linux nftables

As the Linux platform is becoming more popular than the networking platform, the Linux data model is changing. Previously, ACLs in Linux were highly protocol specific, and different utilities were used (iptables, ip6tables, arptables, and ebtables), so each one had a separate data model. Recently, this has changed, and a single utility, nftables, has been developed. With a single application, it has a single data model for firewall filters, and it follows very similarly the ietf-access-control list module proposed in this document. The nftables support input and output ACEs, and each ACE can be defined with match and action.

The example in Section 4.3 can be configured using nftable tool as below.

```
nft add table ip filter
nft add chain filter input
nft add rule ip filter input ip protocol tcp ip saddr \
    192.0.2.1/24 drop
```

The configuration entries added in nftable would be:

```
table ip filter {
    chain input {
        ip protocol tcp ip saddr 192.0.2.1/24 drop
    }
}
```

We can see that there are many similarities between Linux nftables and IETF ACL YANG data models and their extension models. It should be fairly easy to do translation between the ACL YANG model described in this document and Linux nftables.

A.3. Ethertypes

The ACL module is dependent on the definition of Ethertypes. IEEE owns the allocation of those Ethertypes. This model is being included here to enable the definition of those types till such time that IEEE takes up the task of publication of the model that defines those Ethertypes. At that time, this model can be deprecated.

<CODE BEGINS> file "ietf-ethertypes@2019-03-04.yang"

```
module ietf-ethertypes {
    namespace "urn:ietf:params:xml:ns:yang:ietf-ethertypes";
    prefix ethertypes;

    organization
        "IETF NETMOD (Network Modeling) Working Group.";

    contact
        "WG Web:    <https://datatracker.ietf.org/wg/netmod/>
        WG List:    <mailto:netmod@ietf.org>

        Editor:    Mahesh Jethanandani
                   <mjethanandani@gmail.com>";

    description
        "This module contains common definitions for the
```

Ethertype used by different modules. It is a placeholder module, till such time that IEEE starts a project to define these Ethertypes and publishes a standard.

At that time, this module can be deprecated.

Copyright (c) 2019 IETF Trust and the persons identified as the document authors. All rights reserved.

Redistribution and use in source and binary forms, with or without modification, is permitted pursuant to, and subject to the license terms contained in, the Simplified BSD License set forth in Section 4.c of the IETF Trust's Legal Provisions Relating to IETF Documents (<http://trustee.ietf.org/license-info>).

This version of this YANG module is part of RFC 8519; see the RFC itself for full legal notices.";

```
revision 2019-03-04 {
  description
    "Initial revision.";
  reference
    "RFC 8519: YANG Data Model for Network Access Control
      Lists (ACLs).";
}

typedef ethertype {
  type union {
    type uint16;
    type enumeration {
      enum ipv4 {
        value 2048;
        description
          "Internet Protocol version 4 (IPv4) with a
            hex value of 0x0800.";
        reference
          "RFC 791: Internet Protocol.";
      }
    }
    enum arp {
      value 2054;
      description
        "Address Resolution Protocol (ARP) with a
          hex value of 0x0806.";
      reference
        "RFC 826: An Ethernet Address Resolution Protocol: Or
          Converting Network Protocol Addresses to 48.bit
```

```
        Ethernet Address for Transmission on Ethernet
        Hardware.";
    }
    enum wlan {
        value 2114;
        description
            "Wake-on-LAN.  Hex value of 0x0842.";
    }
    enum trill {
        value 8947;
        description
            "Transparent Interconnection of Lots of Links.
            Hex value of 0x22F3.";
        reference
            "RFC 6325: Routing Bridges (RBridges): Base Protocol
            Specification.";
    }
    enum srp {
        value 8938;
        description
            "Stream Reservation Protocol.  Hex value of
            0x22EA.";
        reference
            "IEEE 801.1Q-2011.";
    }
    enum decnet {
        value 24579;
        description
            "DECnet Phase IV.  Hex value of 0x6003.";
    }
    enum rarp {
        value 32821;
        description
            "Reverse Address Resolution Protocol.
            Hex value 0x8035.";
        reference
            "RFC 903: A Reverse Address Resolution Protocol.";
    }
    enum appletalk {
        value 32923;
        description
            "Appletalk (Ethertalk).  Hex value of 0x809B.";
    }
    enum aarp {
        value 33011;
        description
            "Appletalk Address Resolution Protocol.  Hex value
            of 0x80F3.";
```

```
}
enum vlan {
  value 33024;
  description
    "VLAN-tagged frame (IEEE 802.1Q) and Shortest Path
    Bridging IEEE 802.1aq with Network-Network
    Interface (NNI) compatibility. Hex value of
    0x8100.";
  reference
    "IEEE 802.1Q.";
}
enum ipx {
  value 33079;
  description
    "Internetwork Packet Exchange (IPX). Hex value
    of 0x8137.";
}
enum qnx {
  value 33284;
  description
    "QNX Qnet. Hex value of 0x8204.";
}
enum ipv6 {
  value 34525;
  description
    "Internet Protocol Version 6 (IPv6). Hex value
    of 0x86DD.";
  reference
    "RFC 8200: Internet Protocol, Version 6 (IPv6)
    Specification
    RFC 8201: Path MTU Discovery for IP version 6.";
}
enum efc {
  value 34824;
  description
    "Ethernet flow control using pause frames.
    Hex value of 0x8808.";
  reference
    "IEEE 802.1Qbb.";
}
enum esp {
  value 34825;
  description
    "Ethernet Slow Protocol. Hex value of 0x8809.";
  reference
    "IEEE 802.3-2015.";
}
enum cobranet {
```

```
    value 34841;
    description
        "CobraNet. Hex value of 0x8819.";
}
enum mpls-unicast {
    value 34887;
    description
        "Multiprotocol Label Switching (MPLS) unicast traffic.
        Hex value of 0x8847.";
    reference
        "RFC 3031: Multiprotocol Label Switching Architecture.";
}
enum mpls-multicast {
    value 34888;
    description
        "MPLS multicast traffic. Hex value of 0x8848.";
    reference
        "RFC 3031: Multiprotocol Label Switching Architecture.";
}
enum pppoe-discovery {
    value 34915;
    description
        "Point-to-Point Protocol over Ethernet. Used during
        the discovery process. Hex value of 0x8863.";
    reference
        "RFC 2516: A Method for Transmitting PPP Over Ethernet
        (PPPoE).";
}
enum pppoe-session {
    value 34916;
    description
        "Point-to-Point Protocol over Ethernet. Used during
        session stage. Hex value of 0x8864.";
    reference
        "RFC 2516: A Method for Transmitting PPP Over Ethernet
        (PPPoE).";
}
enum intel-ans {
    value 34925;
    description
        "Intel Advanced Networking Services. Hex value of
        0x886D.";
}
enum jumbo-frames {
    value 34928;
    description
        "Jumbo frames or Ethernet frames with more than
        1500 bytes of payload, up to 9000 bytes.";
```

```
}
enum homeplug {
  value 34939;
  description
    "Family name for the various power line
    communications. Hex value of 0x887B.";
}
enum eap {
  value 34958;
  description
    "Ethernet Access Protocol (EAP) over LAN. Hex value
    of 0x888E.";
  reference
    "IEEE 802.1X.";
}
enum profinet {
  value 34962;
  description
    "PROcess FIeld Net (PROFINET). Hex value of 0x8892.";
}
enum hyperscsi {
  value 34970;
  description
    "Small Computer System Interface (SCSI) over Ethernet.
    Hex value of 0x889A.";
}
enum aoe {
  value 34978;
  description
    "Advanced Technology Advancement (ATA) over Ethernet.
    Hex value of 0x88A2.";
}
enum ethercat {
  value 34980;
  description
    "Ethernet for Control Automation Technology (EtherCAT).
    Hex value of 0x88A4.";
}
enum provider-bridging {
  value 34984;
  description
    "Provider Bridging (802.1ad) and Shortest Path Bridging
    (801.1aq). Hex value of 0x88A8.";
  reference
    "IEEE 802.1ad and IEEE 802.1aq.";
}
enum ethernet-powerlink {
  value 34987;
```



```
    description
      "Ethernet Powerlink.  Hex value of 0x88AB.";
  }
  enum goose {
    value 35000;
    description
      "Generic Object Oriented Substation Event (GOOSE).
       Hex value of 0x88B8.";
    reference
      "IEC/ISO 8802-2 and 8802-3.";
  }
  enum gse {
    value 35001;
    description
      "Generic Substation Events.  Hex value of 88B9.";
    reference
      "IEC 61850.";
  }
  enum sv {
    value 35002;
    description
      "Sampled Value Transmission.  Hex value of 0x88BA.";
    reference
      "IEC 61850.";
  }
  enum lldp {
    value 35020;
    description
      "Link Layer Discovery Protocol (LLDP).  Hex value of
       0x88CC.";
    reference
      "IEEE 802.1AB.";
  }
  enum sercos {
    value 35021;
    description
      "Sercos Interface.  Hex value of 0x88CD.";
  }
  enum wsmpp {
    value 35036;
    description
      "WAVE Short Message Protocol (WSMP).  Hex value of
       0x88DC.";
  }
  enum homeplug-av-mme {
    value 35041;
    description
      "HomePlug AV Mobile Management Entity (MME).  Hex value
```

```
        of 88E1.";
    }
    enum mrp {
        value 35043;
        description
            "Media Redundancy Protocol (MRP). Hex value of
            0x88E3.";
        reference
            "IEC 62439-2.";
    }
    enum macsec {
        value 35045;
        description
            "MAC Security. Hex value of 0x88E5.";
        reference
            "IEEE 802.1AE.";
    }
    enum pbb {
        value 35047;
        description
            "Provider Backbone Bridges (PBB). Hex value of
            0x88E7.";
        reference
            "IEEE 802.1ah.";
    }
    enum cfm {
        value 35074;
        description
            "Connectivity Fault Management (CFM). Hex value of
            0x8902.";
        reference
            "IEEE 802.1ag.";
    }
    enum fcoe {
        value 35078;
        description
            "Fiber Channel over Ethernet (FCoE). Hex value of
            0x8906.";
        reference
            "T11 FC-BB-5.";
    }
    enum fcoe-ip {
        value 35092;
        description
            "FCoE Initialization Protocol. Hex value of 0x8914.";
    }
    enum roce {
        value 35093;
```

```
        description
          "RDMA over Converged Ethernet (RoCE).  Hex value of
           0x8915.";
      }
      enum tte {
        value 35101;
        description
          "TTEthernet Protocol Control Frame (TTE).  Hex value
           of 0x891D.";
        reference
          "SAE AS6802.";
      }
      enum hsr {
        value 35119;
        description
          "High-availability Seamless Redundancy (HSR).  Hex
           value of 0x892F.";
        reference
          "IEC 62439-3:2016.";
      }
    }
  }
  description
    "The uint16 type placeholder is defined to enable
     users to manage their own ethertypes not
     covered by the module.  Otherwise, the module contains
     enum definitions for the more commonly used ethertypes.";
}
}
<CODE ENDS>
```

Acknowledgements

Alex Clemm, Andy Bierman, and Lisa Huang started by sketching an initial draft version in several past IETF meetings. That document included an ACL YANG model structure and a rich set of match filters, and it acknowledged contributions by Louis Fourie, Dana Blair, Tula Kraiser, Patrick Gili, George Serpa, Martin Bjorklund, Kent Watsen, and Phil Shafer. Many people have reviewed the various earlier draft versions that made the document that went into IETF charter.

Dean Bogdanovic, Kiran Agrahara Sreenivasa, Lisa Huang, and Dana Blair each evaluated the YANG model in earlier draft versions separately, and then they worked together to create an ACL draft version that was supported by different vendors. That document removed vendor-specific features and gave examples that allowed vendors to extend their own proprietary ACLs. That earlier draft version was superseded with this document and received participation from many vendors.

The authors would like to thank Jason Sterne, Lada Lhotka, Juergen Schoenwalder, David Bannister, Jeff Haas, Kristian Larsson, and Einar Nilsen-Nygaard for their reviews of and suggestions for the document.

Authors' Addresses

Mahesh Jethanandani
VMware

Email: mjethanandani@gmail.com

Sonal Agarwal
Cisco Systems, Inc.

Email: sagarwal12@gmail.com

Lisa Huang

Email: huangyi_99@yahoo.com

Dana Blair

Email: dana@blairhome.com