

Internet Engineering Task Force (IETF)
Request for Comments: 8896
Category: Standards Track
ISSN: 2070-1721

S. Randriamasy
Nokia Bell Labs
Y. Yang
Yale University
Q. Wu
Huawei
L. Deng
China Mobile
N. Schwan
Thales Deutschland
November 2020

Application-Layer Traffic Optimization (ALTO) Cost Calendar

Abstract

This document is an extension to the base Application-Layer Traffic Optimization (ALTO) protocol. It extends the ALTO cost information service so that applications decide not only 'where' to connect but also 'when'. This is useful for applications that need to perform bulk data transfer and would like to schedule these transfers during an off-peak hour, for example. This extension introduces the ALTO Cost Calendar with which an ALTO Server exposes ALTO cost values in JSON arrays where each value corresponds to a given time interval. The time intervals, as well as other Calendar attributes, are specified in the Information Resources Directory and ALTO Server responses.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 7841.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <https://www.rfc-editor.org/info/rfc8896>.

Copyright Notice

Copyright (c) 2020 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of

described in the Simplified BSD License.

Table of Contents

1. Introduction
 - 1.1. Some Recent Known Uses
 - 1.2. Terminology
 2. Requirements Language
 3. Overview of ALTO Cost Calendars and Terminology
 - 3.1. ALTO Cost Calendar Overview
 - 3.2. ALTO Cost Calendar Information Features
 - 3.3. ALTO Calendar Design Characteristics
 - 3.3.1. ALTO Cost Calendar for All Cost Modes
 - 3.3.2. Compatibility with Legacy ALTO Clients
 4. ALTO Calendar Specification: IRD Extensions
 - 4.1. Calendar Attributes in the IRD Resource Capabilities
 - 4.2. Calendars in a Delegate IRD
 - 4.3. Example IRD with ALTO Cost Calendars
 5. ALTO Calendar Specification: Service Information Resources
 - 5.1. Calendar Extensions for Filtered Cost Maps (FCM)
 - 5.1.1. Calendar Extensions in Filtered Cost Map Requests
 - 5.1.2. Calendar Extensions in Filtered Cost Map Responses
 - 5.1.3. Use Case and Example: FCM with a Bandwidth Calendar
 - 5.2. Calendar Extensions in the Endpoint Cost Service
 - 5.2.1. Calendar-Specific Input in Endpoint Cost Requests
 - 5.2.2. Calendar Attributes in the Endpoint Cost Response
 - 5.2.3. Use Case and Example: ECS with a routingcost Calendar
 - 5.2.4. Use Case and Example: ECS with a Multi-cost Calendar for routingcost and owdelay
 6. IANA Considerations
 7. Security Considerations
 8. Operational Considerations
 9. References
 - 9.1. Normative References
 - 9.2. Informative References
- Acknowledgments
- Authors' Addresses

1. Introduction

The base Application-Layer Traffic Optimization (ALTO) protocol specified in [RFC7285] provides guidance to overlay applications that need to select one or several hosts from a set of candidates able to provide a desired resource. This guidance is based on parameters that affect performance and efficiency of the data transmission between the hosts, such as the topological distance. The goal of ALTO is to improve the Quality of Experience (QoE) in the application while optimizing resource usage in the underlying network infrastructure.

The ALTO protocol in [RFC7285] specifies a network map that defines groupings of endpoints in provider-defined network regions identified by Provider-defined Identifiers (PIDs). The Cost Map Service, Endpoint Cost Service (ECS), and Endpoint Ranking Service then provide ISP-defined costs and rankings for connections among the specified endpoints and PIDs and thus incentives for application

clients to connect to ISP-preferred locations, for instance, to reduce their costs. For the reasons outlined in the ALTO problem statement [RFC5693] and requirement AR-14 of [RFC6708], ALTO does not disseminate network metrics that change frequently. In a network, the costs can fluctuate for many reasons having to do with instantaneous traffic load or diurnal patterns of traffic demand or planned events, such as network maintenance, holidays, or highly publicized events. Thus, an ALTO application wishing to use the Cost Map and Endpoint Cost Service at some future time will have to estimate the state of the network at that time, a process that is, at best, fragile and brittle, since the application does not have any visibility into the state of the network. Providing network costs for only the current time thus may not be sufficient, in particular for applications that can schedule their traffic in a span of time, for example, by deferring backups or other background traffic to off-peak hours.

In case the ALTO cost value changes are predictable over a certain period of time and the application does not require immediate data transfer, it can save time to get the whole set of cost values over this period in one single ALTO response. Using this set to schedule data transfers allows optimizing the network resources usage and QoE. ALTO Clients and Servers can also minimize their workload by reducing and accordingly scheduling their data exchanges.

This document extends [RFC7285] to allow an ALTO Server to provide network costs for a given duration of time. A sequence of network costs across a time span for a given pair of network locations is named an "ALTO Cost Calendar". The Filtered Cost Map Service and Endpoint Cost Service are extended to provide Cost Calendars. In addition to this functional ALTO enhancement, we expect to further save network and storage resources by gathering multiple cost values for one cost type into one single ALTO Server response.

In this document, an "ALTO Cost Calendar" is specified in terms of information resource capabilities that are applicable to time-sensitive ALTO metrics. An ALTO Cost Calendar exposes ALTO cost values in JSON arrays, see [RFC8259], where each value corresponds to a given time interval. The time intervals, as well as other Calendar attributes, are specified in the Information Resources Directory (IRD) and in the Server response to allow the ALTO Client to interpret the received ALTO values. Last, the extensions for ALTO Calendars are applicable to any cost mode, and they ensure backwards compatibility with legacy ALTO Clients -- those that only support [RFC7285].

In the rest of this document, Section 3 provides the design characteristics. Sections 4 and 5 define the formal specifications for the IRD and the information resources. IANA, security considerations, and operational considerations are addressed respectively in Sections 6, 7, and 8.

1.1. Some Recent Known Uses

A potential use case is implementing smart network services that allow applications to dynamically build end-to-end, virtual networks

to satisfy given demands with no manual intervention. For example, data-transfer automation applications may need a network service to determine the availability of bandwidth resources to decide when to transfer their data sets. The SENSE project [SENSE] supports such applications by requiring that a network provides services such as the Time-Bandwidth-Product (TBP) service, which informs applications of bandwidth availability during a specific time period. ALTO Calendars can support this service if the Calendar start date and duration cover the period of interest of the requesting application.

The need of future scheduling of large-scale traffic that can be addressed by the ALTO protocol is also motivated by Unicorn, a unified resource orchestration framework for multi-domain, geo-distributed data analytics, see [UNICORN-FGCS].

1.2. Terminology

ALTO transaction:

A request/response exchange between an ALTO Client and an ALTO Server.

Client:

When used with a capital "C", this term refers to an ALTO Client.

Calendar, Cost Calendar, ALTO Calendar:

When used with capitalized words, these terms refer to an ALTO Cost Calendar.

Calendared:

This adjective qualifies information resources providing Cost Calendars and information on costs that are provided in the form of a Cost Calendar.

Endpoint (EP):

An endpoint is defined as in Section 2.1 of [RFC7285]. It can be, for example, a peer, a CDN storage location, a physical server involved in a virtual server-supported application, a party in a resource-sharing swarm such as a computation grid, or an online multi-party game.

ECM:

An abbreviation for Endpoint Cost Map.

FCM:

An abbreviation for Filtered Cost Map.

Server:

When used with a capital "S", this term refers to an ALTO Server.

2. Requirements Language

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in BCP 14 [RFC2119] [RFC8174] when, and only when, they appear in all capitals, as shown here.

When the words appear in lower case, they are to be interpreted with their natural language meanings.

3. Overview of ALTO Cost Calendars and Terminology

This section gives a high-level overview of the design. It assumes the reader is familiar with the ALTO protocol [RFC7285] and its Multi-Cost ALTO extension [RFC8189].

3.1. ALTO Cost Calendar Overview

An ALTO Cost Calendar provided by the ALTO Server provides 2 information items:

- * an array of values for a given metric, where each value specifies the metric corresponding to a time interval, where the value array can sometimes be a cyclic pattern that repeats a certain number of times and
- * attributes describing the time scope of the Calendar, including the size and number of the intervals and the date of the starting point of the Calendar, allowing an ALTO Client to interpret the values properly.

An ALTO Cost Calendar can be used like a "time table" to figure out the best time to schedule data transfers and also to proactively manage application traffic given predictable events, such as an expected spike in traffic due to crowd gathering (concerts, sports, etc.), traffic-intensive holidays, and network maintenance. A Calendar may be viewed as a synthetic abstraction of, for example, real measurements gathered over previous periods on which statistics have been computed. However, like for any schedule, unexpected network incidents may require the current ALTO Calendar to be updated and resent to the ALTO Clients needing it. The "ALTO Incremental Updates Using Server-Sent Events (SSE)" Service [RFC8895] can be used to directly update the Calendar upon value changes if supported by both the Server and the Client.

Most likely, the ALTO Cost Calendar would be used for the Endpoint Cost Service, assuming that a limited set of feasible endpoints for a non-real time application is already identified, and that those endpoints do not need to be accessed immediately and that their access can be scheduled within a given time period. The Filtered Cost Map Service is also applicable as long as the size of the Map allows it.

3.2. ALTO Cost Calendar Information Features

The Calendar attributes are provided in the Information Resources Directory (IRD) and in ALTO Server responses. The IRD announces attributes without date values in its information resources capabilities, whereas attributes with time-dependent values are provided in the "meta" section of Server responses. The ALTO Cost Calendar attributes provide the following information:

- * attributes to describe the time scope of the Calendar value array:
 - "time-interval-size": the applicable time interval size for each Calendar value, defined in seconds, that can cover a wide range of values.
 - "number-of-intervals": the number of intervals provided in the Calendar.
- * "calendar-start-time": specifying when the Calendar starts; that is, to which date the first value of the Cost Calendar is applicable.
- * "repeated": an optional attribute indicating how many iterations of the provided Calendar will have the same values. The Server may use it to allow the Client to schedule its next request and thus save its own workload by reducing processing of similar requests.

Attribute "repeated" may take a very high value if a Calendar represents a cyclic value pattern that the Server considers valid for a long period. In this case, the Server will only update the Calendar values once this period has elapsed or if an unexpected event occurs on the network. See Section 8 for more discussion.

3.3. ALTO Calendar Design Characteristics

The present document uses the notations defined in "Notation" (Section 8.2 of [RFC7285]).

The extensions in this document encode requests and responses using JSON [RFC8259].

In the base protocol [RFC7285], an ALTO cost is specified as a generic JSONValue [RFC8259] to allow extensions. However, that section (Section 11.2.3.6 of [RFC7285]) states:

An implementation of the protocol in this document SHOULD assume that the cost is a JSONNumber and fail to parse if it is not, unless the implementation is using an extension to this document that indicates when and how costs of other data types are signaled.

The present document extends the definition of a legacy cost map given in [RFC7285] to allow a cost entry to be an array of values, with one value per time interval, instead of being just one number, when the Cost Calendar functionality is activated on this cost. Therefore, the implementor of this extension MUST consider that a cost entry is an array of values if this cost has been queried as a Calendar.

Specifically, an implementation of this extension MUST parse the "number-of-intervals" attribute of the Calendar attributes in an IRD entry announcing a service providing a Cost Calendar for a given cost type. The implementation then will know that a cost entry of the service will be an array of values, and the expected size of the

array is that specified by the "number-of-intervals" attribute. The following rules attempt to ensure consistency between the array size announced by the Server and the actual size of the array received by the Client:

- * The size of the array of values conveyed in a Cost Calendar and received by the Client MUST be equal to the value of attribute "number-of-intervals" indicated in the IRD for the requested cost type.
- * When the size of the array received by the Client is different from the expected size, the Client SHOULD ignore the received array.

To realize an ALTO Calendar, this document extends the IRD and the ALTO requests and responses for Cost Calendars.

This extension is designed to be lightweight and to ensure backwards compatibility with base protocol ALTO Clients and with other extensions. It relies on "Parsing of Unknown Fields" (Section 8.3.7 of [RFC7285]), which states: "Extensions may include additional fields within JSON objects defined in this document. ALTO implementations MUST ignore unknown fields when processing ALTO messages."

The Calendar-specific capabilities are integrated in the information resources of the IRD and in the "meta" member of ALTO responses to Cost Calendars requests. A Calendar and its capabilities are associated with a given information resource and within this information resource with a given cost type. This design has several advantages:

- * it does not introduce a new mode,
- * it does not introduce new media types, and
- * it allows an ALTO Server to offer, for a cost type, different Calendars with attributes that are specific to the information resources providing a Calendar for this cost type, instead of being globally specific to the cost type.

The applicable Calendared information resources are:

- * the Filtered Cost Map and
- * the Endpoint Cost Map.

The ALTO Server can choose in which frequency it provides cost Calendars to ALTO Clients. It may either provide Calendar updates starting at the request date or carefully schedule its updates so as to take profit from a potential repetition/periodicity of Calendar values.

Since Calendar attributes are specific to an information resource, a Server may adapt the granularity of the calendared information so as to moderate the volume of exchanged data. For example, suppose a

Server provides a Calendar for cost type name "routingcost". The Server may offer a Calendar in a Cost Map resource, which may be a voluminous resource, as an array of 6 intervals lasting each 4 hours. It may also offer a Calendar in an Endpoint Cost Map resource, which is potentially less voluminous, as a finer-grained array of 24 intervals lasting 1 hour each.

The ALTO Server does not support constraints on Calendars, provided Calendars are requested for numerical values, for two main reasons:

- * Constraints on an array of values may be various. For instance, some Clients may refuse Calendars with one single value violating a constraint, whereas other ones may tolerate Calendars with values violating constraints, for example, at given times. Therefore, expressing constraints in a way that covers all possible Client preferences is challenging.
- * If constraints were to be supported, the processing overhead would be substantial for the Server as it would have to parse all the values of the Calendar array before returning a response.

As providing the constraint functionality in conjunction with the Calendar functionality is not feasible for the reasons described above, the two features are mutually exclusive. The absence of constraints on Filtered Cost Map and Endpoint Cost Map Calendars reflects a divergence from the non-calendared information resources defined in [RFC7285] and extended in [RFC8189], which support optional constraints.

3.3.1. ALTO Cost Calendar for All Cost Modes

An ALTO Cost Calendar is well suited for values encoded in the "numerical" mode. Actually, a Calendar can also represent metrics in other modes considered as compatible with time-varying values. For example, types of cost values (such as JSONBool) can also be calendared (as their value may be 'true' or 'false' depending on given time periods or likewise) values represented by strings, such as "medium", "high", "low", "blue", and "open".

Note also that a Calendar is suitable as well for time-varying metrics provided in the "ordinal" mode if these values are time-varying and the ALTO Server provides updates of cost-value-based preferences.

3.3.2. Compatibility with Legacy ALTO Clients

The ALTO protocol extensions for Cost Calendars have been defined so as to ensure that Calendar-capable ALTO Servers can provide legacy ALTO Clients with legacy information resources as well. That is, a legacy ALTO Client can request resources and receive responses as specified in [RFC7285].

A Calendar-aware ALTO Server MUST implement the base protocol specified in [RFC7285].

A Calendar-aware ALTO Client MUST implement the base protocol

specified in [RFC7285].

As a consequence, when a metric is available as a Calendar array, it also **MUST** be available as a single value, as required by [RFC7285]. The Server, in this case, provides the current value of the metric to either Calendar-aware Clients not interested in future or time-based values or Clients implementing [RFC7285] only.

For compatibility with legacy ALTO Clients specified in [RFC7285], calendared information resources are not applicable for full cost maps for the following reason: a legacy ALTO Client would receive a calendared cost map via an HTTP 'GET' command. As specified in Section 8.3.7 of [RFC7285], it will ignore the Calendar attributes indicated in the "meta" of the responses. Therefore, lacking information on Calendar attributes, it will not be able to correctly interpret and process the values of the received array of Calendar cost values.

Therefore, calendared information resources **MUST** be requested via the Filtered Cost Map Service or the Endpoint Cost Service using a POST method.

4. ALTO Calendar Specification: IRD Extensions

The Calendar attributes in the IRD information resources capabilities carry dateless values. A Calendar is associated with an information resource rather than a cost type. For example, a Server can provide a "routingcost" Calendar for the Filtered Cost Map Service at a granularity of one day and a "routingcost" Calendar for the Endpoint Cost Service at a finer granularity but for a limited number of endpoints. An example IRD with Calendar-specific features is provided in Section 4.3.

4.1. Calendar Attributes in the IRD Resource Capabilities

A Cost Calendar for a given cost type **MUST** be indicated in the IRD by an object of type CalendarAttributes. A CalendarAttributes object is represented by the "calendar-attributes" member of a resource entry. Member "calendar-attributes" is an array of CalendarAttributes objects. Each CalendarAttributes object lists a set of one or more cost types it applies to. A cost type name **MUST NOT** appear more than once in the "calendar-attributes" member of a resource entry; multiple appearances of a cost type name in the CalendarAttributes object of the "calendar-attributes" member **MUST** cause the ALTO Client to ignore any occurrences of this name beyond the first encountered occurrence. The Client **SHOULD** consider the CalendarAttributes object in the array containing the first encountered occurrence of a cost type as the valid one for this cost type. As an alternative, the Client may want to avoid the risks of erroneous guidance associated to the use of potentially invalid Calendar values. In this case, the Client **MAY** ignore the totality of occurrences of CalendarAttributes objects containing the cost type name and query the cost type using [RFC7285].

The encoding format for object CalendarAttributes using JSON [RFC8259] is as follows:

CalendarAttributes calendar-attributes <1..*>;

```
object{
  JSONString cost-type-names <1..*>;
  JSONNumber time-interval-size;
  JSONNumber number-of-intervals;
} CalendarAttributes;
```

"cost-type-names":

An array of one or more elements indicating the cost type names in the IRD entry to which the values of "time-interval-size" and "number-of-intervals" apply.

"time-interval-size":

The duration of an ALTO Calendar time interval in a unit of seconds. A "time-interval-size" value contains a non-negative JSONNumber. Example values are 300 and 7200, meaning that each Calendar value applies on a time interval that lasts 5 minutes and 2 hours, respectively. Since an interval size (e.g., 100 ms) can be smaller than the unit, the value specified may be a floating point (e.g., 0.1). Both ALTO Clients and Servers should be aware of potential precision issues caused by using floating point numbers; for example, the floating number 0.1 cannot be represented precisely using a finite number of binary bits. To improve interoperability and be consistent with [RFC7285] on the use of floating point numbers, the Server and the Client SHOULD use IEEE 754 double-precision floating point [IEEE.754.2019] to store this value.

"number-of-intervals":

A strictly positive integer (greater or equal to 1) that indicates the number of values of the Cost Calendar array.

- * An ALTO Server SHOULD specify the "time-interval-size" in the IRD as the smallest it is able to provide. A Client that needs a longer interval can aggregate multiple cost values to obtain it.
- * Attribute "cost-type-names" is associated with "time-interval-size" and "number-of-intervals", because multiple cost types may share the same values for attributes "time-interval-size" and "number-of-intervals". To avoid redundancies, cost type names sharing the same values for "time-interval-size" and "number-of-intervals" are grouped in the "cost-type-names" attribute. In the example IRD provided in Section 4.3, the information resource "filtered-cost-map-calendar" provides a Calendar for cost type names "num-routingcost", "num-throughputrating", and "string-servicestatus". Cost type names "num-routingcost" and "num-throughputrating" are grouped in the "cost-type-names" attribute because they share the same values for "time-interval-size" and "number-of-intervals", which are respectively 7200 and 12.
- * Multiplying "time-interval-size" by "number-of-intervals" provides the duration of the provided Calendar. For example, an ALTO Server may provide a Calendar for ALTO values changing every "time-interval-size" equal to 5 minutes. If "number-of-intervals"

has the value 12, then the duration of the provided Calendar is 1 hour.

4.2. Calendars in a Delegate IRD

It may be useful to distinguish IRD resources supported by the base ALTO protocol from resources supported by its extensions. To achieve this, one option is that a "root" ALTO Server implementing [RFC7285] resources and running at a given domain delegates "specialized" information resources, such as the ones providing Cost Calendars, to another ALTO Server running in a subdomain. The "root" ALTO Server can provide a Calendar-specific resource entry that has a media-type of "application/alto-directory+json" and that specifies the URI allowing to retrieve the location of a Calendar-aware Server and discover its resources. This option is described in "Delegation Using IRD" (Section 9.2.4 of [RFC7285]).

This document provides an example where a "root" ALTO Server runs in a domain called "alto.example.com". It delegates the announcement of Calendars capabilities to an ALTO Server running in a subdomain called "custom.alto.example.com". The location of the "delegate Calendar IRD" is assumed to be indicated in the "root" IRD by the resource entry: "custom-calendared-resources".

Another benefit of delegation is that some cost types for some resources may be more advantageous as Cost Calendars, and it makes little sense to get them as a single value. For example, if a cost type has predictable and frequently changing values calendared in short time intervals, such as a minute, it saves time and network resources to track the cost values via a focused delegate Server rather than the more general "root" Server.

4.3. Example IRD with ALTO Cost Calendars

This section provides an example ALTO Server IRD that supports various cost metrics and cost modes. In particular, since [RFC7285] makes it mandatory, the Server uses metric "routingcost" in the "numerical" mode.

For illustrative purposes, this section introduces 3 other fictitious example metrics and modes that should be understood as examples and should not be used or considered as normative.

The cost type names used in the example IRD are as follows:

"num-routingcost":

Refers to metric "routingcost" in the numerical mode, as defined in [RFC7285] and registered with IANA.

"num-owdelay":

Refers to fictitious performance metric "owdelay" in the "numerical" mode to reflect the one-way packet transmission delay on a path. A related performance metric is currently under definition in [ALTO_METRICS].

"num-throughputrating":

Refers to fictitious metric "throughputrating" in the "numerical" mode to reflect the provider preference in terms of end-to-end throughput.

"string-servicestatus":

Refers to fictitious metric "servicestatus" containing a string to reflect the availability, defined by the provider, of, for instance, path connectivity.

The example IRD includes 2 particular URIs providing Calendars:

"https://custom.alto.example.com/calendar/costmap/filtered":

A Filtered Cost Map in which Calendar capabilities are indicated for cost type names "num-routingcost", "num-throughputrating", and "string-servicestatus" and

"https://custom.alto.example.com/calendar/endpointcost/lookup":

An Endpoint Cost Map in which Calendar capabilities are indicated for cost type names "num-routingcost", "num-owdelay", "num-throughputrating", and "string-servicestatus".

The design of the Calendar capabilities allows some Calendars with the same cost type name to be available in several information resources with different Calendar attributes. This is the case for Calendars on "num-routingcost", "num-throughputrating", and "string-servicestatus", available in both the Filtered Cost Map and Endpoint Cost Service but with different time interval sizes for "num-throughputrating" and "string-servicestatus".

--- Client to Server request for IRD -----

GET /calendars-directory HTTP/1.1

Host: custom.alto.example.com

Accept: application/alto-directory+json,application/alto-error+json

--- Server response to Client -----

HTTP/1.1 200 OK

Content-Length: 2622

Content-Type: application/alto-directory+json

```
{
  "meta" : {
    "default-alto-network-map" : "my-default-network-map",
    "cost-types": {
      "num-routingcost": {
        "cost-mode" : "numerical",
        "cost-metric" : "routingcost"
      },
      "num-owdelay": {
        "cost-mode" : "numerical",
        "cost-metric": "owdelay"
      },
      "num-throughputrating": {
        "cost-mode" : "numerical",
        "cost-metric": "throughputrating"
      }
    }
  }
}
```

```

    },
    "string-servicestatus": {
        "cost-mode" : "string",
        "cost-metric": "servicestatus"
    }
}
},
"resources" : {
    "filtered-cost-map-calendar" : {
        "uri" :
            "https://custom.alto.example.com/calendar/costmap/filtered",
        "media-type" : "application/alto-costmap+json",
        "accepts" : "application/alto-costmapfilter+json",
        "capabilities" : {
            "cost-constraints" : true,
            "cost-type-names" : [ "num-routingcost",
                                   "num-throughputrating",
                                   "string-servicestatus" ],
            "calendar-attributes" : [
                { "cost-type-names" : [ "num-routingcost",
                                         "num-throughputrating" ],
                  "time-interval-size" : 7200,
                  "number-of-intervals" : 12
                },
                { "cost-type-names" : [ "string-servicestatus" ],
                  "time-interval-size" : 1800,
                  "number-of-intervals" : 48
                }
            ]
        },
        "uses": [ "my-default-network-map" ]
    },
    "endpoint-cost-map-calendar" : {
        "uri" :
            "https://custom.alto.example.com/calendar/endpointcost/lookup",
        "media-type" : "application/alto-endpointcost+json",
        "accepts" : "application/alto-endpointcostparams+json",
        "capabilities" : {
            "cost-constraints" : true,
            "cost-type-names" : [ "num-routingcost",
                                   "num-owdelay",
                                   "num-throughputrating",
                                   "string-servicestatus" ],
            "calendar-attributes" : [
                { "cost-type-names" : [ "num-routingcost" ],
                  "time-interval-size" : 3600,
                  "number-of-intervals" : 24
                },
                { "cost-type-names" : [ "num-owdelay" ],
                  "time-interval-size" : 300,
                  "number-of-intervals" : 12
                },
                { "cost-type-names" : [ "num-throughputrating" ],
                  "time-interval-size" : 60,
                  "number-of-intervals" : 60
                }
            ],
        },
    },

```


5.1. Calendar Extensions for Filtered Cost Maps (FCM)

A legacy ALTO Client requests and gets Filtered Cost Map responses, as specified in [RFC7285].

5.1.1. Calendar Extensions in Filtered Cost Map Requests

The input parameters of a "legacy" request for a Filtered Cost Map, defined by object ReqFilteredCostMap in Section 11.3.2 of [RFC7285], are augmented with one additional member. The same augmentation applies to object ReqFilteredCostMap defined in Section 4.1.2 of [RFC8189].

A Calendar-aware ALTO Client requesting a Calendar on a given cost type for a Filtered Cost Map resource having Calendar capabilities MUST add the following field to its input parameters:

JSONBoolean calendared<1..*>;

This field is an array of 1 to N boolean values, where N is the number of requested metrics. N is greater than 1 when the Client and the Server also implement [RFC8189].

Each entry corresponds to the requested metric at the same array position. Each boolean value indicates whether or not the ALTO Server should provide the values for this cost type as a Calendar. The array MUST contain exactly N boolean values, otherwise, the Server returns an error.

This field MUST NOT be included if no member "calendar-attributes" is specified in this information resource.

If a value of field "calendared" is 'true' for a cost type name for which no Calendar attributes have been specified, an ALTO Server, whether it implements the extensions of this document or only implements [RFC7285], MUST ignore it and return a response with a single cost value, as specified in [RFC7285].

If this field is not present, it MUST be assumed to have only values equal to 'false'.

A Calendar-aware ALTO Client that supports requests for only one cost type at a time and wants to request a Calendar MUST provide an array of 1 element:

"calendared" : [true],

A Calendar-aware ALTO Client that supports requests for more than one cost type at a time, as specified in [RFC8189], MUST provide an array of N values set to 'true' or 'false', depending whether it wants the applicable cost type values as a single or calendared value.

5.1.2. Calendar Extensions in Filtered Cost Map Responses

In a calendared ALTO Filtered Cost Map, a cost value between a source

and a destination is a JSON array of JSON values. An ALTO Calendar values array has a number of values equal to the value of member "number-of-intervals" of the Calendar attributes that are indicated in the IRD. These attributes will be conveyed as metadata in the Filtered Cost Map response. Each element of the array is valid for the time interval that matches its array position.

The FCM response conveys metadata, among which:

- * some are not specific to Calendars and ensure compatibility with [RFC7285] and [RFC8189] and
- * some are specific to Calendars.

The non-Calendar-specific "meta" fields of a calendared Filtered Cost Map response MUST include at least:

- * if the ALTO Client requests cost values for one cost type at a time, only the "meta" fields specified in [RFC7285] for these information service responses:
 - "dependent-vtags" and
 - "cost-type" field.
- * if the ALTO Client implements the Multi-Cost ALTO extension specified in [RFC8189] and requests cost values for several cost types at a time, the "meta" fields specified in [RFC8189] for these information service responses:
 - "dependent-vtags",
 - "cost-type" field with value set to '{}', for backwards compatibility with [RFC7285], and
 - "multi-cost-types" field.

If the Client request does not provide member "calendared" or if it provides it with a value equal to 'false' for all the requested cost types, then the ALTO Server response is exactly as specified in [RFC7285] and [RFC8189].

If the value of member "calendared" is equal to 'false' for a given requested cost type, the ALTO Server MUST return, for this cost type, a single cost value as specified in [RFC7285].

If the value of member "calendared" is equal to 'true' for a given requested cost type, the ALTO Server returns, for this cost type, a cost value Calendar, as specified above in this section. In addition to the above cited non-Calendar-specific "meta" members, the Server MUST provide a Calendar-specific metadata field.

The Calendar-specific "meta" field that a calendared Filtered Cost Map response MUST include is a member called "calendar-response-attributes", which describes properties of the Calendar and where:

- * member "calendar-response-attributes" is an array of one or more objects of type "CalendarResponseAttributes",
- * each "CalendarResponseAttributes" object in the array is specified for one or more cost types for which the value of member "calendared", in object ReqFilteredCostMap provided in the Client request, is equal to 'true' and for which a Calendar is provided for the requested information resource, and
- * the "CalendarResponseAttributes" object that applies to a cost type name has a corresponding "CalendarAttributes" object defined for this cost type name in the IRD capabilities of the requested information resource. This object is the entry in the "calendar-attributes" array member of the IRD resource entry, which includes the name of the requested cost type. This corresponding "CalendarAttributes" object has the same values as object "CalendarResponseAttributes" for members "time-interval-size" and "number-of-intervals". The members of the "CalendarResponseAttributes" object include all the members of the corresponding "CalendarAttributes" object.

The format of member "CalendarResponseAttributes" is defined as follows:

CalendarResponseAttributes calendar-response-attributes <1..*>;

```
object{
  [JSONString cost-type-names <1..*>;]
  JSONString calendar-start-time;
  JSONNumber time-interval-size;
  JSONNumber number-of-intervals;
  [JSONNumber repeated;]
} CalendarResponseAttributes;
```

Object CalendarResponseAttributes has the following attributes:

"cost-type-names":

An array of one or more cost type names to which the value of the other members of CalendarResponseAttributes apply and for which a Calendar has been requested. The value of this member is a subset of the "cost-type-names" member of the abovementioned corresponding "CalendarAttributes" object in the "calendar-attributes" array member in the IRD. This member MUST be present when Cost Calendars are provided for more than one cost type.

"calendar-start-time":

Indicates the date at which the first value of the Calendar applies. The value is a string that, as specified in Section 5, contains an HTTP "Date" header field using the IMF-fixdate format specified in Section 7.1.1.1 of [RFC7231]. The value provided for attribute "calendar-start-time" SHOULD NOT be later than the request date.

"time-interval-size":

As specified in Section 4.1 and with the same value as in the abovementioned corresponding "CalendarAttributes" object.

"number-of-intervals":

As specified in Section 4.1 and with the same value as in the abovementioned corresponding "CalendarAttributes" object.

"repeated":

An optional field provided for Calendars. It is an integer N greater or equal to '1' that indicates how many iterations of the Calendar value array starting at the date indicated by "calendar-start-time" have the same values. The number N includes the iteration provided in the returned response.

For example, suppose the "calendar-start-time" member has value "Mon, 30 Jun 2019 00:00:00 GMT", the "time-interval-size" member has value '3600', the "number-of-intervals" member has value '24', and the value of member "repeated" is equal to '4'. This means that the Calendar values are the same on Monday, Tuesday, Wednesday, and Thursday on a period of 24 hours starting at 00:00:00 GMT. The ALTO Client thus may use the same Calendar for the next 4 days starting at "calendar-start-time" and will only need to request a new one for Friday, July 4th at 00:00:00 GMT.

Attribute "repeated" may take a very high value if a Calendar represents a cyclic value pattern that the Server considers valid for a long period and hence will only update once this period has elapsed or if an unexpected event occurs on the network. In the latter case, the Client will be notified if it uses the "ALTO Incremental Updates Using Server-Sent Events (SSE)" Service, specified in [RFC8895]. To this end, it is RECOMMENDED that ALTO Servers providing ALTO Calendars also provide the "ALTO Incremental Updates Using Server-Sent Events (SSE)" Service, which is specified in [RFC8895]. Likewise, ALTO Clients capable of using ALTO Calendars SHOULD also use the SSE Service. See also discussion in Section 8 "Operational Considerations".

5.1.3. Use Case and Example: FCM with a Bandwidth Calendar

An example of non-real-time information that can be provisioned in a Calendar is the expected path throughput. While the transmission rate can be measured in real time by end systems, the operator of a data center is in the position of formulating preferences for given paths at given time periods to avoid traffic peaks due to diurnal usage patterns. In this example, we assume that an ALTO Client requests a Calendar of network-provider-defined throughput ratings as specified in the IRD to schedule its bulk data transfers as described in the use cases.

In the example IRD, Calendars for cost type name "num-throughputrating" are available for the information resources "filtered-cost-calendar-map" and "endpoint-cost-map-calendar". The ALTO Client requests a Calendar for "num-throughputrating" via a POST request for a Filtered Cost Map.

We suppose in the present example that the ALTO Client sends its request on Tuesday, July 1st 2019 at 13:15. The Server returns Calendars with arrays of 12 numbers for each source and destination

pair. The values for metric "throughputrating", in this example, are assumed to be encoded in 2 digits.

```
POST /calendar/costmap/filtered HTTP/1.1
Host: custom.alto.example.com
Content-Length: 217
Content-Type: application/alto-costmapfilter+json
Accept: application/alto-costmap+json,application/alto-error+json
```

```
{
  "cost-type" : {"cost-mode" : "numerical",
                 "cost-metric" : "throughputrating"},
  "calendared" : [true],
  "pids" : {
    "srcs" : [ "PID1", "PID2" ],
    "dsts" : [ "PID1", "PID2", "PID3" ]
  }
}
```

```
HTTP/1.1 200 OK
Content-Length: 1043
Content-Type: application/alto-costmap+json
```

```
{
  "meta" : {
    "dependent-vtags" : [
      { "resource-id": "my-default-network-map",
        "tag": "3ee2cb7e8d63d9fab71b9b34cbf764436315542e"
      }
    ],
    "cost-type" : {"cost-mode" : "numerical",
                   "cost-metric" : "throughputrating"},
    "calendar-response-attributes" : [
      { "calendar-start-time" : "Tue, 1 Jul 2019 13:00:00 GMT",
        "time-interval-size" : 7200,
        "number-of-intervals" : 12}
    ]
  },
  "cost-map" : {
    "PID1": { "PID1": [ 1, 12, 14, 18, 14, 14,
                       14, 18, 19, 20, 11, 12 ],
              "PID2": [ 13, 4, 15, 16, 17, 18,
                       19, 20, 11, 12, 13, 14 ],
              "PID3": [ 20, 20, 18, 14, 12, 12,
                       14, 14, 12, 12, 14, 16 ] },
    "PID2": { "PID1": [ 17, 18, 19, 10, 11, 12,
                       13, 14, 15, 16, 17, 18 ],
              "PID2": [ 20, 20, 18, 16, 14, 14,
                       14, 16, 16, 16, 14, 16 ],
              "PID3": [ 20, 20, 18, 14, 12, 12,
                       14, 14, 12, 12, 14, 16 ] }
  }
}
```

5.2. Calendar Extensions in the Endpoint Cost Service

This document extends the Endpoint Cost Service, as defined in Section 11.5.1 of [RFC7285], by adding new input parameters and capabilities and by returning JSONArrays instead of JSONNumbers as the cost values. The media type (Section 11.5.1.1 of [RFC7285]) and HTTP method (Section 11.5.1.2 of [RFC7285]) are unchanged.

5.2.1. Calendar-Specific Input in Endpoint Cost Requests

The extensions to the requests for calendared Endpoint Cost Maps are the same as for the Filtered Cost Map Service, specified in Section 5.1.1 of this document. Likewise, the rules defined around the extensions to ECM requests are the same as those defined in Section 5.1.1 for FCM requests.

The ReqEndpointCostMap object for a calendared ECM request will have the following format:

```
object {  
  [CostType cost-type;]  
  [CostType multi-cost-types<1..*>;]  
  [JSONBoolean calendared<1..*>;]  
  EndpointFilter endpoints;  
} ReqEndpointCostMap;  
  
object {  
  [TypedEndpointAddr srcs<0..*>;]  
  [TypedEndpointAddr dsts<0..*>;]  
} EndpointFilter;
```

Member "cost-type" is optional because, in the ReqEndpointCostMap object definition of this document, it is jointly present with member "multi-cost-types" to ensure compatibility with [RFC8189]. In [RFC8189], members "cost-type" and "multi-cost-types" are both optional and have to obey the rule specified in Section 4.1.2 of [RFC8189] stating that "the Client MUST specify either "cost-type" or "multi-cost-types" but MUST NOT specify both".

The interpretation of member "calendared" is the same as for the ReqFilteredCostMap object defined in Section 5.1.1 of this document. The interpretation of the other members is the same as for object ReqEndpointCostMap defined in [RFC7285] and [RFC8189]. The type TypedEndpointAddr is defined in Section 10.4.1 of [RFC7285].

For the reasons explained in Section 3.3, a Calendar-aware ALTO Server does not support constraints. Therefore, member "[constraints]" is not present in the ReqEndpointCostMap object, and member "constraints" MUST NOT be present in the input parameters of a request for an Endpoint Cost Calendar. If this member is present, the Server MUST ignore it.

5.2.2. Calendar Attributes in the Endpoint Cost Response

The "meta" field of a calendared Endpoint Cost response MUST include at least:

- * if the ALTO Client supports cost values for one cost type at a

time only, the "meta" fields specified in Section 11.5.1.6 of [RFC7285] for the Endpoint Cost response:

- "cost-type" field.
- * if the ALTO Client supports cost values for several cost types at a time, as specified in [RFC8189], the "meta" fields specified in [RFC8189] for the Endpoint Cost response:
 - "cost-type" field with value set to '{}', for backwards compatibility with [RFC7285].
 - "multi-cost-types" field.

If the Client request does not provide member "calendared" or if it provides it with a value equal to 'false', for all the requested cost types, then the ALTO Server response is exactly as specified in [RFC7285] and [RFC8189].

If the ALTO Client provides member "calendared" in the input parameters with a value equal to 'true' for given requested cost types, the "meta" member of a calendared Endpoint Cost response MUST include, for these cost types, an additional member "calendar-response-attributes", the contents of which obey the same rules as for the Filtered Cost Map Service, specified in Section 5.1.2. The Server response is thus changed as follows, with respect to [RFC7285] and [RFC8189]:

- * the "meta" member has one additional field "CalendarResponseAttributes", as specified for the Filtered Cost Map Service, and
- * the calendared costs are JSONArrays instead of the JSONNumbers format used by legacy ALTO implementations. All arrays have a number of values equal to 'number-of-intervals'. Each value corresponds to the cost in that interval.

If the value of member "calendared" is equal to 'false' for a given requested cost type, the ALTO Server MUST return, for this cost type, a single cost value as specified in [RFC7285].

5.2.3. Use Case and Example: ECS with a routingcost Calendar

Let us assume an Application Client is located in an end system with limited resources and has access to the network that is either intermittent or provides an acceptable quality in limited but predictable time periods. Therefore, it needs to schedule both its resource-greedy networking activities and its ALTO transactions.

The Application Client has the choice to trade content or resources with a set of endpoints and needs to decide with which one it will connect and at what time. For instance, the endpoints are spread in different time zones or have intermittent access. In this example, the 'routingcost' is assumed to be time-varying, with values provided as ALTO Calendars.

The ALTO Client associated with the Application Client queries an ALTO Calendar on 'routingcost' and will get the Calendar covering the 24-hour time period "containing" the date and time of the ALTO Client request.

For cost type "num-routingcost", the solicited ALTO Server has defined 3 different daily patterns, each represented by a Calendar to cover the week of Monday, June 30th at 00:00 to Sunday, July 6th 23:59:

- * C1 for Monday, Tuesday, Wednesday, and Thursday (weekdays)
- * C2 for Saturday and Sunday (weekends)
- * C3 for Friday (maintenance outage on July 4, 2019 from 02:00:00 GMT to 04:00:00 GMT or a big holiday that is widely celebrated and generates a large number of connections).

In the following example, the ALTO Client sends its request on Tuesday, July 1st 2019 at 13:15.

The "routingcost" values are assumed to be encoded in 3 digits.

```
POST /calendar/endpointcost/lookup HTTP/1.1
Host: custom.alto.example.com
Content-Length: 304
Content-Type: application/alto-endpointcostparams+json
Accept: application/alto-endpointcost+json,
        application/alto-error+json
```

```
{
  "cost-type" : {"cost-mode" : "numerical",
                 "cost-metric" : "routingcost"},
  "calendared" : [true],
  "endpoints" : {
    "srcs": [ "ipv4:192.0.2.2" ],
    "dsts": [
      "ipv4:192.0.2.89",
      "ipv4:198.51.100.34",
      "ipv4:203.0.113.45",
      "ipv6:2001:db8::10"
    ]
  }
}
```

HTTP/1.1 200 OK

```
Content-Length: 1351
Content-Type: application/alto-endpointcost+json
```

```
{
  "meta" : {
    "cost-type" : {"cost-mode" : "numerical",
                   "cost-metric" : "routingcost"},
    "calendar-response-attributes" : [
      {"calendar-start-time" : "Mon, 30 Jun 2019 00:00:00 GMT",
```

```

        "time-interval-size" : 3600,
        "number-of-intervals" : 24,
        "repeated": 4
    }
]
},
"endpoint-cost-map" : {
    "ipv4:192.0.2.2": {
        "ipv4:192.0.2.89"      : [100, 100, 100, 100, 100, 150,
                                200, 300, 300, 300, 300, 250,
                                250, 300, 300, 300, 300, 300,
                                400, 250, 250, 200, 150, 150],
        "ipv4:198.51.100.34" : [ 80,  80,  80,  80, 150, 150,
                                250, 400, 400, 450, 400, 200,
                                200, 350, 400, 400, 400, 350,
                                500, 200, 200, 200, 100, 100],
        "ipv4:203.0.113.45"  : [300, 400, 250, 250, 200, 150,
                                150, 100, 100, 100, 100, 100,
                                100, 100, 100, 100, 100, 150,
                                200, 300, 300, 300, 300, 250],
        "ipv6:2001:db8::10"   : [200, 250, 300, 300, 300, 300,
                                250, 300, 300, 300, 300, 350,
                                300, 400, 250, 150, 100, 100,
                                100, 150, 200, 250, 250, 300]
    }
}
}

```

When the Client gets the Calendar for "routingcost", it sees that the "calendar-start-time" is Monday at 00h00 GMT and member "repeated" is equal to '4'. It understands that the provided values are valid until Thursday and will only need to get a Calendar update on Friday.

5.2.4. Use Case and Example: ECS with a Multi-cost Calendar for routingcost and owdelay

In this example, it is assumed that the ALTO Server implements multi-cost capabilities, as specified in [RFC8189]. That is, an ALTO Client can request and receive values for several cost types in one single transaction. An illustrating use case is a path selection done on the basis of 2 metrics: routingcost and owdelay.

As in the previous example, the IRD indicates that the ALTO Server provides "routingcost" Calendars in terms of 24 time intervals of 1 hour (3600 seconds) each.

For metric "owdelay", the IRD indicates that the ALTO Server provides Calendars in terms of 12 time interval values lasting 5 minutes (300 seconds) each.

In the following example transaction, the ALTO Client sends its request on Tuesday, July 1st 2019 at 13:15.

This example assumes that the values of metric "owdelay" and "routingcost" are encoded in 3 digits.

POST calendar/endpointcost/lookup HTTP/1.1
Host: custom.alto.example.com
Content-Length: 390
Content-Type: application/alto-endpointcostparams+json
Accept: application/alto-endpointcost+json,
application/alto-error+json

```
{
  "cost-type" : {},
  "multi-cost-types" : [
    {"cost-mode" : "numerical", "cost-metric" : "routingcost"},
    {"cost-mode" : "numerical", "cost-metric" : "owdelay"}
  ],
  "calendared" : [true, true],
  "endpoints" : {
    "srcs": [ "ipv4:192.0.2.2" ],
    "dsts": [
      "ipv4:192.0.2.89",
      "ipv4:198.51.100.34",
      "ipv4:203.0.113.45",
      "ipv6:2001:db8::10"
    ]
  }
}
```

HTTP/1.1 200 OK
Content-Length: 2165
Content-Type: application/alto-endpointcost+json

```
{
  "meta" : {
    "multi-cost-types" : [
      {"cost-mode" : "numerical", "cost-metric" : "routingcost"},
      {"cost-mode" : "numerical", "cost-metric" : "owdelay"}
    ],
    "calendar-response-attributes" : [
      {"cost-type-names" : [ "num-routingcost" ],
        "calendar-start-time" : "Mon, 30 Jun 2019 00:00:00 GMT",
        "time-interval-size" : 3600,
        "number-of-intervals" : 24,
        "repeated": 4 },
      {"cost-type-names" : [ "num-owdelay" ],
        "calendar-start-time" : "Tue, 1 Jul 2019 13:00:00 GMT",
        "time-interval-size" : 300,
        "number-of-intervals" : 12}
    ]
  },
  "endpoint-cost-map" : {
    "ipv4:192.0.2.2": {
      "ipv4:192.0.2.89" : [[100, 100, 100, 100, 100, 150,
                           200, 300, 300, 300, 300, 250,
                           250, 300, 300, 300, 300, 300,
                           400, 250, 250, 200, 150, 150],
                          [ 20, 400, 20, 80, 80, 90,
                           100, 90, 60, 40, 30, 20]],
      "ipv4:198.51.100.34" : [[ 80, 80, 80, 80, 150, 150,
```



```

                250, 400, 400, 450, 400, 200,
                200, 350, 400, 400, 400, 350,
                500, 200, 200, 200, 100, 100],
        [ 20, 20, 50, 30, 30, 30,
          30, 40, 40, 30, 20, 20]],
    "ipv4:203.0.113.45" : [[300, 400, 250, 250, 200, 150,
                          150, 100, 100, 100, 100, 100,
                          100, 100, 100, 100, 100, 150,
                          200, 300, 300, 300, 300, 250],
      [100, 90, 80, 60, 50, 50,
        40, 40, 60, 90, 100, 80]],
    "ipv6:2001:db8::10" : [[200, 250, 300, 300, 300, 300,
                          250, 300, 300, 300, 300, 350,
                          300, 400, 250, 150, 100, 100,
                          100, 150, 200, 250, 250, 300],
      [ 40, 40, 40, 40, 50, 50,
        50, 20, 10, 15, 30, 40]]
  }
}

```

When receiving the response, the Client sees that the Calendar values for metric "routingcost" are repeated for 4 iterations. Therefore, in its next requests until the "routingcost" Calendar is expected to change, the Client will only need to request a Calendar for "owdelay".

Without the ALTO Calendar extensions, the ALTO Client would have no clue on the dynamicity of the metric value change and would spend needless time requesting values at an inappropriate pace. In addition, without the Multi-Cost ALTO capabilities, the ALTO Client would duplicate this waste of time as it would need to send one request per cost metric.

6. IANA Considerations

This document has no IANA actions.

7. Security Considerations

As an extension of the base ALTO protocol [RFC7285], this document fits into the architecture of the base protocol and hence the security considerations (Section 15 of [RFC7285]) fully apply when this extension is provided by an ALTO Server. For example, the same authenticity and integrity considerations (Section 15.1 of [RFC7285]) still fully apply; the same considerations for the privacy of ALTO users (Section 15.4 of [RFC7285]) also still fully apply.

The calendaring information provided by this extension requires additional considerations on three security considerations discussed in [RFC7285]: potential undesirable guidance to Clients (Section 15.2 of [RFC7285]), confidentiality of ALTO information (Section 15.3 of [RFC7285]), and availability of ALTO (Section 15.5 of [RFC7285]). For example, by providing network information in the future in a Calendar, this extension may improve availability of ALTO when the ALTO Server is unavailable but related information is already

provided in the Calendar.

For confidentiality of ALTO information, an operator should be cognizant that this extension may introduce a new risk, a malicious ALTO Client may get information for future events that are scheduled through Calendaring. Possessing such information, the malicious Client may use it to generate massive connections to the network at times where its load is expected to be high.

To mitigate this risk, the operator should address the risk of ALTO information being leaked to malicious Clients or third parties. As specified in "Protection Strategies" (Section 15.3.2 of [RFC7285]), the ALTO Server should authenticate ALTO Clients and use the Transport Layer Security (TLS) protocol so that man-in-the-middle (MITM) attacks to intercept an ALTO Calendar are not possible. "Authentication and Encryption" (Section 8.3.5 of [RFC7285]) ensures the availability of such a solution. It specifies that "ALTO Server implementations as well as ALTO Client implementations MUST support the "https" URI scheme of [RFC2818] and Transport Layer Security (TLS) of [RFC5246]".

Section 1 of TLS 1.3 [RFC8446] states: "While TLS 1.3 is not directly compatible with previous versions, all versions of TLS incorporate a versioning mechanism which allows Clients and Servers to interoperably negotiate a common version if one is supported by both peers". ALTO Clients and Servers SHOULD support both TLS 1.3 [RFC8446] and TLS 1.2 [RFC5246] and MAY support and use newer versions of TLS as long as the negotiation process succeeds.

The operator should be cognizant that the preceding mechanisms do not address all security risks. In particular, they will not help in the case of "malicious Clients" possessing valid authentication credentials. The threat here is that legitimate Clients have become subverted by an attacker and are now 'bots' being asked to participate in a DDoS attack. The Calendar information now becomes valuable in knowing exactly when to perpetrate a DDoS attack. A mechanism, such as a monitoring system that detects abnormal behaviors, may still be needed.

To avoid malicious or erroneous guidance from ALTO information, an ALTO Client should be cognizant that using calendaring information can have risks: (1) Calendar values, especially in "repeated" Calendars, may be only statistical and (2) future events may change. Hence, a more robust ALTO Client should adapt and extend protection strategies specified in Section 15.2 of [RFC7285]. For example, to be notified immediately when a particular ALTO value that the Client depends on changes, it is RECOMMENDED that both the ALTO Client and ALTO Server using this extension support "Application-Layer Traffic Optimization (ALTO) Incremental Updates Using Server-Sent Events (SSE)" [RFC8895].

Another risk of erroneous guidance appears when the Server exposes an occurrence of a same cost type name in different elements of the Calendar objects array associated to an information resource. In this case, there is no way for the Client to figure out which Calendar object in the array is valid. The specification in this

document recommends, in this case, that the Client uses the first encountered Calendar object occurrence containing the cost type name. However, the Client may want to avoid the risks of erroneous guidance associated to the use of potentially invalid Calendar values. To this end, as an alternative to the recommendation in this document, the Client MAY ignore the totality of occurrences of CalendarAttributes objects containing the cost type name and query this cost type using [RFC7285].

8. Operational Considerations

It is important that both the operator of the network and the operator of the applications consider both the feedback aspect and the prediction-based (uncertainty) aspect of using the Cost Calendar.

First, consider the feedback aspect and consider the Cost Calendar as a traffic-aware map service (e.g., Google Maps). Using the service without considering its own effect, a large fleet can turn a not-congested road into a congested one; a large number of individual cars each choosing a road with light traffic ("cheap link") can also result in congestion or result in a less-optimal global outcome (e.g., the Braess' Paradox [BRAESS_PARADOX]).

Next, consider the prediction aspect. Conveying ALTO Cost Calendars tends to reduce the on-the-wire data exchange volume compared to multiple single-cost ALTO transactions. An application using Calendars has a set of time-dependent values upon which it can plan its connections in advance with no need for the ALTO Client to query information at each time. Additionally, the Calendar response attribute "repeated", when provided, saves additional data exchanges in that it indicates that the ALTO Client does not need to query Calendars during a period indicated by this attribute. The preceding is true only when "accidents" do not happen.

Although individual network operators and application operators can choose their own approaches to address the aforementioned issues, this document recommends the following considerations. First, a typical approach to reducing instability and handling uncertainty is to ensure timely update of information. The SSE Service, as discussed in Section 7, can handle updates if supported by both the Server and the Client. Second, when a network operator updates the Cost Calendar and when an application reacts to the update, they should consider the feedback effects. This is the best approach even though there is theoretical analysis [SELFISH_RTG_2002] and Internet-based evaluation [SELFISH_RTG_2003] showing that uncoordinated behaviors do not always cause substantial suboptimal results.

High-resolution intervals may be needed when values change, sometimes during very small time intervals but in a significant manner. A way to avoid conveying too many entries is to leverage on the "repeated" feature. A Server can smartly set the Calendar start time and number of intervals so as to declare them "repeated" for a large number of periods until the Calendar values change and are conveyed to requesting Clients.

The newer JSON Data Interchange Format specification [RFC8259] used

in ALTO Calendars replaces the older one [RFC7159] used in the base ALTO protocol [RFC7285]. The newer JSON mandates UTF-8 text encoding to improve interoperability. Therefore, ALTO Clients and Servers implementations using UTF-^{16,32} need to be cognizant of the subsequent interoperability risks and MUST switch to UTF-8 encoding if they want to interoperate with Calendar-aware Servers and Clients.

9. References

9.1. Normative References

- [IEEE.754.2019]
IEEE, "IEEE Standard for Floating-Point Arithmetic",
IEEE 754-2019, DOI 10.1109/IEEESTD.2019.8766229, June
2019, <<https://doi.org/10.1109/IEEESTD.2019.8766229>>.
- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate
Requirement Levels", BCP 14, RFC 2119,
DOI 10.17487/RFC2119, March 1997,
<<https://www.rfc-editor.org/info/rfc2119>>.
- [RFC5246] Dierks, T. and E. Rescorla, "The Transport Layer Security
(TLS) Protocol Version 1.2", RFC 5246,
DOI 10.17487/RFC5246, August 2008,
<<https://www.rfc-editor.org/info/rfc5246>>.
- [RFC7231] Fielding, R., Ed. and J. Reschke, Ed., "Hypertext Transfer
Protocol (HTTP/1.1): Semantics and Content", RFC 7231,
DOI 10.17487/RFC7231, June 2014,
<<https://www.rfc-editor.org/info/rfc7231>>.
- [RFC7285] Alimi, R., Ed., Penno, R., Ed., Yang, Y., Ed., Kiesel, S.,
Previdi, S., Roome, W., Shalunov, S., and R. Woundy,
"Application-Layer Traffic Optimization (ALTO) Protocol",
RFC 7285, DOI 10.17487/RFC7285, September 2014,
<<https://www.rfc-editor.org/info/rfc7285>>.
- [RFC8174] Leiba, B., "Ambiguity of Uppercase vs Lowercase in RFC
2119 Key Words", BCP 14, RFC 8174, DOI 10.17487/RFC8174,
May 2017, <<https://www.rfc-editor.org/info/rfc8174>>.
- [RFC8189] Randriamasy, S., Roome, W., and N. Schwan, "Multi-Cost
Application-Layer Traffic Optimization (ALTO)", RFC 8189,
DOI 10.17487/RFC8189, October 2017,
<<https://www.rfc-editor.org/info/rfc8189>>.
- [RFC8259] Bray, T., Ed., "The JavaScript Object Notation (JSON) Data
Interchange Format", STD 90, RFC 8259,
DOI 10.17487/RFC8259, December 2017,
<<https://www.rfc-editor.org/info/rfc8259>>.
- [RFC8446] Rescorla, E., "The Transport Layer Security (TLS) Protocol
Version 1.3", RFC 8446, DOI 10.17487/RFC8446, August 2018,
<<https://www.rfc-editor.org/info/rfc8446>>.
- [RFC8895] Roome, W. and Y. Yang, "Application-Layer Traffic

Optimization (ALTO) Incremental Updates Using Server-Sent Events (SSE)", RFC 8895, DOI 10.17487/RFC8895, November 2020, <<https://www.rfc-editor.org/info/rfc8895>>.

9.2. Informative References

[ALTO_METRICS]

Wu, Q., Yang, Y. R., Dhody, D., Randriamasy, S., and L. M. Contreras, "ALTO Performance Cost Metrics", Work in Progress, Internet-Draft, draft-ietf-alto-performance-metrics-09, 9 March 2020, <<https://tools.ietf.org/html/draft-ietf-alto-performance-metrics-09>>.

[BRAESS_PARADOX]

Steinberg, R. and W. Zangwill, "The Prevalence of Braess' Paradox", Transportation Science Vol. 17, No. 3, DOI 10.1287/trsc.17.3.301, 1 August 1983, <<https://doi.org/10.1287/trsc.17.3.301>>.

[RFC2818] Rescorla, E., "HTTP Over TLS", RFC 2818, DOI 10.17487/RFC2818, May 2000, <<https://www.rfc-editor.org/info/rfc2818>>.

[RFC5693] Seedorf, J. and E. Burger, "Application-Layer Traffic Optimization (ALTO) Problem Statement", RFC 5693, DOI 10.17487/RFC5693, October 2009, <<https://www.rfc-editor.org/info/rfc5693>>.

[RFC6708] Kiesel, S., Ed., Previdi, S., Stiernerling, M., Woundy, R., and Y. Yang, "Application-Layer Traffic Optimization (ALTO) Requirements", RFC 6708, DOI 10.17487/RFC6708, September 2012, <<https://www.rfc-editor.org/info/rfc6708>>.

[RFC7159] Bray, T., Ed., "The JavaScript Object Notation (JSON) Data Interchange Format", RFC 7159, DOI 10.17487/RFC7159, March 2014, <<https://www.rfc-editor.org/info/rfc7159>>.

[SELFISH_RTG_2002]

Roughgarden, T., "Selfish Routing", Dissertation Thesis, Cornell, May 2002.

[SELFISH_RTG_2003]

Qiu, L., Yang, Y., Zhang, Y., and S. Shenker, "Selfish Routing in Internet-Like Environments", Proceedings of SIGCOMM '03, DOI 10.1145/863955.863974, August 2003, <<https://doi.org/10.1145/863955.863974>>.

[SENSE] Department of Energy Office of Science Advanced Scientific Computing Research (ASCR) Program, "SDN for End-to-End Networked Science at the Exascale (SENSE)", <<http://sense.es.net/overview>>.

[UNICORN-FGCS]

Xiang, Q., Wang, T., Zhang, J., Newman, H., Yang, Y., and Y. Liu, "Unicorn: Unified resource orchestration for multi-domain, geo-distributed data analytics", Future

Acknowledgments

The authors would like to thank Fred Baker, Li Geng, Diego Lopez, He Peng, and Haibin Song for fruitful discussions and feedback on earlier draft versions. Dawn Chan, Kai Gao, Vijay Gurbani, Yichen Qian, Jürgen Schönwälder, Brian Weis, and Jensen Zhang provided substantial review feedback and suggestions to the protocol design.

Authors' Addresses

Sabine Randriamasy
Nokia Bell Labs
Route de Villejust
91460 Nozay
France

Email: Sabine.Randriamasy@nokia-bell-labs.com

Y. Richard Yang
Yale University
51 Prospect St.
New Haven, CT 06520
United States of America

Email: yry@cs.yale.edu

Qin Wu
Huawei
Yuhua District
101 Software Avenue
Nanjing
Jiangsu, 210012
China

Email: sunseawq@huawei.com

Lingli Deng
China Mobile
China

Email: denglingli@chinamobile.com

Nico Schwan
Thales Deutschland
Lorenzstrasse 10
70435 Stuttgart
Germany

Email: nico.schwan@thalesgroup.com